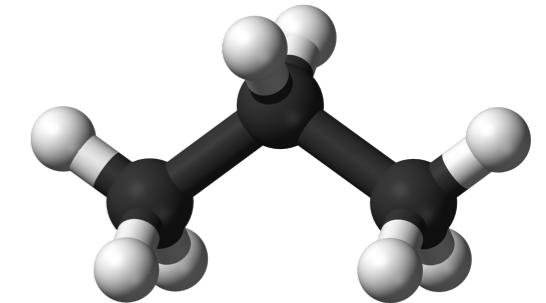
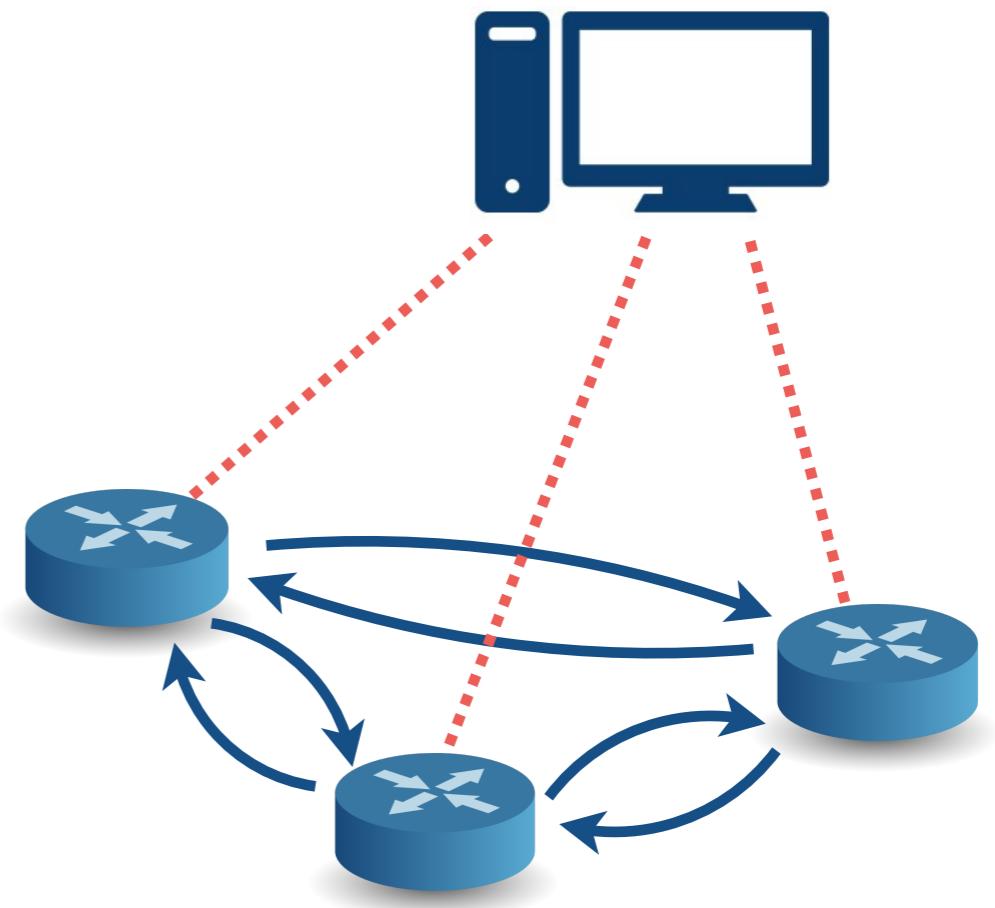


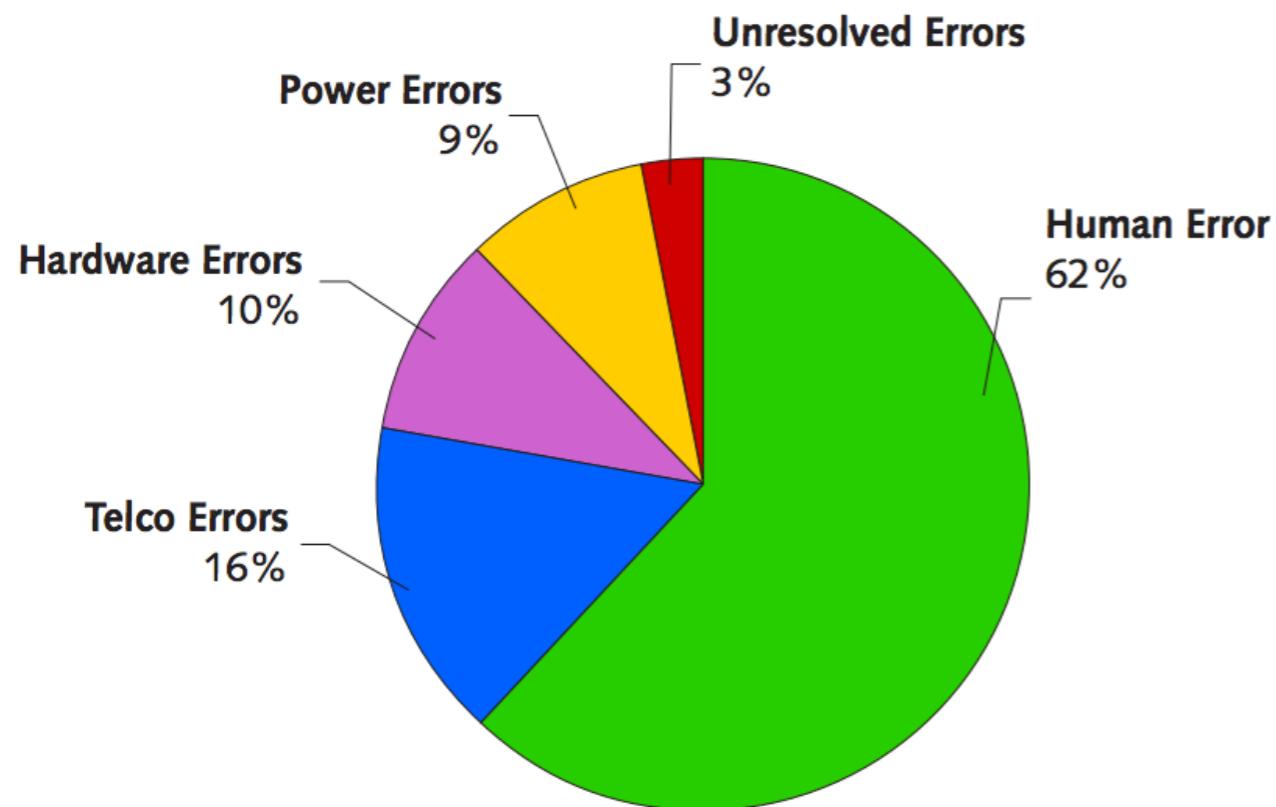
Propane: Programming Distributed Control Planes



Ryan Beckett (Princeton, MSR)
Ratul Mahajan (MSR)
Todd Millstein (UCLA)
Jitu Padhye (MSR)
David Walker (Princeton)

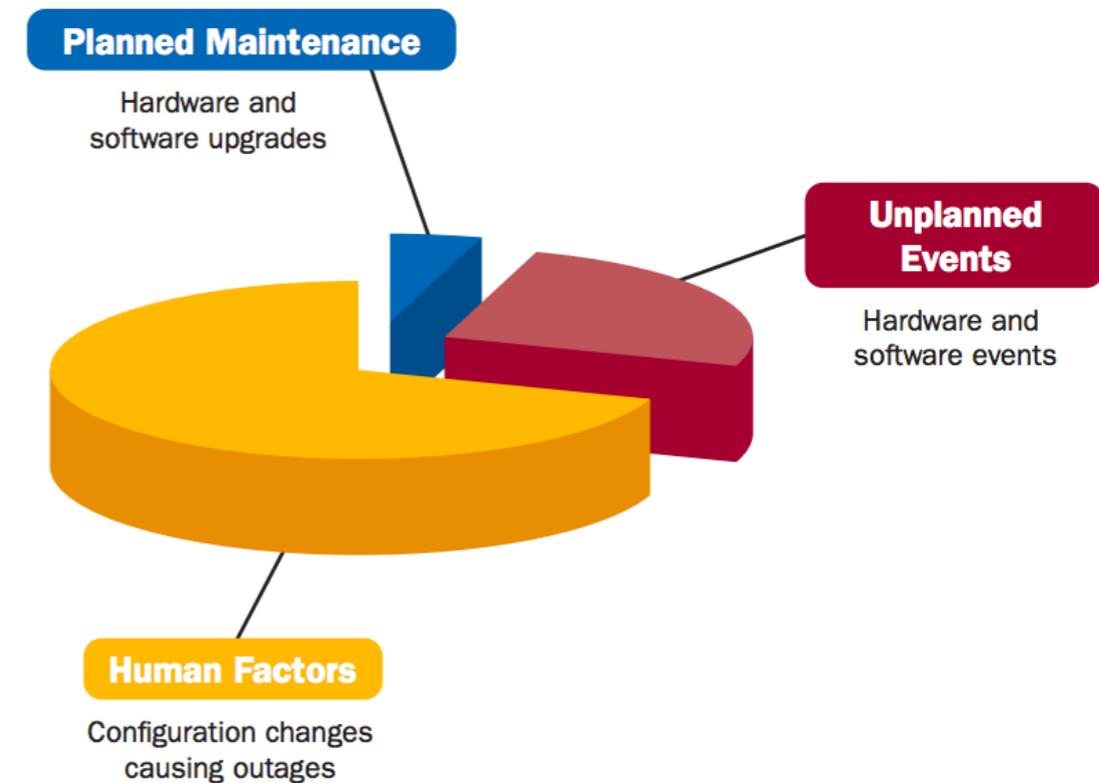


Configuring Networks is Error-Prone



**~60% of network downtime
is caused by human error**

-Yankee group 2002



**50-80% of outages are
the result of human error**

-Juniper 2008

Configuring Networks is Error-Prone

Understanding BGP Misconfiguration

Ratul Mahajan David Wetherall Tom Anderson

(ratul.djw.tom)@cs.washington.edu
Computer Science and Engineering
University of Washington
Seattle, WA 98195-2350

ABSTRACT
It is well-known that simple, accidental BGP configuration errors can disrupt Internet connectivity. Yet little is known about the frequency of misconfiguration or its causes, except for the few spectacular incidents of widespread outages. In this paper, we present the first quantitative study of BGP misconfiguration. Over a three week period, we analyzed routing table advertisements from 23 vantage points across the Internet backbone to detect incidents of misconfiguration. For each incident we polled the ISP operators involved to verify whether it was a misconfiguration, and to learn the cause of the incident. We also actively probed the Internet to determine the impact of misconfiguration on connectivity.

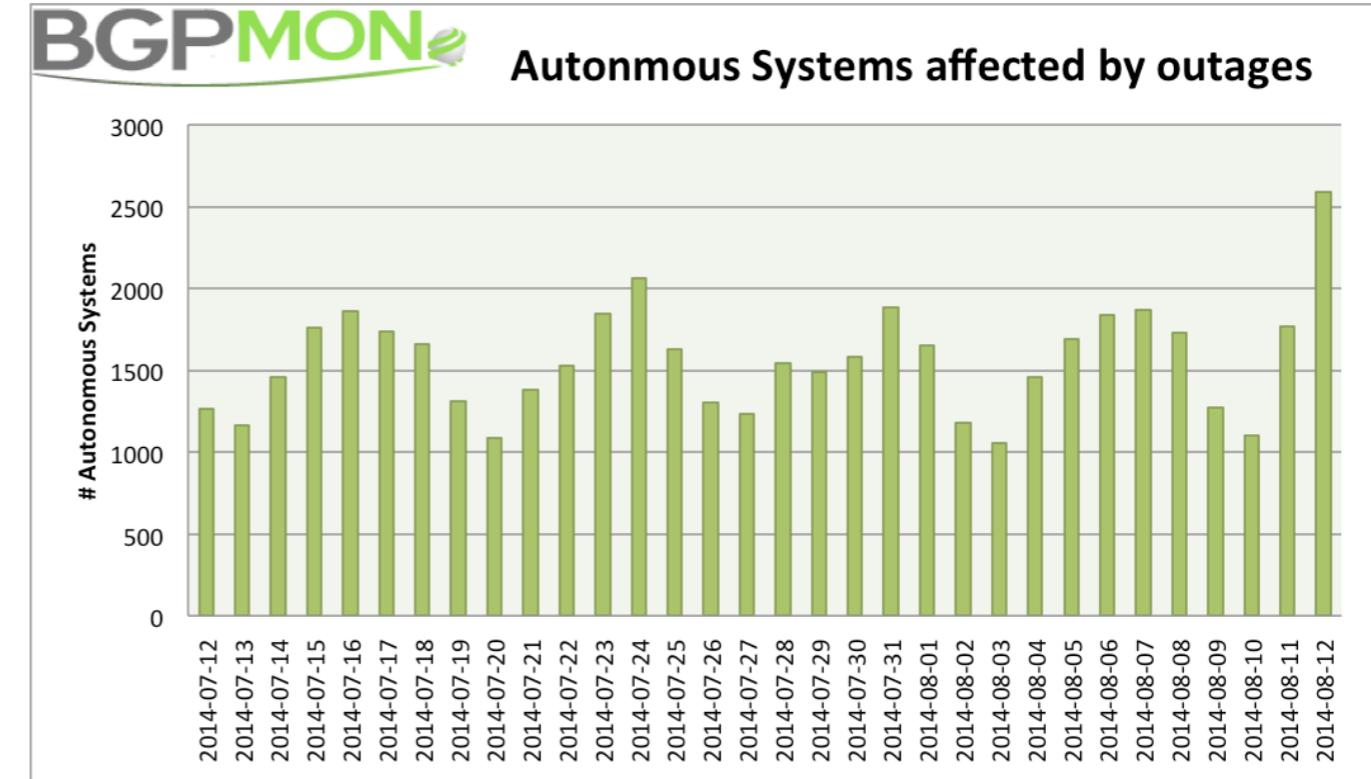
Surprisingly, we find that configuration errors are pervasive, with 200-1200 prefixes (0.2-1.0% of the BGP table size) suffering from misconfiguration each day. Close to 3 in 4 of all new prefix advertisements were results of misconfiguration. Fortunately, the connectivity seen by end users is surprisingly robust to misconfigurations. While misconfigurations can substantially increase the update load on routers, only one in twenty five affects connectivity. While the causes of misconfiguration are diverse, we argue that most could be prevented through better router design.

Categories and Subject Descriptors
C.2.3 [Communication Networks]: Operations—management;
C.4 [Computer Systems]: Performance—reliability, availability, and serviceability

General Terms
Human Factors, Management, Reliability

1. INTRODUCTION
As the Internet's inter-domain routing protocol, the Border Gateway Protocol (BGP) [34] is crucial to the overall reliability of the Internet. Faults in BGP implementations or mistakes in the way it is used have been known to disrupt large regions of the Internet. Recent studies have examined several kinds of BGP problems, including excessive churn due to implementation deficiencies [26].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.
SIGCOMM'02, August 19–23, 2002, Pittsburgh, Pennsylvania, USA.
Copyright 2002 ACM 1-58113-570-X/02/0008 ... \$5.00.



“Close to 3 in 4 of all new prefix advertisements were results of misconfiguration”

~1500 ASes affected by outages every day

Configuring Networks is Error-Prone

2/5/2016

Log in | Sign up

<https://www.thousandeyes.com>

Product Solutions About Login (<https://www.thousandeyes.com/login>) Sign Up ([https://www.thousandeeyes.com/signup](https://www.thousandeyes.com/signup))

← Blog Home (/)

 Time Warner Cable Outage Impacts

Posted by [Pete Anderson](https://blog.thousandeyes.com/pete-anderson) (<https://blog.thousandeyes.com/pete-anderson>)

By now a lot of you have probably read about the Time Warner outage. This morning I was greeted with a slew of alarms, names to check, and a lack of home office employees without any Internet access. The initial alarm was occurring and quickly opened up the Time Warner Outage page.

The alerts started coming in a little before 930 UTC (04:30 EST). Numerous inaccessible websites, DNS names failing to resolve, and a general lack of Internet. Companies that peer with Time Warner experienced problems with services such as supply chain portals (Figure 1).

24h 7d 14d

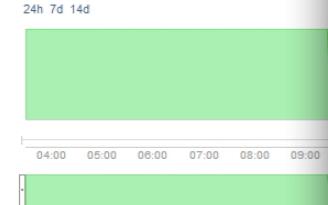


Figure 1: Supply chain portal with limited availability. The chart shows a green bar from 04:00 to 09:00, indicating a period of low availability. Below the chart is a world map with several red dots representing affected locations.

I could see right away that users and networks that connected to Time Warner via its supply chain portal, indicating the issue was in the Time Warner network. A brief service interruption while all the traffic re-routed through AT&T. Availability issues continued for the entire duration of the outage. I took a look at the path visualization view to figure out exactly what was going on. Normally, two locations (Tokyo and Dallas) transit Road Runner (Time Warner) to reach this supply chain port, while the rest go through AT&T (Figure 2).

China routing snafu briefly mangles interweb • The Register

Cash'n'Carrion | Whitepapers | The Channel | The Next Platform

GET YOU ONLY YOU CAN

2/5/2016

Internet-Wide Catastrophe—Last

 DATA CENTER SOFTWARE NETWORKS

[Networks](#) ▶ [Broadband](#)

China routing snafu briefly mangles interweb

Cockup, not conspiracy

9 Apr 2010 at 12:24, John Leyden

Bad routing information sourced from China has caused a temporary Internet-wide catastrophe. Global BGP (Border Gateway Routing) lookup tables have been hijacked by China Telecom, a state-owned Telecommunication, apparently accidentally by mistake. Telecommunications, IDG reports. ISPs including Qwest and Telefonica accepted ill-thought out routing information from China Telecom. BGP is a core routing protocol which maps out the Internet. Several routing options are normally available for a given connection, and it's equivalent of TomTom publishing routes via SatNav between London and Paris. IDC China Telecommunication published ill-conceived routing tables for parts of the Internet about 10 per cent of the net - instead of the normal 1 per cent. It has caused many viable routing options by many service providers to disappear. (This happened at 04:30 UTC time) after China Telecommunications republished them. It's not clear if other parts of Asia would have been more likely to adopt the same routing tables. The incident were recorded all over the world. BGPmon.net, a BGP monitoring service, has described as a prefix hijack. [here](#).

Although it seems they [IDC China Telecom] were responsible for the routing tables, about 10 per cent of these prefixes probably include prefixes for popular websites such as www.rapidshare.com and www.geocities.com. A large number of networks impacted the Internet-wide catastrophe, including some popular Chinese websites such as www.huanqiu.com, www.tianya.cn and www.sohu.com.

A cock-up is suspected, rather than a conspiracy. Given the large number of prefixes and servers involved, it's likely a prefix hijack. Most likely it's because of configuration errors or misconfiguration.

The practical consequences of the screw-up are not clear. It could result in dropped connections or, worse, traffic routed through a wrong path. This is one of the clearest illustrations of the security risks of the Internet's most nonethless important network protocol. The China BGP global routing represents a major challenge for network management. For example, just two weeks ago, TTNet in Turkey (AS9121) announced to route all the Internet traffic through a DNS (Domain Name System) server. This was a major mortem by internet monitoring firm Renesys [here](http://www.theregister.co.uk/2010/04/09/china_bgp).

« Previous Story [Next Story](#)

able to reach a large number of sites. We can take a look at what happened in the intervening time.

Morning 2004, TTNet (AS9121) started announcing

<http://research.dyn.com/2005/12/internetwide-nearcatastrophela/>

Sign In | Register



YouTube/Pakistan incident: Could something similar whack your site?

Configuring BGP properly is key to avoidance, 'Net registry official says



By Carolyn Duffy Marsan

Network World | Mar 10, 2008 1:00 AM PT

In light of Pakistan Telecom/YouTube incident, Internet registry official explains how you can avoid having your web site victimized by such an attack.

When Pakistan Telecom blocked YouTube's traffic one Sunday evening in February, the ISP created an international incident that wreaked havoc on the popular video site for more than two hours.

RIPE NCC, the European registry for Internet addresses, has conducted an analysis of what happened during Pakistan Telecom's hijacking of YouTube's traffic and the steps that YouTube took to stop the attack.

We posed some questions to RIPE NCC's Chief Scientist Daniel Karrenberg about the YouTube incident. Here's what he had to say:

How frequently do hijacking incidents like the Pakistan Telecom/YouTube incident happen?

Misconfigurations of iBGP (internal BGP, the protocol used between the routers in the same Autonomous System) happen regularly and are usually the result of an error. One such misconfiguration caused the Pakistan Telecom/YouTube incident. It appears that the Pakistan Telecom/YouTube incident was not an "attack" as some have labeled it, but a configuration error. (See [Columnist Johnna Till Johnson's take on the topic](#).)

What is significant about the YouTube incident?

TTNet (AS9121) started announcing

1/5



Objectives: Network-wide

- Prefer traffic go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Ensure traffic goes through X
- Adhere to policies even when **failures** occur



Objectives: Network-wide

- Prefer traffic go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Ensure traffic goes through X
- Adhere to policies even when **failures** occur





Objectives: Network-wide

- Prefer traffic go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Ensure traffic goes through X
- Adhere to policies even when **failures** occur



Mechanisms: Device-by-Device

- Local decisions made independently on each device
- Several device-level actions together imply some higher-level behavior
- Failures interact with local decision-making algorithms in complex ways

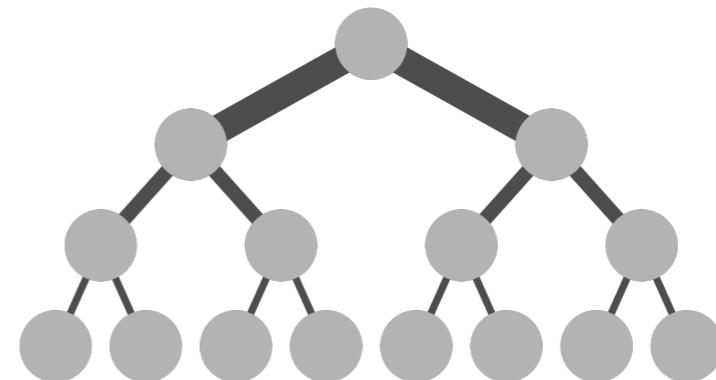
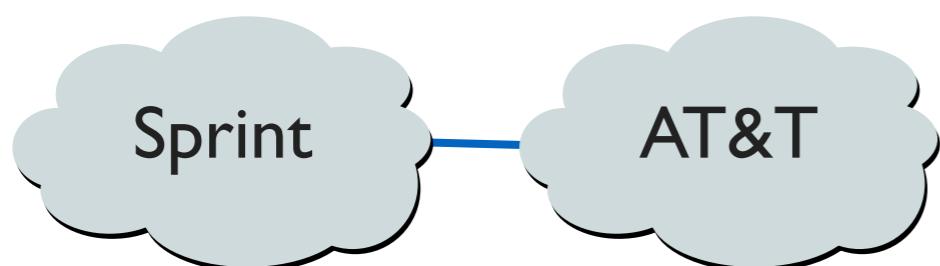
Why use BGP?

Advantages

- Allows for *local* policy among nodes
- Massively *scalable* — hundreds of thousands of IP prefixes
- Fully *distributed* — does not require global knowledge of network state

BGP in the Wild

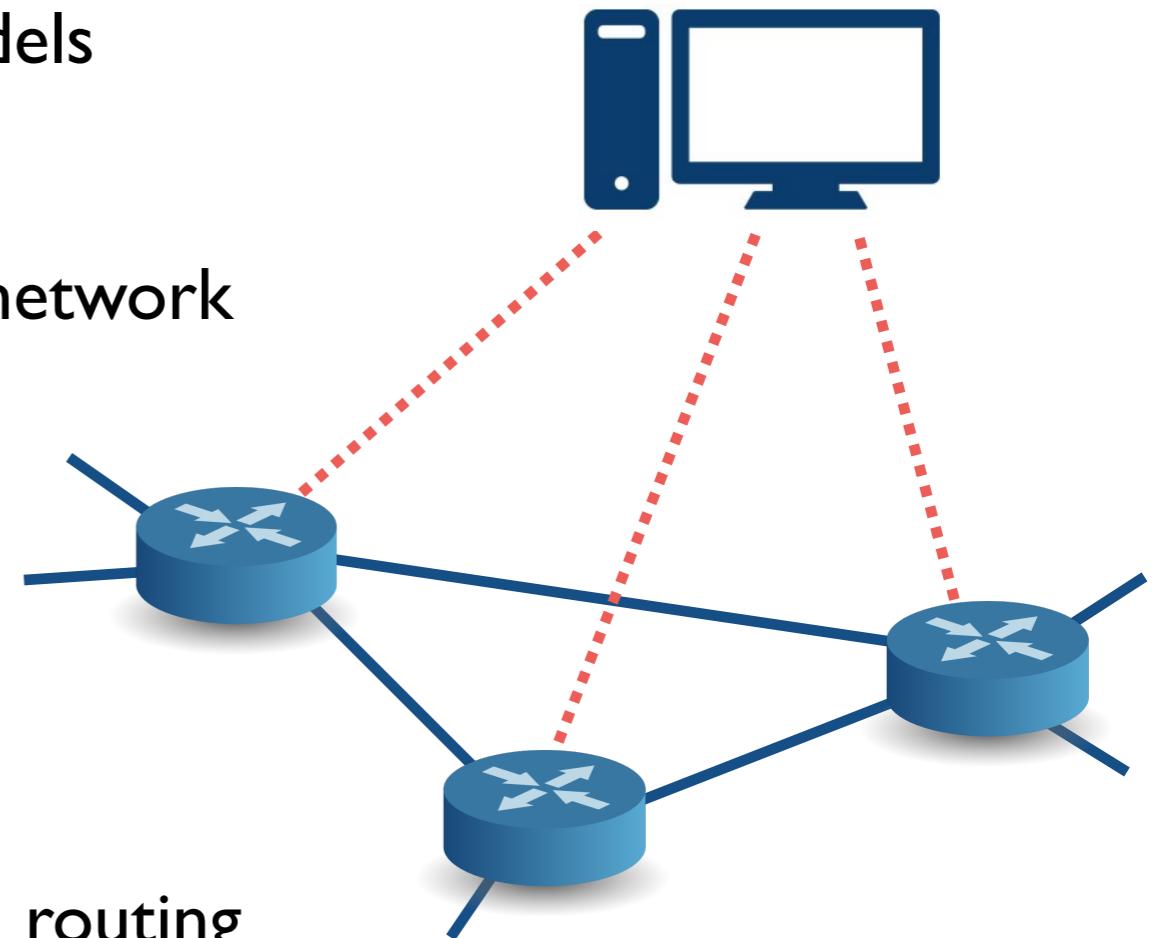
- Inter-domain routing between Autonomous Systems (AT&T, Sprint)
- Intra-domain routing in data centers



What about SDN?

Software Defined Networks

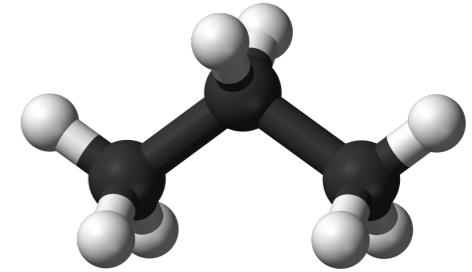
- Simpler, centralized programming models
- Network-wide abstractions like paths
- Centralized controller programs the network



But they are not a panacea:

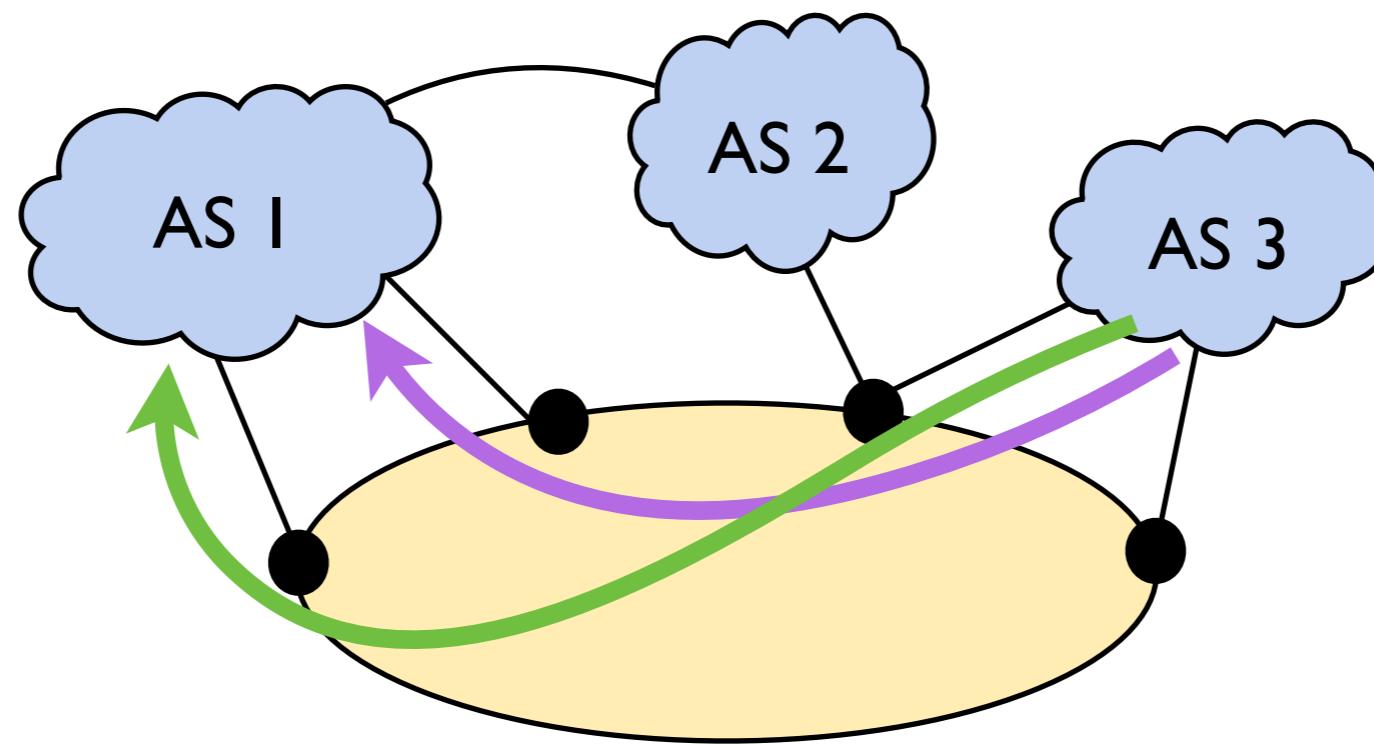
- Do not usually help with **inter-domain** routing
- **Latency**, and **scalability** can be problematic for large networks
- Require careful design and engineering to be robust to **failures**
- Their implementations don't exploit **existing** network infrastructure

Propane: Programming a Distributed Control Plane

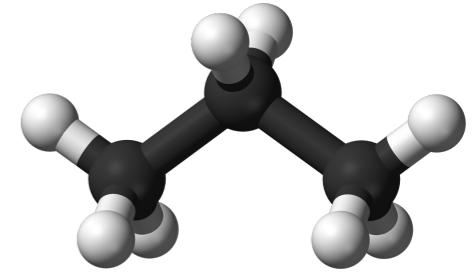


I) Language for expressing high-level operator objectives with:

- Network-wide programming abstraction
- Uniform abstractions for intra- and inter-domain routing
- Paths constraints and relative preferences with fall-backs in case of failures

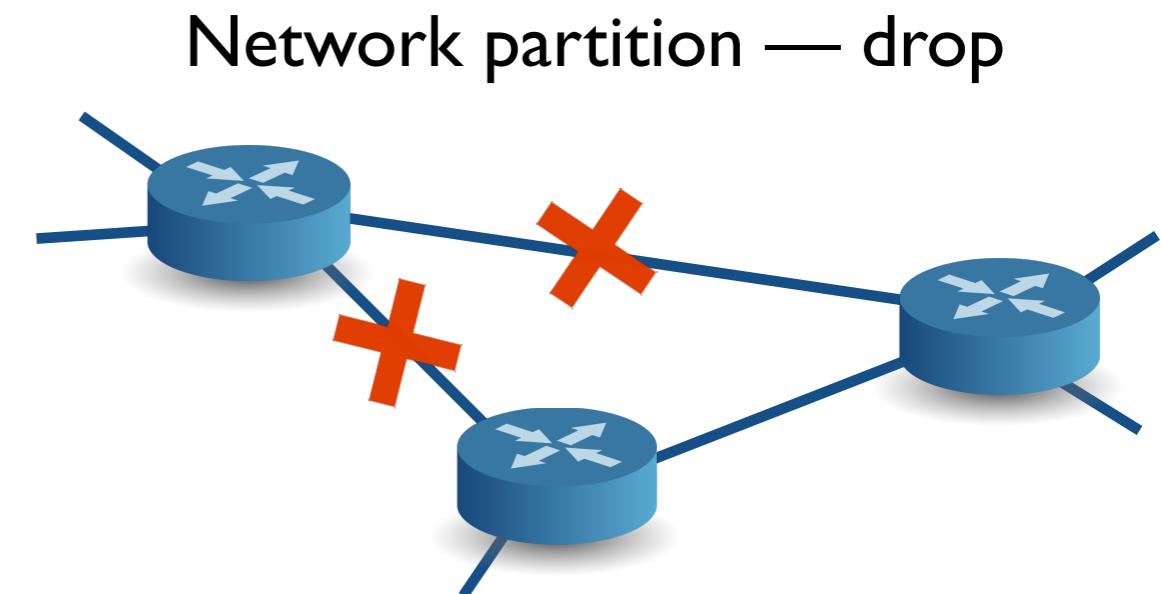
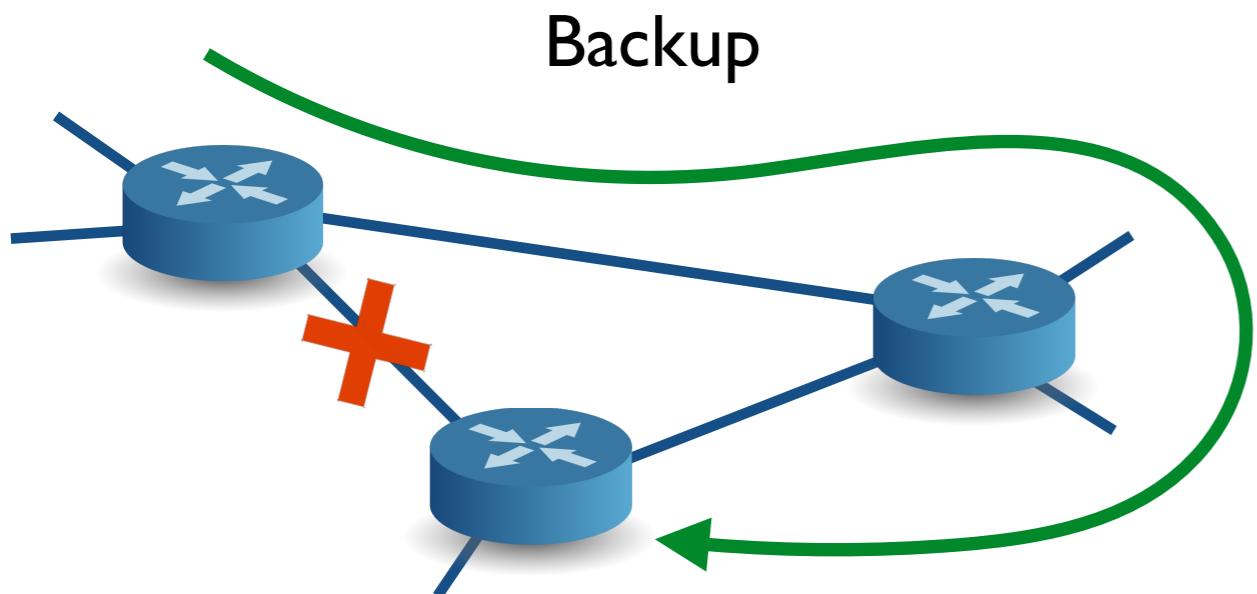


Propane: Programming a Distributed Control Plane



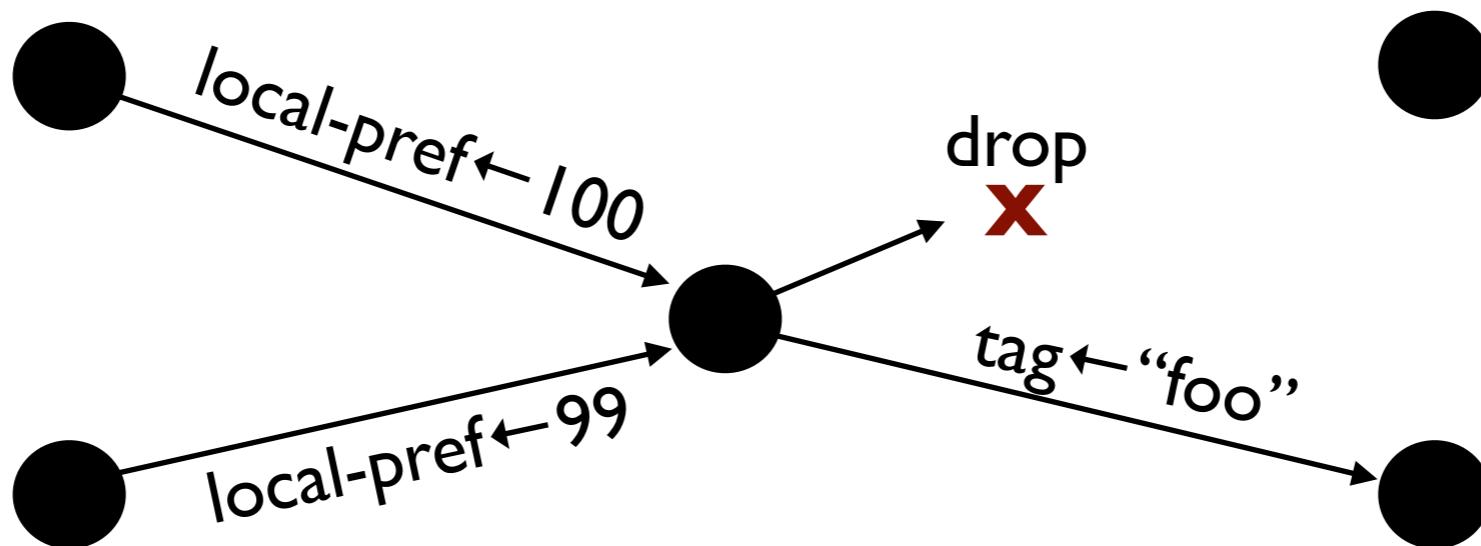
2) Compiler to generate a low-level distributed implementation:

- Efficient algorithms to synthesize a set of *policy-compliant* BGP configs
- Static analysis guarantees policy compliance under *all* failures

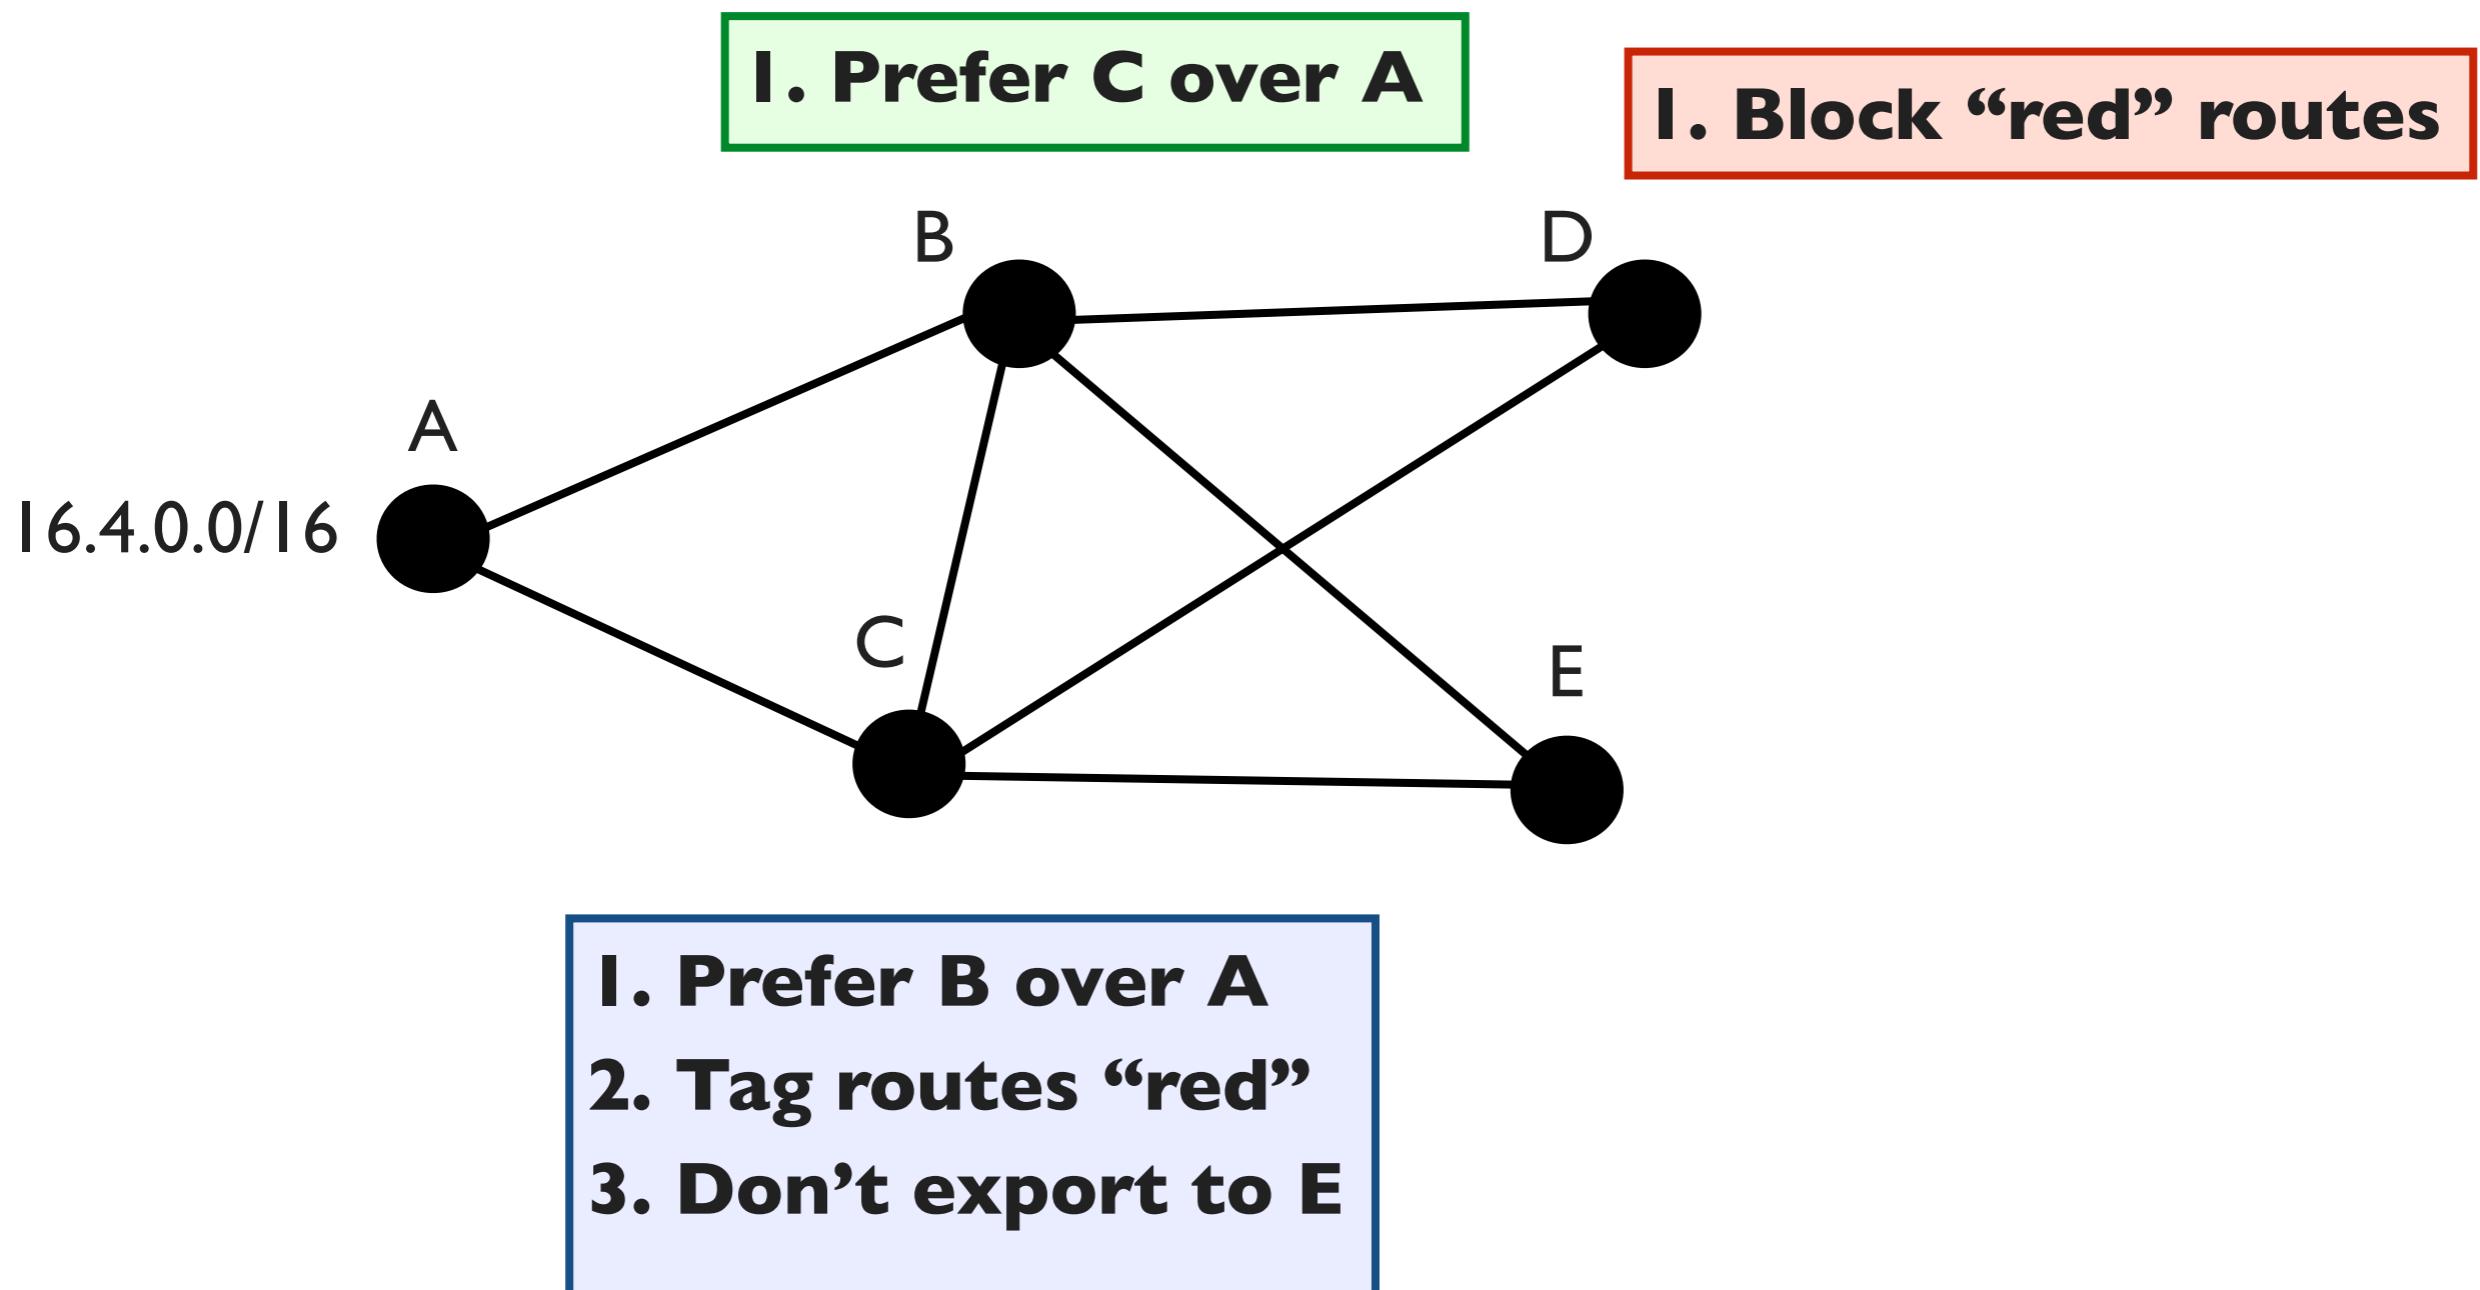


Border Gateway Protocol (BGP)

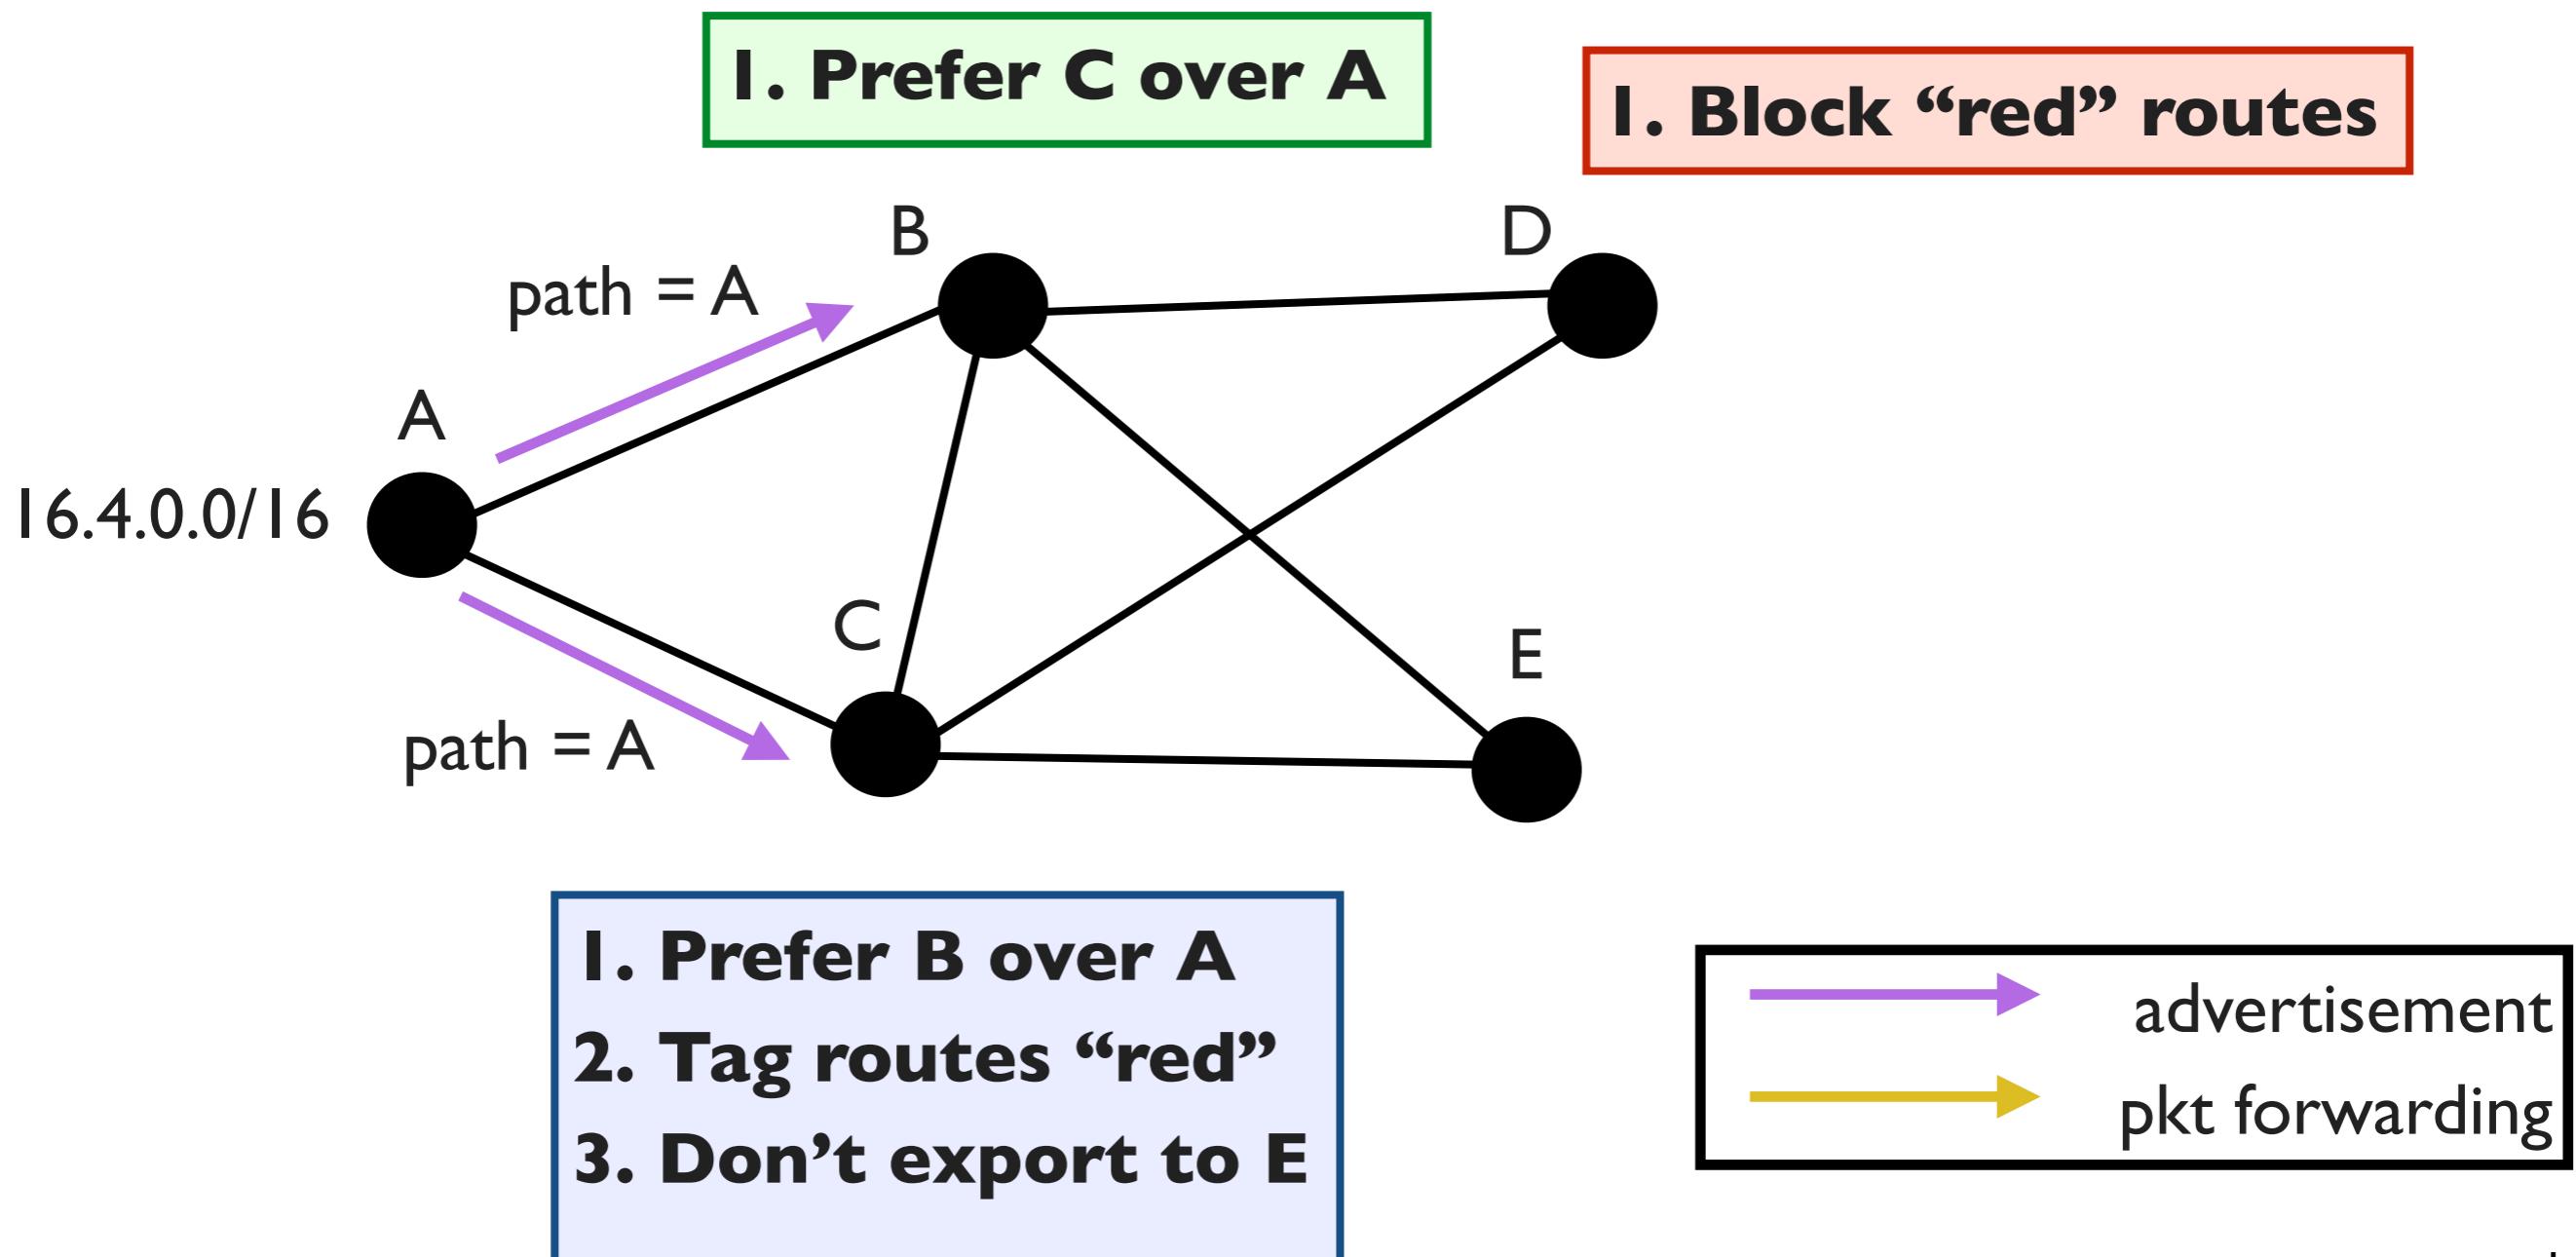
- Routers send advertisements for a destination (e.g., 16.4.0.0/16)
- Advertisements contain the AS path traffic would take
- Each router applies **local policy** to choose best route
 - 1.Import filters drop advertisements or modify attributes
 - 2.Highest preference is chosen. Path length as a tie breaker
 - 3.Export filters drop advertisements or modify attributes



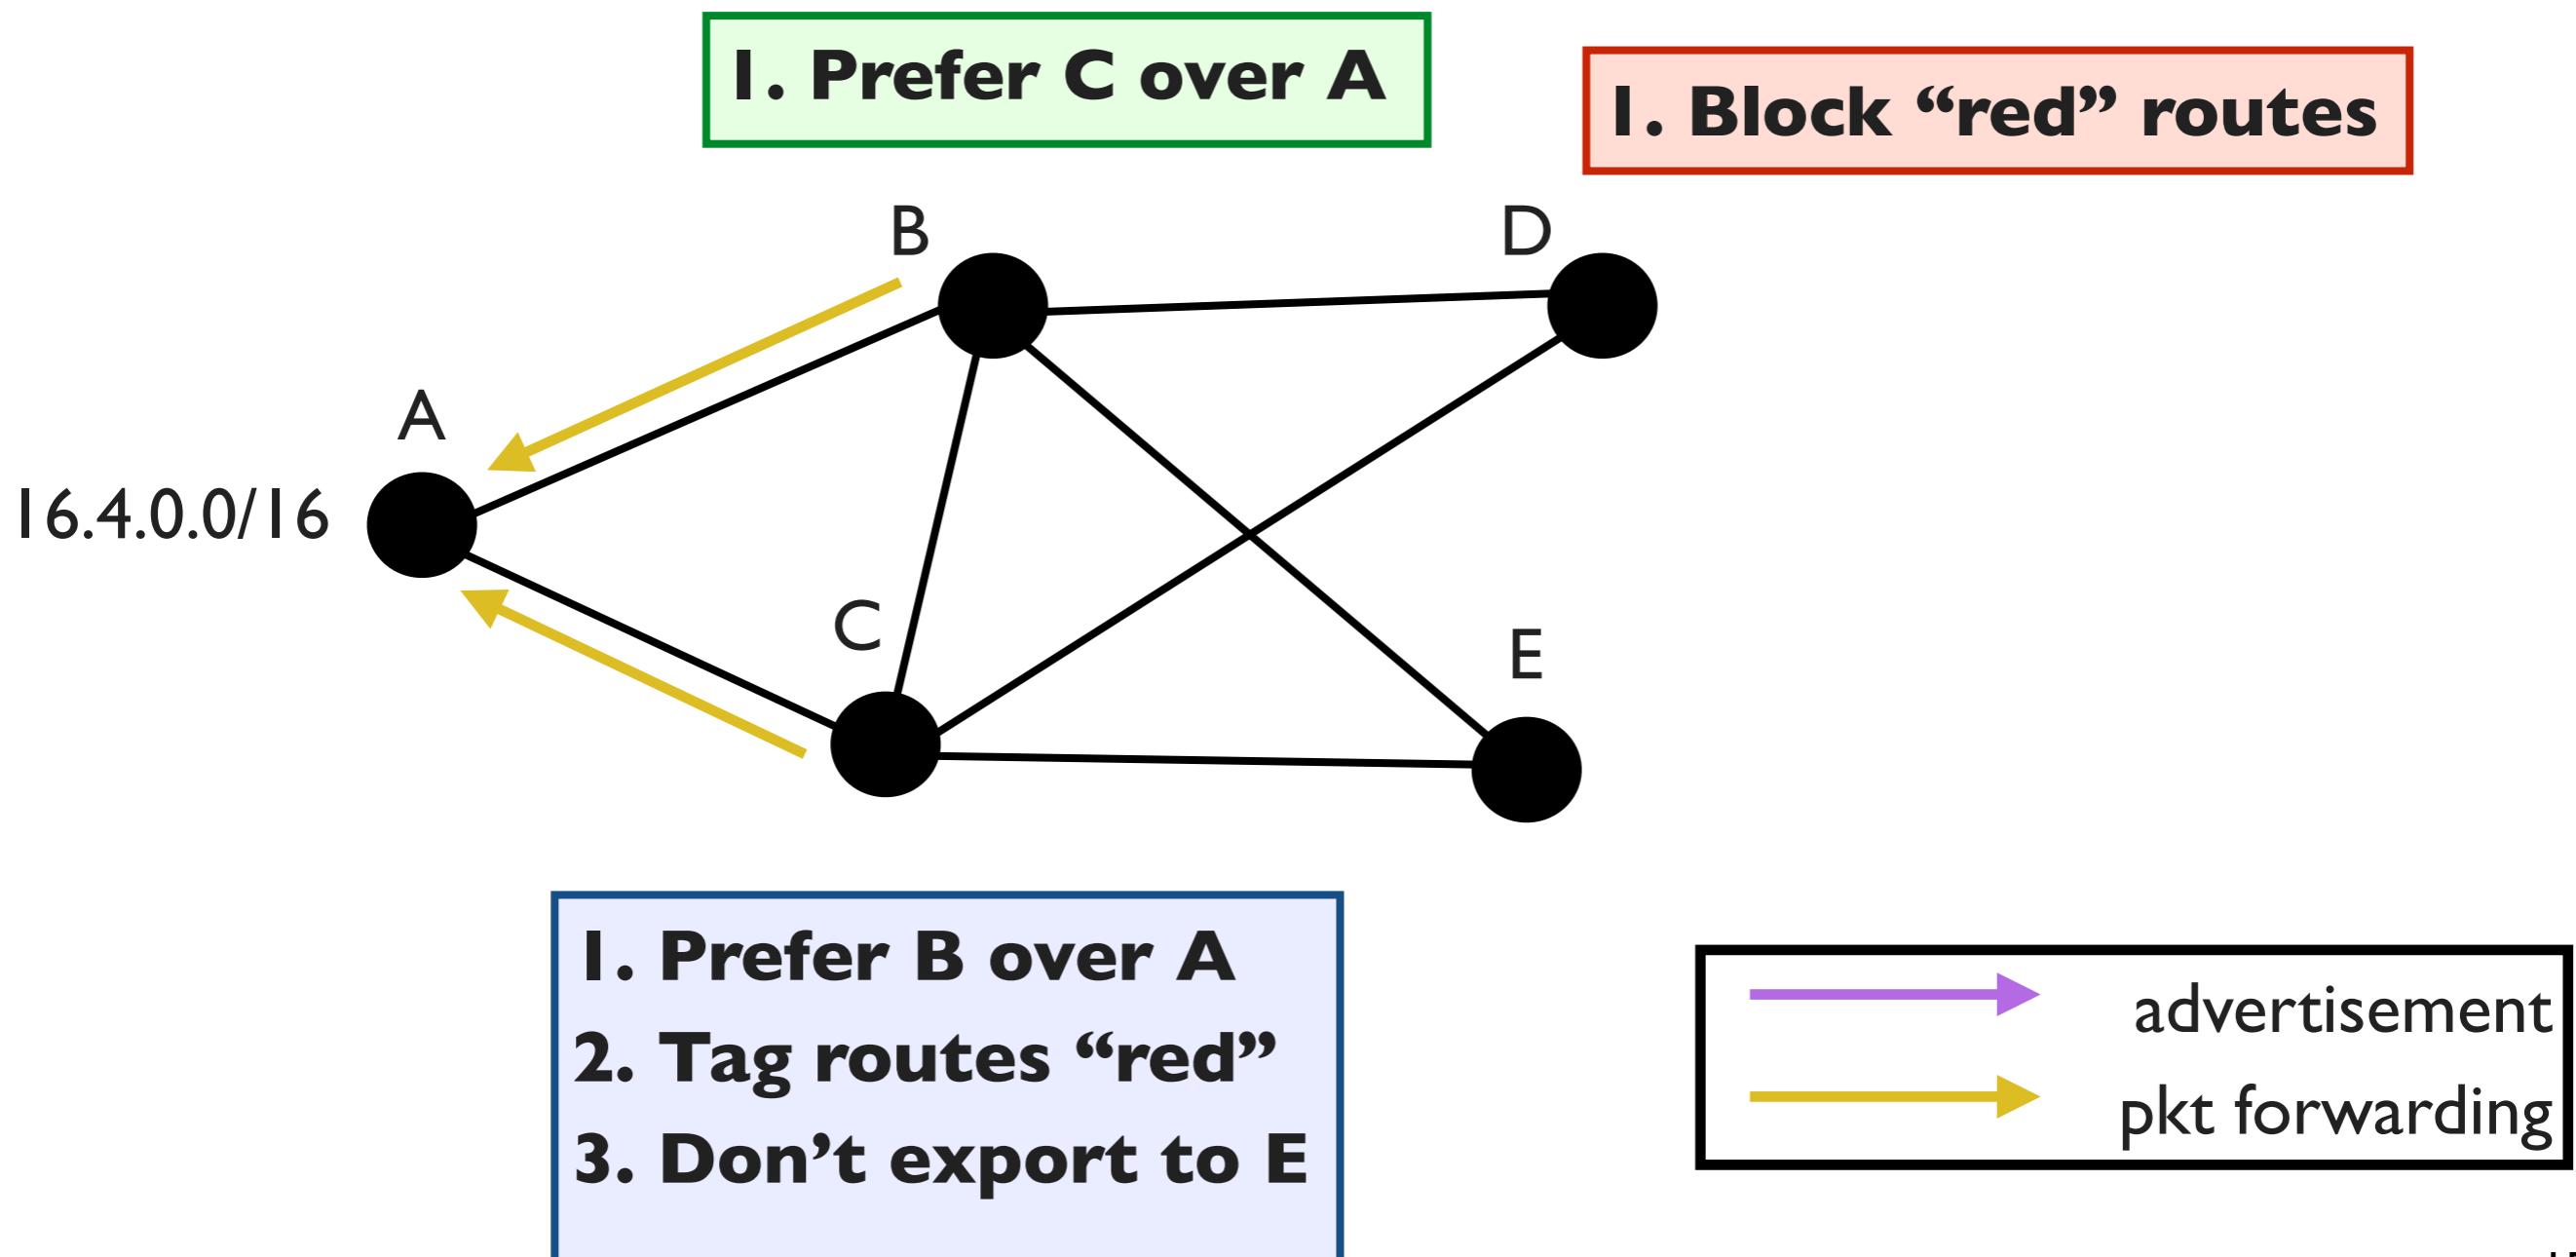
Border Gateway Protocol (BGP)



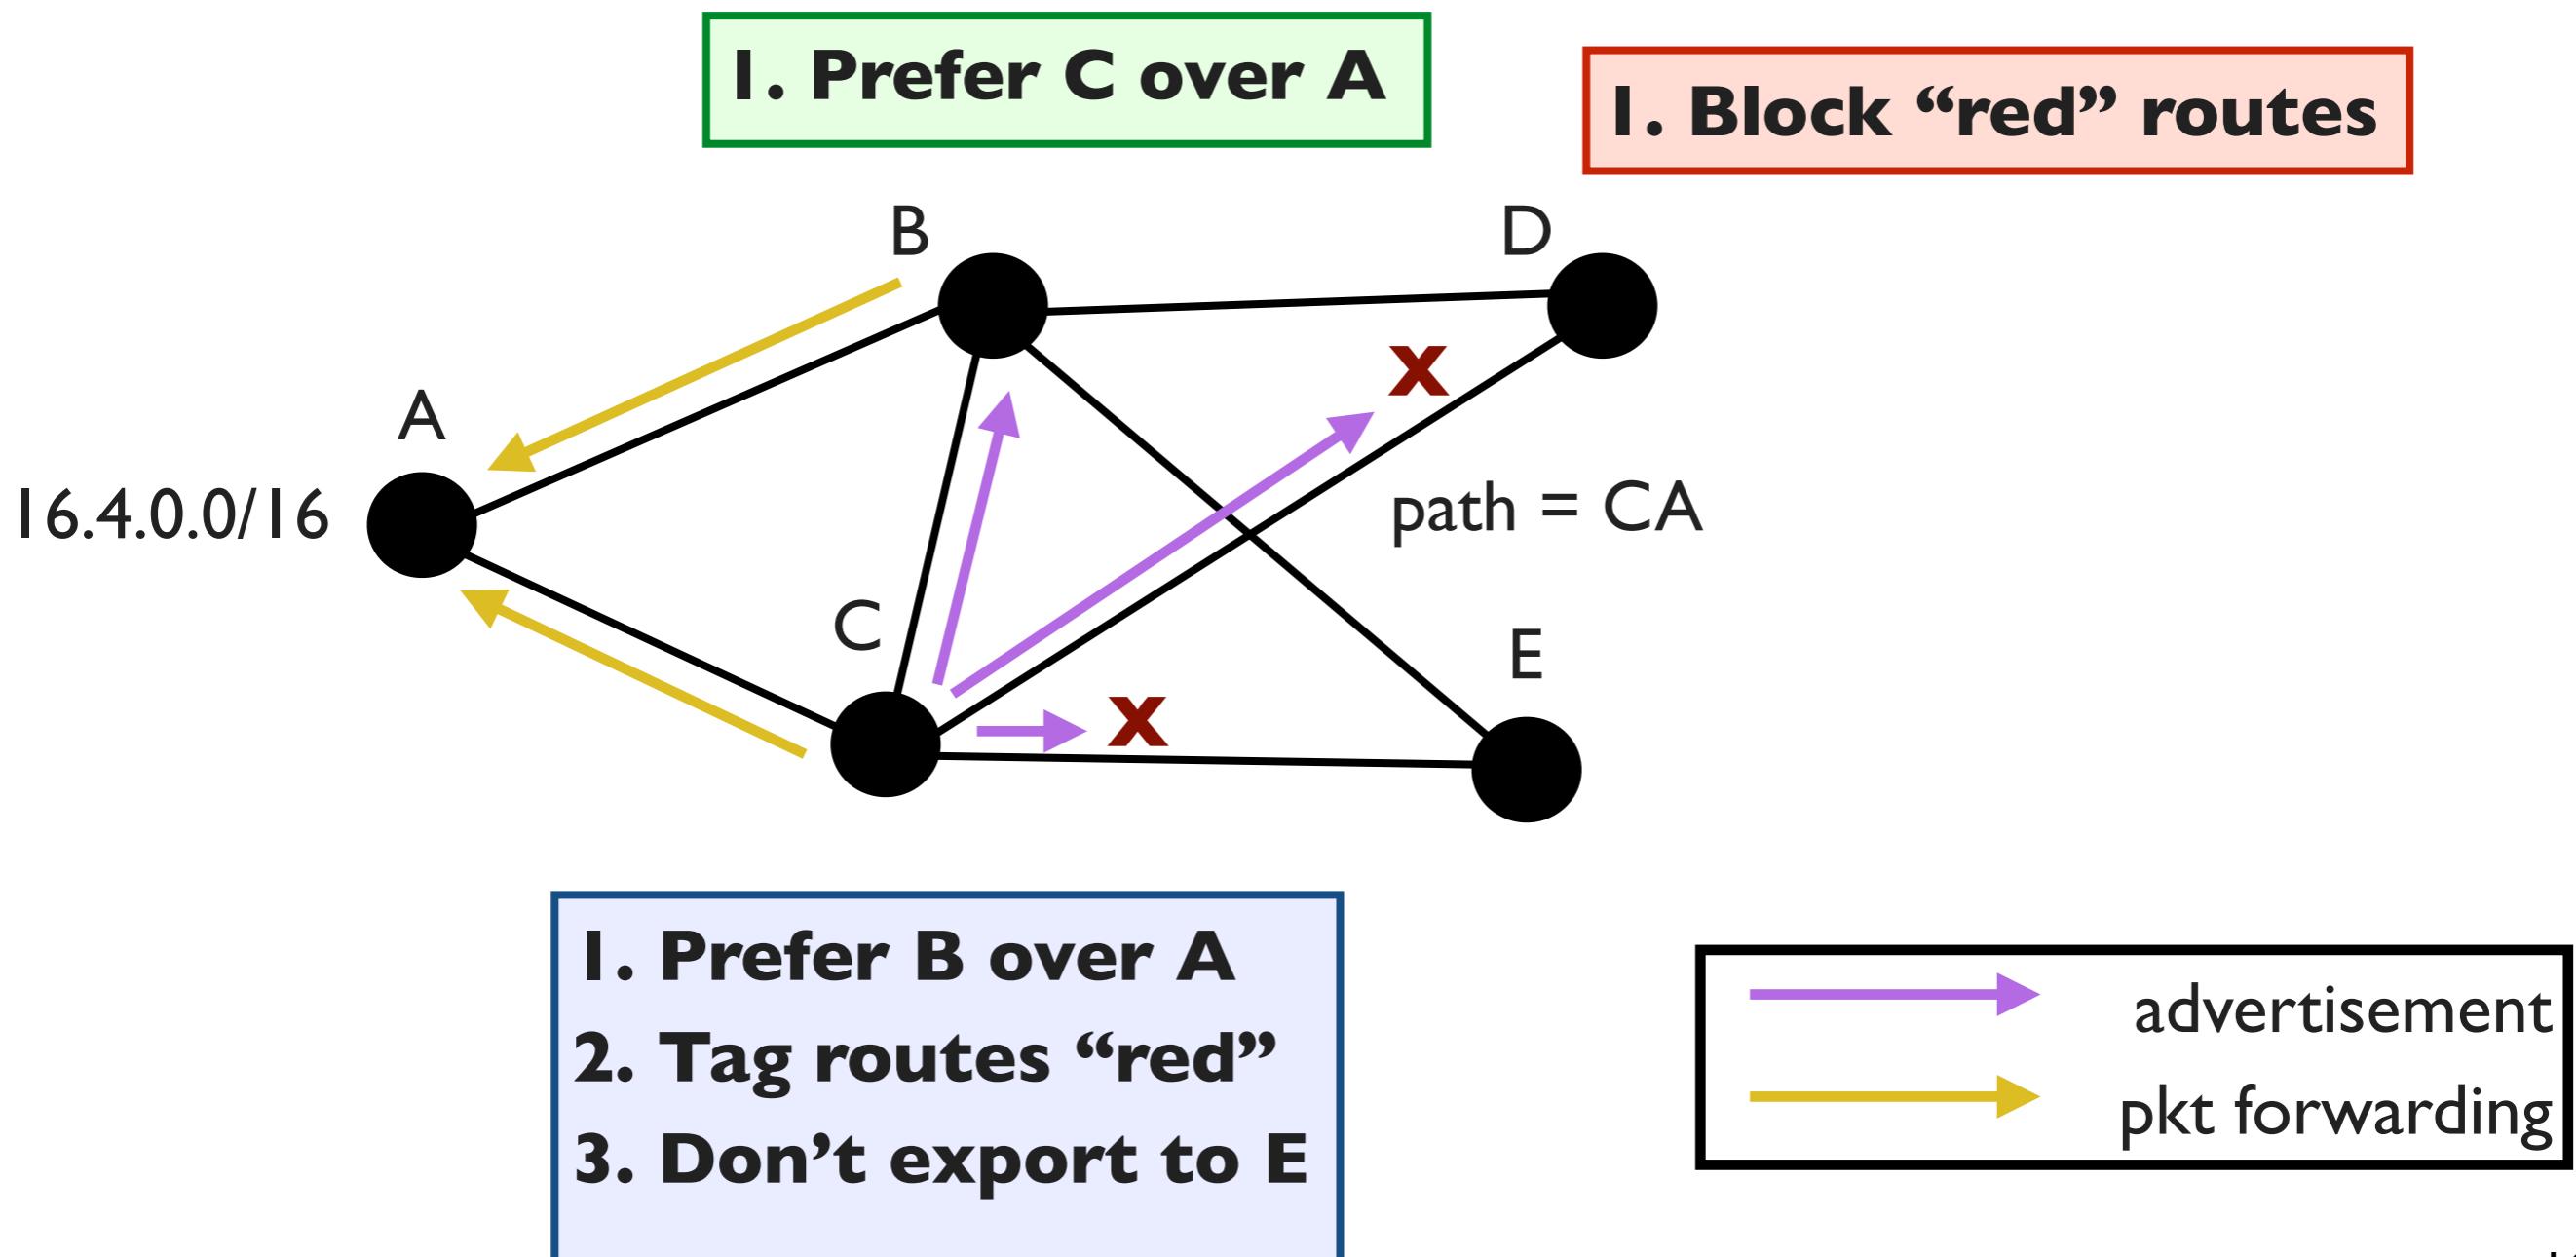
Border Gateway Protocol (BGP)



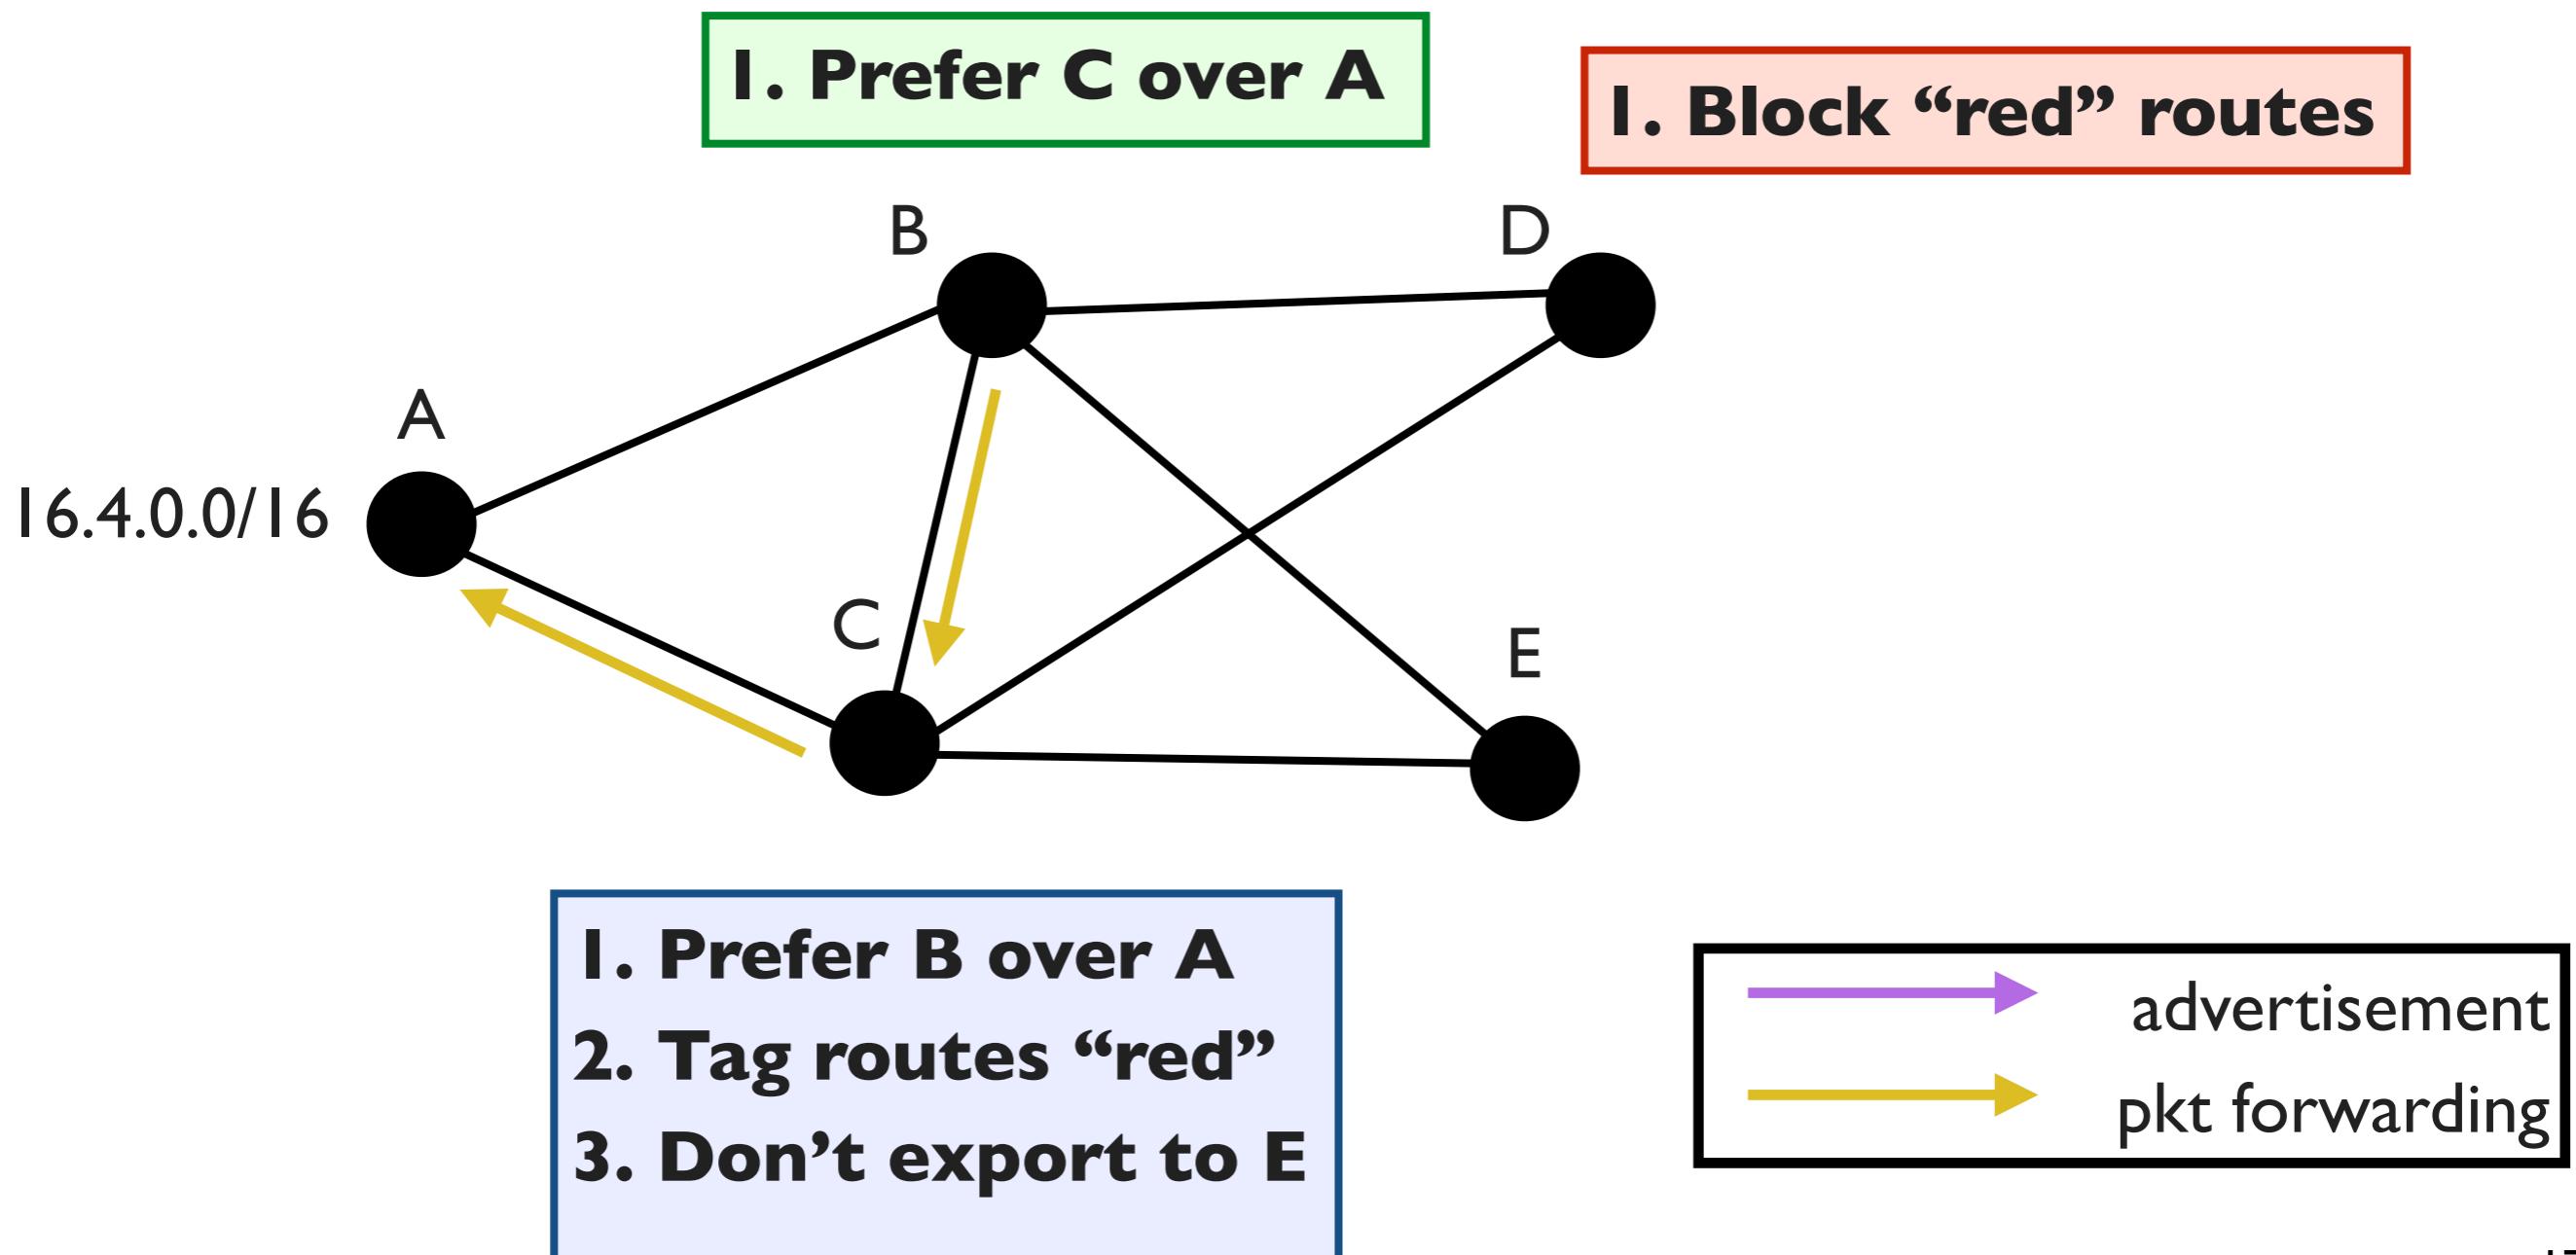
Border Gateway Protocol (BGP)



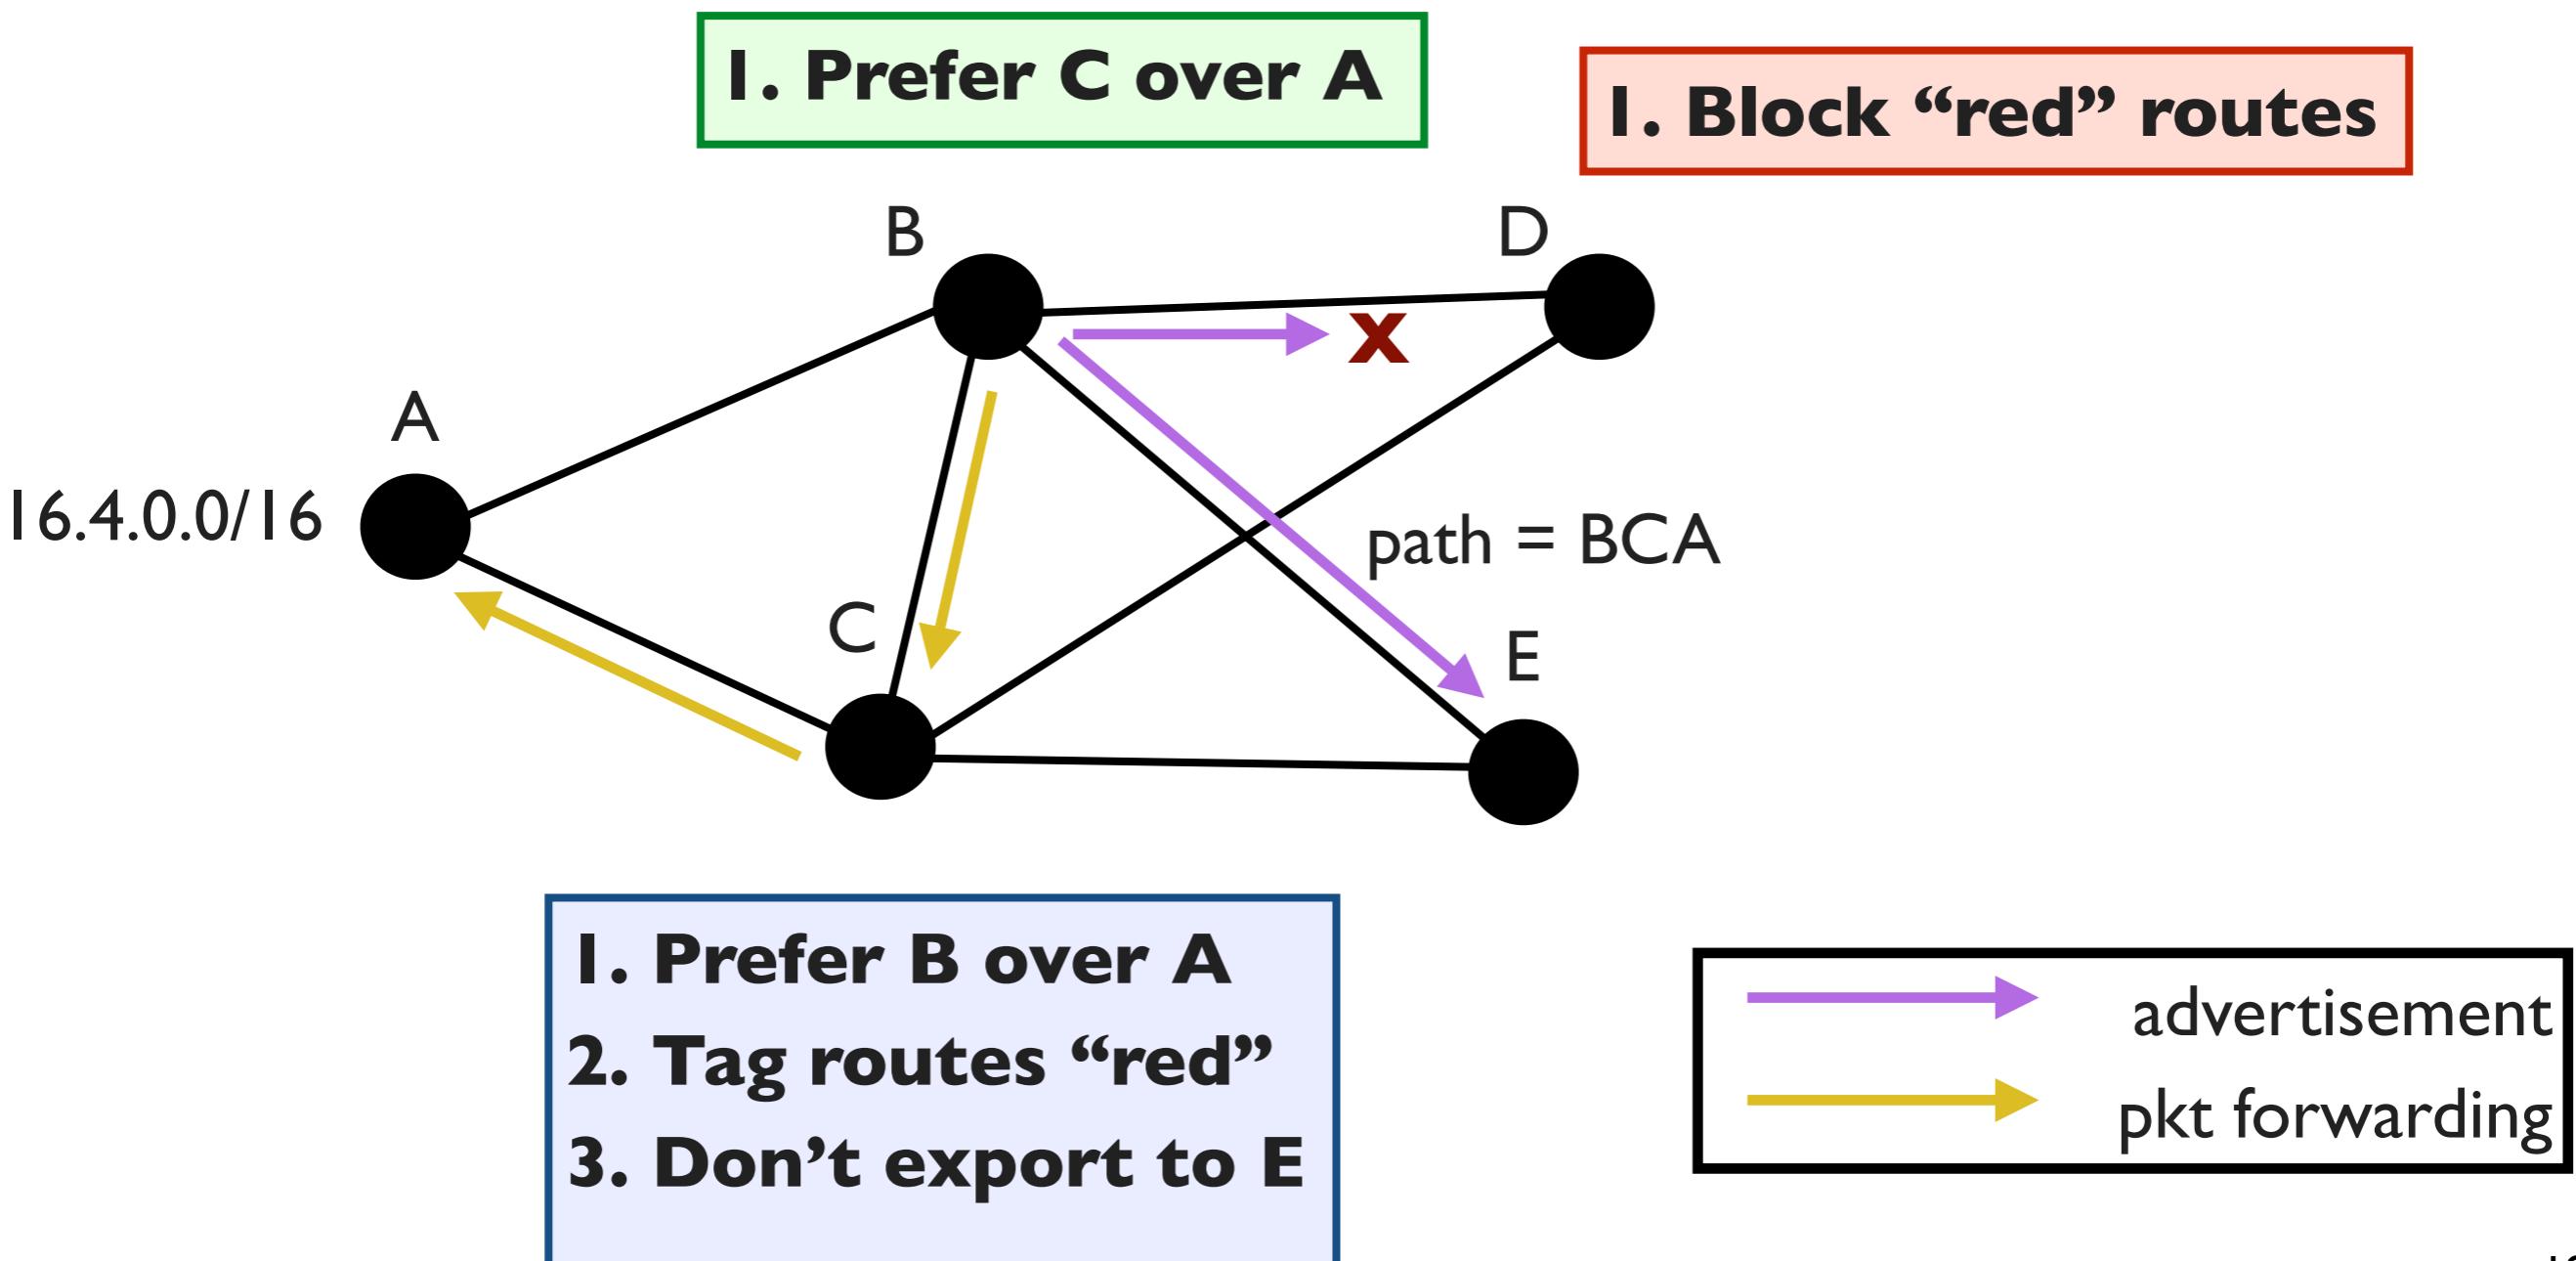
Border Gateway Protocol (BGP)



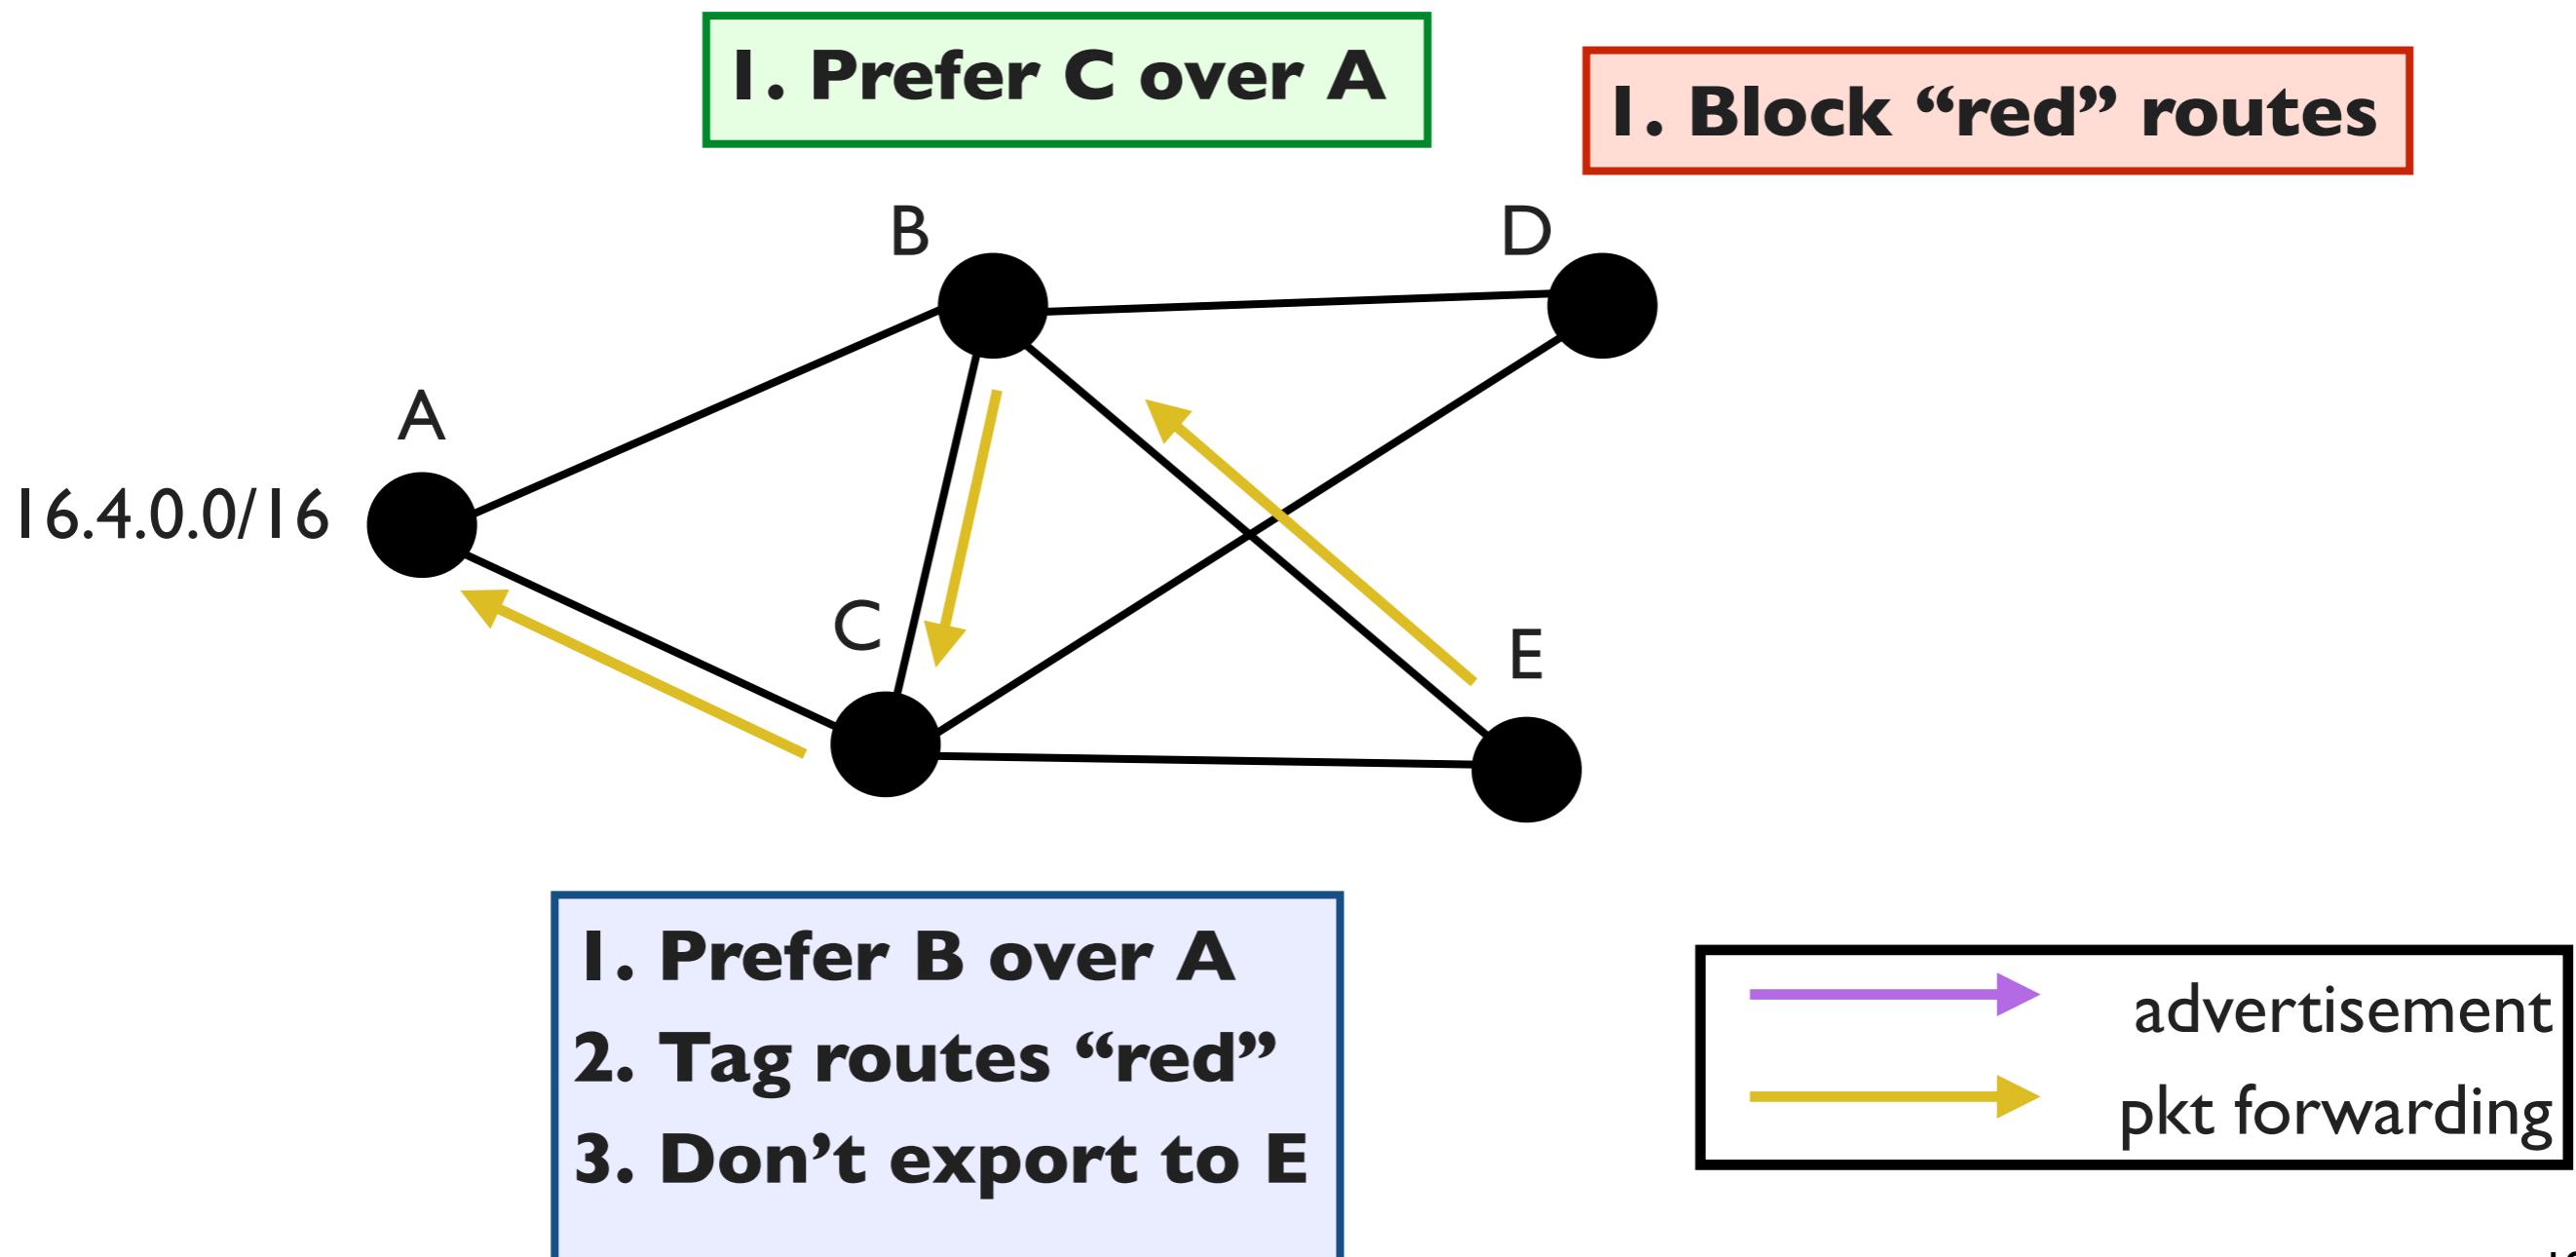
Border Gateway Protocol (BGP)



Border Gateway Protocol (BGP)



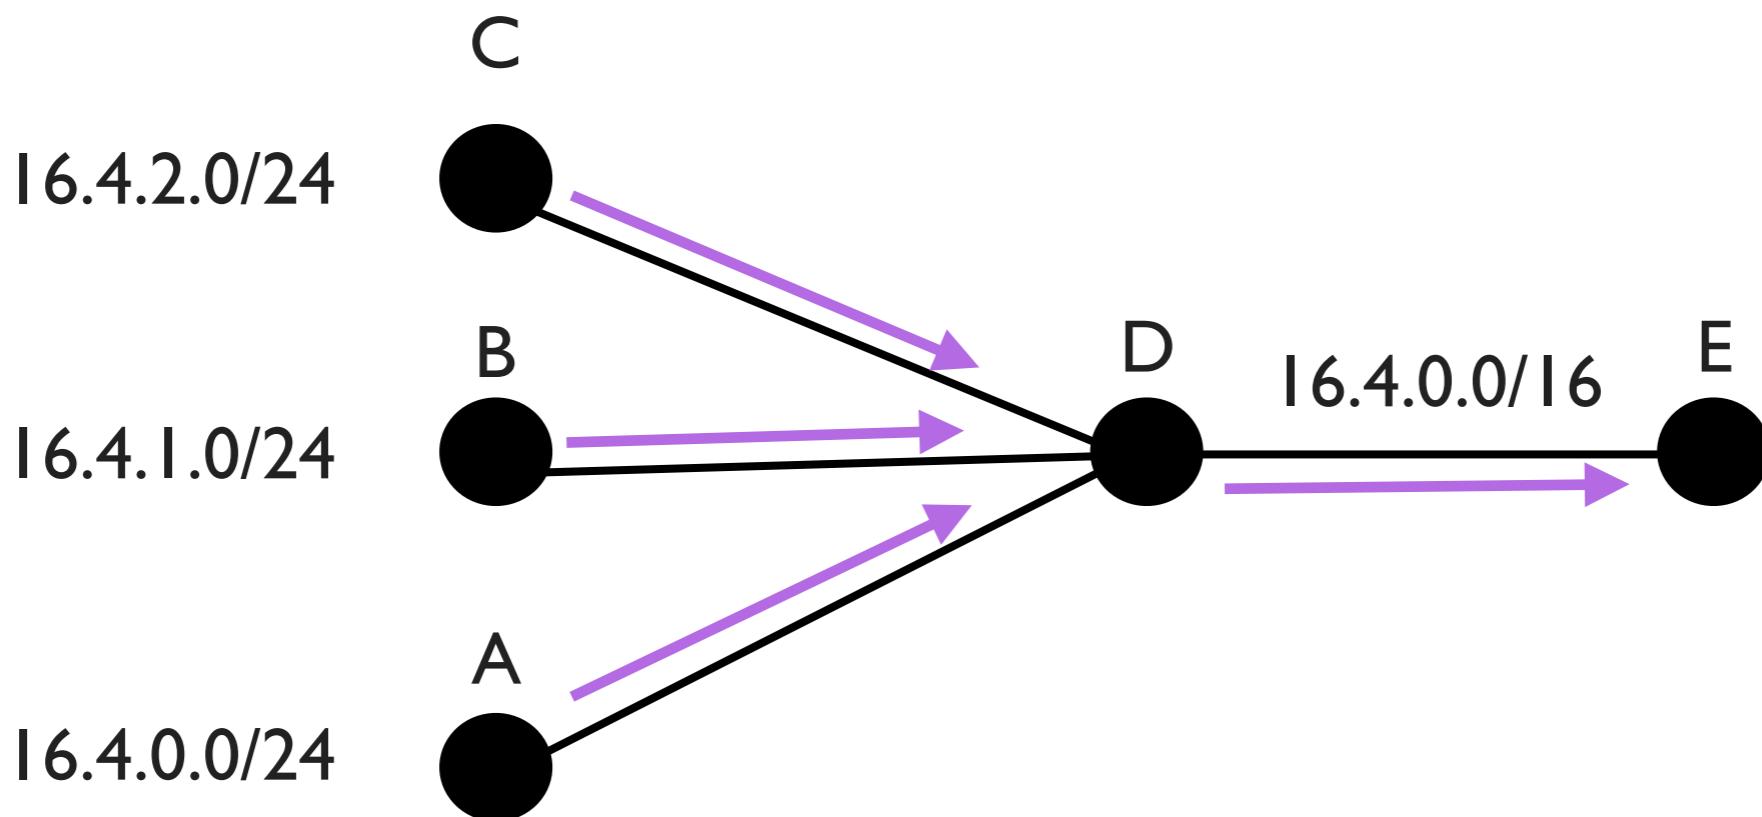
Border Gateway Protocol (BGP)



Border Gateway Protocol (BGP)

Aggregation (Route summarization)

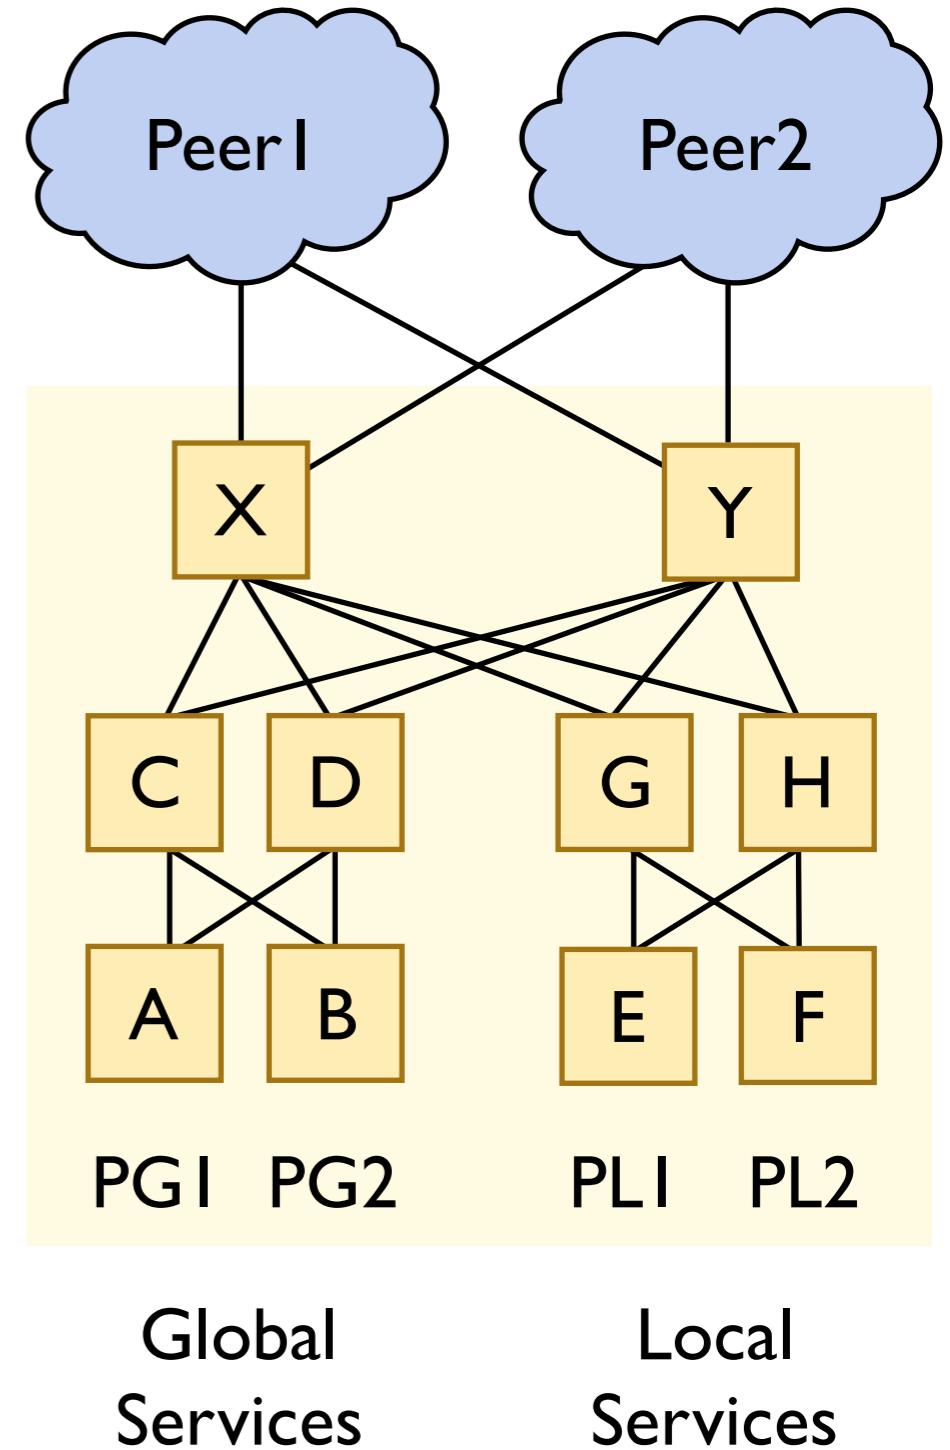
- Advertise a more general subnetwork
- Information hiding for routers
- Reduces router table size and improves stability



Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers



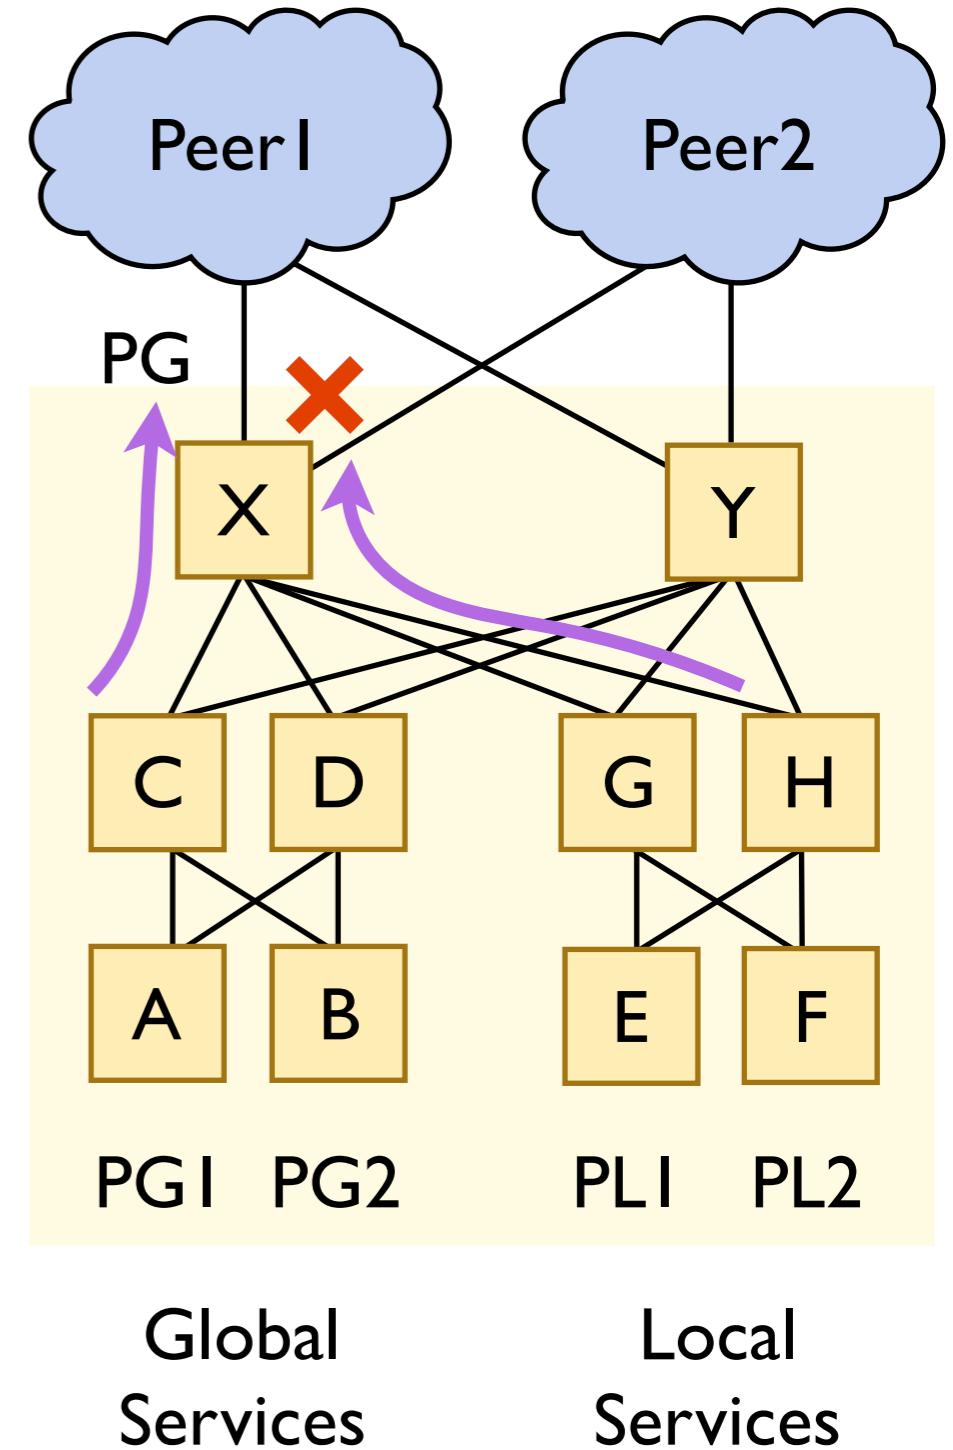
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG



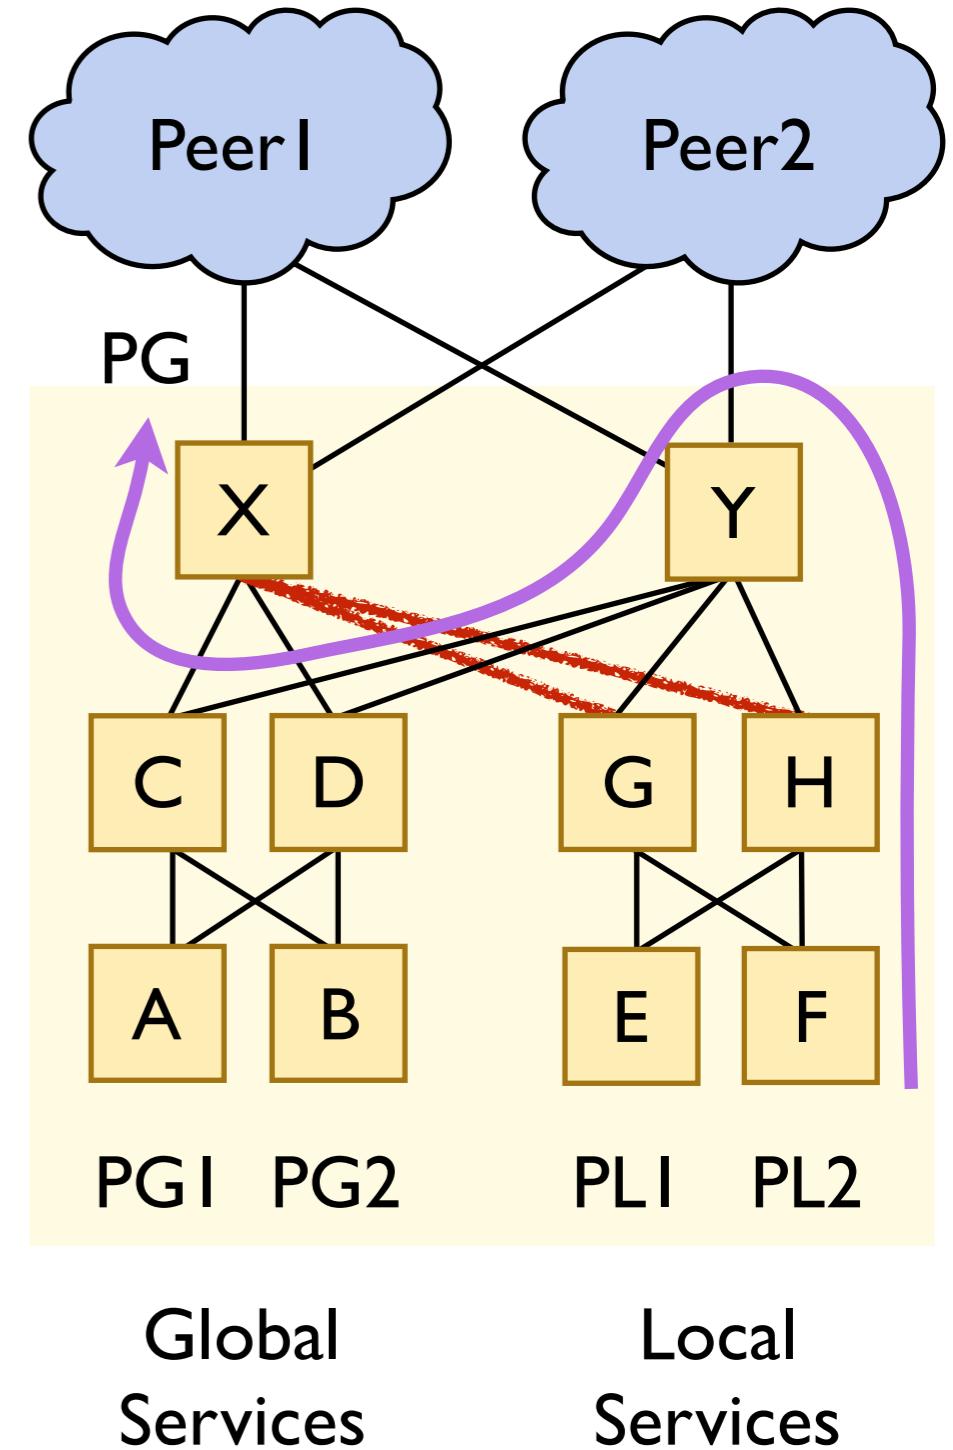
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG



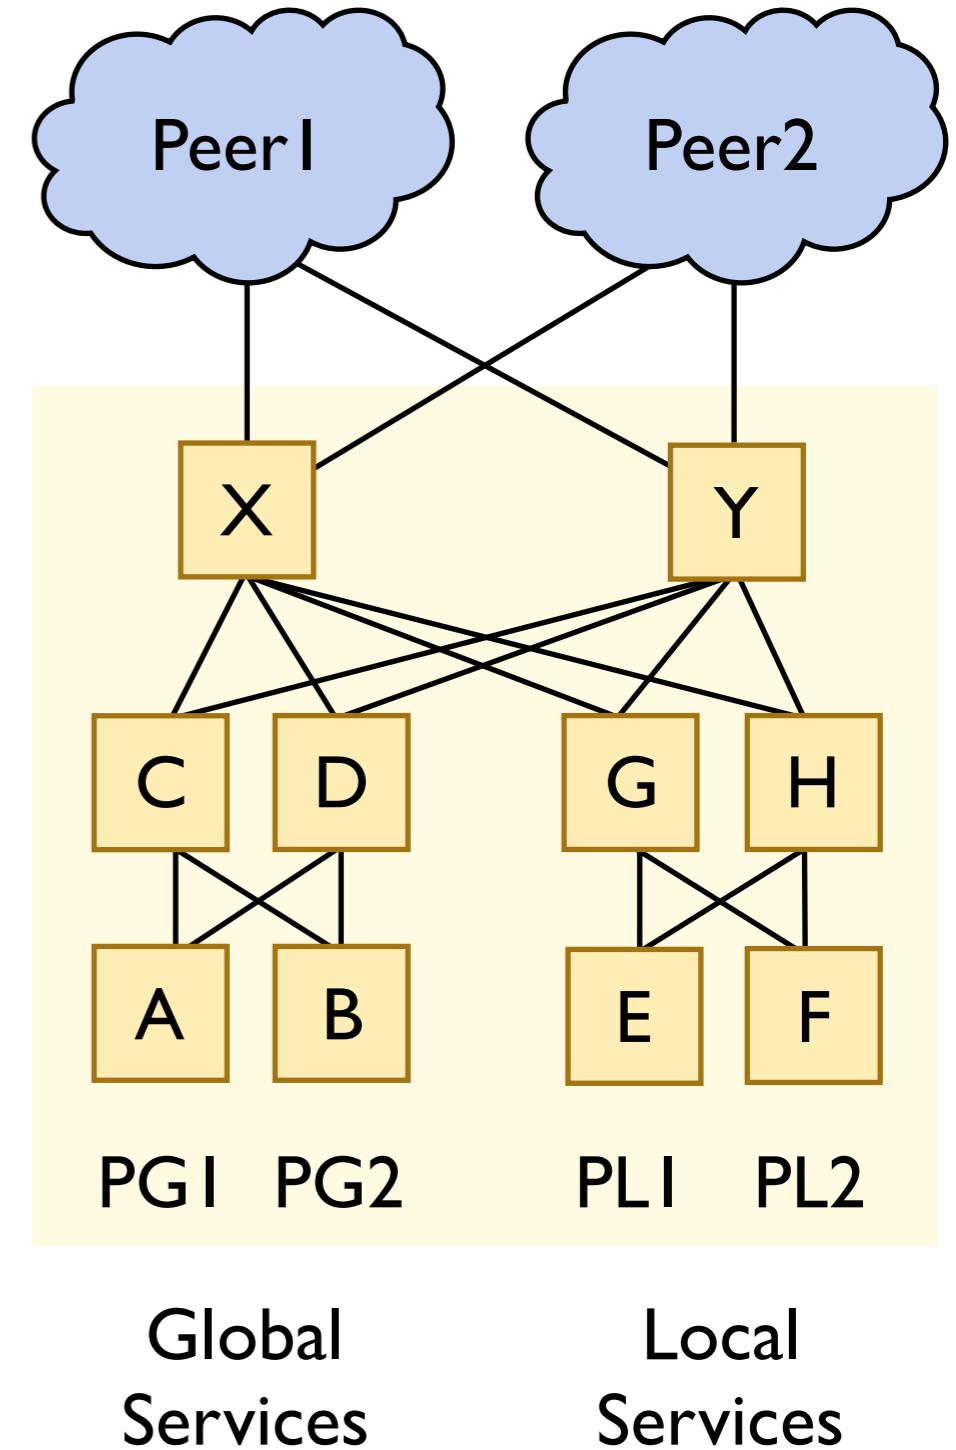
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG
- X,Y block routes through each other



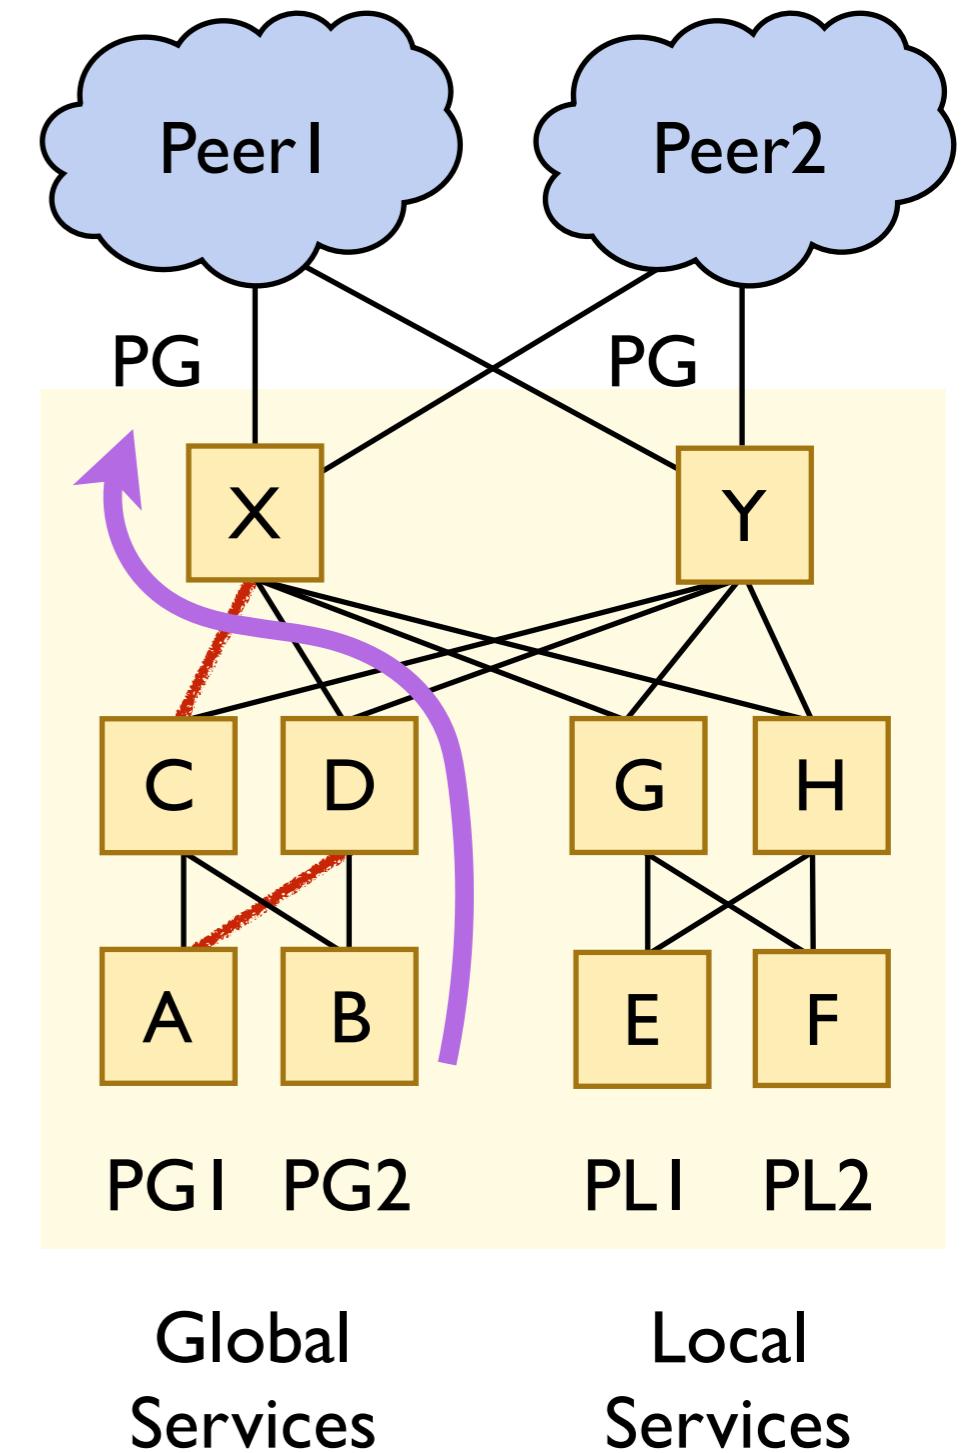
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG
- X,Y block routes through each other



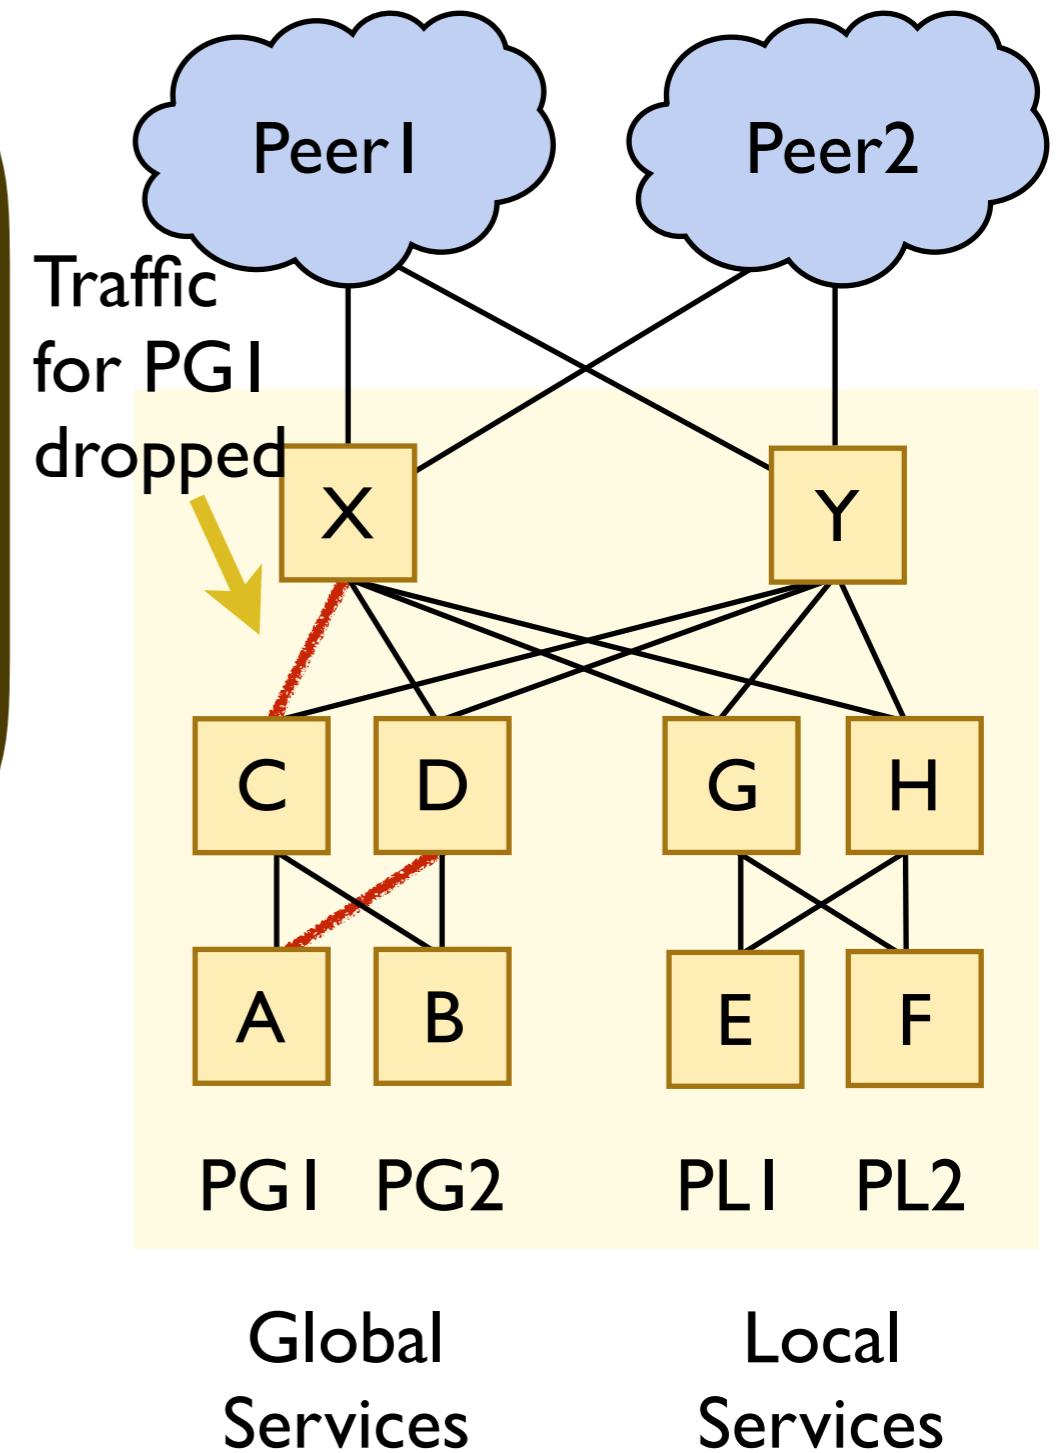
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG
- X,Y block routes through each other



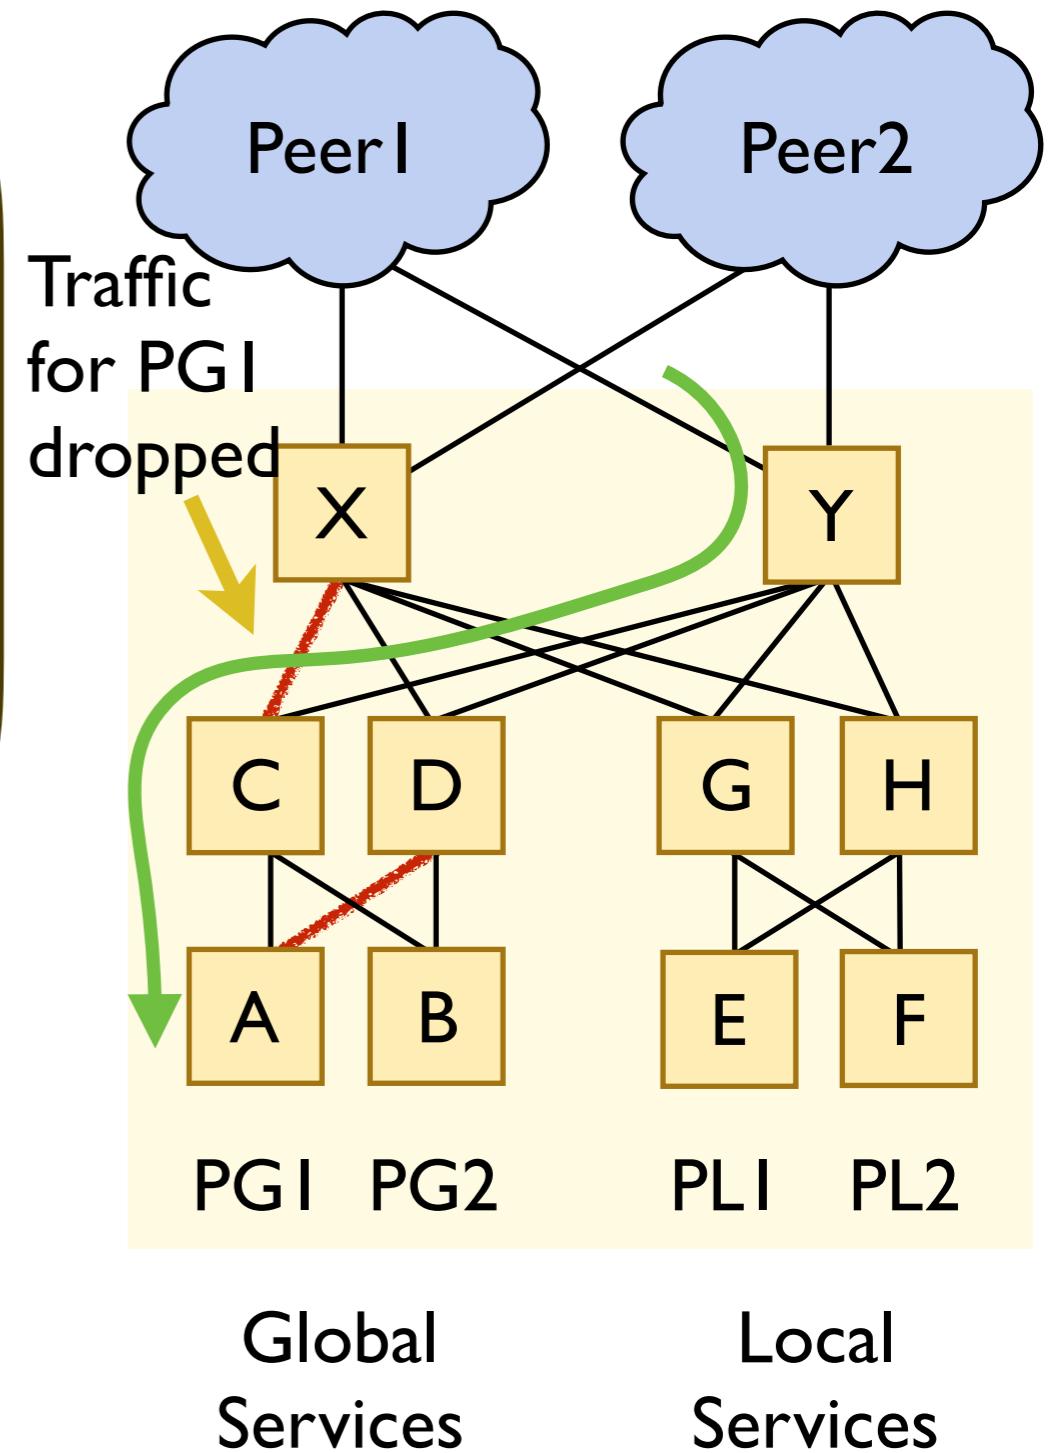
Example I: A Data Center Network

Goals

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- Don't export from G, H to external
- Aggregate externally as PG
- X, Y block routes through each other

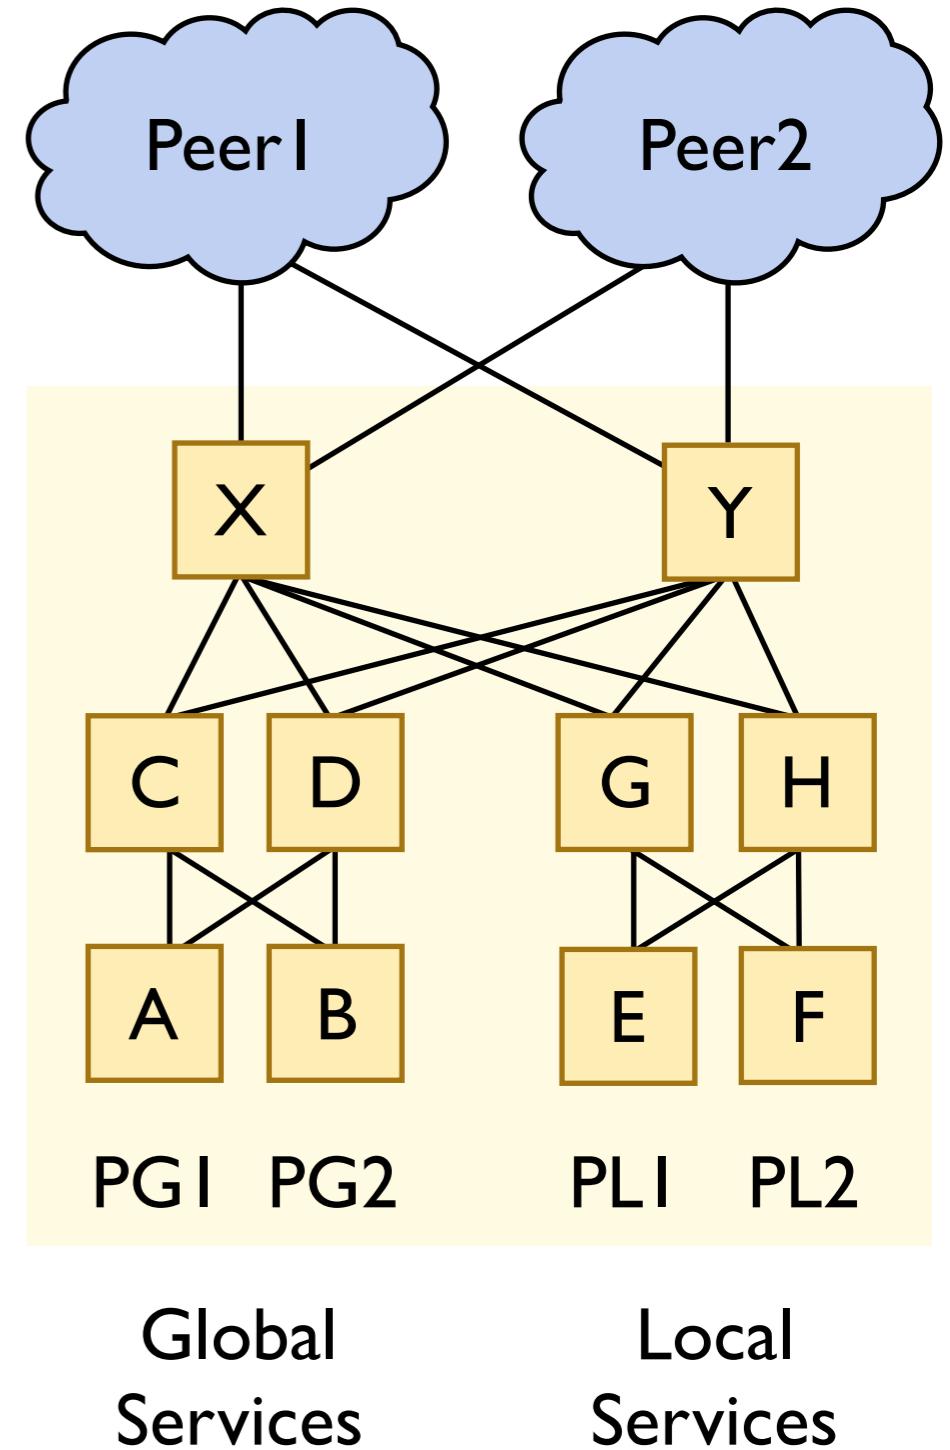


Aggregation-Induced Black Hole!

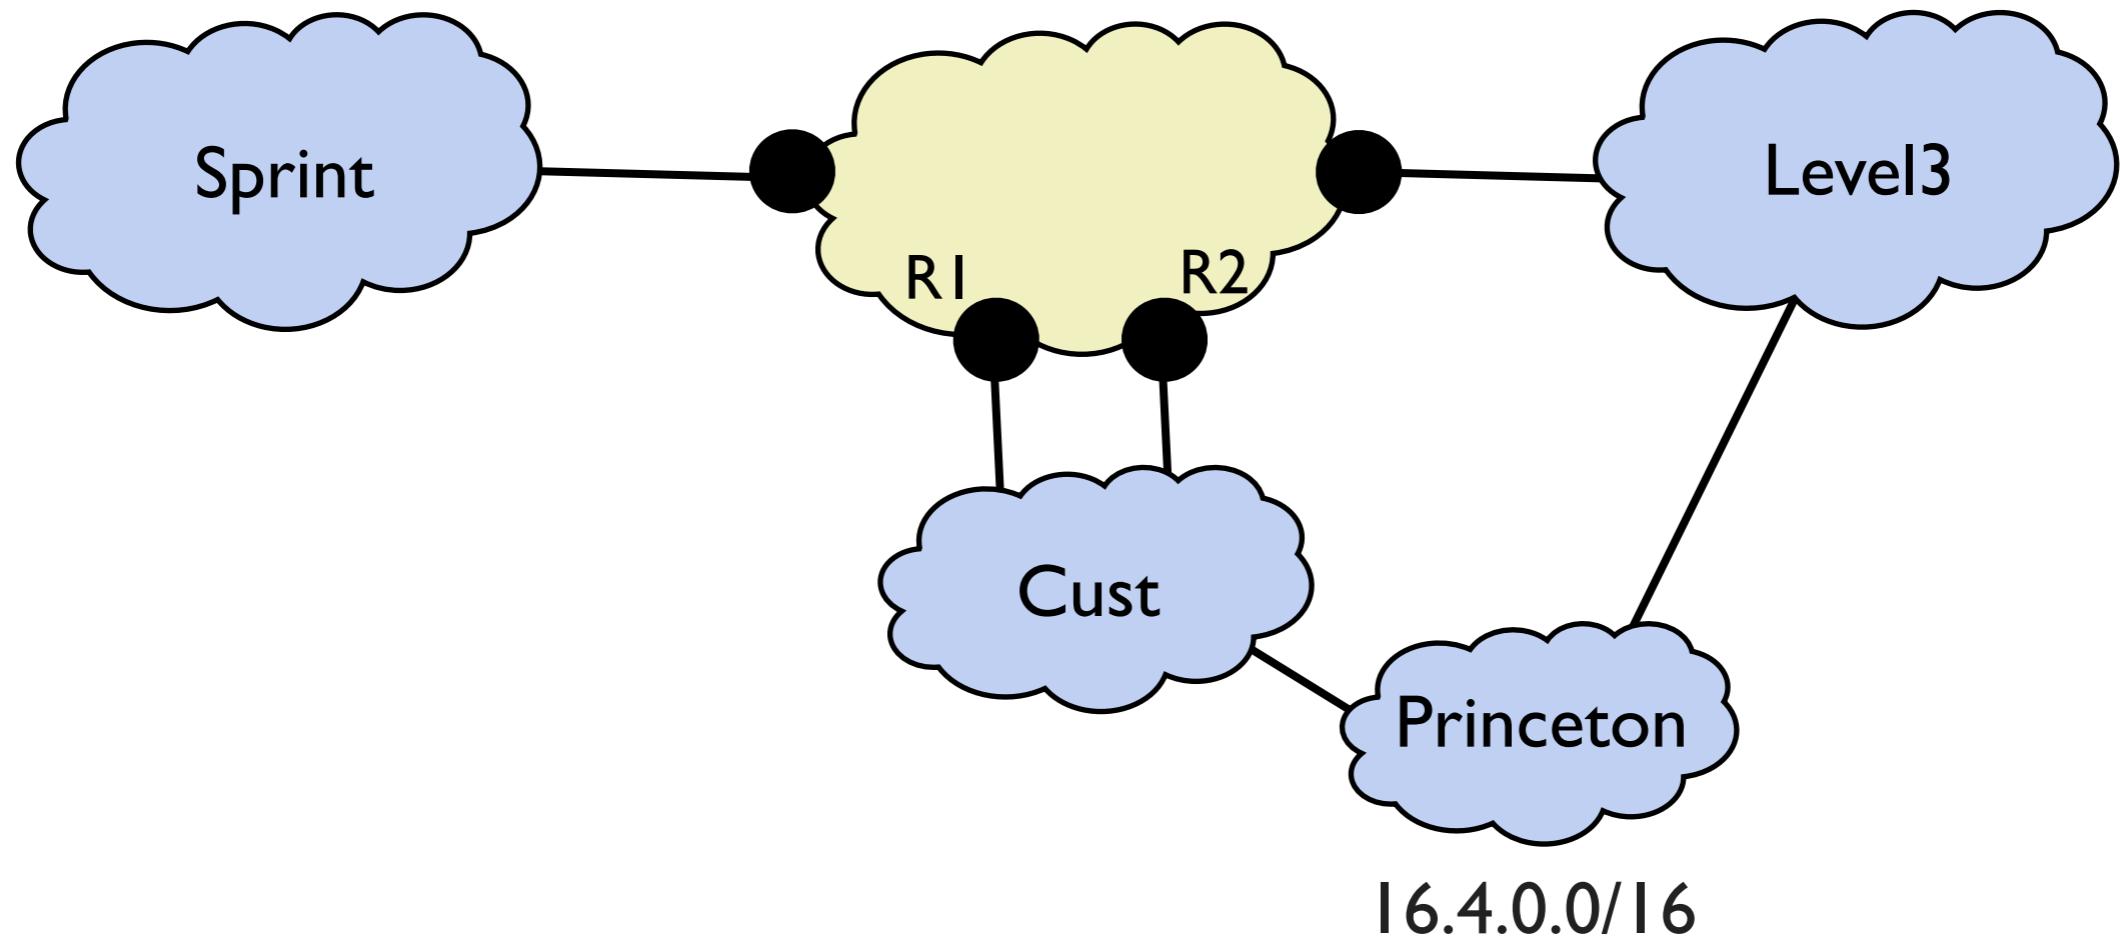
Example I: A Data Center Network

Goals

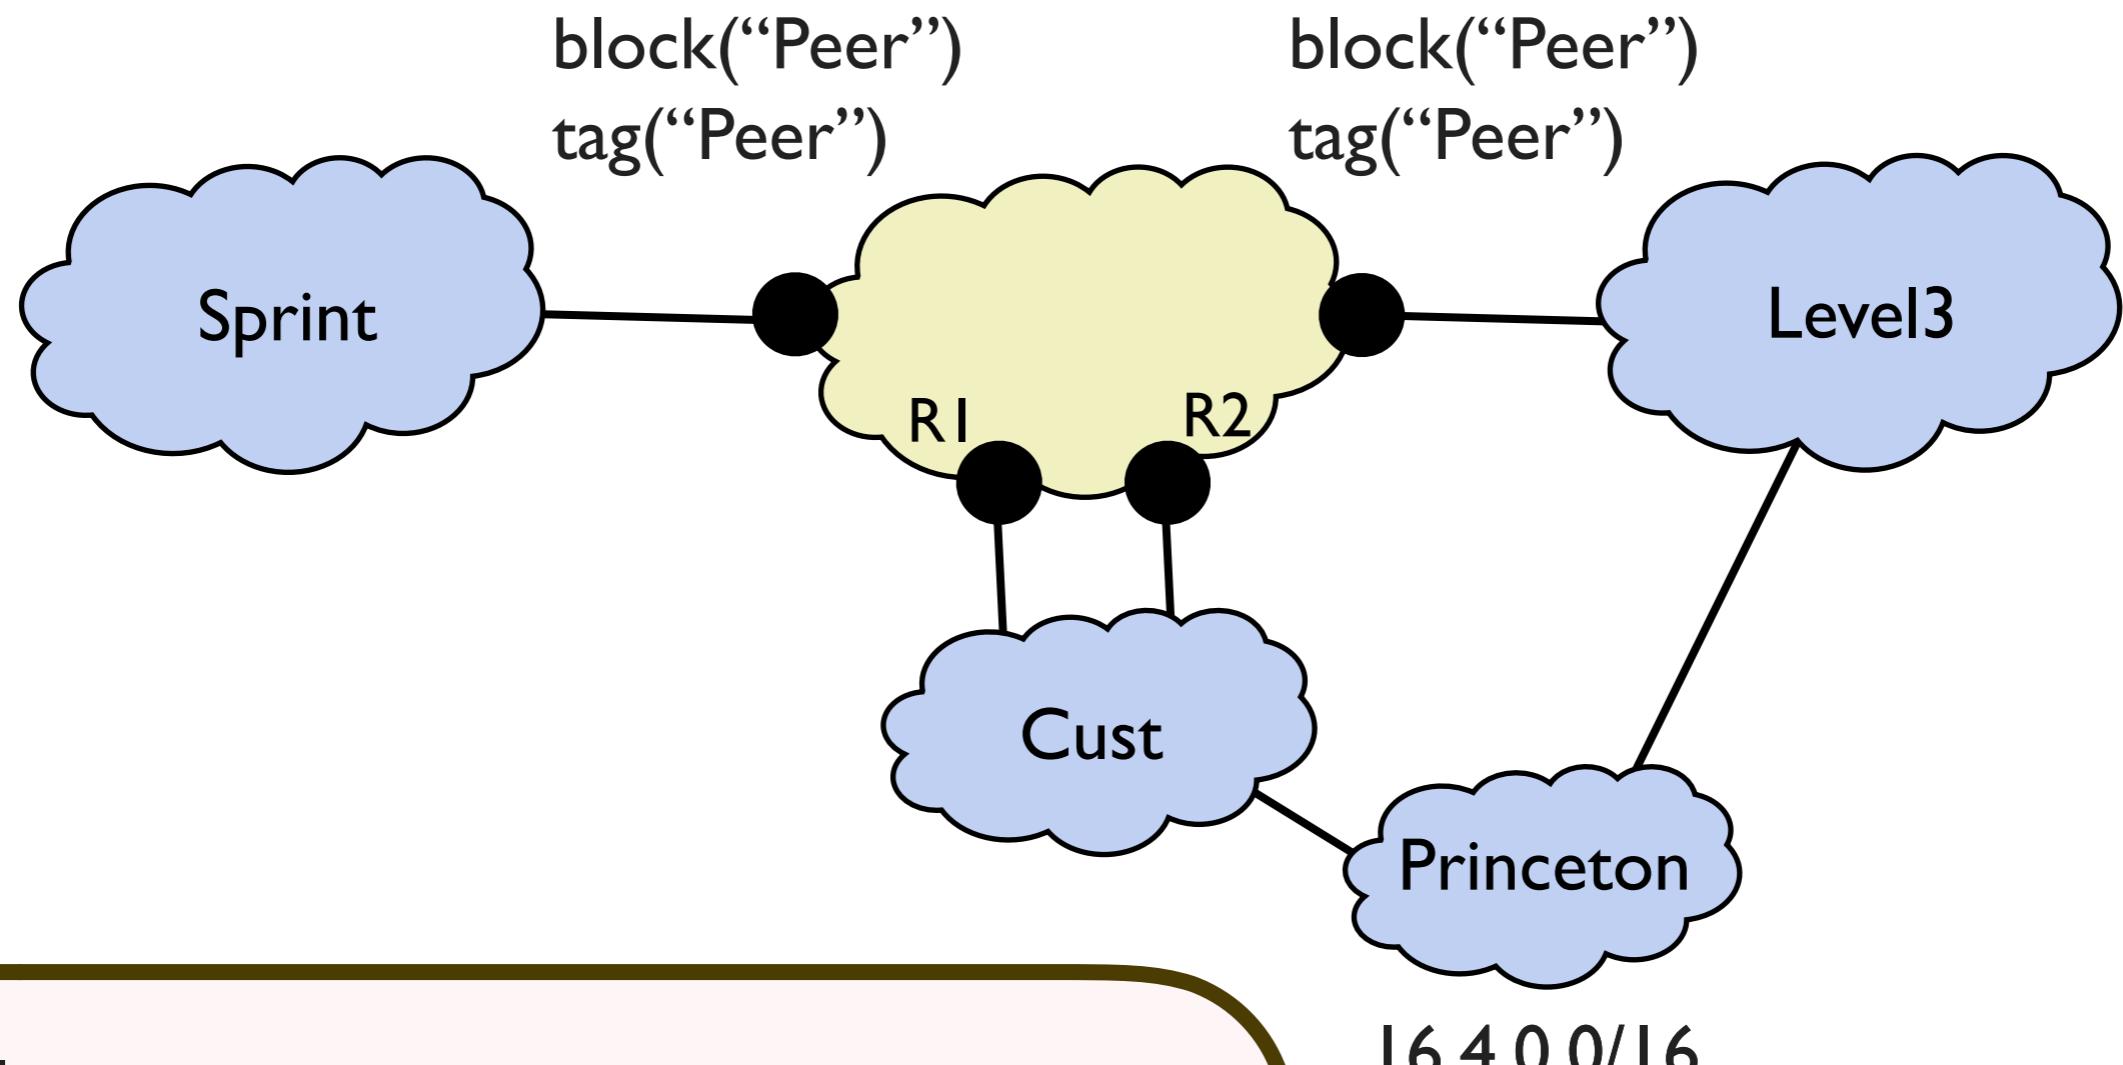
- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers



Example 2: Backbone Network



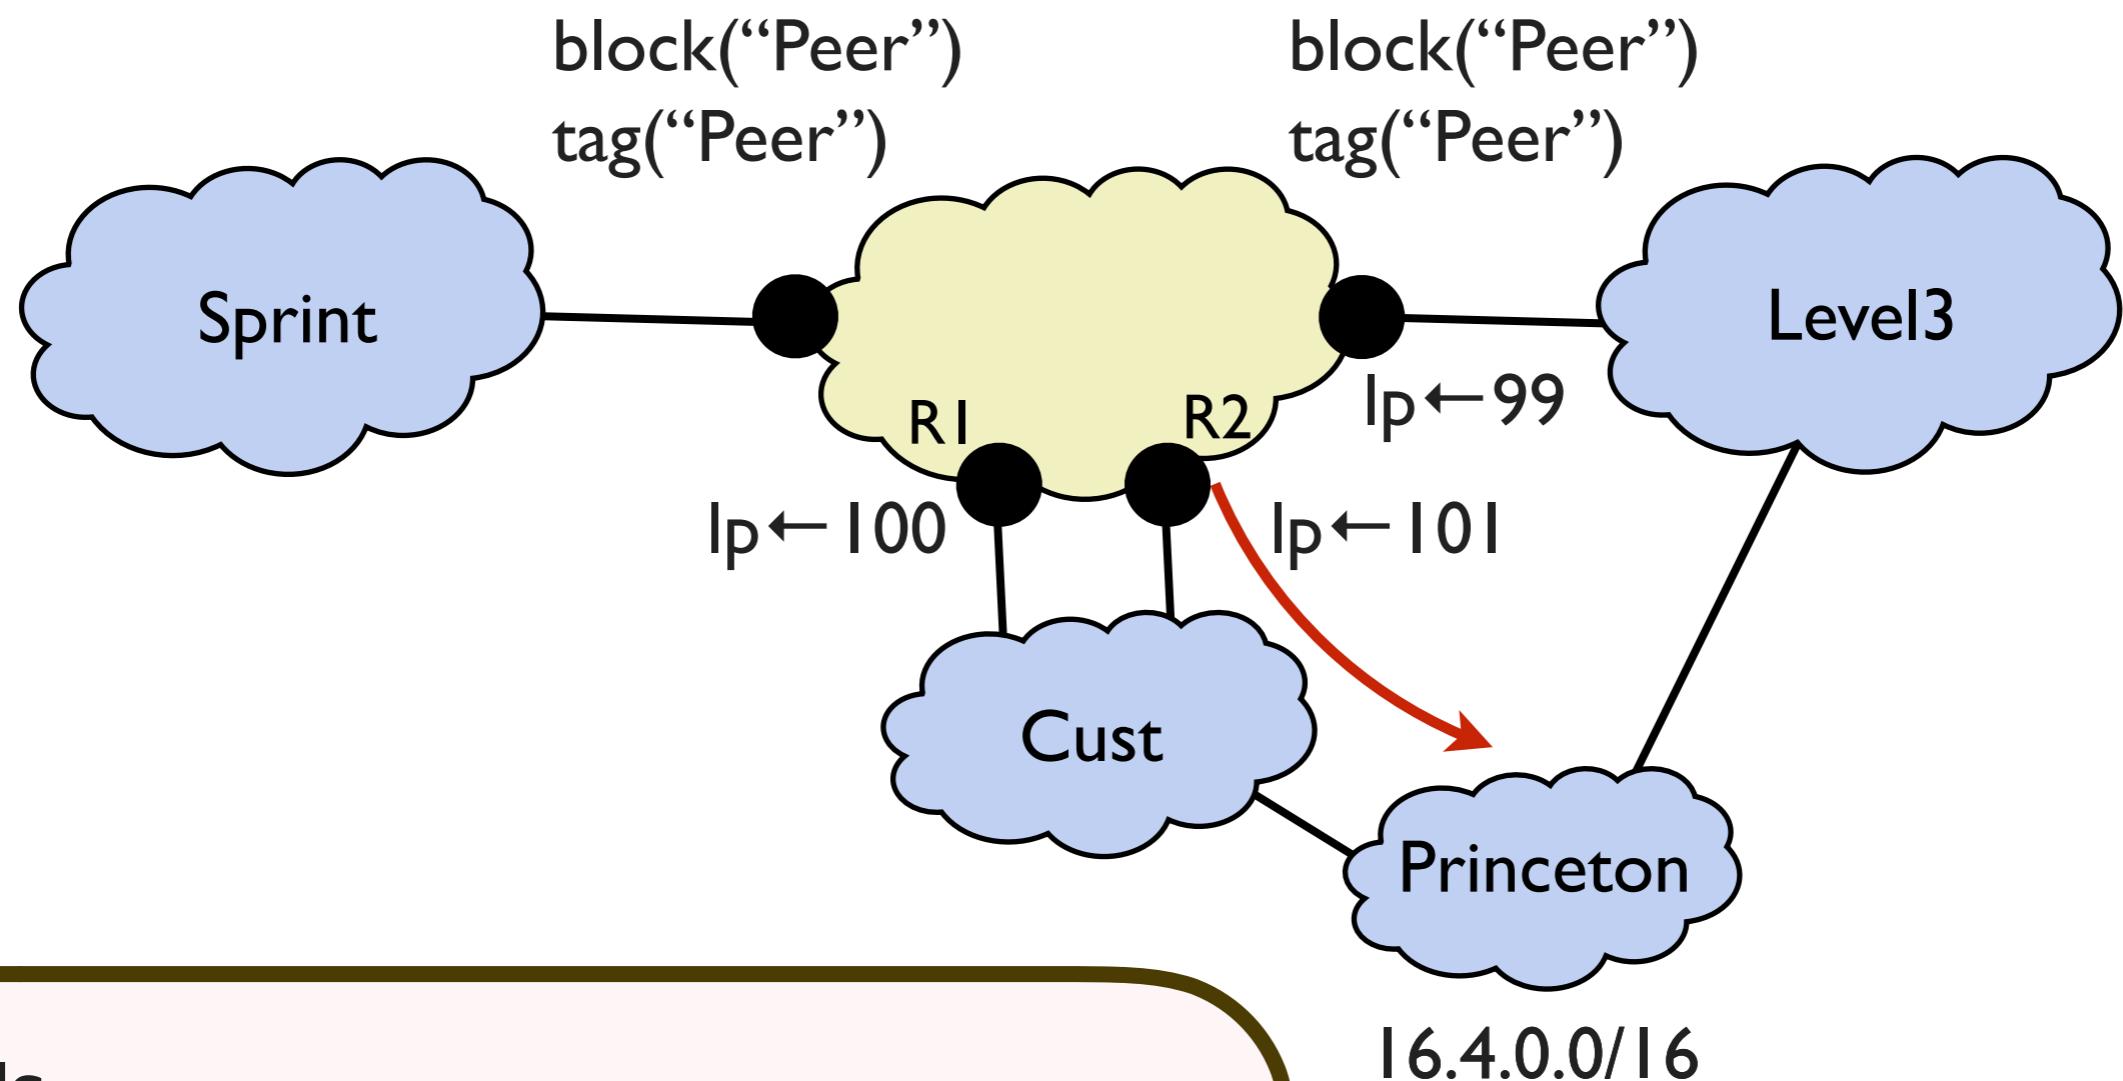
Example 2: Backbone Network



Goals

- Prevent transit between peers (\$\$\$)

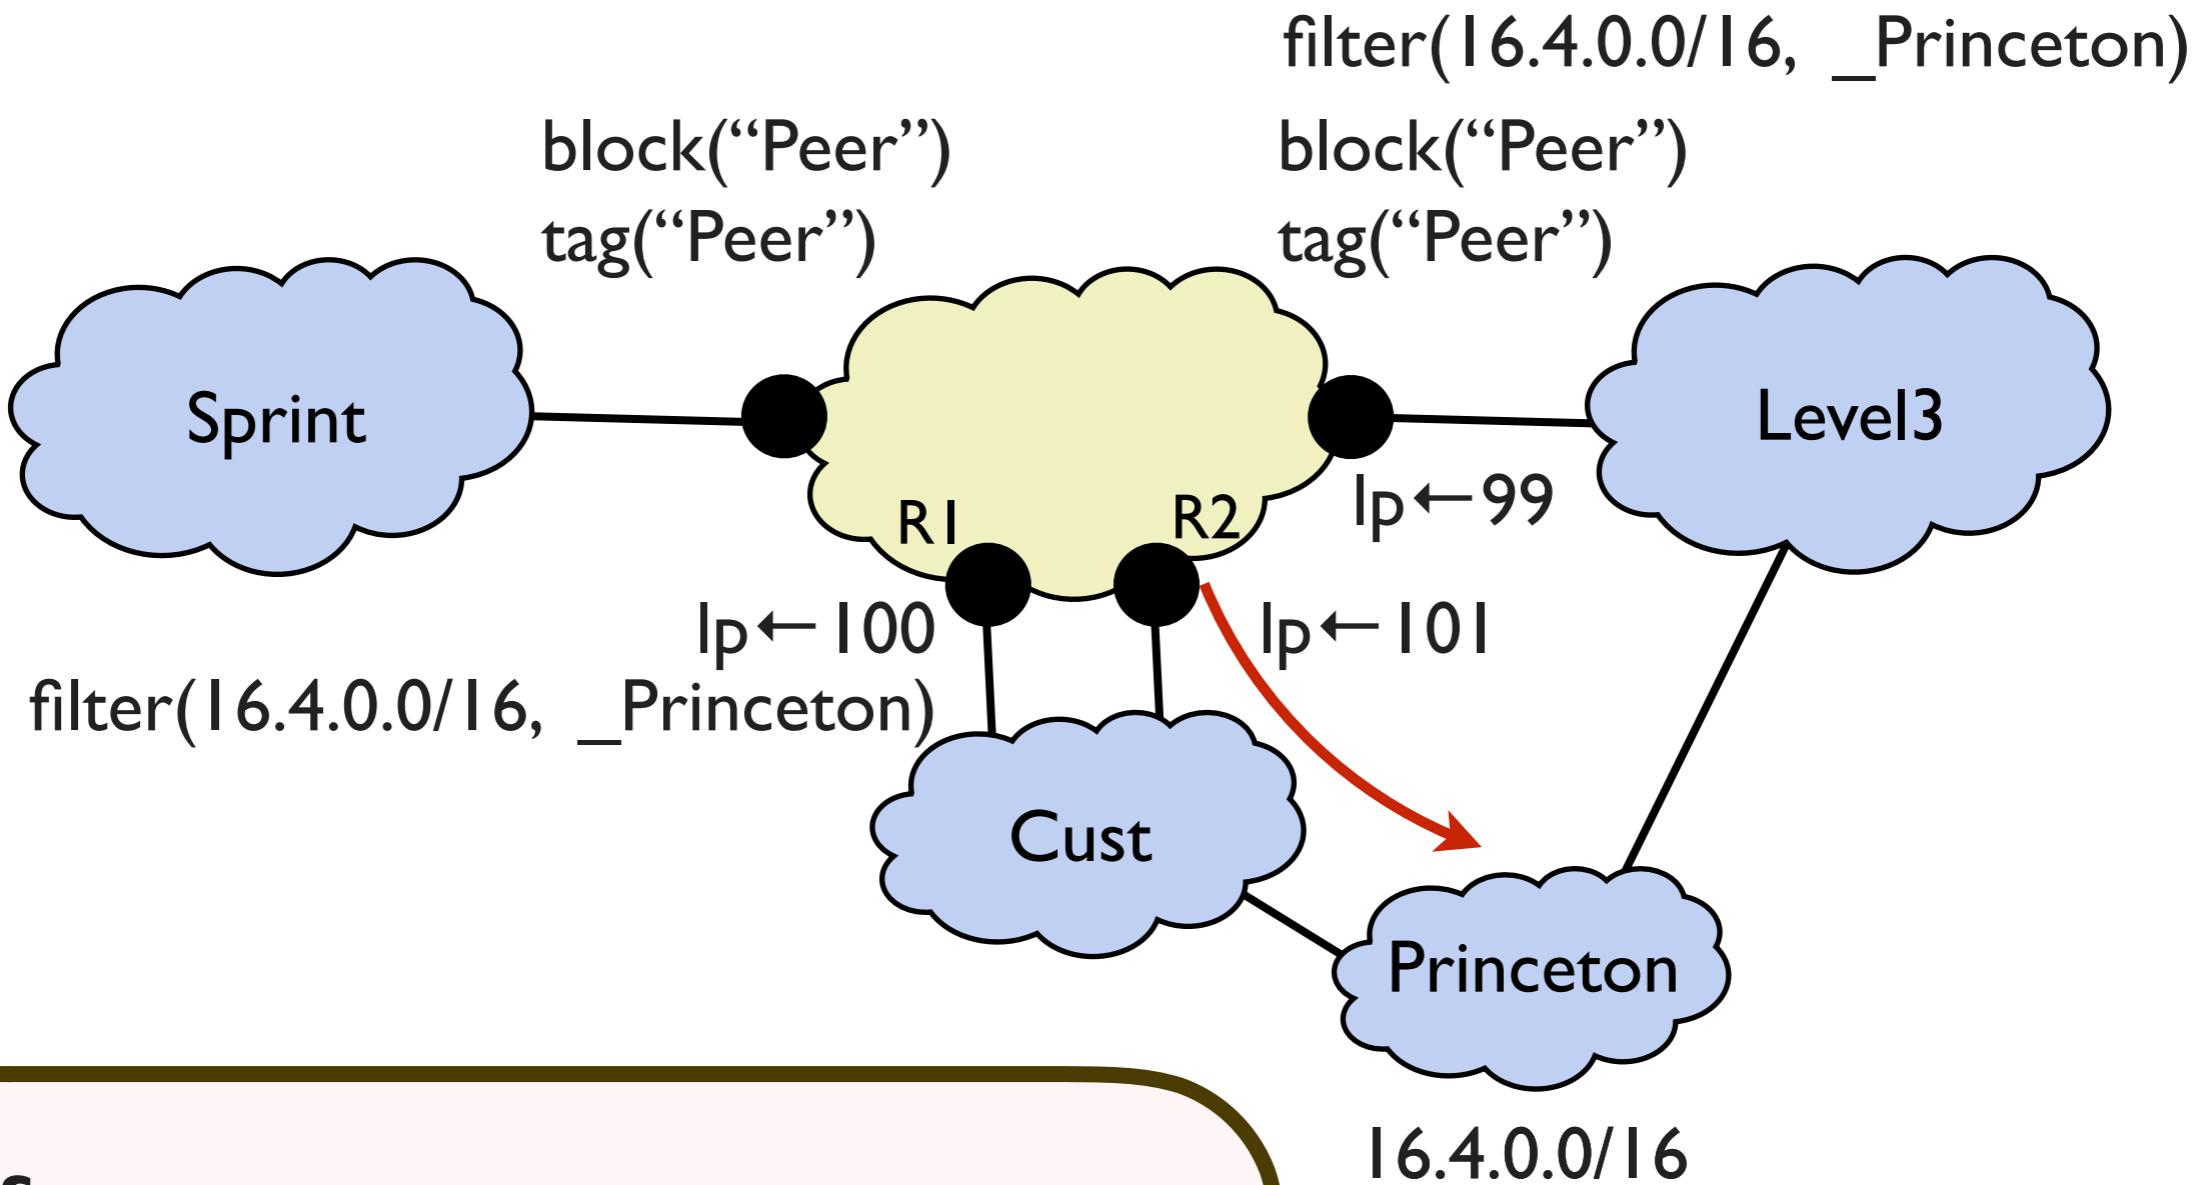
Example 2: Backbone Network



Goals

- Prevent transit between peers (\$\$\$\$)
- Prefer R2 > R1 > Peer (\$\$\$\$)

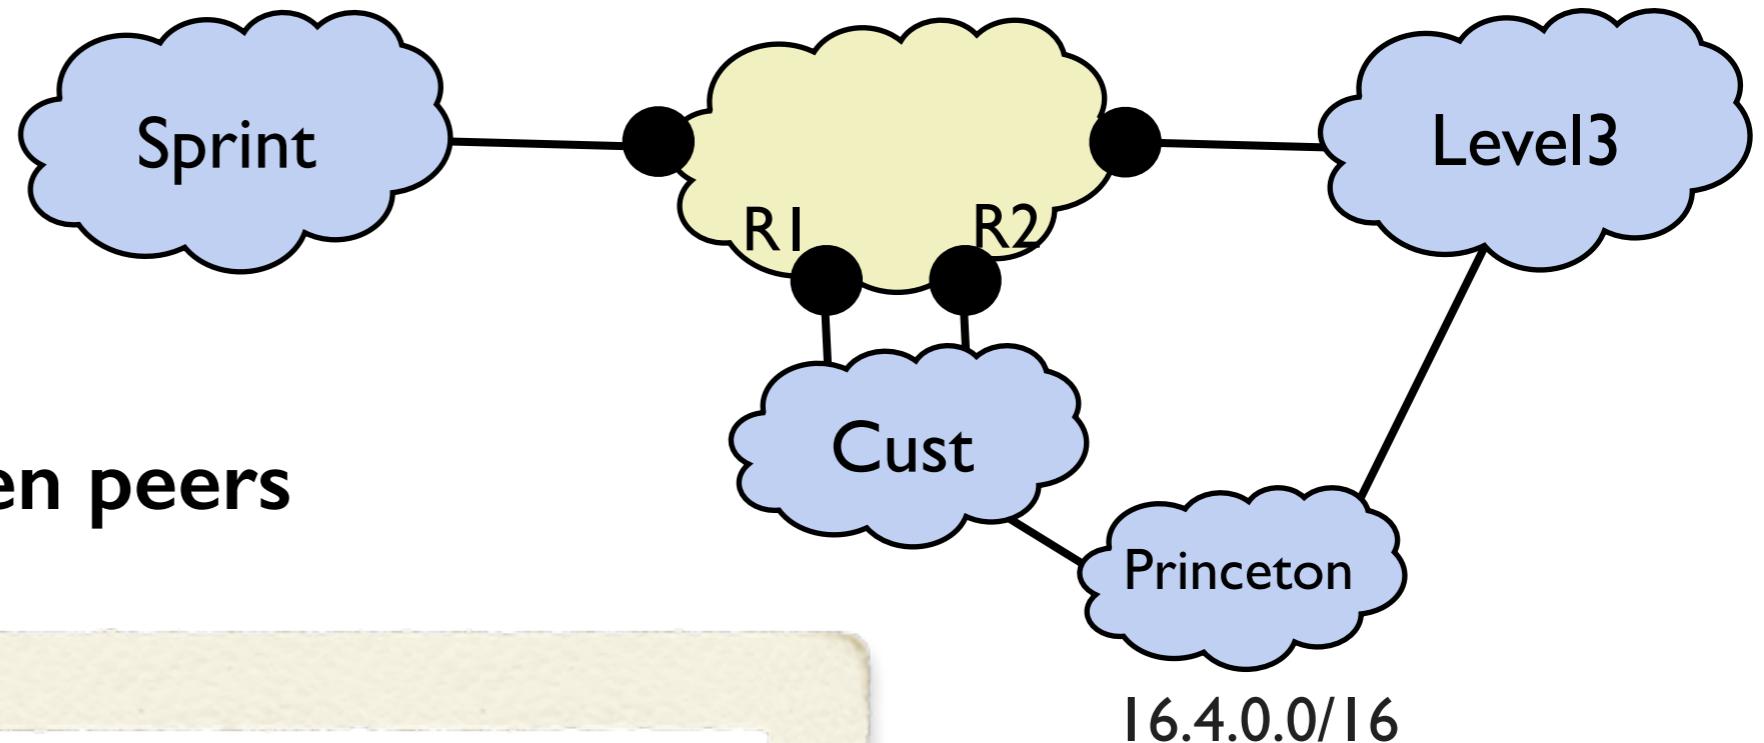
Example 2: Backbone Network



Goals

- Prevent transit between peers (\$\$\$)
- Prefer R2 > RI > Peer (\$\$\$)
- Filter customer by prefix (security)

Example 2: Backbone Network



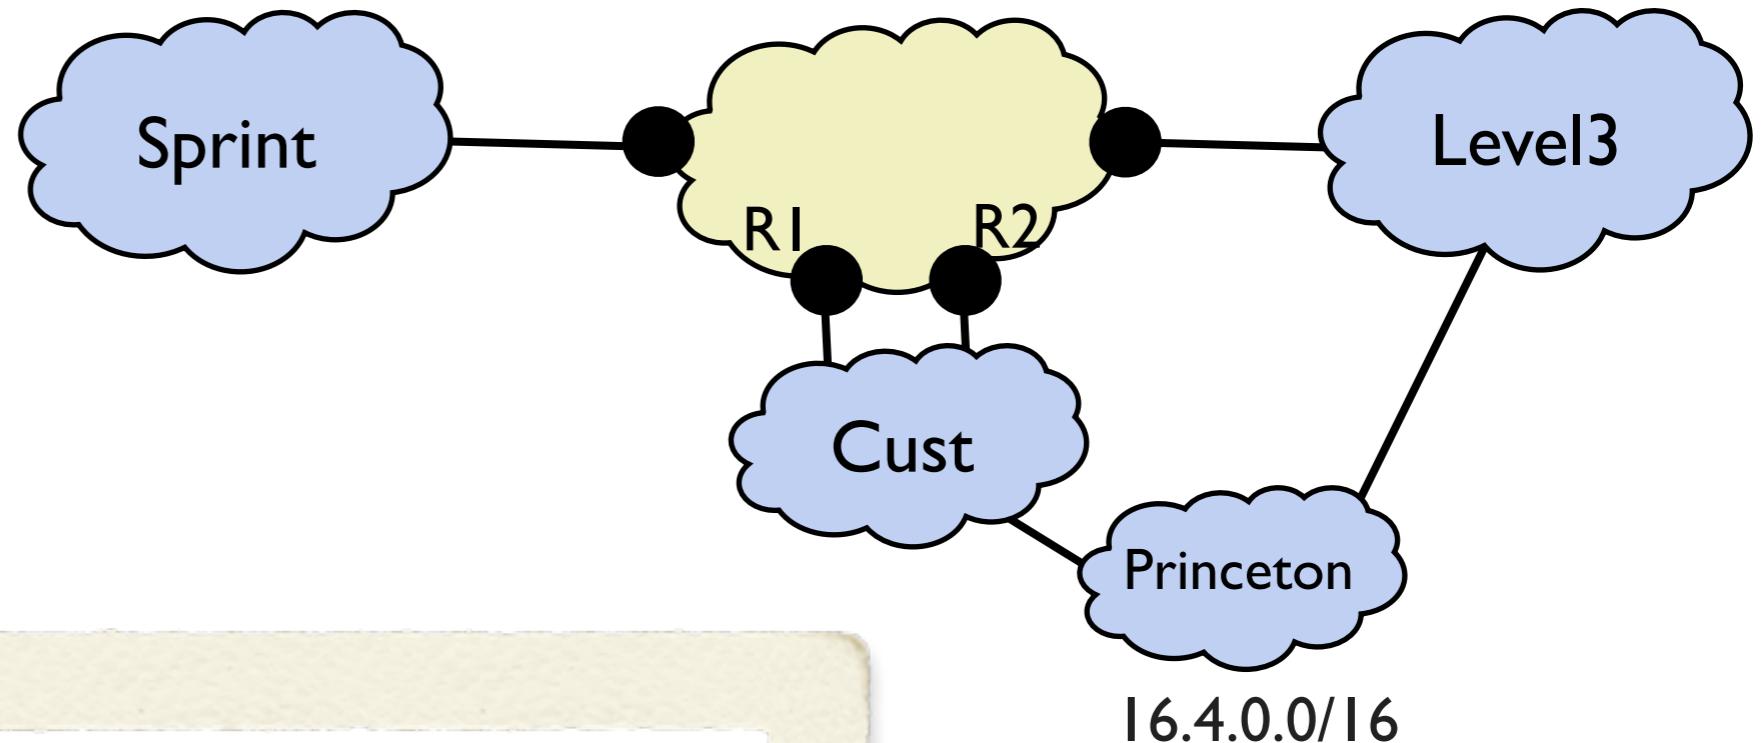
Prevent transit between peers

```
define Peer = {Sprint, Level3}
```

```
define NoTransit = {
    true => !transit(Peer,Peer)
}
```

Example 2: Backbone Network

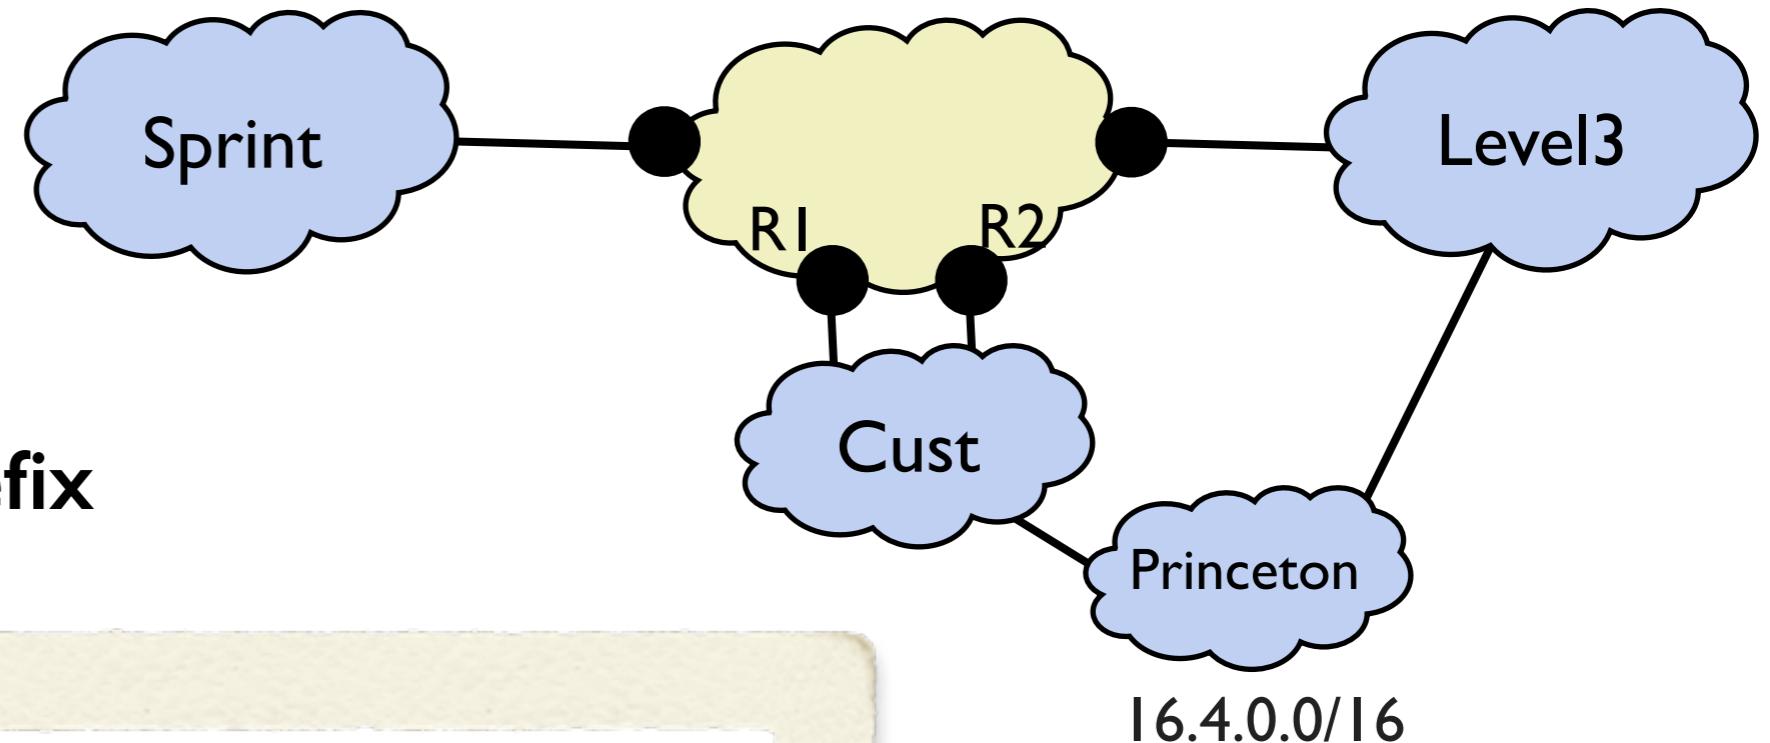
Prefer R2 > RI > Peer



...

```
define Preferences = {  
    true => exit(R2 >> RI >> Peer)  
}
```

Example 2: Backbone Network



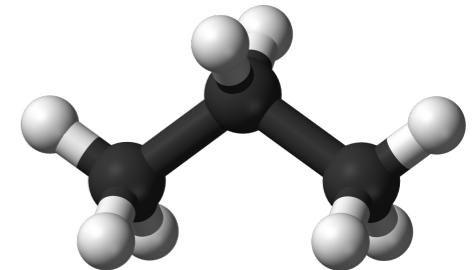
Filter customer by prefix

...

```
define Ownership = {  
    16.4.0.0/16 => end(Princeton),  
}
```

```
define Main =  
    Preferences & Ownership & NoTransit
```

Propane: Summary



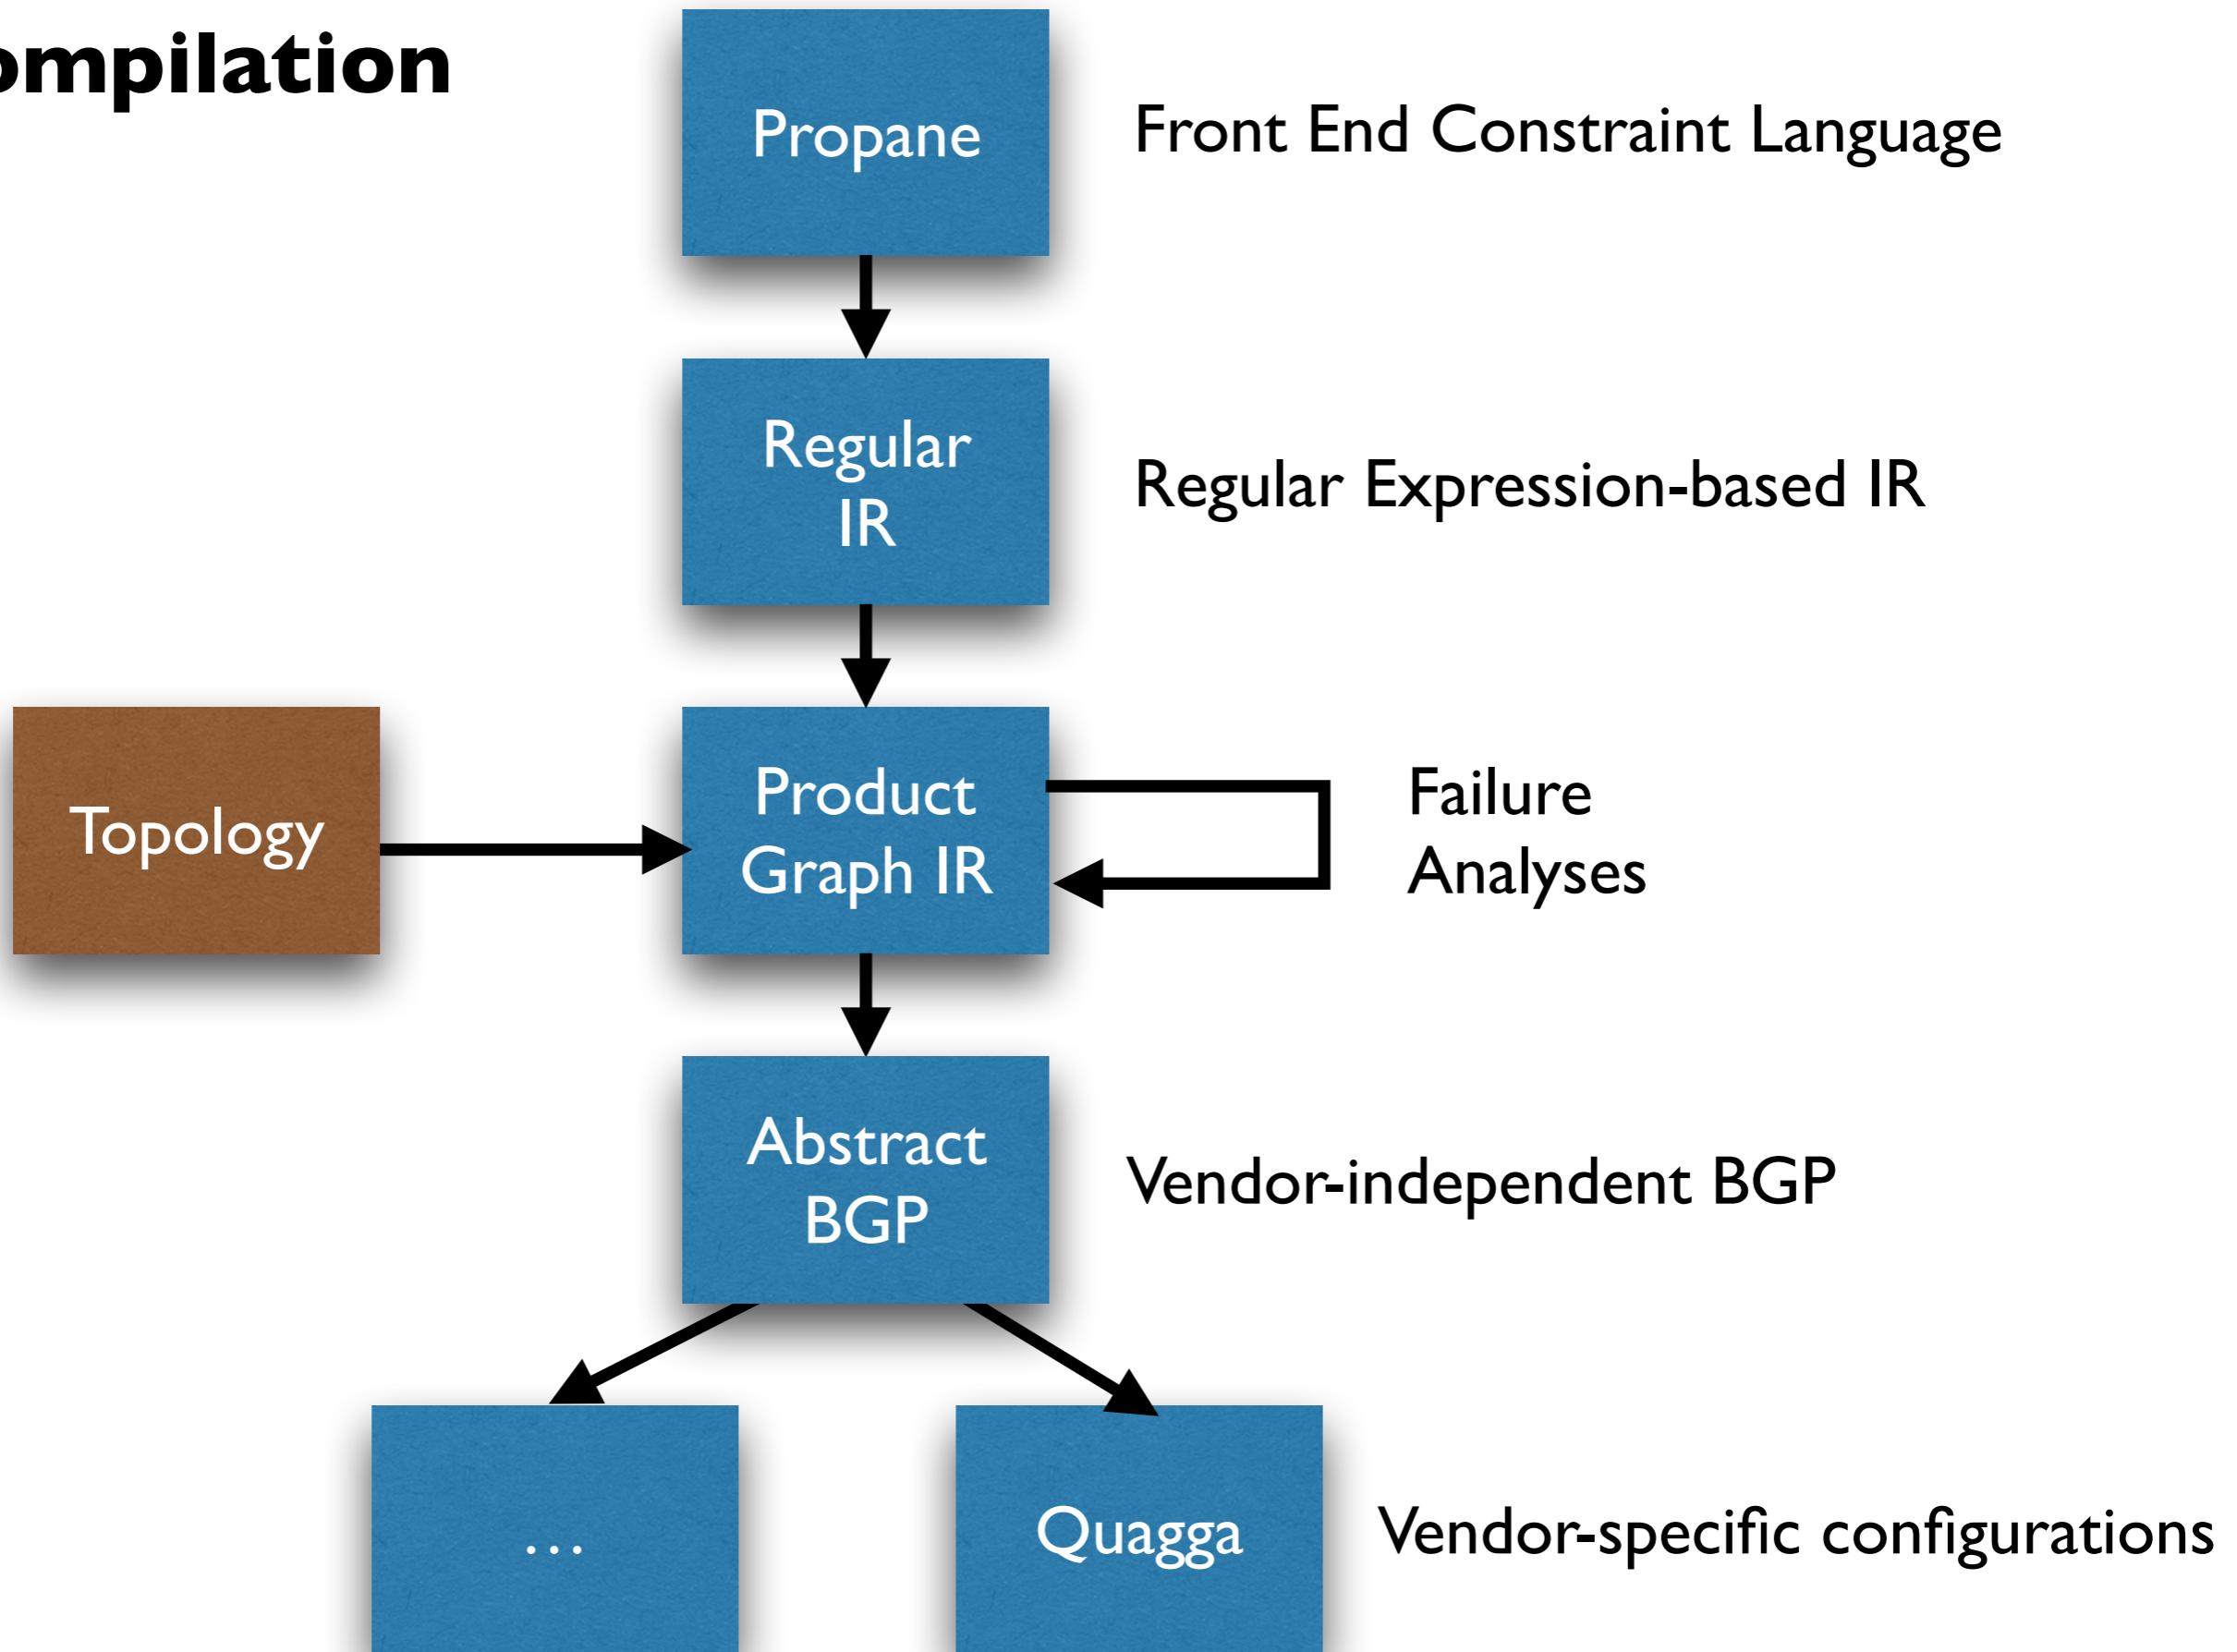
High-level language

- Single, network-wide policy
- Program using constraints and preferences
- Create new, reusable abstractions (e.g., transit traffic)
- Configure the network modularly

Compiler

- Automatically generates filters, preferences, community values, etc
- Static analysis guarantees policy compliance under all failures
- Lower bound on failures for aggregation-induced black holes

Compilation



Propane Regular IR

Propane

Step I: Combine modular constraints

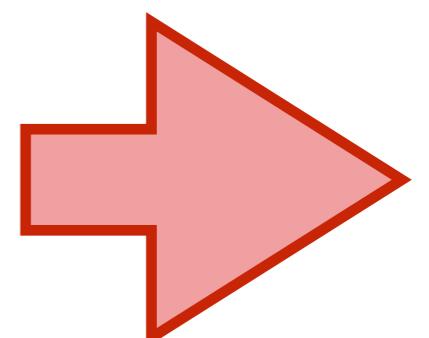
```
define Ownership =
{ PG1 => end(A)
  PG2 => end(B)
  PL1 => end(E)
  PL2 => end(F)
  true => exit(Peer1 >> Peer2) }
```

```
define NoTransit =
{ true => !transit(Peer, Peer) }
```

```
define Locality =
{ PL1 | PL2 => always(in) }
```

```
define Main =
  Ownership & Locality & NoTransit
  & agg(PG, in -> out)
```

Regular
IR



Propane Regular IR

Propane



Regular
IR

Prefix-by-prefix intersection of constraints

```
PG1 => !transit(Peer, Peer) & end(A)
PG2 => !transit(Peer, Peer) & end(B)
PL1 => !transit(Peer, Peer) & always(in) & end(E)
PL2 => !transit(Peer, Peer) & always(in) & end(F)
true => !transit(Peer, Peer) & exit(Peer1 >> Peer2)
```

Propane Regular IR

Propane

Step 2: Expand constraints in to regular expressions

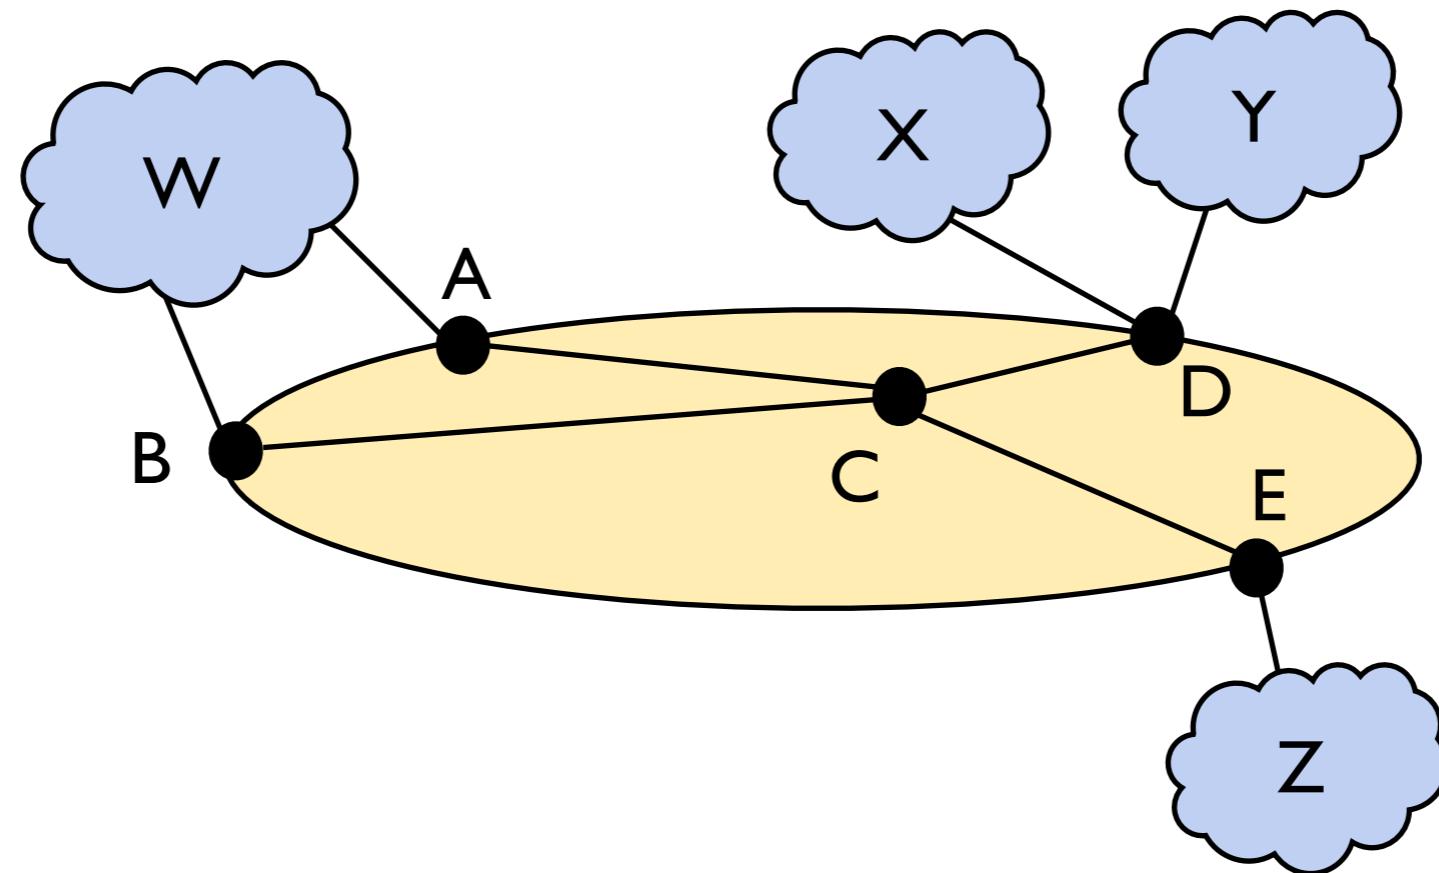
Regular
IR

```
any = out*.in+.out*
end(X) = (\Sigma*.X)
always(X) = (X)*
exit(X) = (out*.in*(X ∩ in).out+) |
           (out*.in+(X ∩ out).out*)
start(X) = (X.\Sigma*)
avoid(X) = (!X)*
waypoint(X) = (\Sigma*.X.\Sigma*)
```

Step 3: Reduced syntax

```
true => A.(X >> Y).out*
true => (A.X.out*) >> (A.Y.out*)
```

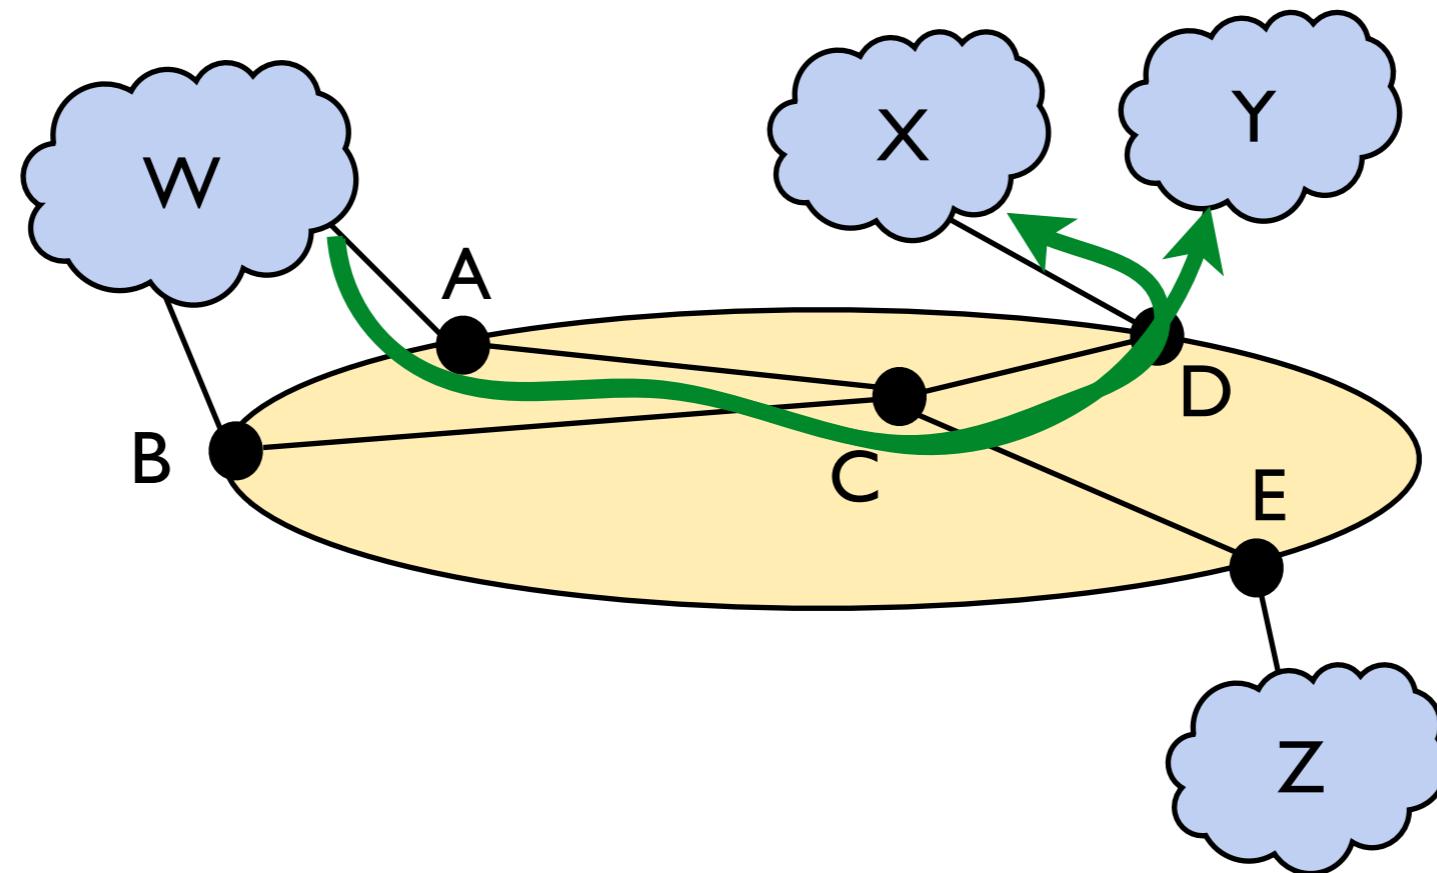
Compilation: An Idealized Example



Policy: `enter(A) & exit(D) >> enter(B) & exit(out)`

Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

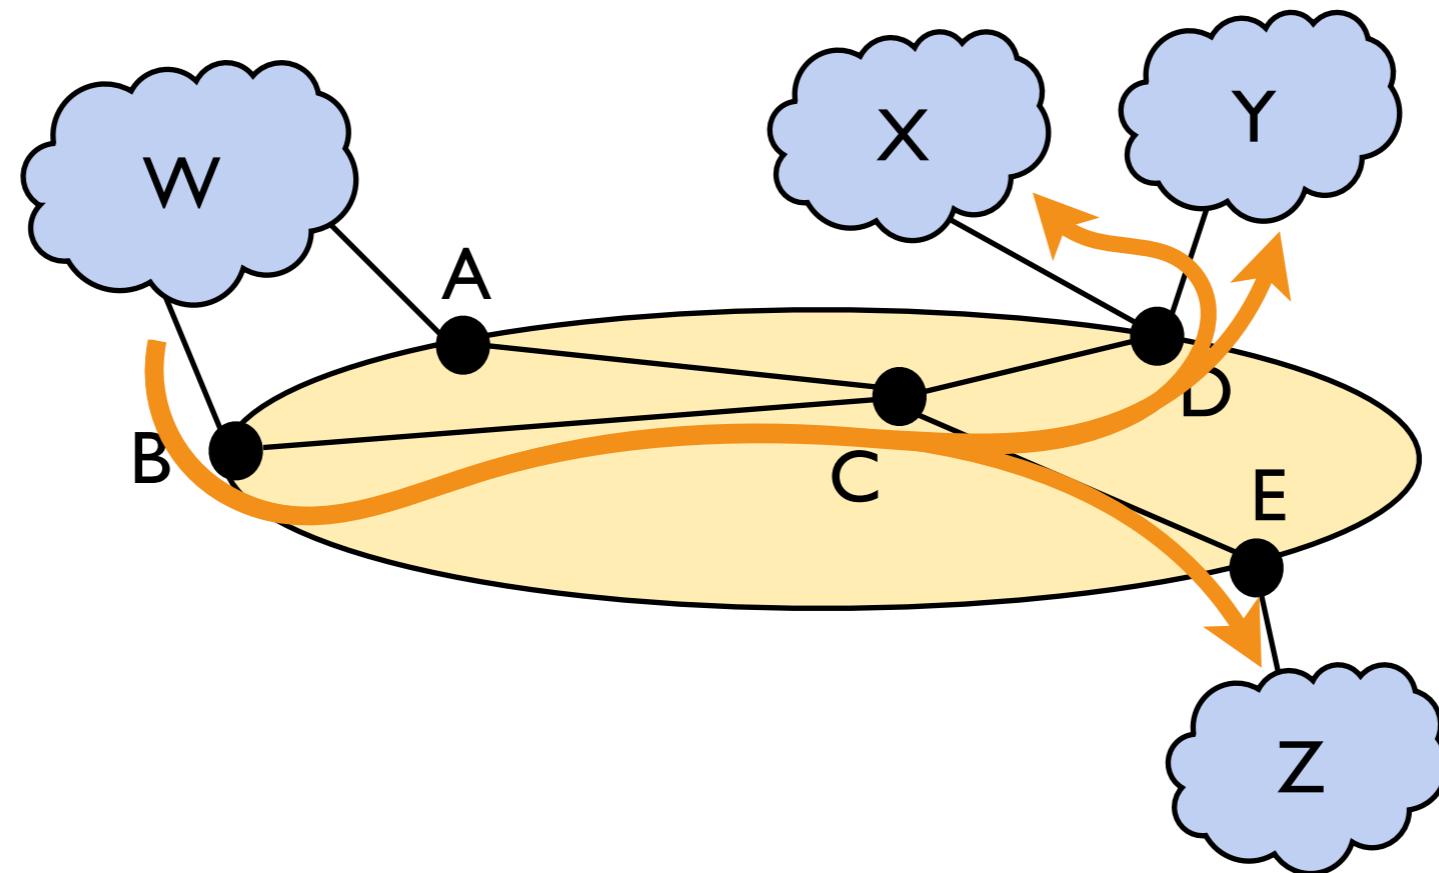
Compilation: An Idealized Example



Policy: `enter(A) & exit(D) >> enter(B) & exit(out)`

Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

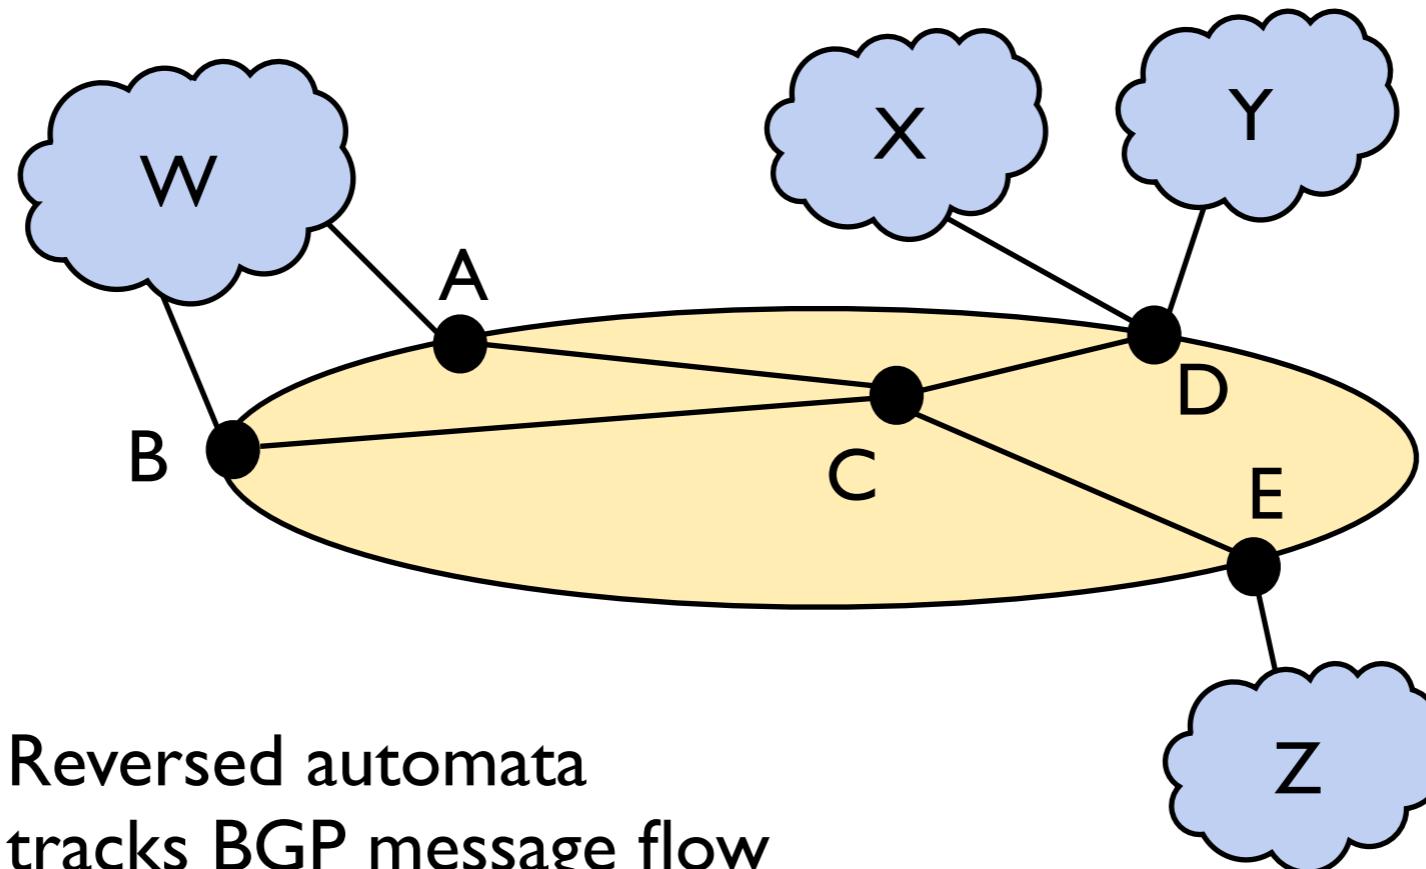
Compilation: An Idealized Example



Policy: `enter(A) & exit(D) >> enter(B) & exit(out)`

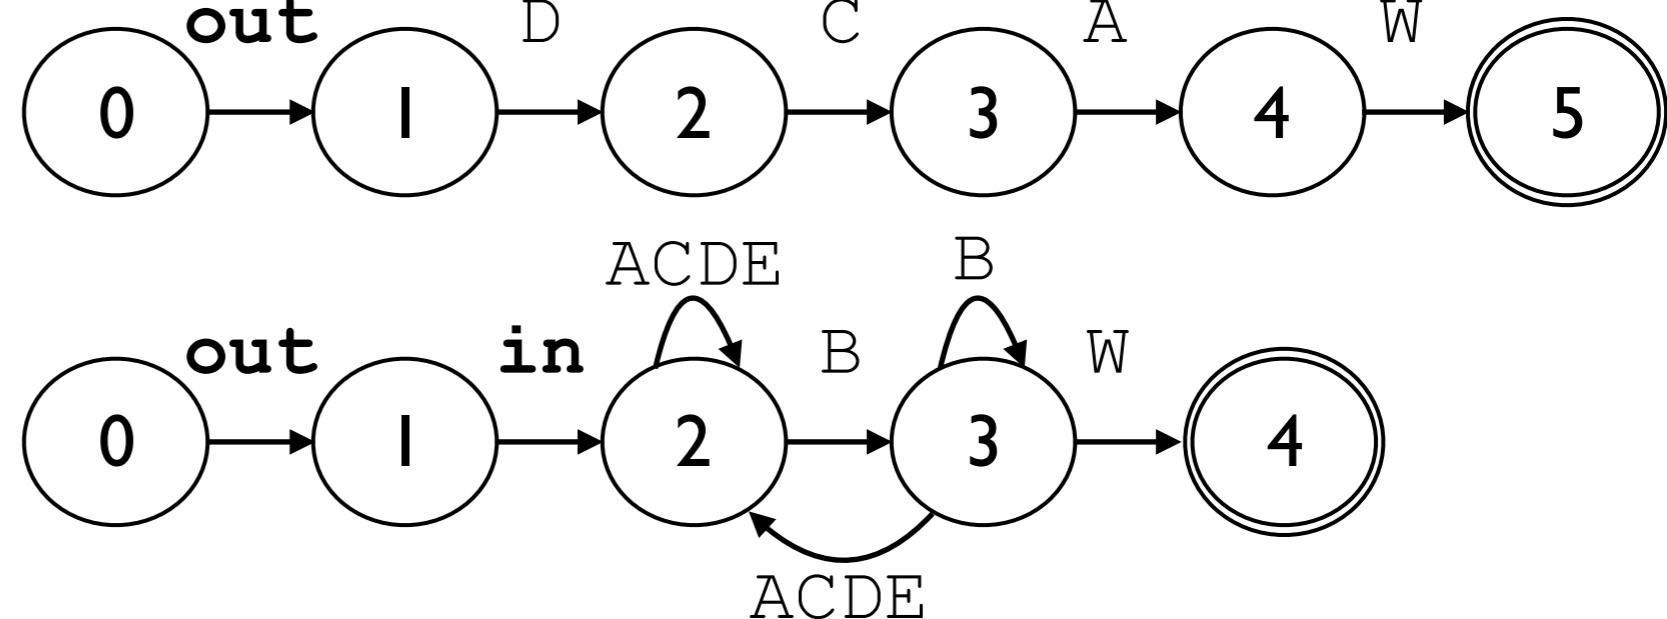
Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

Reversed Automata from Policies



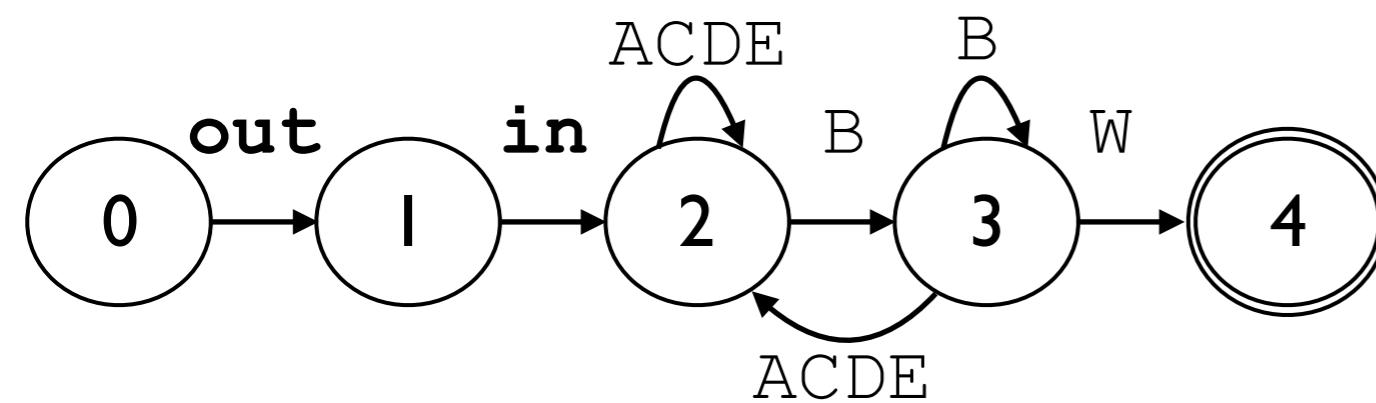
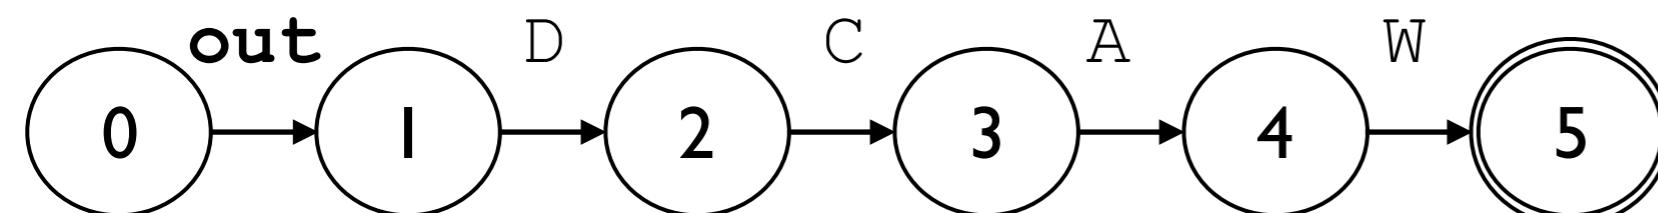
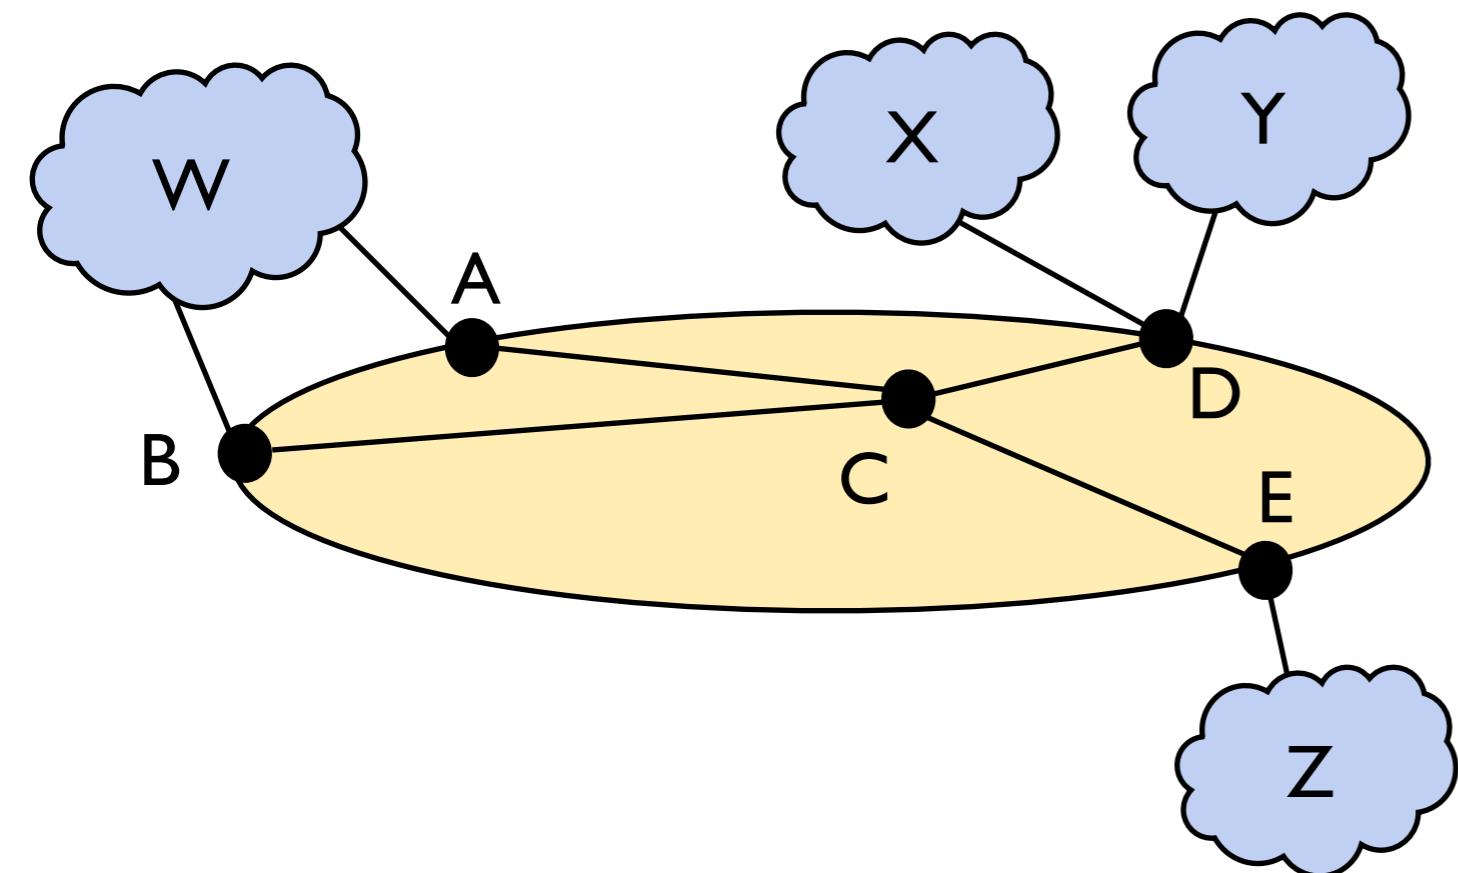
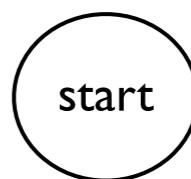
Policy:

1. (W.A.C.D.out)
2. (W.B.in+.out)



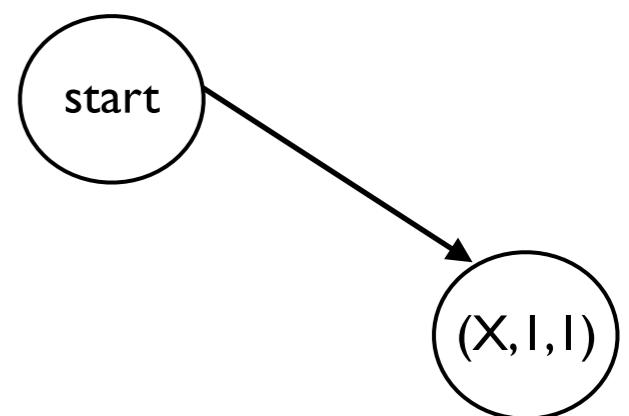
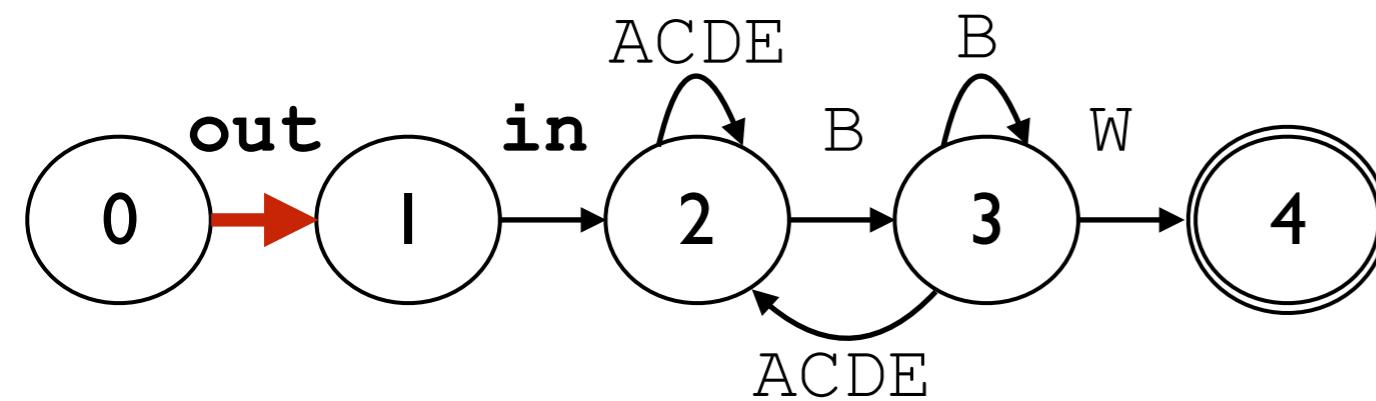
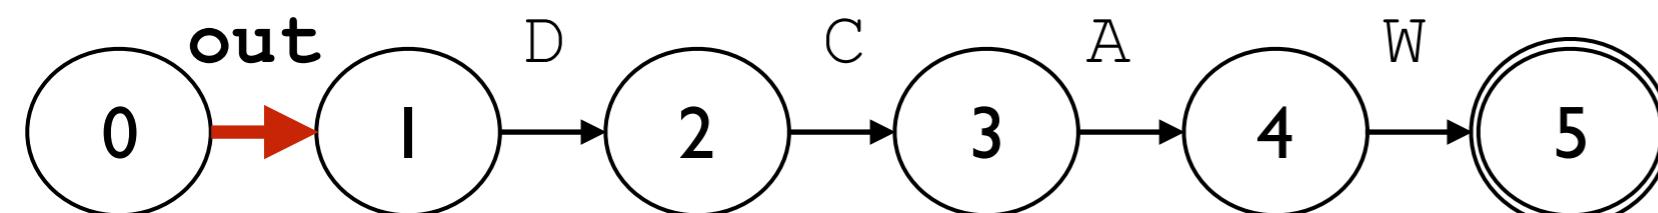
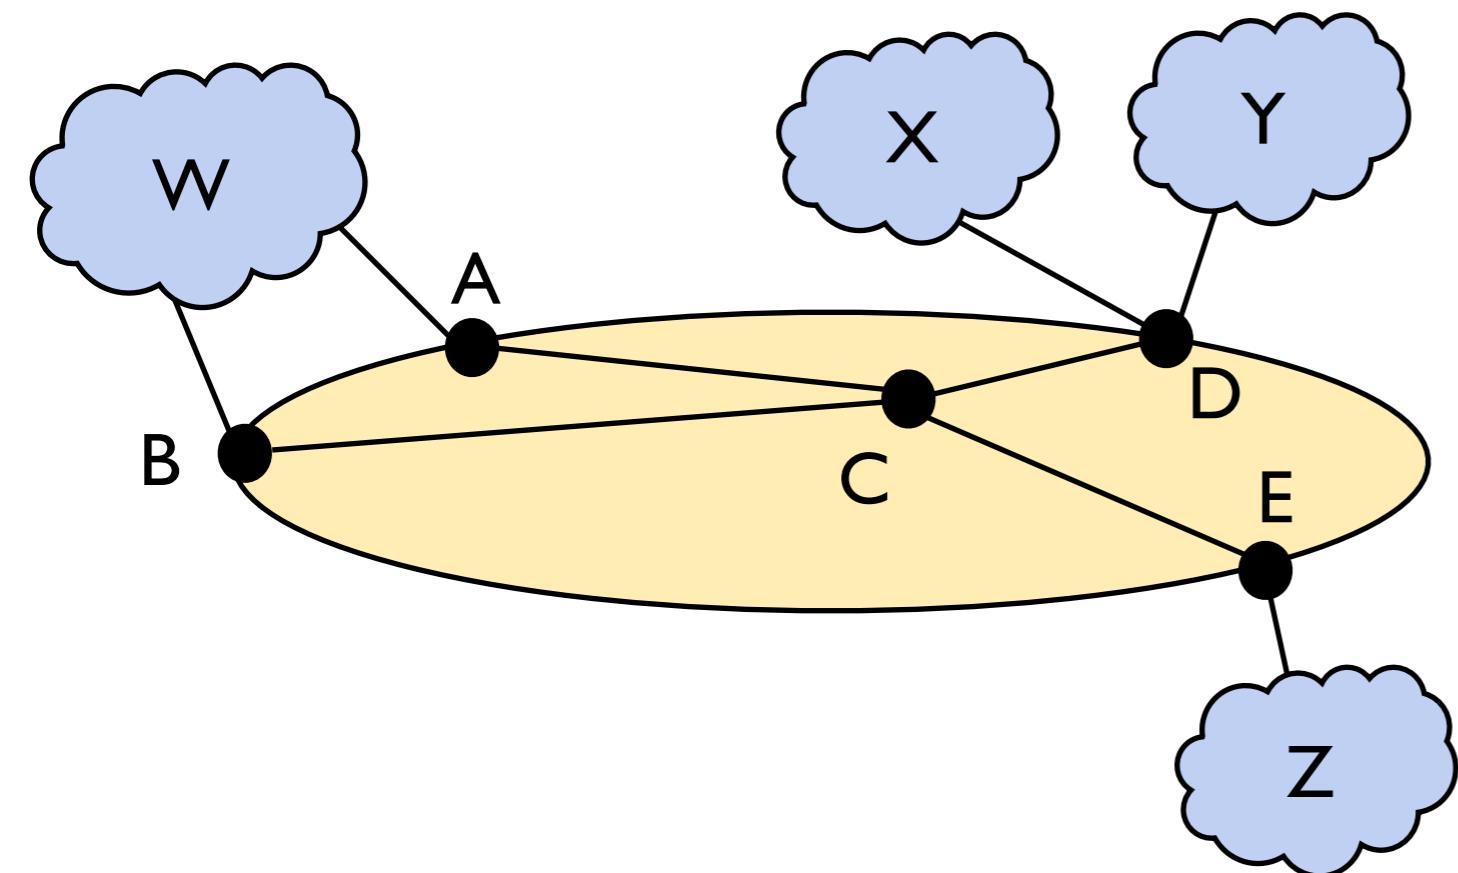
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



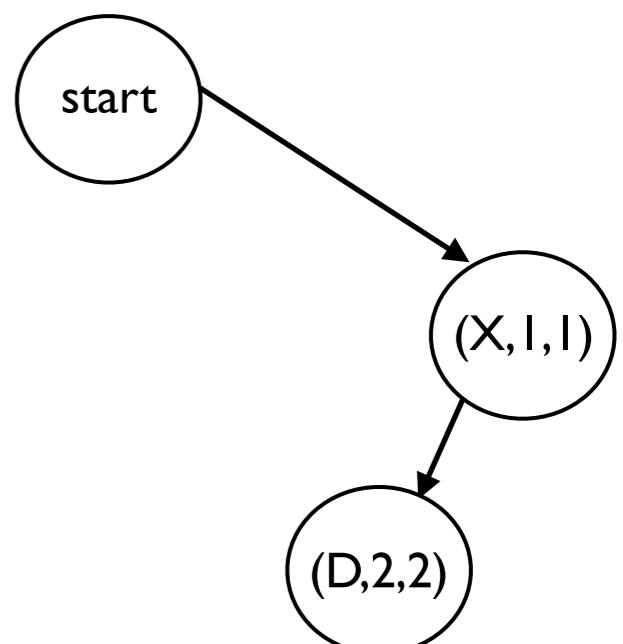
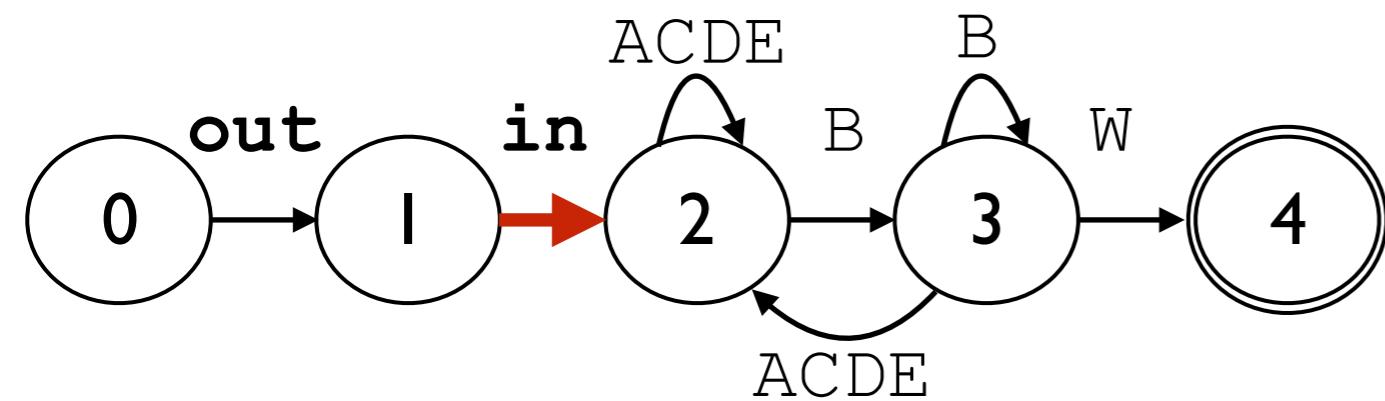
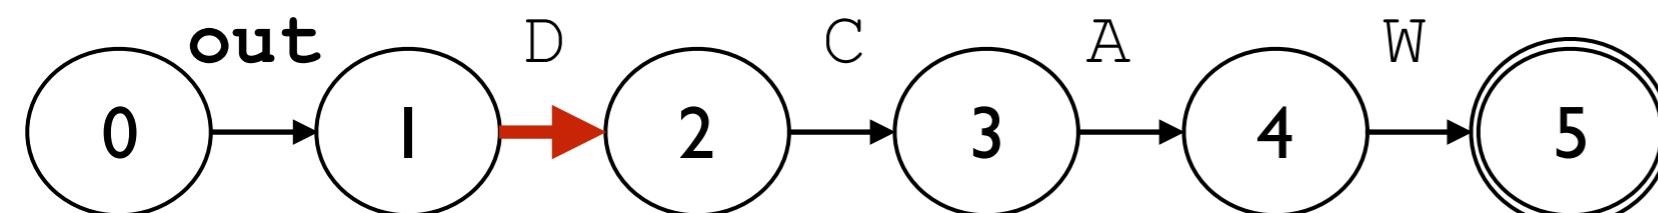
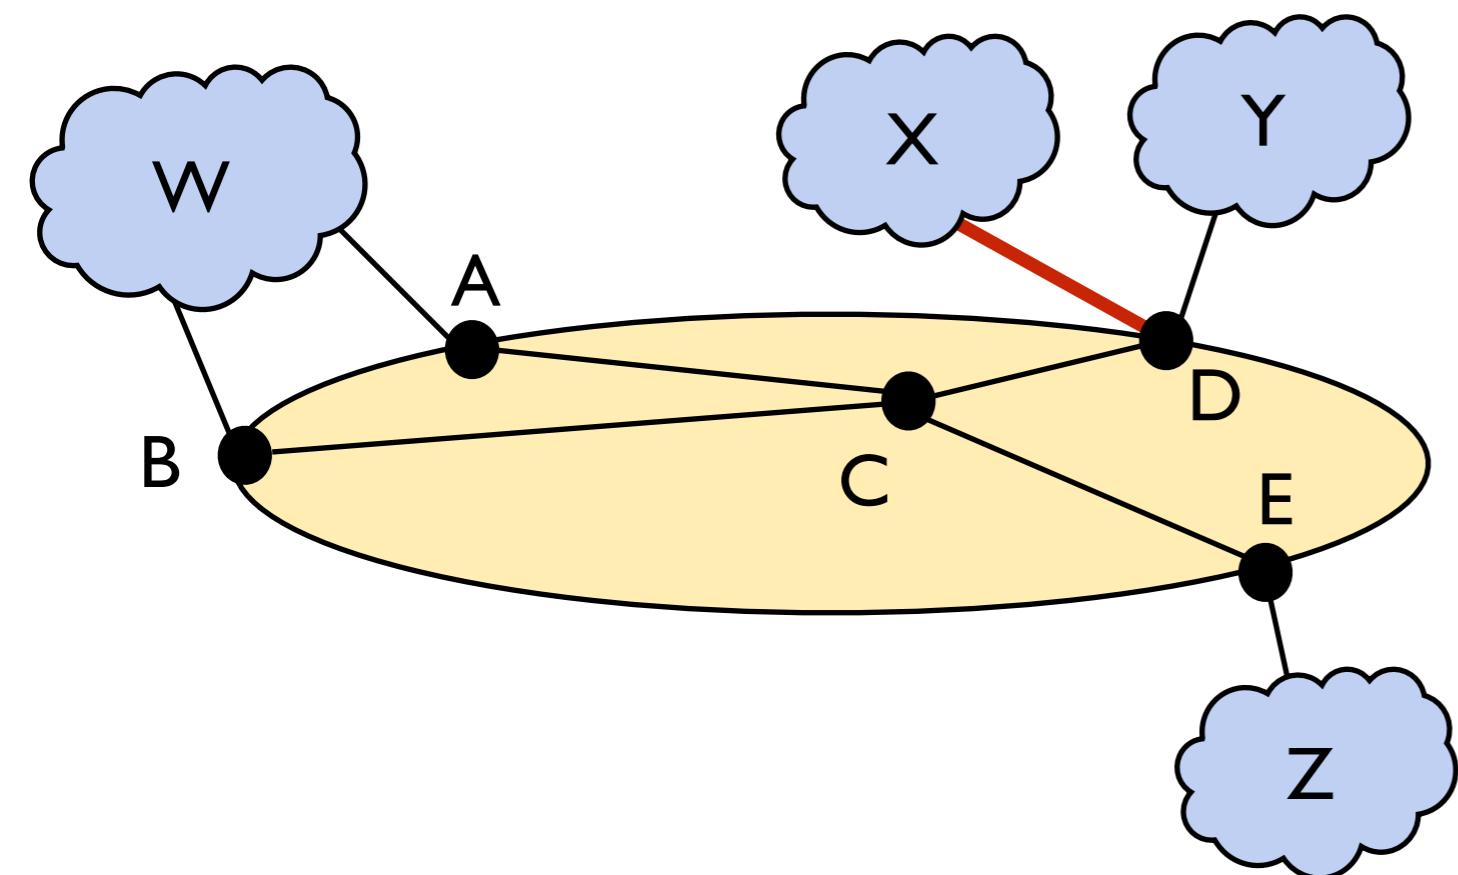
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



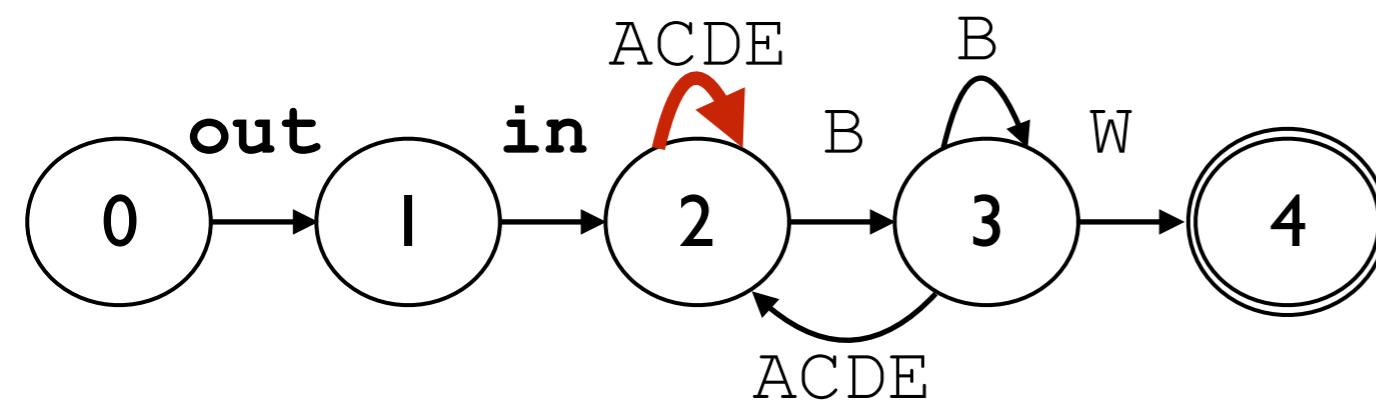
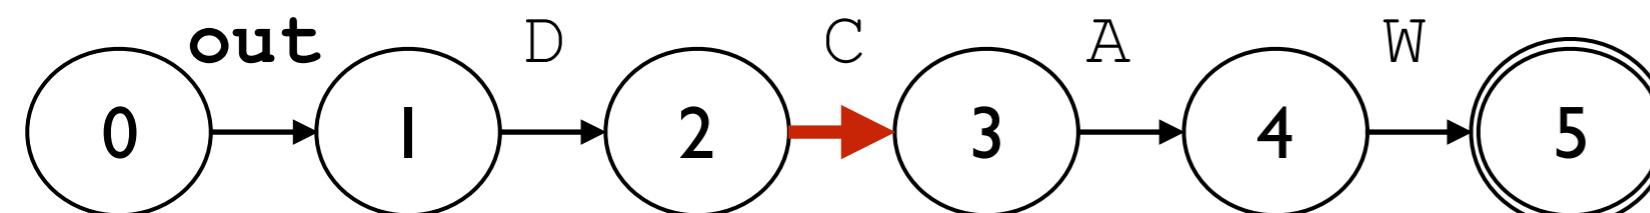
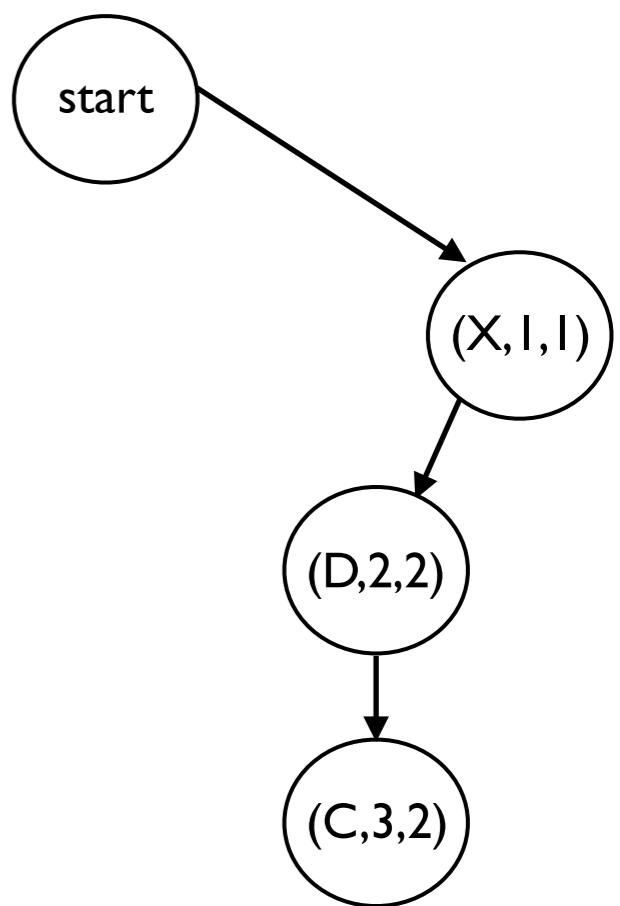
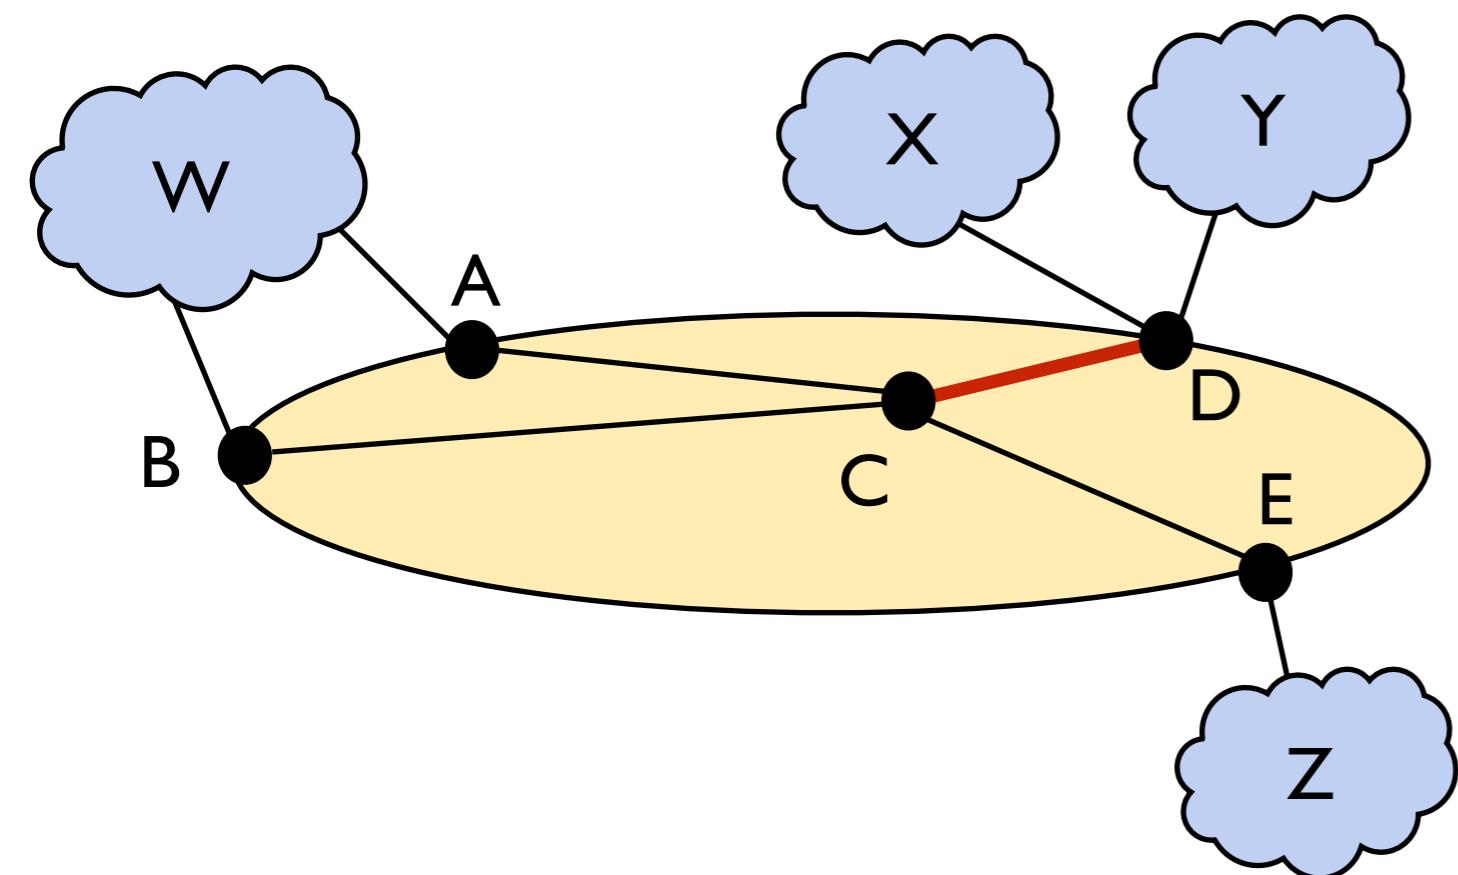
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



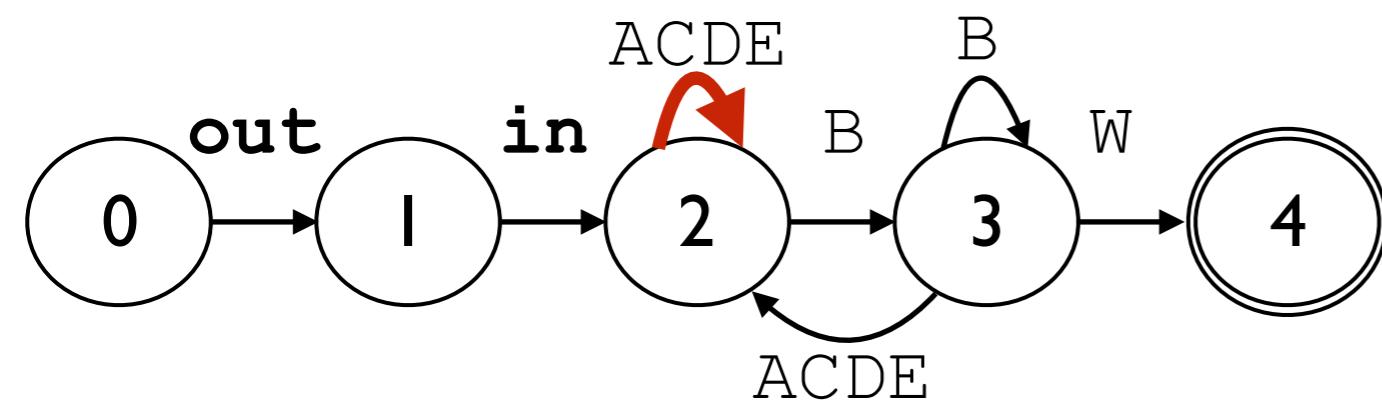
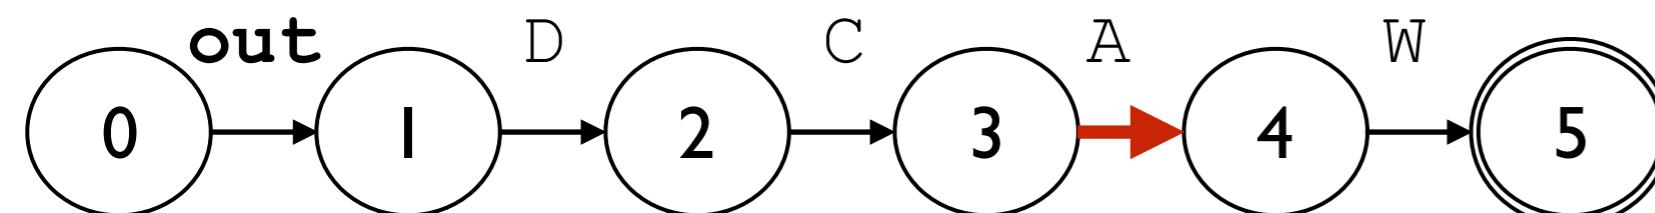
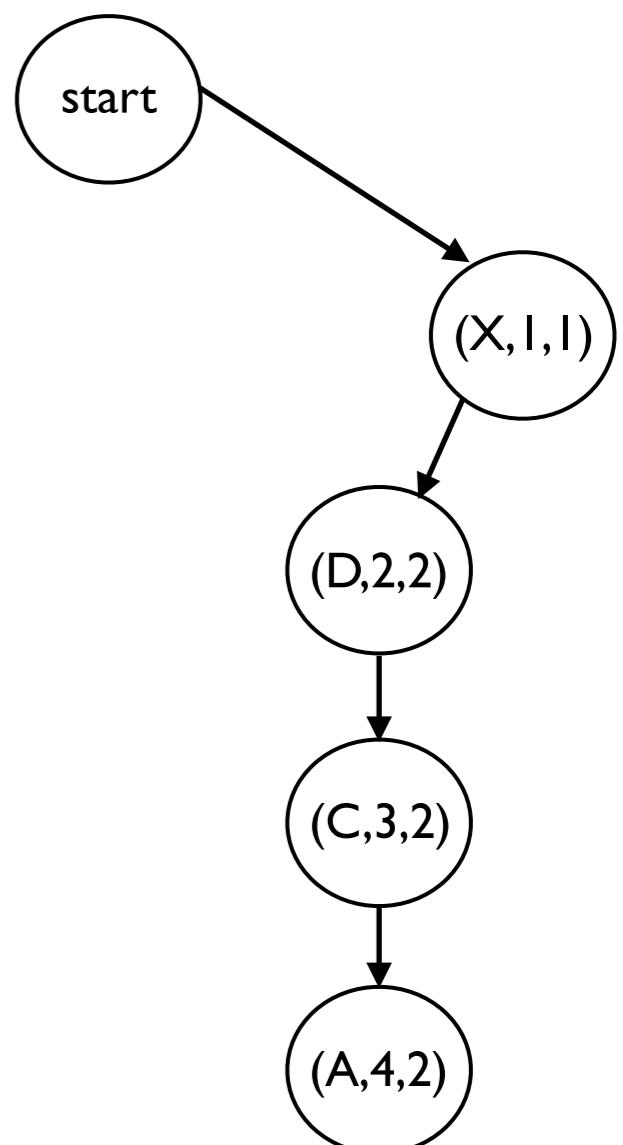
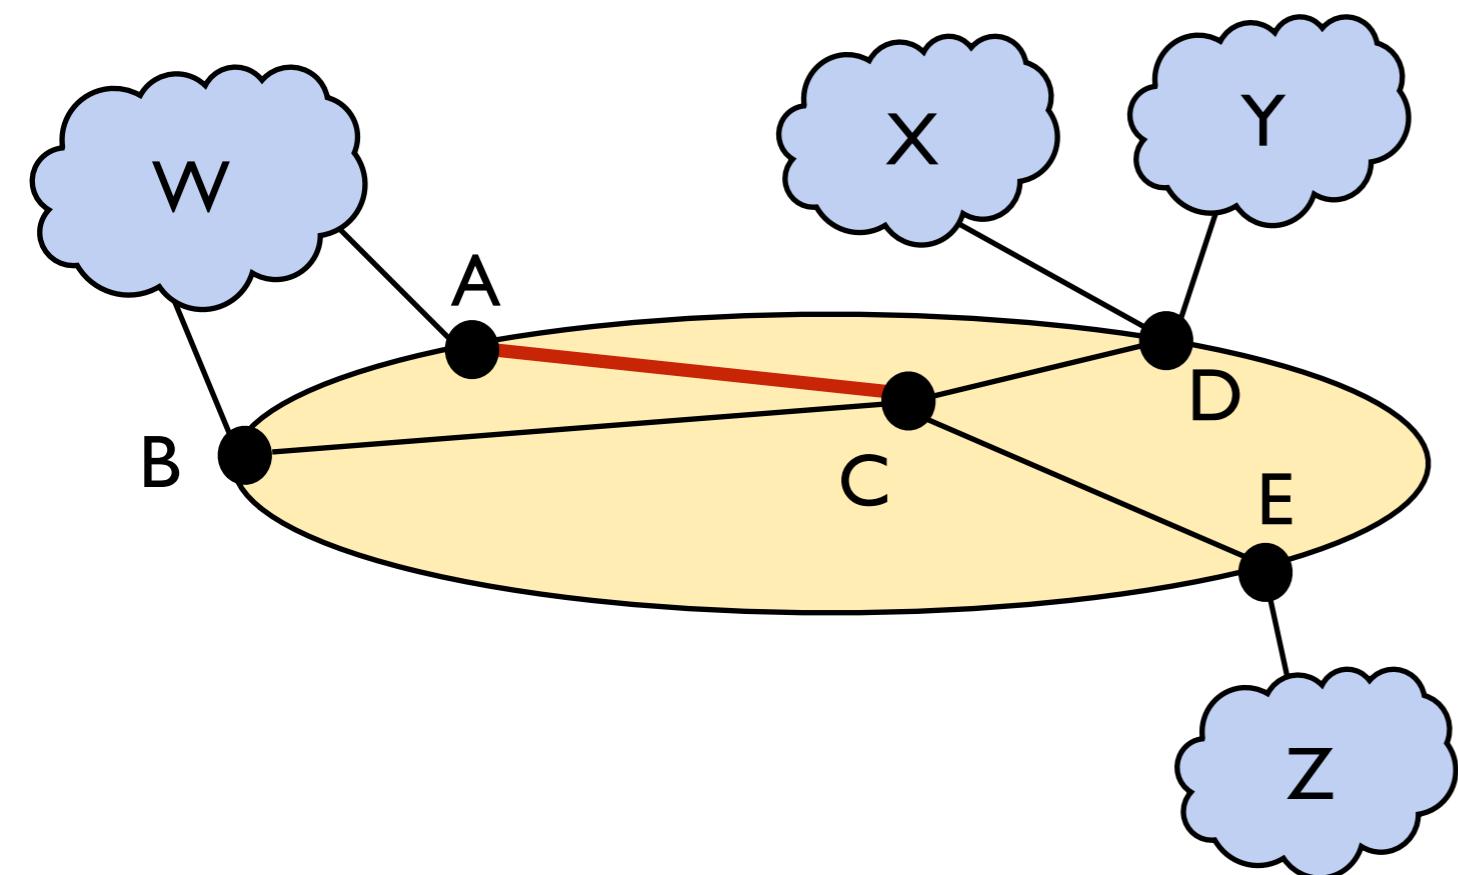
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



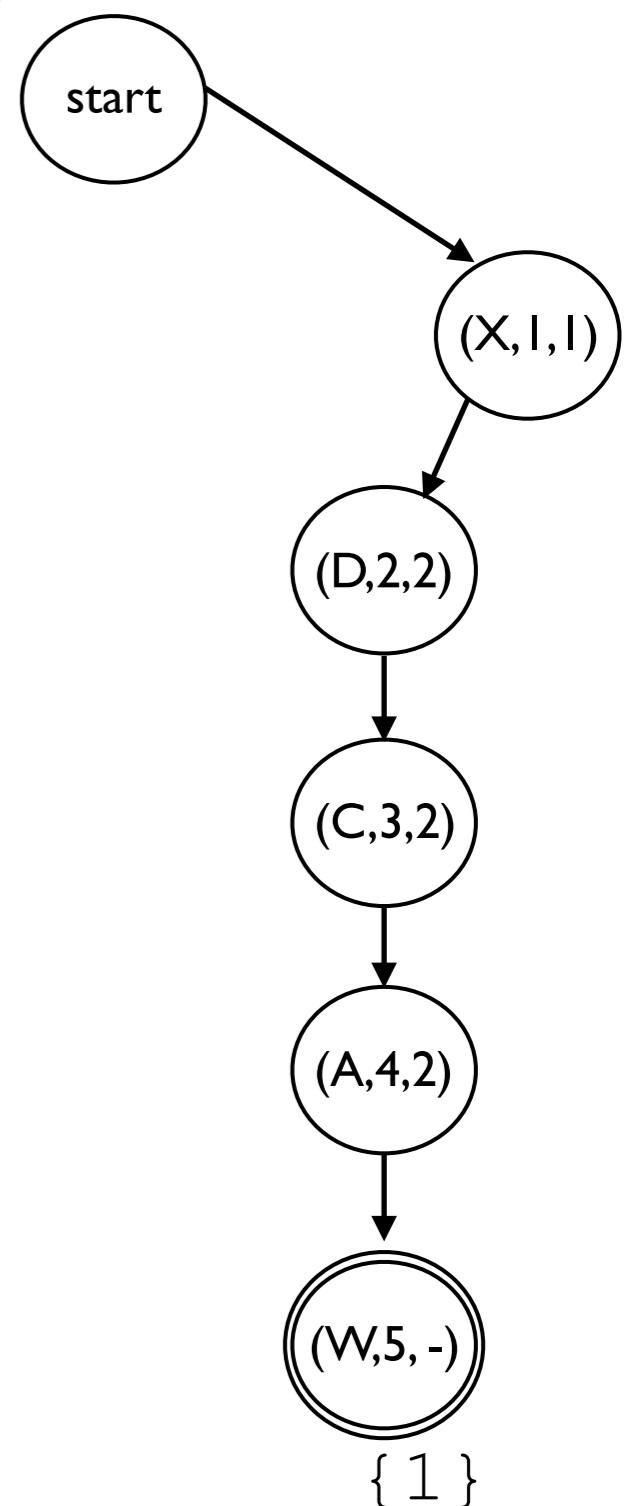
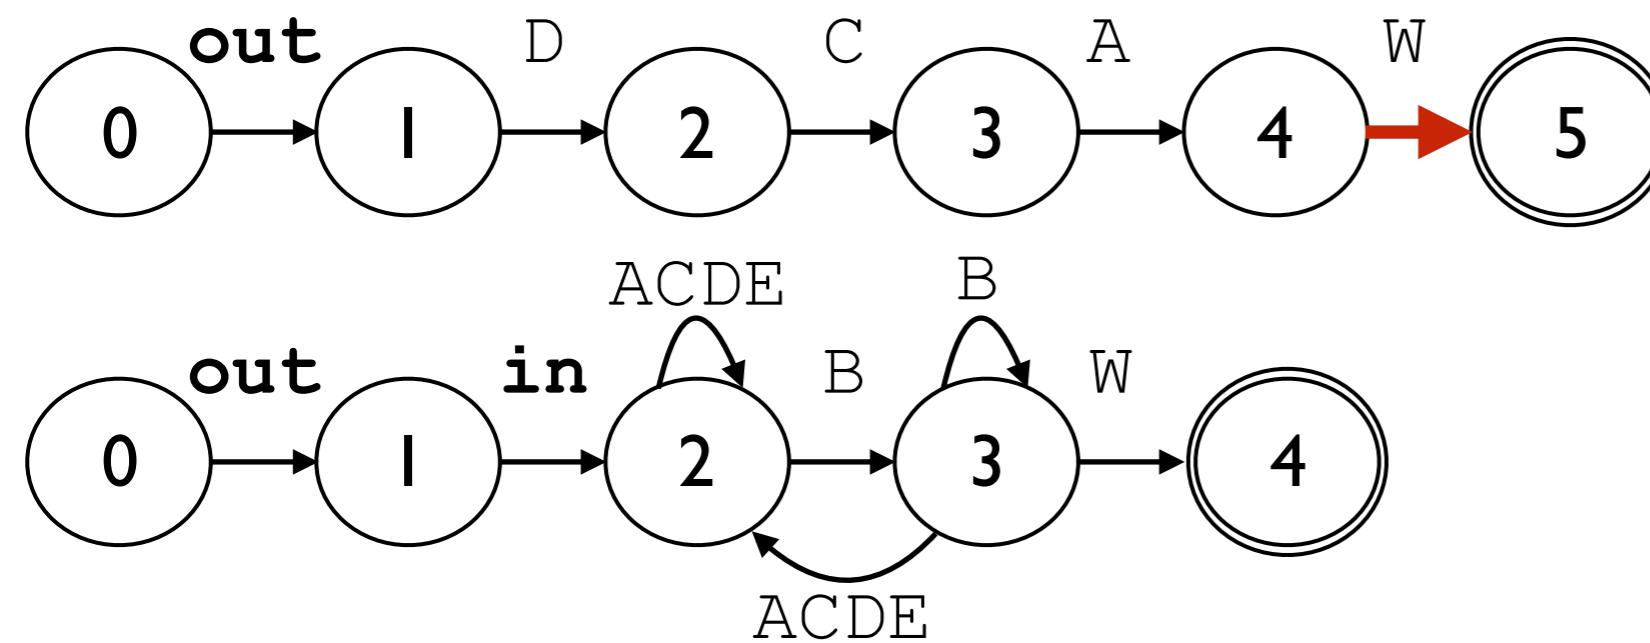
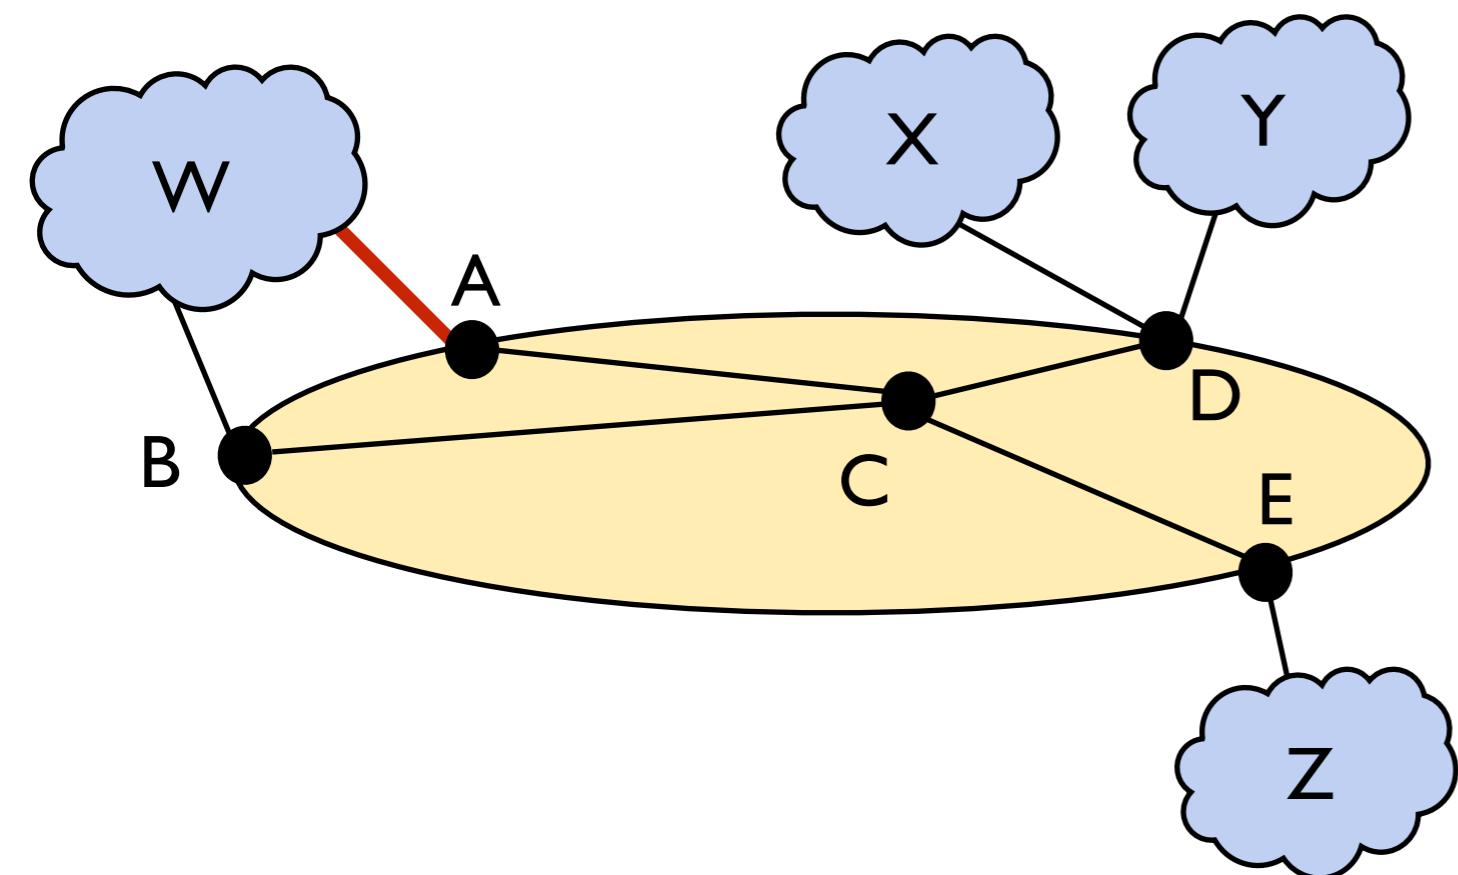
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



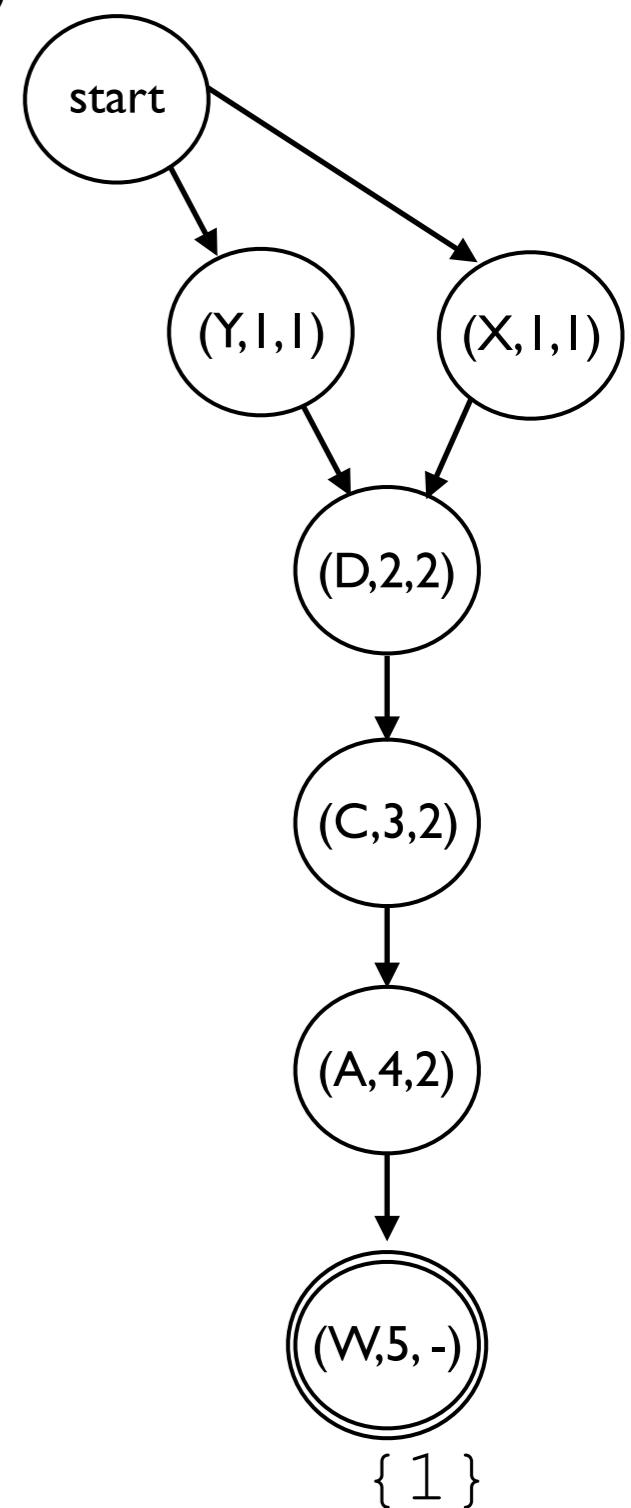
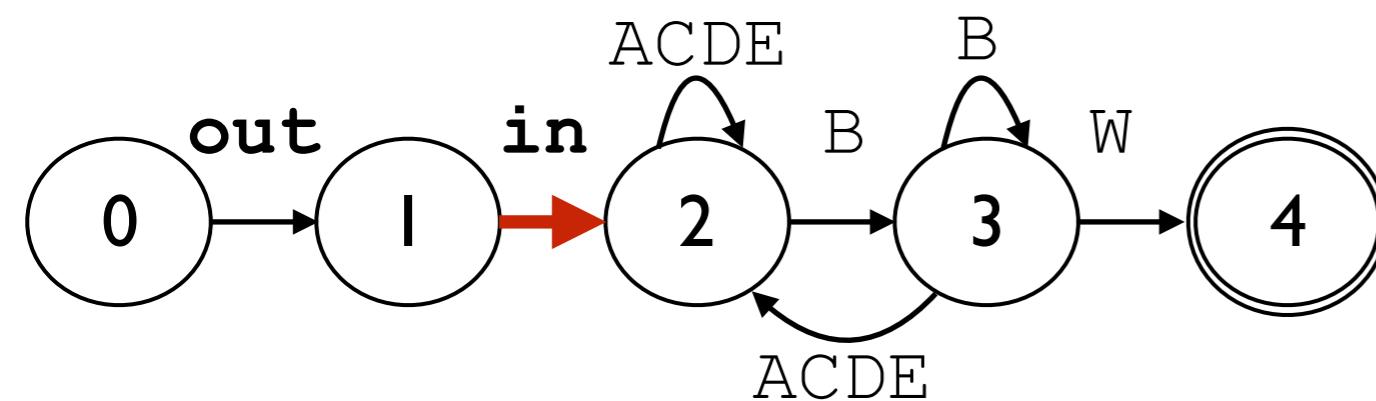
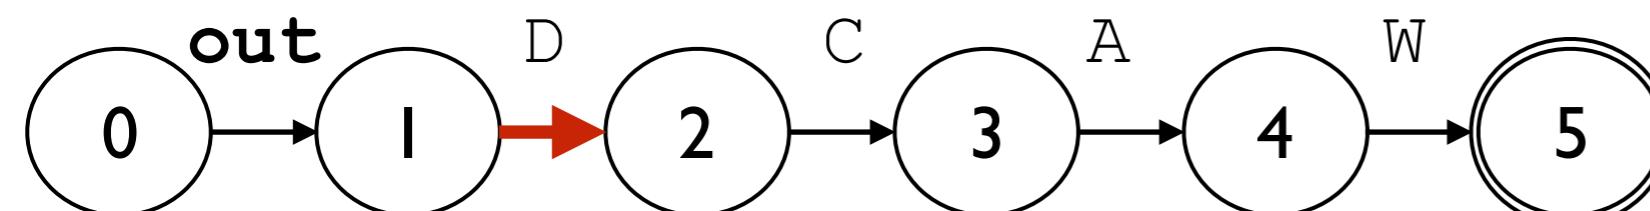
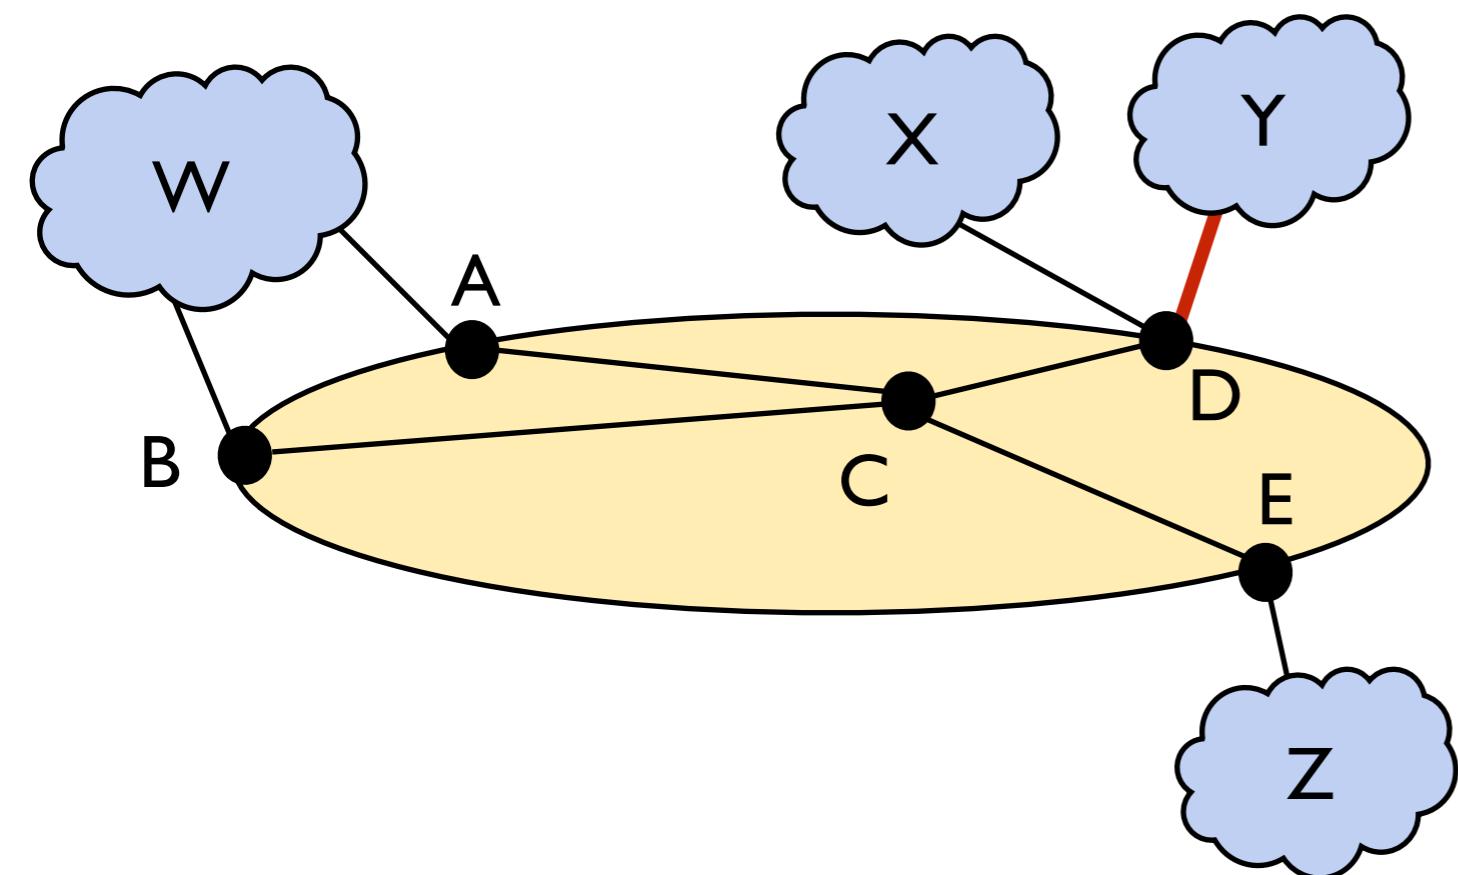
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



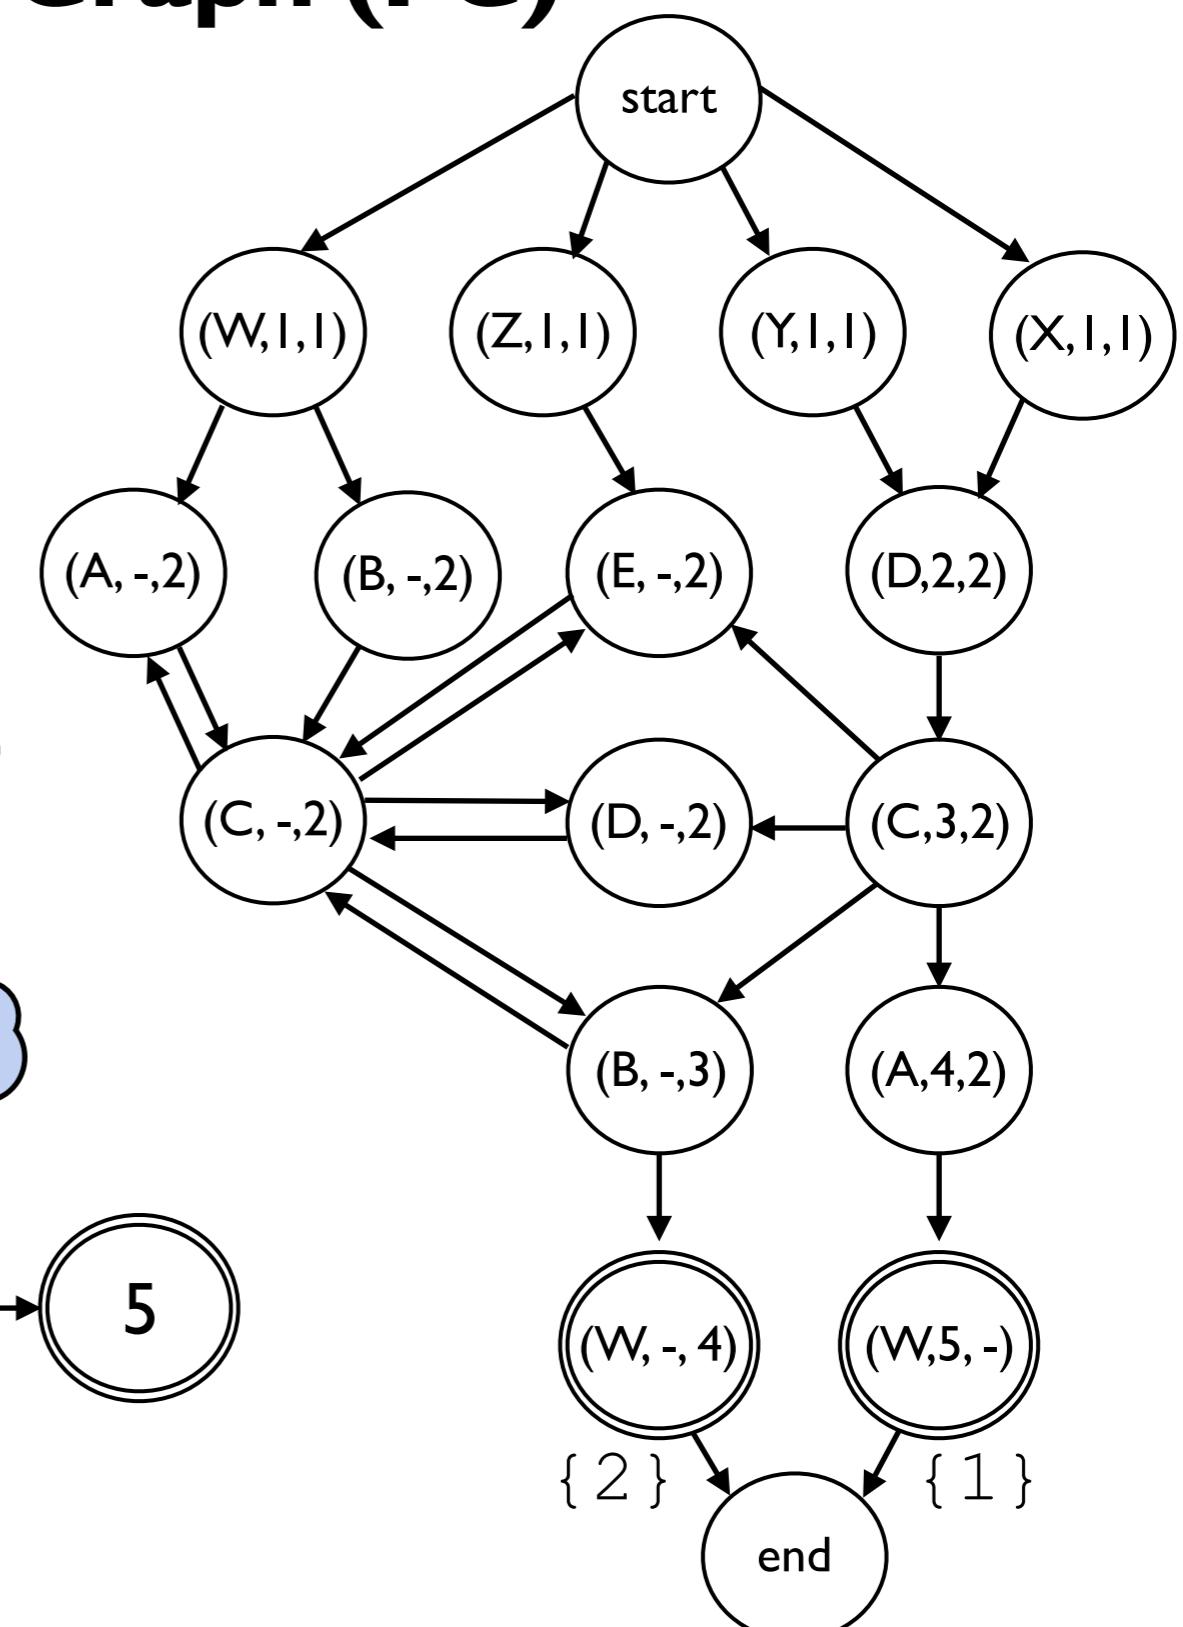
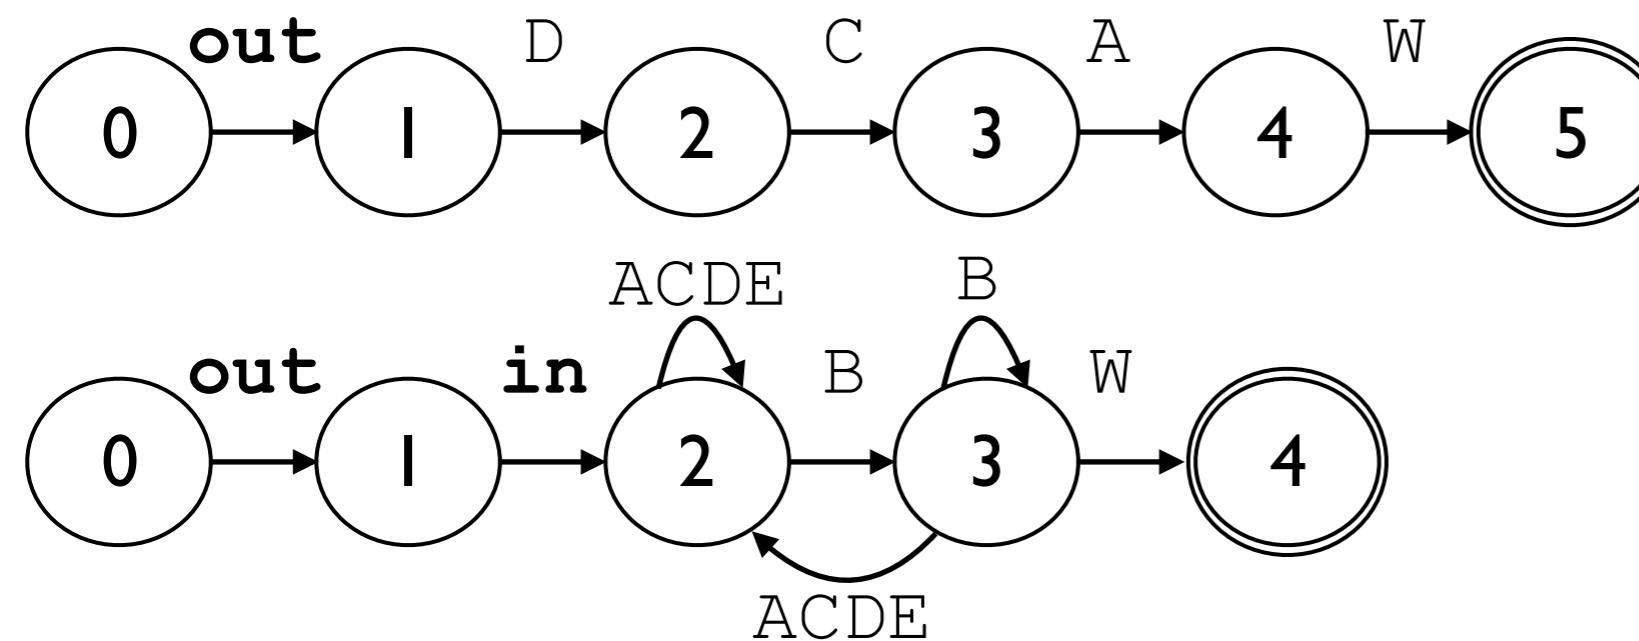
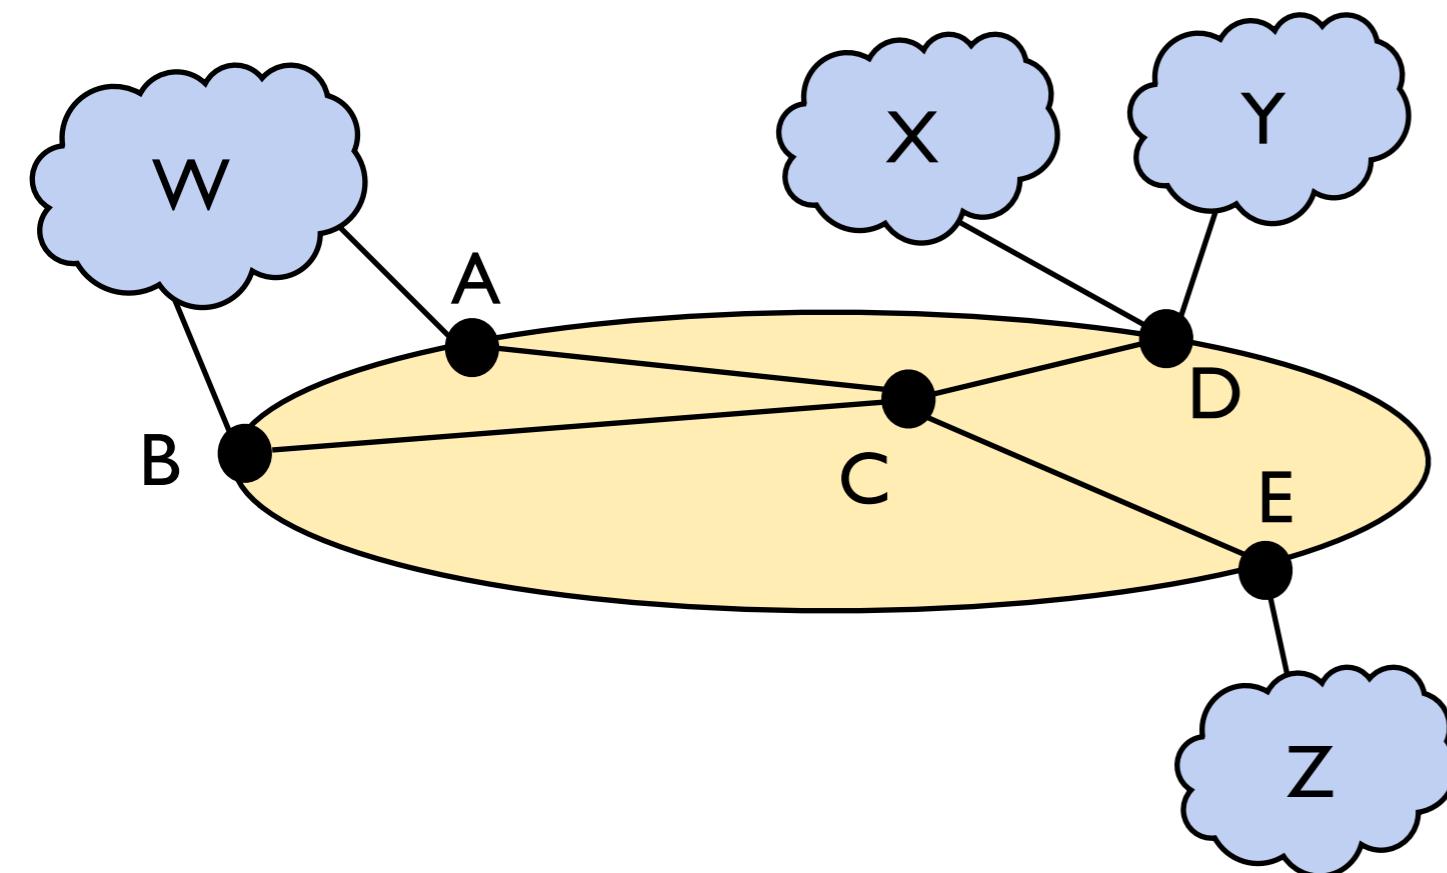
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



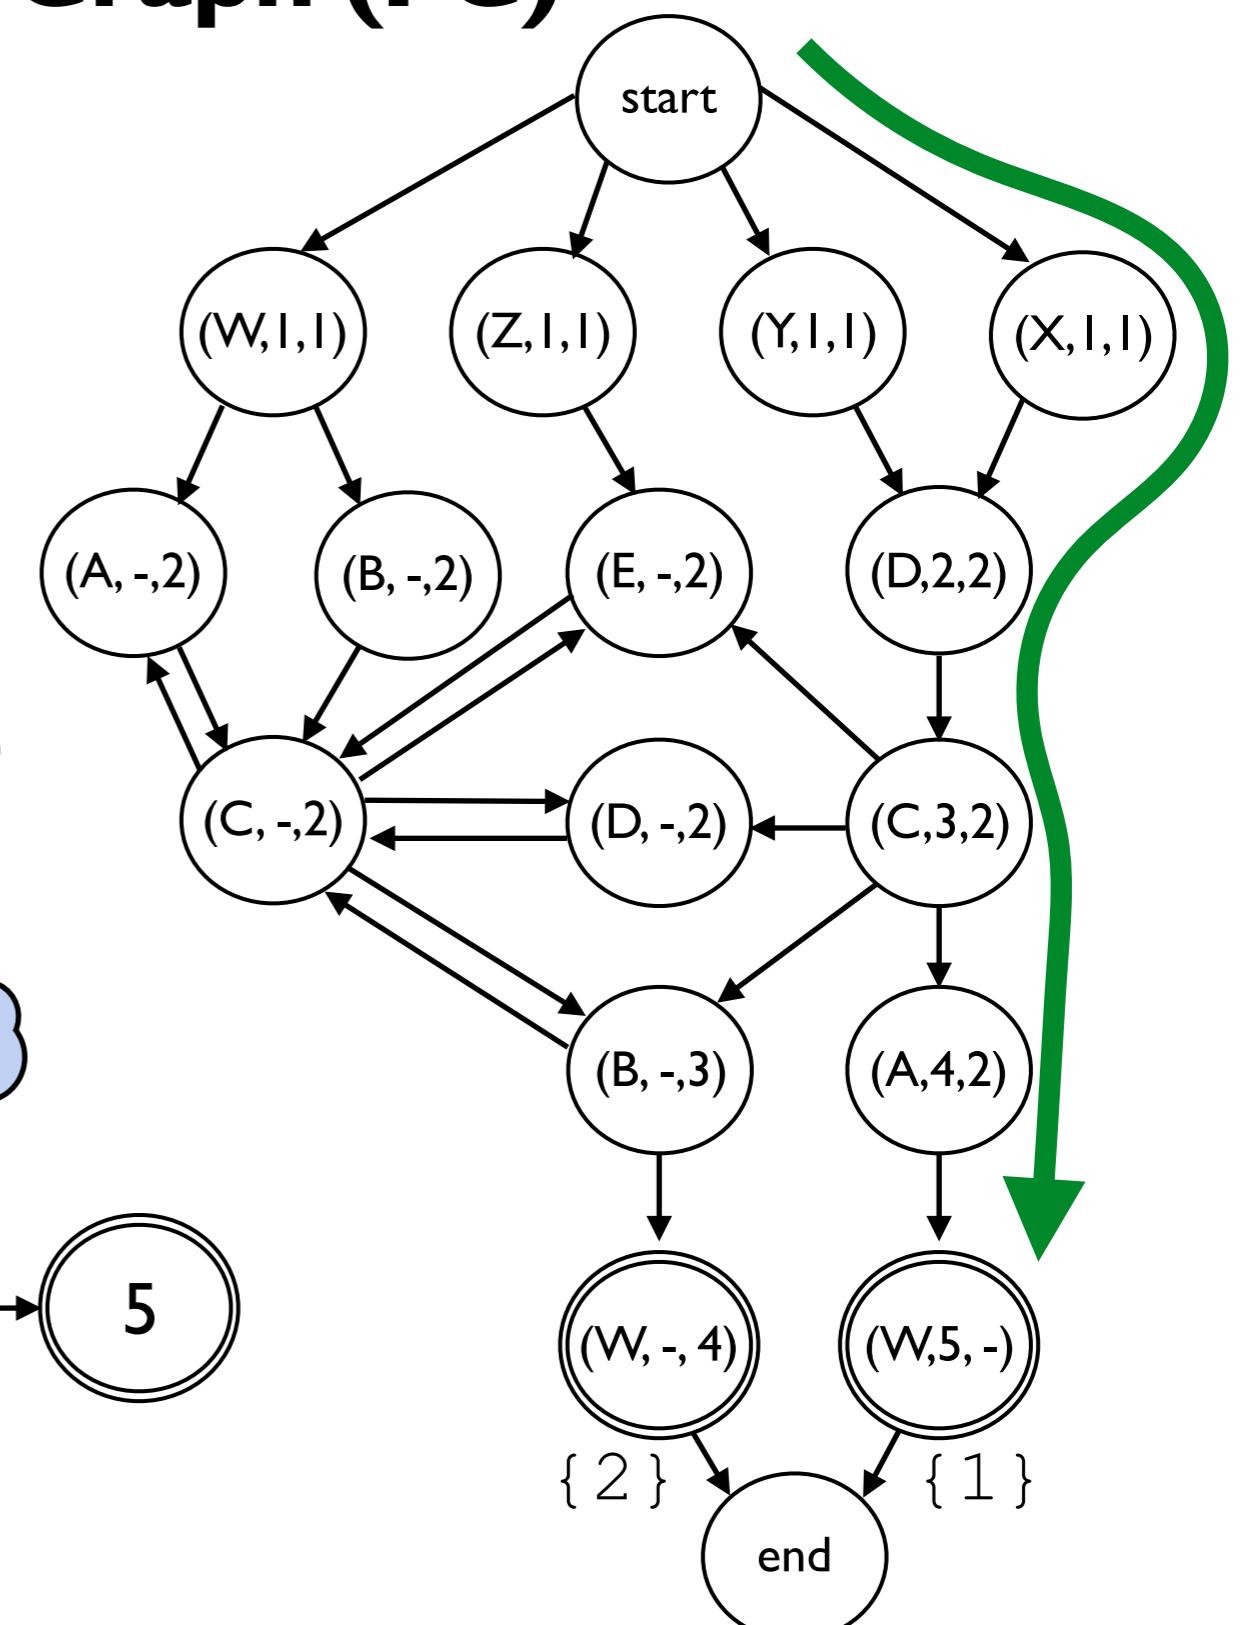
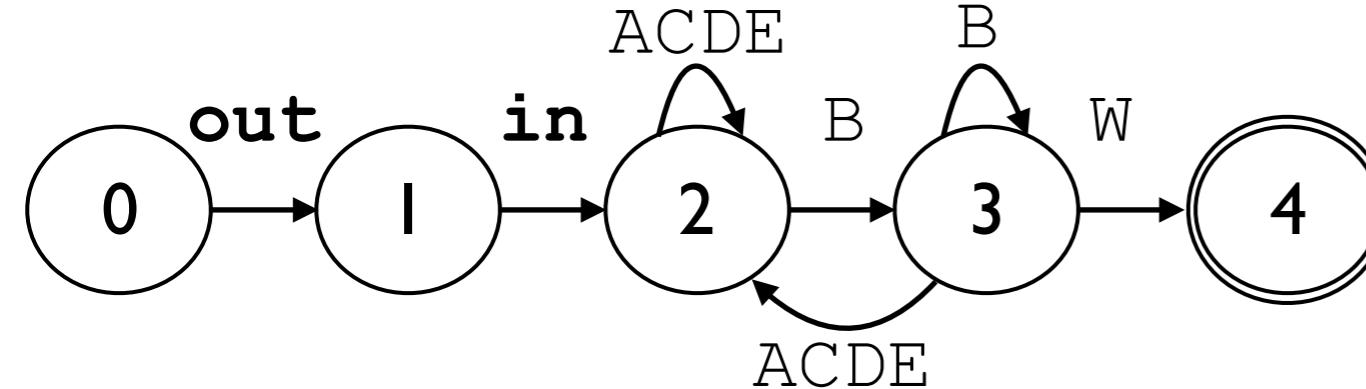
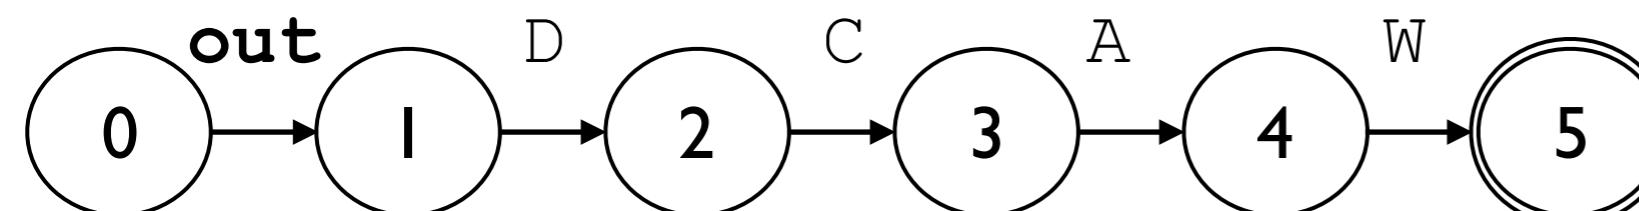
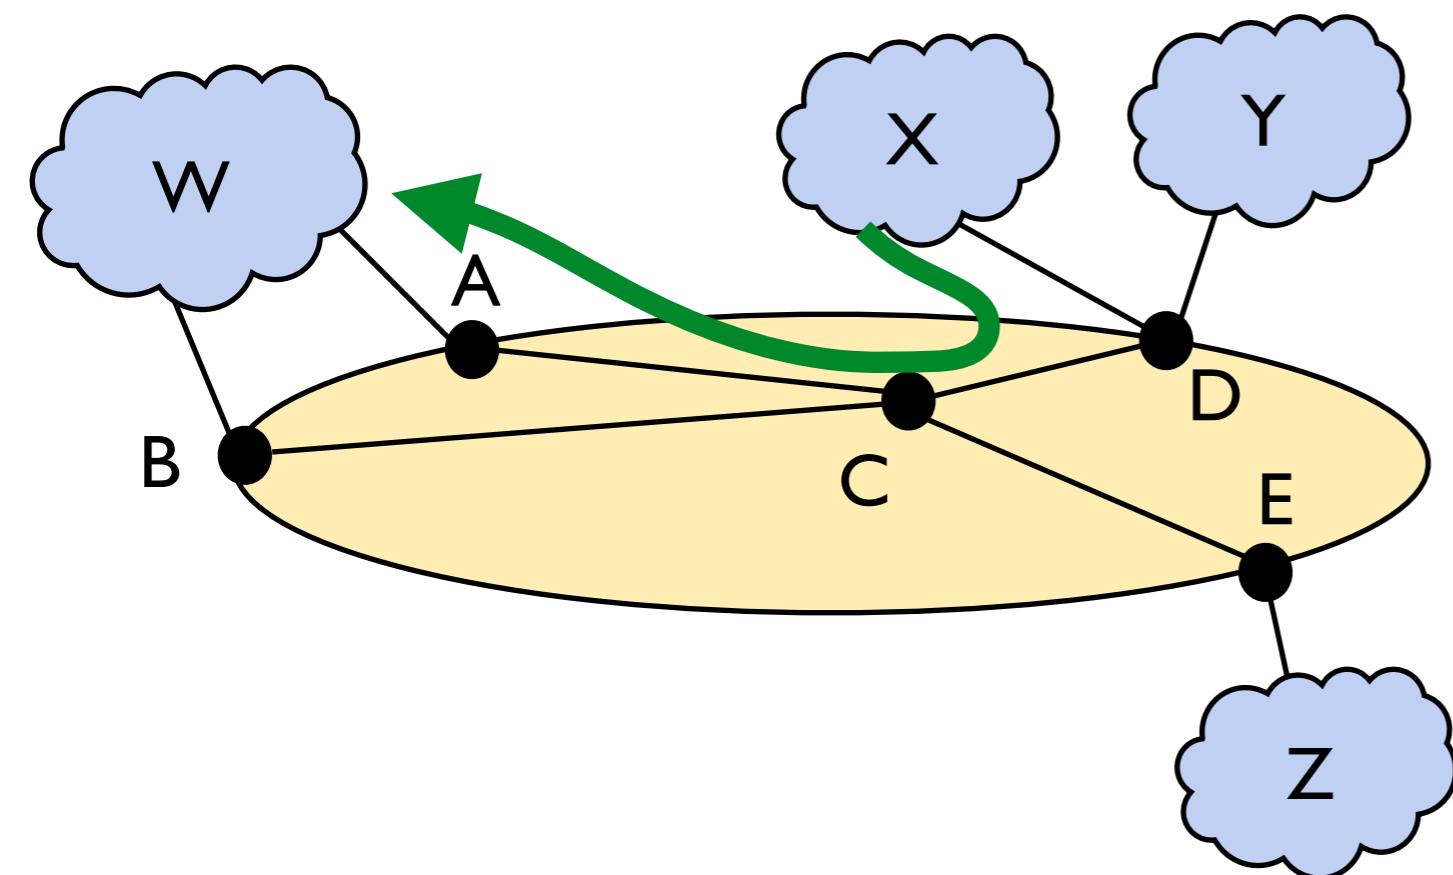
Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$



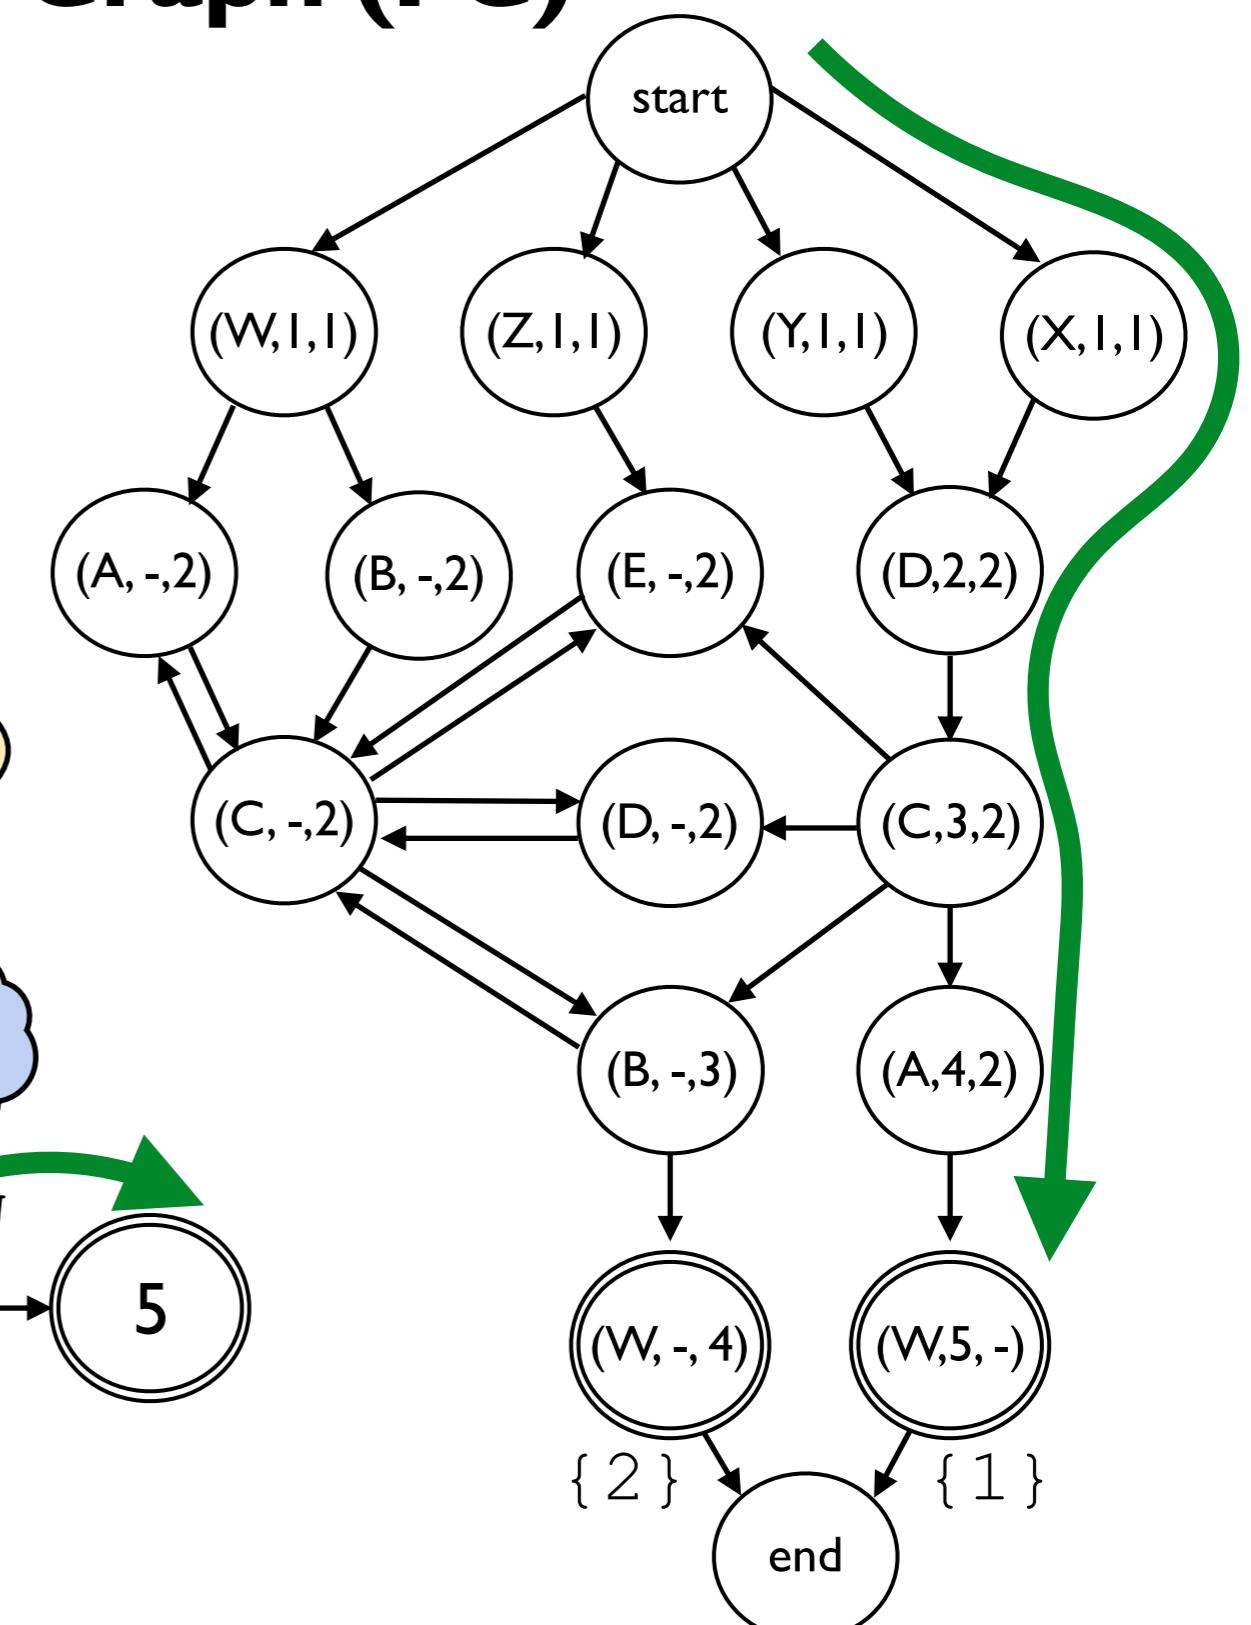
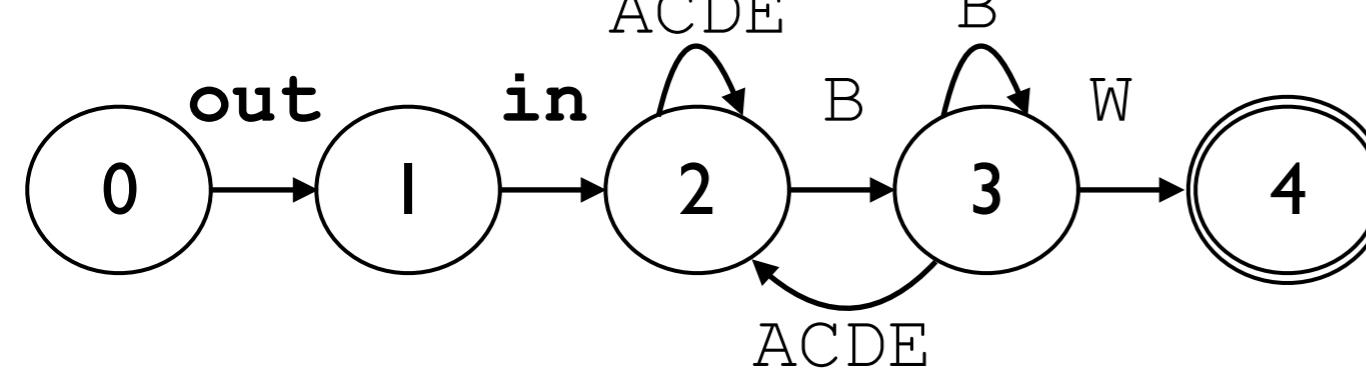
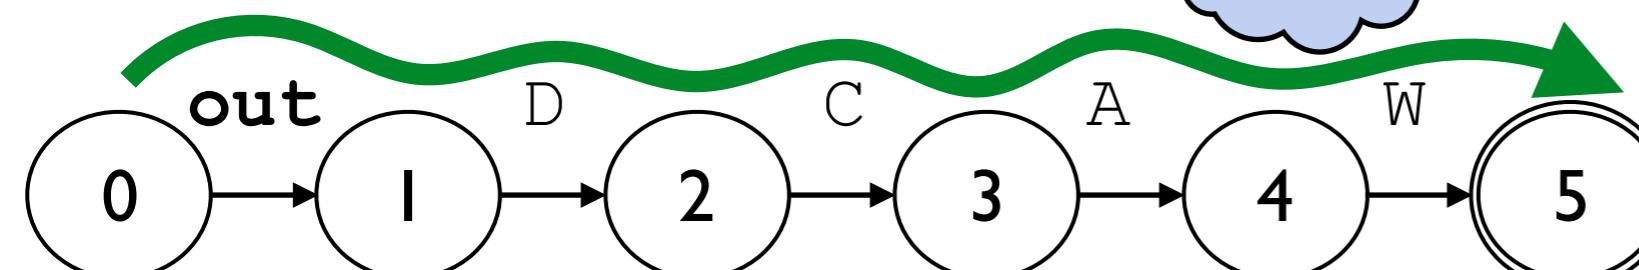
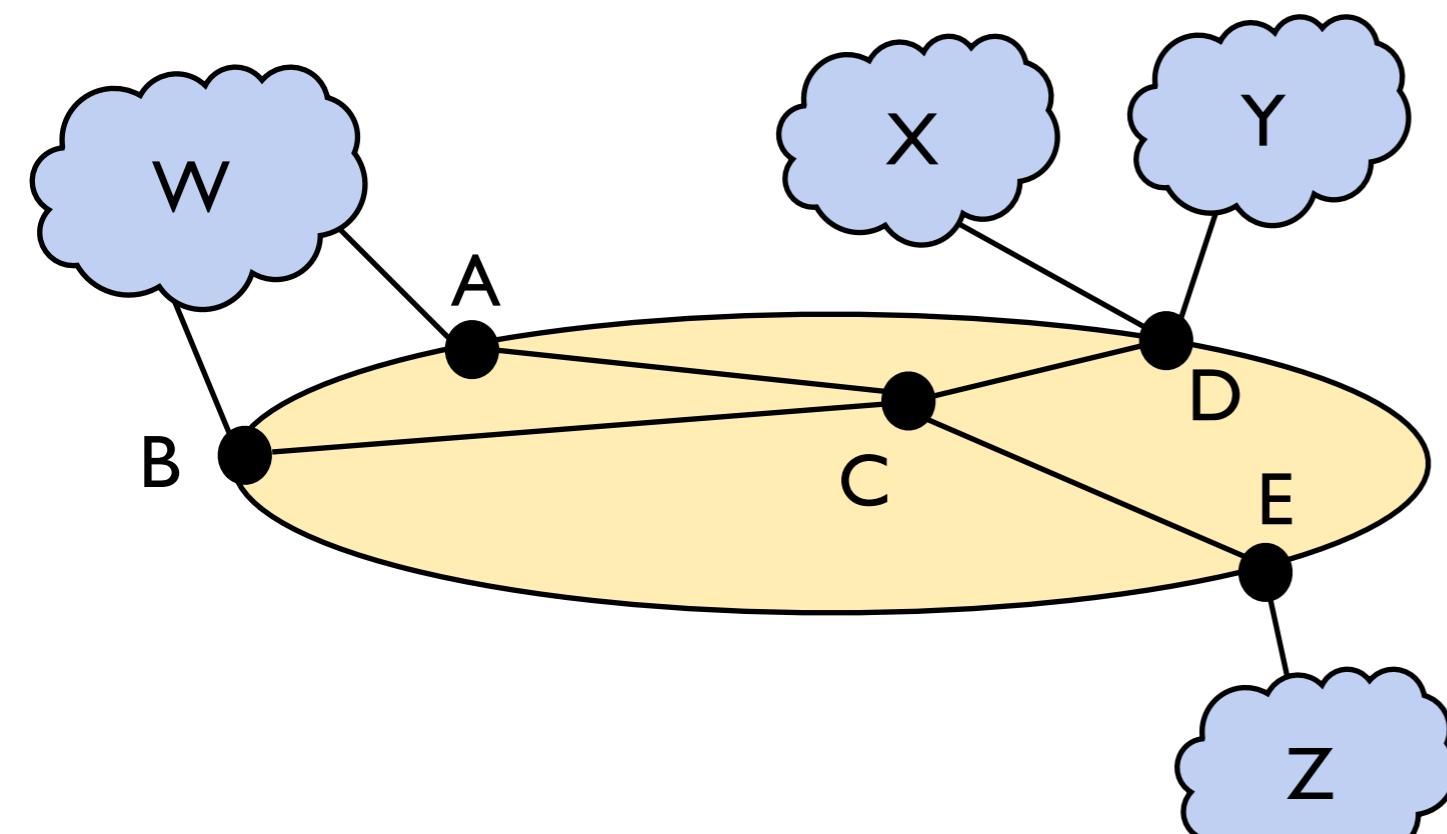
Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$

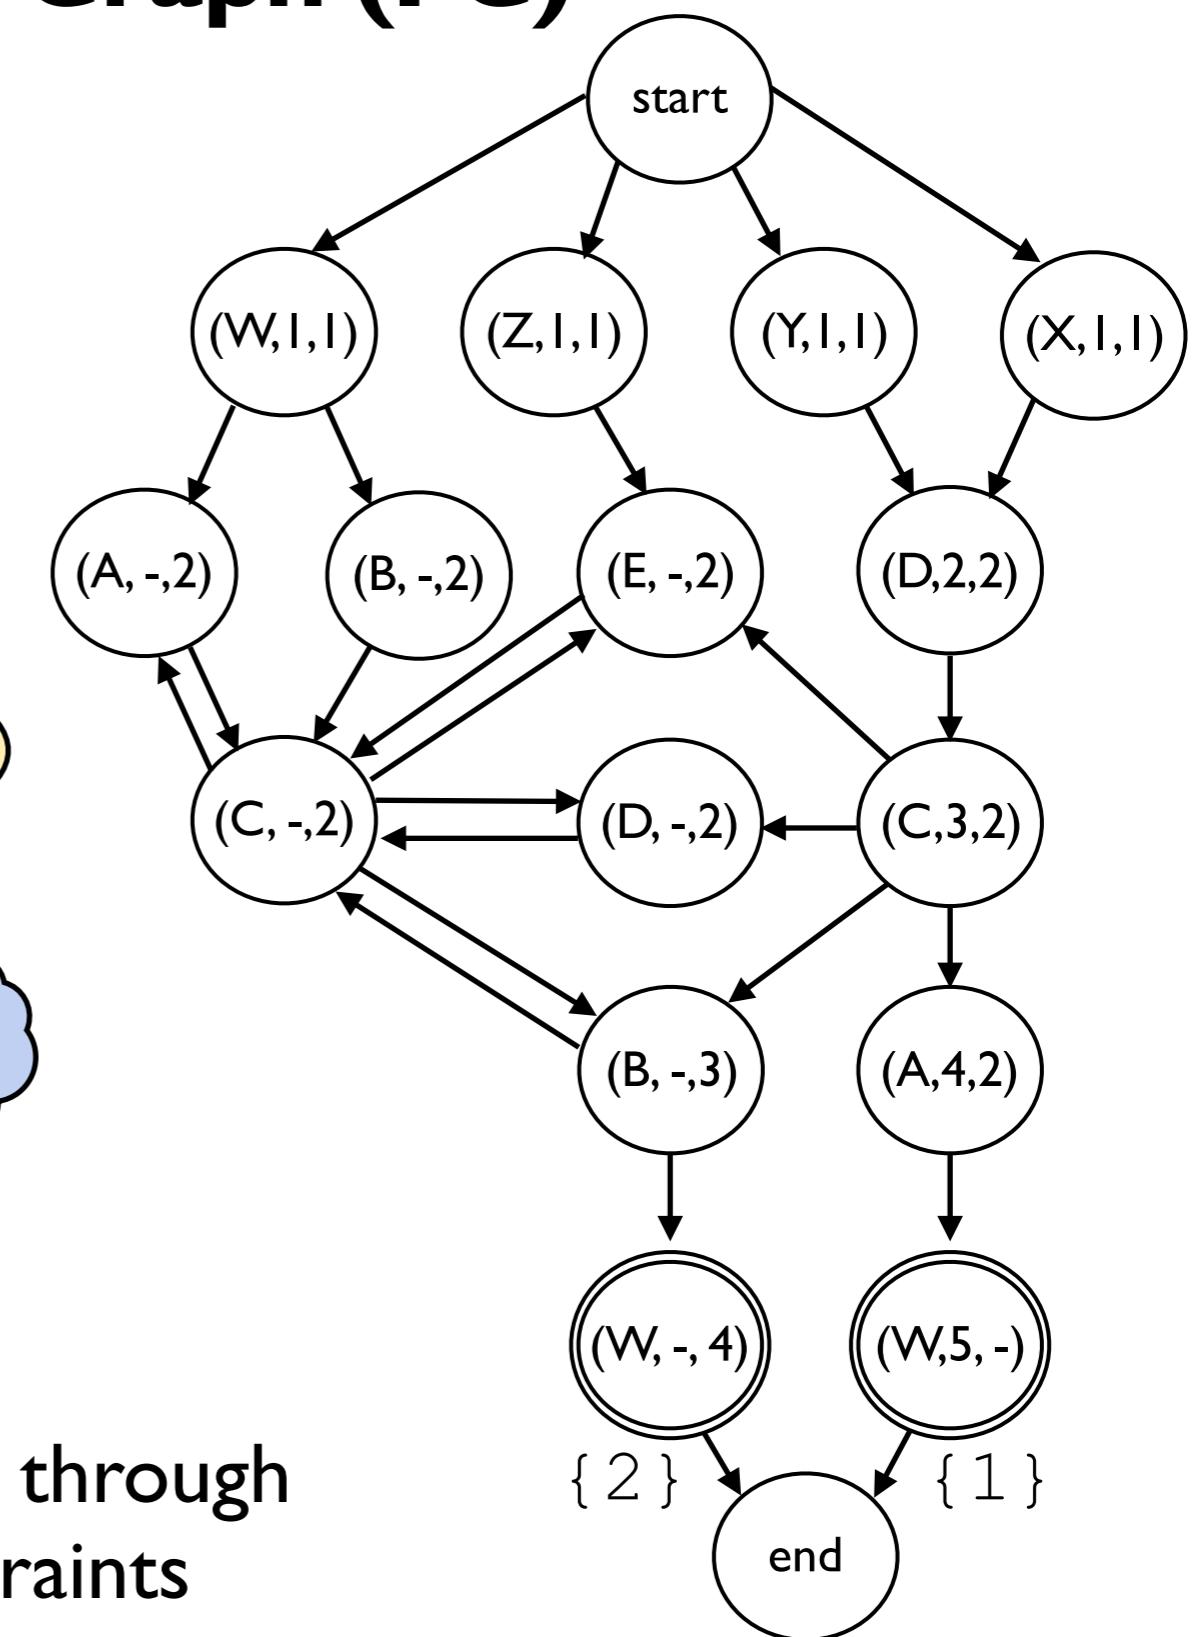
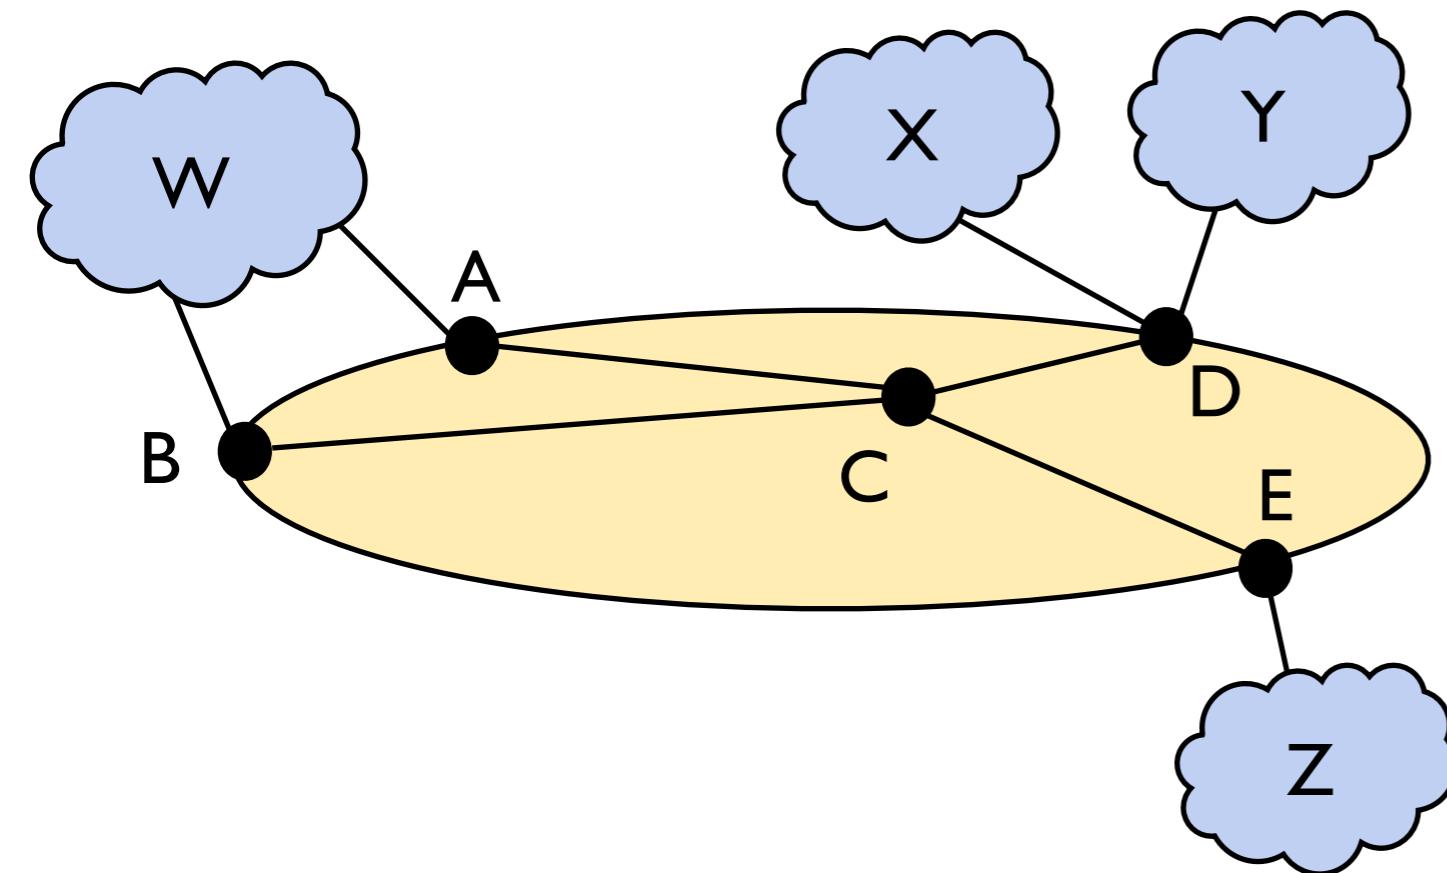


Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$

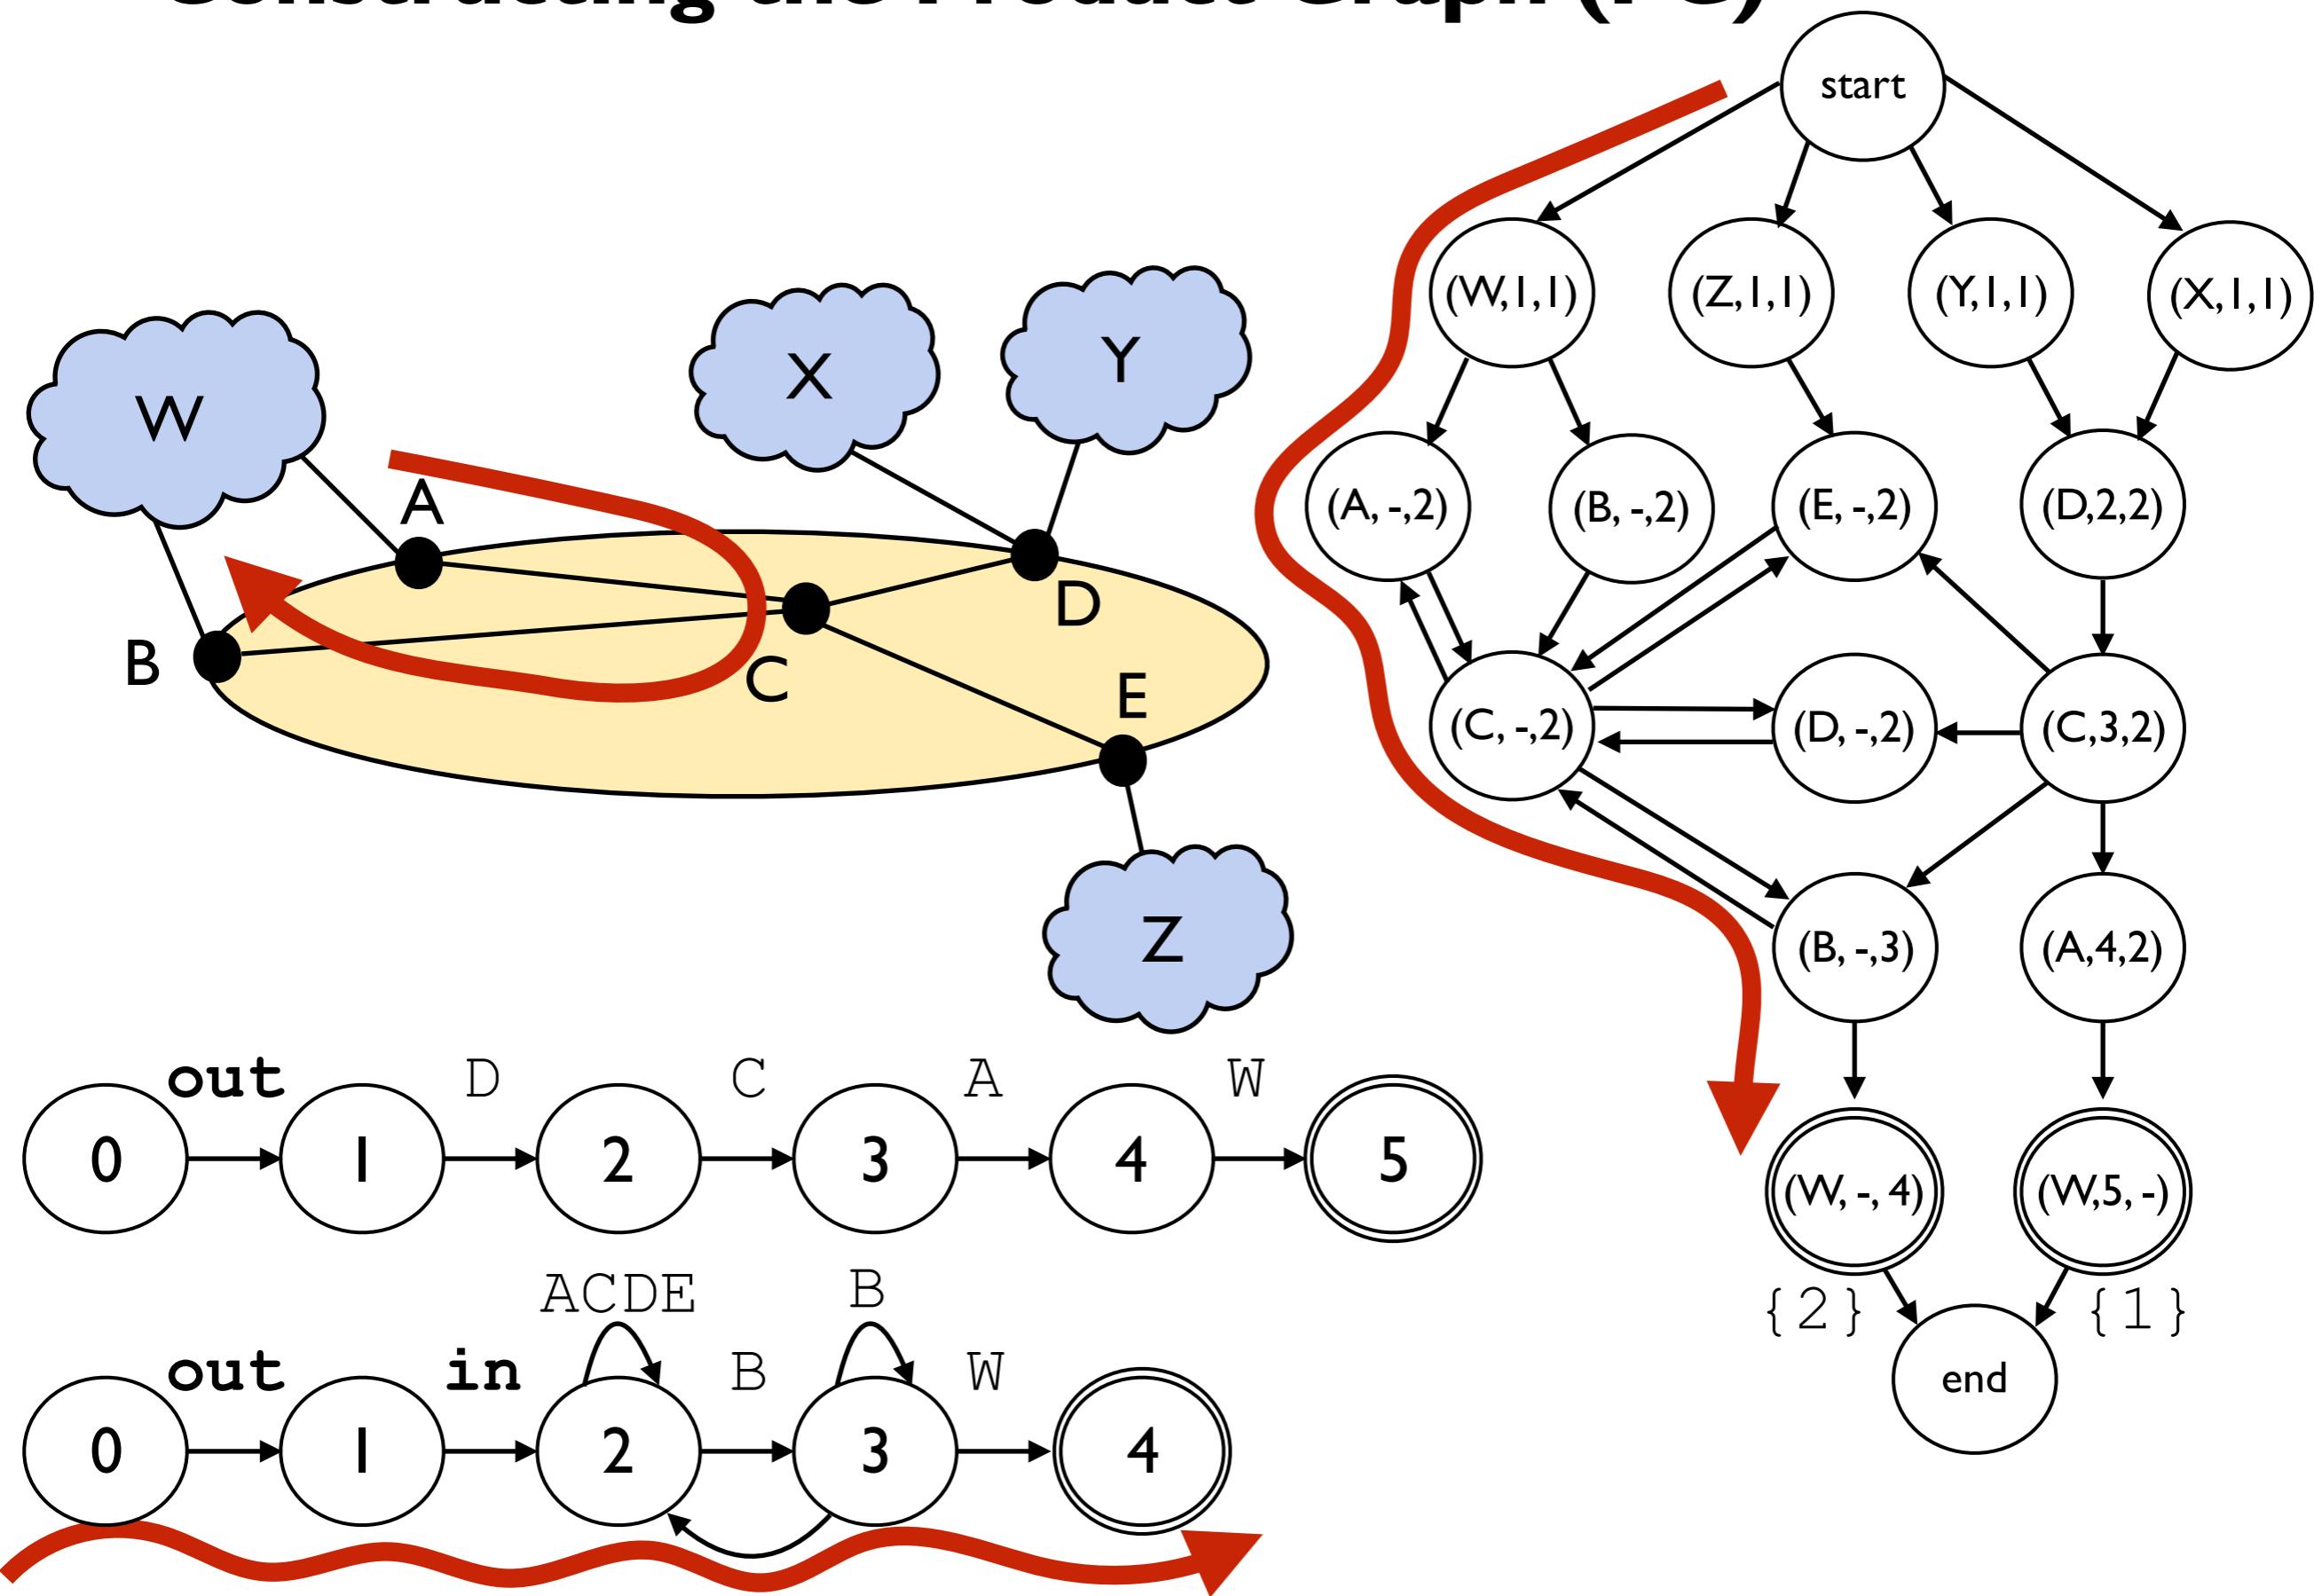


Constructing the Product Graph (PG)

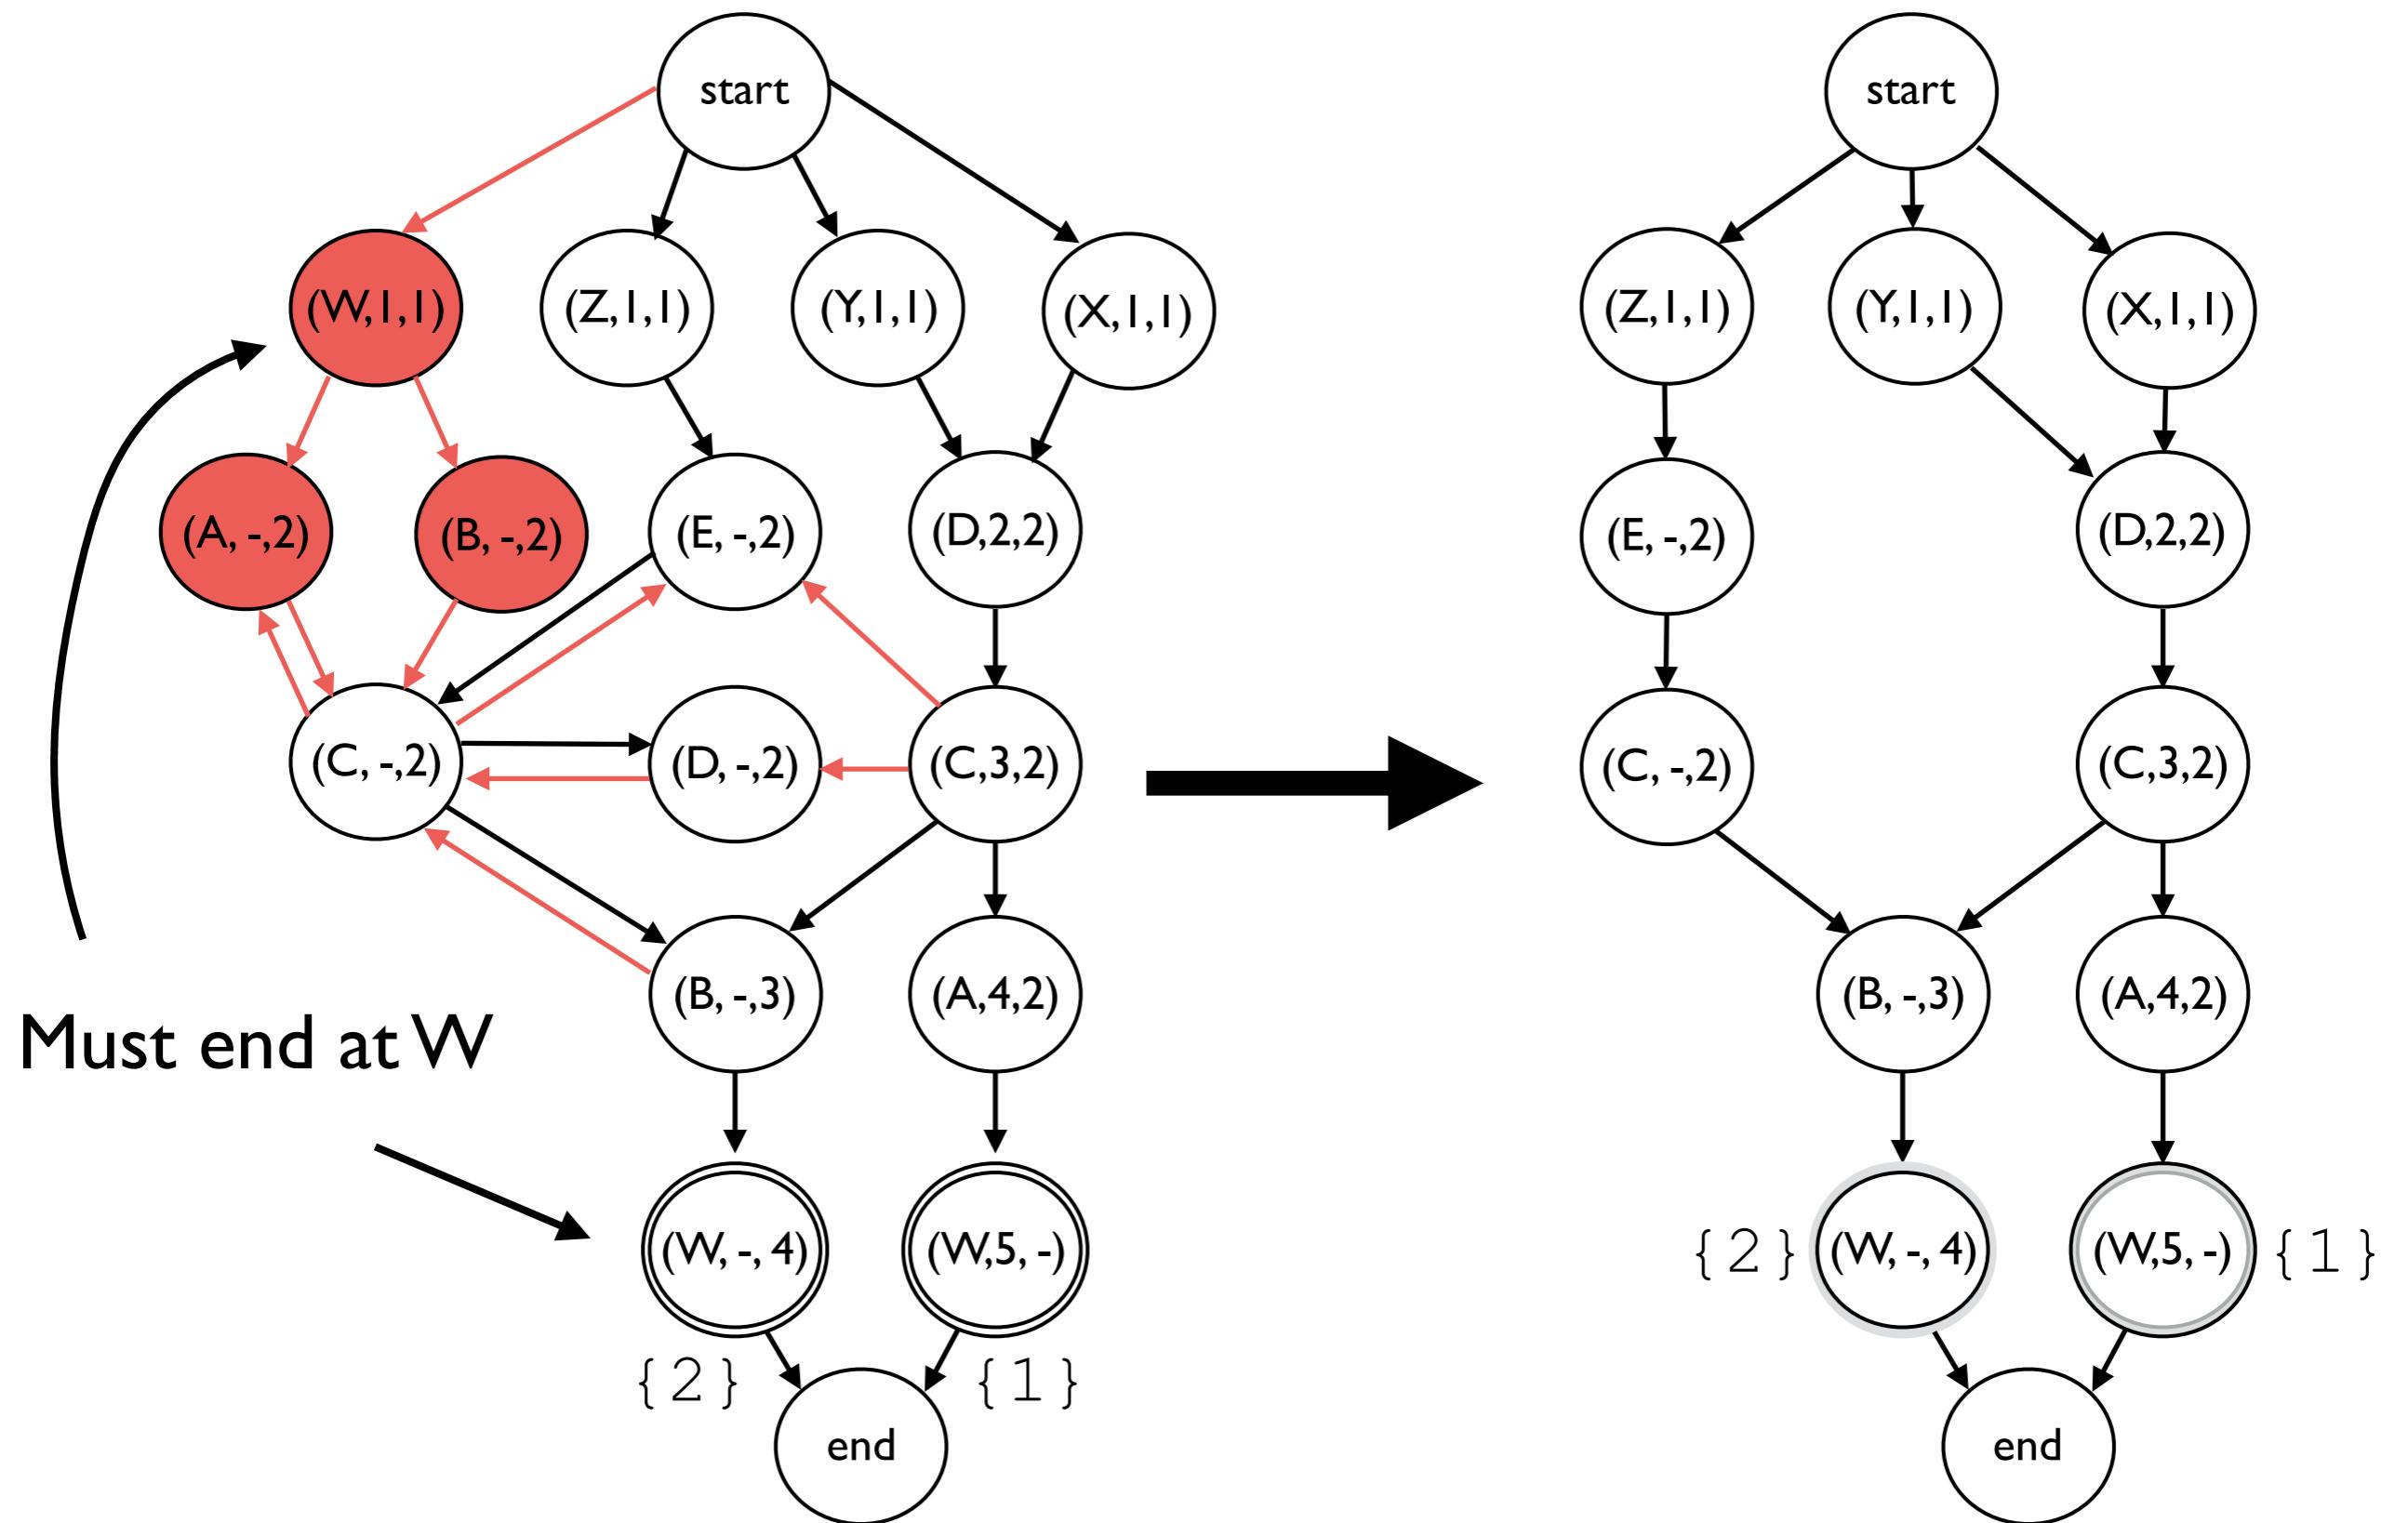


Compact representation of all paths through
the topology subject to policy constraints

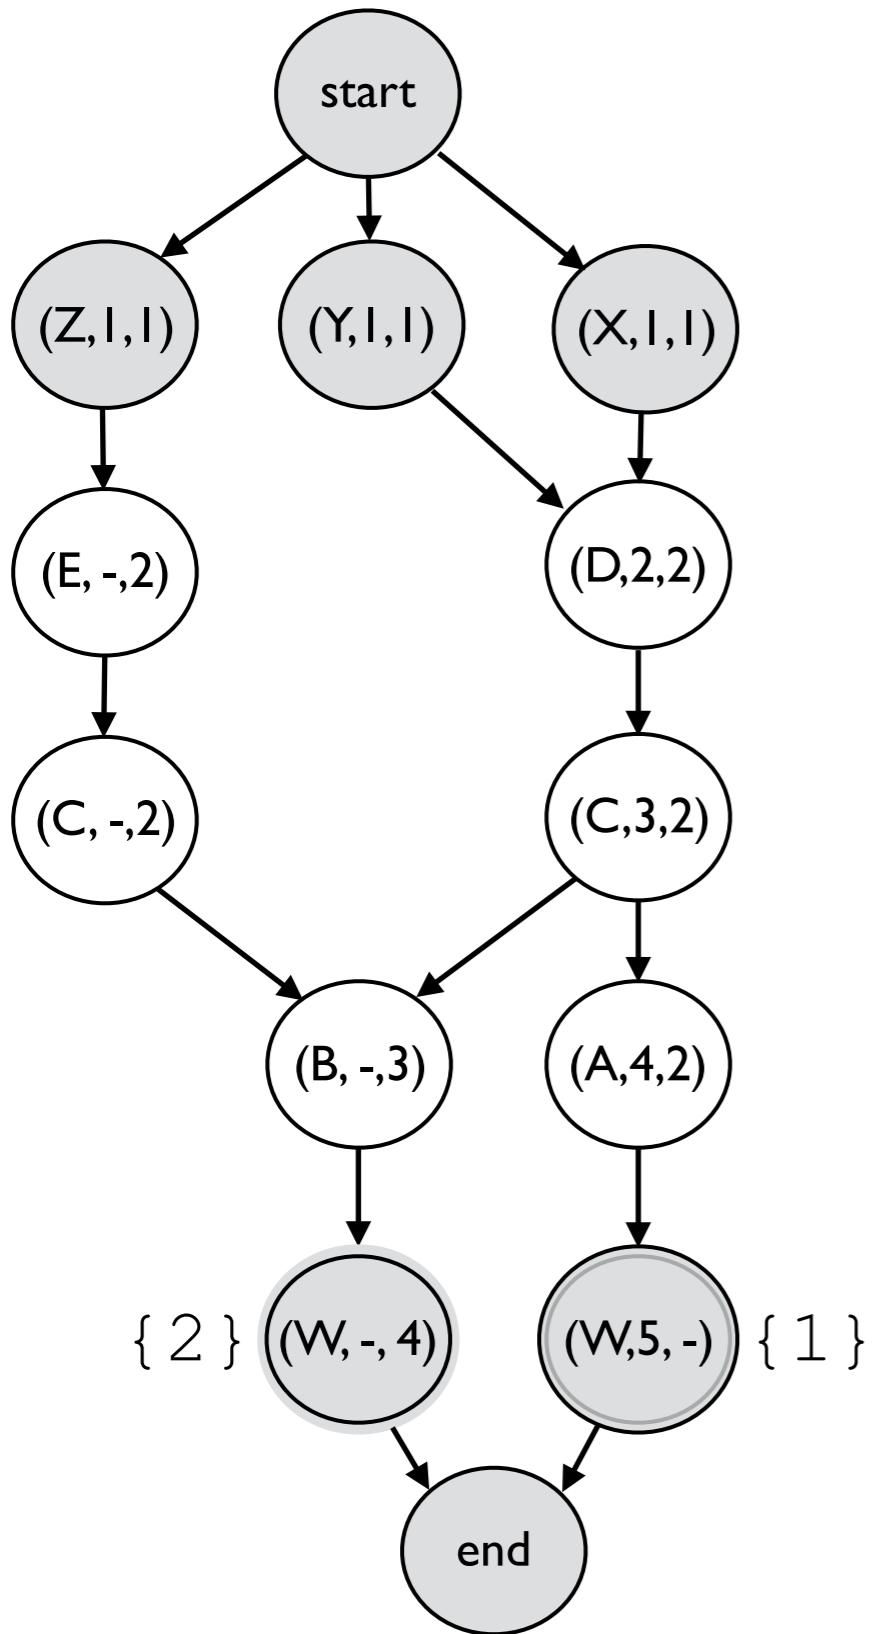
Constructing the Product Graph (PG)



PG Minimization (Loop analysis)

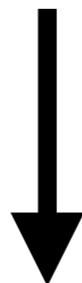


Compilation to BGP:



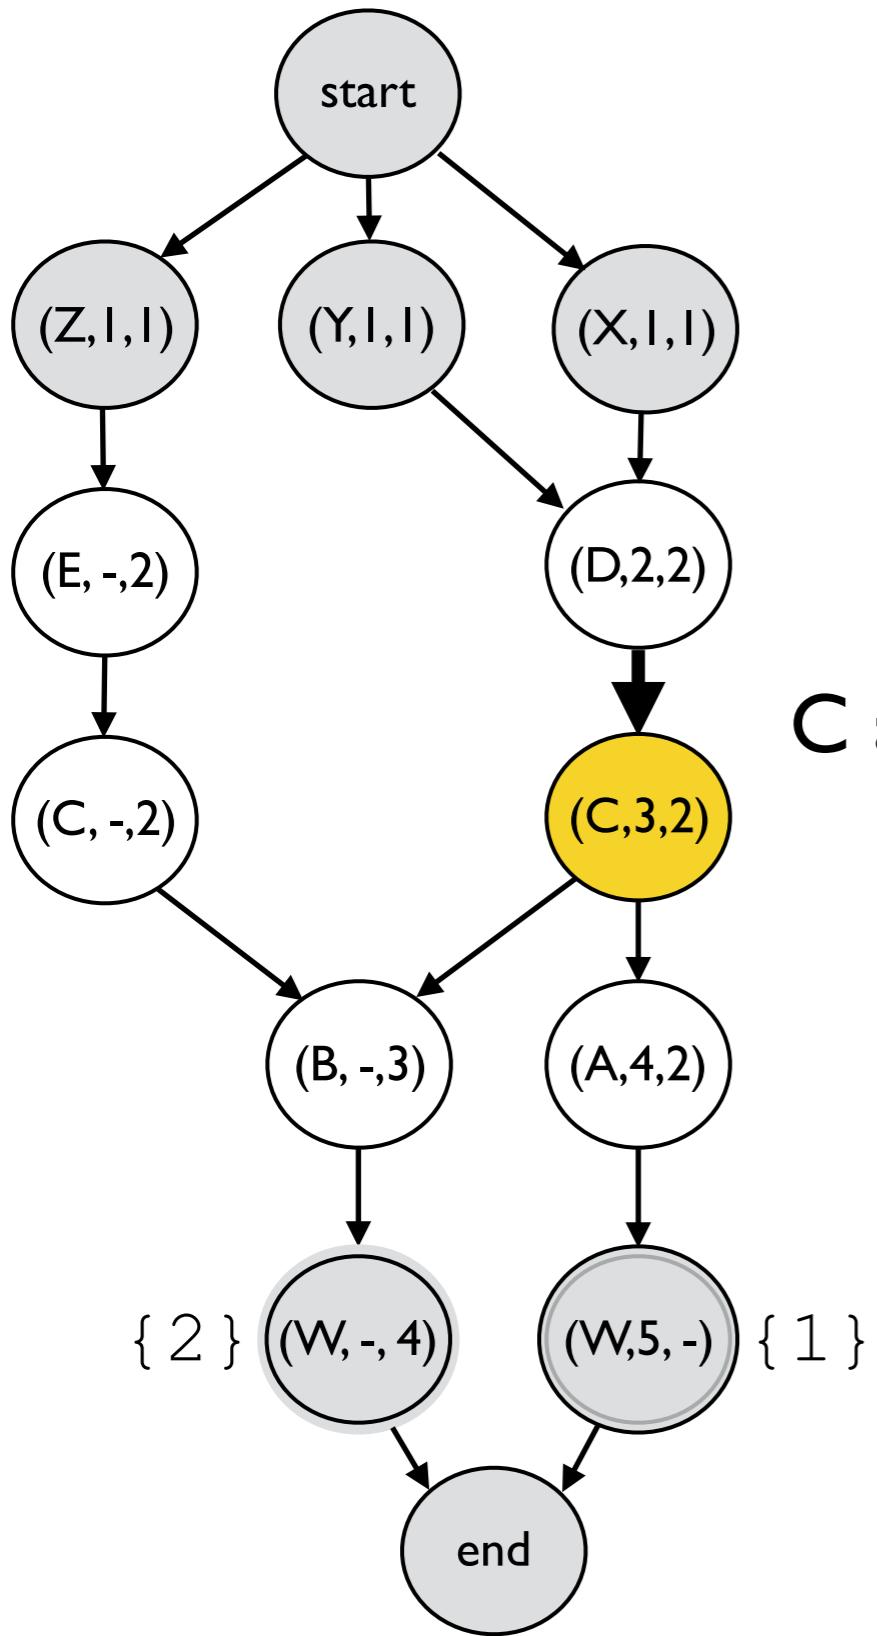
Idea I: Restrict advertisements to PG edges

- Encode PG state in community tag
- Incoming edges — import filters
- Outgoing edges — export filters

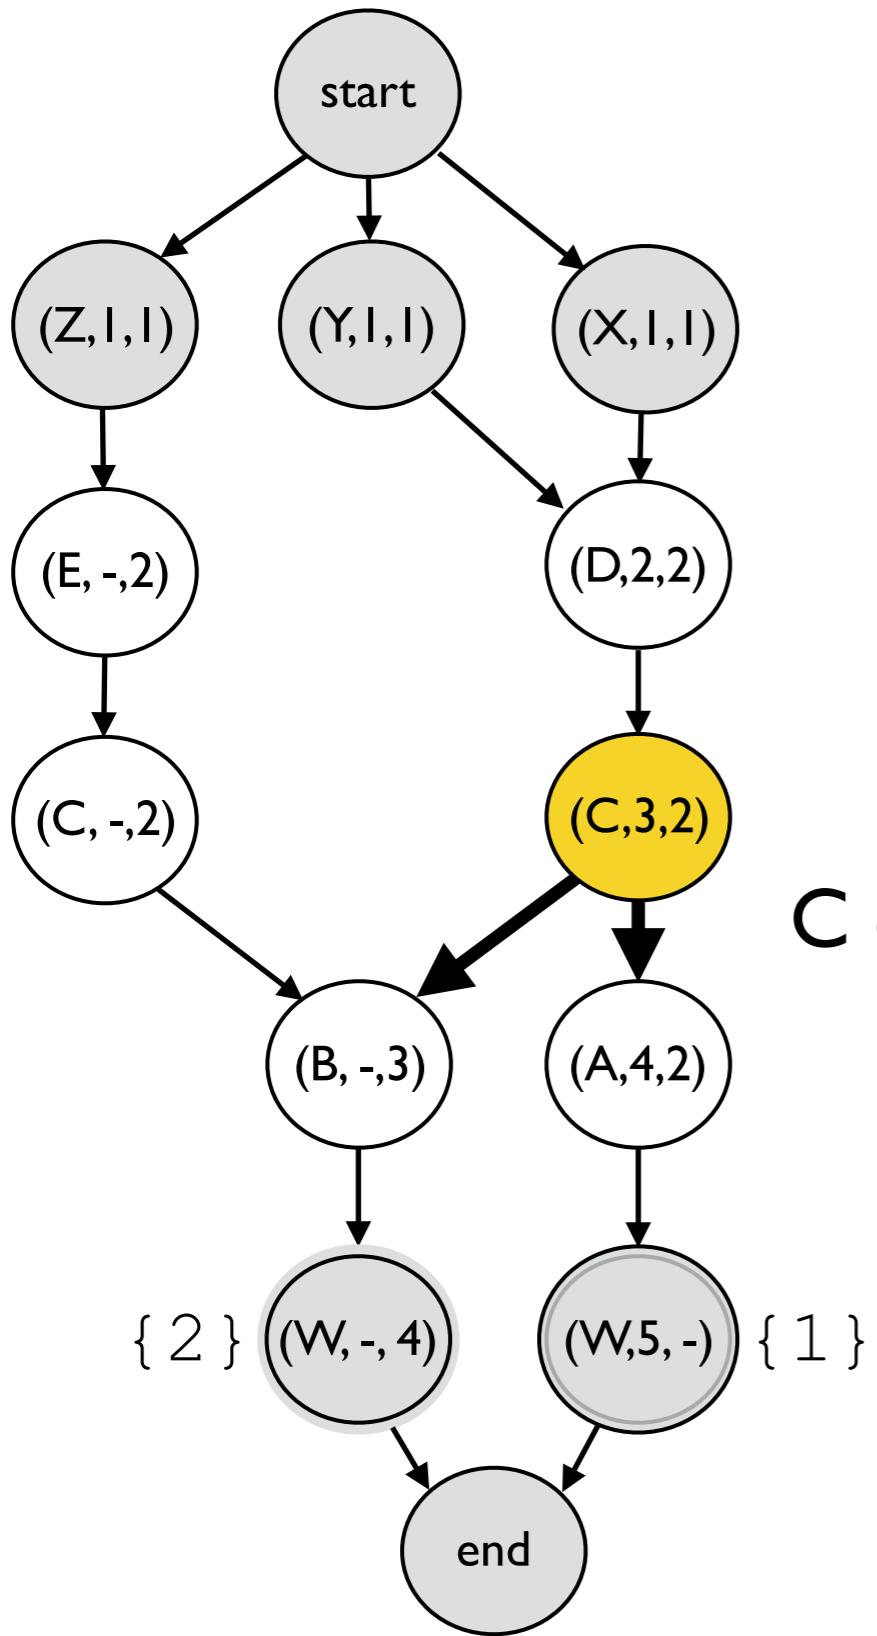


Let BGP find **some allowed** path dynamically

Compilation to BGP:

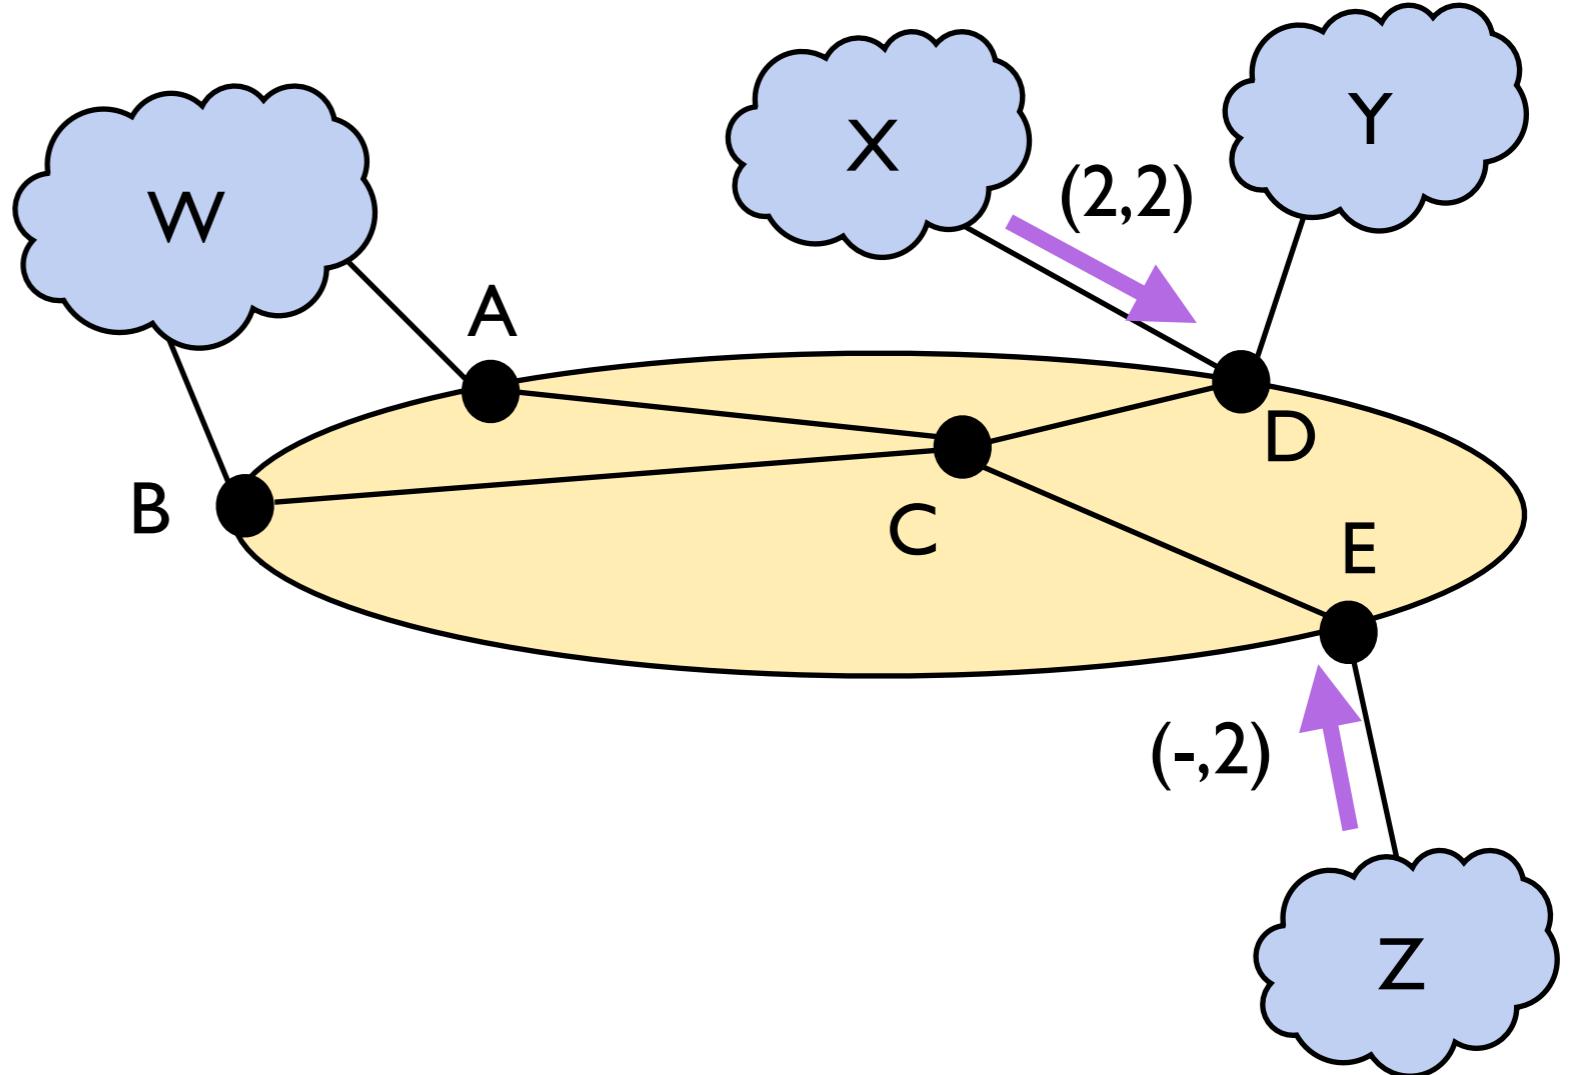
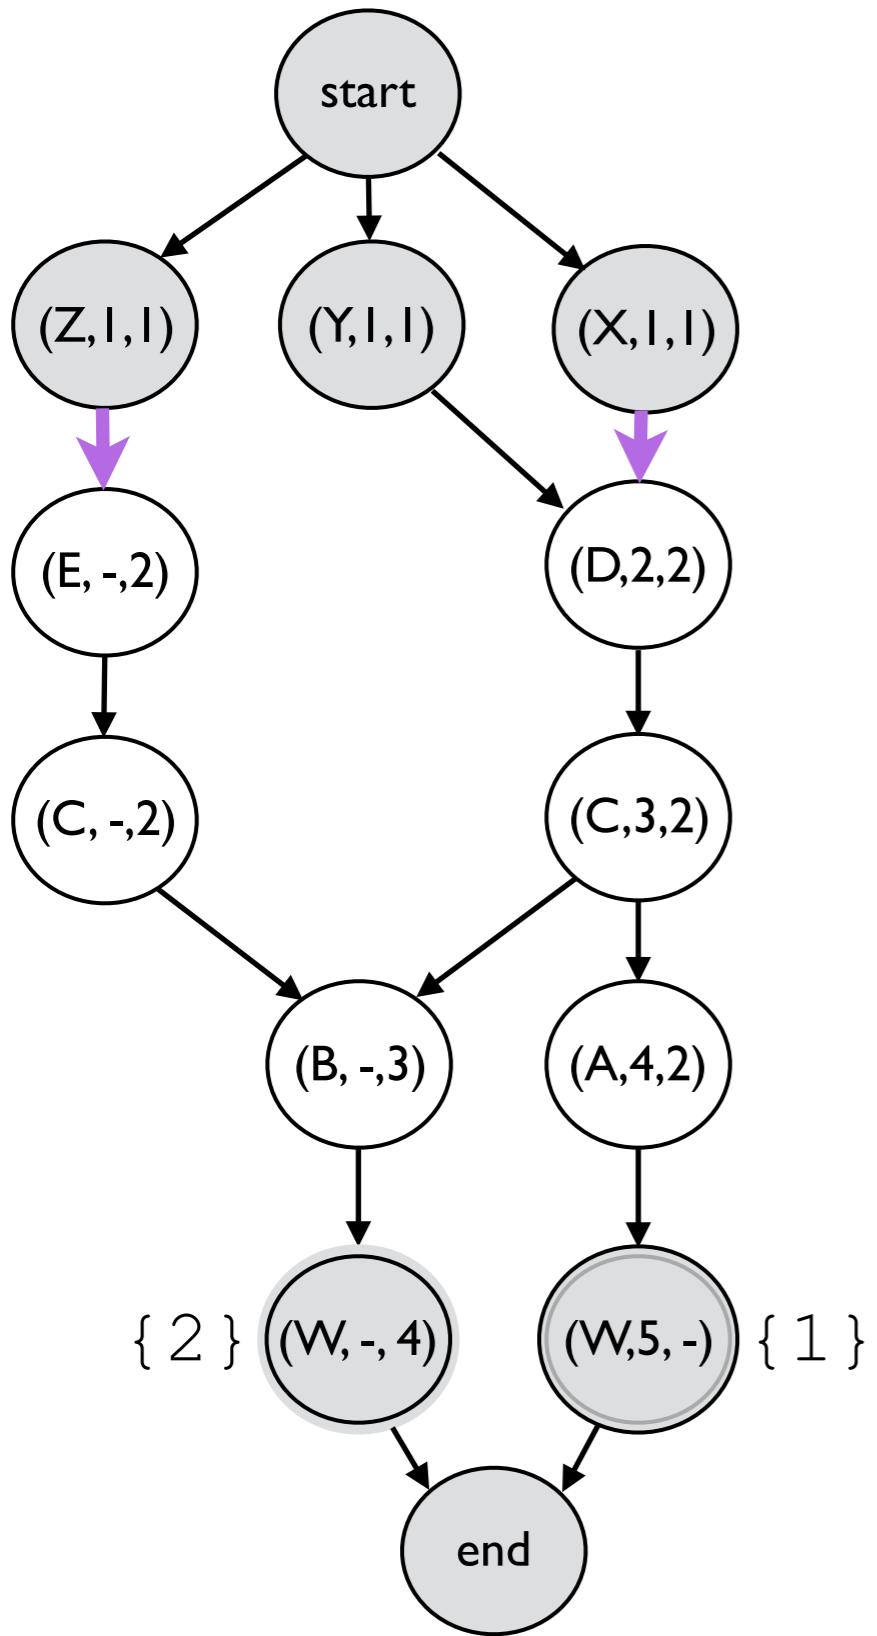


Compilation to BGP:



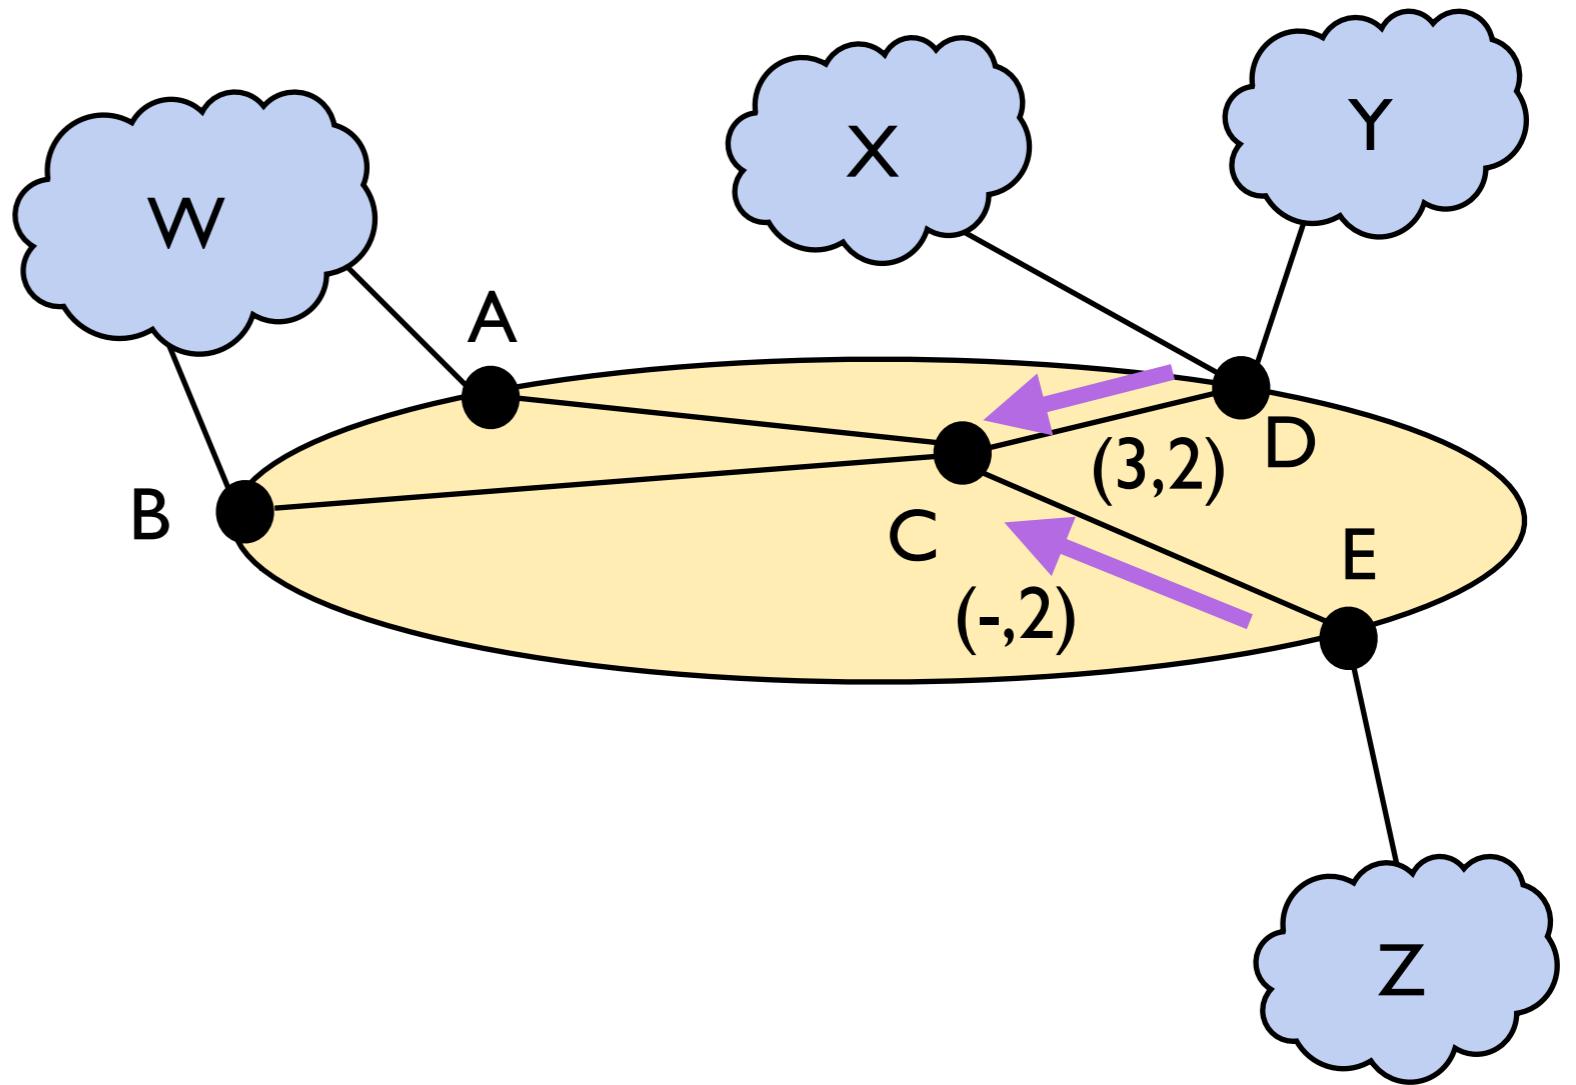
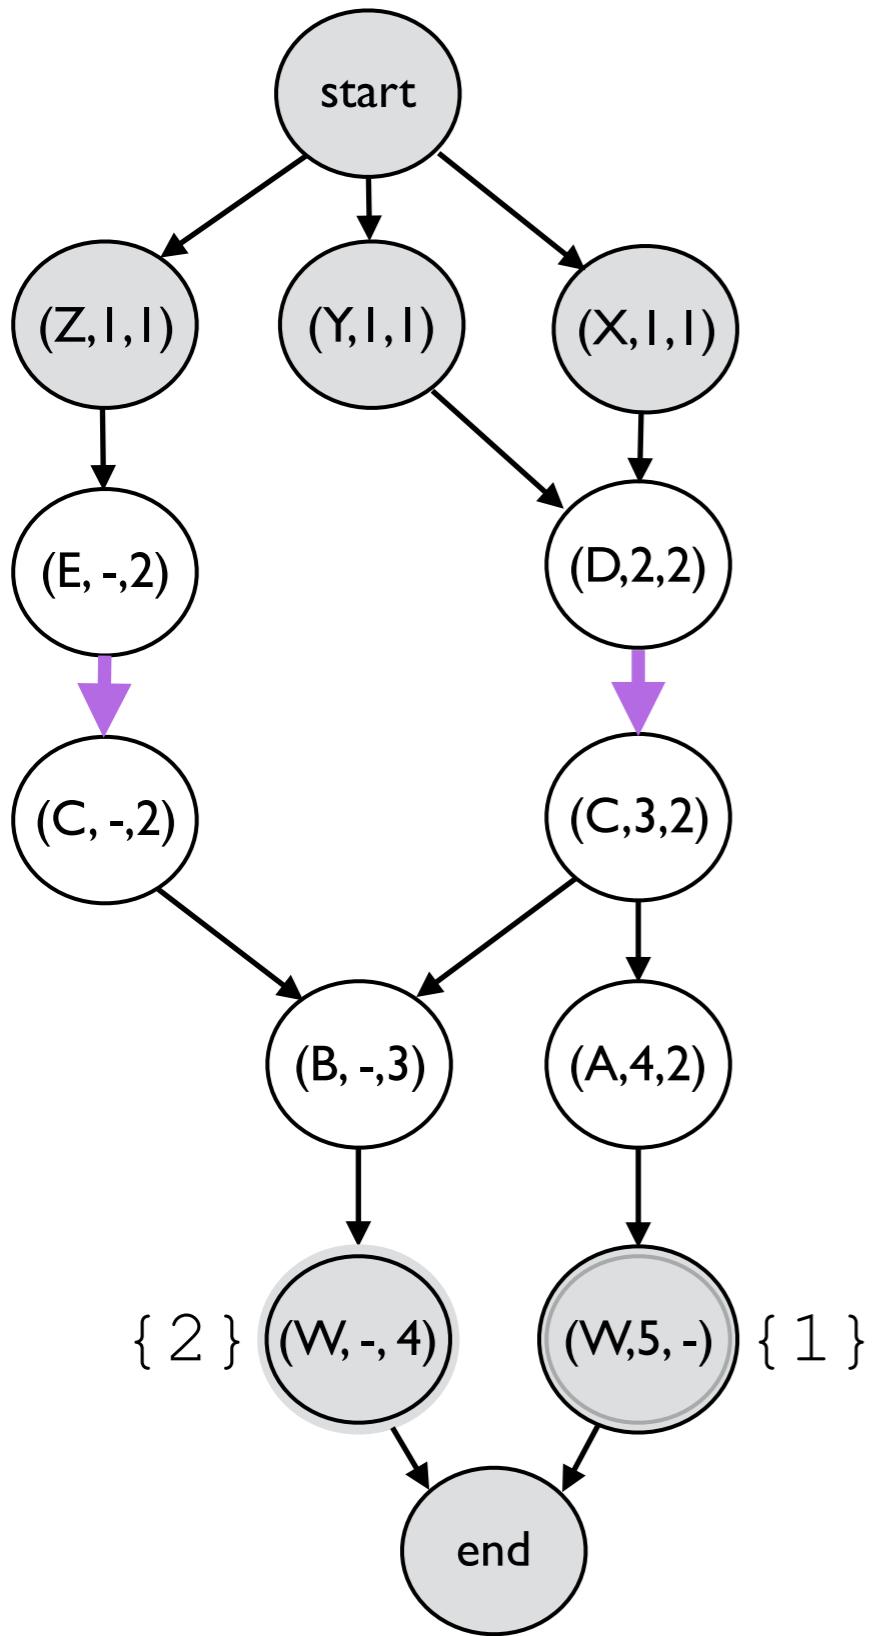
C exports to A,B with tag (3,2)

Compilation to BGP:



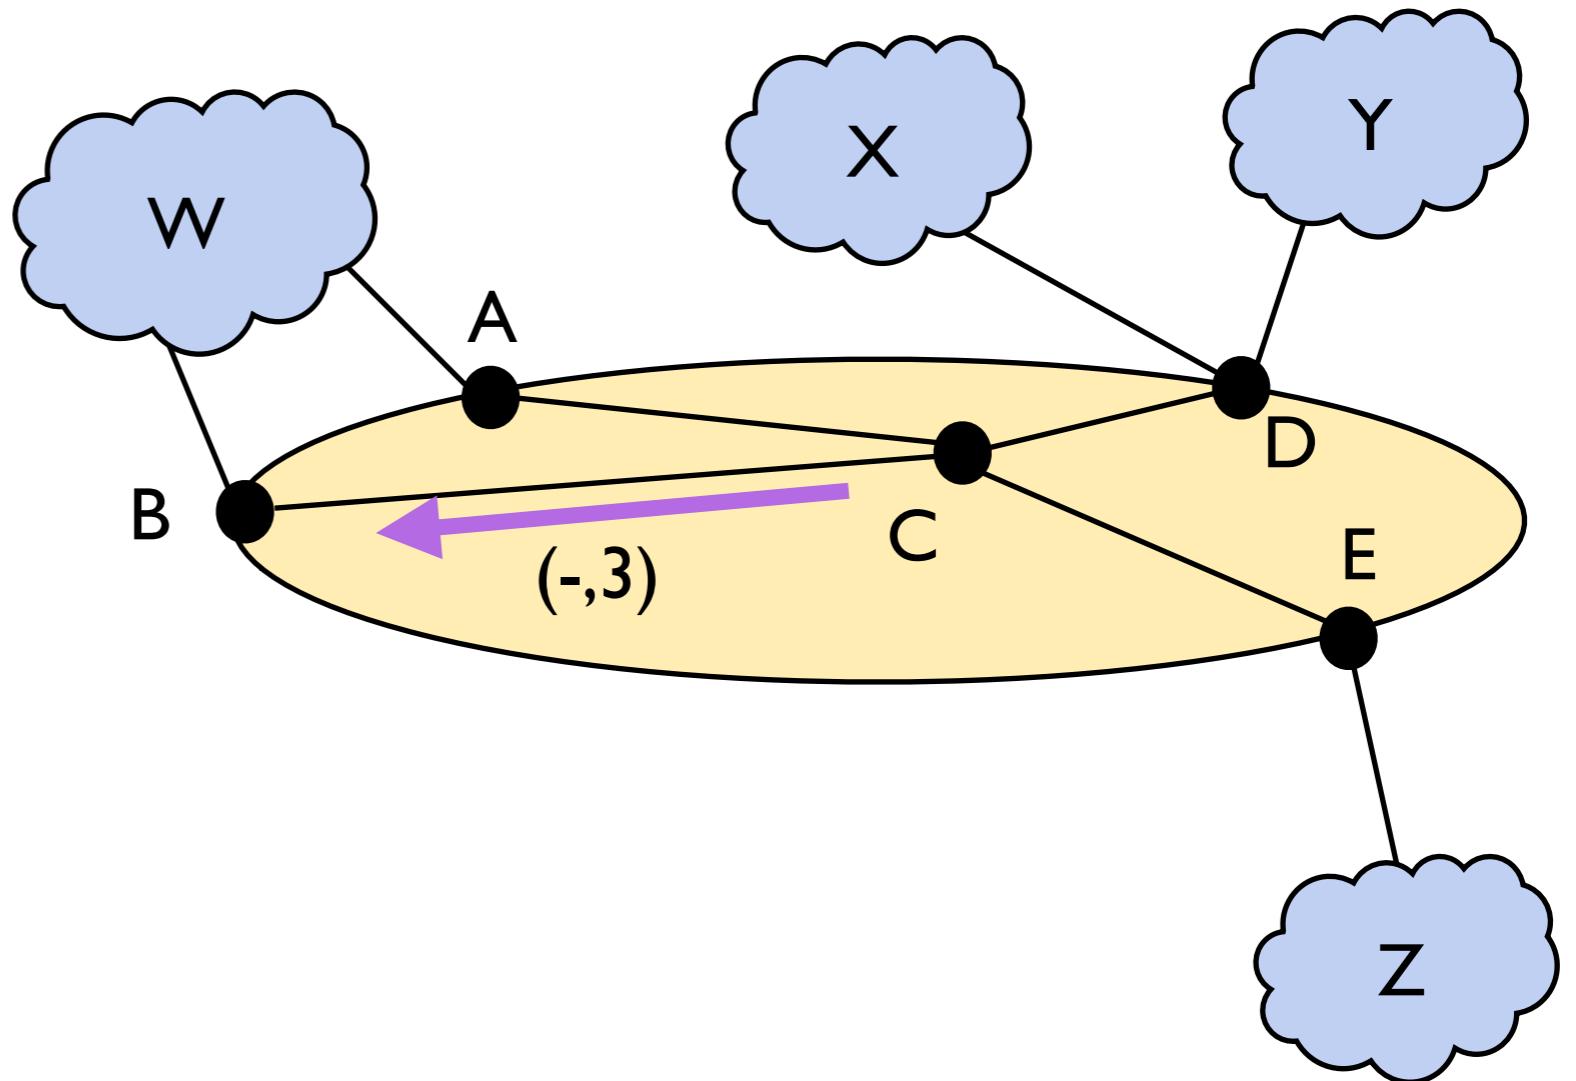
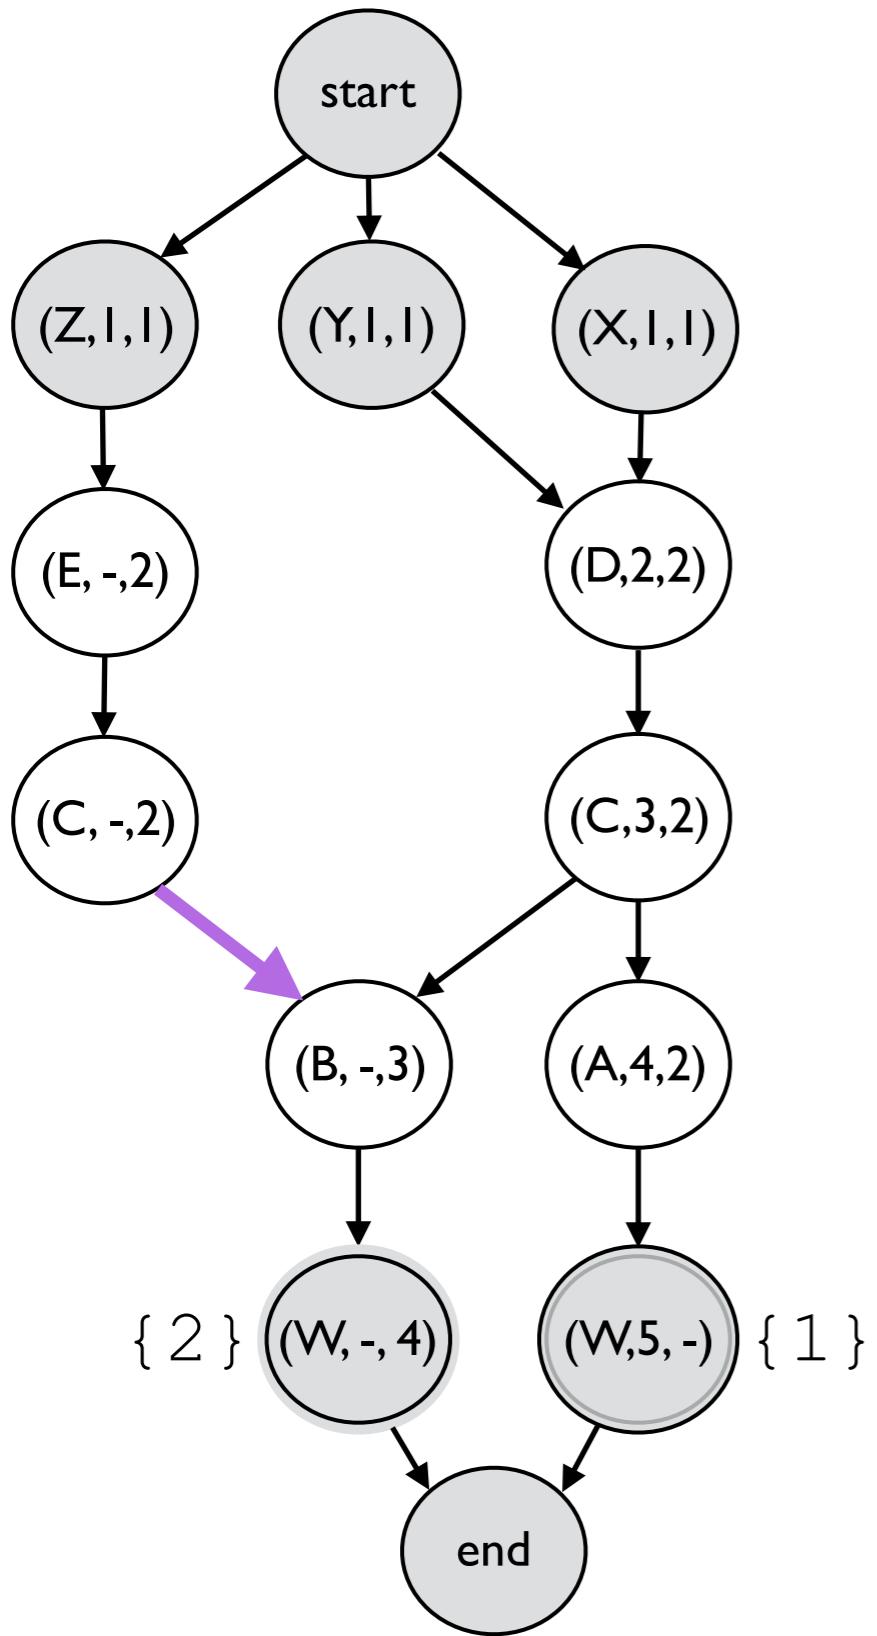
Policy: $(W.A.C.D.out) \gg (W.B.in+out)$

Compilation to BGP:



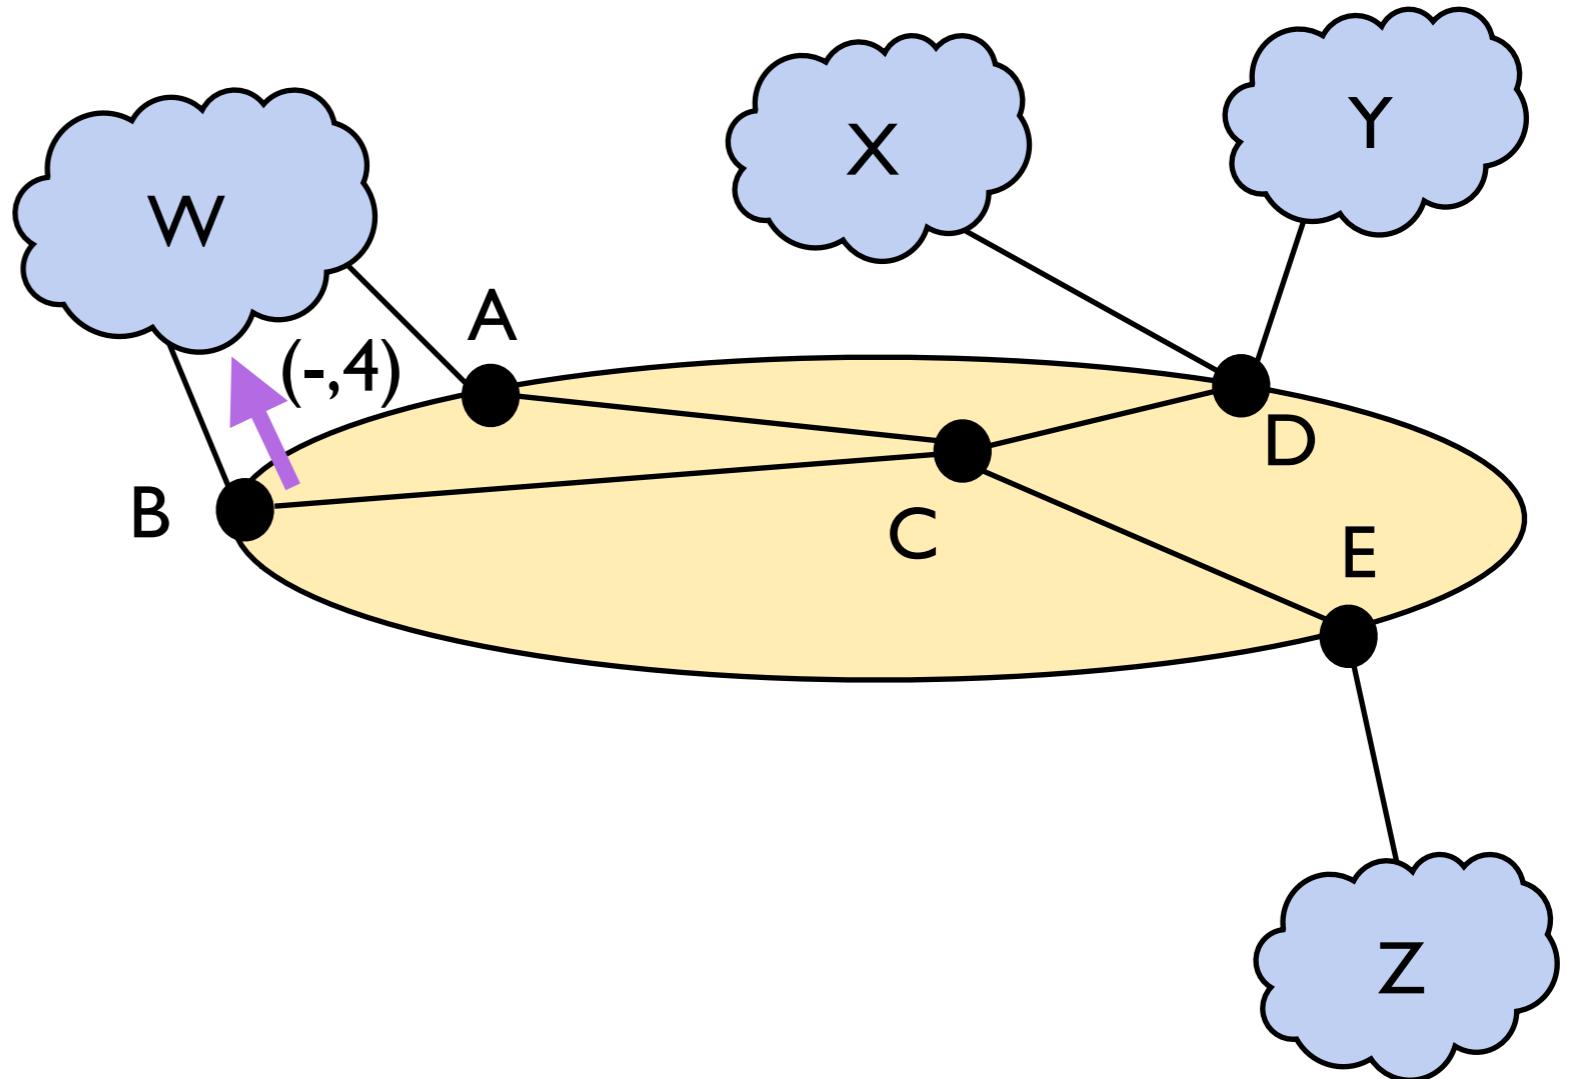
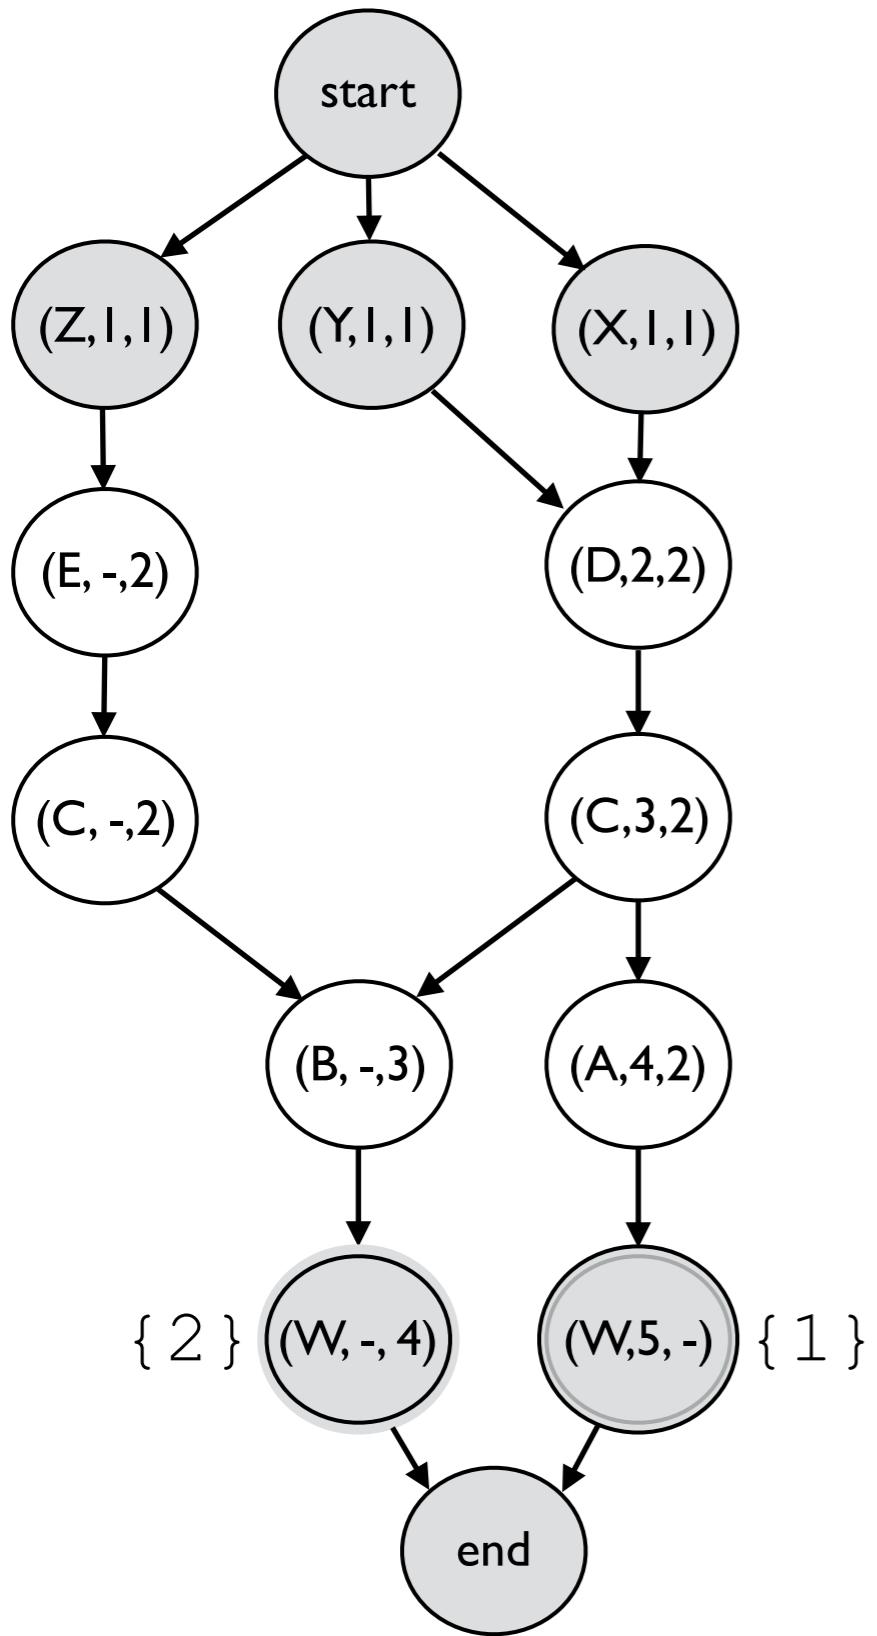
Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

Compilation to BGP:



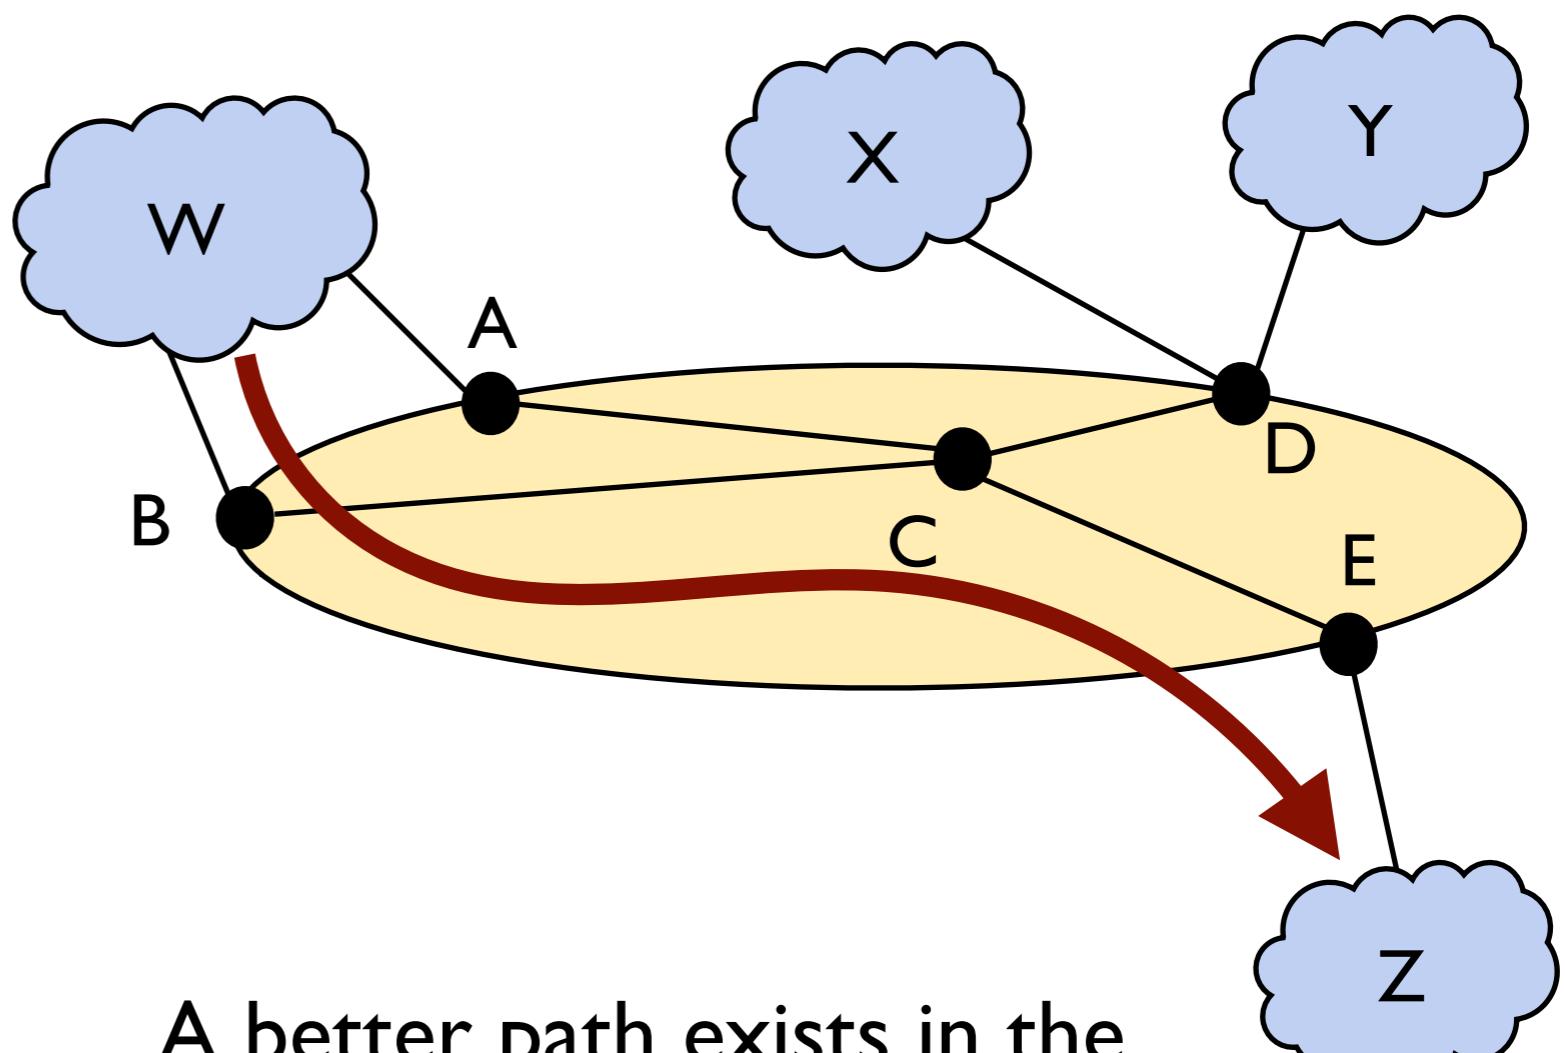
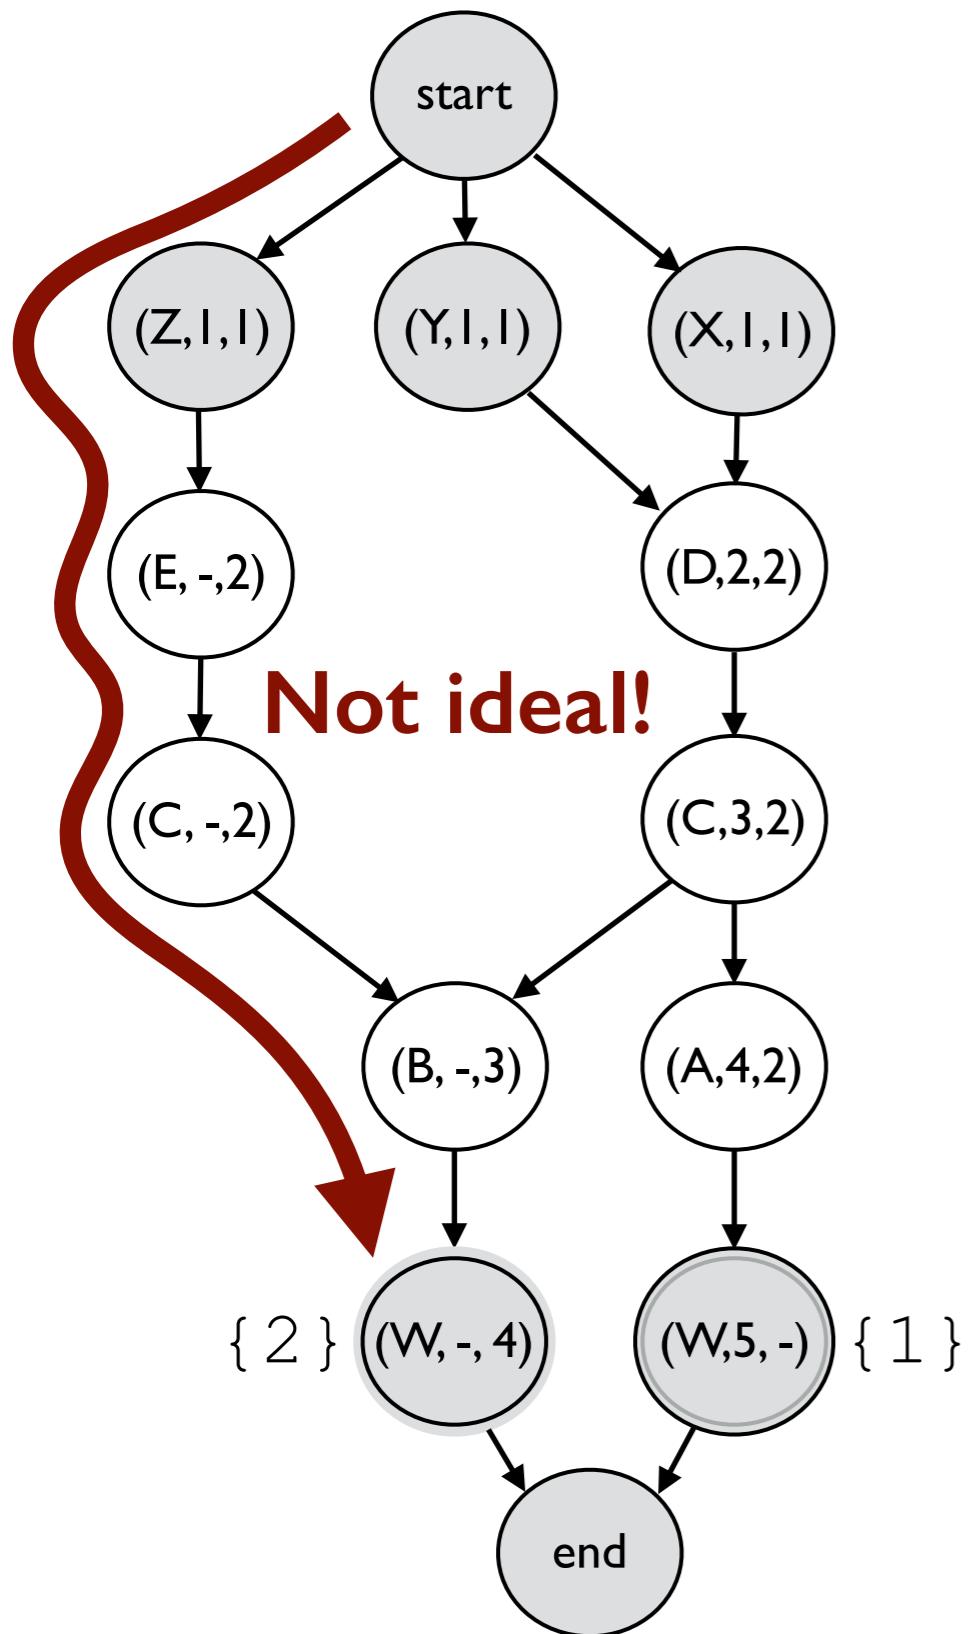
Policy: $(W.A.C.D.out) \gg (W.B.in+out)$

Compilation to BGP:



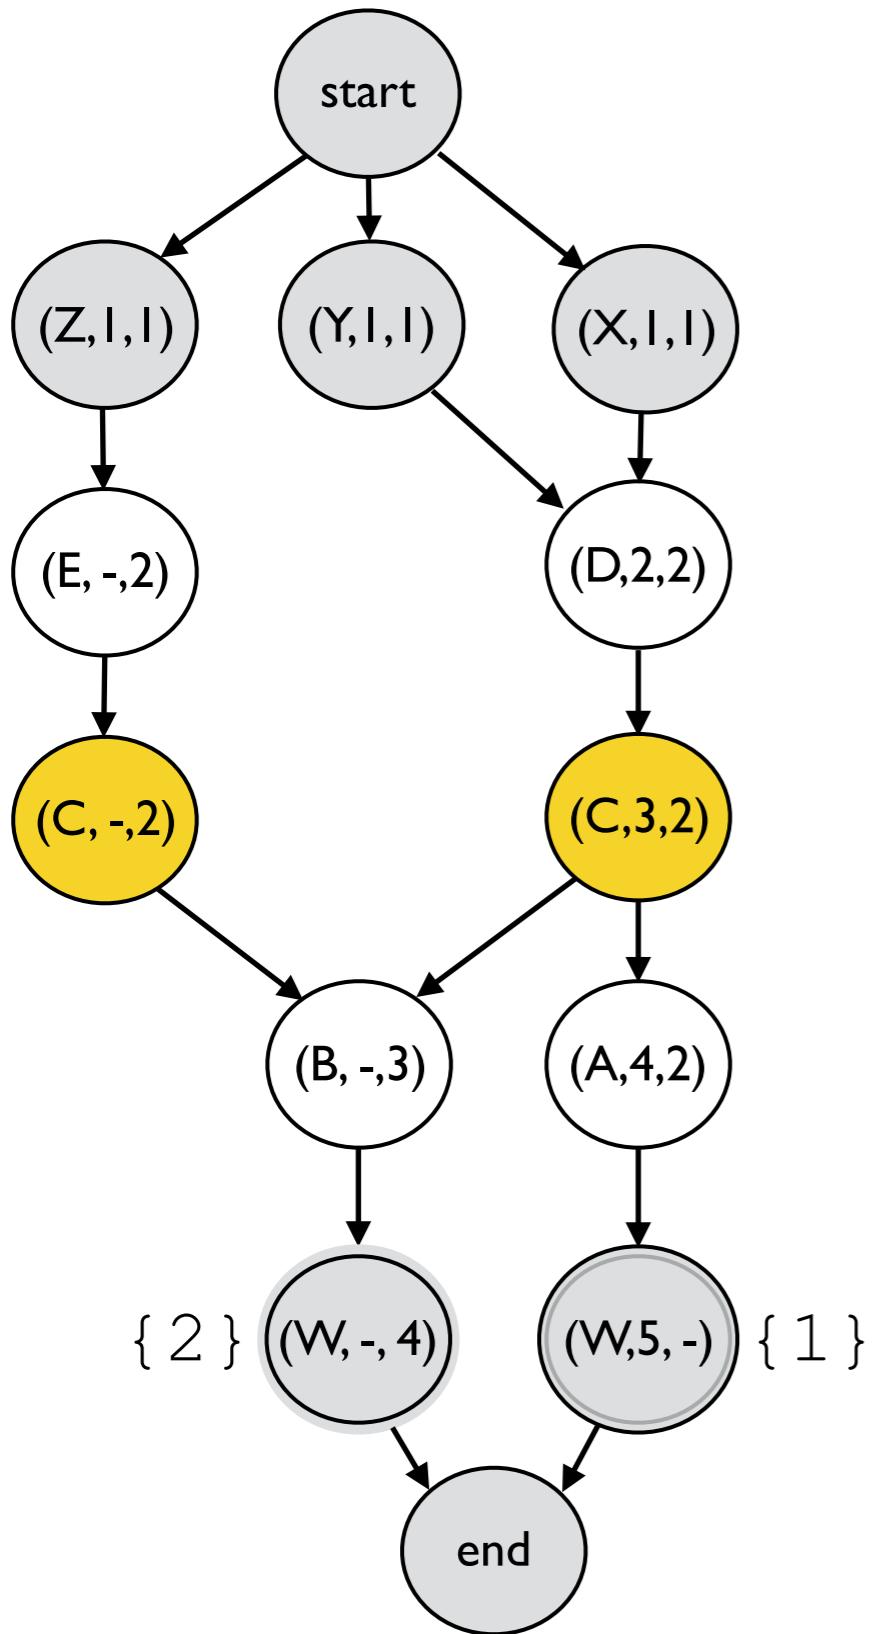
Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

Compilation to BGP:



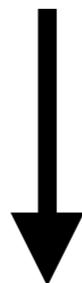
A better path exists in the network, but is not used!

Compilation to BGP:



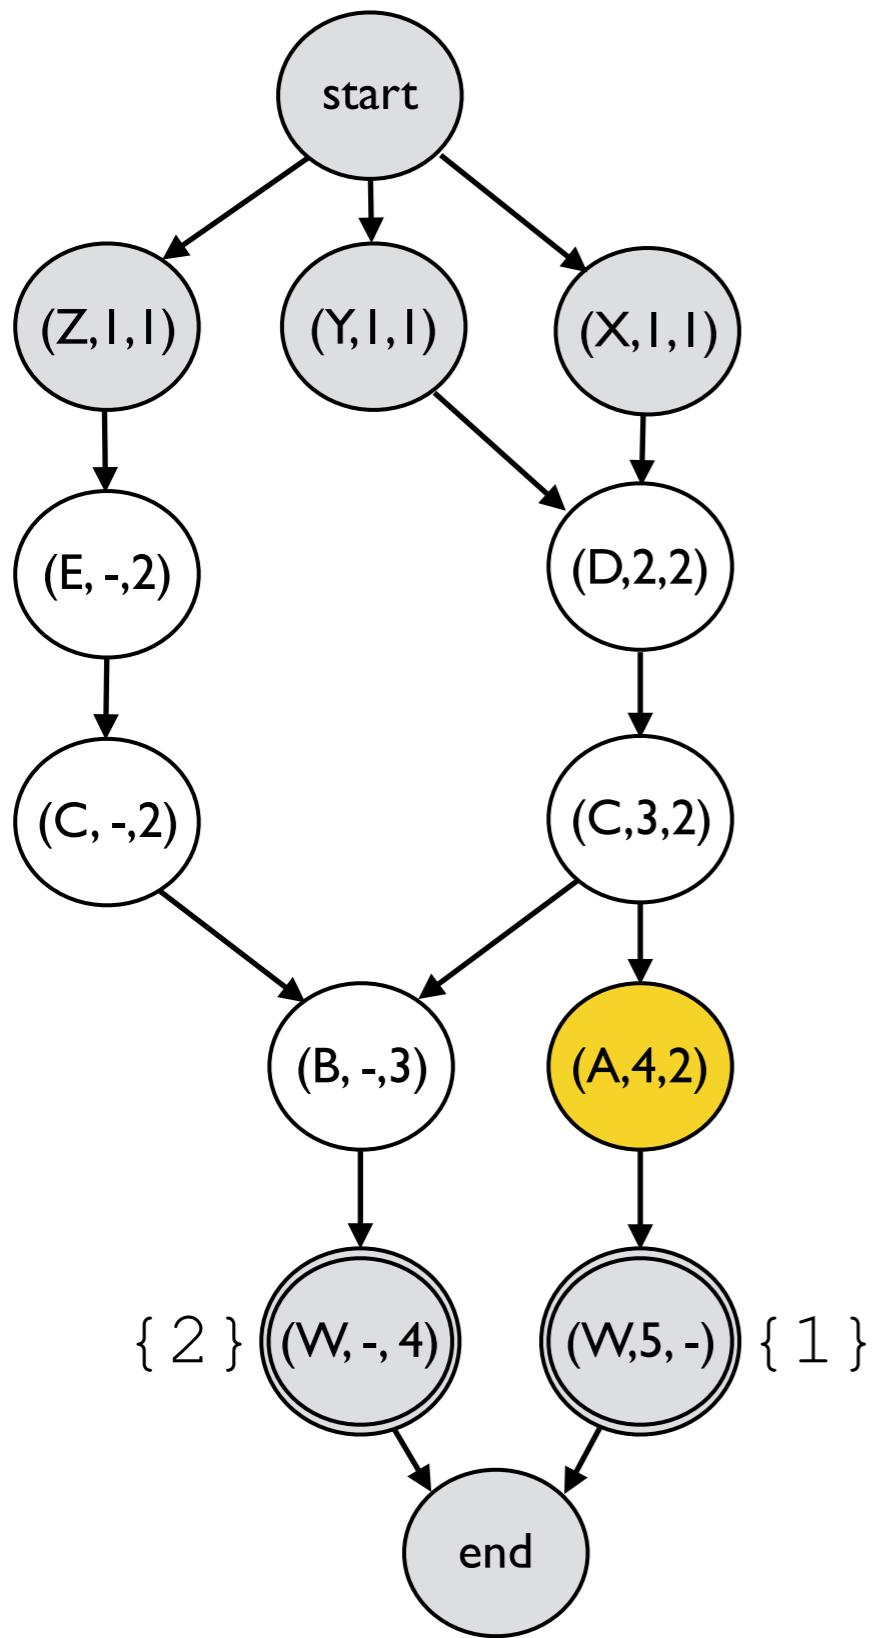
Idea 2: Find local preferences

- Direct BGP towards best path
- Under all combinations of failures
- Frame in terms of graph algorithms



Let BGP find **the best allowed** path dynamically

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

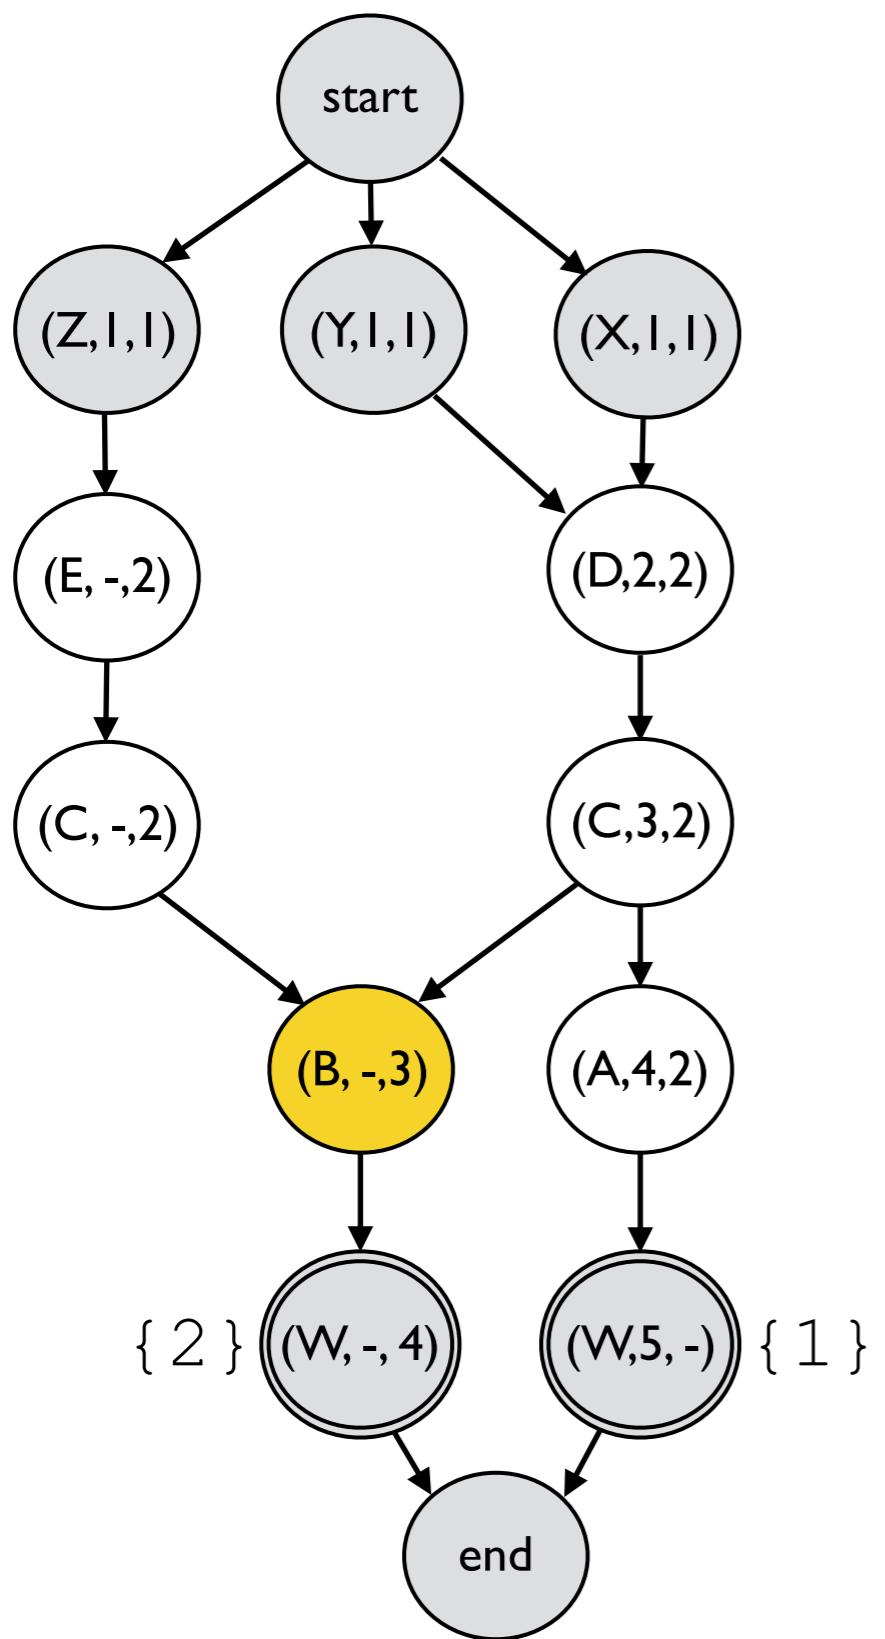
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```

match peer=C comm=(3,2)
export peer←W, comm←(4,2),
      comm←noexport, MED←80
  
```

Router B

```

match peer=C
export peer←W, comm←(-,3),
      comm←noexport, MED←81
  
```

Router C

```

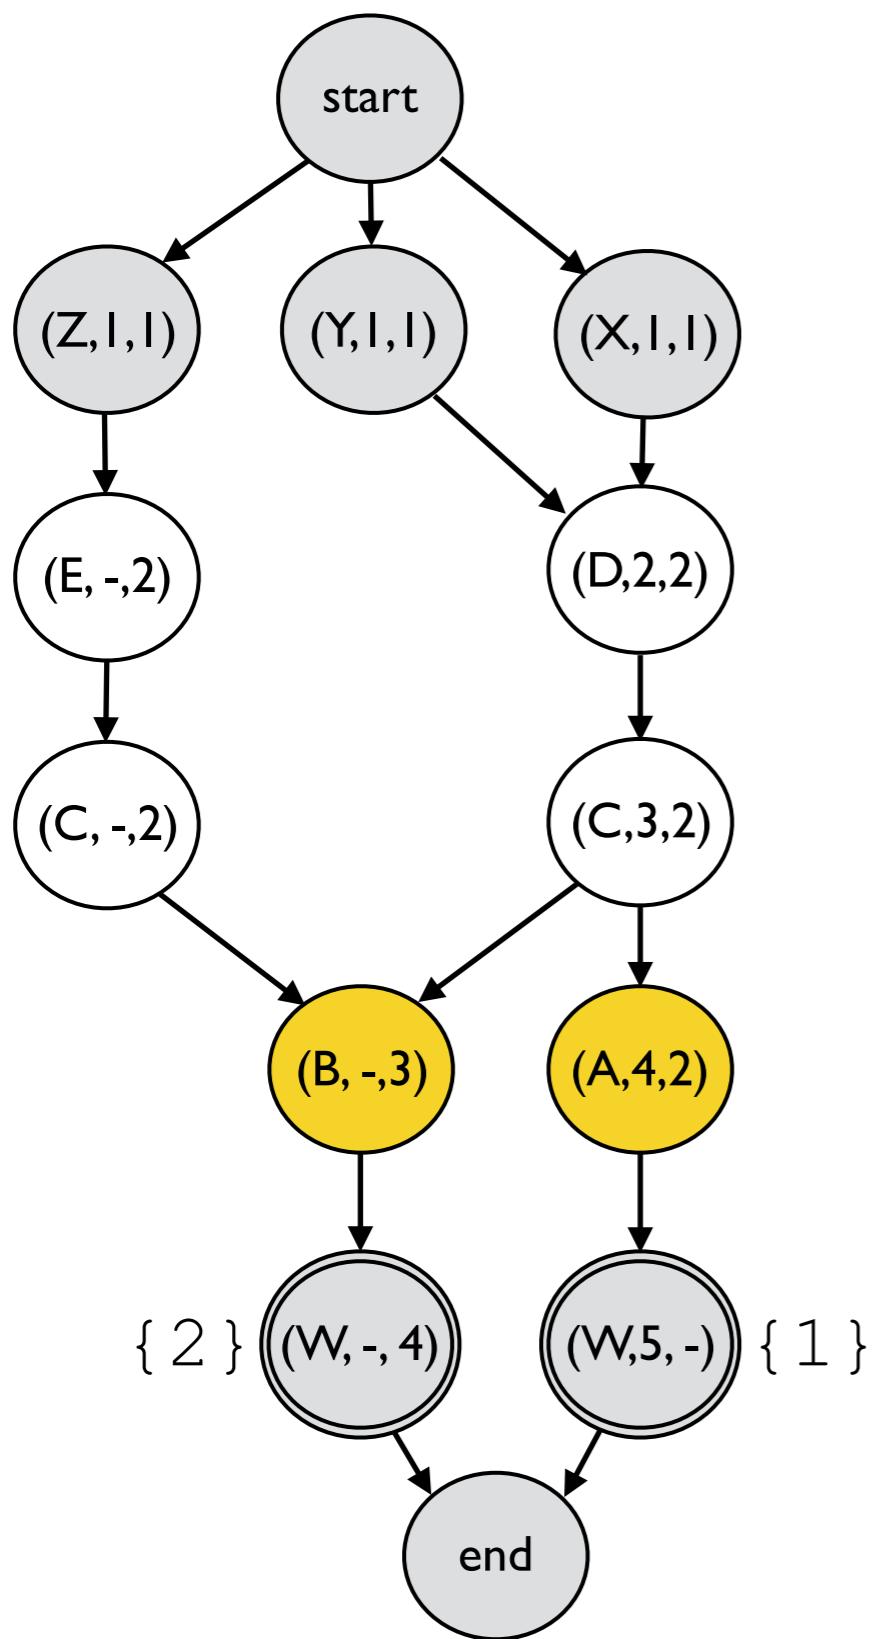
match[lp=99] peer=E, comm=(-,2)
export peer←B, comm←(-,2)
match[lp=100] peer=D, comm=(2,2)
export peer←A,B, comm←(3,2)
  
```

Router D

```

match regex=(X + Y)
export peer←C, comm←(2,2)
...
  
```

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

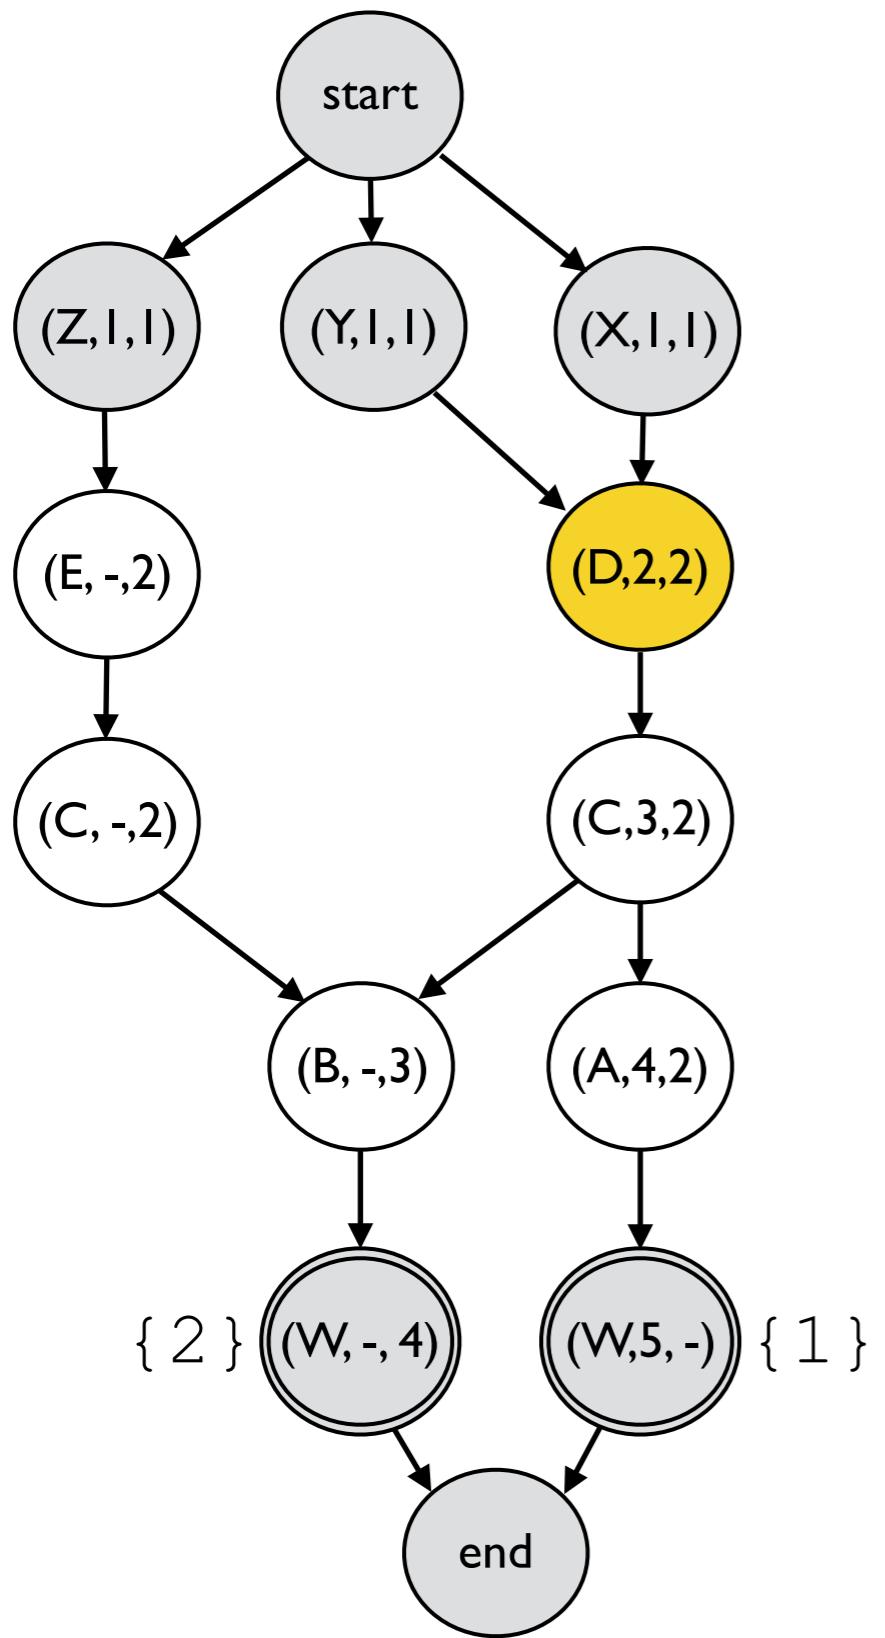
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

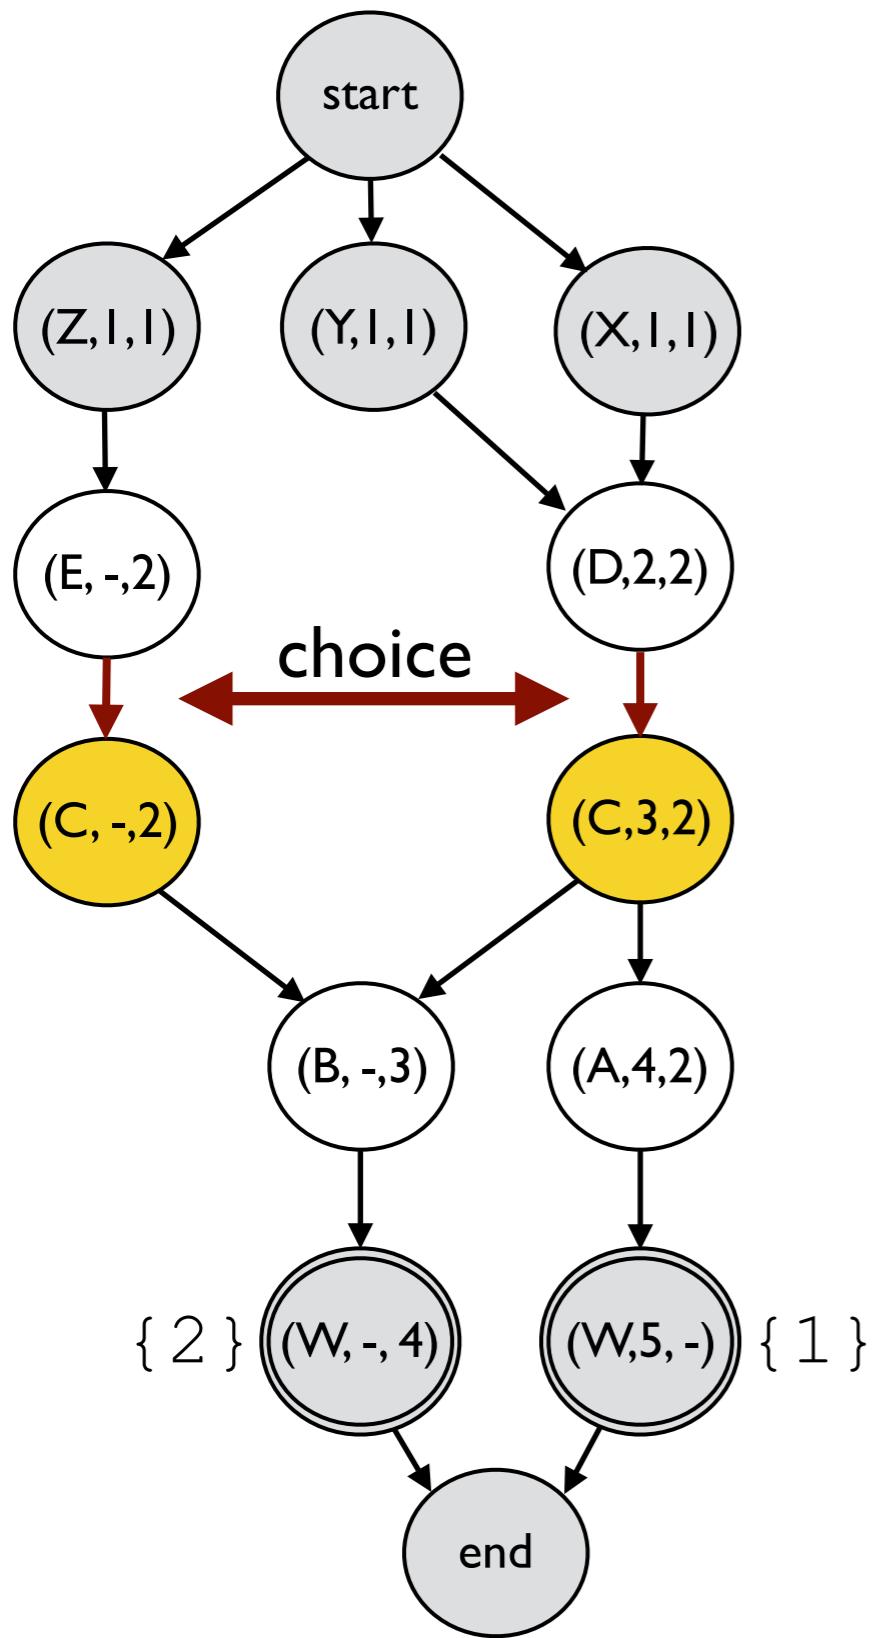
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

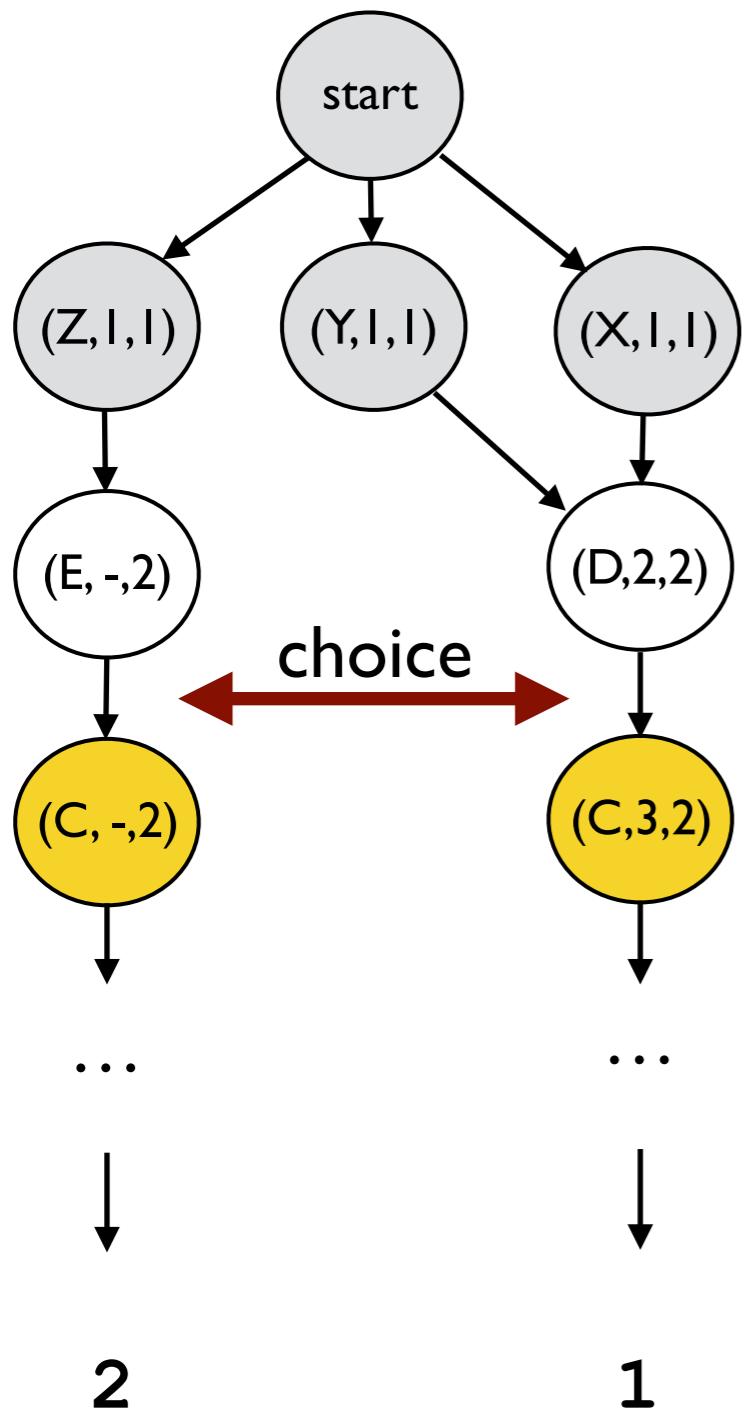
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
      comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
      comm←noexport, MED←81
```

Router C

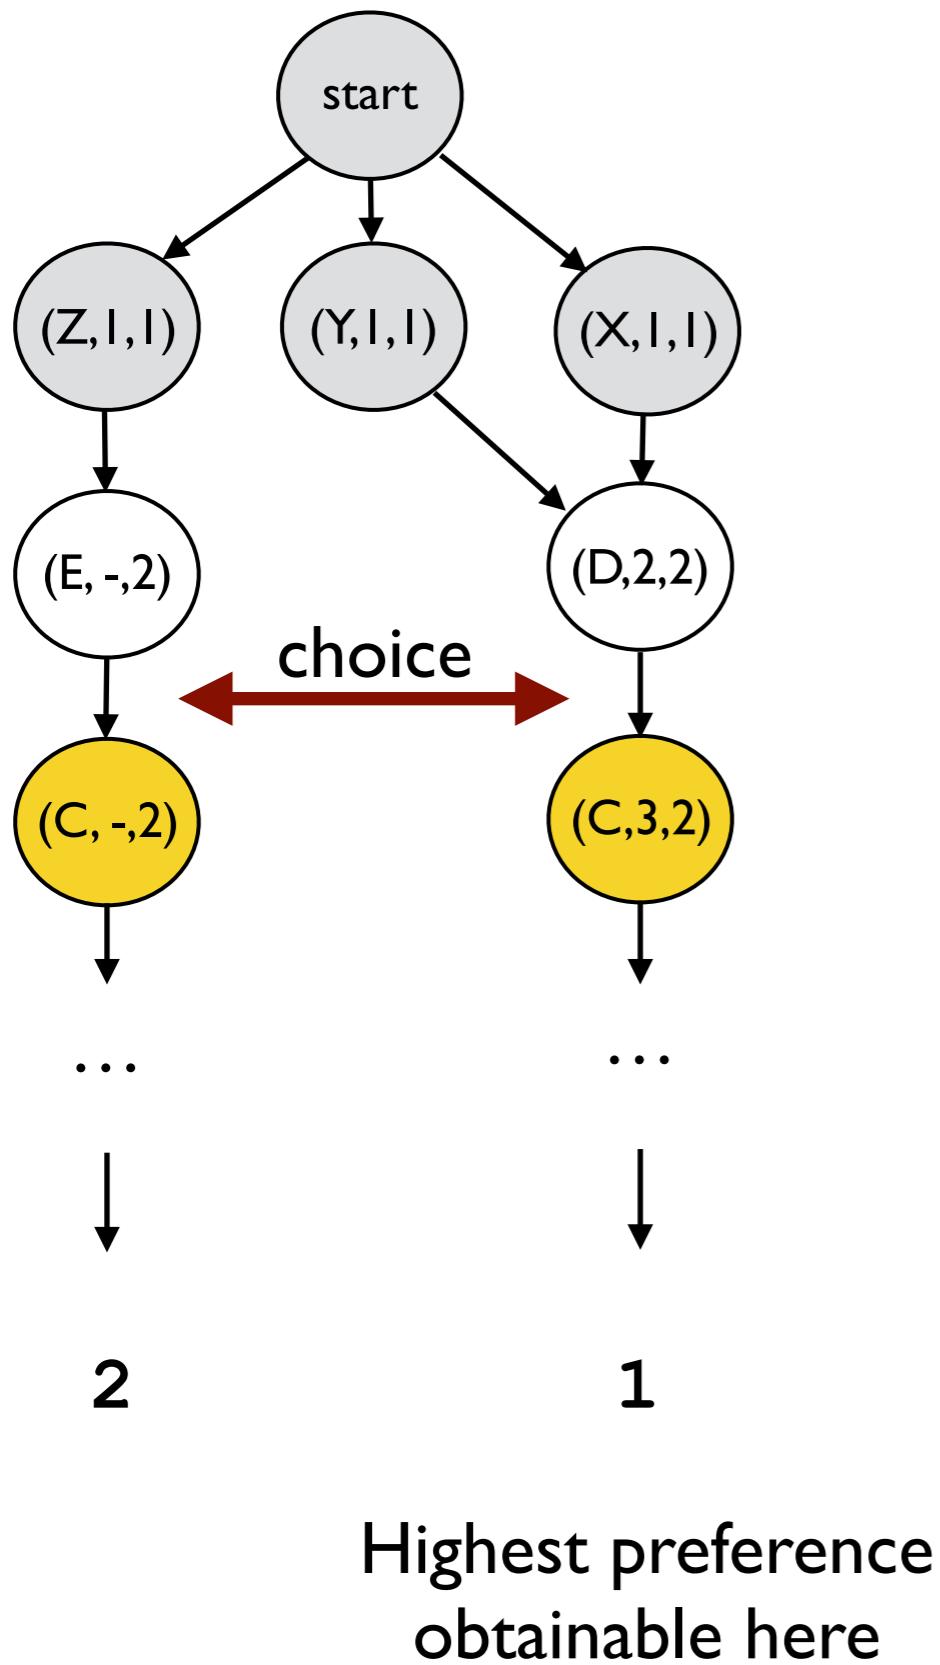
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-3),  
comm←noexport, MED←81
```

Router C

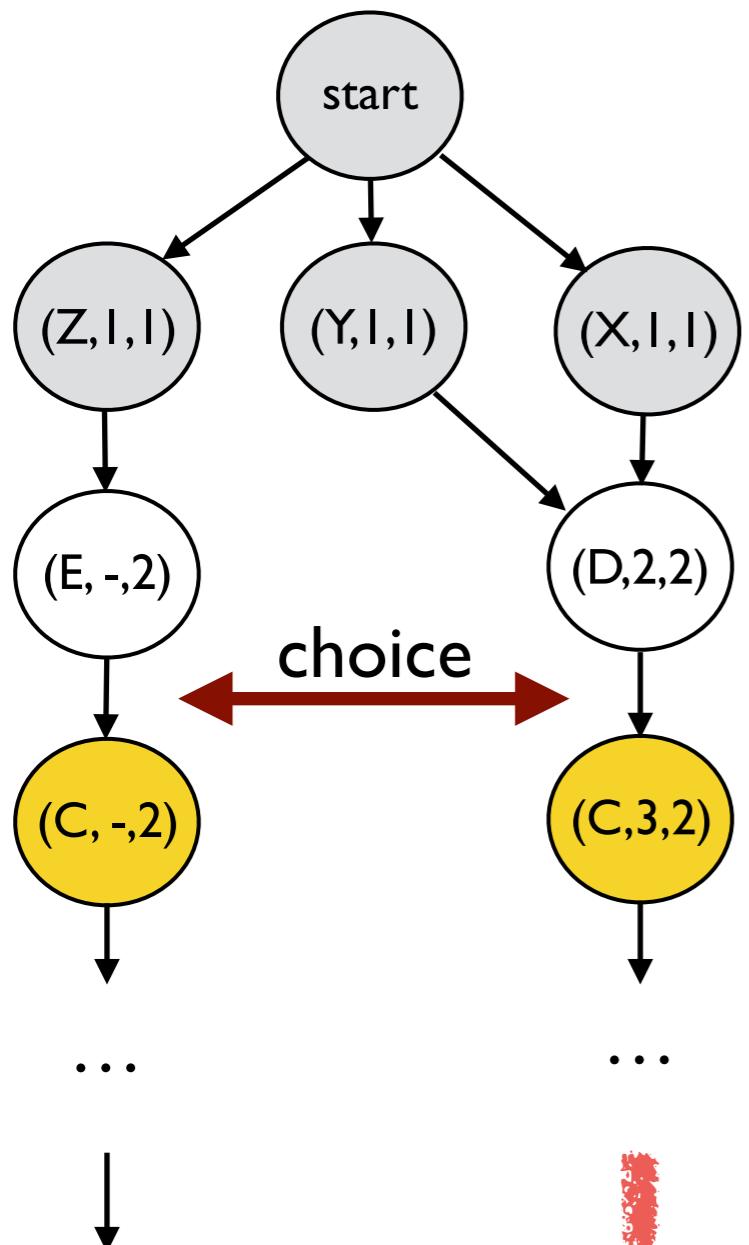
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



But there
could be a
failure!

Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

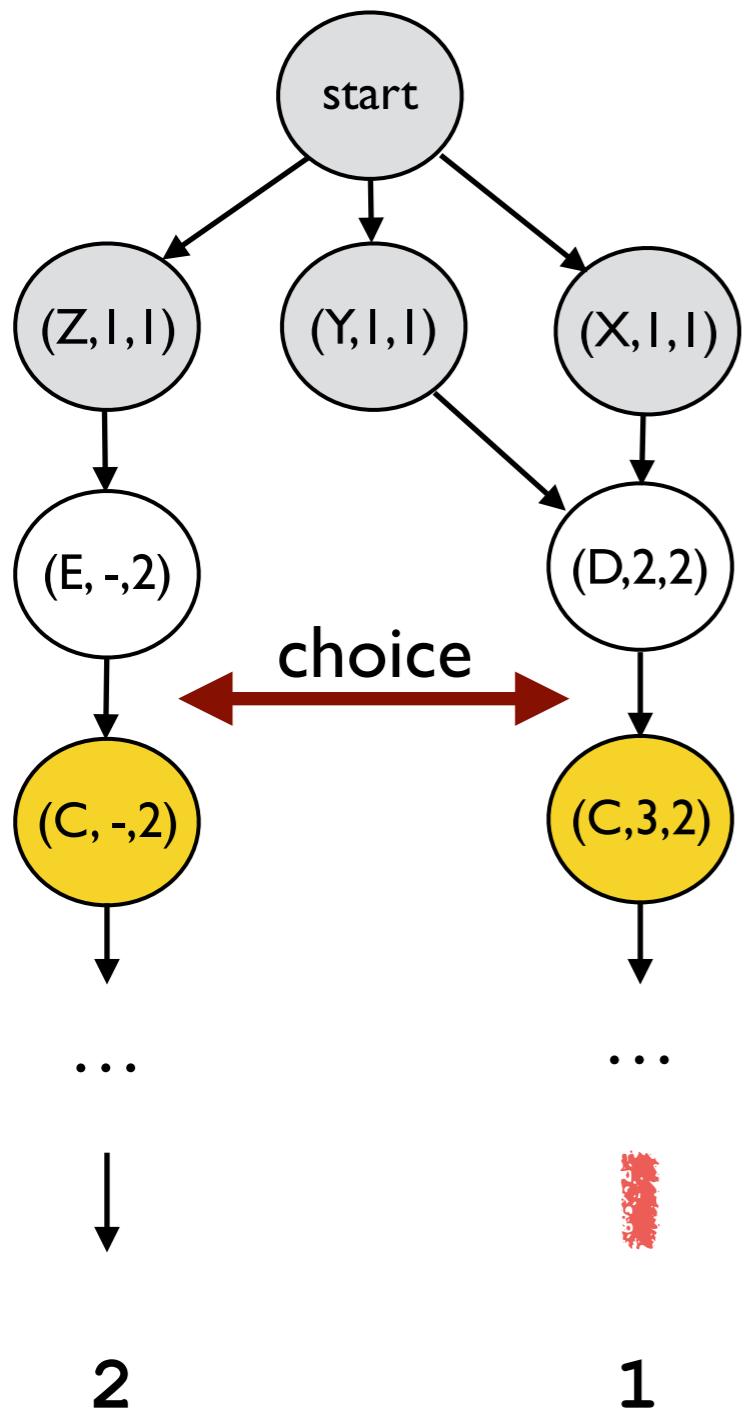
Router C

```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

Compilation to BGP:



*Sometimes there is
no best choice*

Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
      comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
      comm←noexport, MED←81
```

Router C

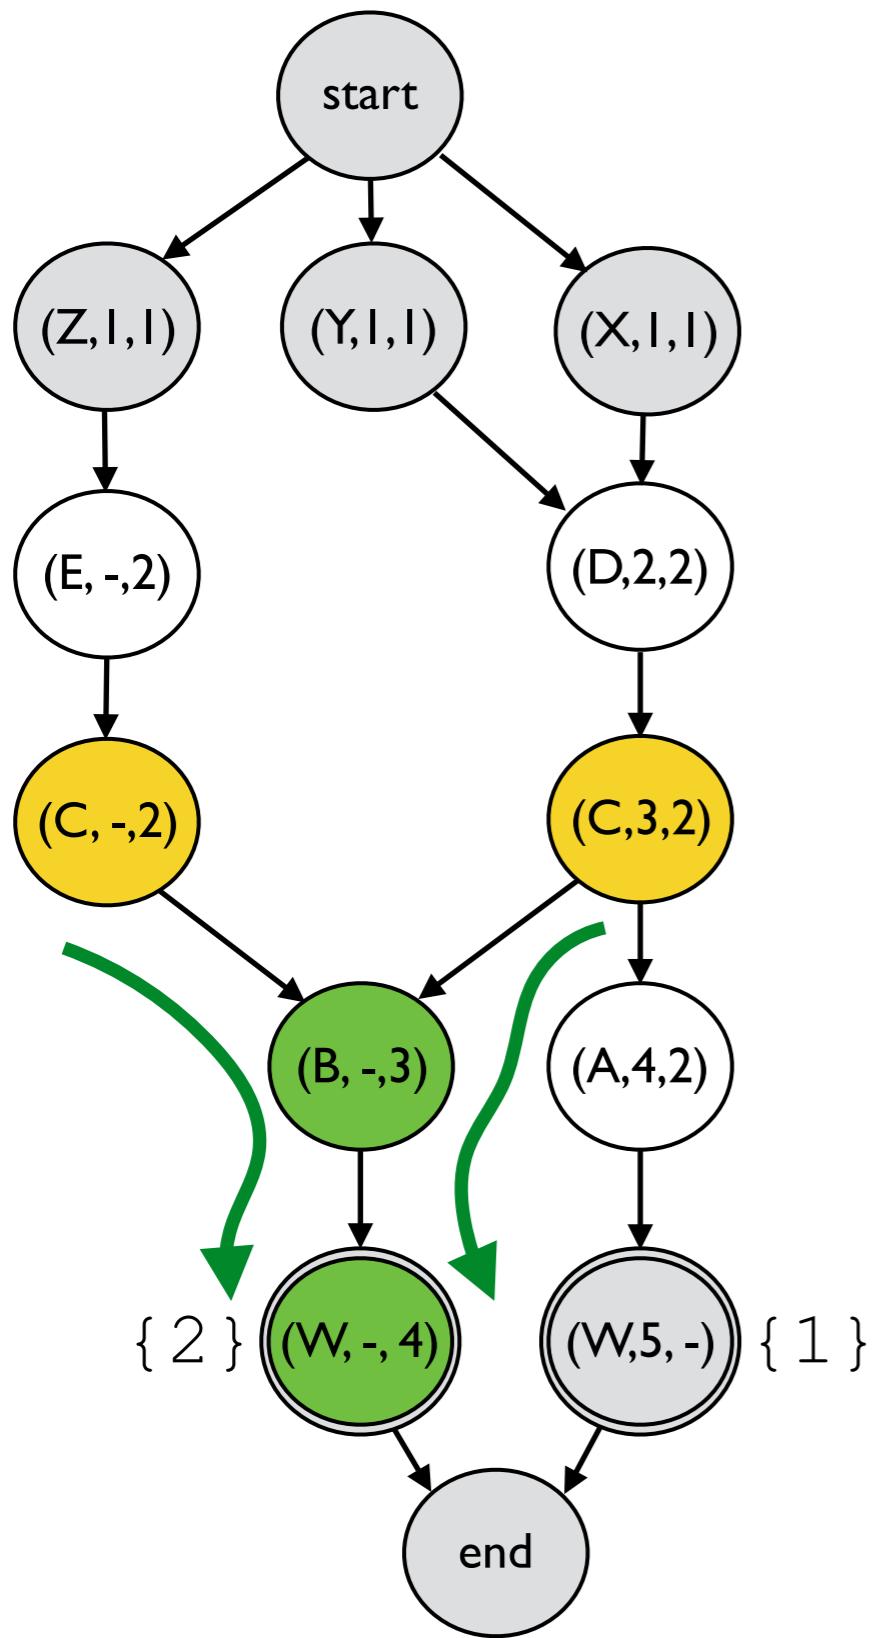
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

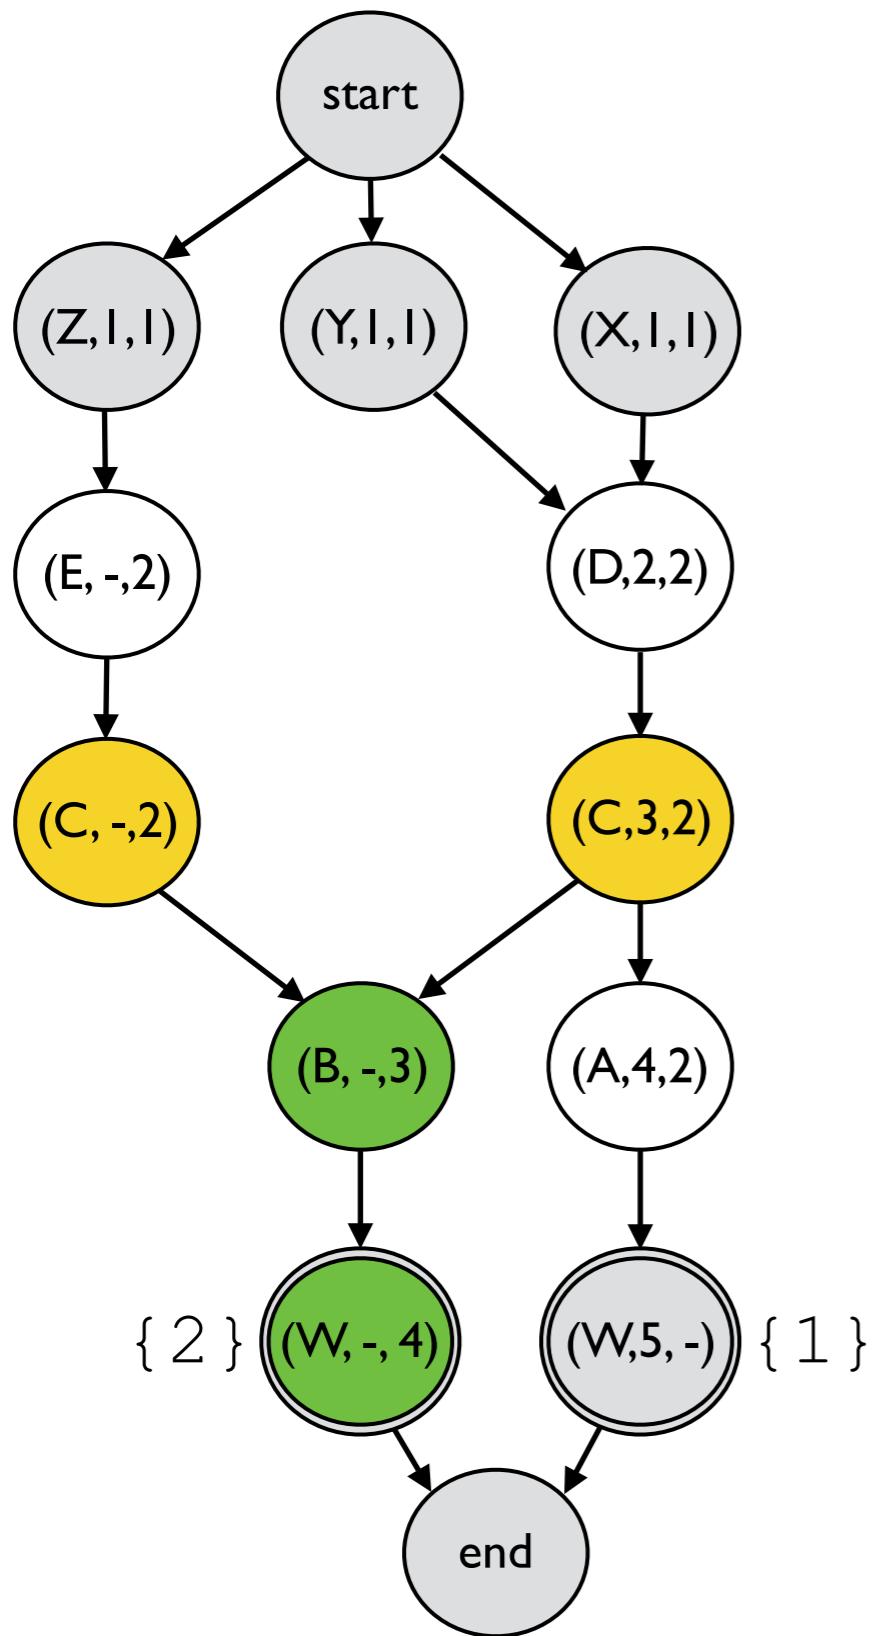
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```

match peer=C comm=(3,2)
export peer←W, comm←(4,2),
          comm←noexport, MED←80
  
```

Router B

```

match peer=C
export peer←W, comm←(-,3),
          comm←noexport, MED←81
  
```

Router C

```

match[lp=99] peer=E, comm=(-,2)
export peer←B, comm←(-,2)
match[lp=100] peer=D, comm=(2,2)
export peer←A,B, comm←(3,2)
  
```

Prefer D's announce

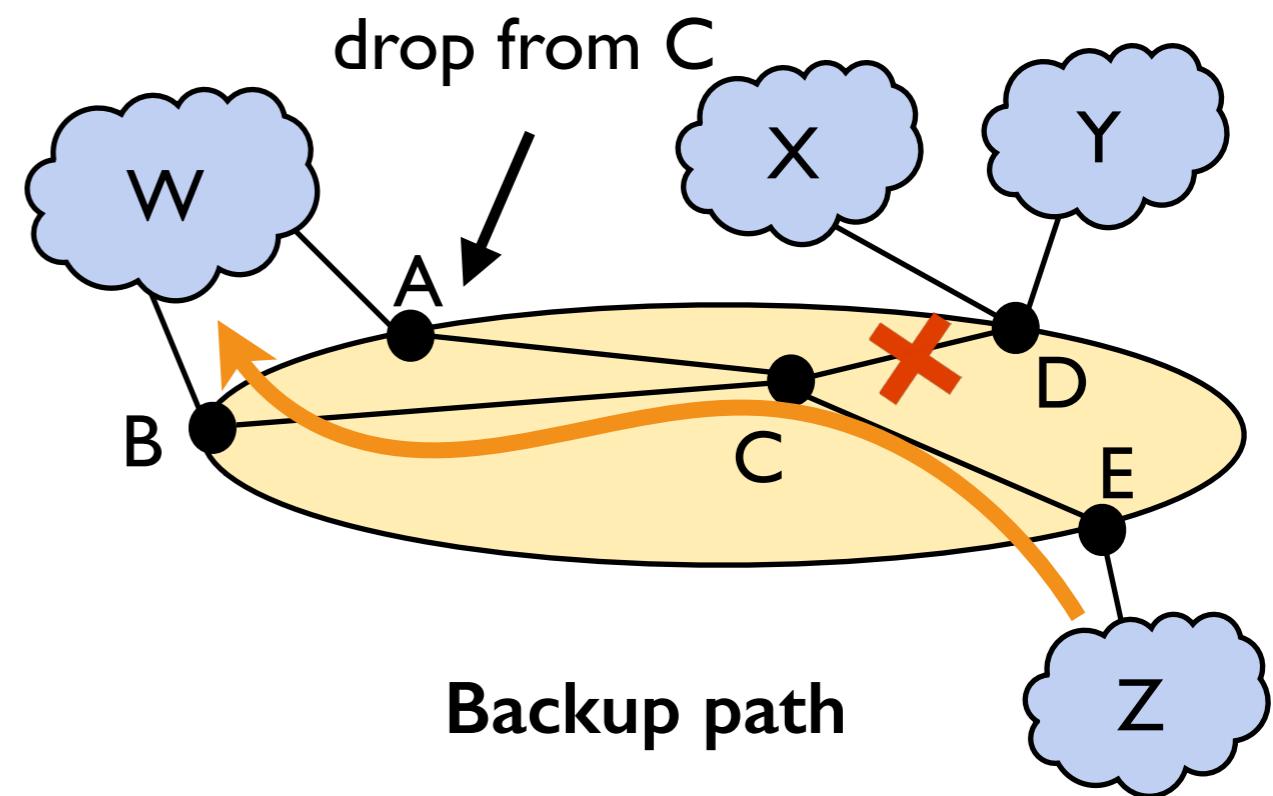
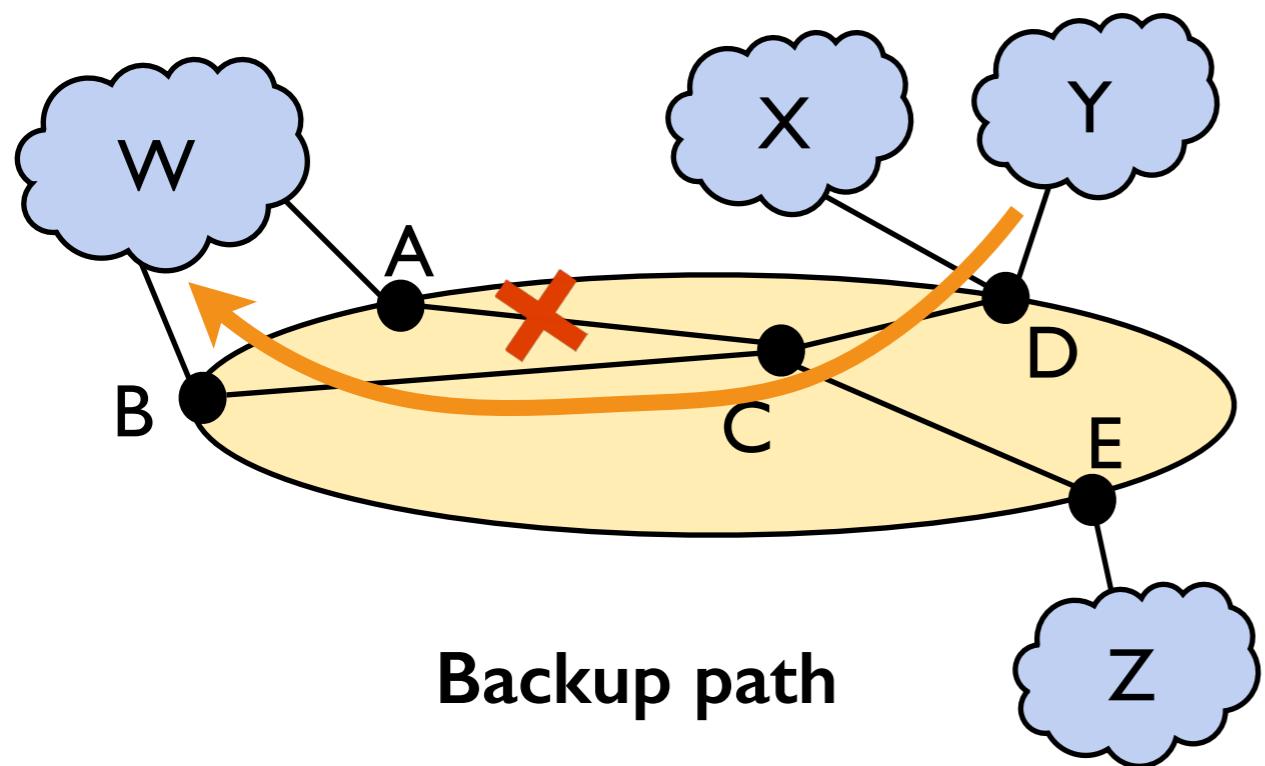
Router D

```

match regex=(X + Y)
export peer←C, comm←(2,2)
...
  
```

Compilation to BGP:

Policy: **enter(A) & exit(D) >> enter(B) & exit(out)**



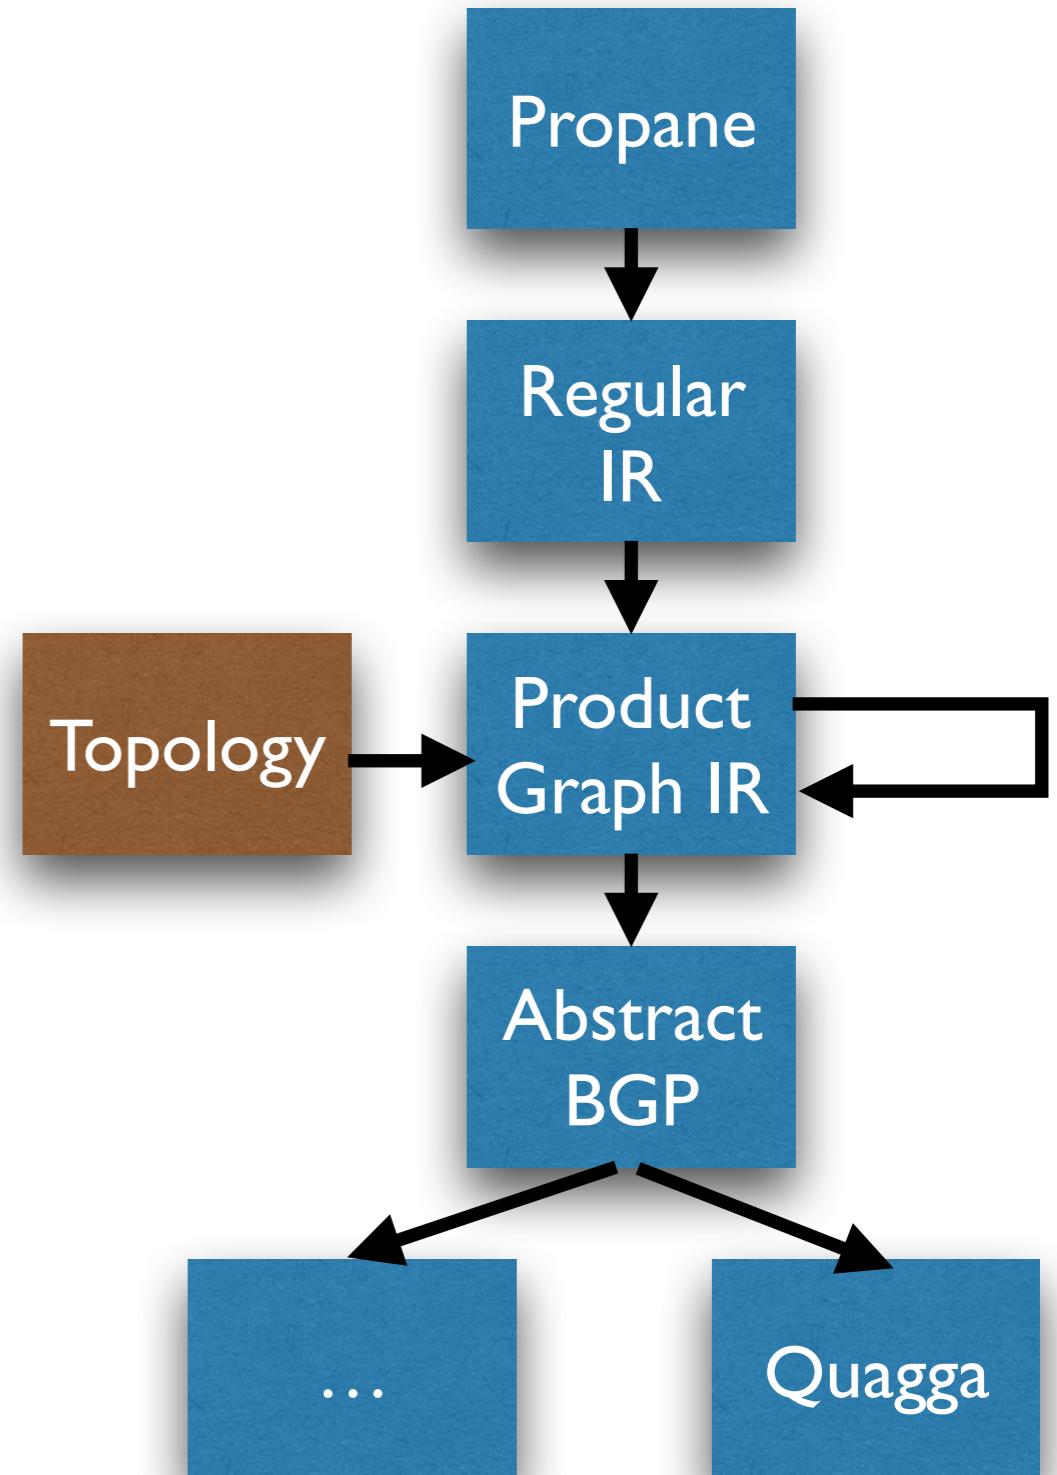
Implementation: C prefers D to E and tags accordingly

The implementation always uses the best available path

Implementation

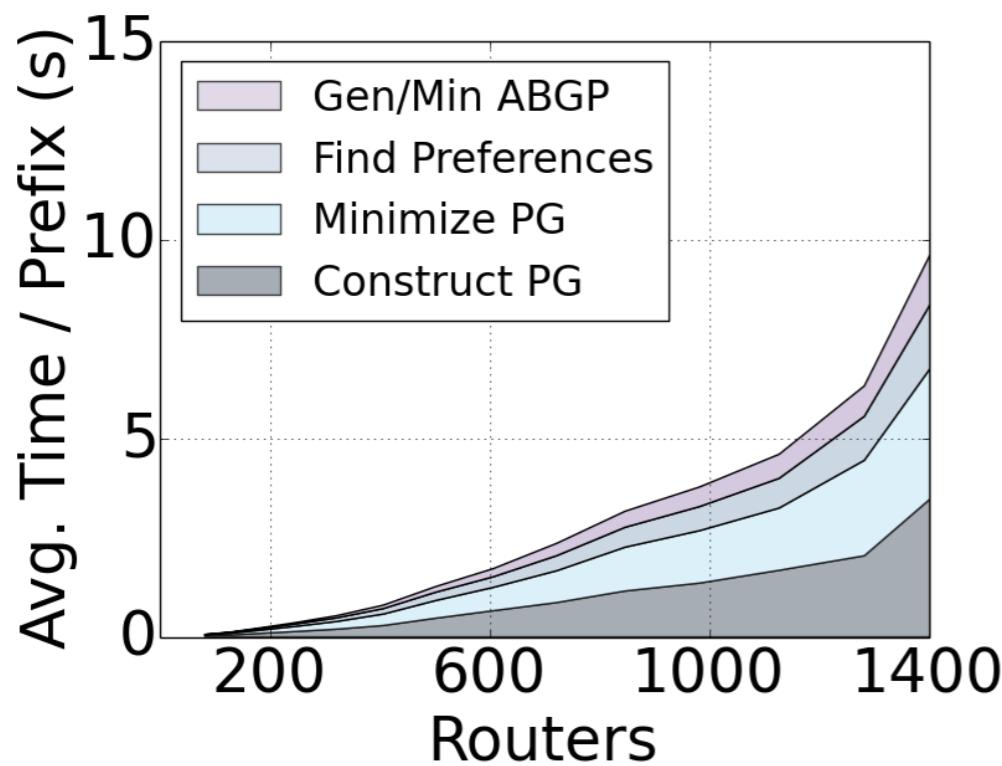
Benchmarks:

- Network configurations from Microsoft's networks (Policy from English docs)
- Data center policies (~1400 routers)
- Backbone policies (~200 routers, many peers/router)
- Ignoring prefix, customer group and ownership definitions:
 - 31 lines for data center
 - 43 lines for backbone

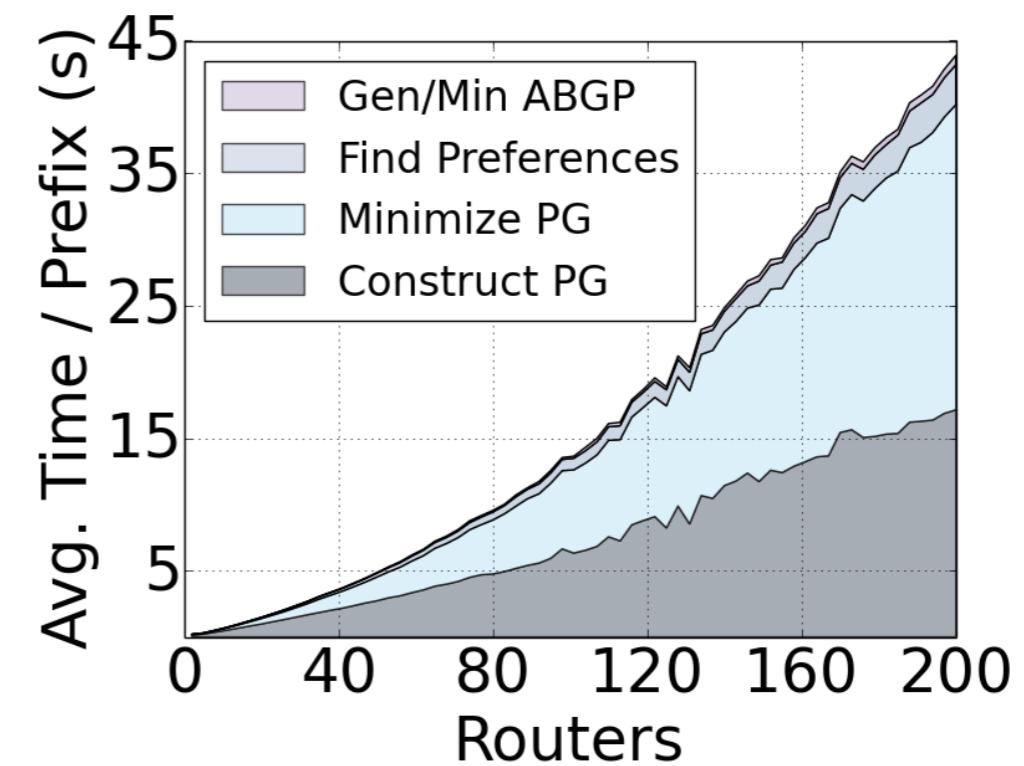


Compilation Time

- Parallelize compilation by prefix
- Use artificial topology to explore scalability



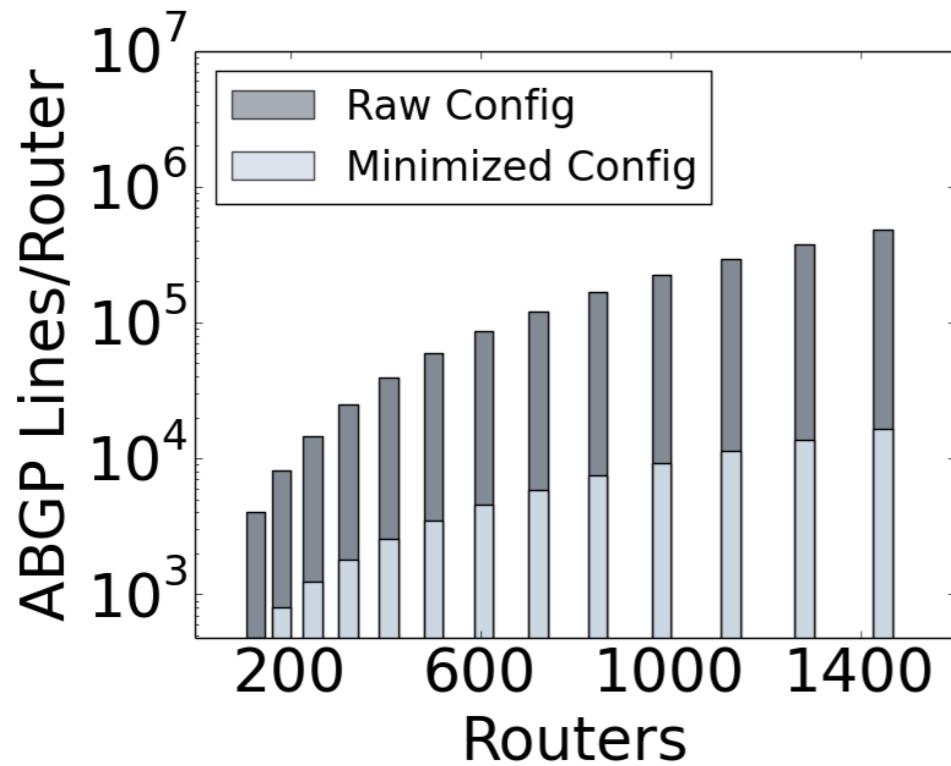
Data center (< 9 min)



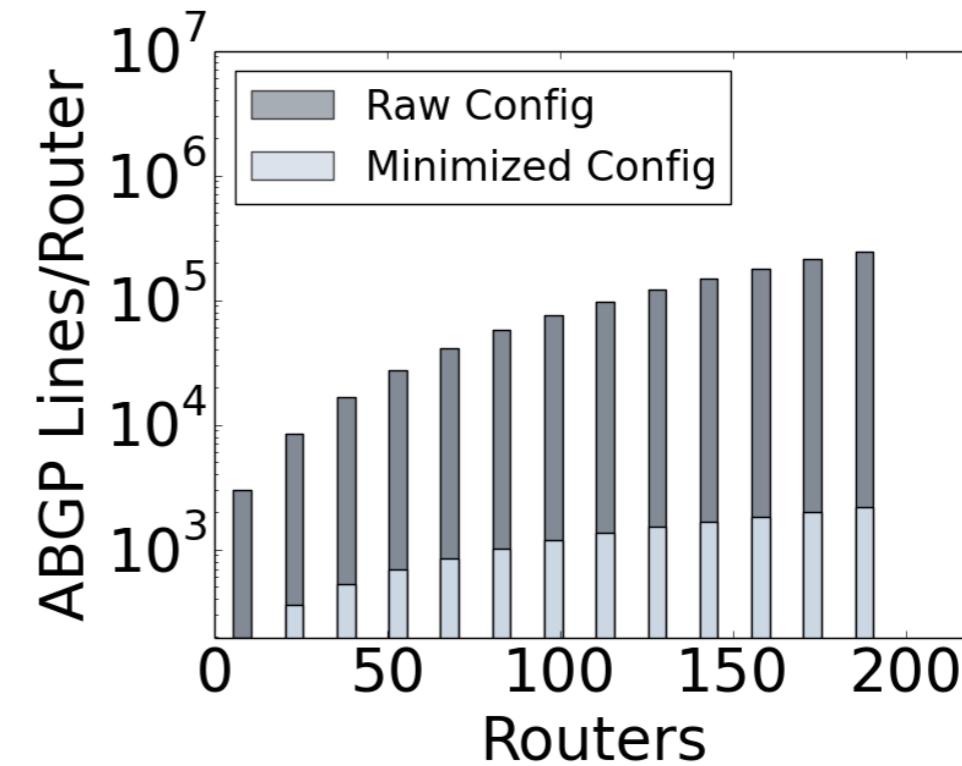
Backbone (< 3 min)

Configuration Size

- Perform configuration minimization during generation
- Avoid using community tags when choices are unambiguous
- Fall-through elimination of route maps



Data center



Backbone