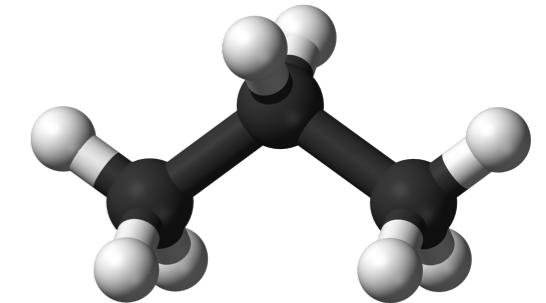
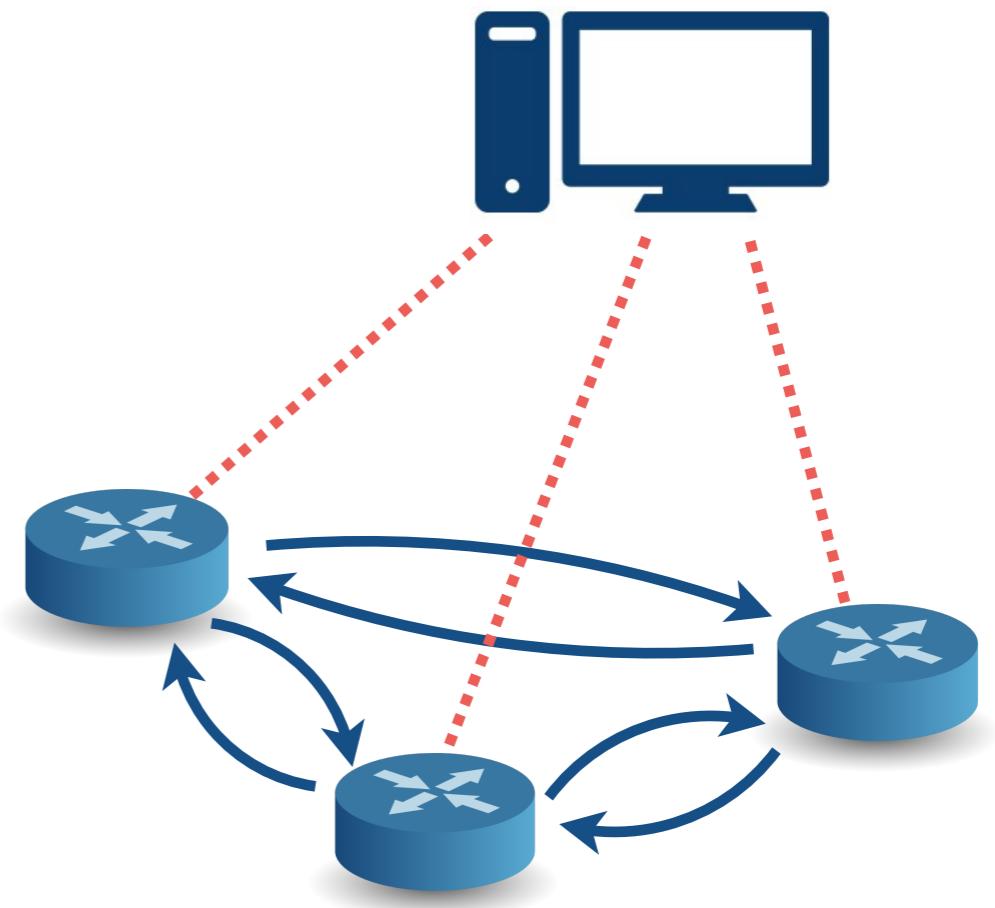


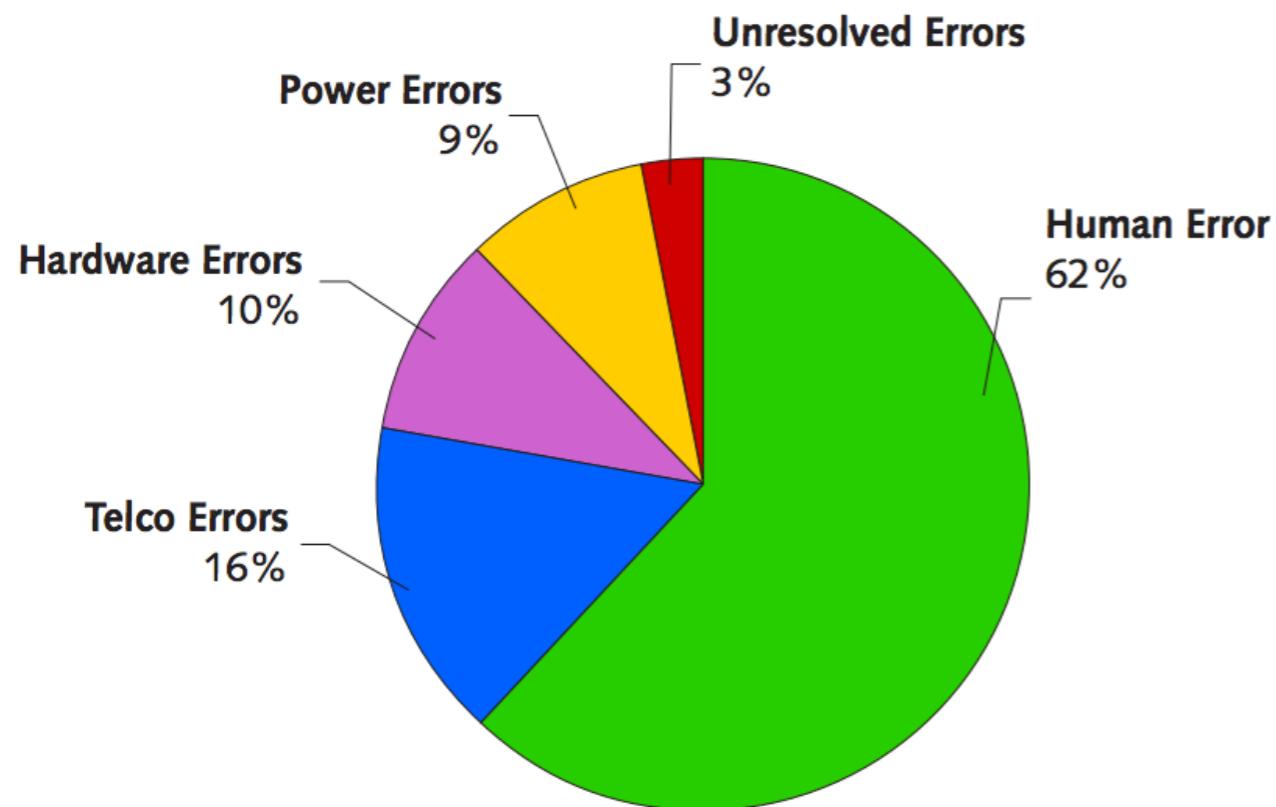
Propane: Programming Distributed Control Planes



Ryan Beckett (Princeton, MSR)
Ratul Mahajan (MSR)
Todd Millstein (UCLA)
Jitu Padhye (MSR)
David Walker (Princeton)

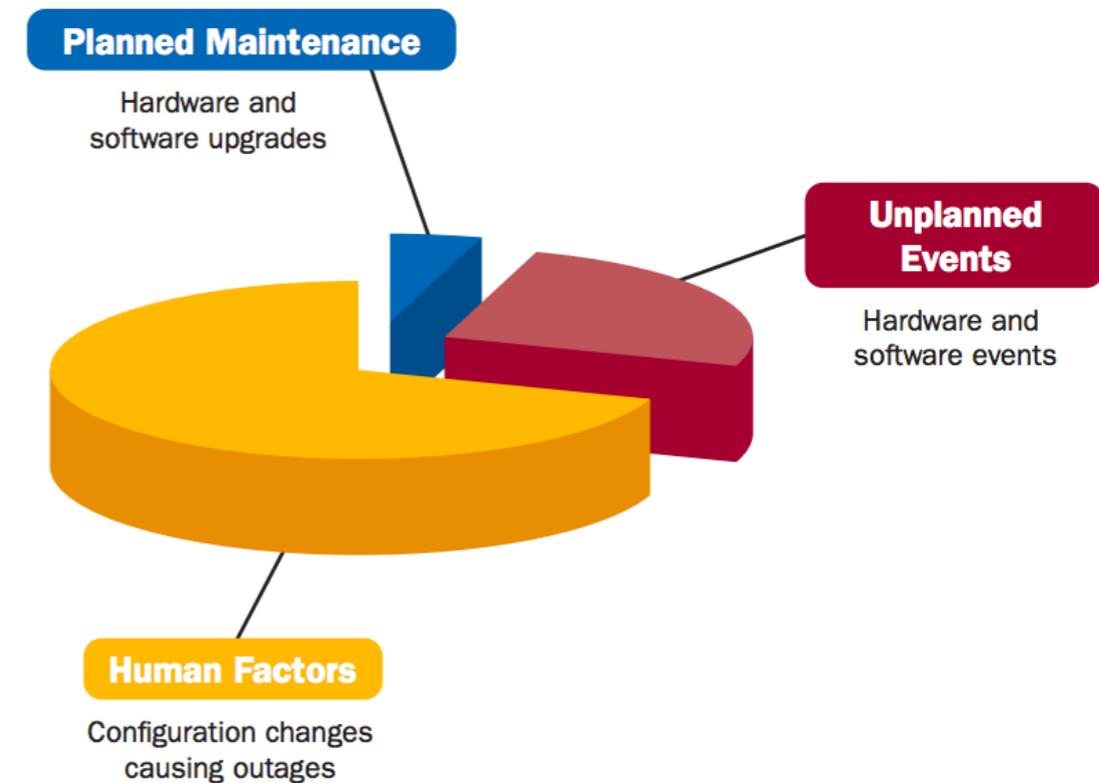


Configuring Networks is Error-Prone



**~60% of network downtime
is caused by human error**

-Yankee group 2002



**50-80% of outages are
the result of human error**

-Juniper 2008

Configuring Networks is Error-Prone

Understanding BGP Misconfiguration

Ratul Mahajan David Wetherall Tom Anderson

(ratul.djw.tom)@cs.washington.edu
Computer Science and Engineering
University of Washington
Seattle, WA 98195-2350

ABSTRACT
It is well-known that simple, accidental BGP configuration errors can disrupt Internet connectivity. Yet little is known about the frequency of misconfiguration or its causes, except for the few spectacular incidents of widespread outages. In this paper, we present the first quantitative study of BGP misconfiguration. Over a three week period, we analyzed routing table advertisements from 23 vantage points across the Internet backbone to detect incidents of misconfiguration. For each incident we polled the ISP operators involved to verify whether it was a misconfiguration, and to learn the cause of the incident. We also actively probed the Internet to determine the impact of misconfiguration on connectivity.

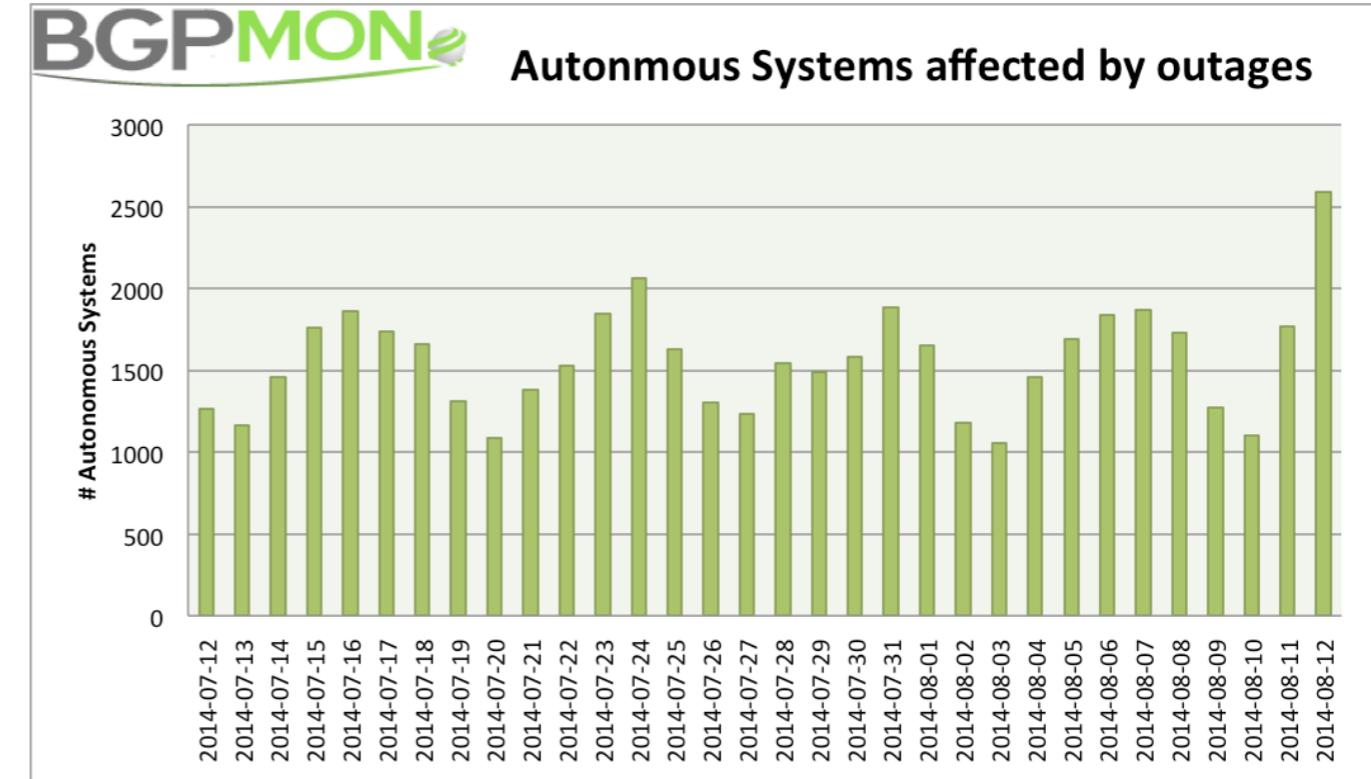
Surprisingly, we find that configuration errors are pervasive, with 200-1200 prefixes (0.2-1.0% of the BGP table size) suffering from misconfiguration each day. Close to 3 in 4 of all new prefix advertisements were results of misconfiguration. Fortunately, the connectivity seen by end users is surprisingly robust to misconfigurations. While misconfigurations can substantially increase the update load on routers, only one in twenty five affects connectivity. While the causes of misconfiguration are diverse, we argue that most could be prevented through better router design.

Categories and Subject Descriptors
C.2.3 [Communication Networks]: Operations—management;
C.4 [Computer Systems]: Performance—reliability, availability, and serviceability

General Terms
Human Factors, Management, Reliability

1. INTRODUCTION
As the Internet's inter-domain routing protocol, the Border Gateway Protocol (BGP) [34] is crucial to the overall reliability of the Internet. Faults in BGP implementations or mistakes in the way it is used have been known to disrupt large regions of the Internet. Recent studies have examined several kinds of BGP problems, including excessive churn due to implementation deficiencies [26].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.
SIGCOMM'02, August 19–23, 2002, Pittsburgh, Pennsylvania, USA.
Copyright 2002 ACM 1-58113-570-X/02/0008 ... \$5.00.



“Close to 3 in 4 of all new prefix advertisements were results of misconfiguration”

~1500 ASes affected by outages every day

Configuring Networks is Error-Prone

2/5/2016

China routing snafu briefly mangles interweb • The Register

Log in | Sign up

Cash'n'Carrion | Whitepaper | The Channel | The News Platform

GET YOUR
ONLY YOU CAN
REACH

2/5/2016

Internet-Wide Catastro

DATA CENTER SOFTWARE NETWORKS

Networks > Broadband

China routing snafu brief

Cockup, not conspiracy

9 Apr 2010 at 12:24, John Leyden

Bad routing information sourced from China has caused a temporary Internet-wide routing

Global BGP (Border Gateway Routing) lookup has been broken since about 10:00 UTC on Saturday. Telecommunication, apparently accidentally but intentionally, has been advertising itself as the best route to the Internet. Companies that peer with Time Warner experienced problems with their services such as supply chain portals (Figure 1).

Time Warner Outage

The alerts started coming in a little before 930 UTC (04:30 UTC) on Saturday, with many users unable to access websites, DNS names failing to resolve, and some parts of the Internet. Companies that peer with Time Warner experienced problems with their services such as supply chain portals (Figure 1).

24h 7d 14d

Figure 1: Supply chain portal with limited availability.

I could see right away that users and networks that connected to the supply chain portal, indicating the issue was in the Time Warner network. The service interruption while all the traffic re-routed through AT&T was a temporary one. Availability issues continued for the entire duration of the outage. I took a look at the path visualization view to figure out exactly what was going on. Normally, two locations (Tokyo and Dallas) transit Road Runner (Time Warner) to reach this supply chain port. While the rest go through AT&T (Figure 2).

2/5/2016

Internet-Wide Catastro

DATA CENTER SOFTWARE NETWORKS

Networks > Broadband

Dyn Research

THE NEW HOME OF .REN

HOME TOPICS PRESENTATION

DECEMBER 24, 2005 COMMENTS (0)

ENGINEERING TODD UNDERWOOD

Internet-Wide Catastrophe—Last

Given the large number of prefixes and short life spans of hijack. Most likely it's because of configuration errors or misconfigurations.

The practical consequences of the screw-up are not clear. It could result in dropped connections or, worse, traffic routed through a network that is not the best. One of the clearest illustrations of the security risks of BGP is the Pakistan Telecom/YouTube incident. As far as anyone knows, it was a mistake, not a malicious attack. The consequences were far from benign: for sever

« Previous Story

able to reach a large number of sites can take a look at what happened during the intervening time.

morning 2004, TTNet (AS9121) started announcing

Sign In | Register

≡

NETWORKWORLD

YouTube/Pakistan incident: Could something similar whack your site?

Configuring BGP properly is key to avoidance, 'Net registry official says

By Carolyn Duffy Marsan

Network World | Mar 10, 2008 1:00 AM PT

In light of Pakistan Telecom/YouTube incident, Internet registry official explains how you can avoid having your web site victimized by such an attack.

When Pakistan Telecom blocked YouTube's traffic one Sunday evening in February, the ISP created an international incident that wreaked havoc on the popular video site for more than two hours.

RIPE NCC, the European registry for Internet addresses, has conducted an analysis of what happened during Pakistan Telecom's hijacking of YouTube's traffic and the steps that YouTube took to stop the attack.

We posed some questions to RIPE NCC's Chief Scientist Daniel Karrenberg about the YouTube incident. Here's what he had to say:

How frequently do hijacking incidents like the Pakistan Telecom/YouTube incident happen?

Misconfigurations of iBGP (internal BGP, the protocol used between the routers in the same Autonomous System) happen regularly and are usually the result of an error. One such misconfiguration caused the Pakistan Telecom/YouTube incident. It appears that the Pakistan Telecom/YouTube incident was not an "attack" as some have labeled it, but a configuration error. (See Columnist Johnna Till Johnson's take on the topic.)

What is significant about the YouTube incident?



Objectives: Network-wide

- Prefer traffic to go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Adhere to policies even when **failures** occur



Objectives: Network-wide

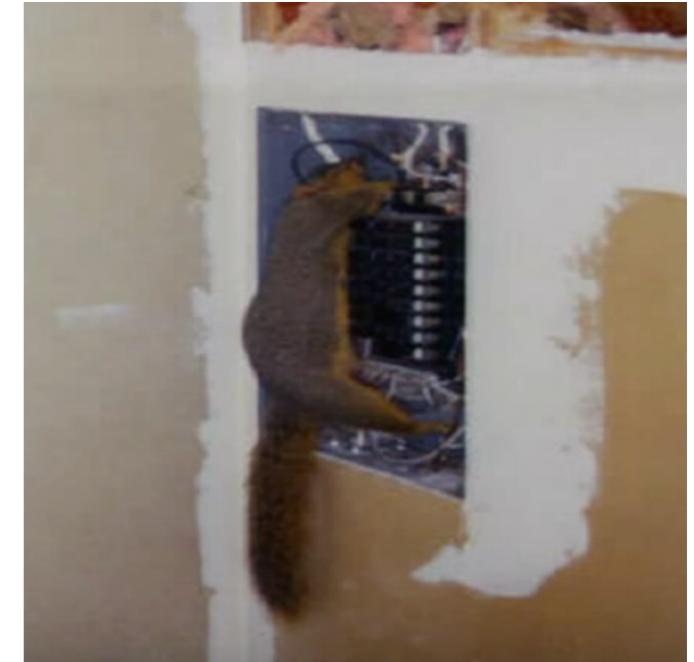
- Prefer traffic to go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Adhere to policies even when **failures** occur





Objectives: Network-wide

- Prefer traffic to go through AT&T over Sprint
- Don't use our network as transit between A and B
- Traffic must stay within national boundaries
- Adhere to policies even when **failures** occur



Mechanisms: Device-by-Device

- Local decisions made independently on each device
- Several device-level actions together imply some higher-level behavior

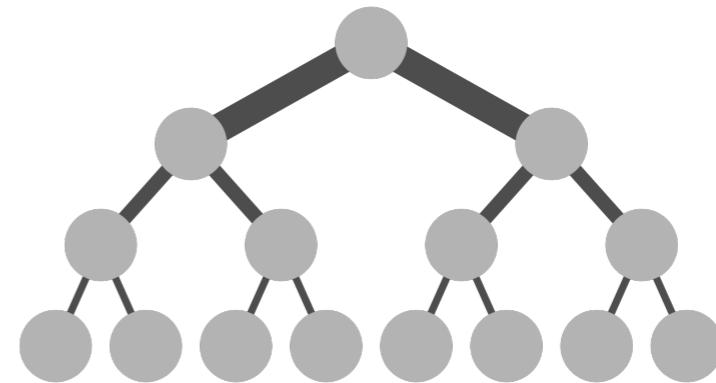
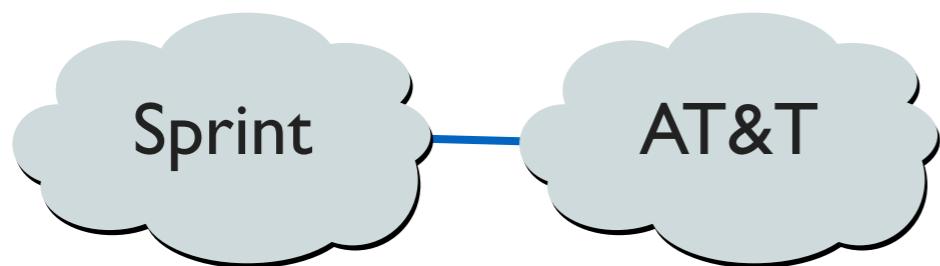
“a large backbone network may have a thousand different policies configured across hundreds of routers” [Fteamster 2005]

- Failures interact with local decision-making algorithms in complex ways

Why use BGP?

BGP in the Wild

- Inter-domain routing between Autonomous Systems (AT&T, Sprint)
- Intra-domain routing in data centers



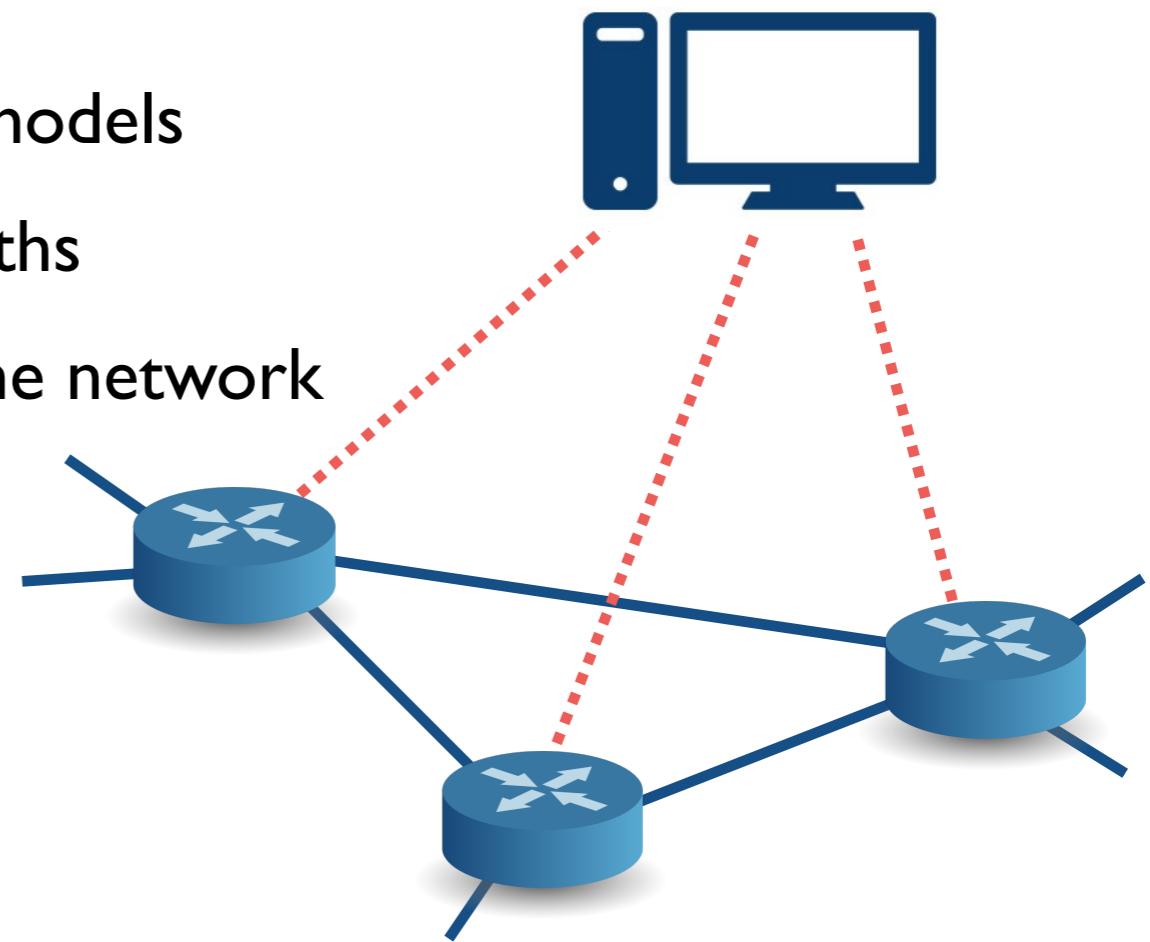
Advantages

- Allows for *local* policy among nodes
- Massively *scalable* — the only routing protocol that scales to hundreds of thousands of IP prefixes
- Fully *distributed* — does not require global knowledge of network state

What about SDN?

Software Defined Networks

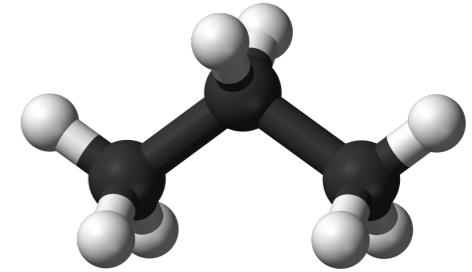
- Simpler, centralized programming models
- Network-wide abstractions like paths
- Centralized controller programs the network



But they are not a panacea:

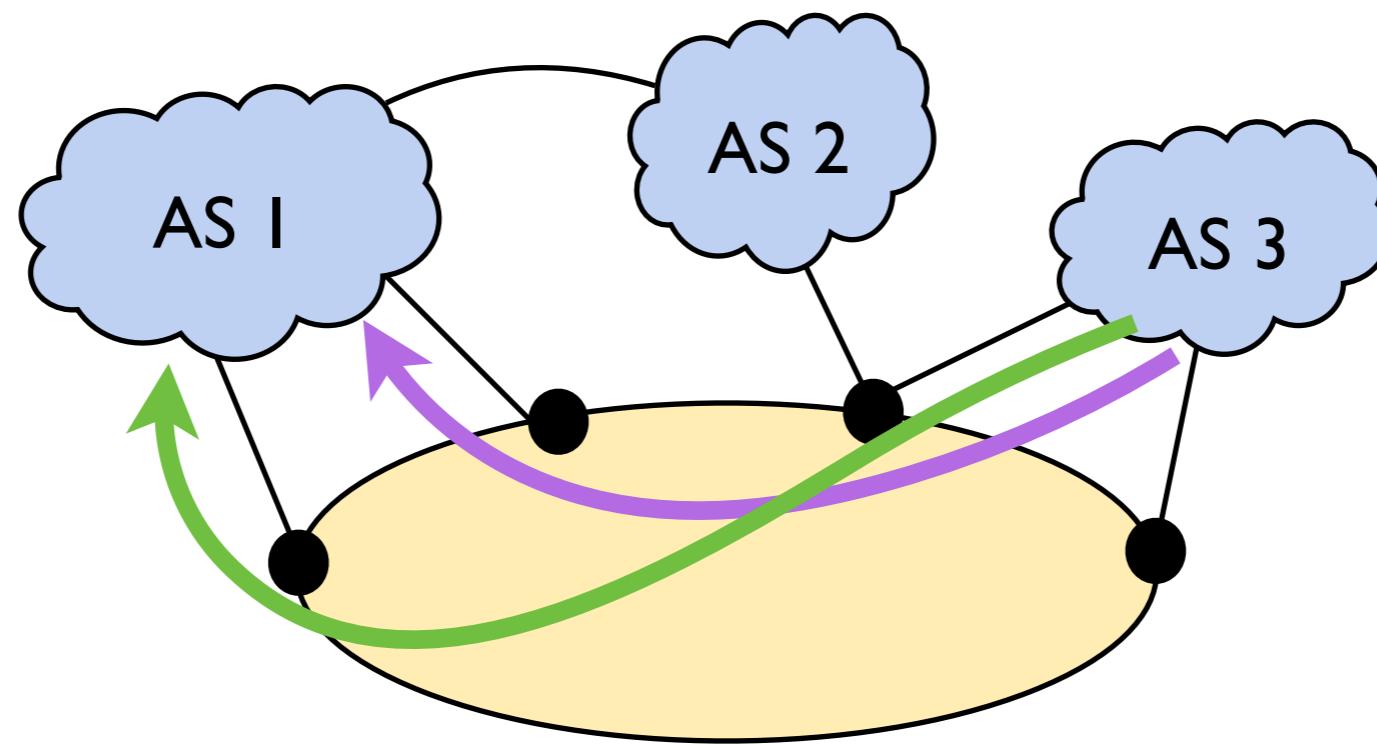
- Do not usually help with **inter-domain** routing
- **Latency**, and **scalability** can be problematic for large networks — often requires bringing back distributed control planes.
- Require careful design and engineering to be robust to **failures**
- Their implementations don't exploit **existing** network infrastructure

Propane: Programming a Distributed Control Plane

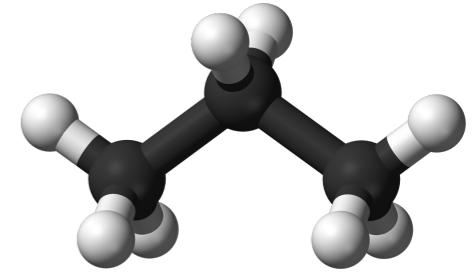


I) Language for expressing high-level operator objectives with:

- Network-wide programming abstraction
- Uniform abstractions for intra- and inter-domain routing
- Paths constraints and relative preferences with fall-backs in case of failures

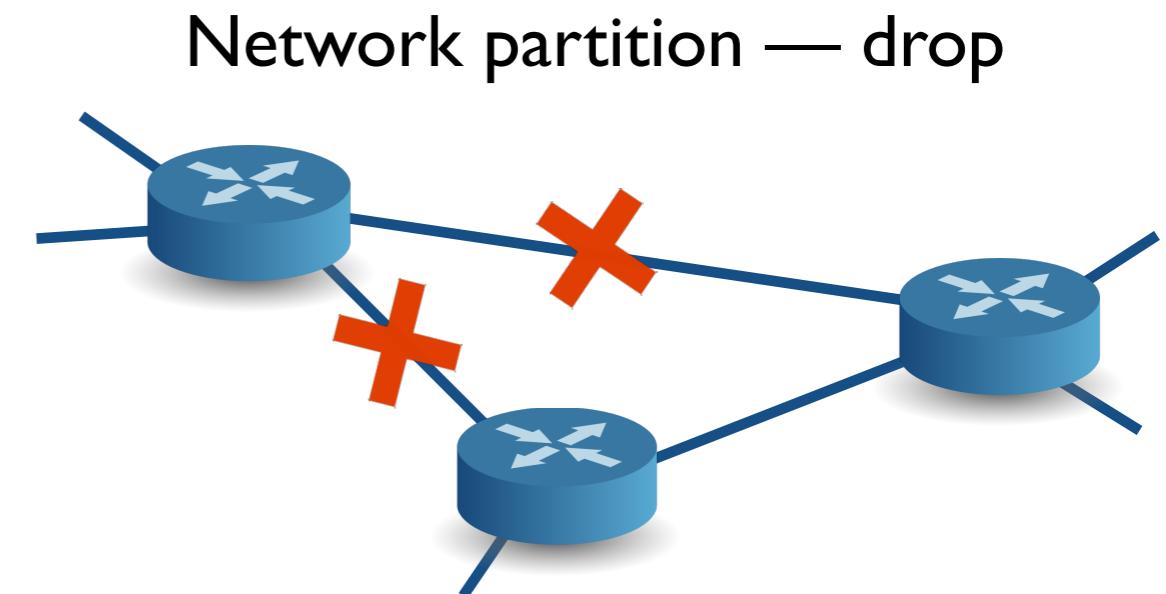
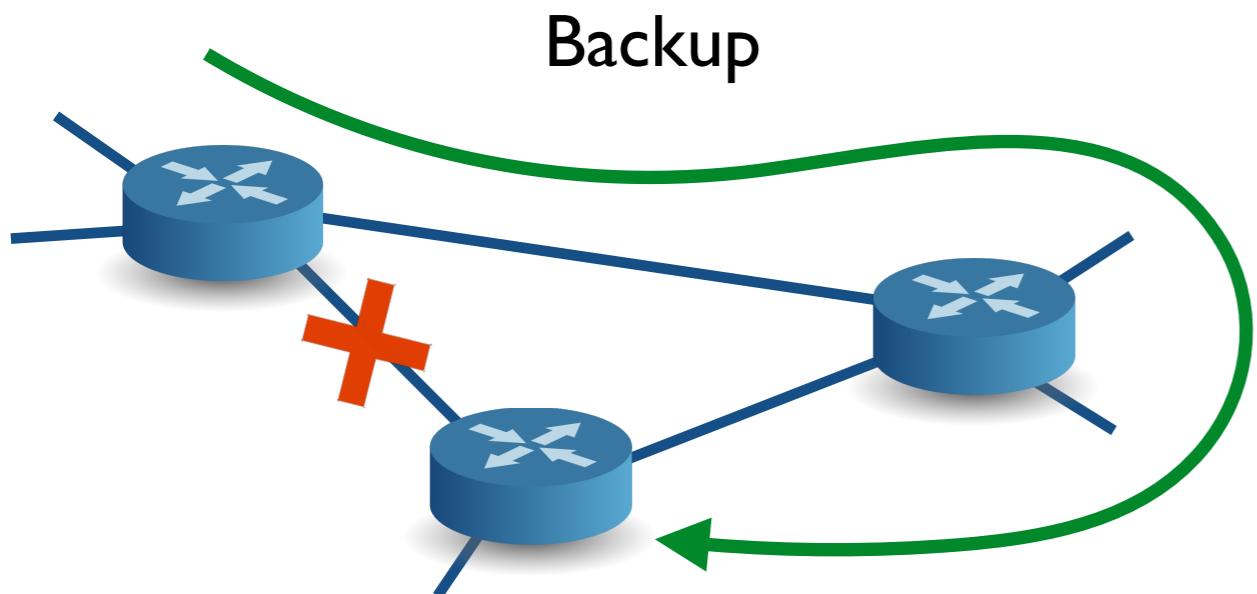


Propane: Programming a Distributed Control Plane



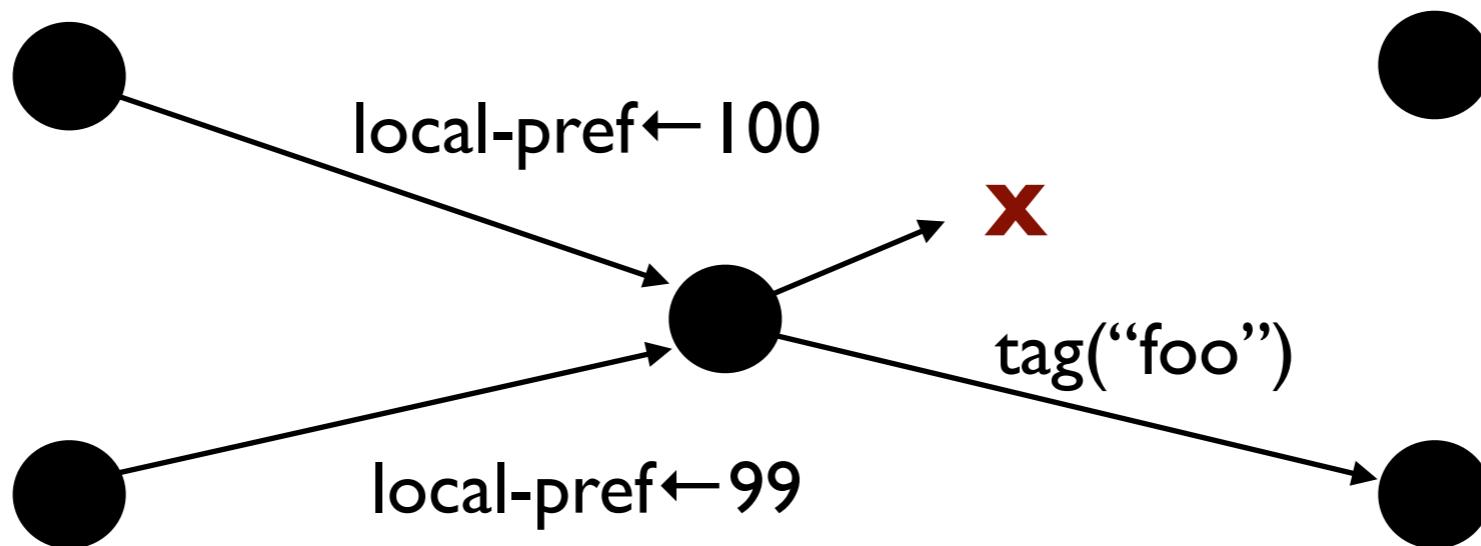
2) Compiler to generate a low-level distributed implementation:

- Efficient algorithms to synthesize a set of *policy-compliant* BGP configs
- Static analysis guarantees policy compliance under *all* failures



Border Gateway Protocol (BGP)

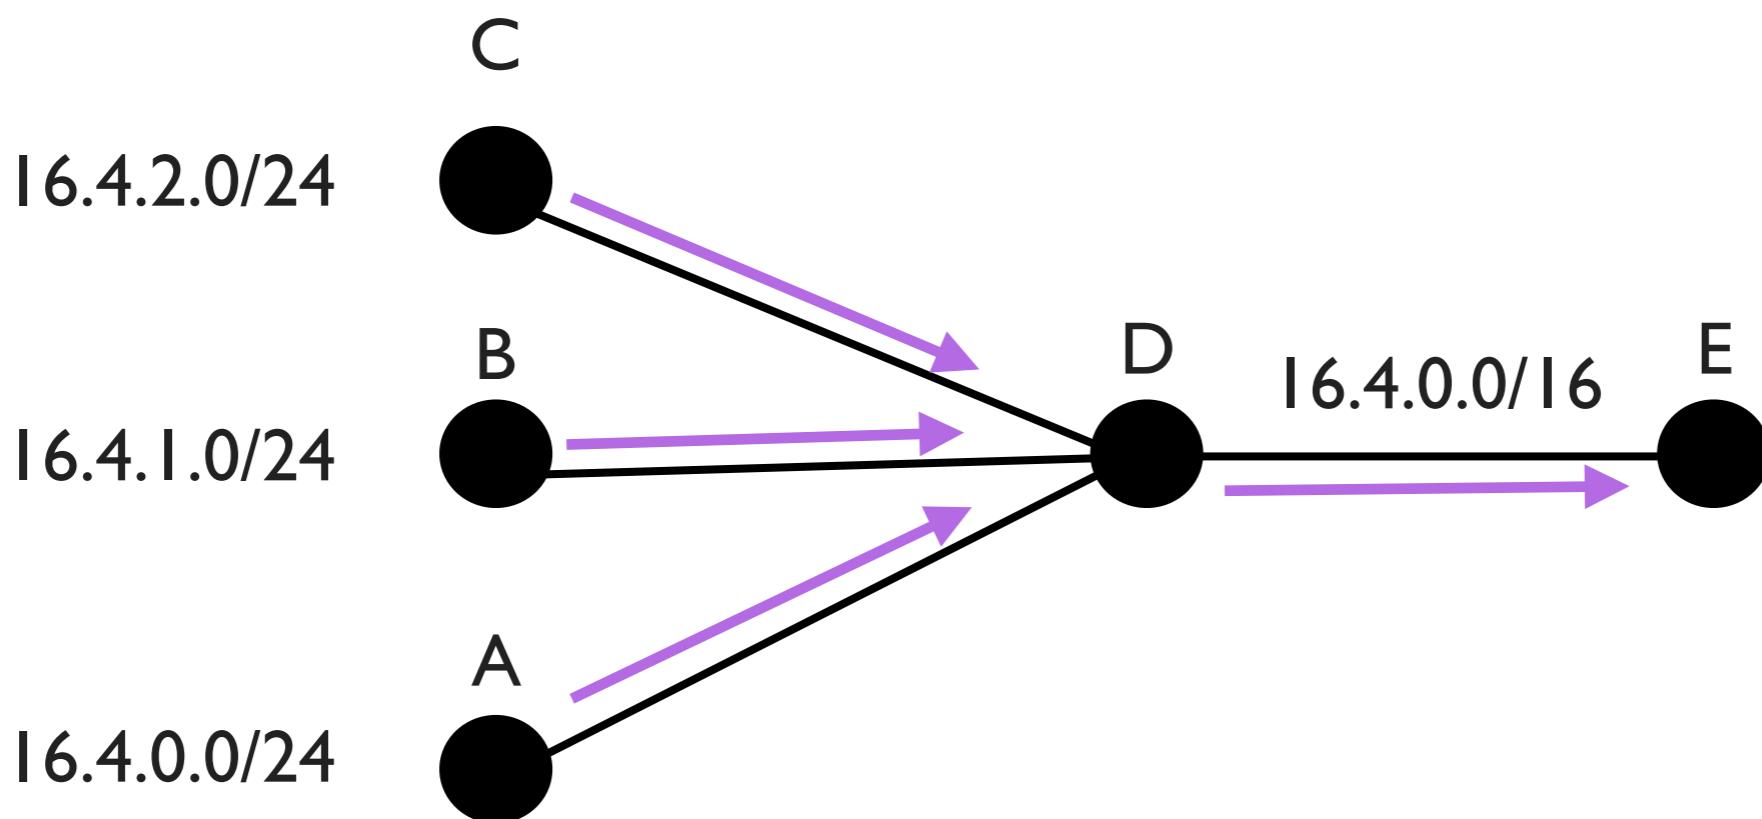
- Routers send advertisements for a given prefix (e.g., 16.4.0.0/16)
- Advertisements contain the AS path traffic would take
- Each router applies **local policy** to choose best route
 - 1.Import filters drop advertisements or modify attributes
 - 2.Highest preference is chosen. Path length as a tie breaker
 - 3.Export filters drop advertisements or modify attributes



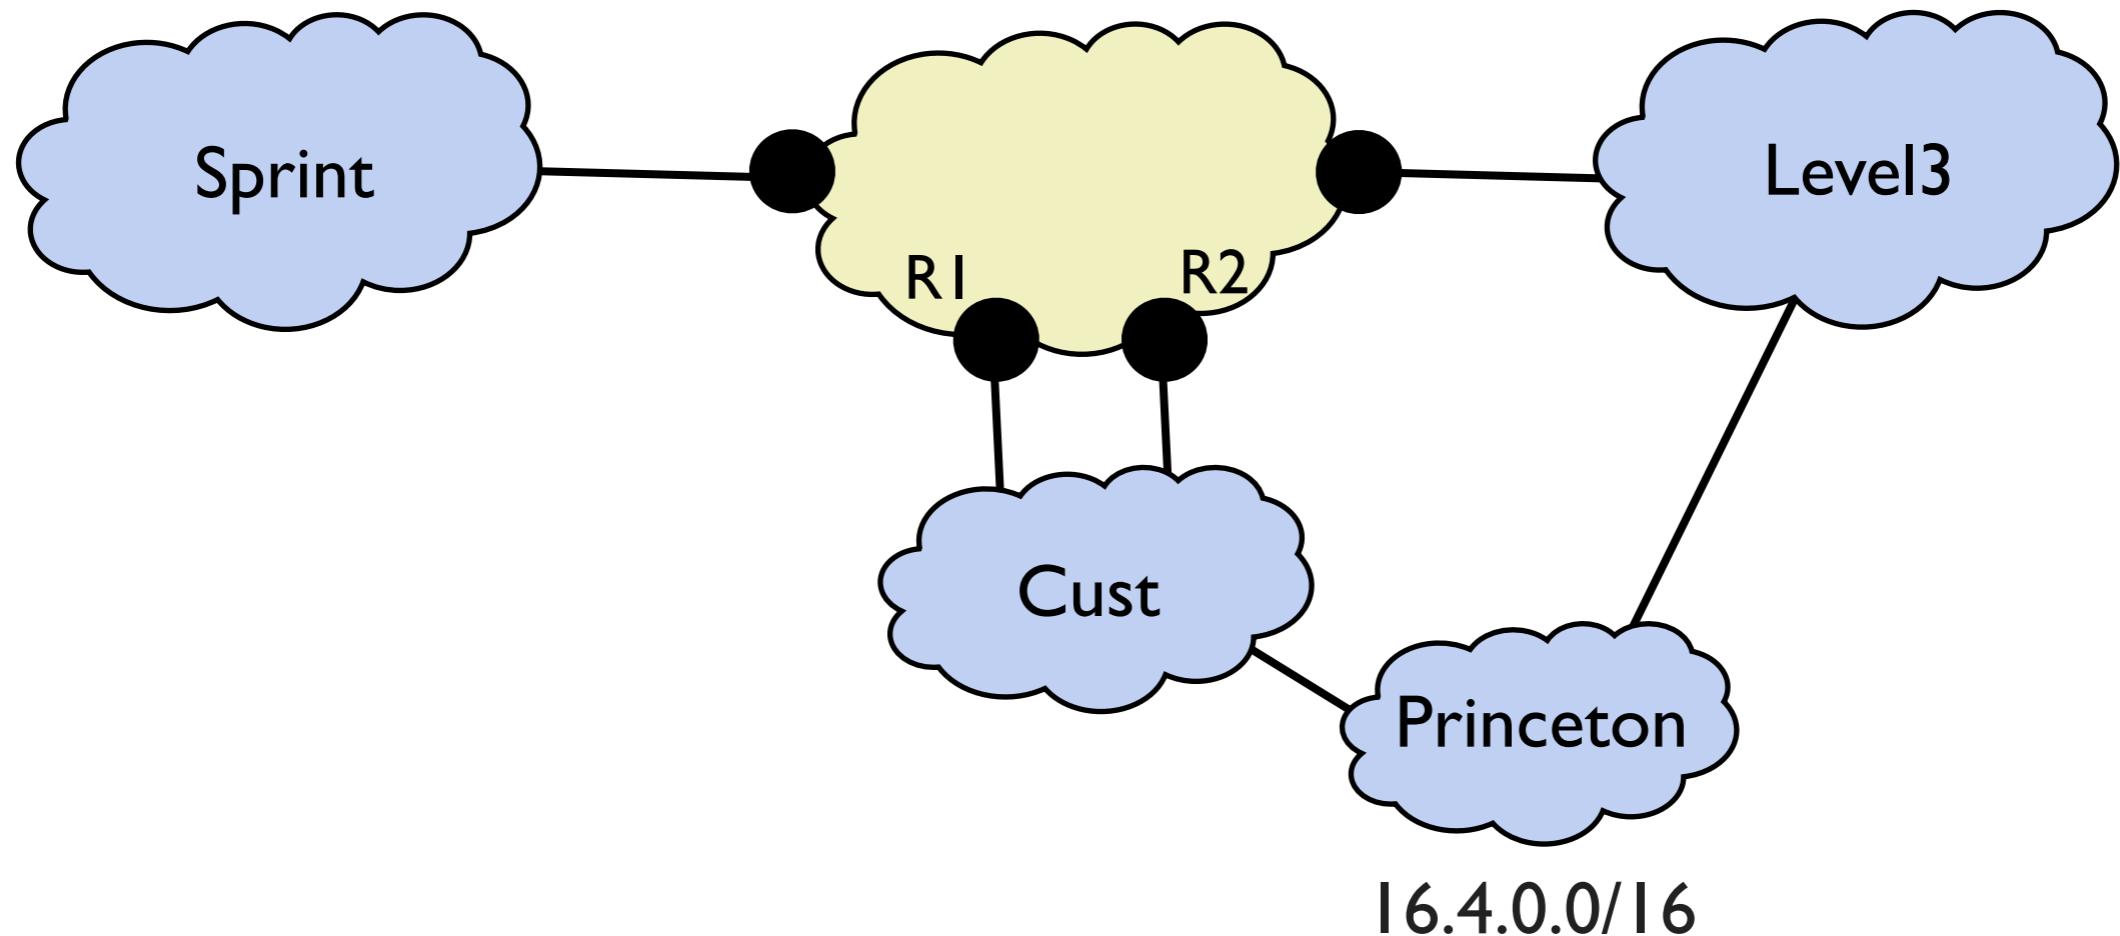
Border Gateway Protocol (BGP)

Aggregation (Route summarization)

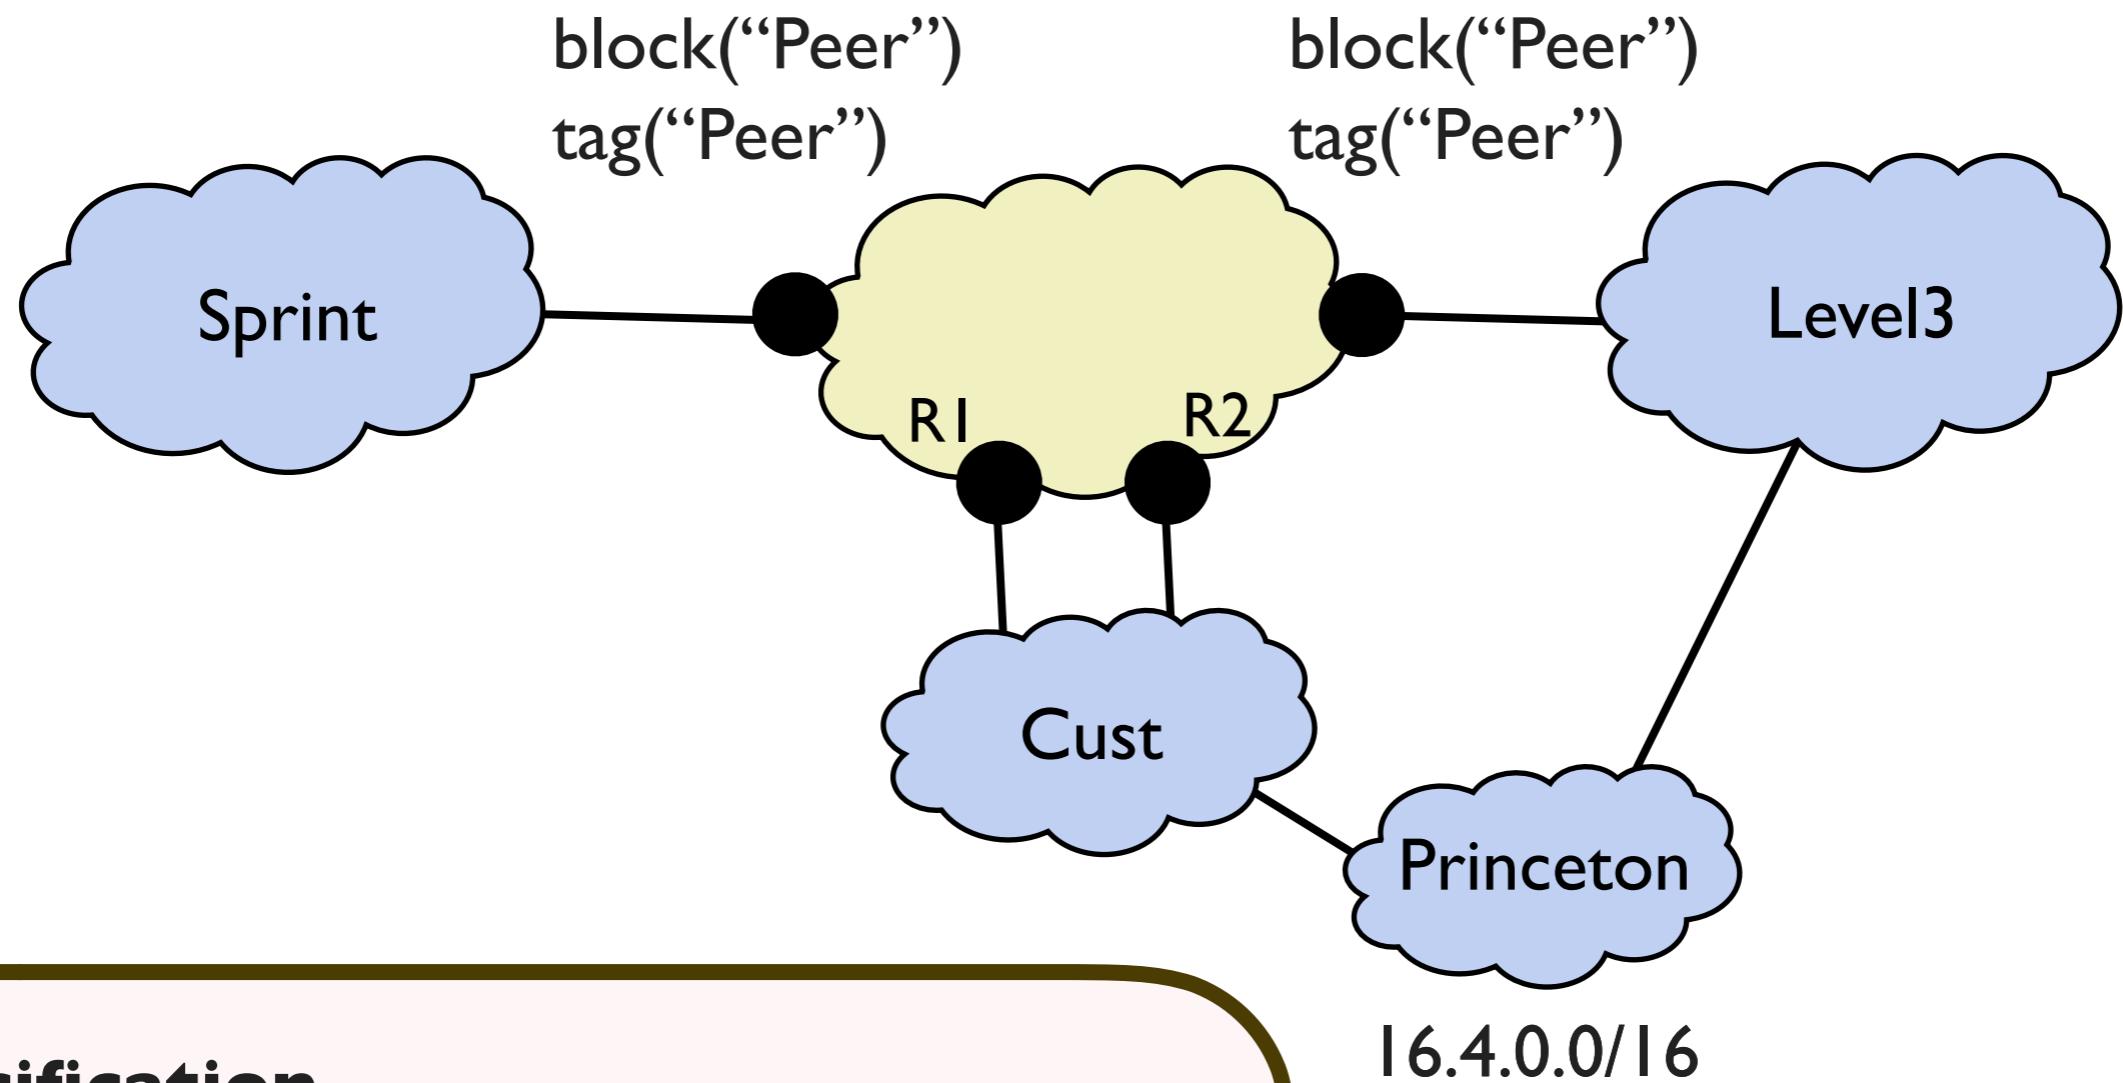
- Advertise a more general subnetwork
- Information hiding for routers
- Reduces router table size and improves network stability



Example 1: Backbone Network



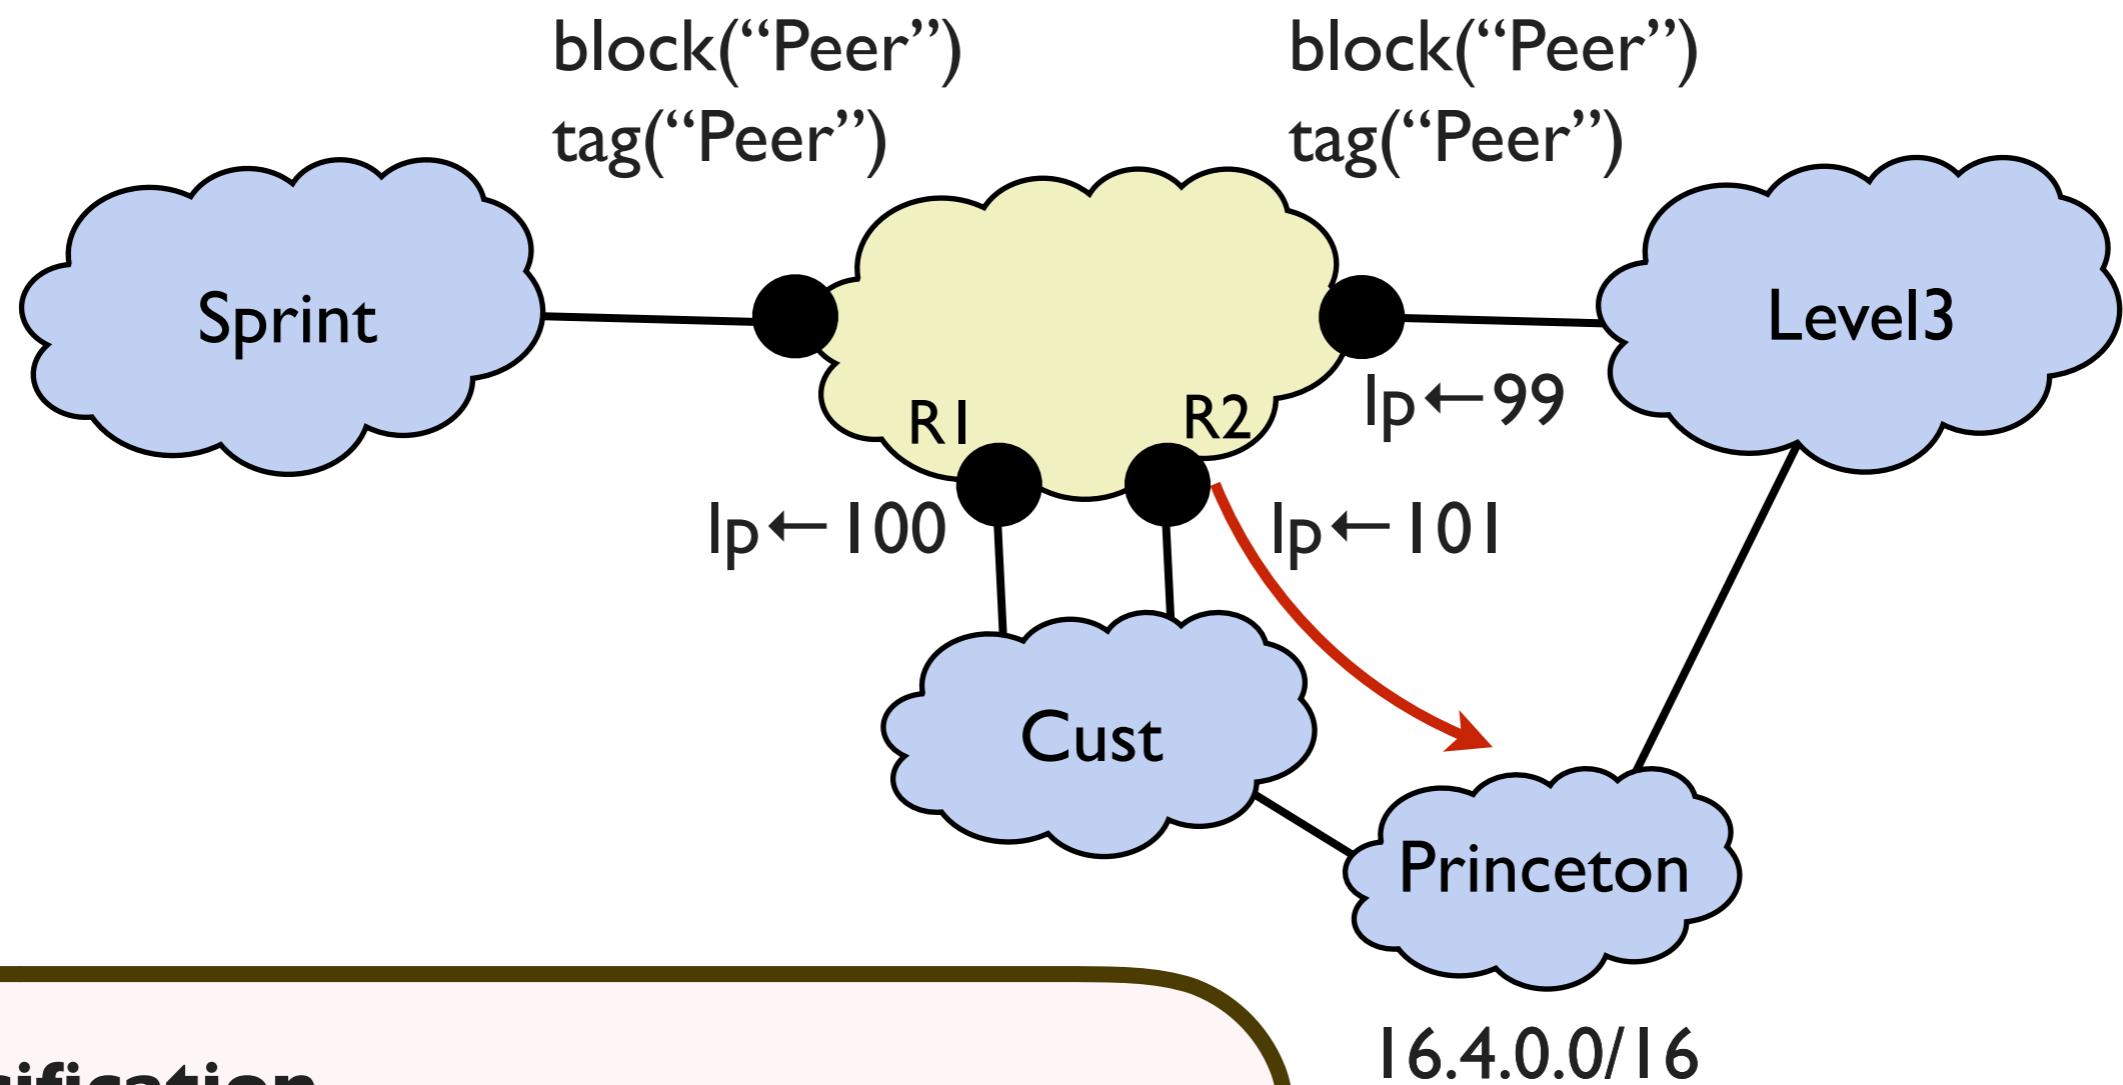
Example 1: Backbone Network



Specification

- Prevent transit between peers (\$\$\$)

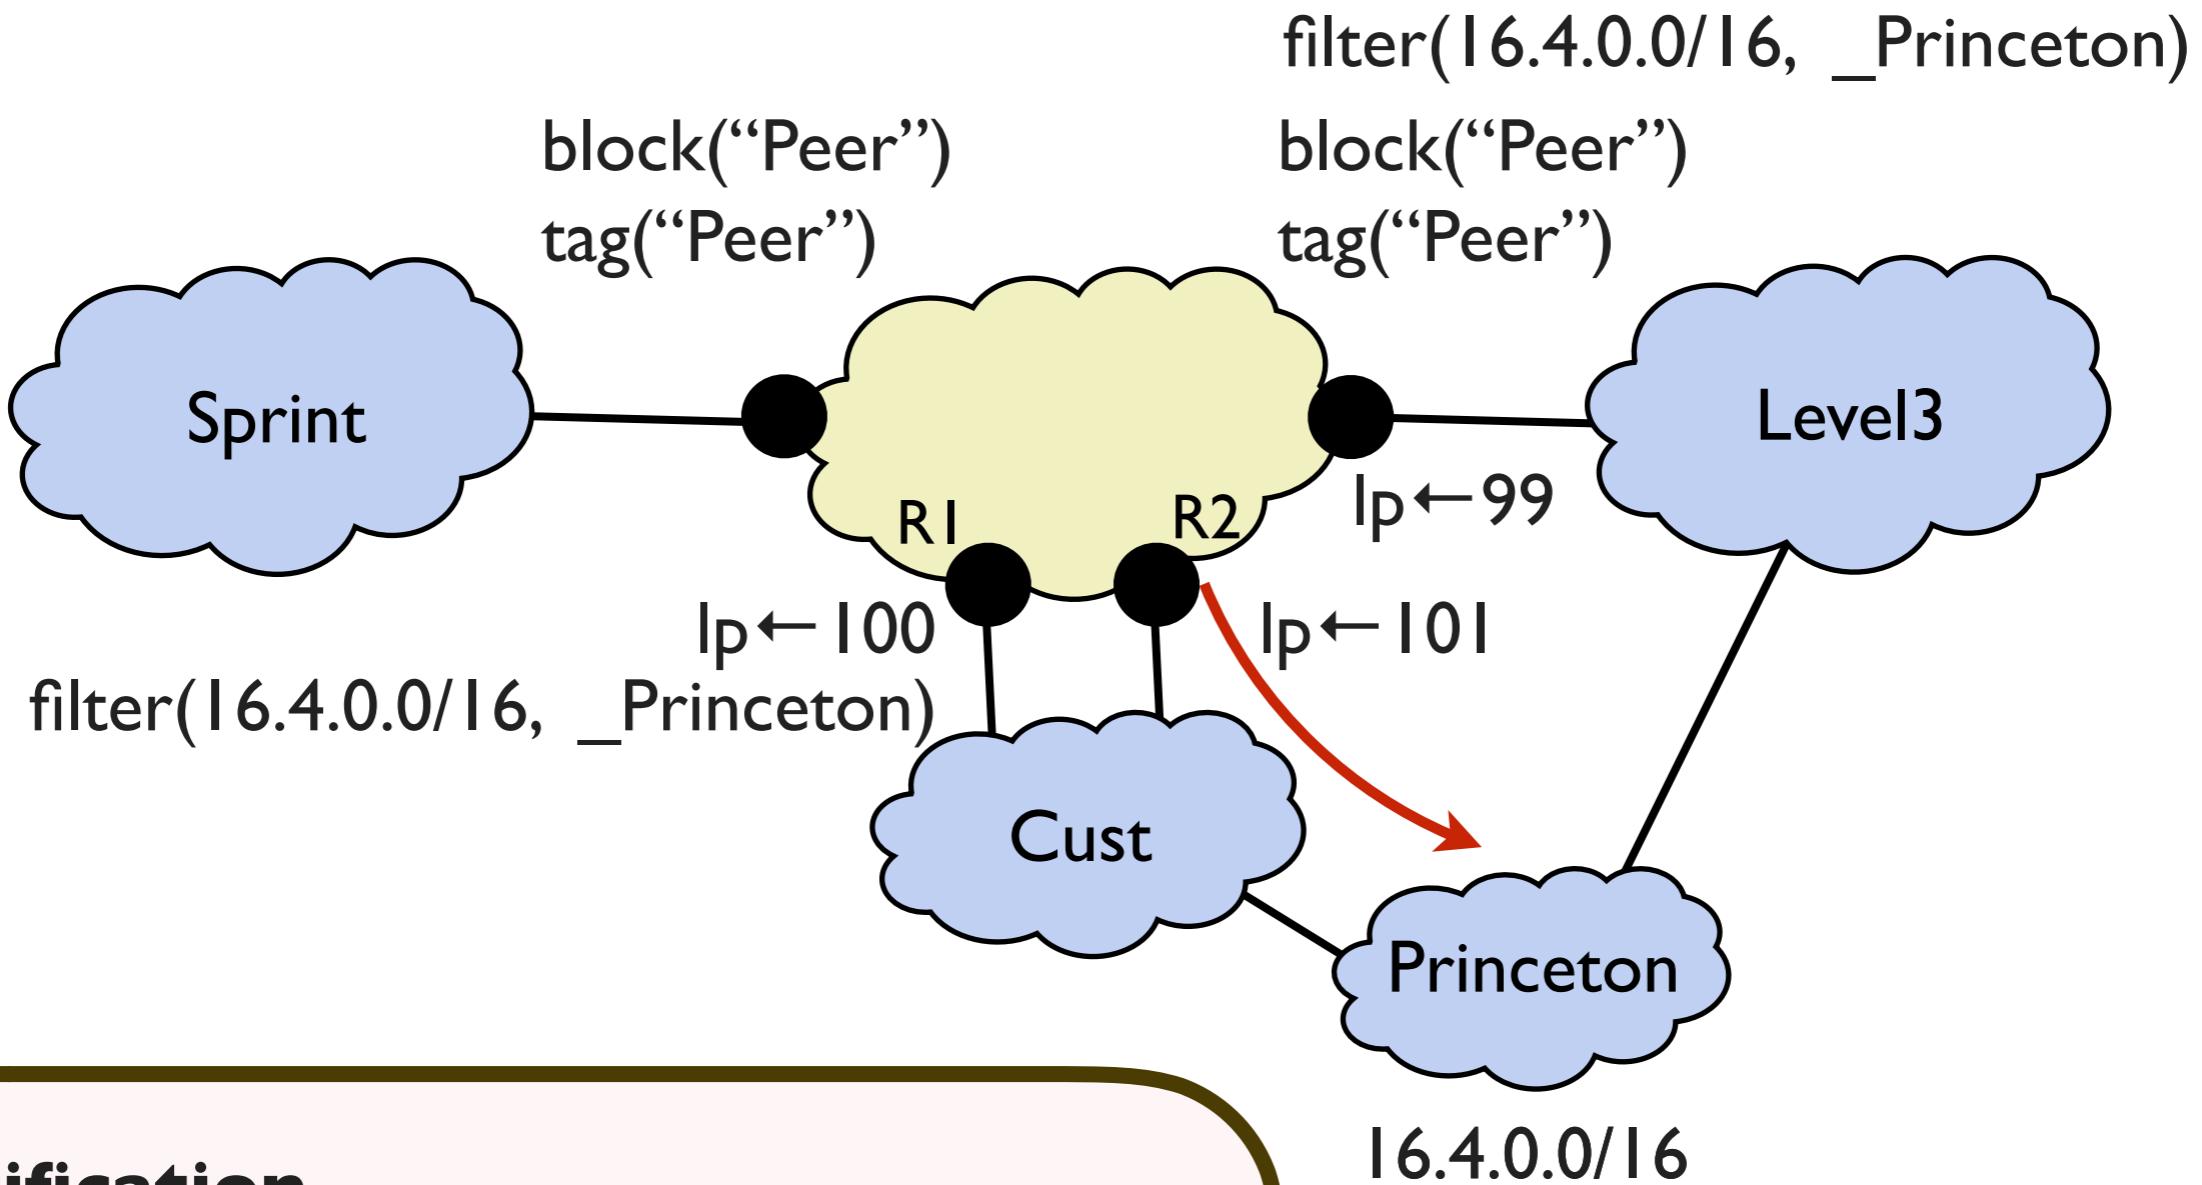
Example I: Backbone Network



Specification

- Prevent transit between peers (\$\$\$)
- Prefer R2 > RI > Peer (\$\$\$)

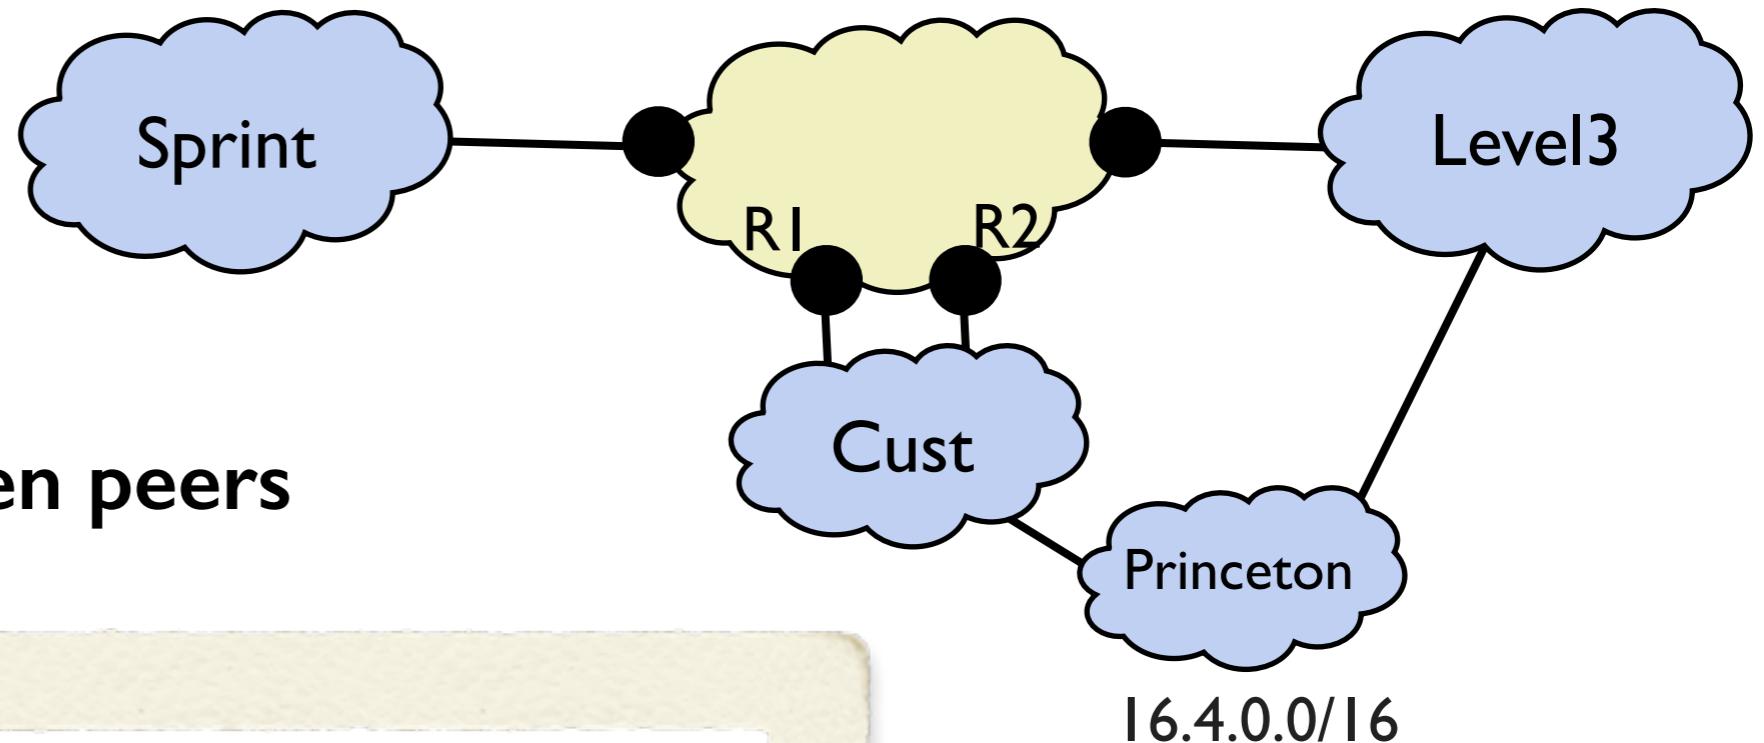
Example I: Backbone Network



Specification

- Prevent transit between peers (\$\$\$)
- Prefer R2 > RI > Peer (\$\$\$)
- Filter customer by prefix (security)

Example 1: Backbone Network



Prevent transit between peers

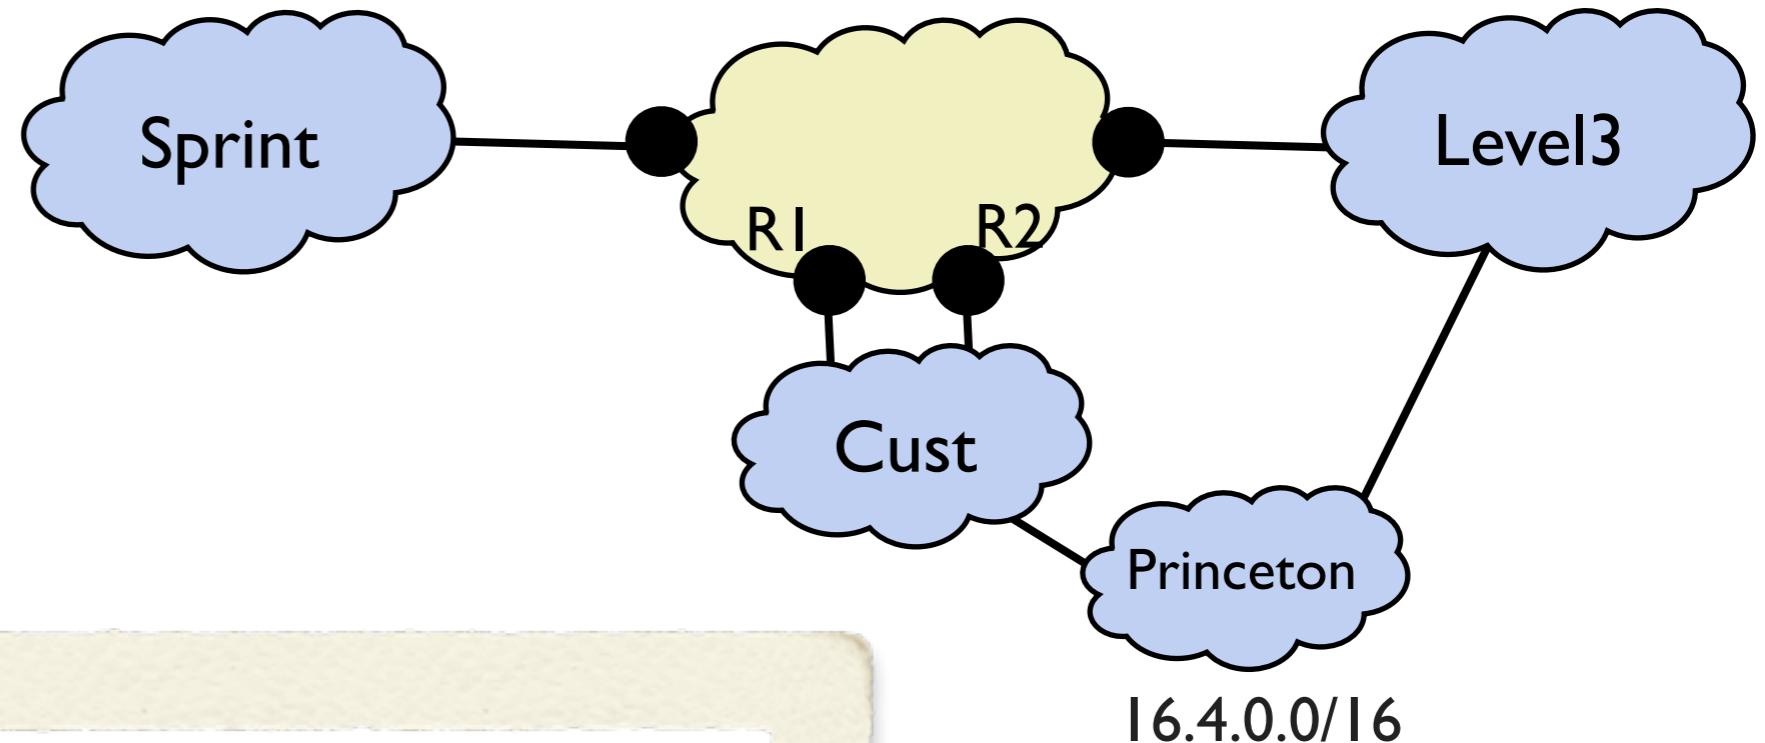
```
define Peer = {Sprint, Level3}
```

```
define transit(X,Y) = enter(X) and exit(Y)
```

```
define NoTransit = {
    true => !transit(Peer,Peer)
}
```

Example 1: Backbone Network

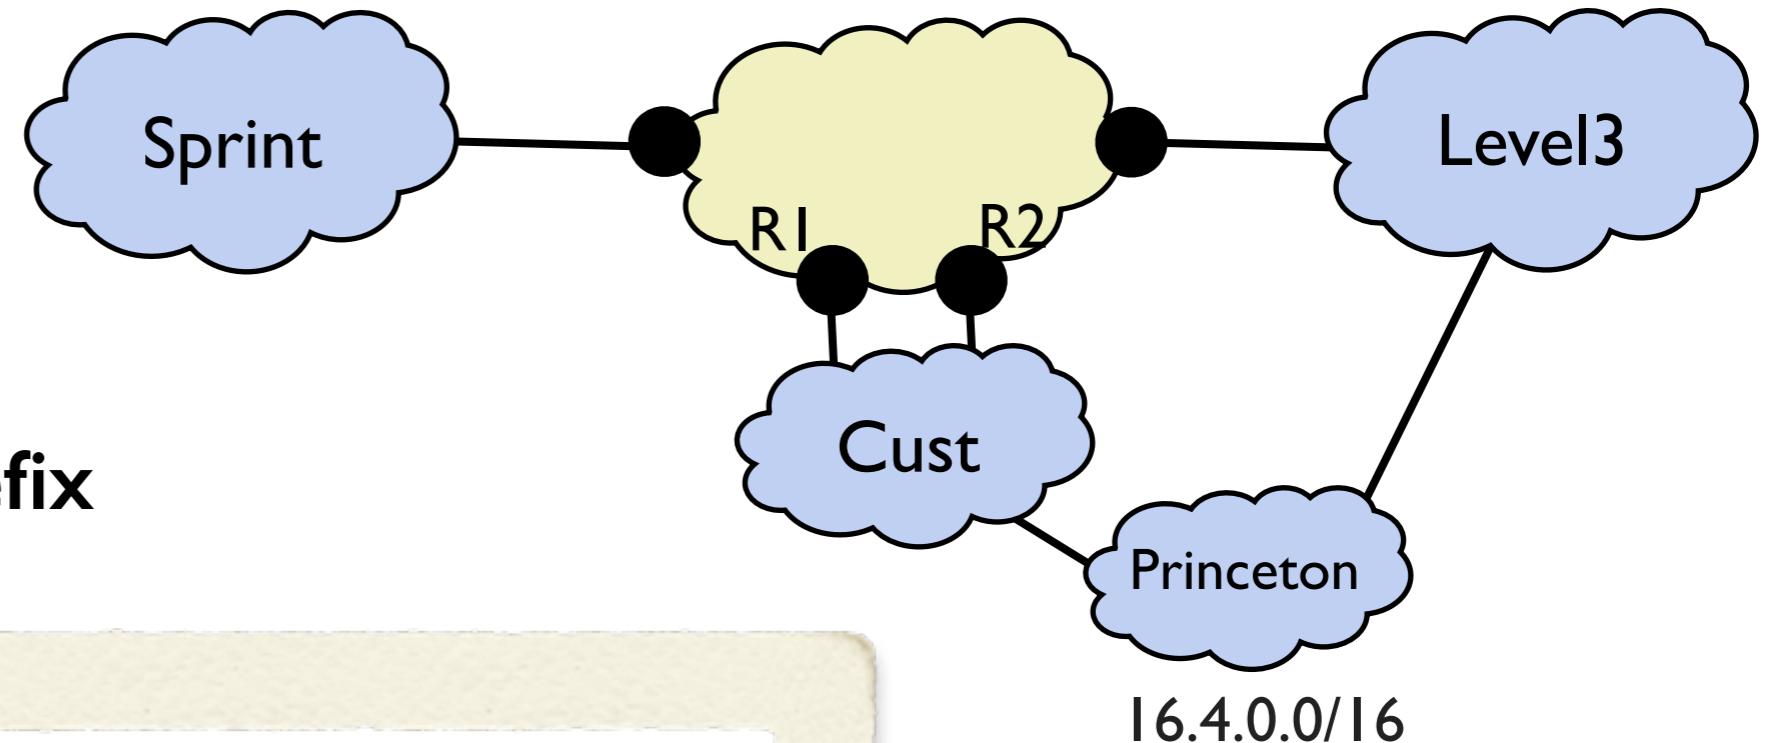
Prefer R2 > RI > Peer



...

```
define Preferences = {  
    true => exit(R2 >> RI >> Peer)  
}
```

Example 1: Backbone Network



Filter customer by prefix

...

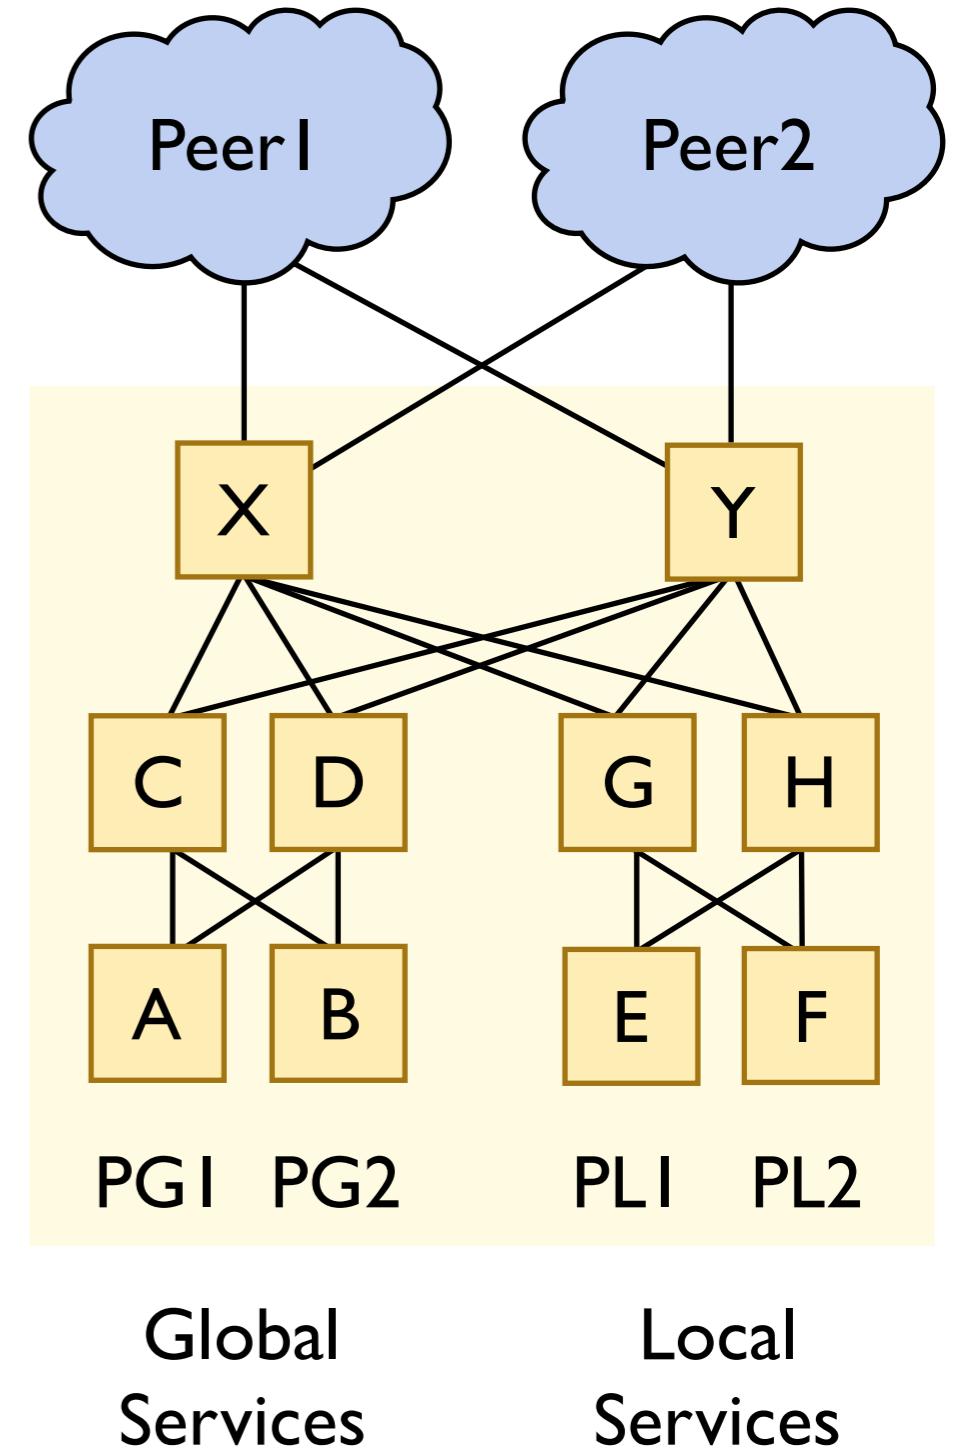
```
define Ownership = {  
    16.4.0.0/16 => end(Princeton),  
}
```

```
define Main =  
    Preferences & Ownership & NoTransit
```

Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers



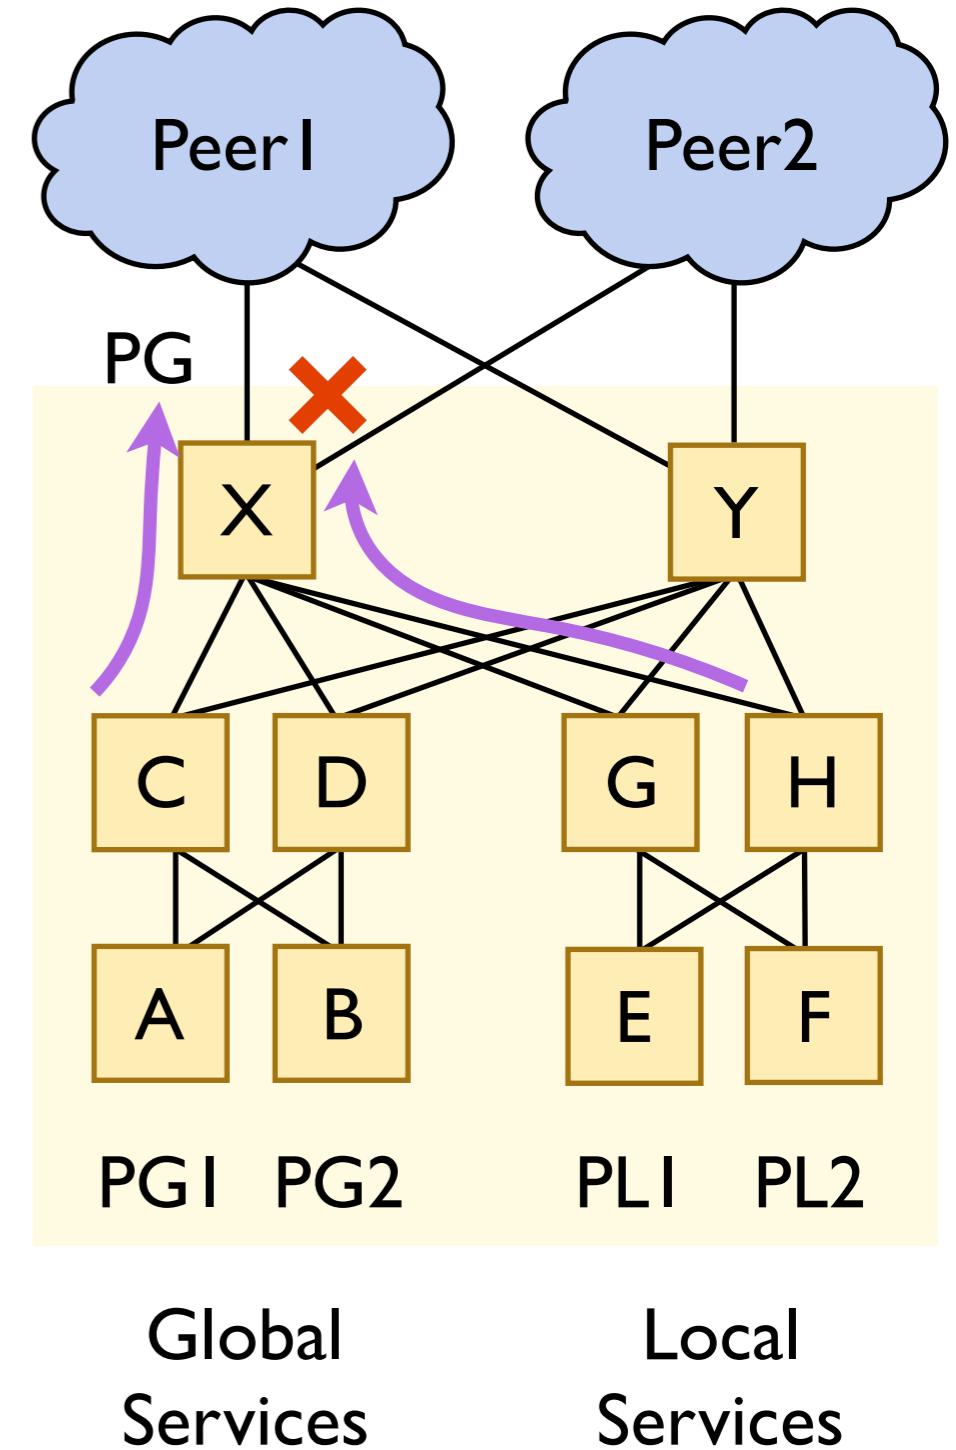
Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- X and Y allow adv. from C,D to peers
- X and Y do not allow adv. from G,H to peers
- Aggregate to PG externally



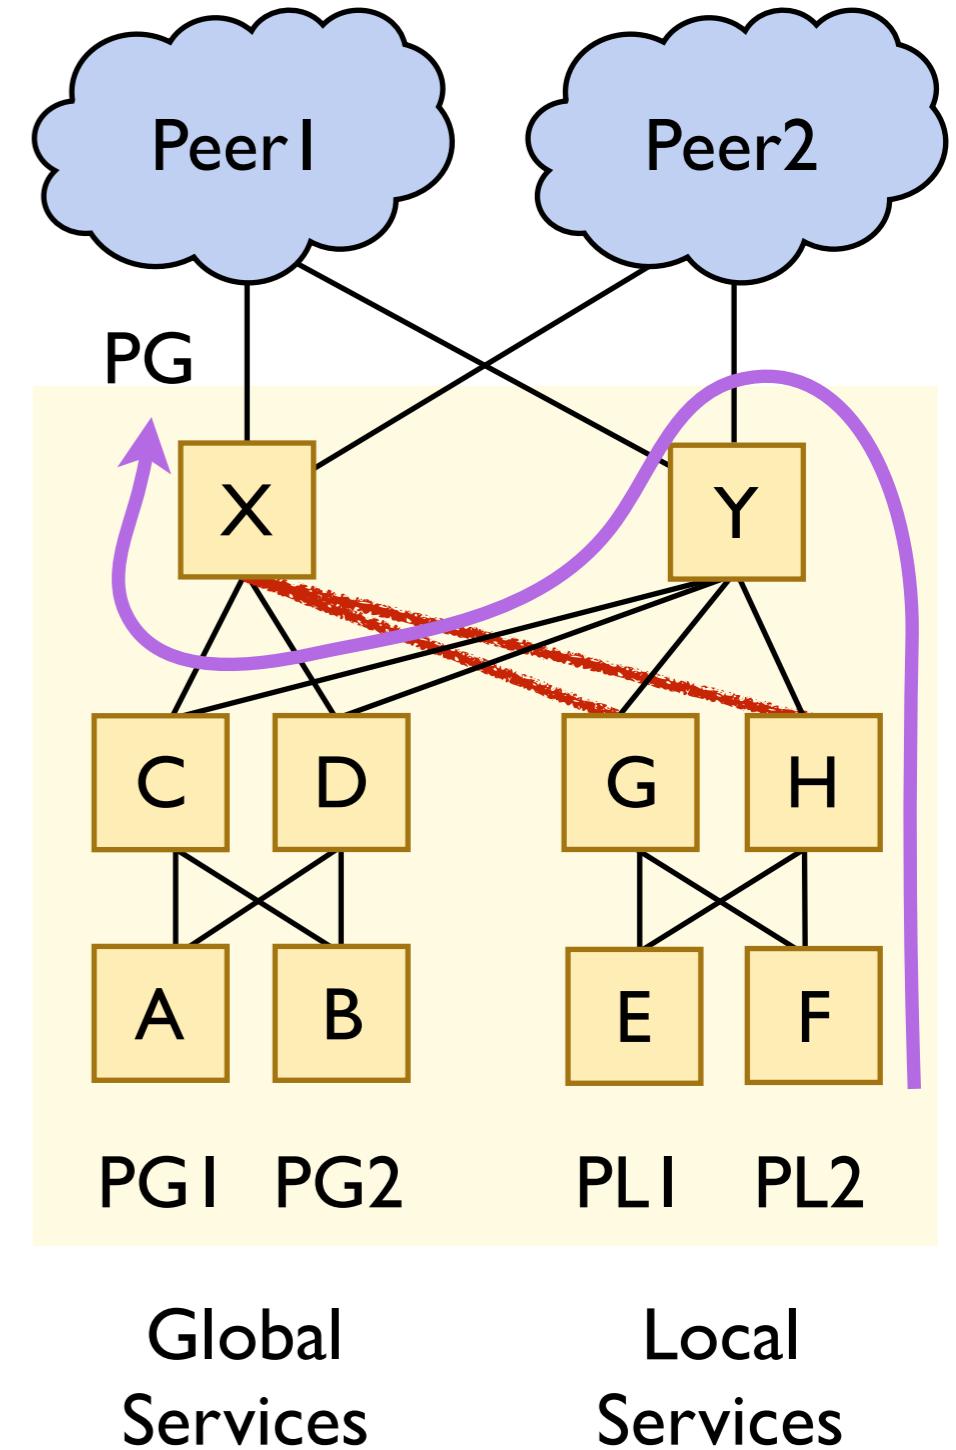
Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- X and Y allow adv. from C,D to peers
- X and Y do not allow adv. from G,H to peers
- Aggregate to PG externally



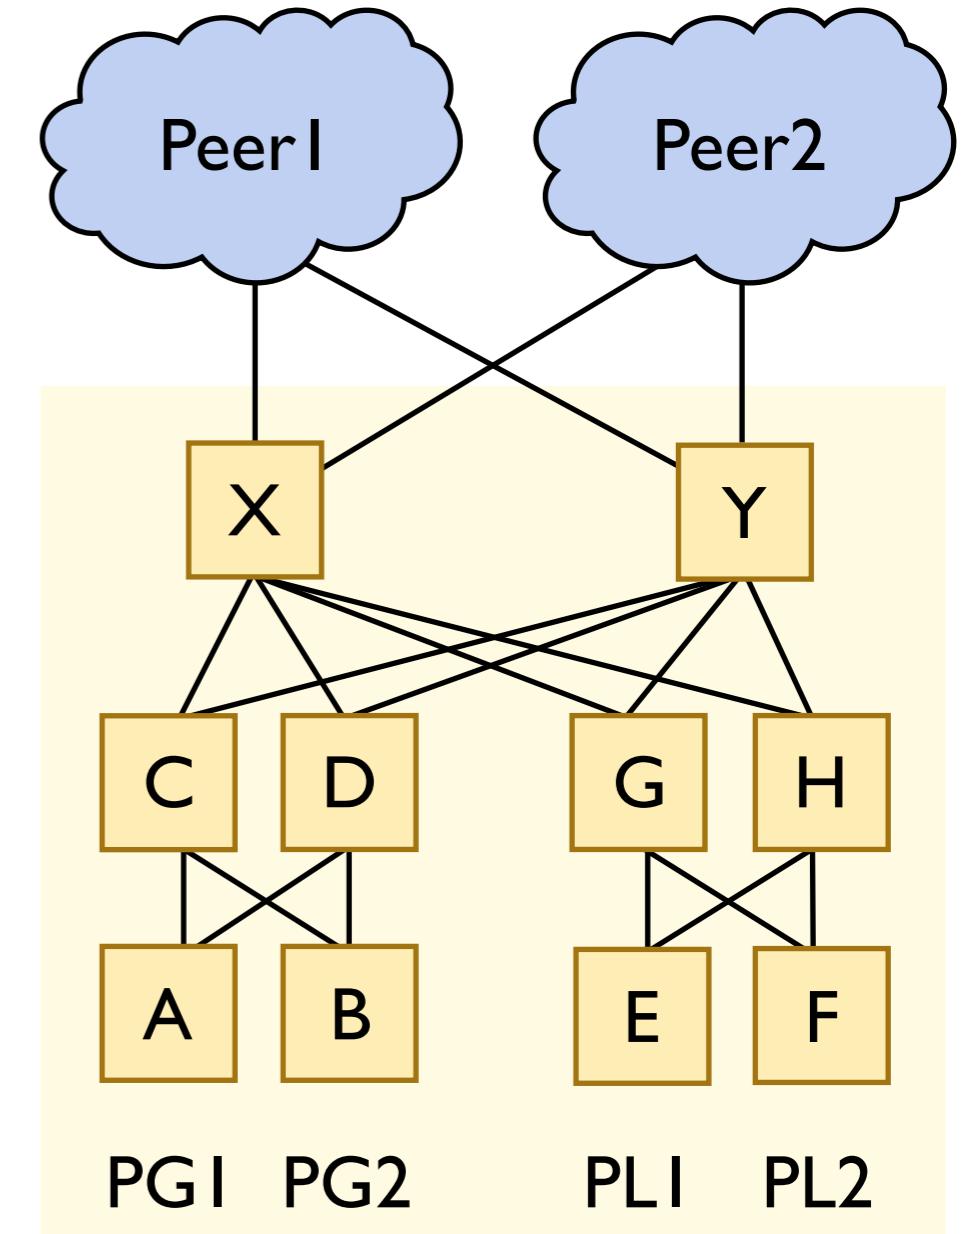
Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- X and Y allow adv. from C,D to peers
- X and Y do not allow adv. from G,H to peers
- X and Y drop routes through each other
- Aggregate to PG externally



Global
Services

Local
Services

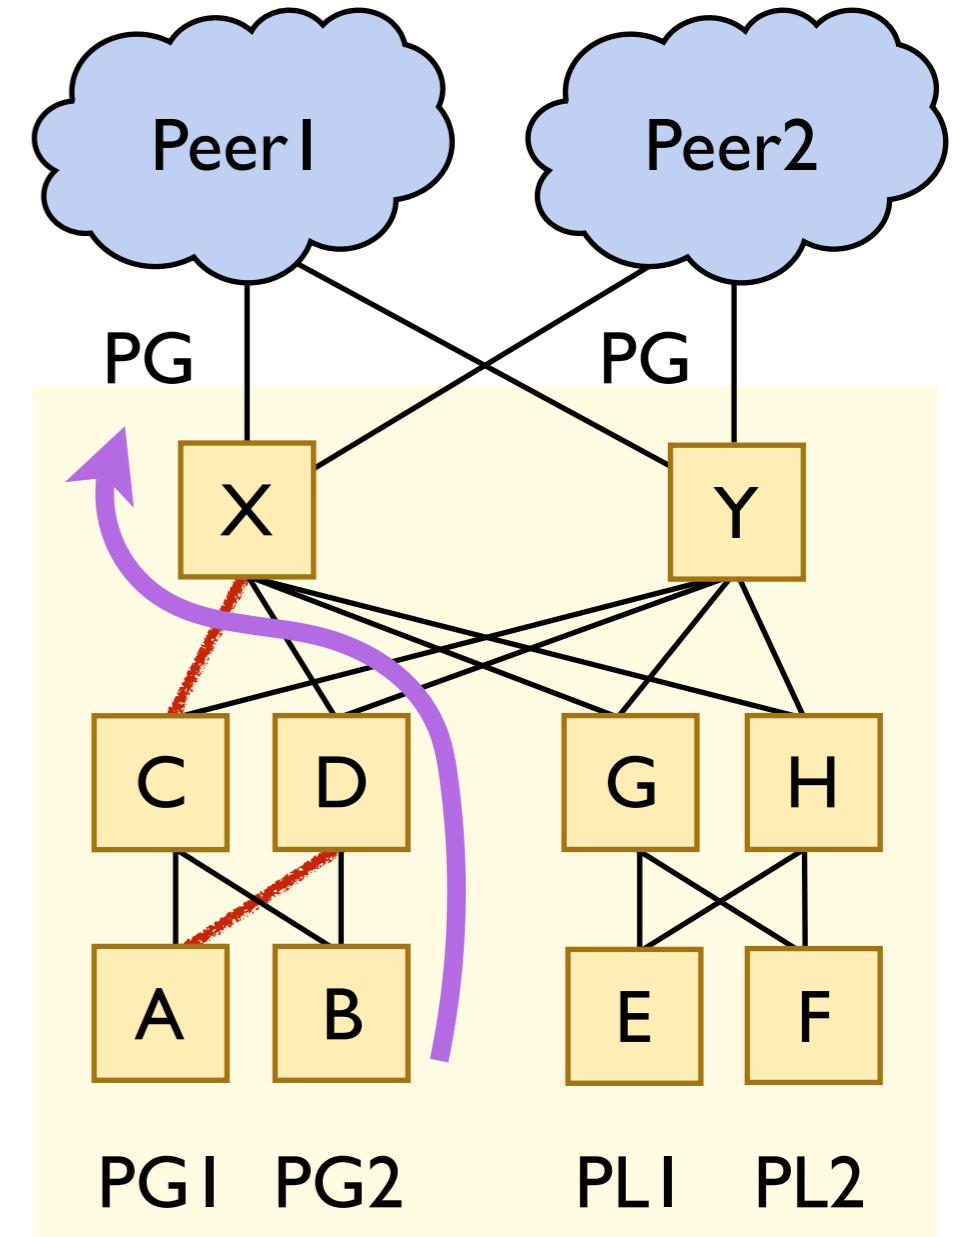
Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

- X and Y allow adv. from C,D to peers
- X and Y do not allow adv. from G,H to peers
- X and Y drop routes through each other
- Aggregate to PG externally



Global
Services

Local
Services

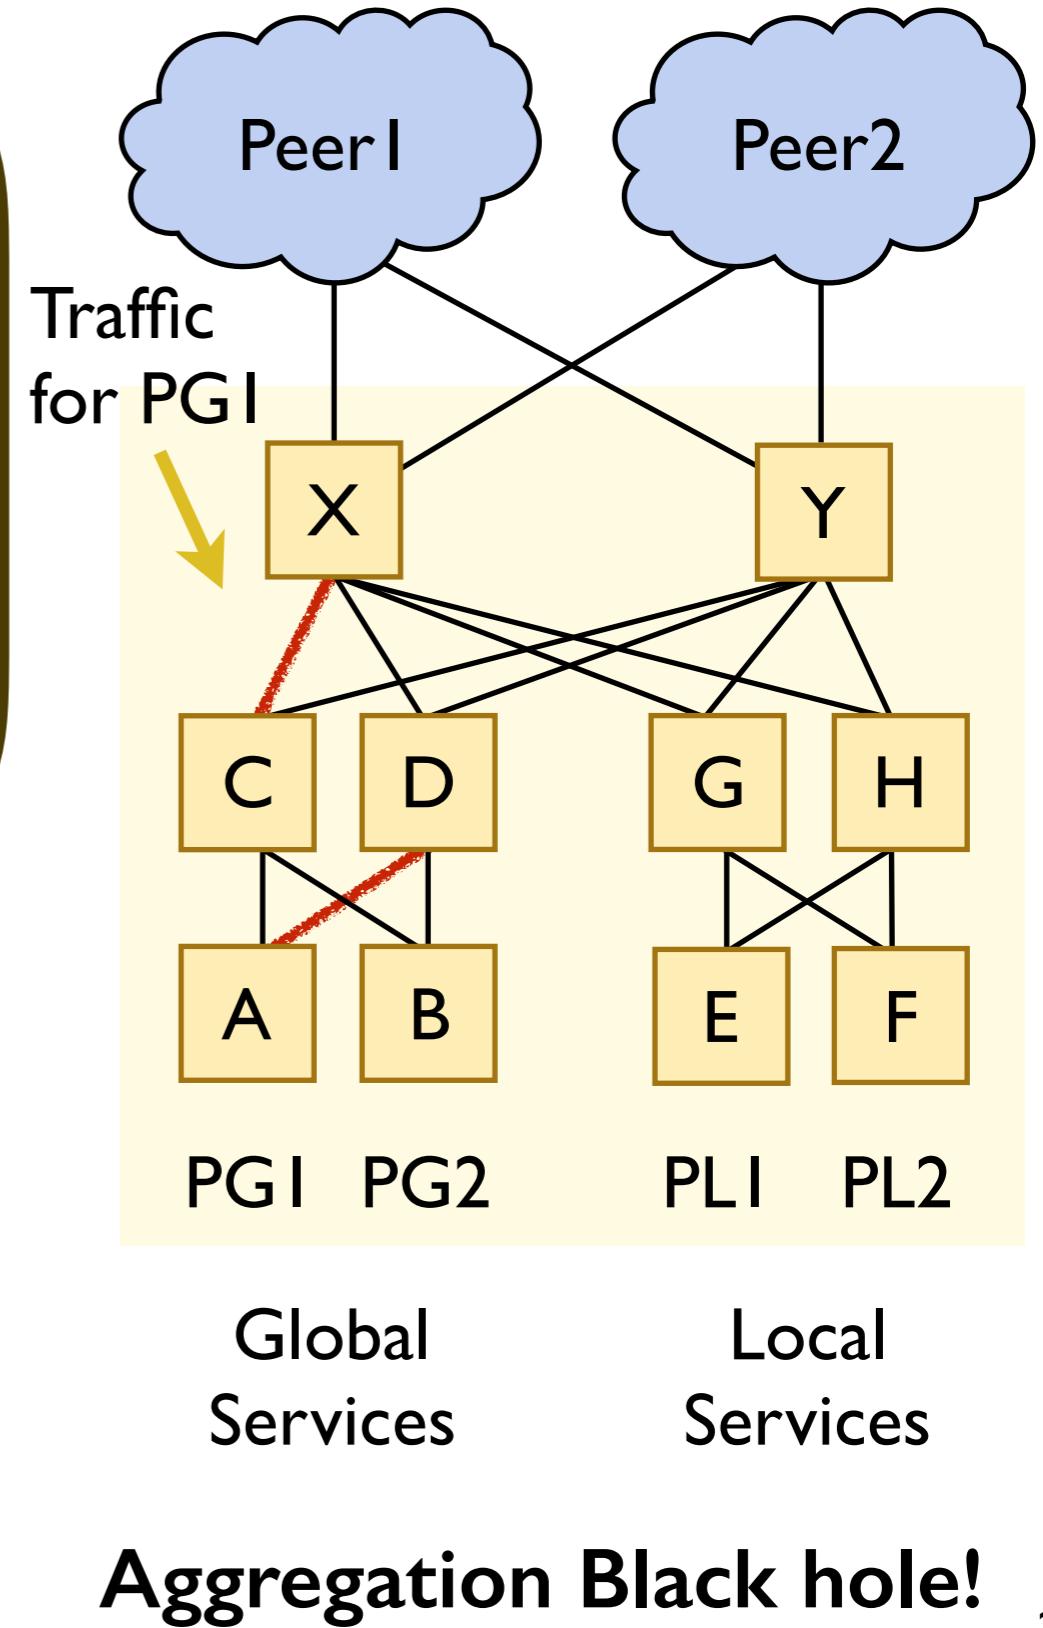
Example 2: A Data Center Network

Specification

- Local prefixes reachable only internally
- Global prefixes reachable externally
- Aggregate global prefixes as PG
- Prefer leaving through Peer1 over Peer2
- Prevent transit traffic between peers

Attempt (I)

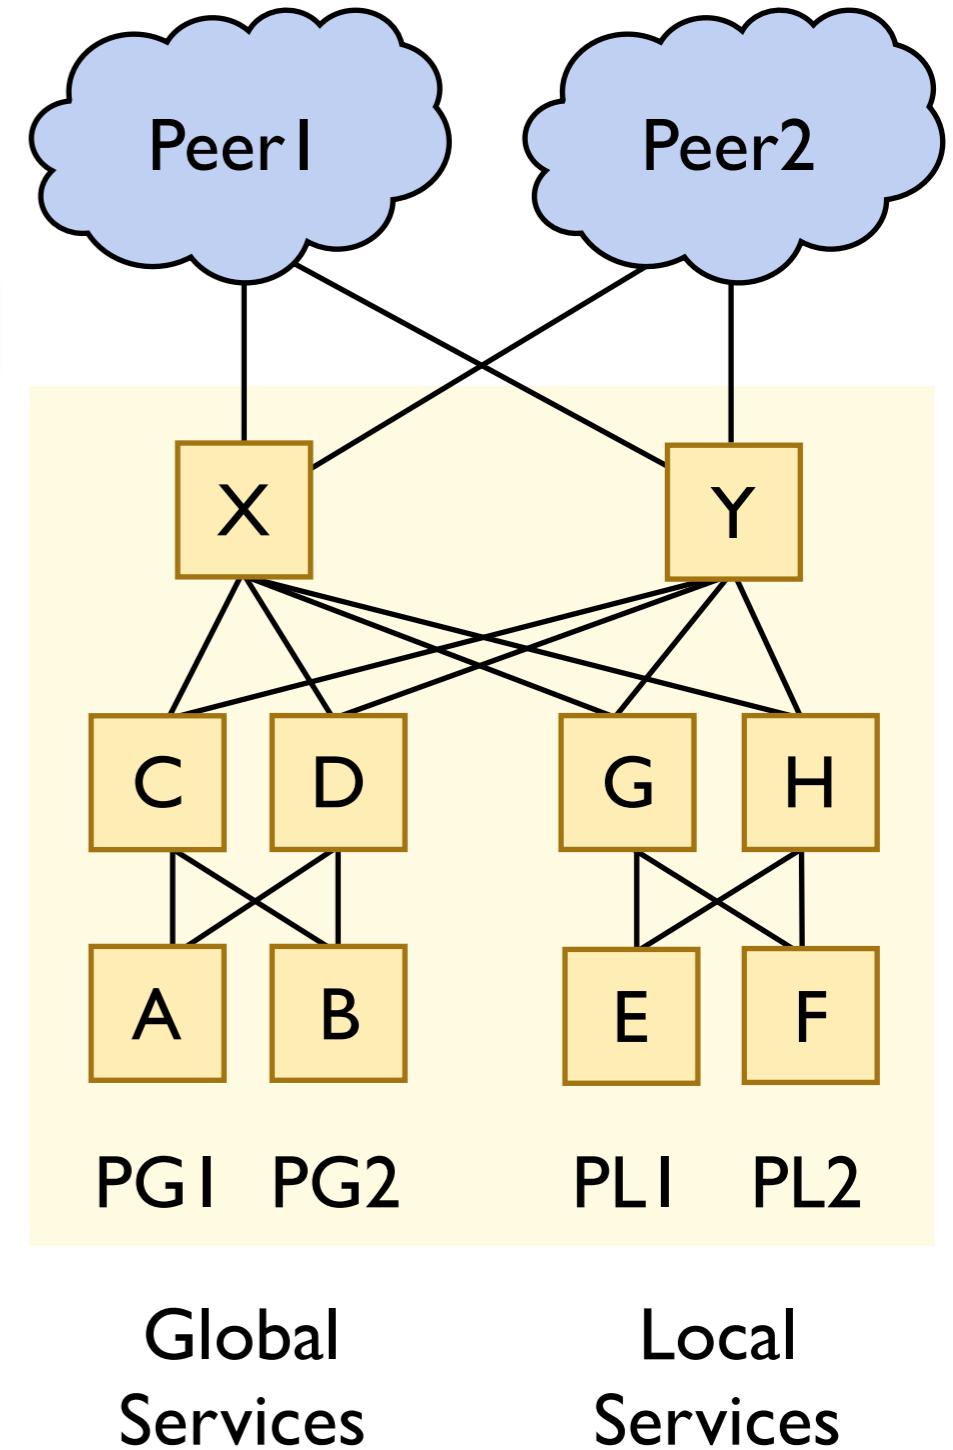
- X and Y allow adv. from C,D to peers
- X and Y do not allow adv. from G,H to peers
- X and Y drop routes through each other
- Aggregate to PG externally



Example 2: A Data Center Network

Basic reachability for prefixes

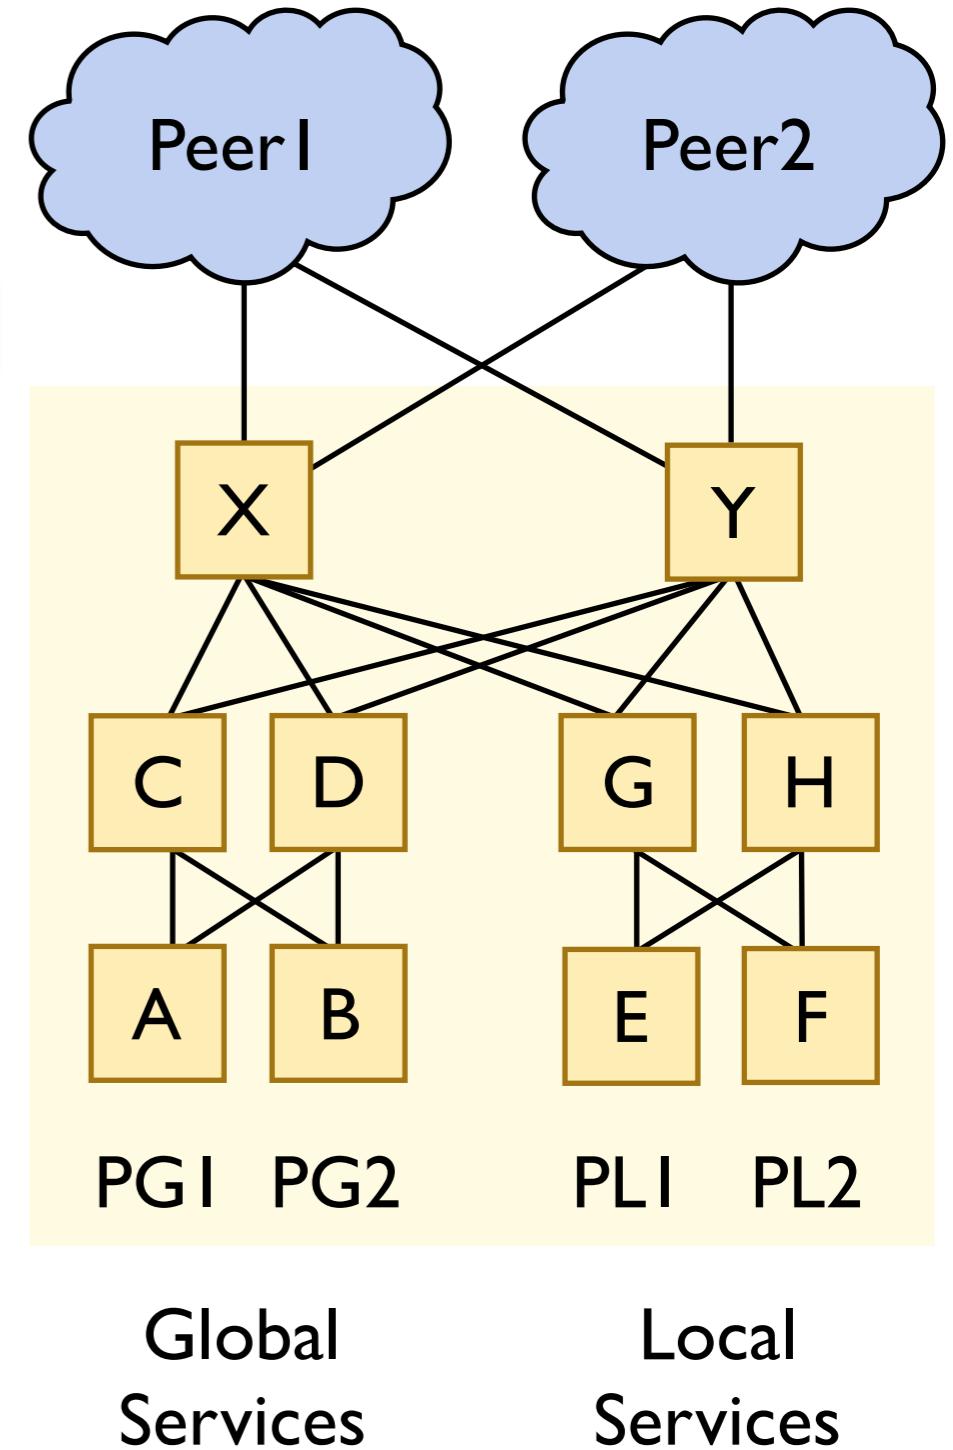
```
define PLI = 1.0.0.0/24  
...  
  
define Ownership = {  
    PGI => end(A),  
    PG2 => end(B),  
    PLI => end(E),  
    PL2 => end(F),  
    true => exit(Peer1) >> exit(Peer2)  
}
```



Example 2: A Data Center Network

Local services stay internal

```
...
# define in = A | B | C | D | E | F | G | H
define Locality = {
    PL1 or PL2 => always(in)
}
```



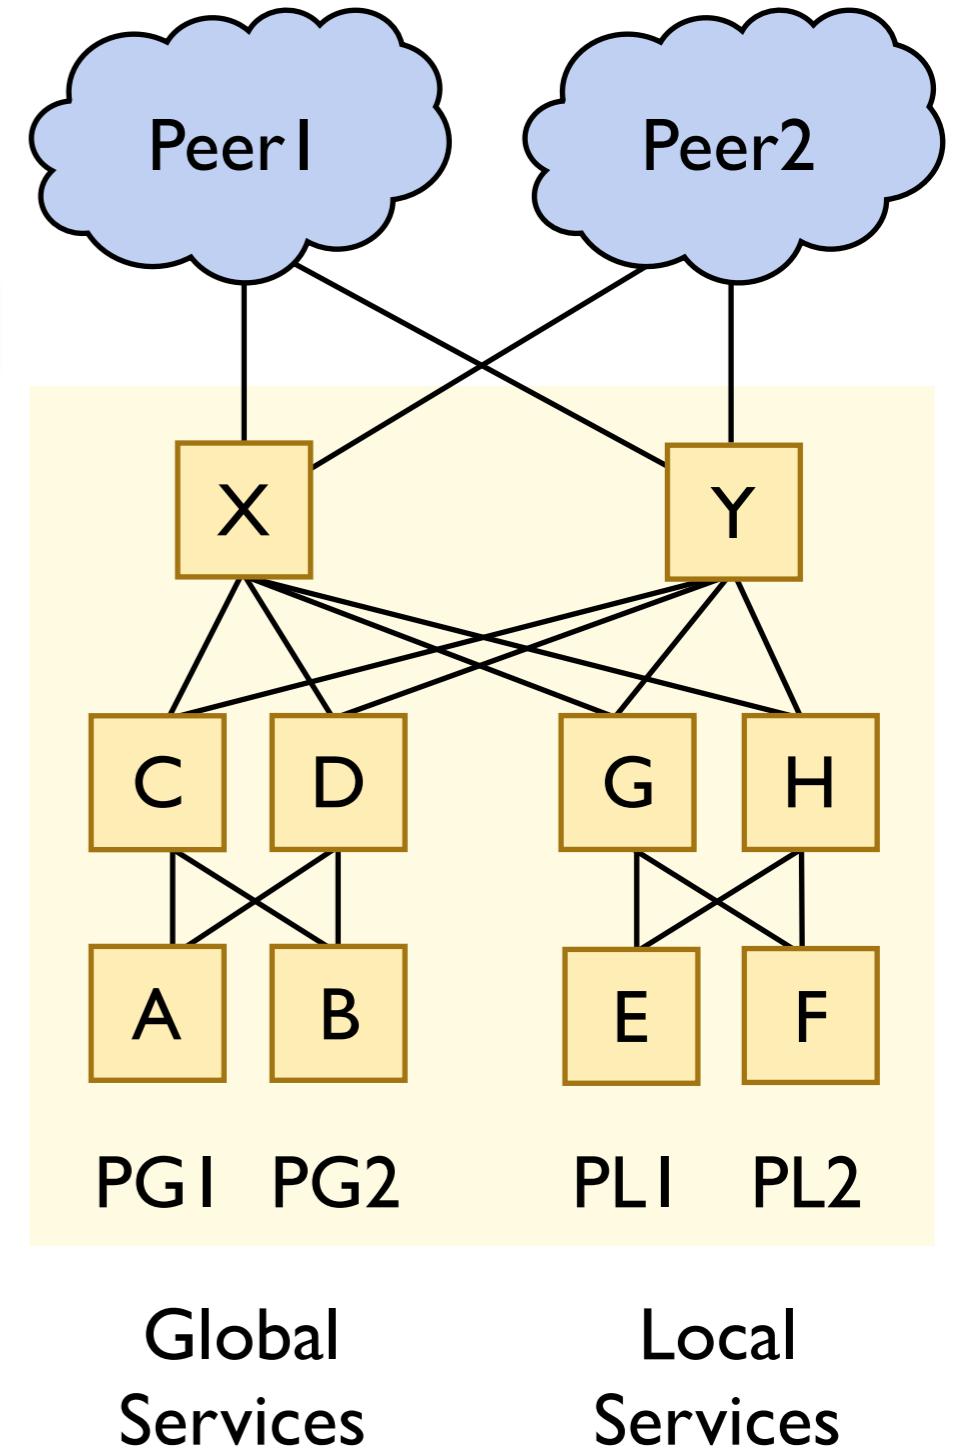
Global
Services

Local
Services

Example 2: A Data Center Network

Prevent transit traffic

```
...
# define in = A | B | C | D | E | F | G | H
define Locality = {
    PL1 or PL2 => always(in)
}
define NoTransit = {
    true => ! transit(Peer,Peer)
}
```

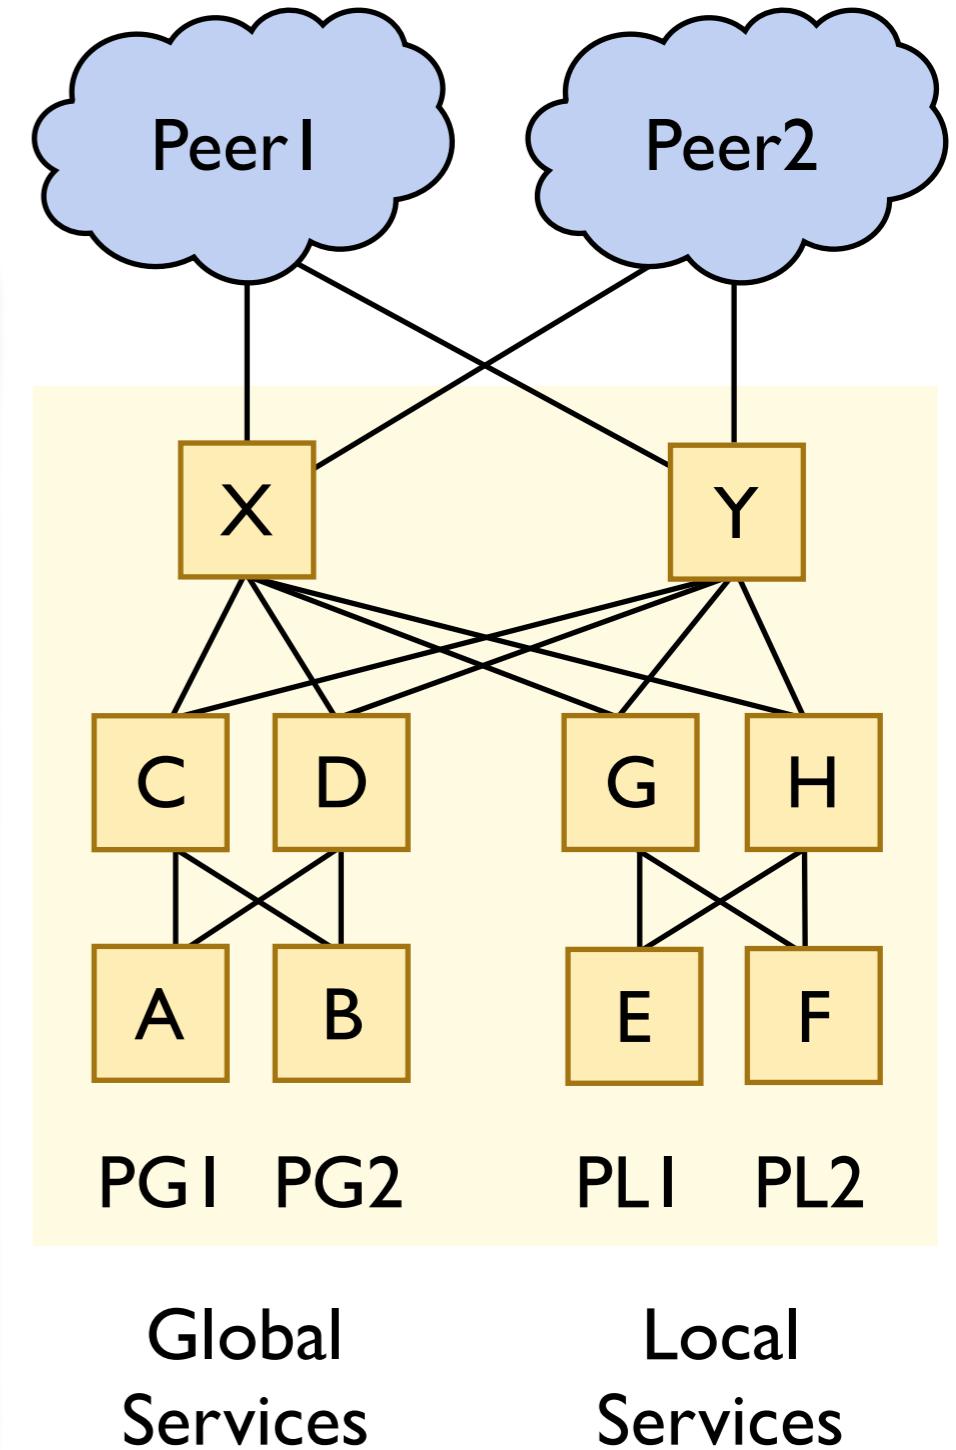


Example 2: A Data Center Network

Prevent transit traffic

...

```
define Main =  
  Ownership & Locality & NoTransit  
  & agg(PG, in → out)
```



Global
Services

Local
Services

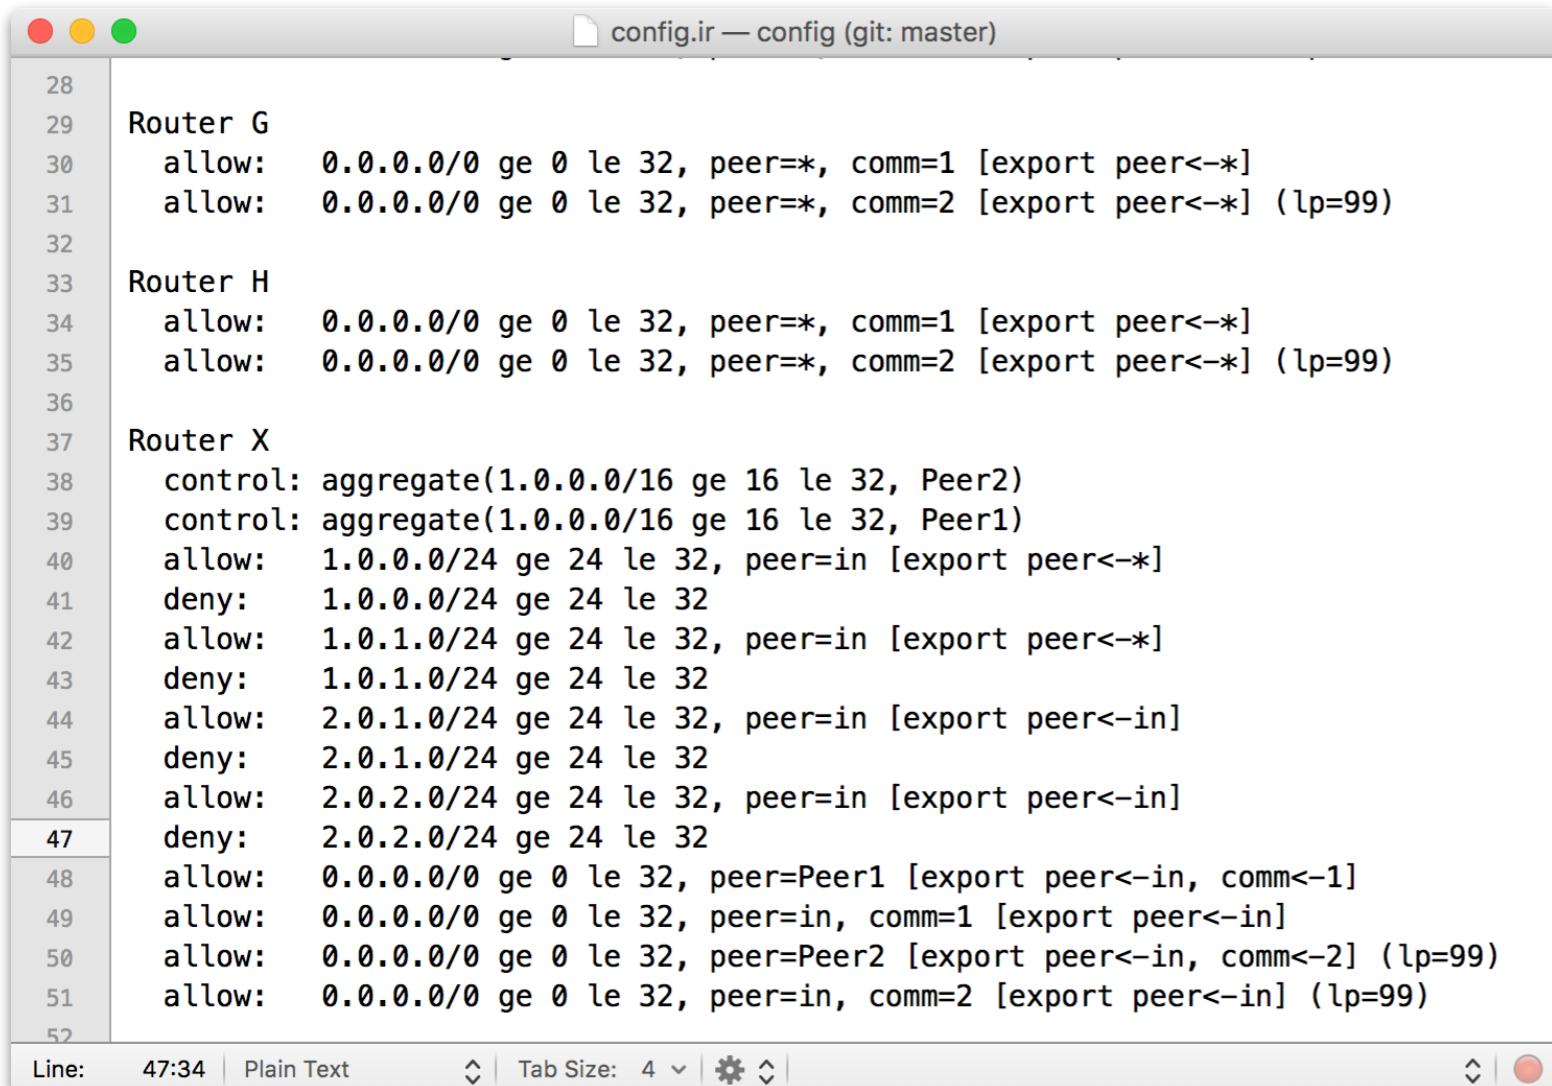
Propane Compiler

The screenshot shows the Propane Compiler interface. On the left is a code editor window titled "dc.pro — dc (git: master)". It contains Propane source code with line numbers from 1 to 31. The code defines various network parameters and routing rules. On the right is a terminal window titled "dc — -bash — 82x21". It shows the command "propane -pol:dc.pro -topo:dc.xml -out:config -parallel:on" being run, followed by a warning message about aggregation black-hole safety, and ends with a prompt "beckett:dc ryanbeckett\$".

```
1 define PL1 = 2.0.1.0/24
2 define PL2 = 2.0.2.0/24
3 define PG1 = 1.0.0.0/24
4 define PG2 = 1.0.1.0/24
5 define PG  = 1.0.0.0/16
6
7 define Peer = {Peer1, Peer2}
8
9 define routing = {
10    PG1 => end(A),
11    PG2 => end(B),
12    PL1 => end(E),
13    PL2 => end(F),
14    true => exit(Peer1) >> exit(Peer2)
15 }
16
17 define locality = {
18    PL1 or PL2 => always(in)
19 }
20
21 define transit(X,Y) = enter(X) and exit(Y)
22
23 define notransit = {
24    true => not transit(Peer, Peer)
25 }
26
27 define main = routing and locality and notransit
28
29 control {
30   aggregate(PG, in -> out)
31 }
```

Line: 31:2 | Propane | Tab Size: 4 | ☰ | ☱ | ☲

Propane Compiler



The screenshot shows a terminal window with the title "config.ir — config (git: master)". The window contains the following text:

```
28
29 Router G
30   allow: 0.0.0.0/0 ge 0 le 32, peer=*, comm=1 [export peer<-*]
31   allow: 0.0.0.0/0 ge 0 le 32, peer=*, comm=2 [export peer<-*] (lp=99)
32
33 Router H
34   allow: 0.0.0.0/0 ge 0 le 32, peer=*, comm=1 [export peer<-*]
35   allow: 0.0.0.0/0 ge 0 le 32, peer=*, comm=2 [export peer<-*] (lp=99)
36
37 Router X
38   control: aggregate(1.0.0.0/16 ge 16 le 32, Peer2)
39   control: aggregate(1.0.0.0/16 ge 16 le 32, Peer1)
40   allow: 1.0.0.0/24 ge 24 le 32, peer=in [export peer<-*]
41   deny: 1.0.0.0/24 ge 24 le 32
42   allow: 1.0.1.0/24 ge 24 le 32, peer=in [export peer<-*]
43   deny: 1.0.1.0/24 ge 24 le 32
44   allow: 2.0.1.0/24 ge 24 le 32, peer=in [export peer<-in]
45   deny: 2.0.1.0/24 ge 24 le 32
46   allow: 2.0.2.0/24 ge 24 le 32, peer=in [export peer<-in]
47   deny: 2.0.2.0/24 ge 24 le 32
48   allow: 0.0.0.0/0 ge 0 le 32, peer=Peer1 [export peer<-in, comm<-1]
49   allow: 0.0.0.0/0 ge 0 le 32, peer=in, comm=1 [export peer<-in]
50   allow: 0.0.0.0/0 ge 0 le 32, peer=Peer2 [export peer<-in, comm<-2] (lp=99)
51   allow: 0.0.0.0/0 ge 0 le 32, peer=in, comm=2 [export peer<-in] (lp=99)
52
```

At the bottom of the terminal window, there are status indicators: "Line: 47:34 | Plain Text" and "Tab Size: 4".

Propane Compiler

The image shows two terminal windows side-by-side. The left window, titled 'config.ir — config (git: master)', displays configuration rules for three routers: Router G, Router H, and Router X. Router G and Router H have identical 'allow' rules. Router X has more complex rules involving aggregates, communities, and AS-path access lists. The right window, titled 'as109.quagga — quagga (git: master)', shows the generated Quagga configuration file 'as109.conf'. It includes route maps for exporting routes to AS 109, defining various prefix and community lists, and applying these to specific routes based on their origin and AS-path.

```
Router G
  allow: 0.0.0.0/0 ge 0 le 32, peer=
  allow: 0.0.0.0/0 ge 0 le 32, peer=

Router H
  allow: 0.0.0.0/0 ge 0 le 32, peer=
  allow: 0.0.0.0/0 ge 0 le 32, peer=

Router X
  control: aggregate(1.0.0.0/16 ge 16
  control: aggregate(1.0.0.0/16 ge 16
  allow: 1.0.0.0/24 ge 24 le 32, peer=
  deny: 1.0.0.0/24 ge 24 le 32
  allow: 1.0.1.0/24 ge 24 le 32, peer=
  deny: 1.0.1.0/24 ge 24 le 32
  allow: 2.0.1.0/24 ge 24 le 32, peer=
  deny: 2.0.1.0/24 ge 24 le 32
  allow: 2.0.2.0/24 ge 24 le 32, peer=
  deny: 2.0.2.0/24 ge 24 le 32
  allow: 0.0.0.0/0 ge 0 le 32, peer=
  allow: 0.0.0.0/0 ge 0 le 32, peer=
  allow: 0.0.0.0/0 ge 0 le 32, peer=
  allow: 0.0.0.0/0 ge 0 le 32, peer=
```

```
neighbor 12.12.12.1 route-map rm-export-1 out
!
ip prefix-list pl-1 permit 1.0.0.0/24 ge 24 le 32
ip prefix-list pl-2 deny 1.0.0.0/24 ge 24 le 32
ip prefix-list pl-3 permit 1.0.1.0/24 ge 24 le 32
ip prefix-list pl-4 deny 1.0.1.0/24 ge 24 le 32
ip prefix-list pl-5 permit 2.0.1.0/24 ge 24 le 32
ip prefix-list pl-6 deny 2.0.1.0/24 ge 24 le 32
ip prefix-list pl-7 permit 2.0.2.0/24 ge 24 le 32
ip prefix-list pl-8 deny 2.0.2.0/24 ge 24 le 32
ip prefix-list pl-9 permit 0.0.0.0/0 ge 0 le 32
!
ip community-list standard cl-1 permit 100:1
ip community-list standard cl-2 permit 100:2
ip community-list standard cl-3 permit 100:3
ip community-list standard cl-4 permit 100:4
!
ip as-path access-list path-1 permit (^107_ | ^106_ | ^103_ | ^102_)
ip as-path access-list path-2 permit (^111_ | ^110_)
ip as-path access-list path-3 permit ^110_
ip as-path access-list path-4 permit ^111_
!
route-map rm-in permit 10
  match ip address prefix-list pl-1
  match as-path path-1
    set community additive 200:1
  !
route-map rm-in permit 20
  match ip address prefix-list pl-2
  !
route-map rm-in permit 30
  match ip address prefix-list pl-3
  match as-path path-1
    set community additive 200:1
  !
```

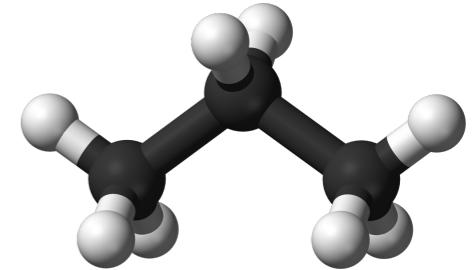
Propane Compiler

```
ryanbeckett — bash — 86x24
[beckett:~ ryanbeckett$ propane

Usage: propane [options]
-pol          Policy file
-topo         Topology file
-out          Output file
-failures:anyIn Failure safety for aggregation (default any)
-anycast:onloff Allow anycast (default off)
-med:onloff   Use MED attribute (default off)
-prepending:onloff Use AS path prepending (default off)
-no-export:onloff Use no-export community (default off)
-ibgp:onloff  Assume IBGP configuration (default off)
-minimize:onloff Minimize configuration (default on)
-parallel:onloff Parallelize compilation (default on)
-target:none|irl|templ Compilation target
-stats:onloff  Display performance statistics to stdout (default off)
-checkenter:onloff Check traffic entry conditions to network
-debug:on:off   Log/Save Debugging information (default off)
-debug-dir     Debugging directory (default 'debug')
-test          Run unit tests
-bench         Generate benchmark files
-help          Display this message

beckett:~ ryanbeckett$ ]
```

Propane: Summary



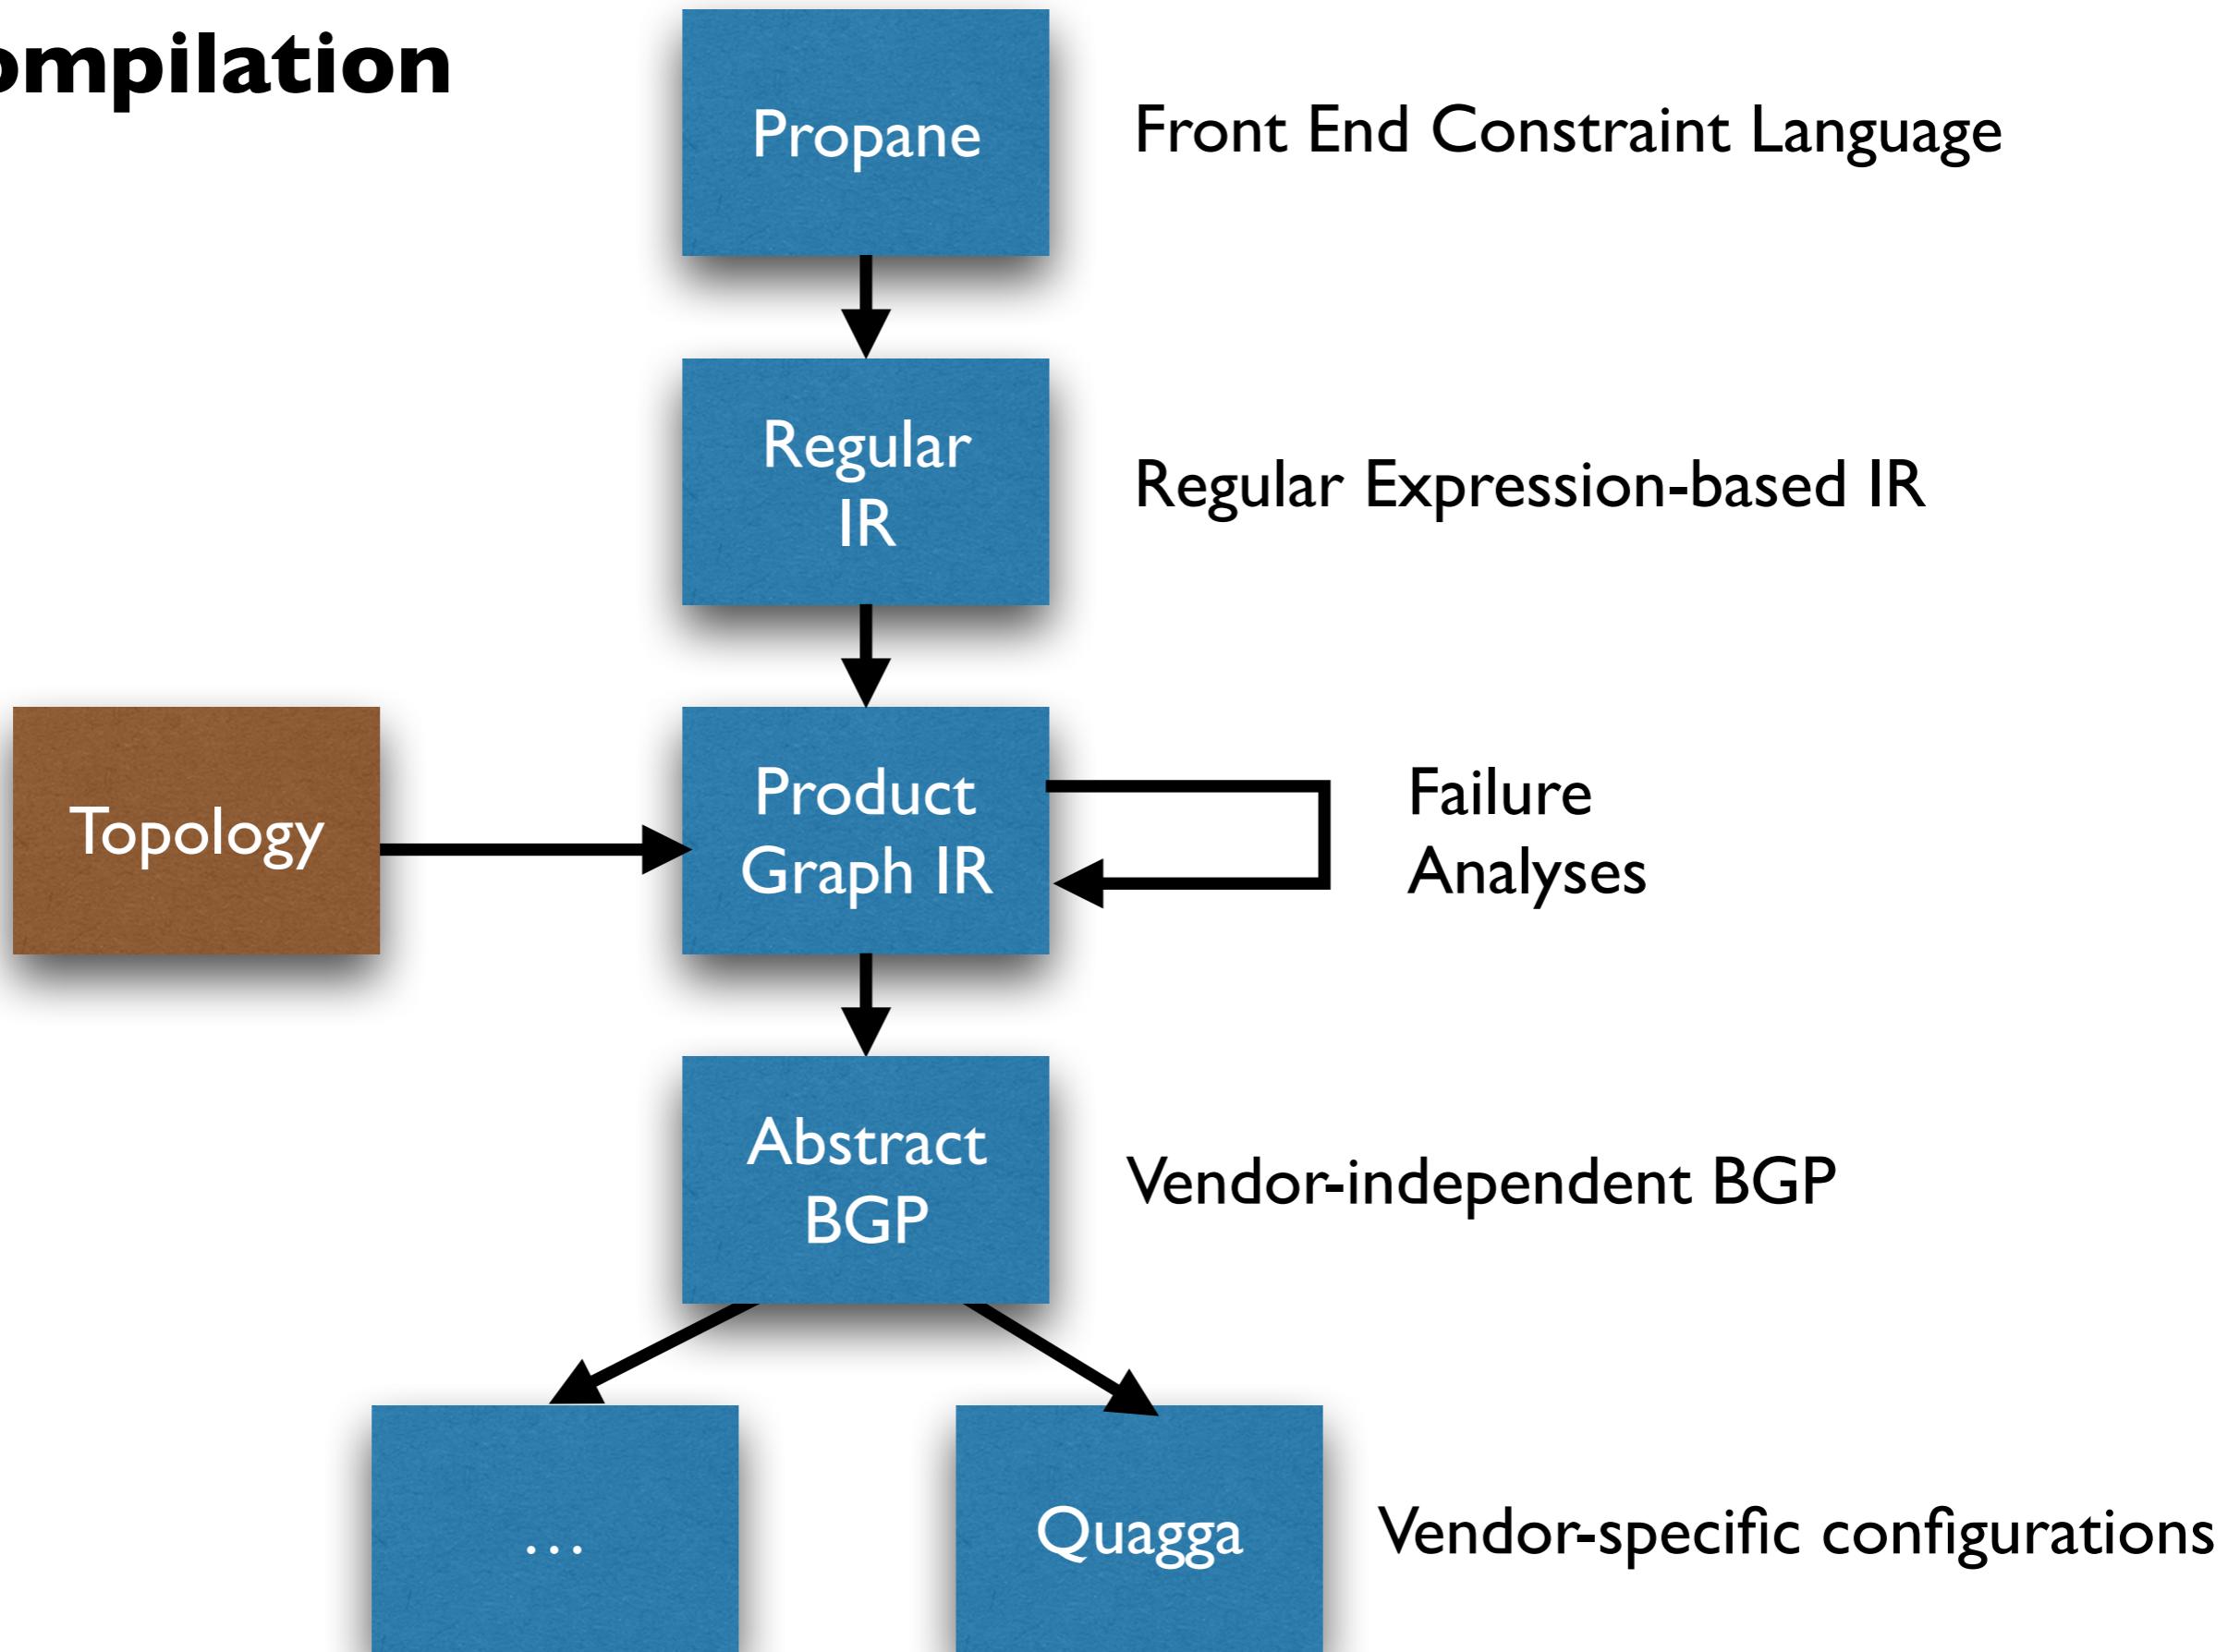
High-level language

- Single, network-wide policy
- Program using constraints and preferences
- Create new, reusable abstractions (e.g., transit traffic)
- Configure the network modularly

Compiler

- Automatically generates filters, preferences, community values, etc
- Static analysis guarantees policy compliance under all failures
- Lower bound on failures for aggregation-induced black holes

Compilation



Propane Regular IR

Propane

Step I: Combine modular constraints

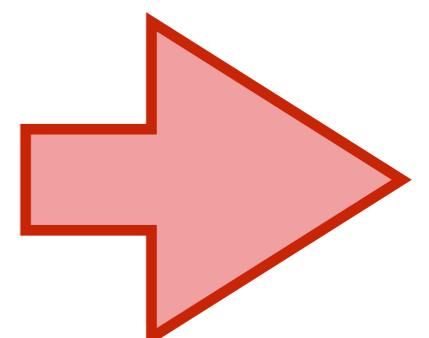
```
define Ownership =
{ PG1 => end(A)
  PG2 => end(B)
  PL1 => end(E)
  PL2 => end(F)
  true => exit(Peer1 >> Peer2) }
```

```
define NoTransit =
{ true => !transit(Peer, Peer) }
```

```
define Locality =
{ PL1 | PL2 => always(in) }
```

```
define Main =
  Ownership & Locality & NoTransit
  & agg(PG, in -> out)
```

Regular
IR



Propane Regular IR

Propane



Regular
IR

Prefix-by-prefix intersection of constraints

```
PG1 => !transit(Peer, Peer) & end(A)
PG2 => !transit(Peer, Peer) & end(B)
PL1 => !transit(Peer, Peer) & always(in) & end(E)
PL2 => !transit(Peer, Peer) & always(in) & end(F)
true => !transit(Peer, Peer) & exit(Peer1 >> Peer2)
```

Propane Regular IR

Propane

Step 2: Expand constraints in to regular expressions

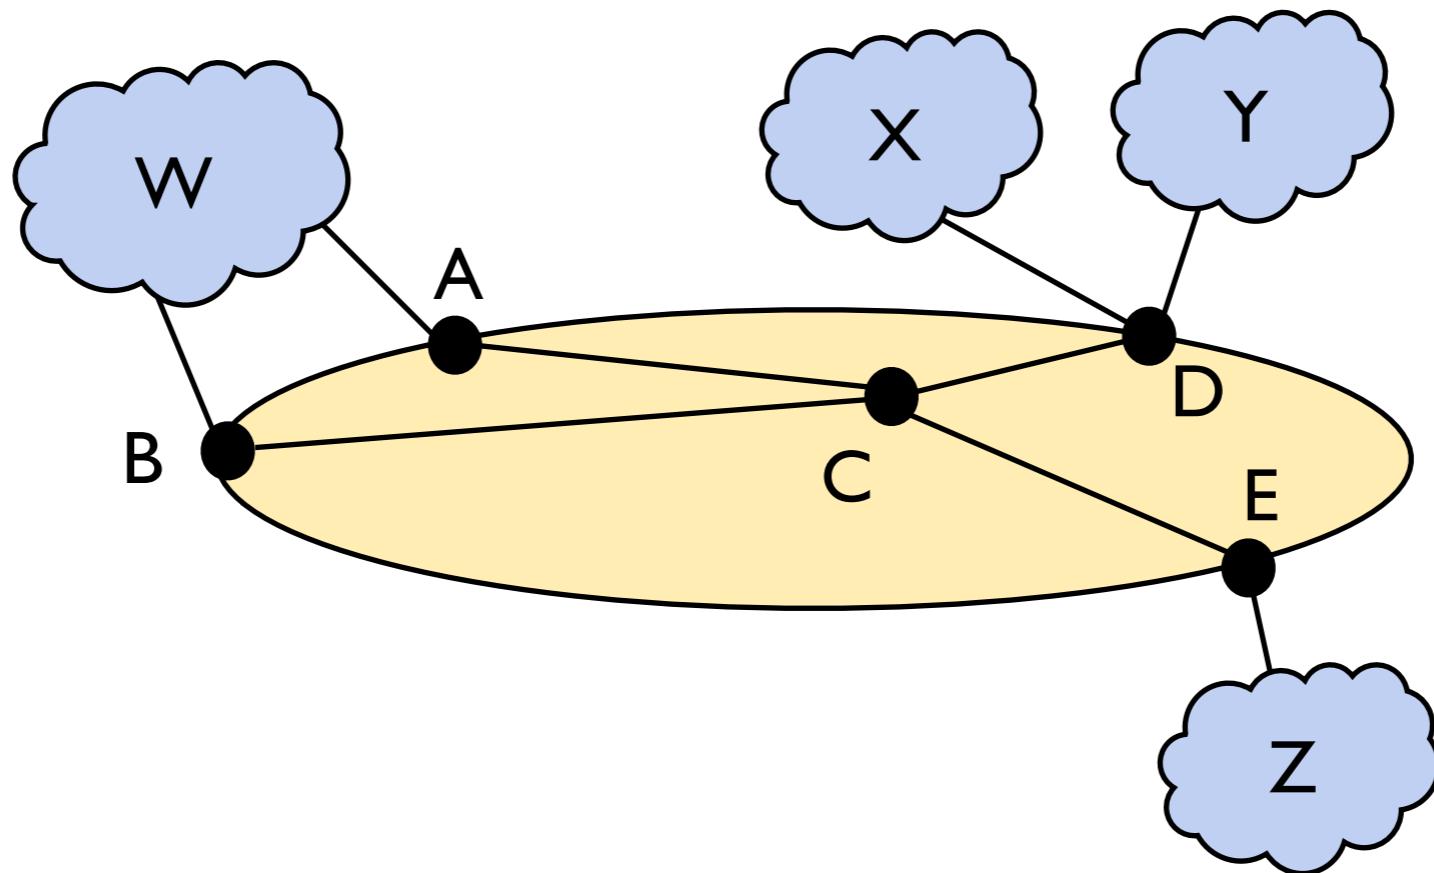
```
any = out*.in+.out*
end(X) = any ∩ (Σ*.X)
always(X) = any ∩ (X)*
exit(X) = (out*.in*. (X ∩ in).out+) |
           (out*.in+. (X ∩ out).out*)
end(X) = ...
avoid(X) = ...
waypoint(X) = ...
```

Regular
IR

Step 3: Reduced syntax

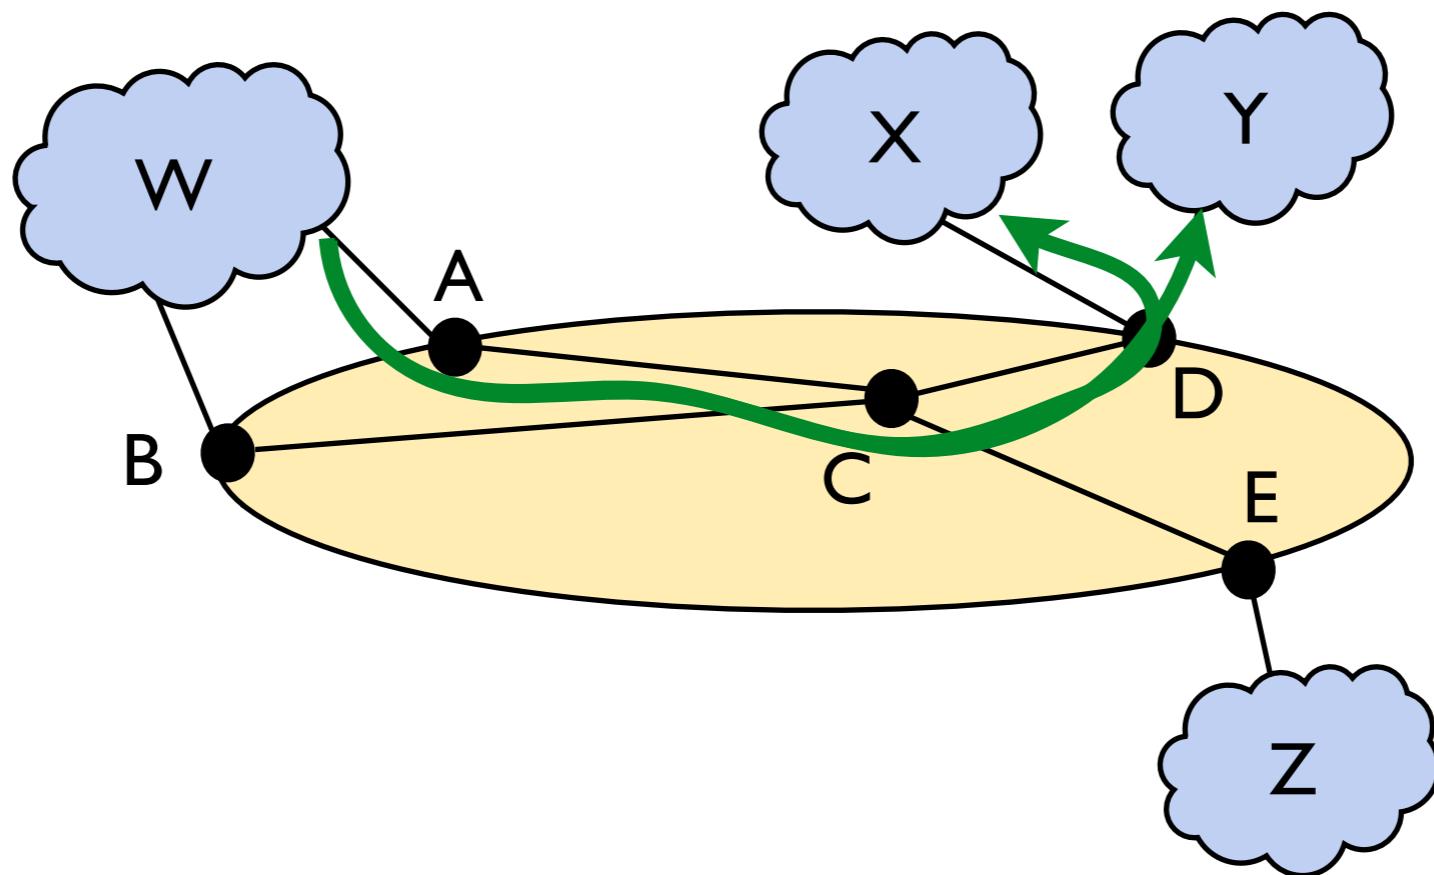
```
true => A.(X >> Y).out*
true => (A.X.out*) >> (A.Y.out*)
```

Compilation: An Idealized Example



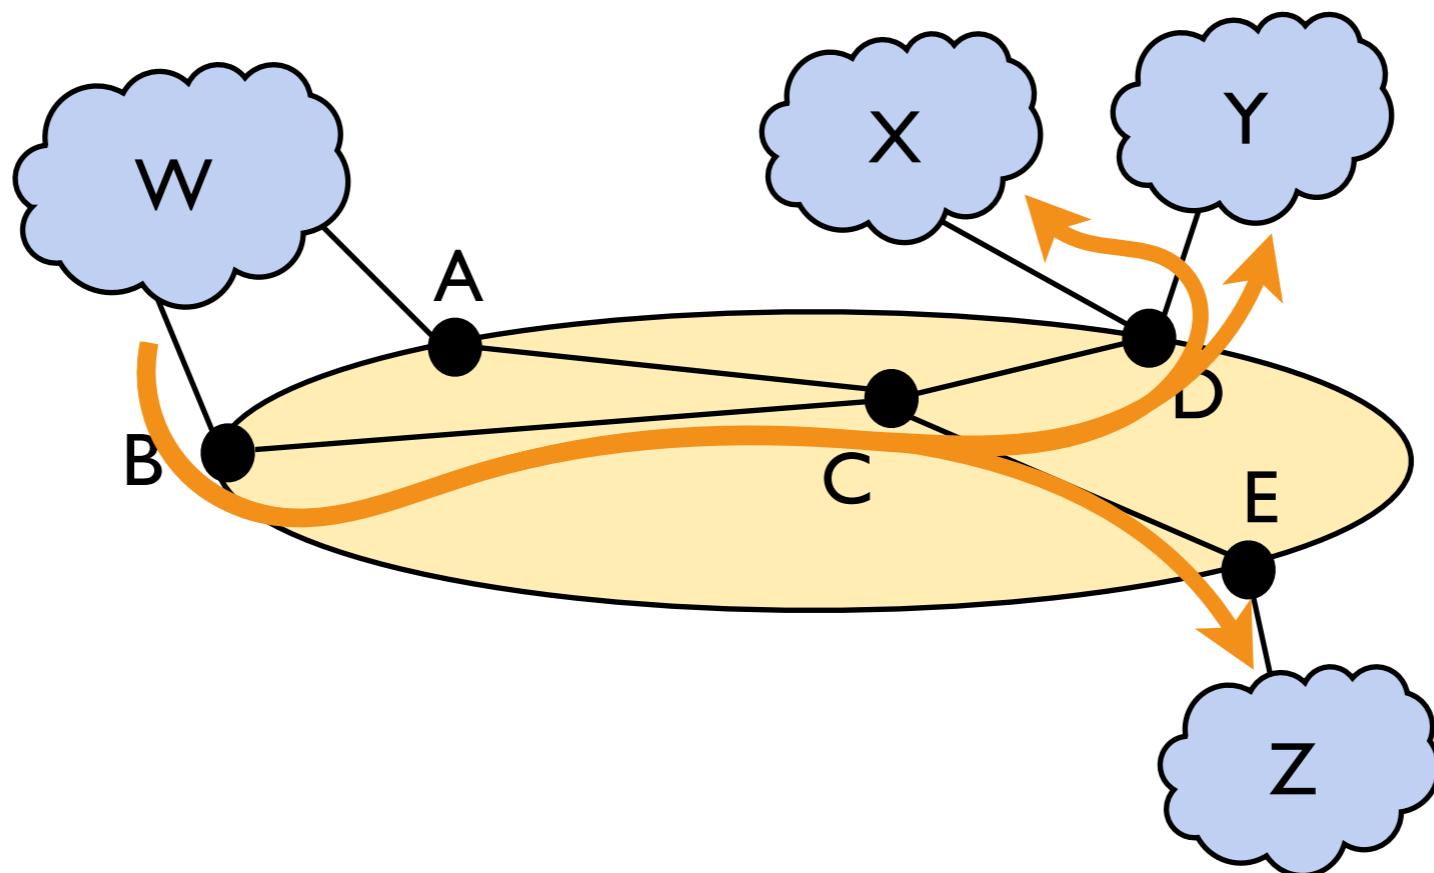
Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

Compilation: An Idealized Example



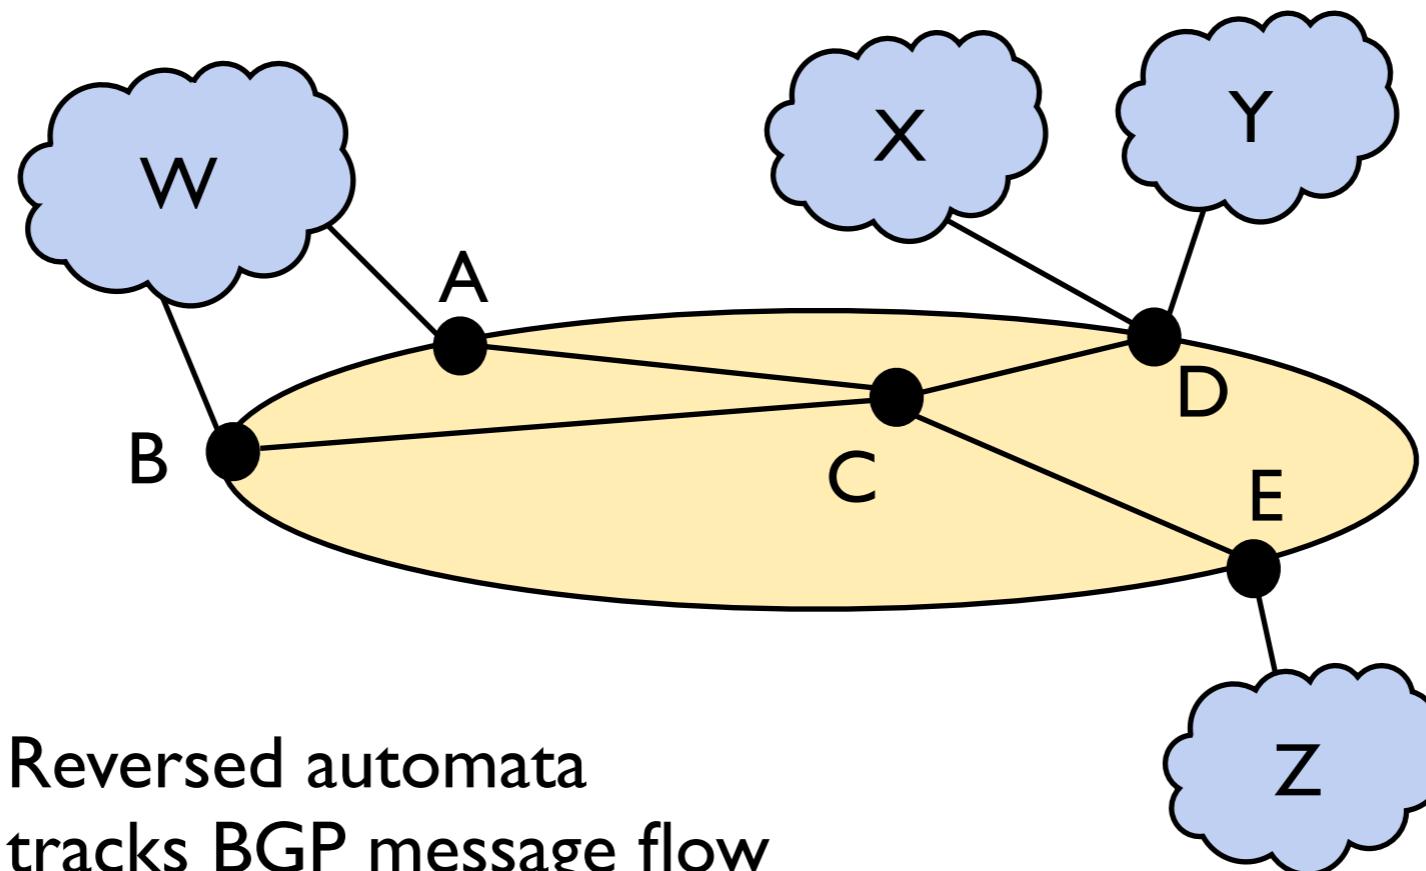
Policy: `(W.A.C.D.out) >> (W.B.in+.out)`

Compilation: An Idealized Example



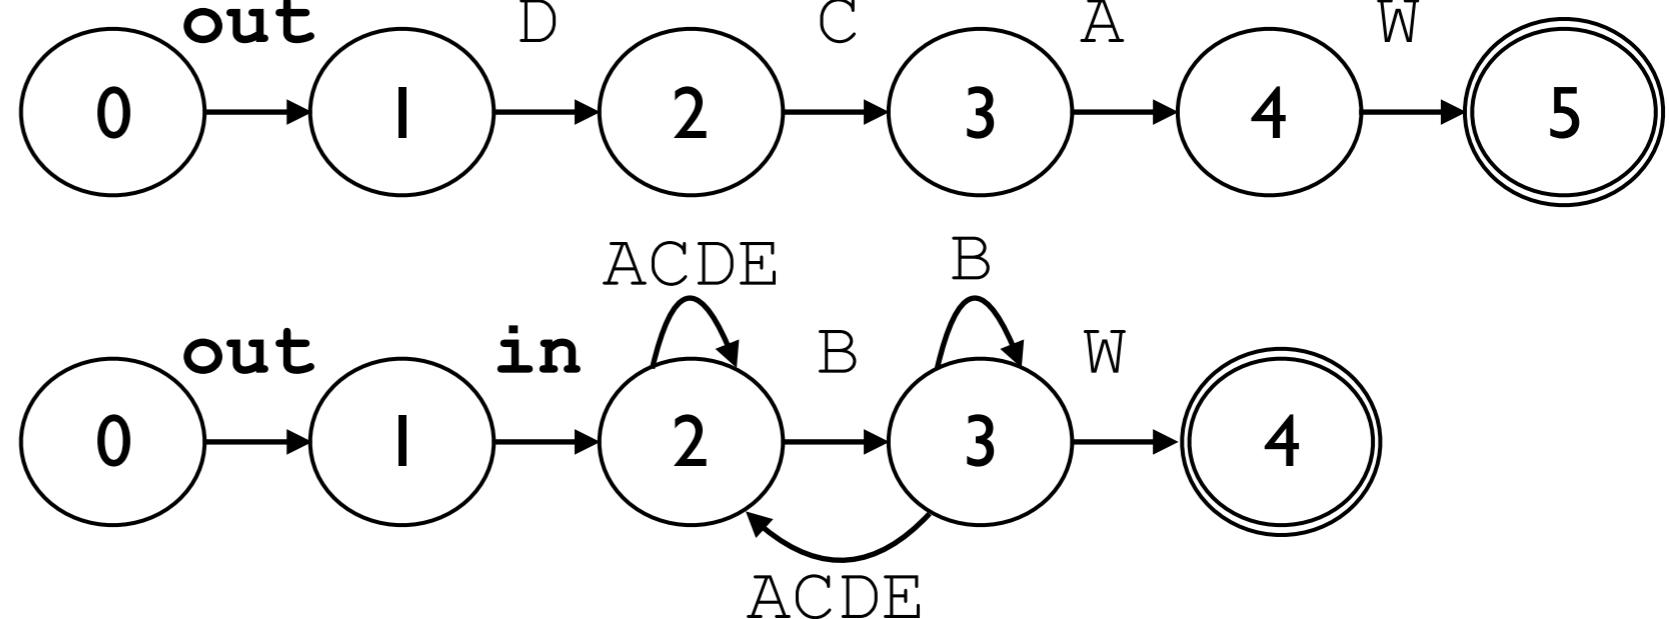
Policy: $(W.A.C.D.out) \gg (W.B.in+out)$

Reversed Automata from Policies



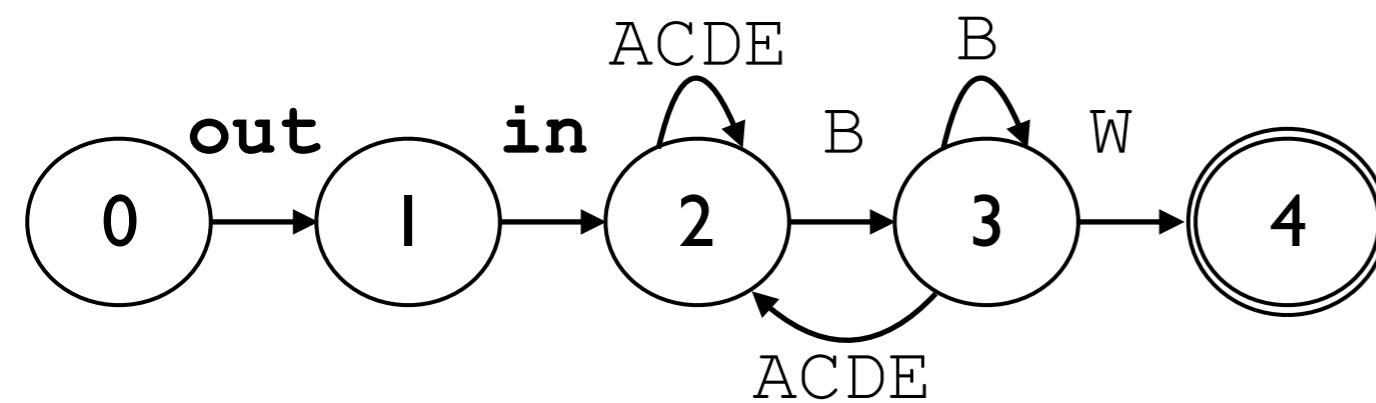
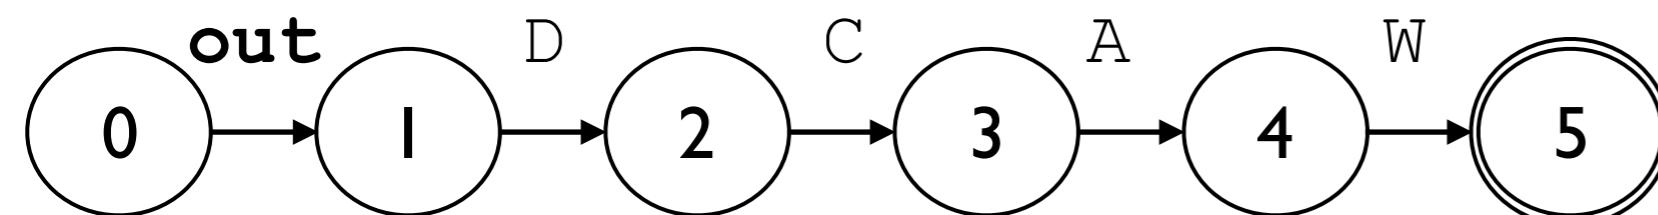
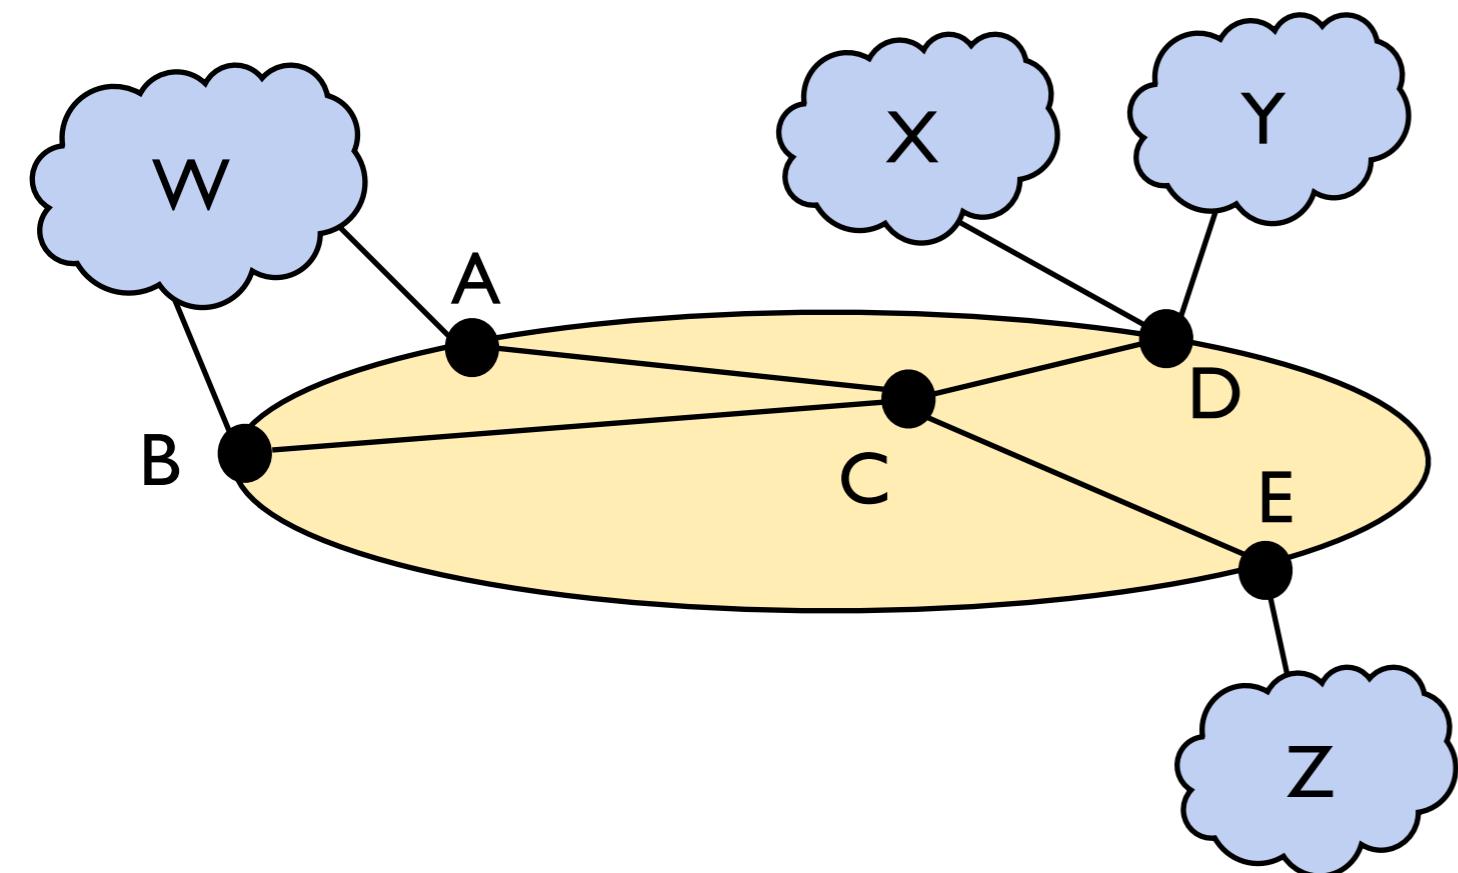
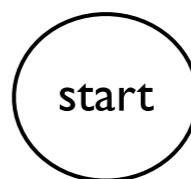
Policy:

1. (W.A.C.D.out)
2. (W.B.in+.out)



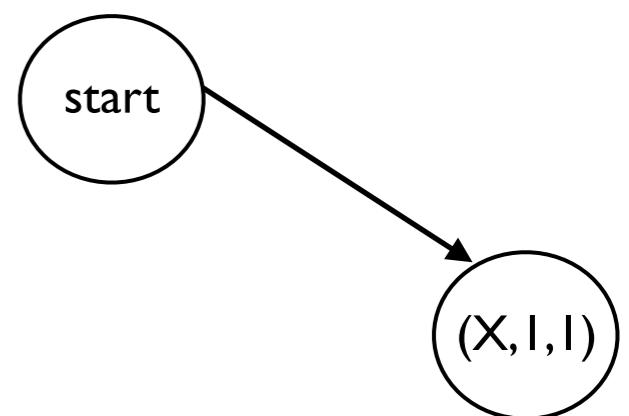
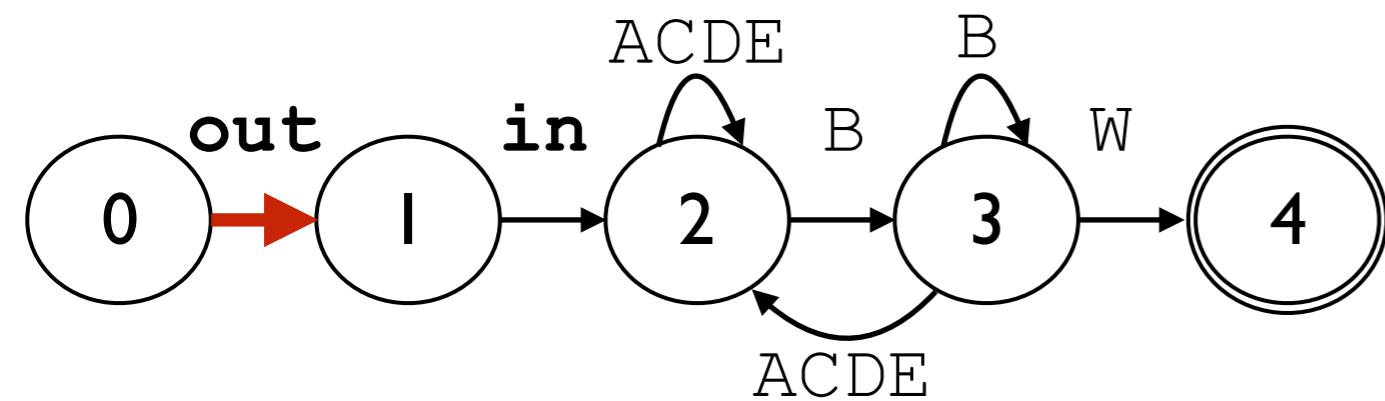
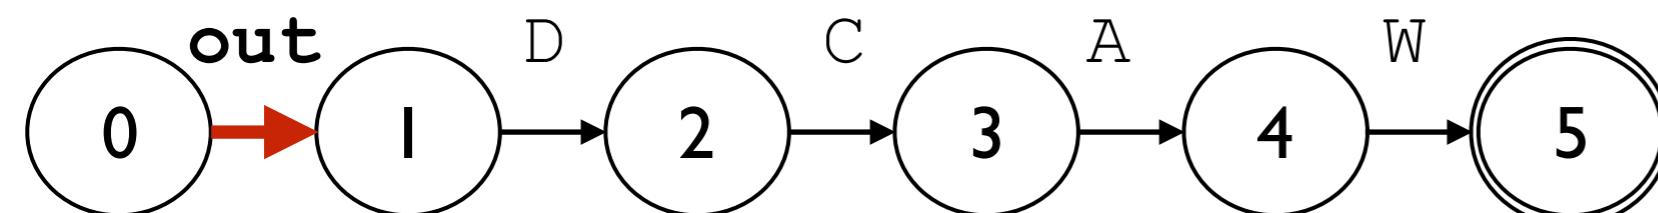
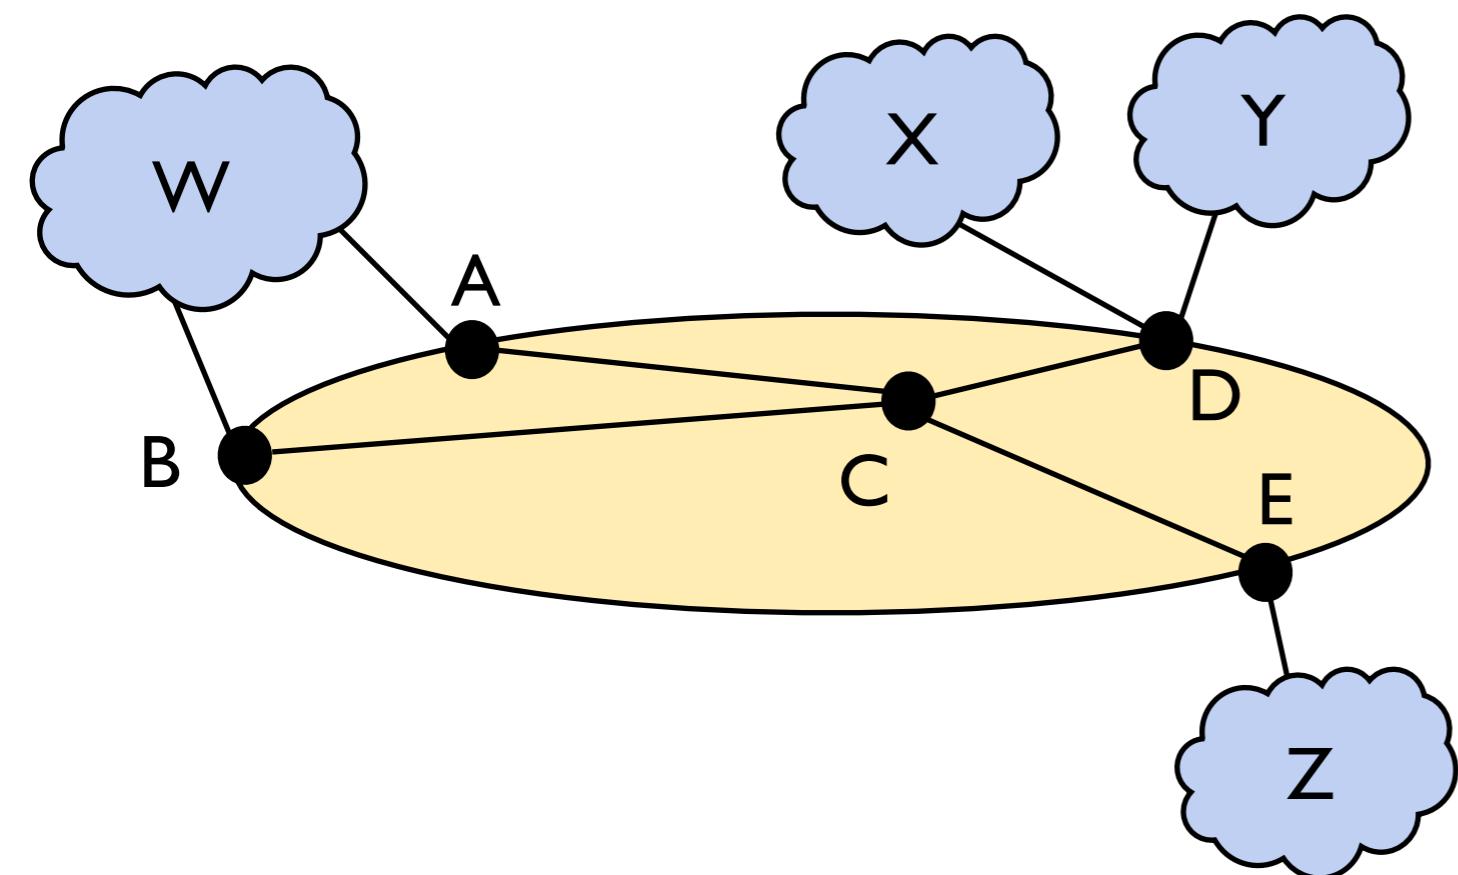
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



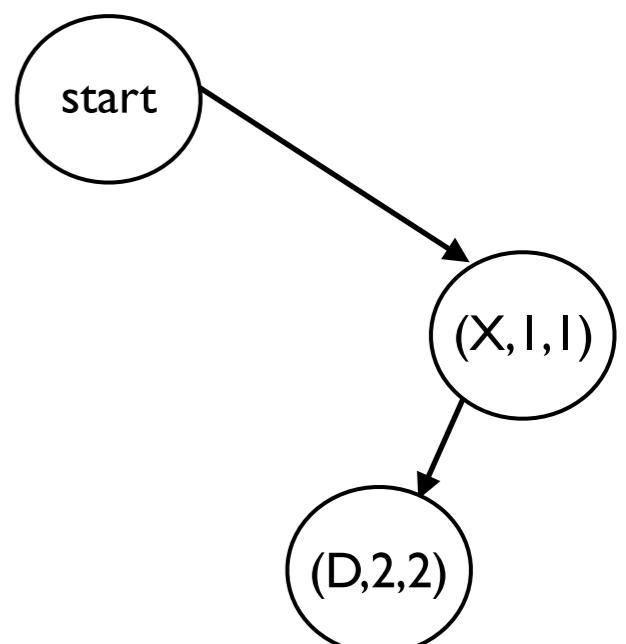
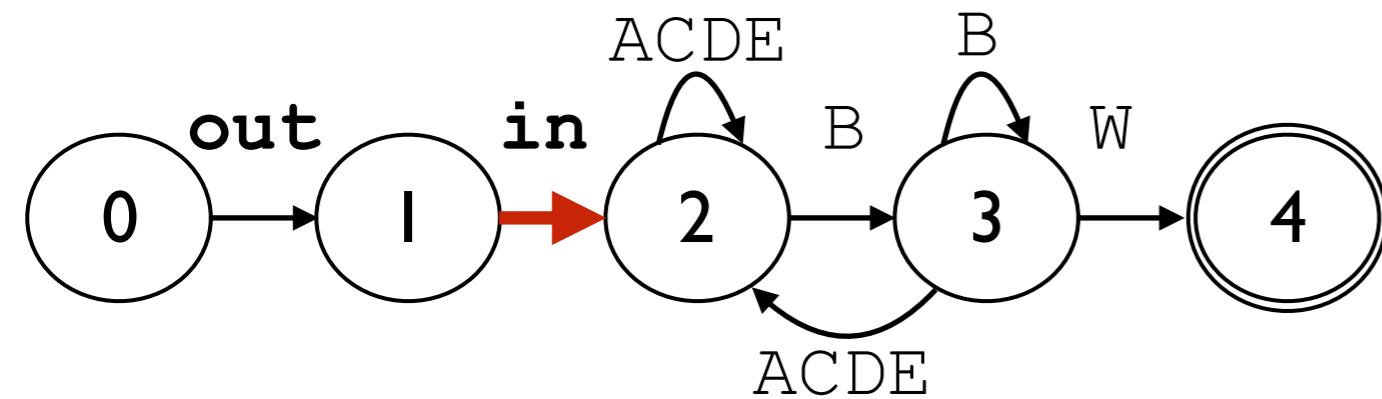
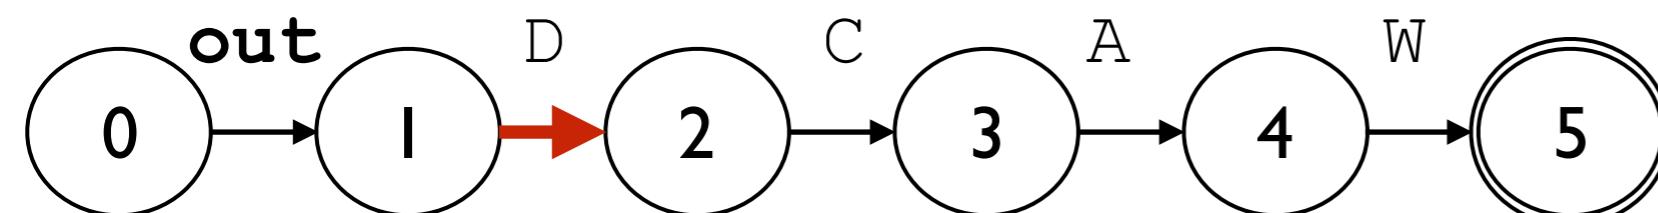
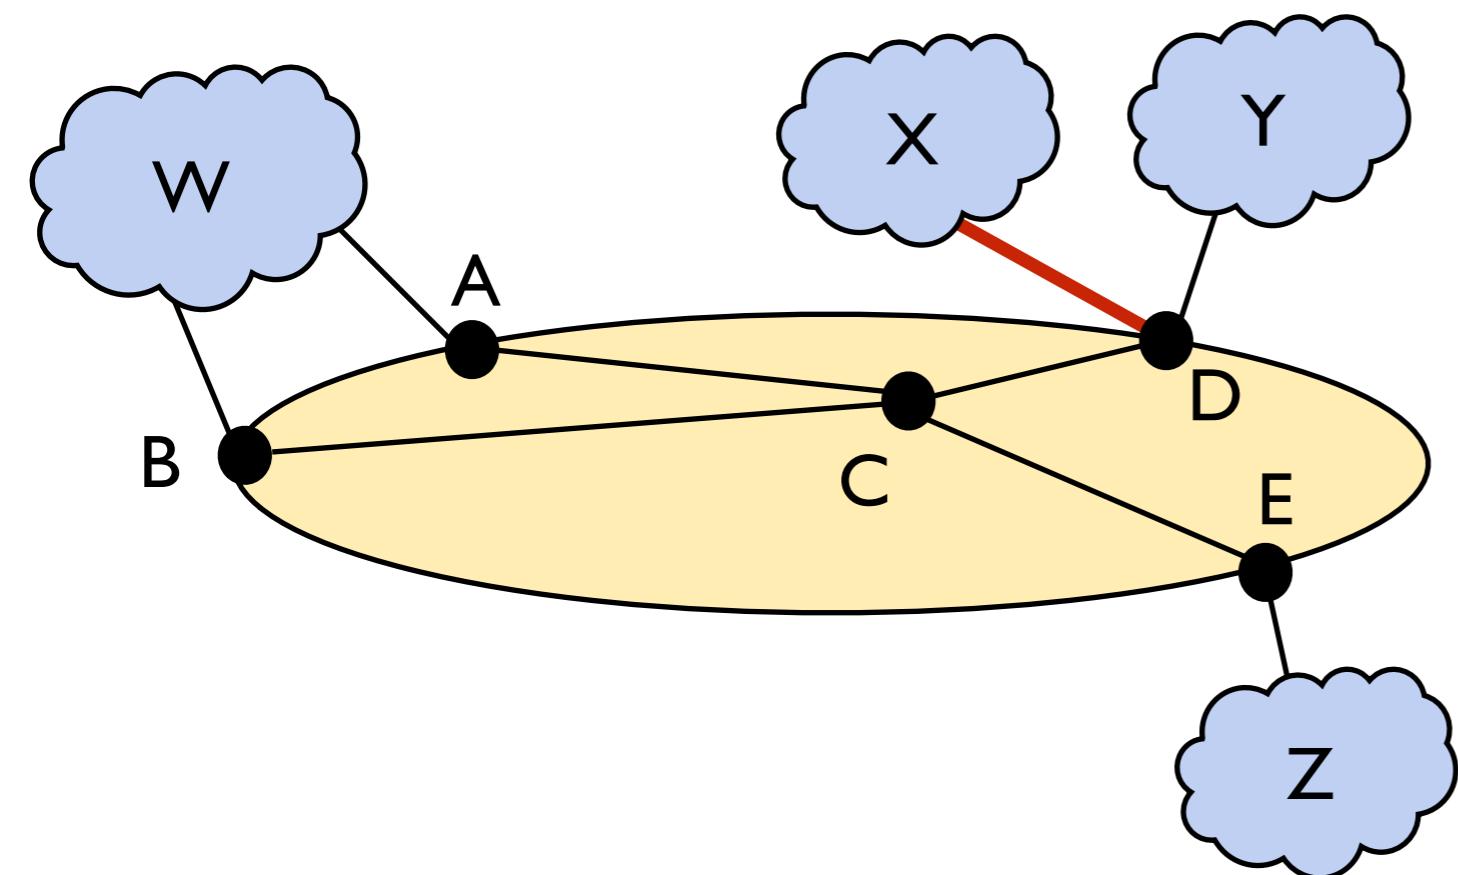
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



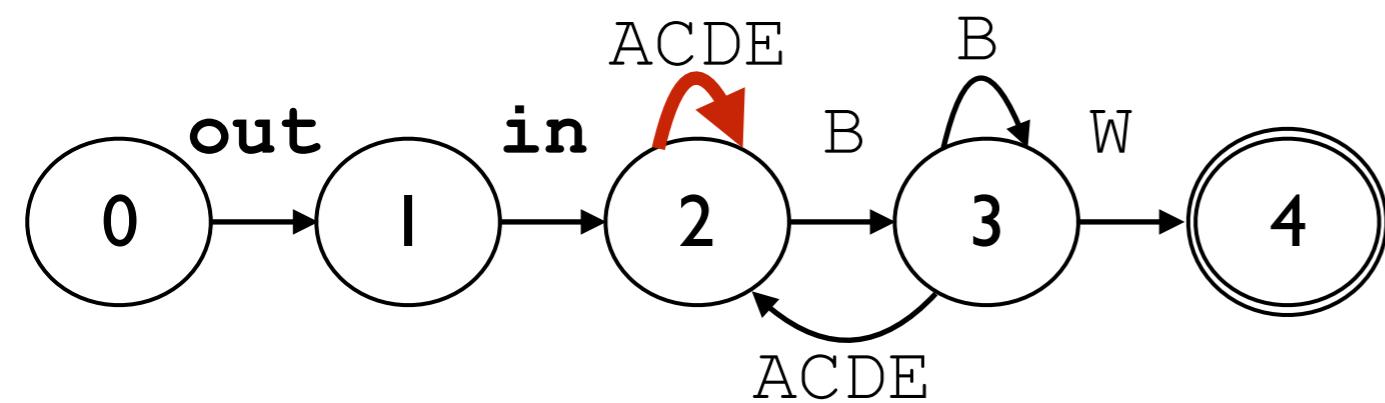
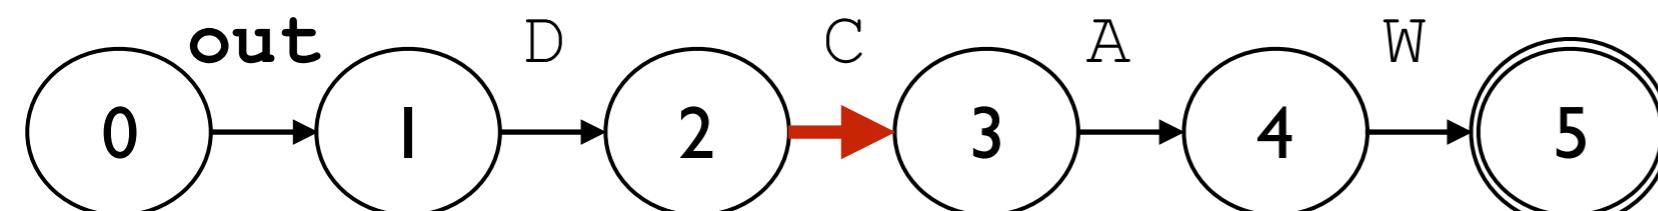
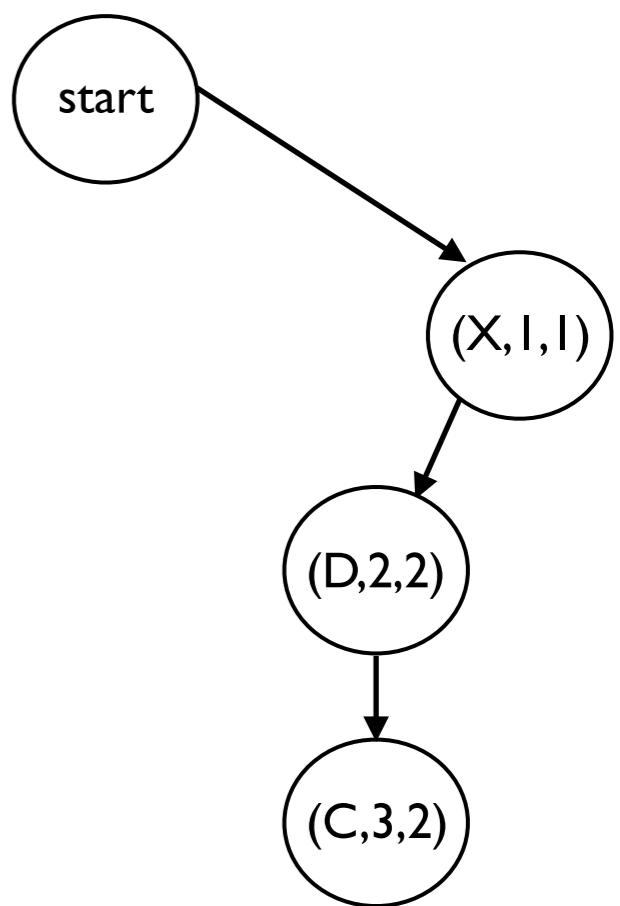
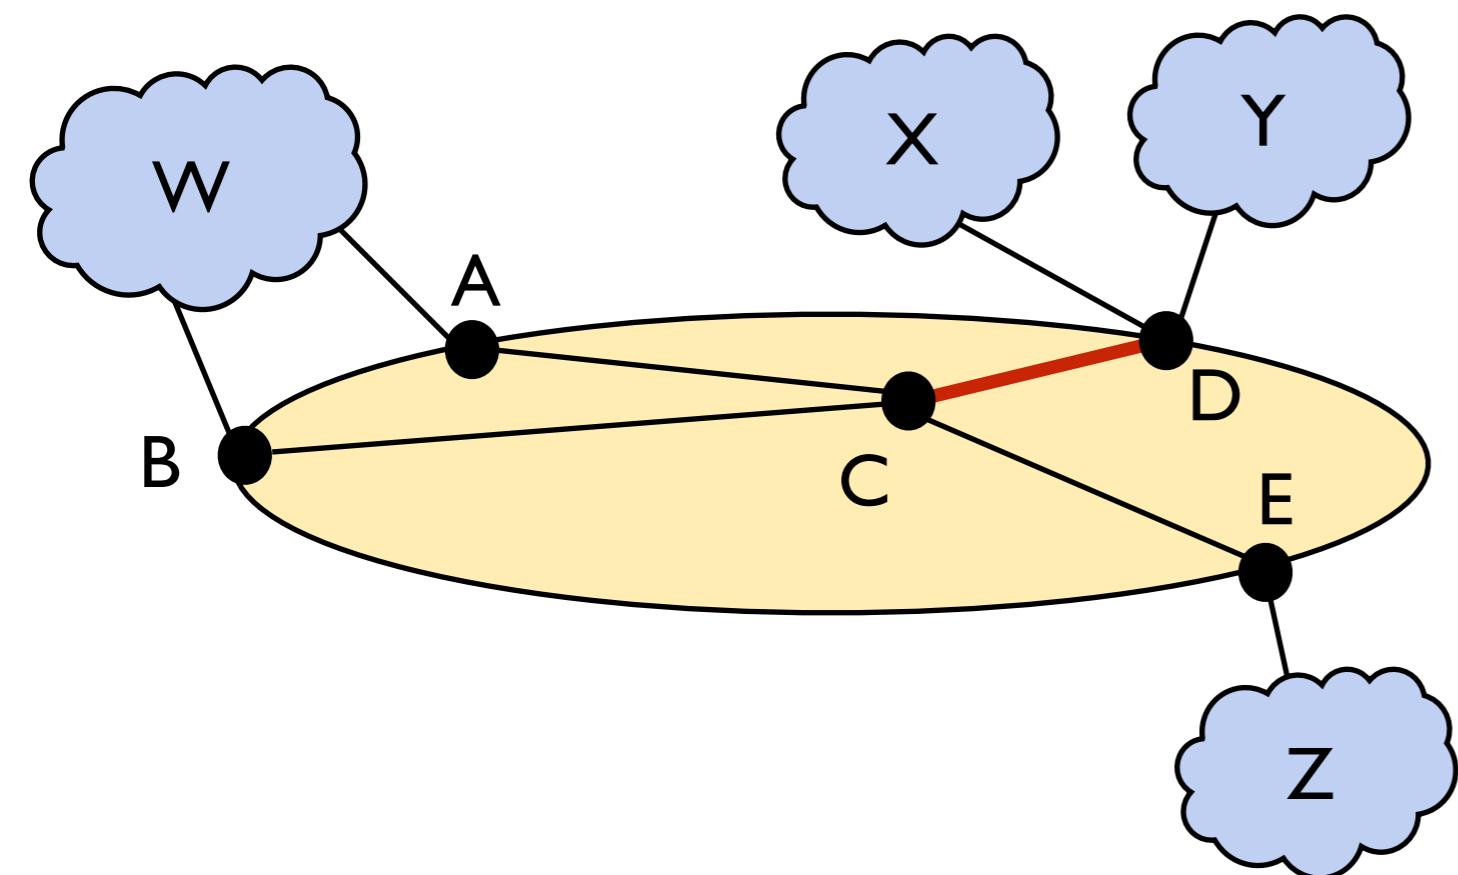
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



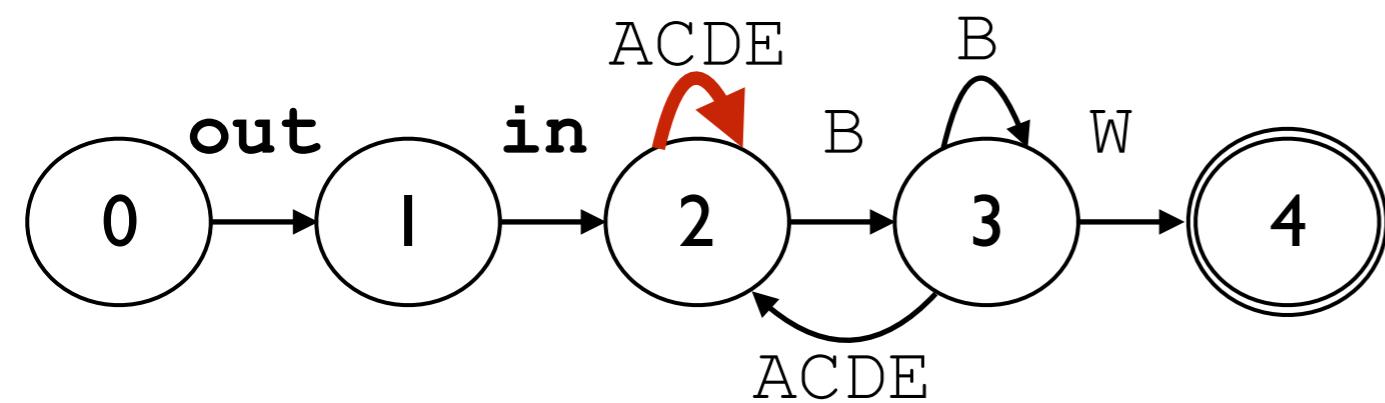
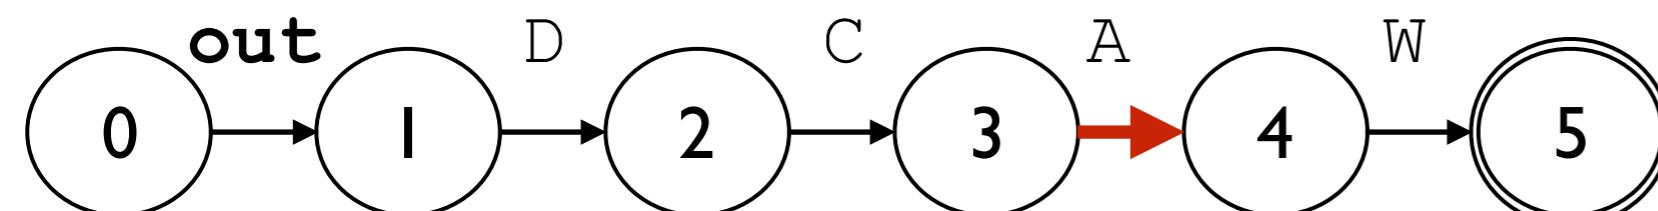
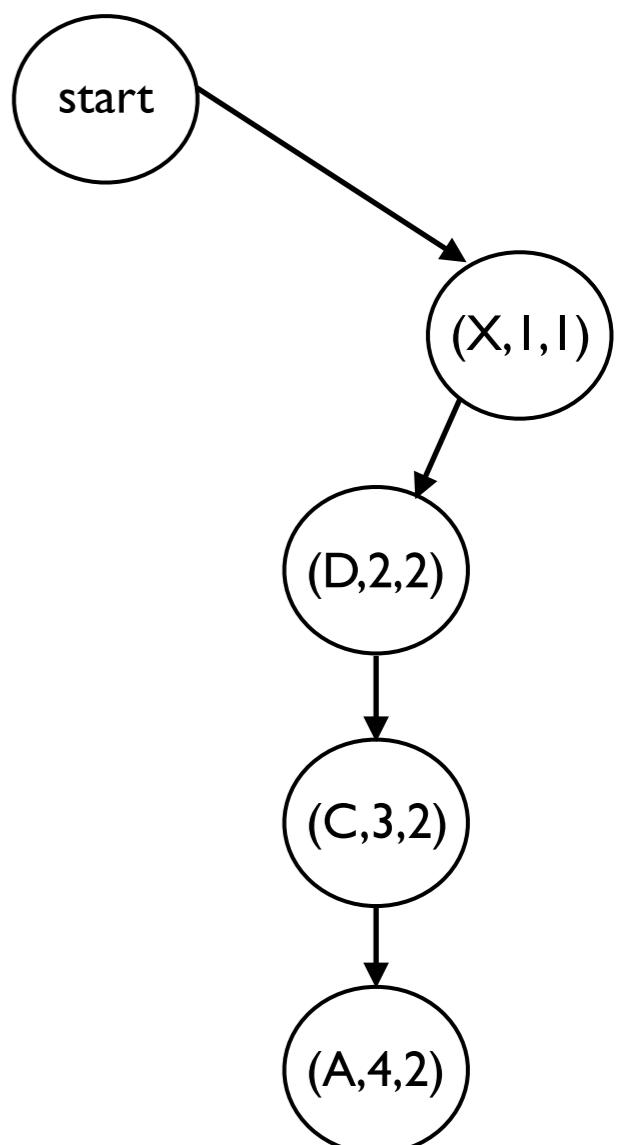
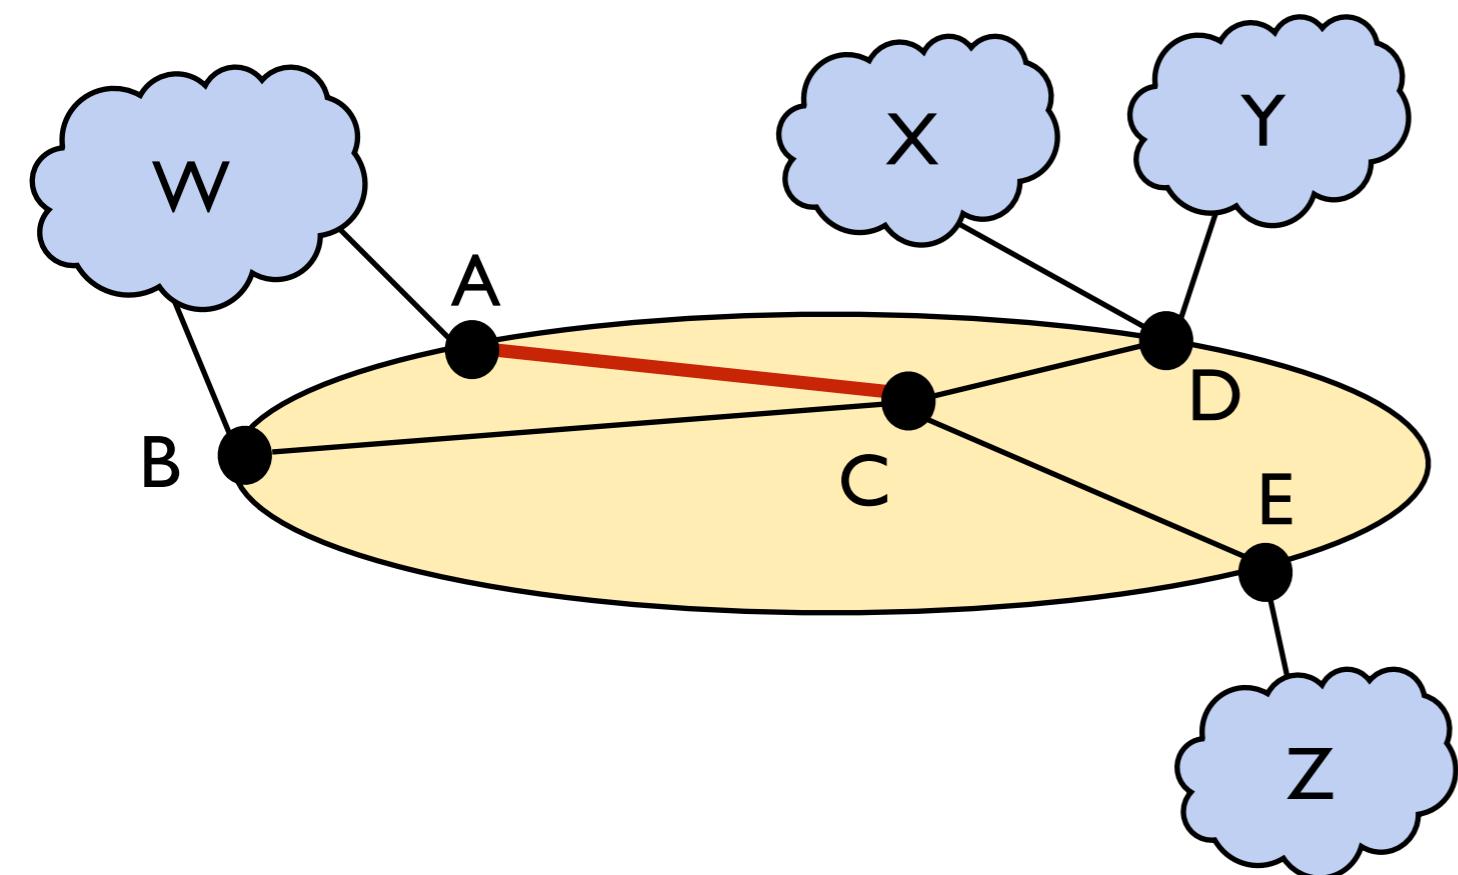
Constructing the Product Graph (PG)

(W.A.C.D.out) >> (W.B.in+.out)



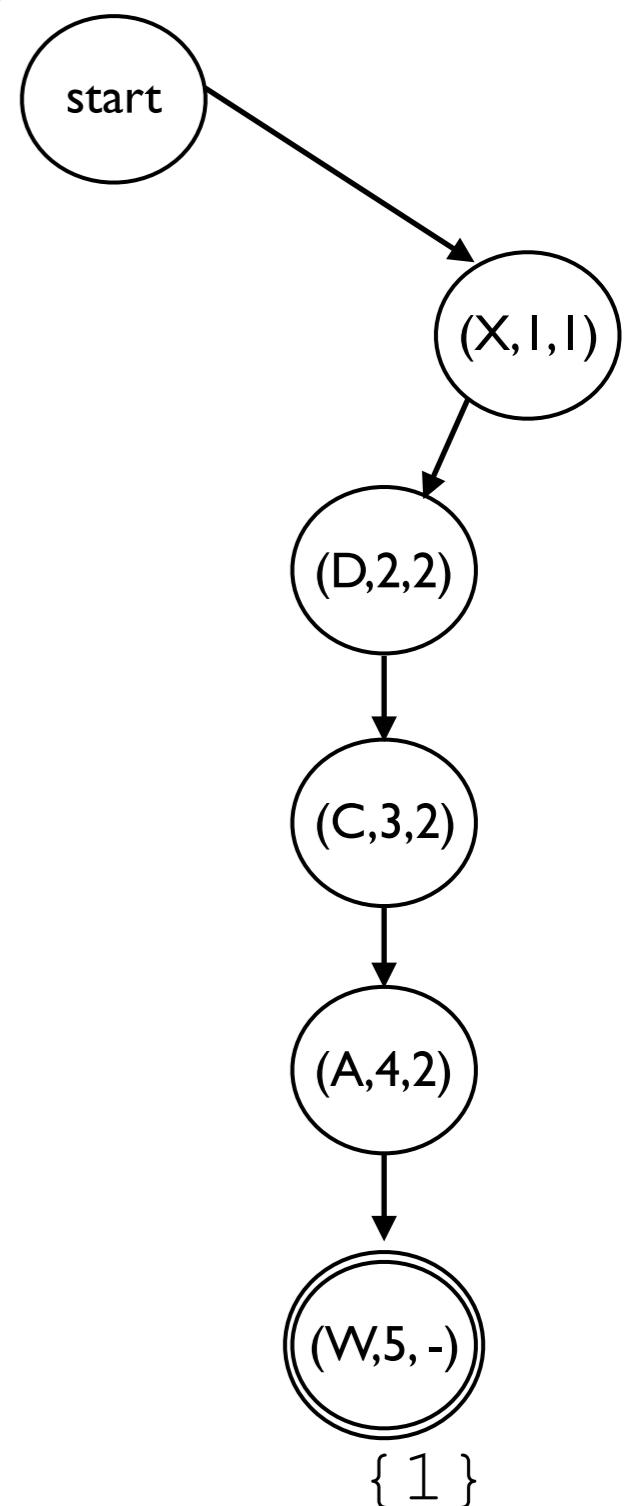
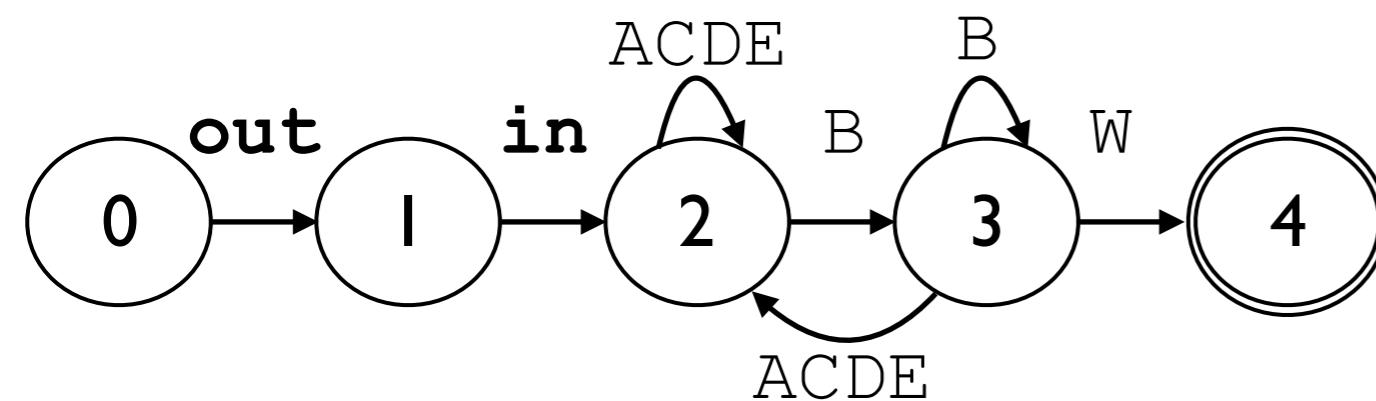
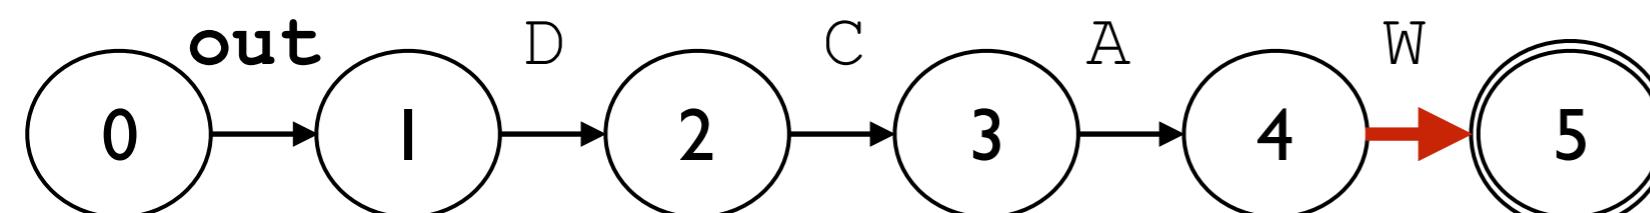
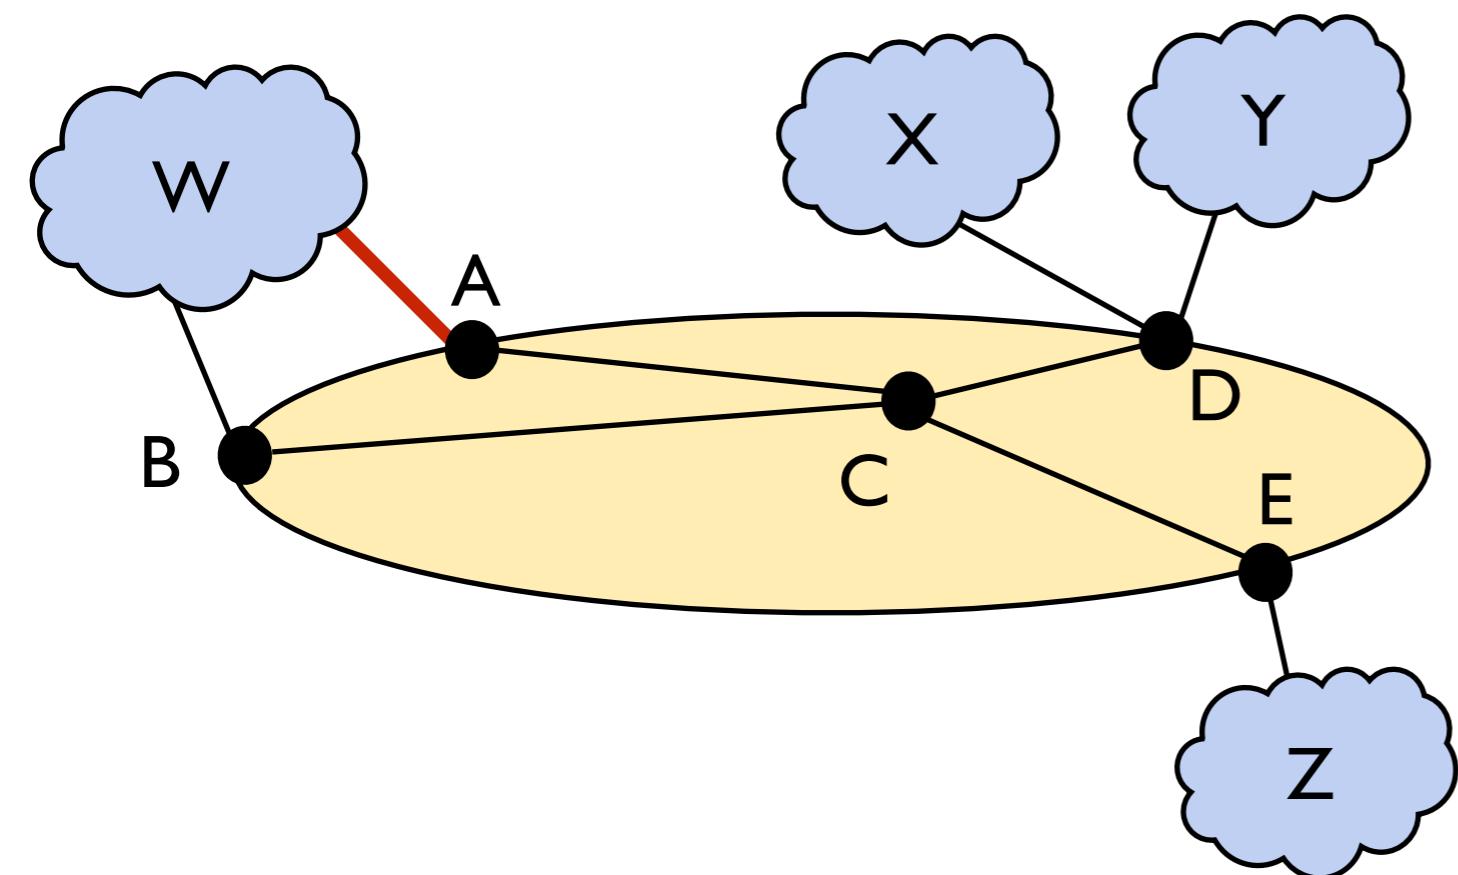
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



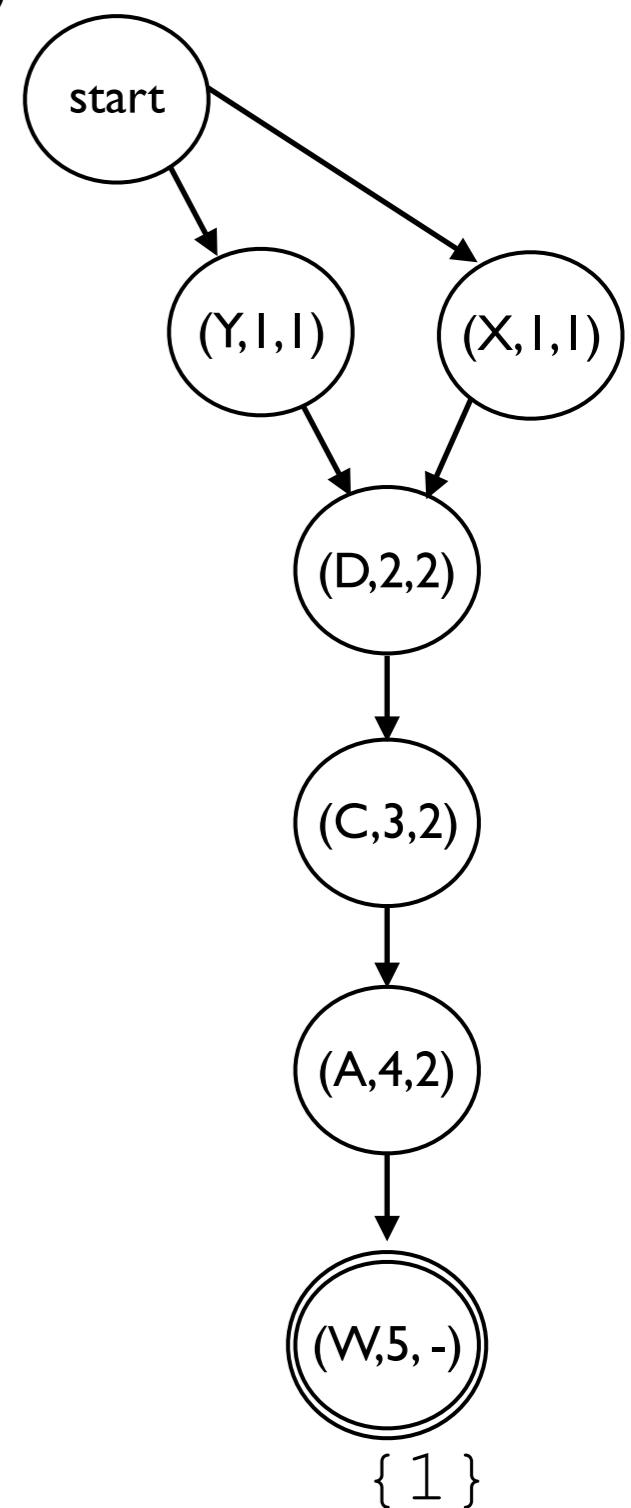
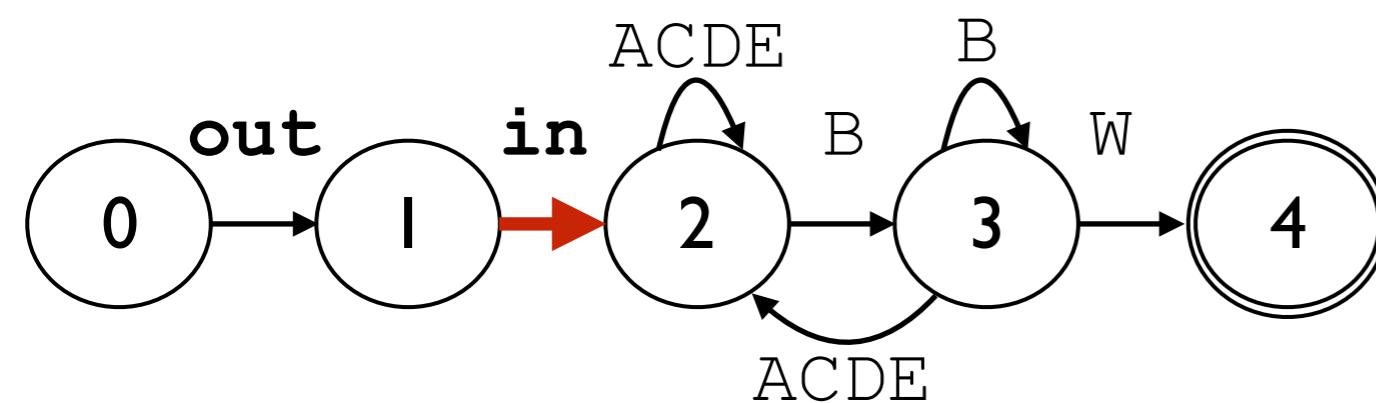
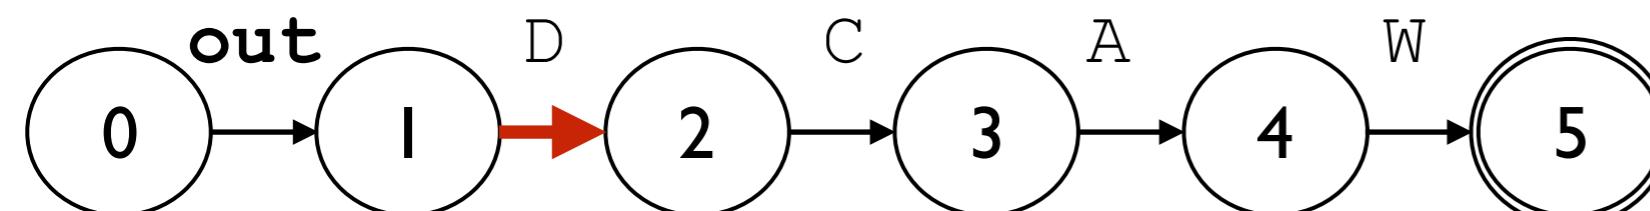
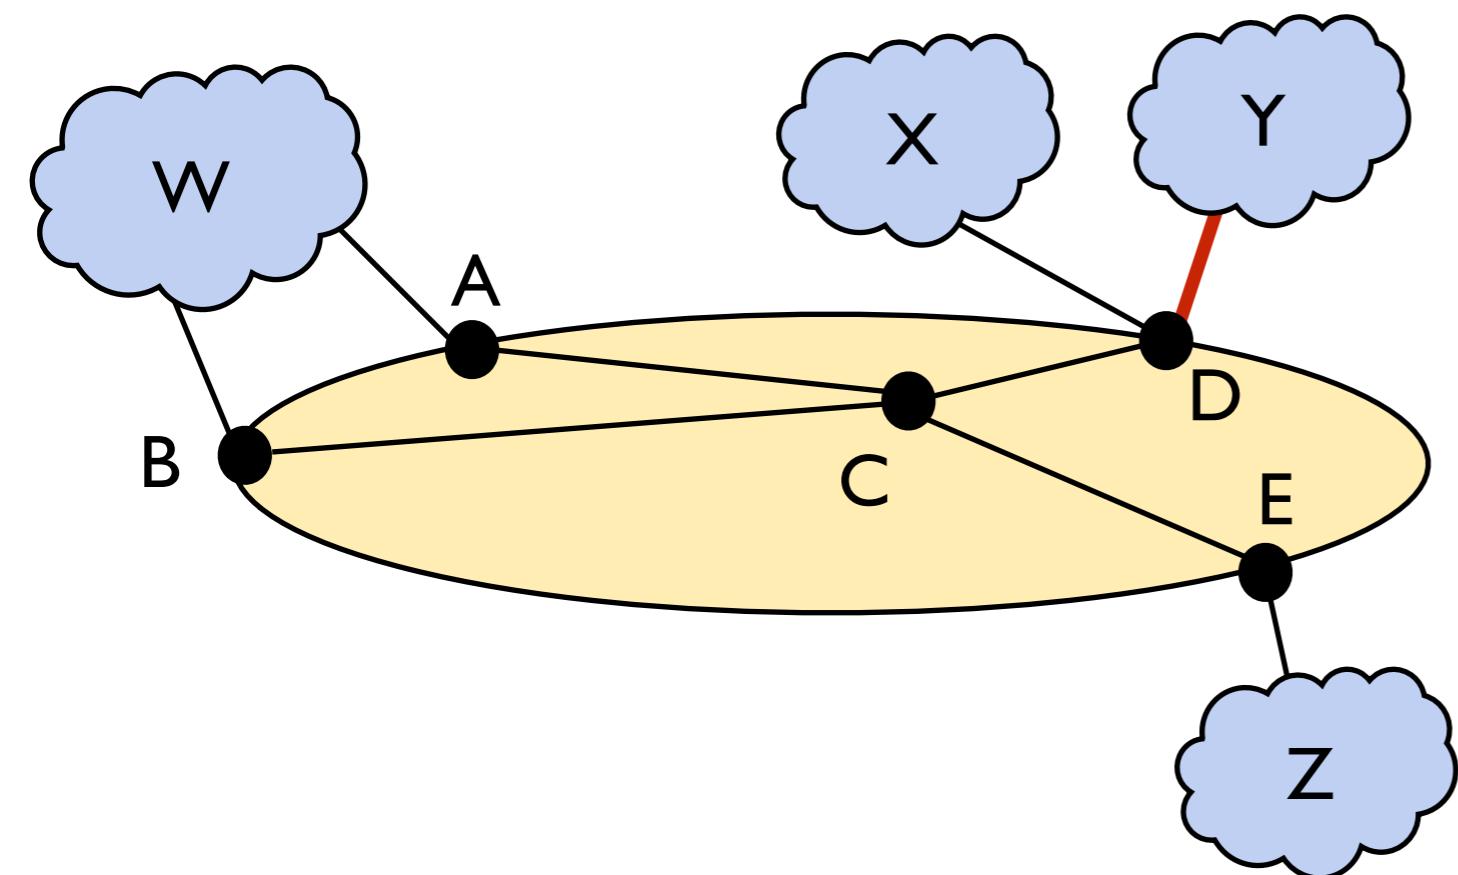
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



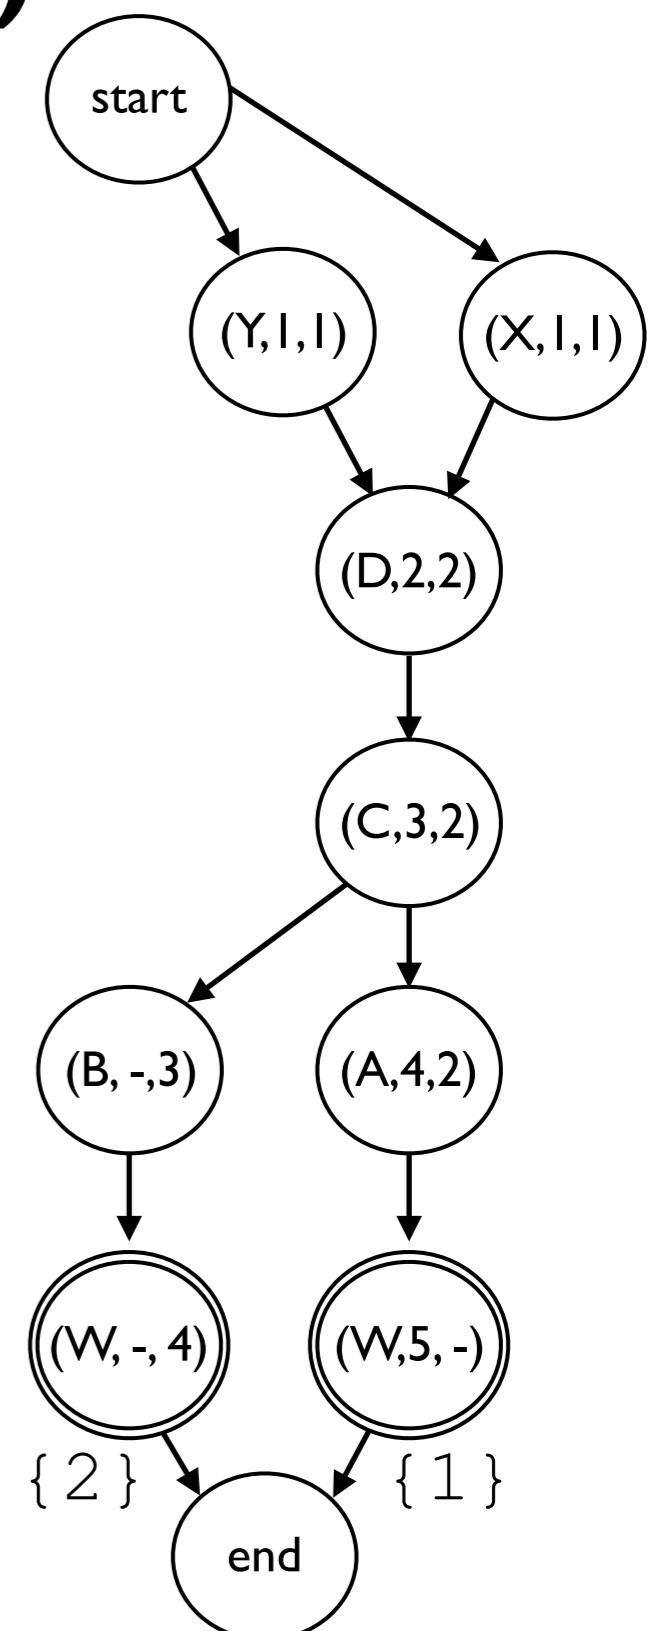
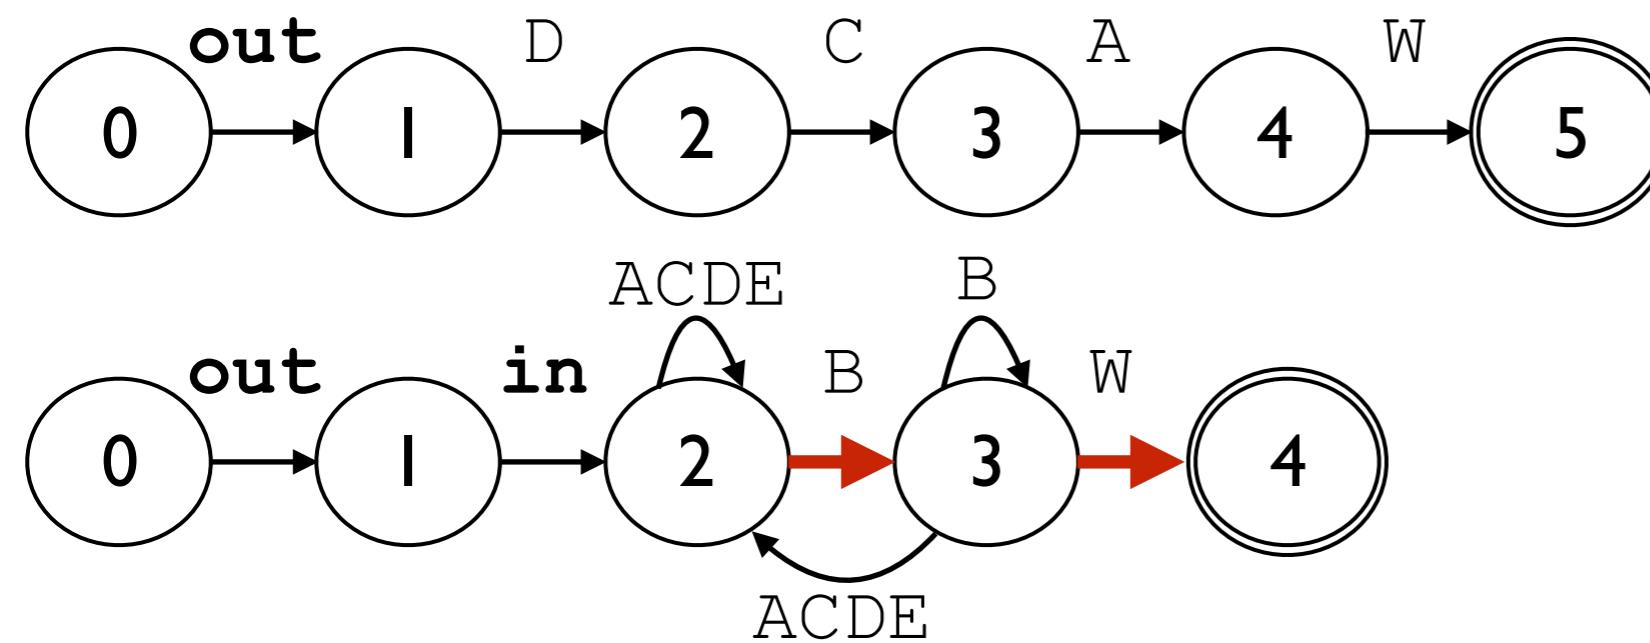
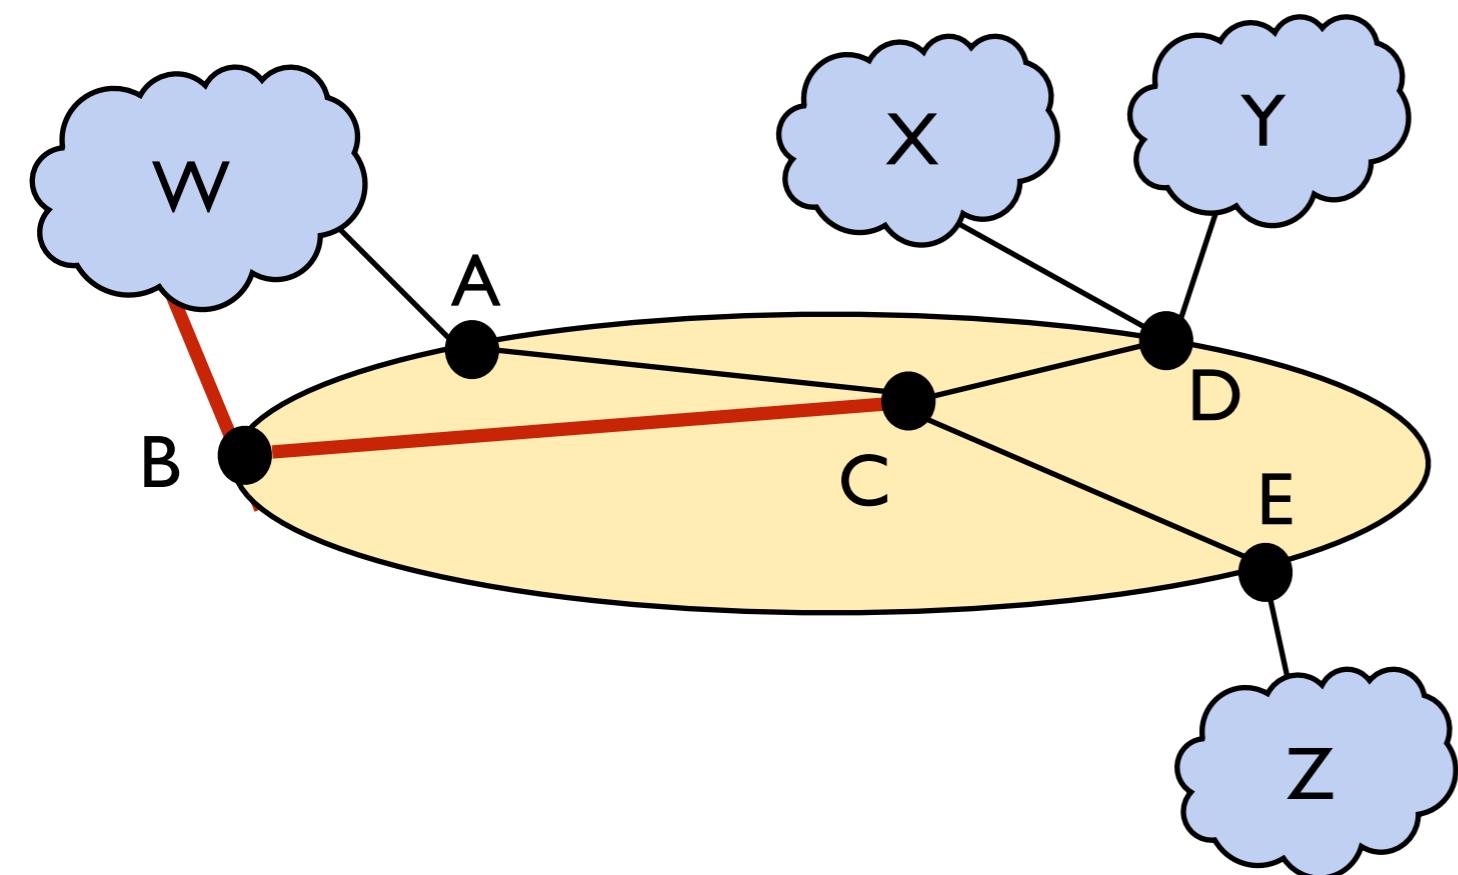
Constructing the Product Graph (PG)

`(W.A.C.D.out) >> (W.B.in+.out)`



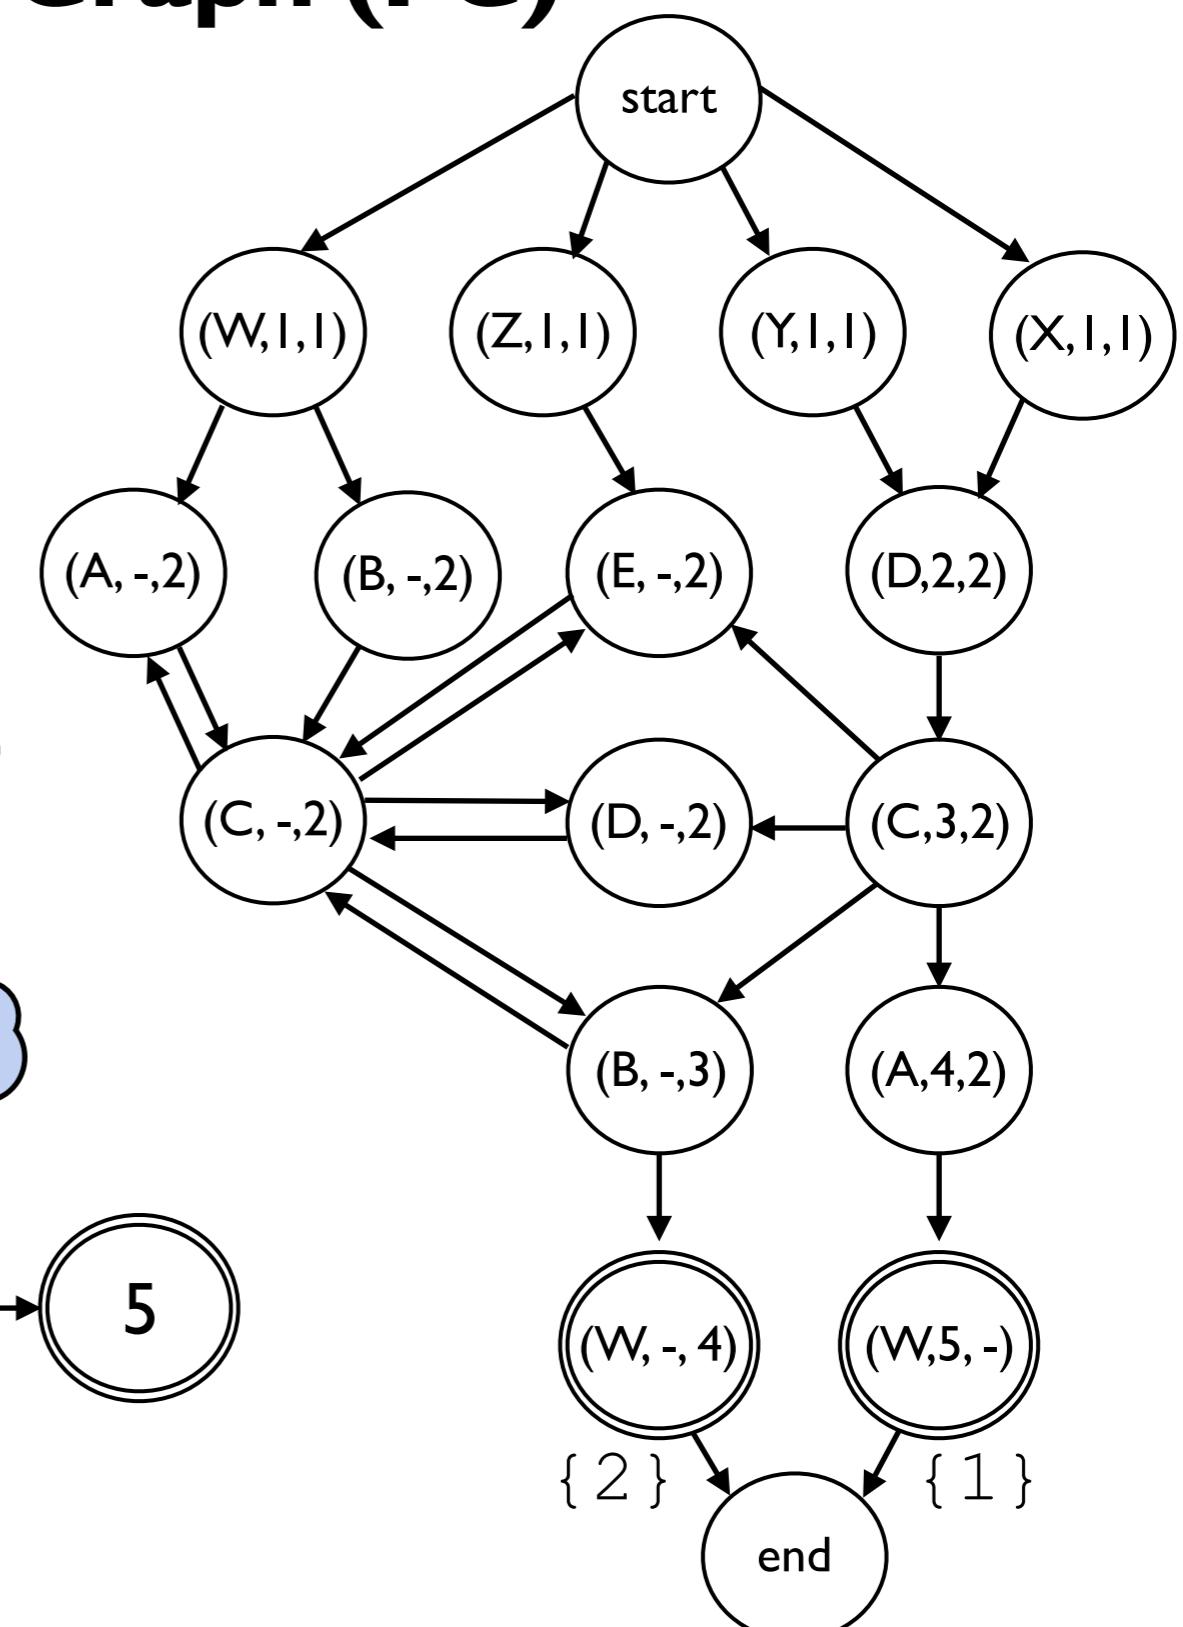
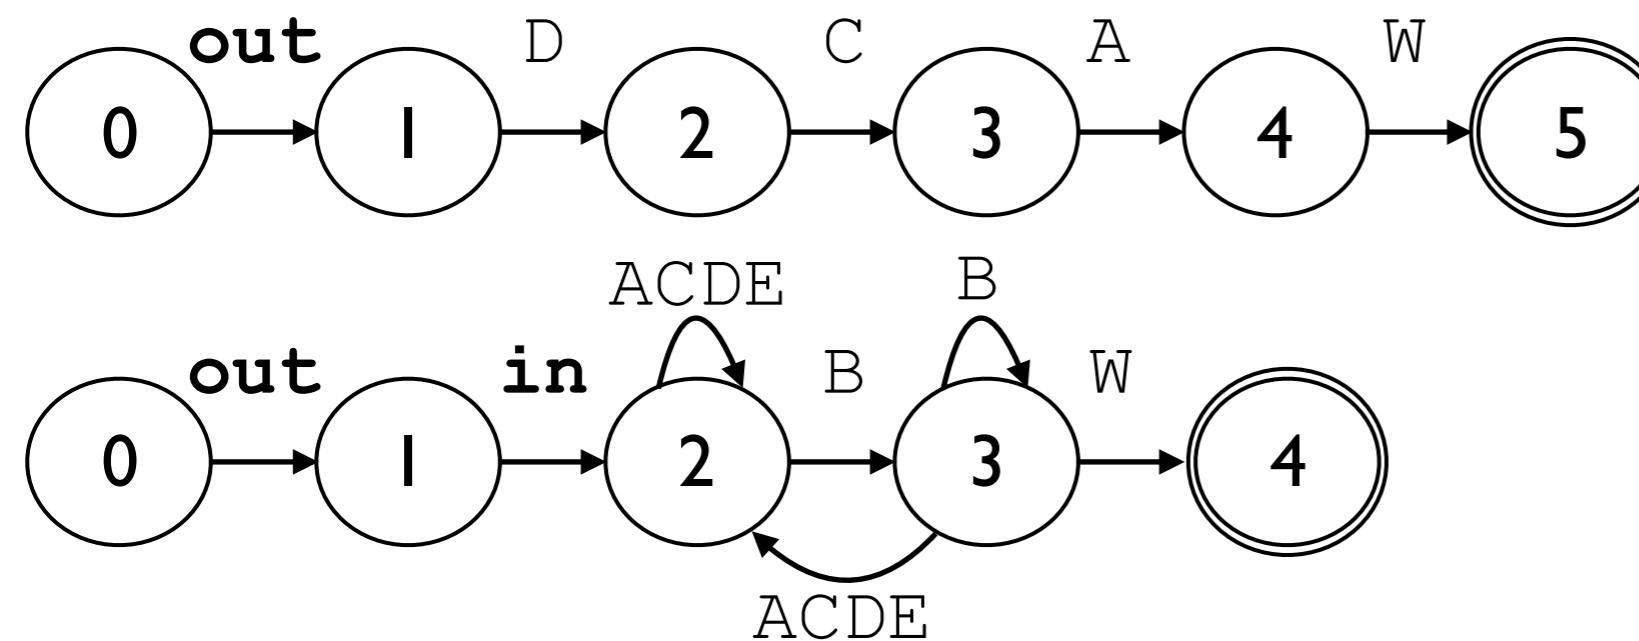
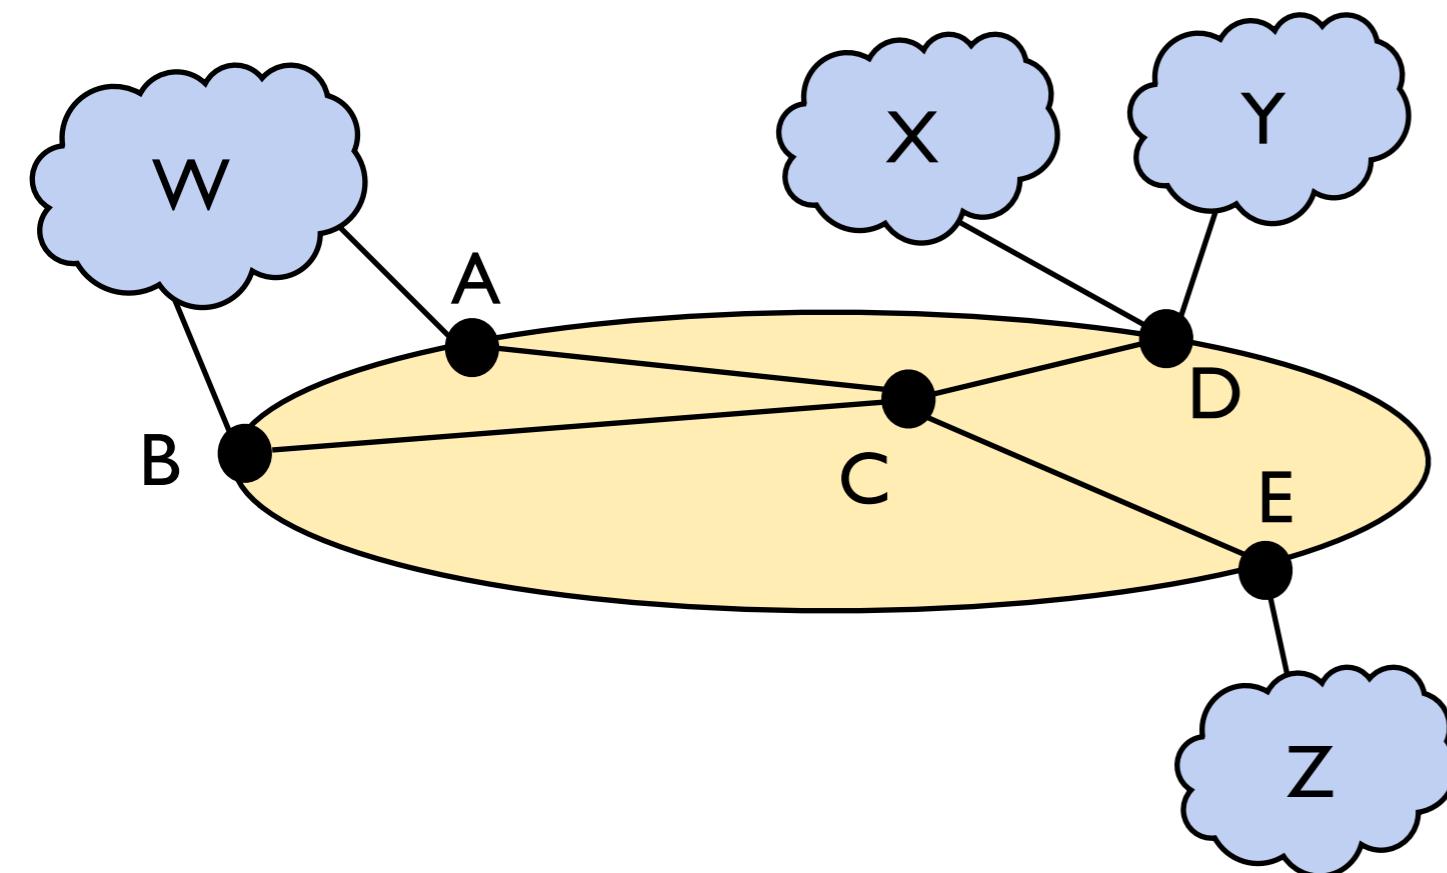
Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$



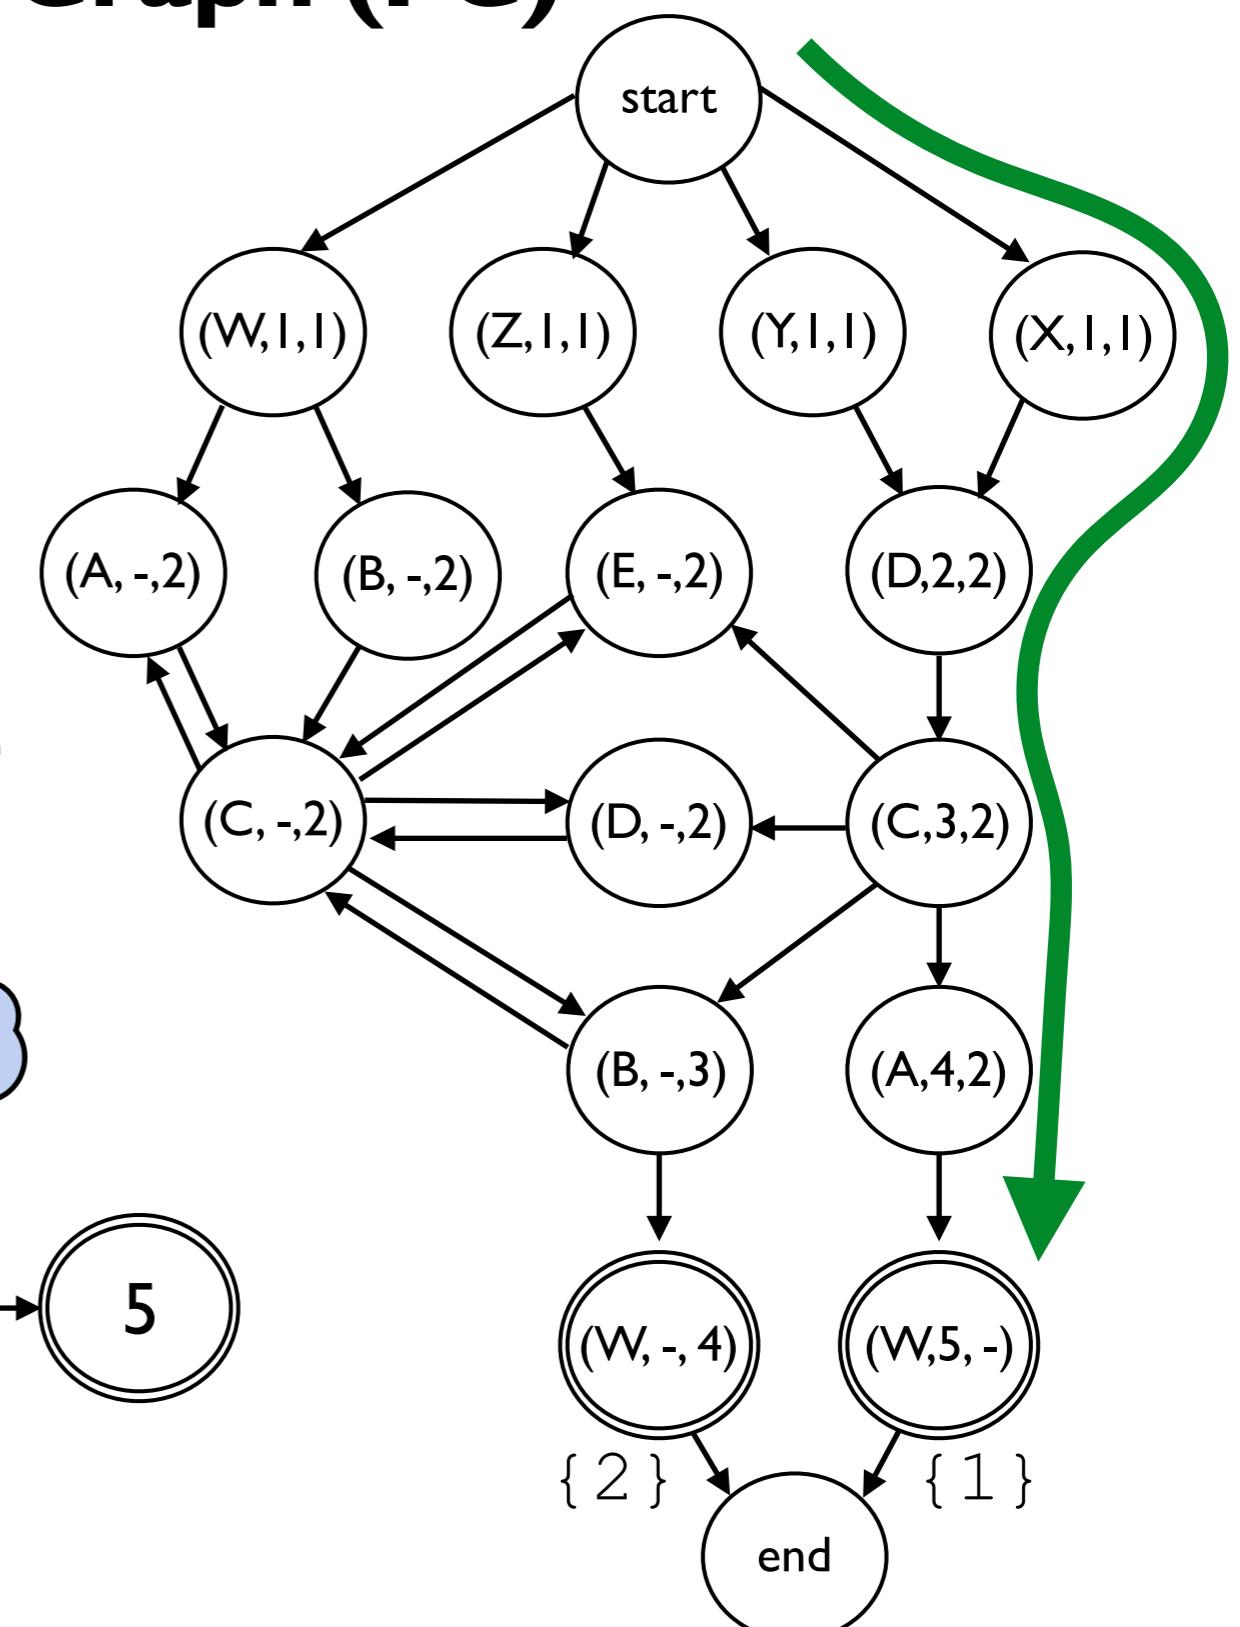
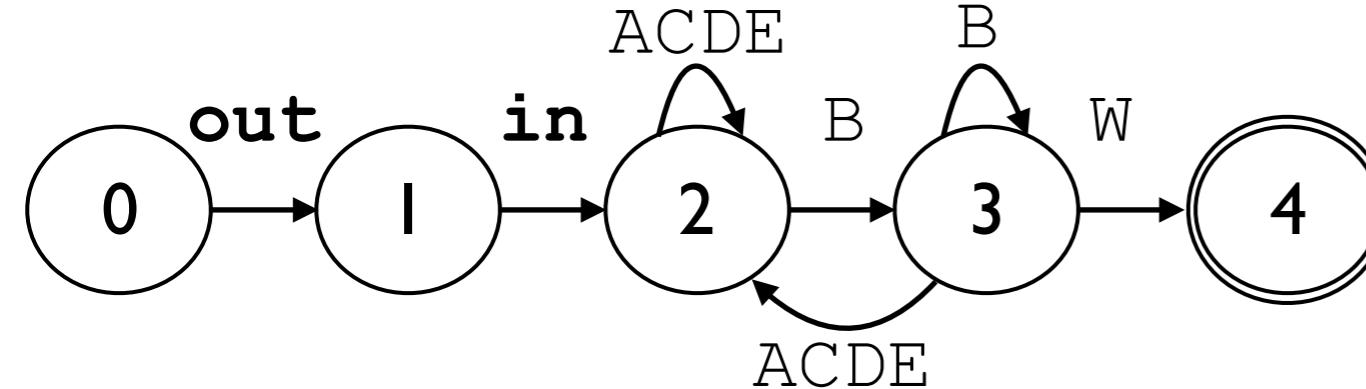
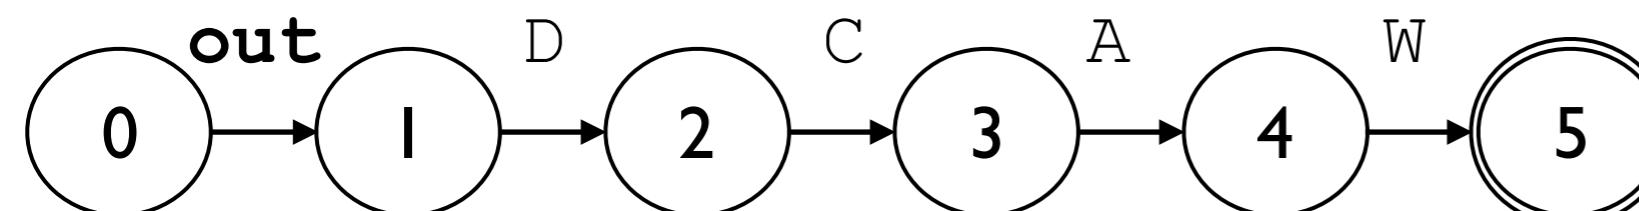
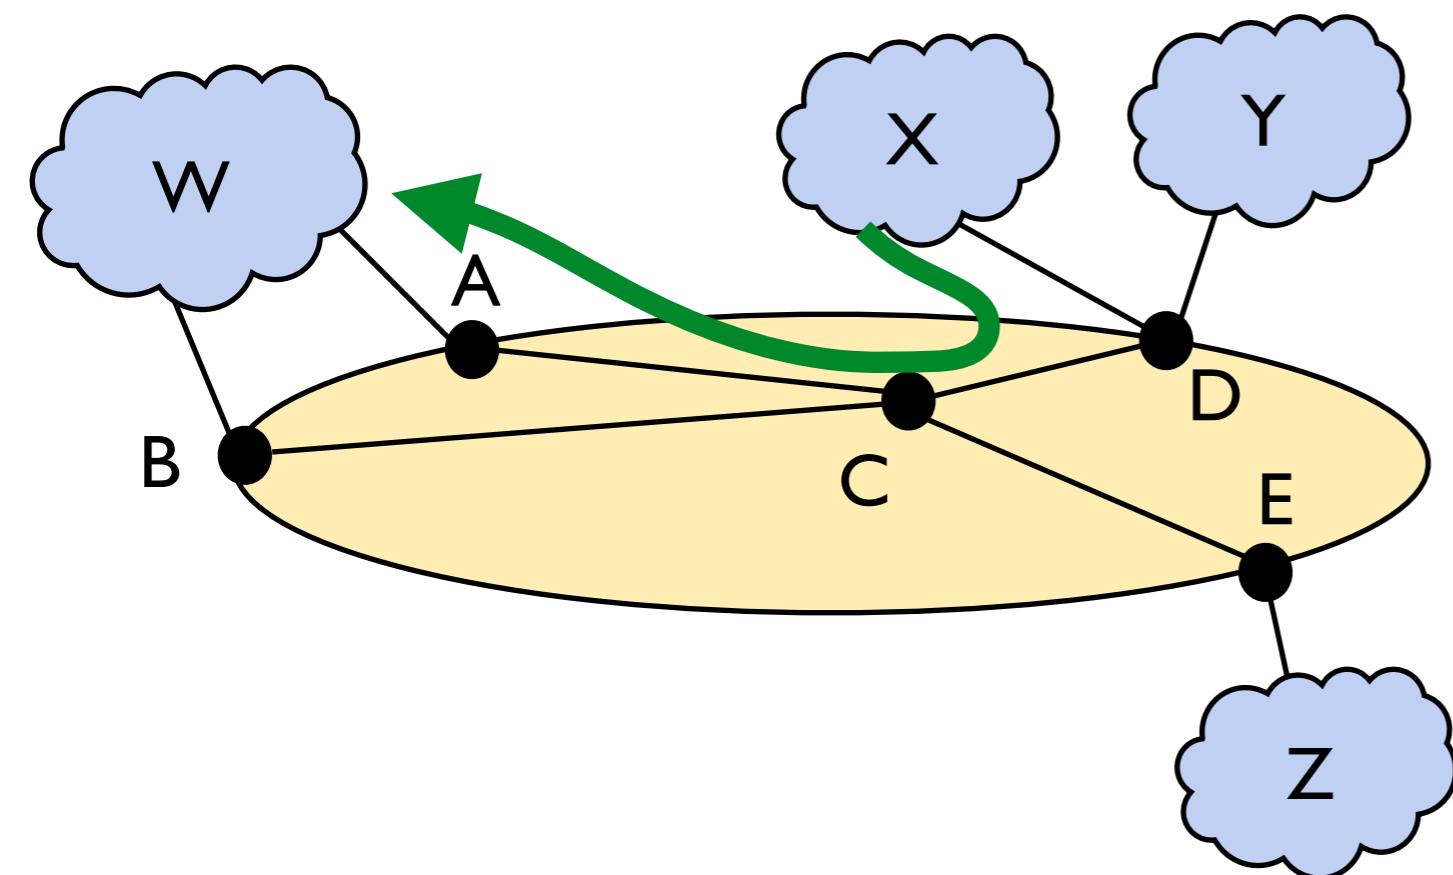
Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$



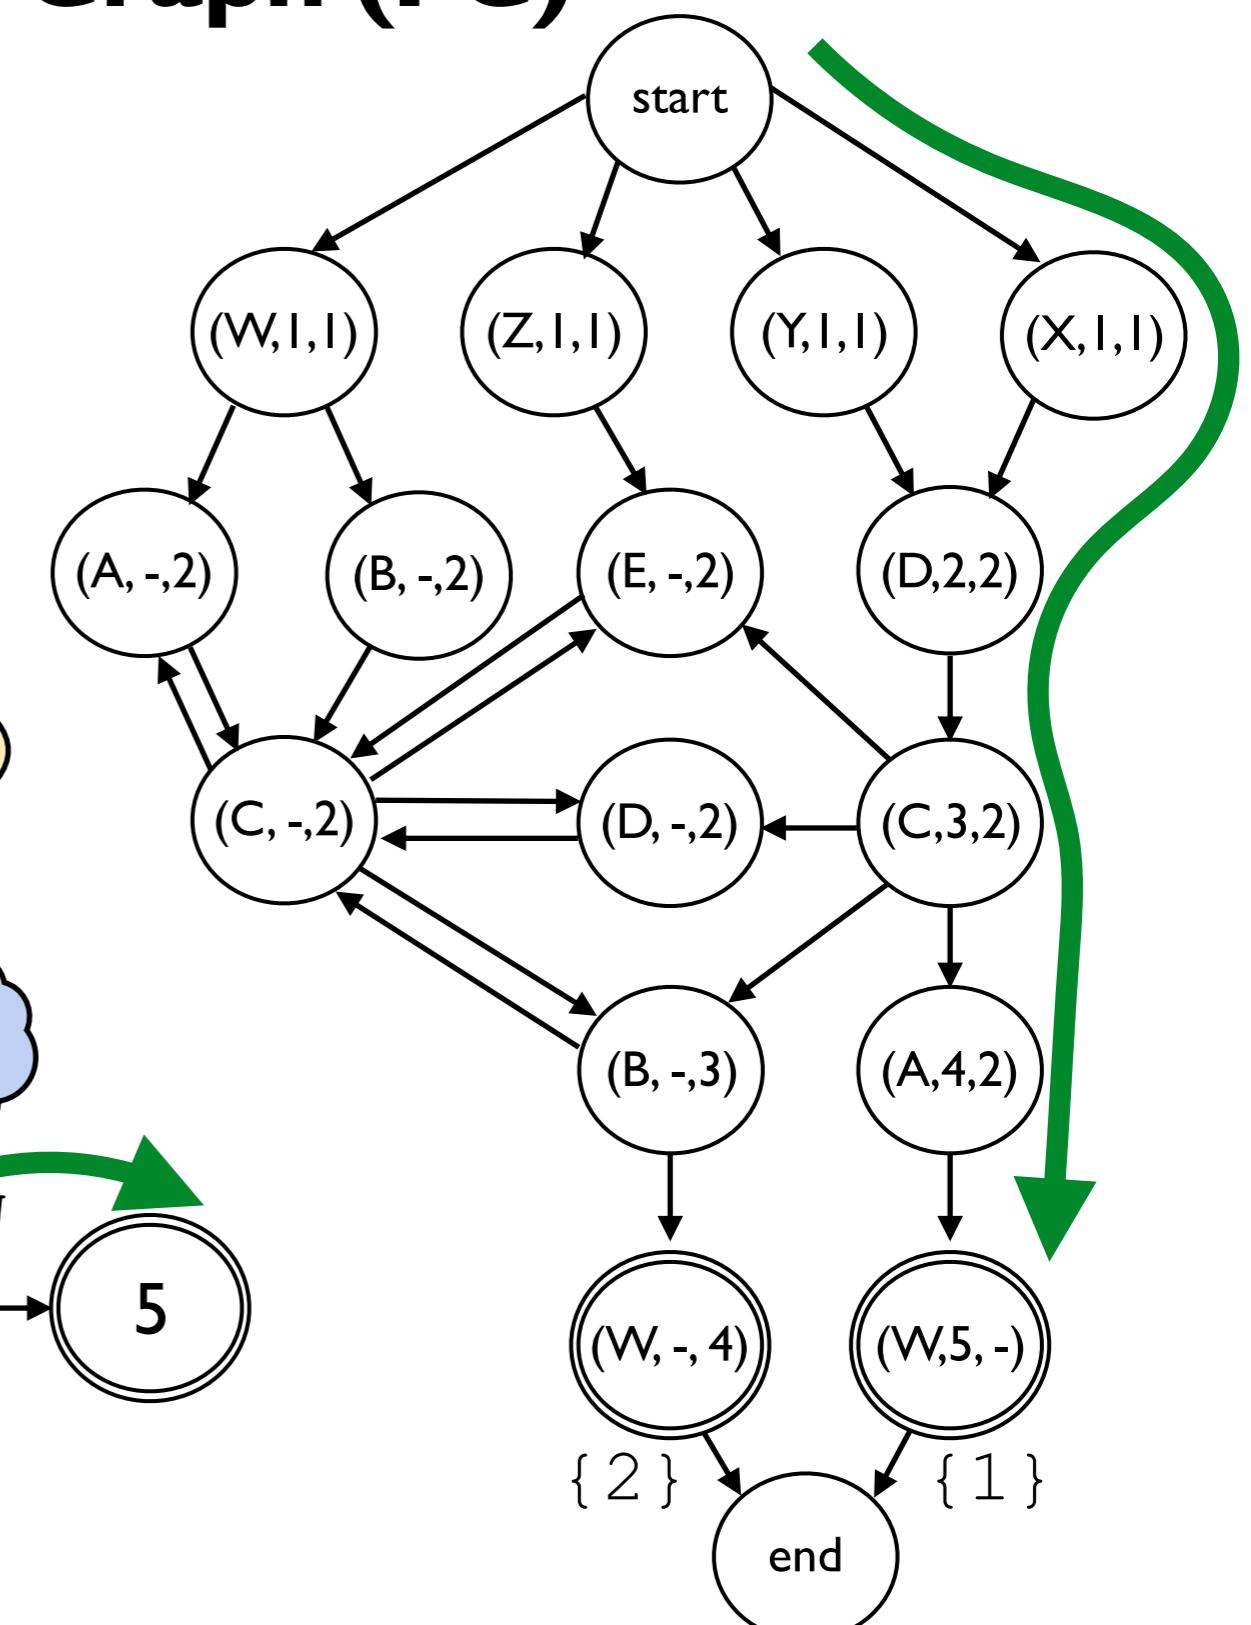
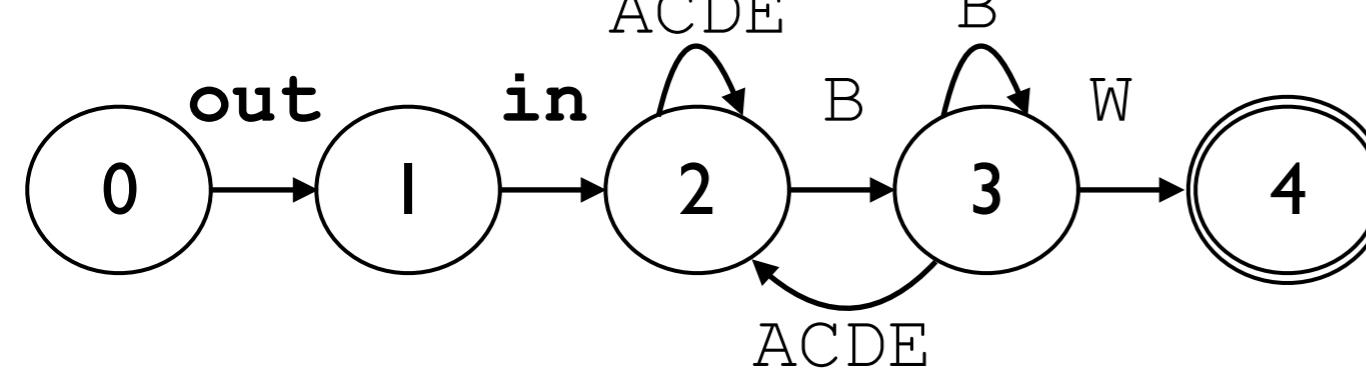
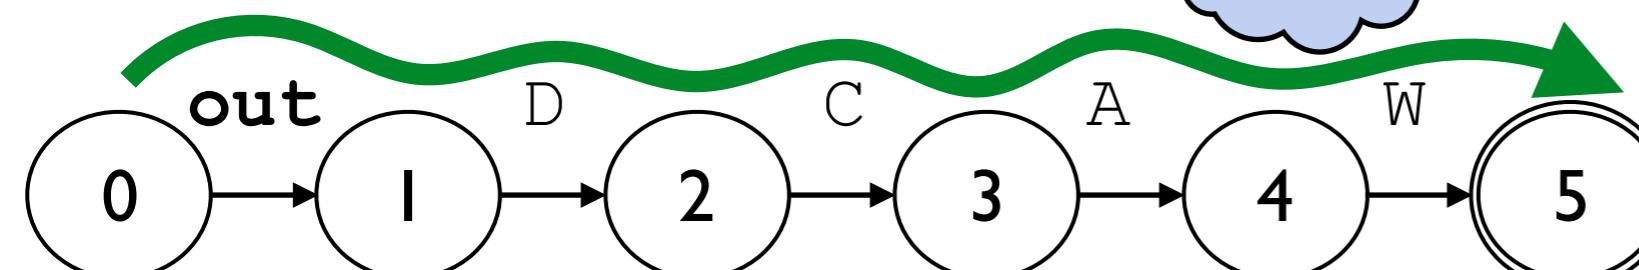
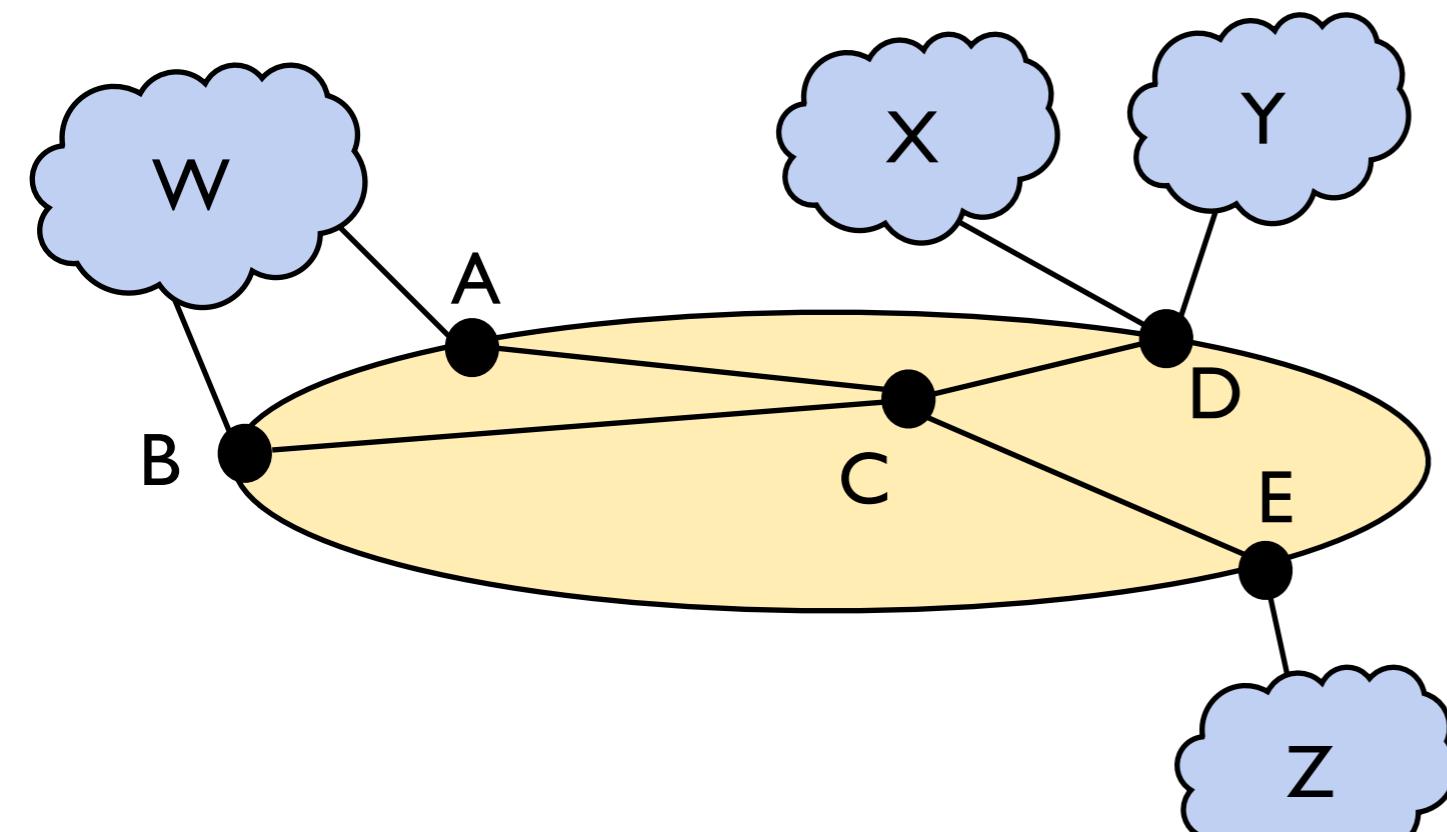
Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$

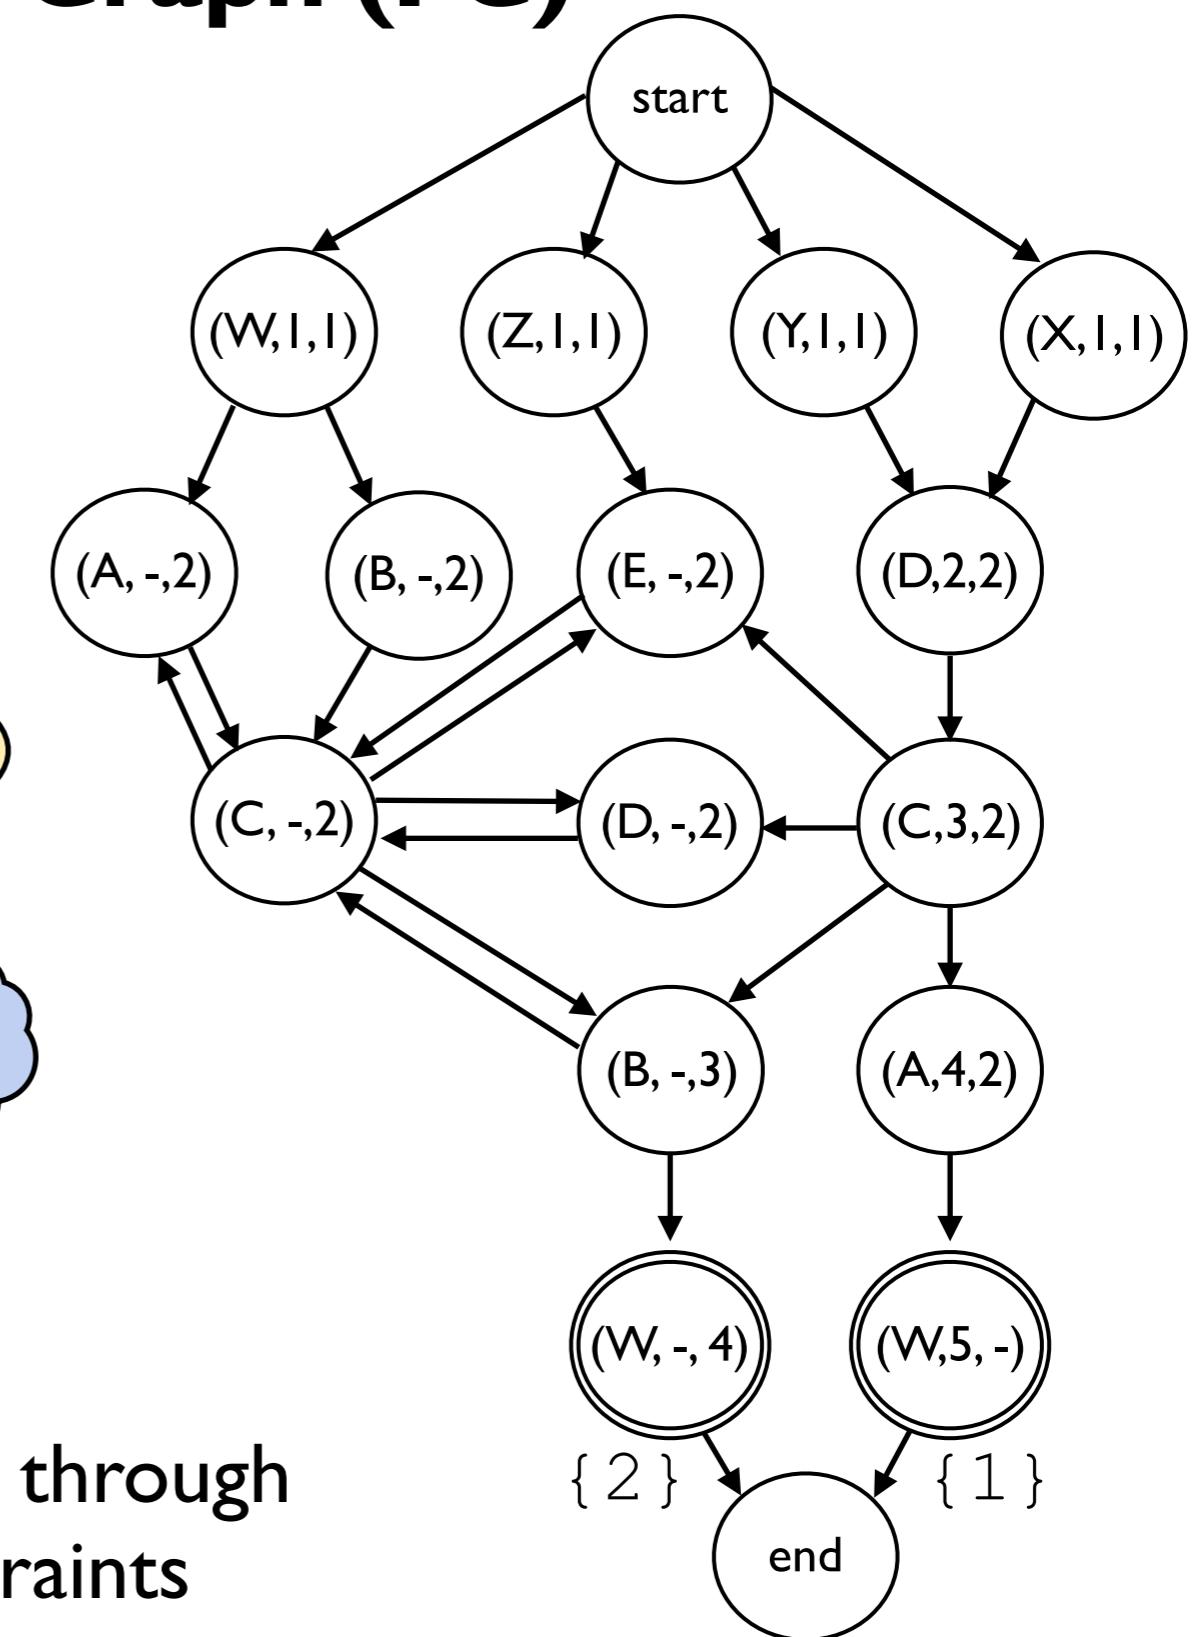
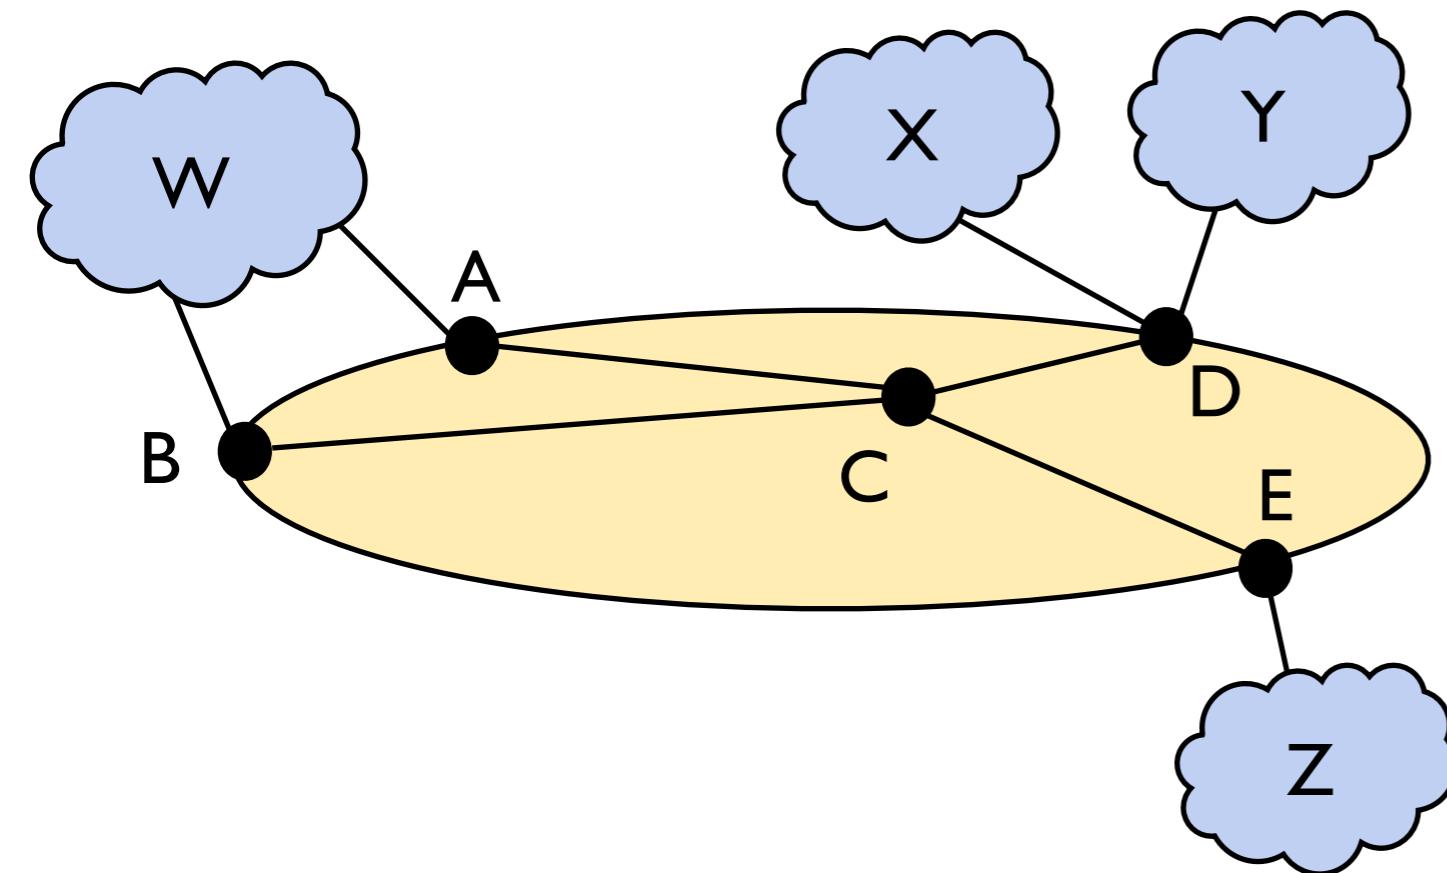


Constructing the Product Graph (PG)

$(W.A.C.D.out) \gg (W.B.in+out)$

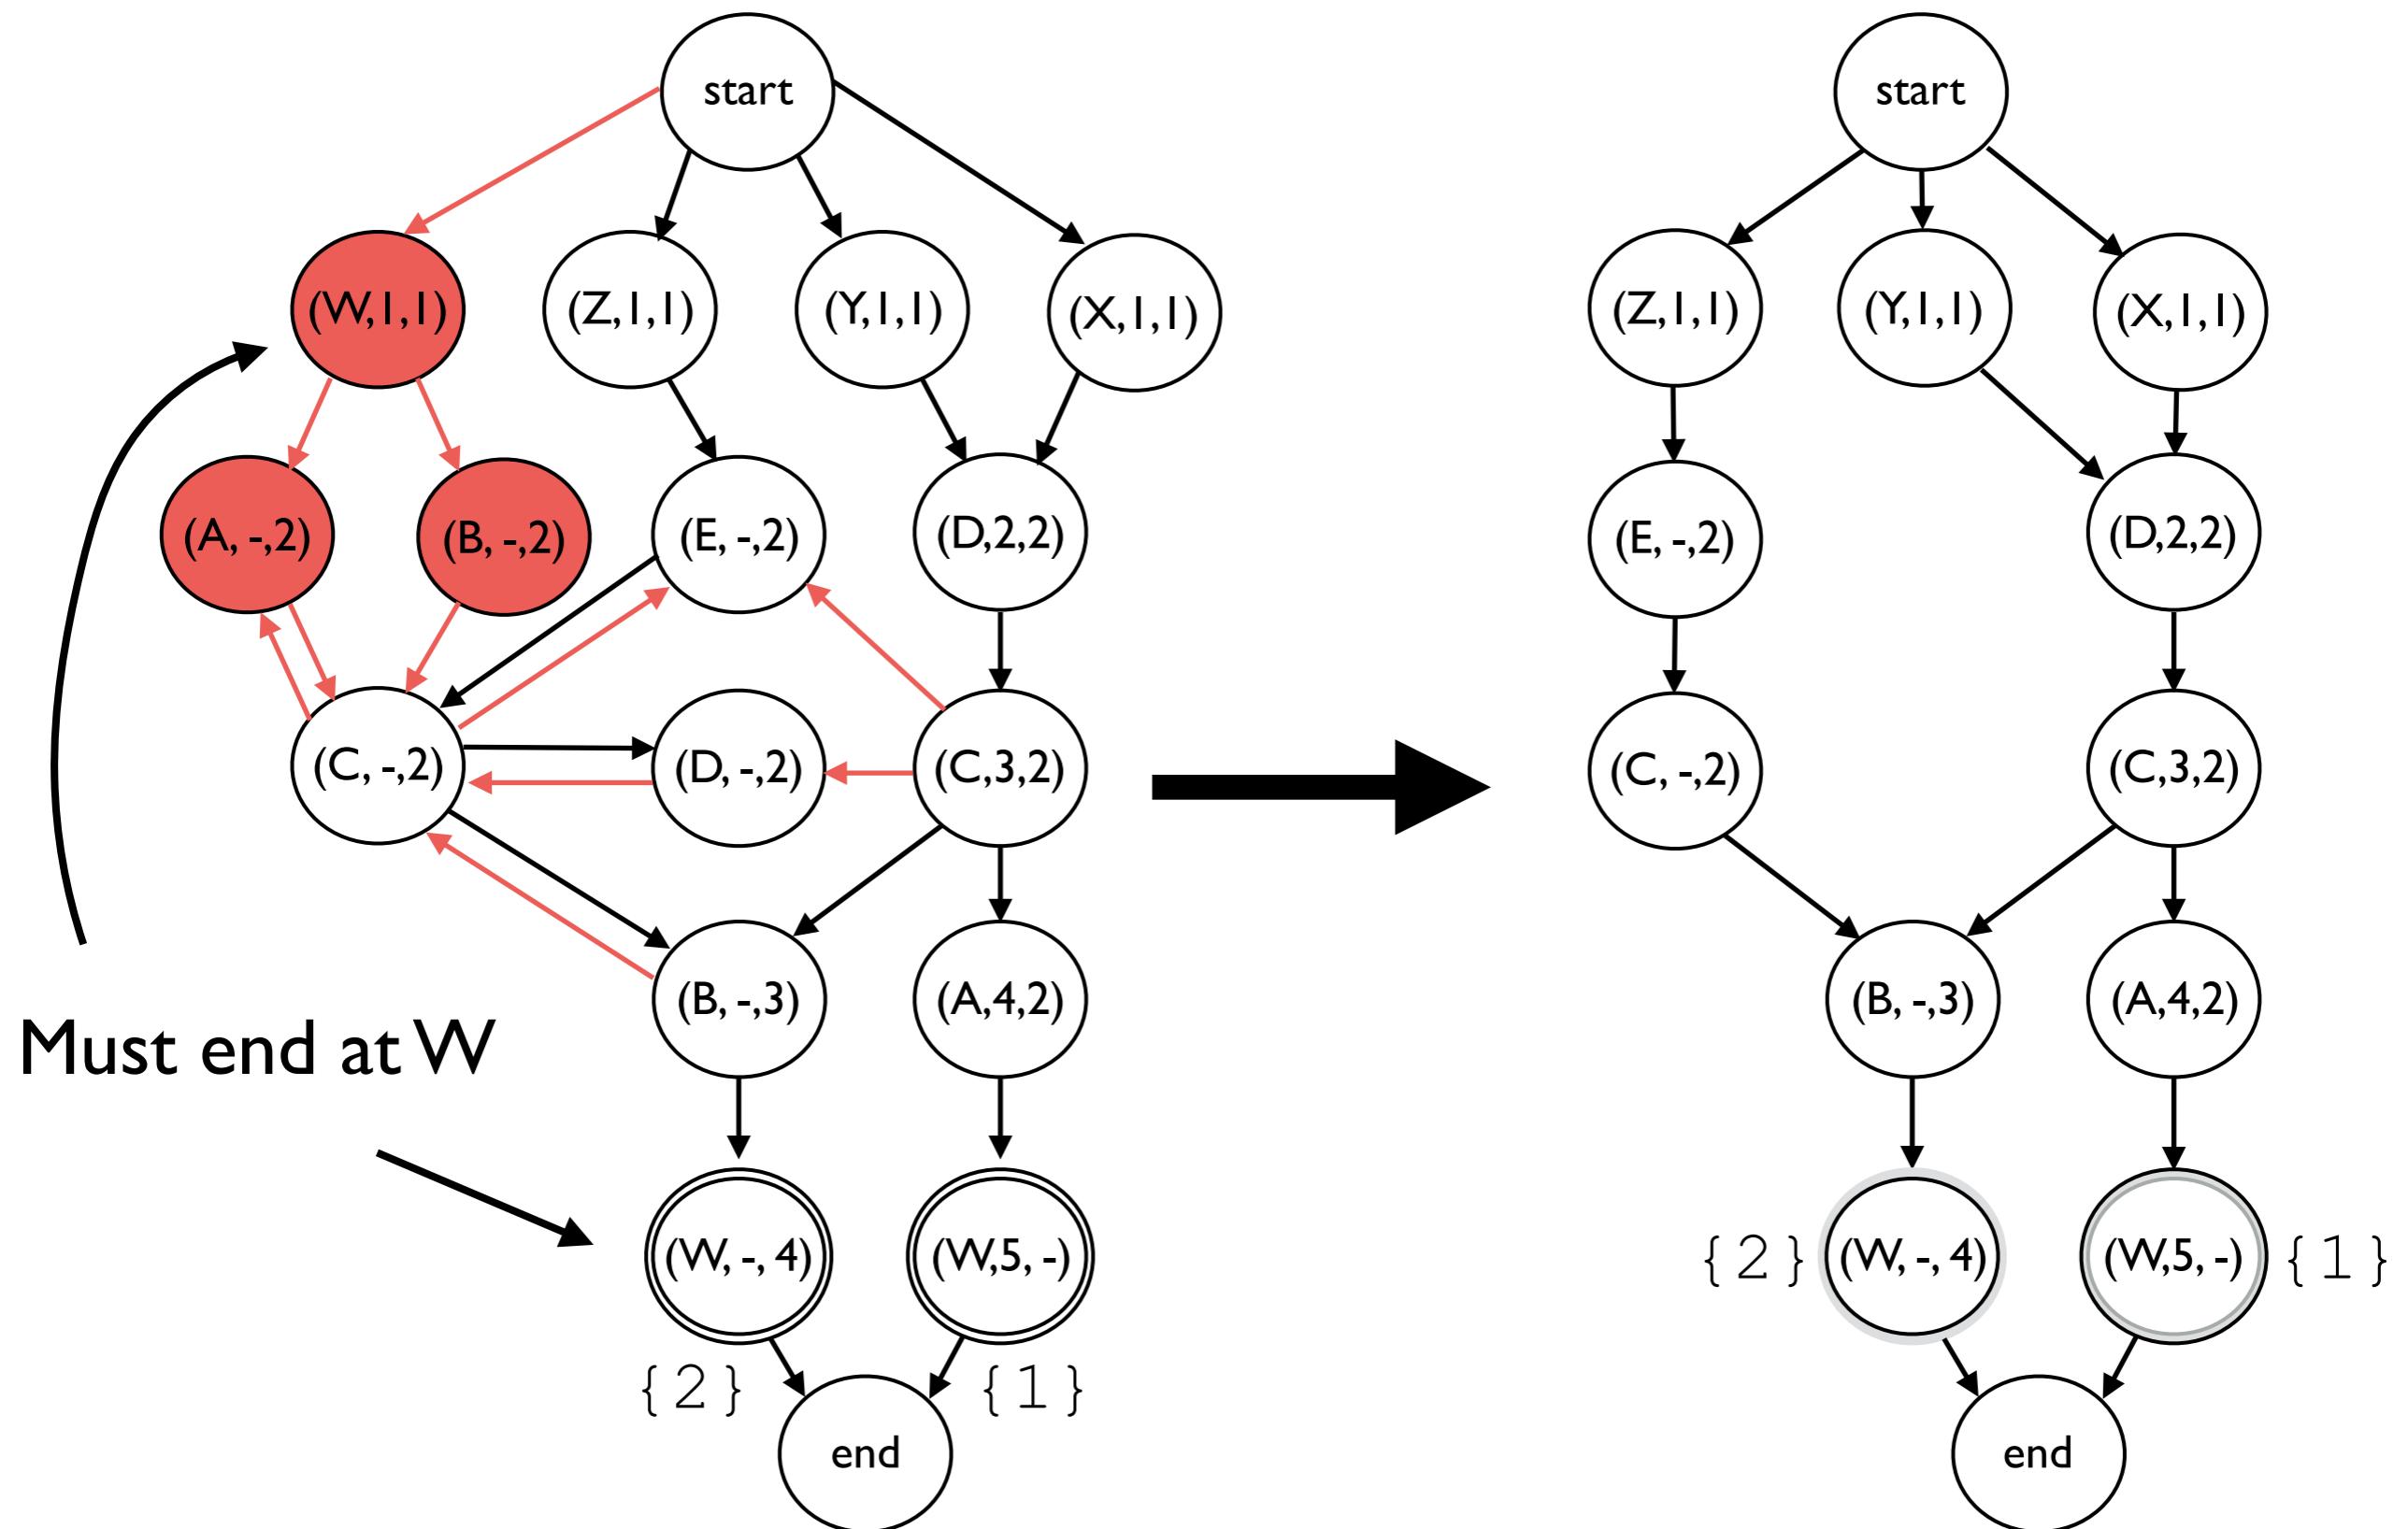


Constructing the Product Graph (PG)

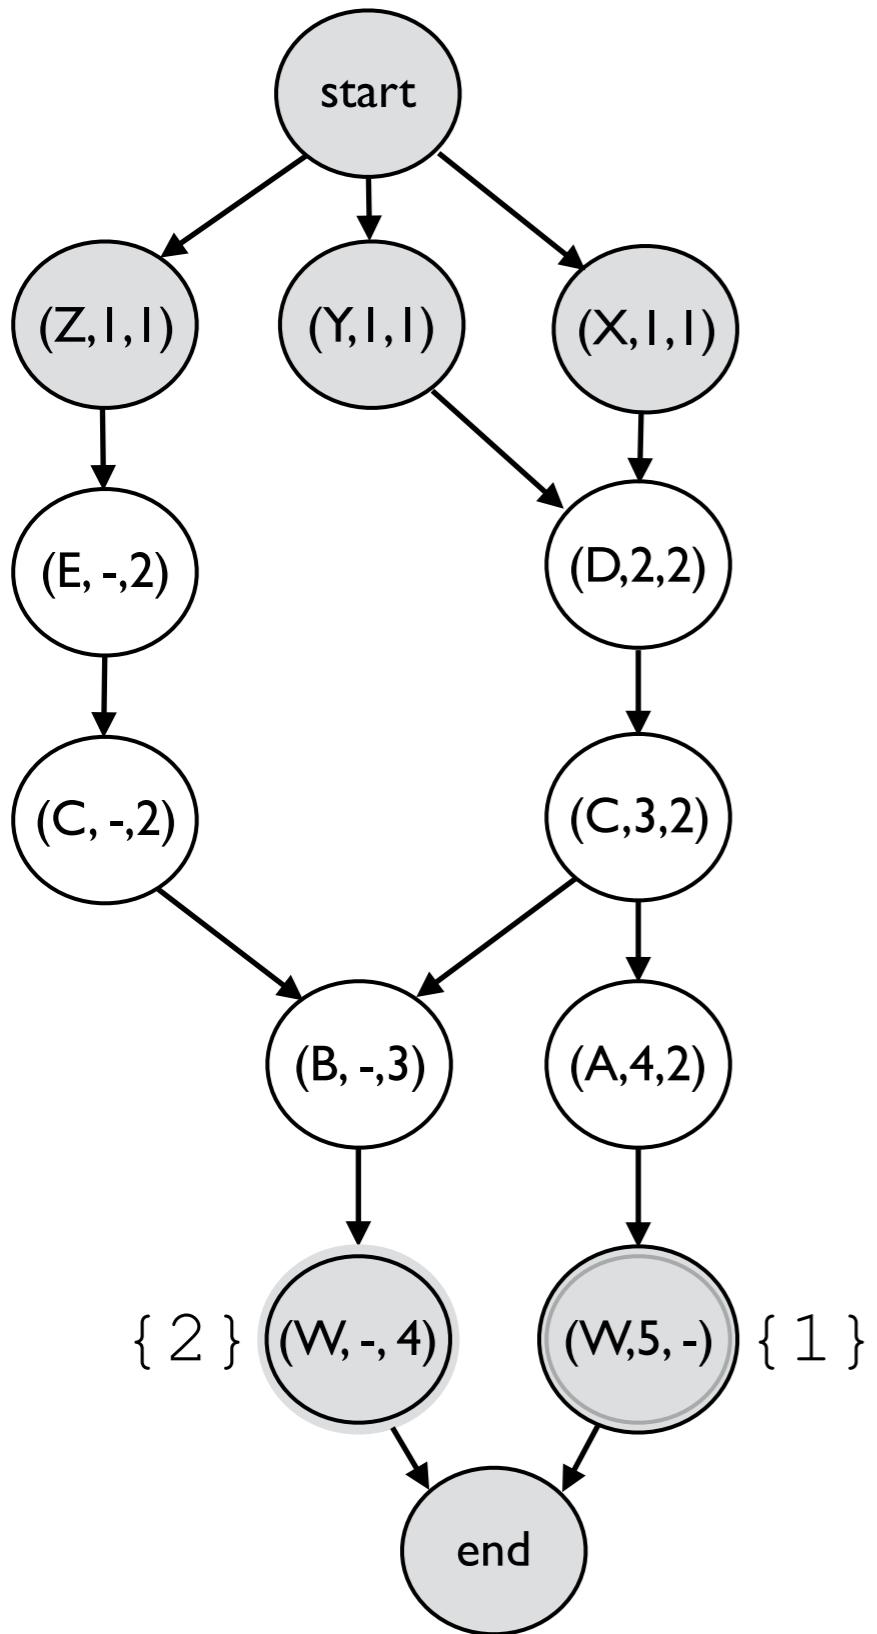


Compact representation of all paths through
the topology subject to policy constraints

PG Minimization (Loop analysis)

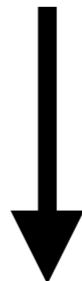


Compilation to BGP:



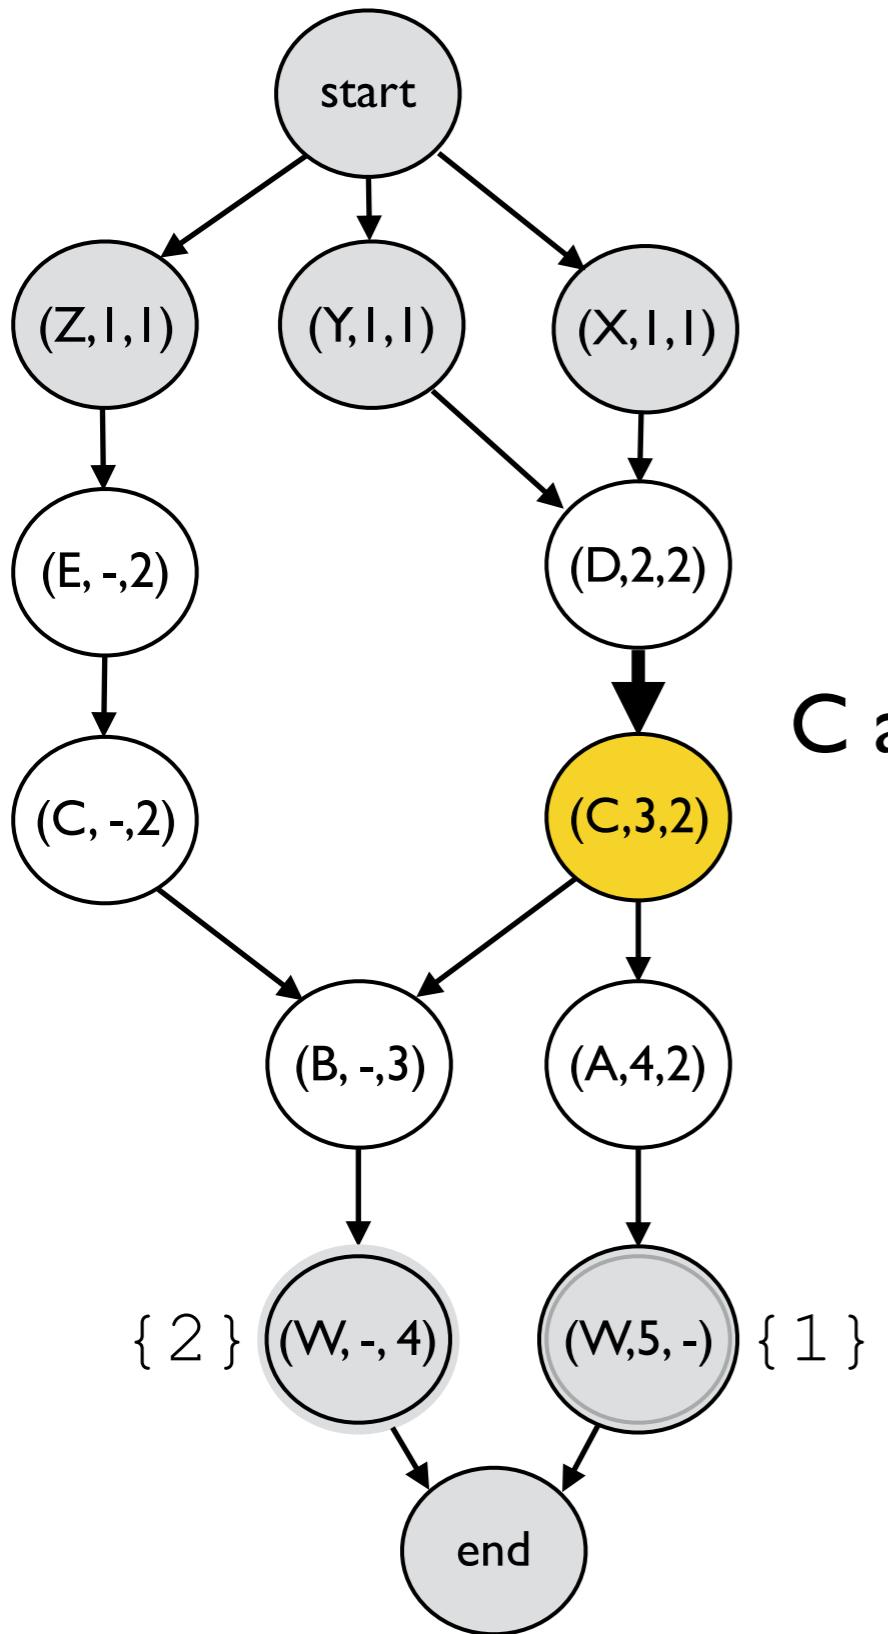
Idea I: Restrict advertisements to PG edges

- Encode PG state in community tag
- Incoming edges — import filters
- Outgoing edges — export filters



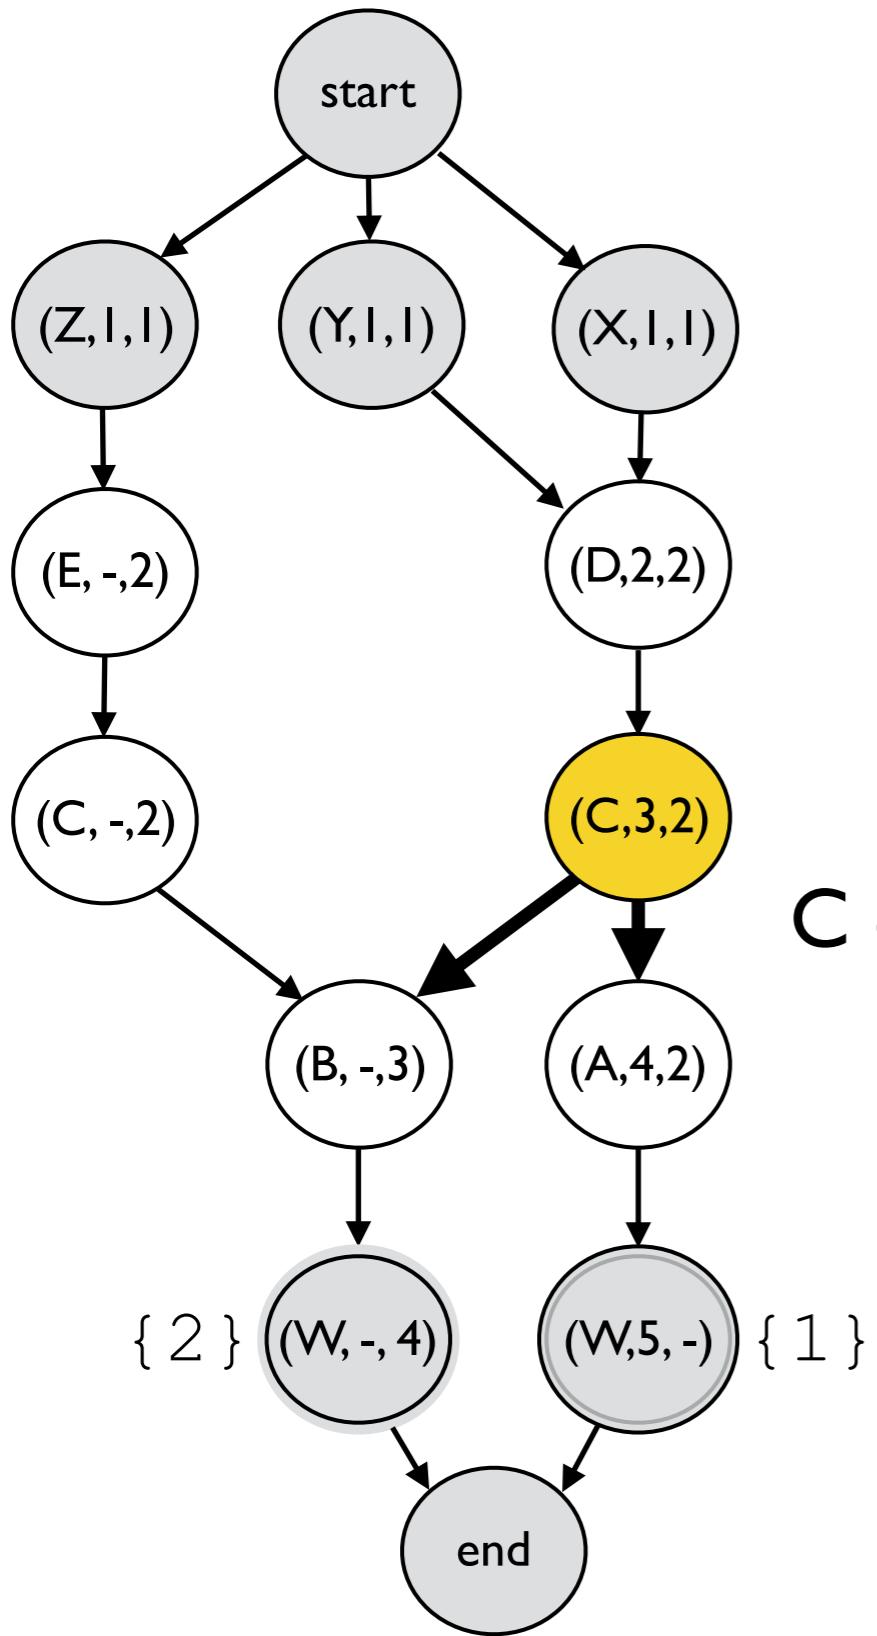
Let BGP find **some allowed** path dynamically

Compilation to BGP:



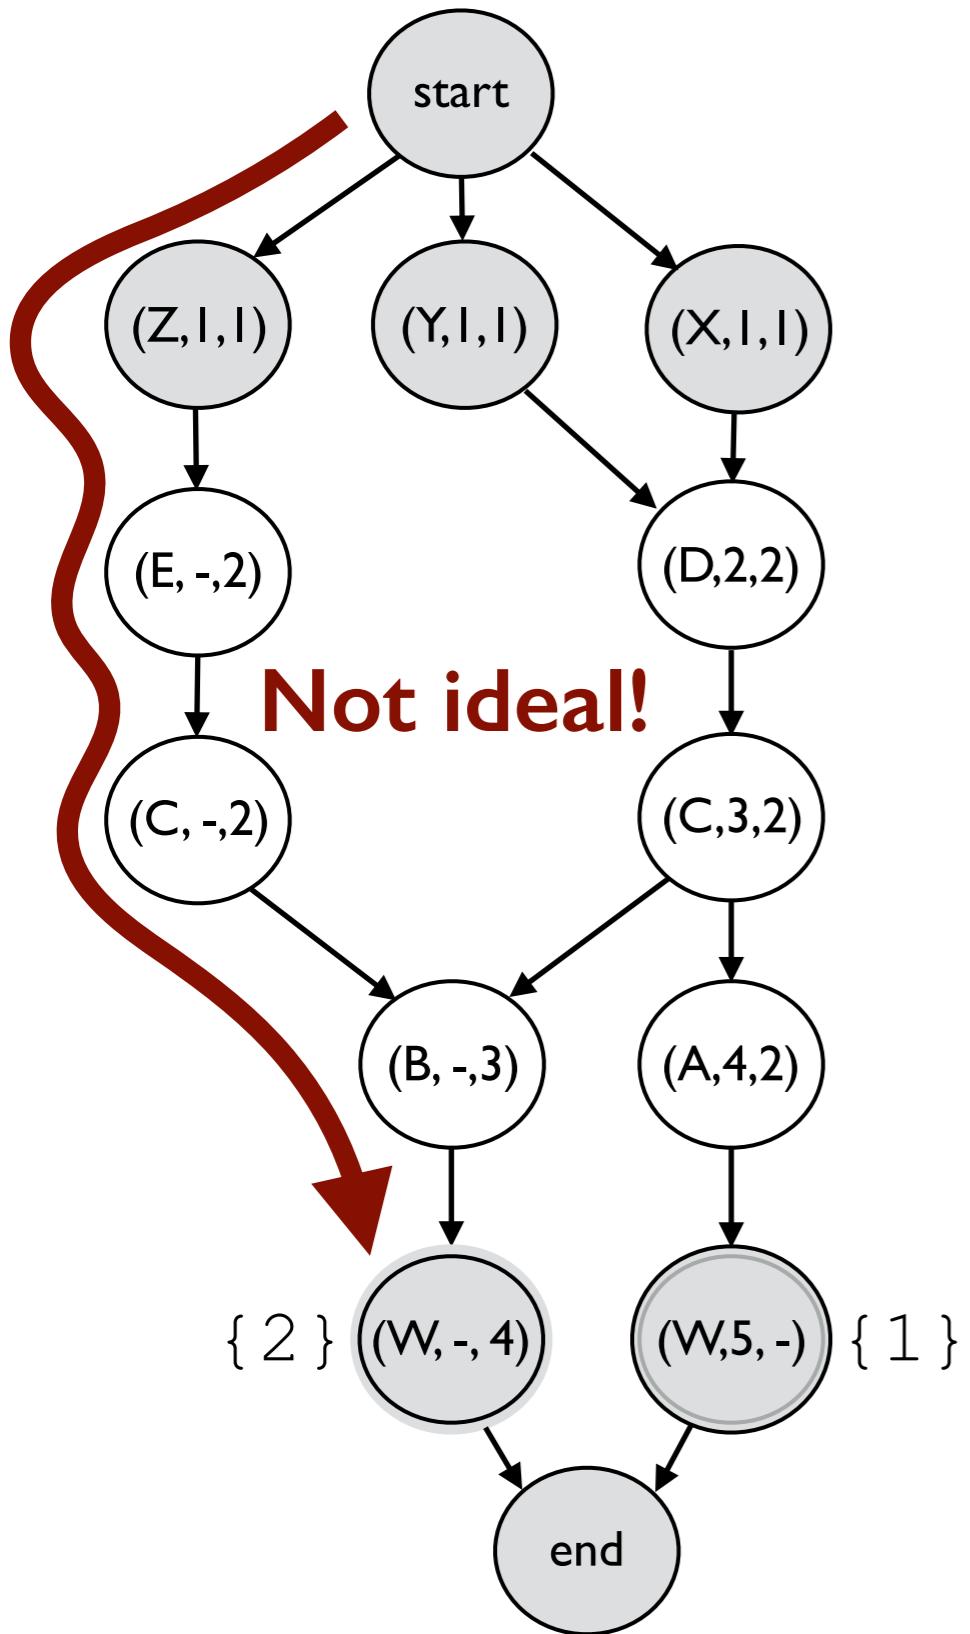
C allows import from D with tag (2,2)

Compilation to BGP:



C exports to A,B with tag (3,2)

Compilation to BGP:



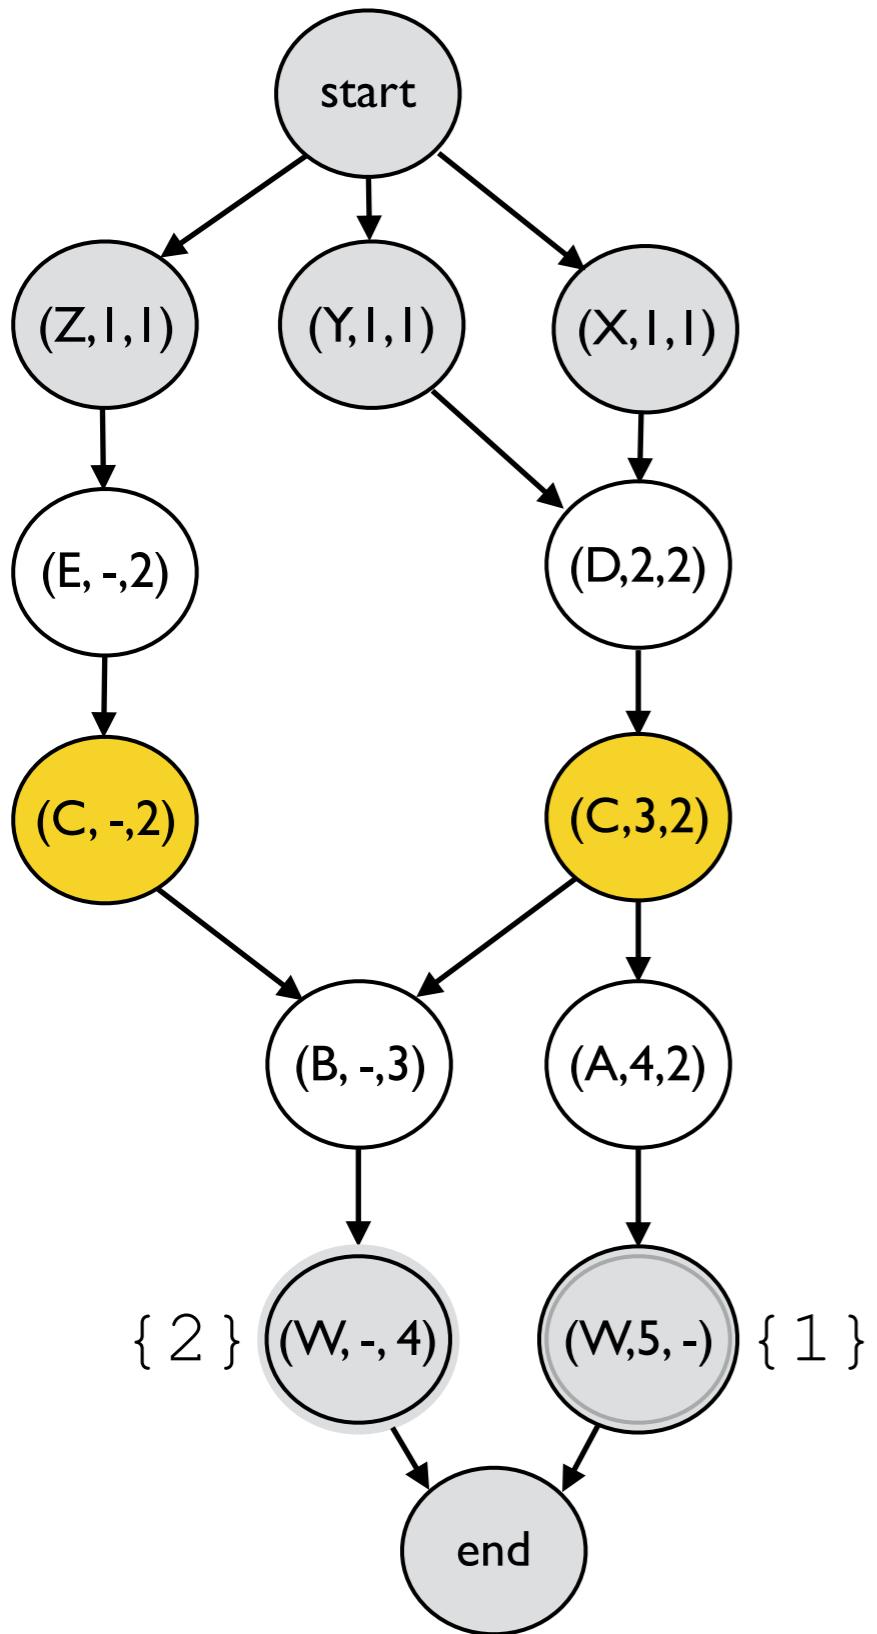
Idea I: Restrict advertisements to PG edges

- Encode PG state in community tag
- Incoming edges — import filters
- Outgoing edges — export filters



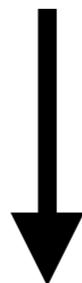
Let BGP find **some allowed** path dynamically

Compilation to BGP:



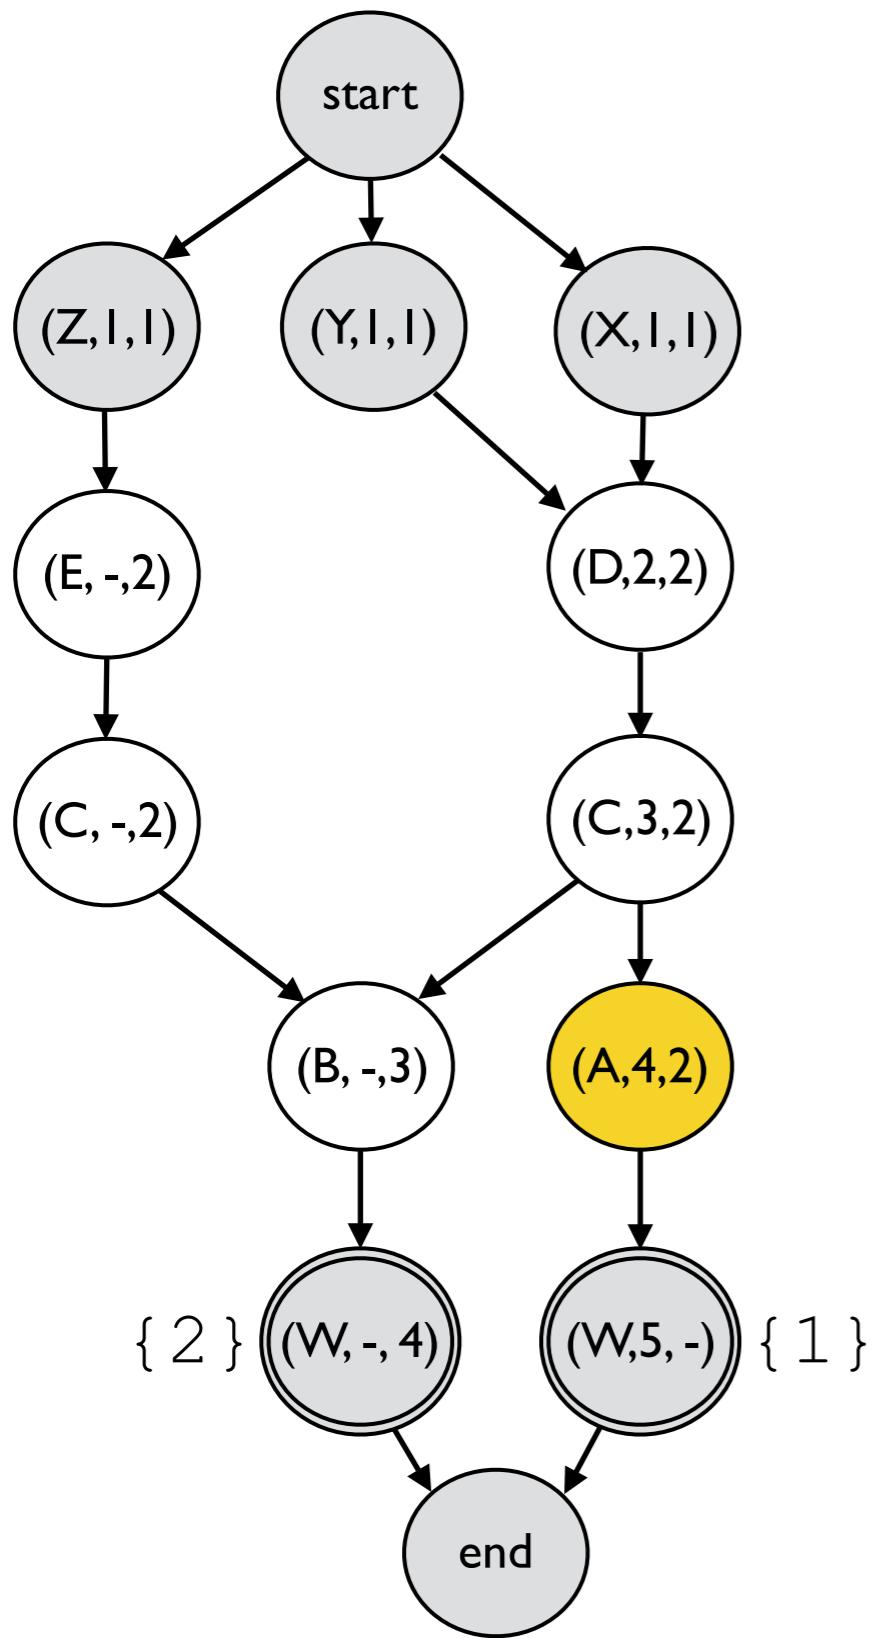
Idea 2: Find local preferences

- Direct BGP towards best path
- Under all combinations of failures
- Frame in terms of graph algorithms



Let BGP find **the best allowed** path dynamically

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

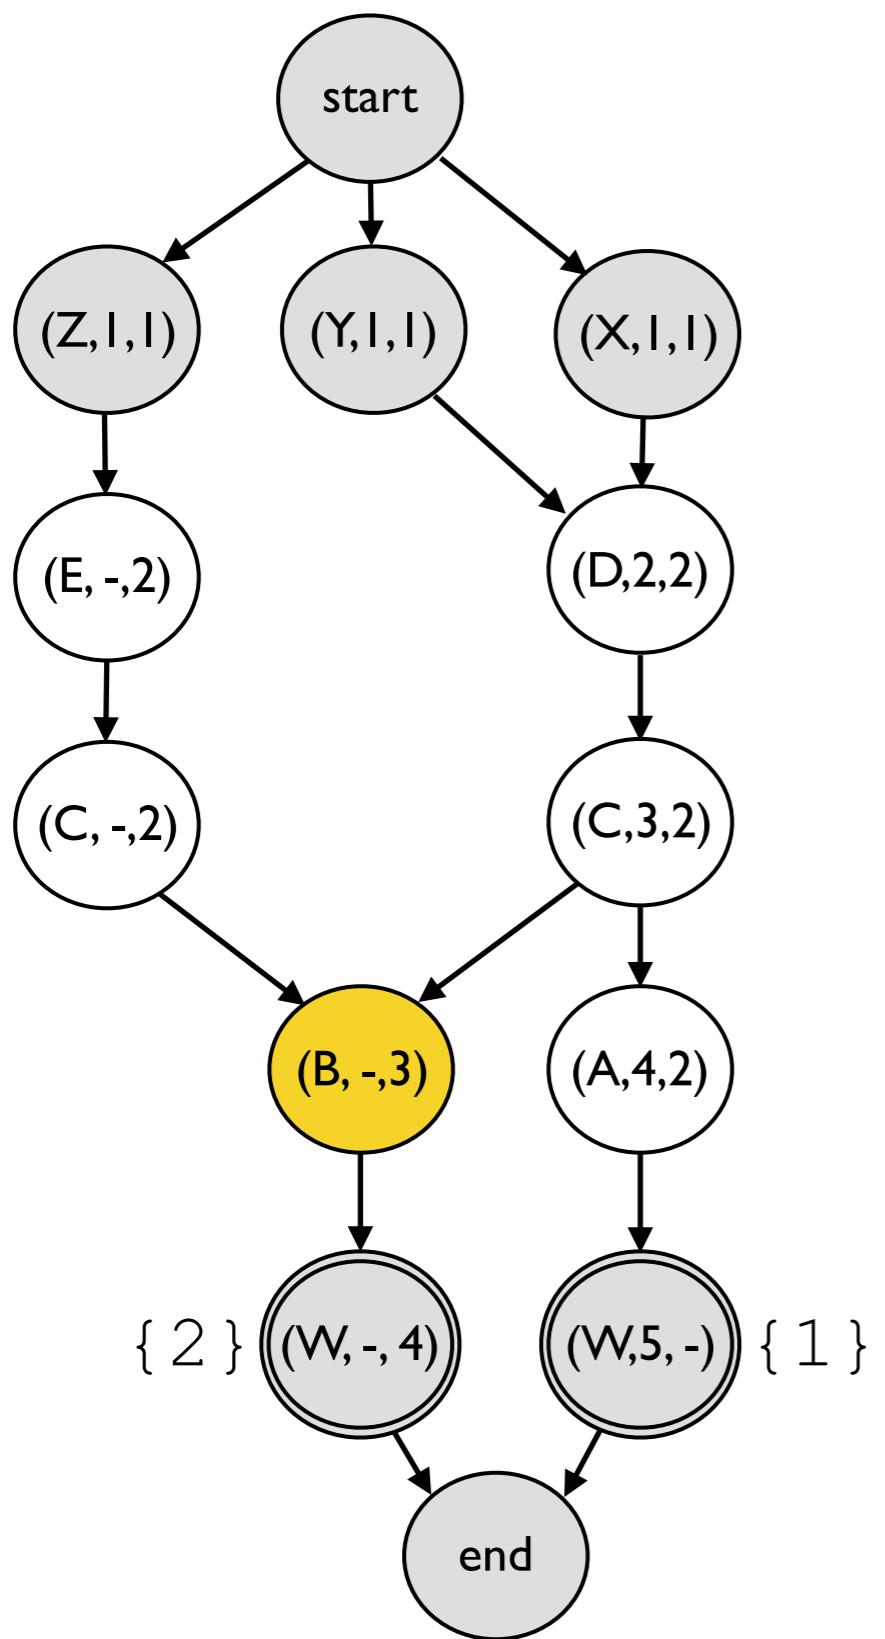
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

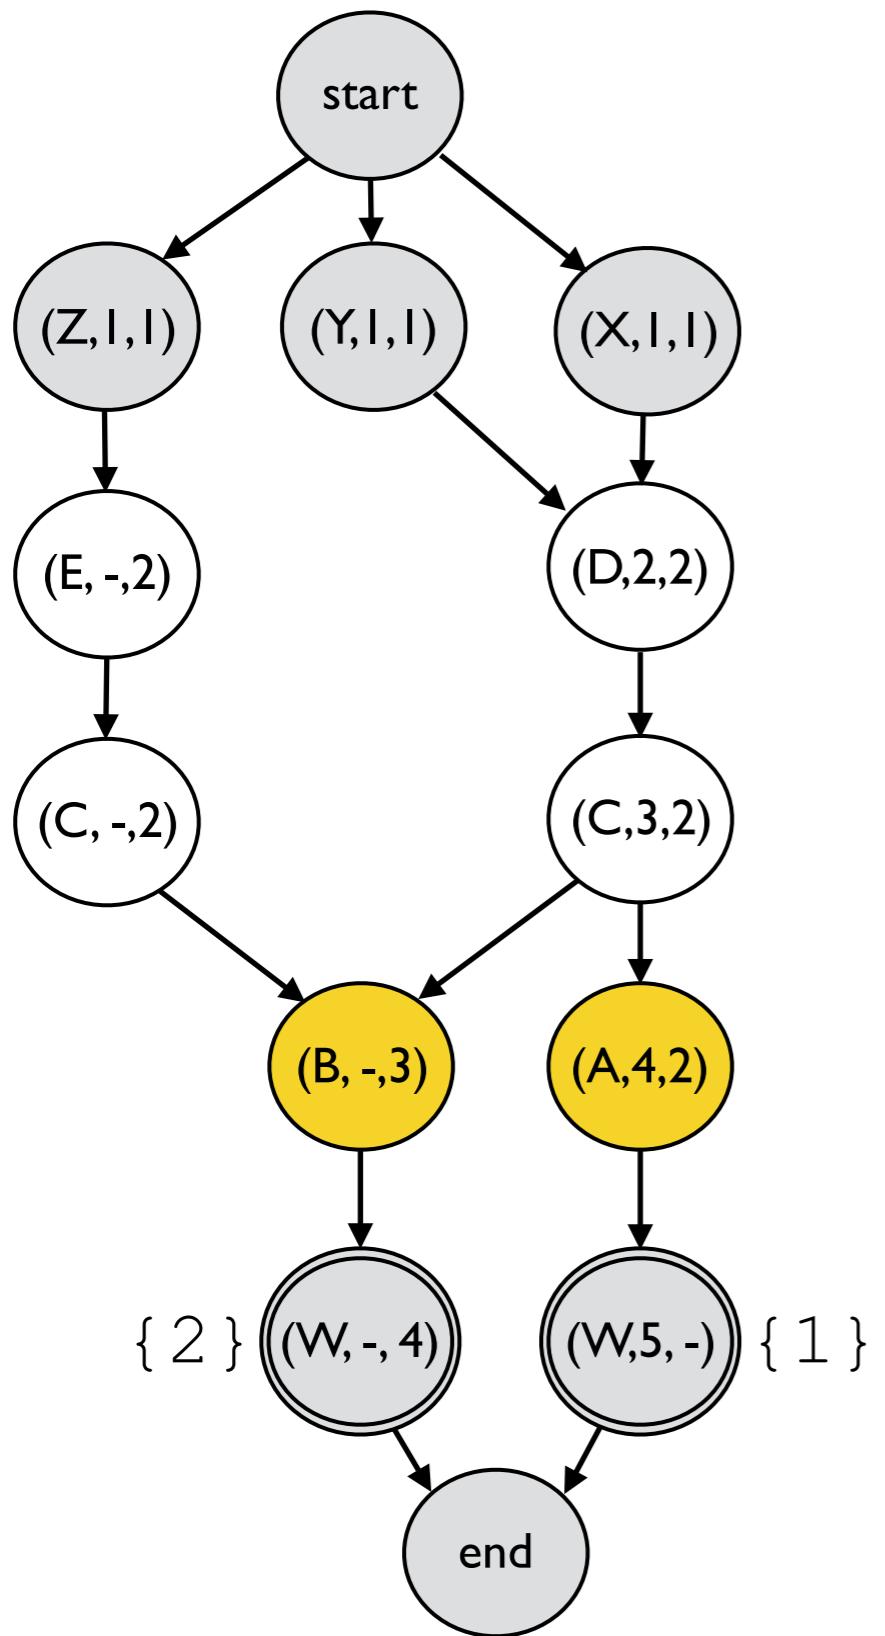
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

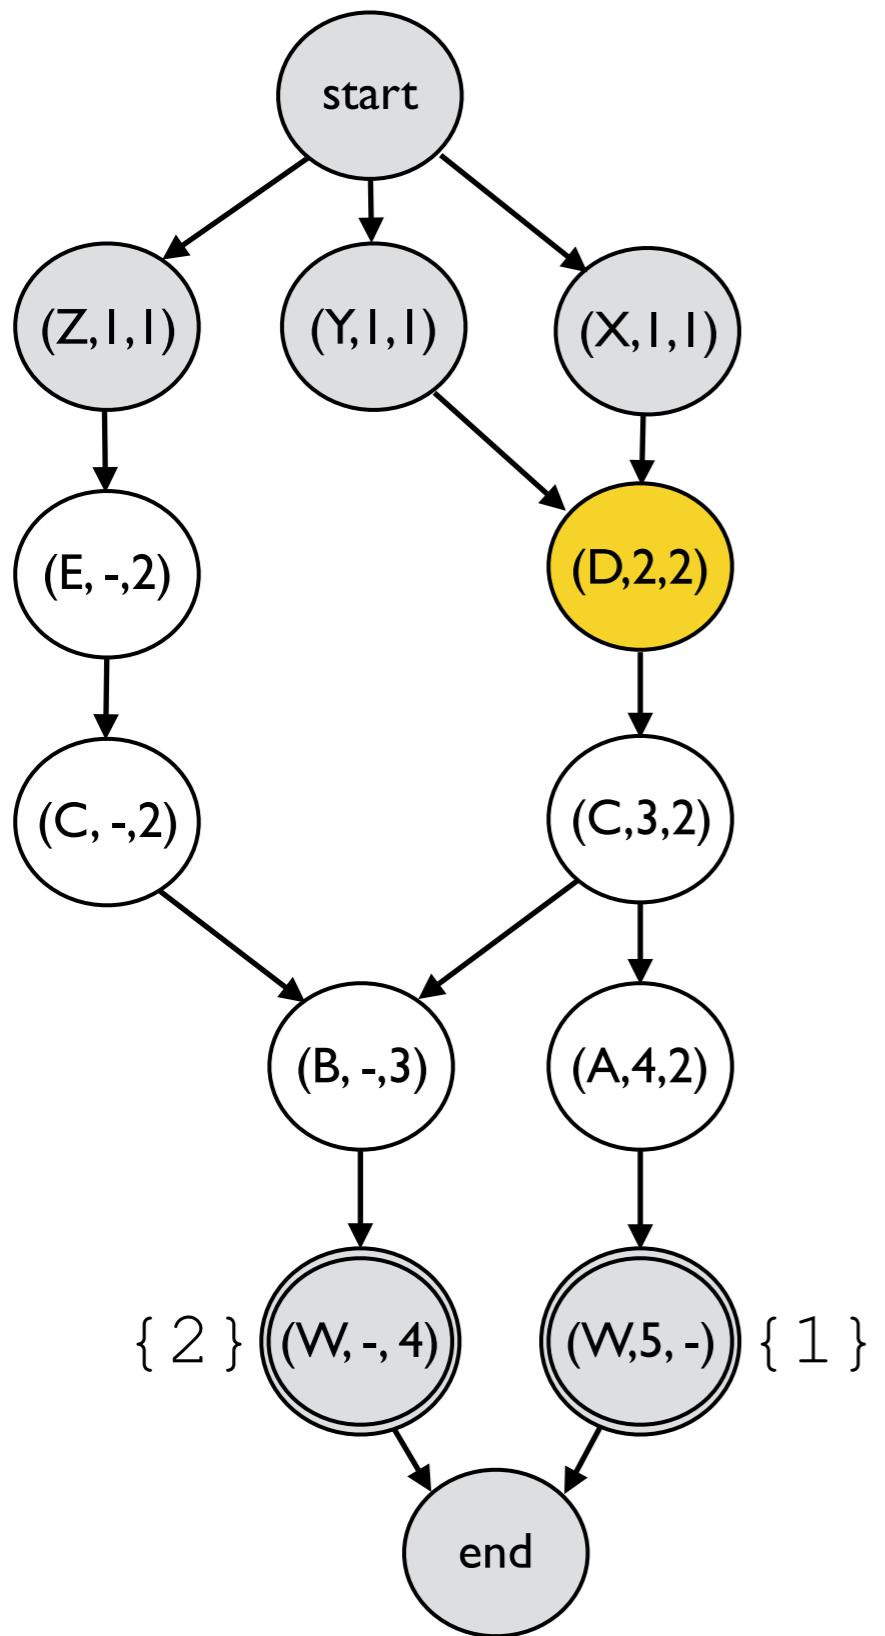
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

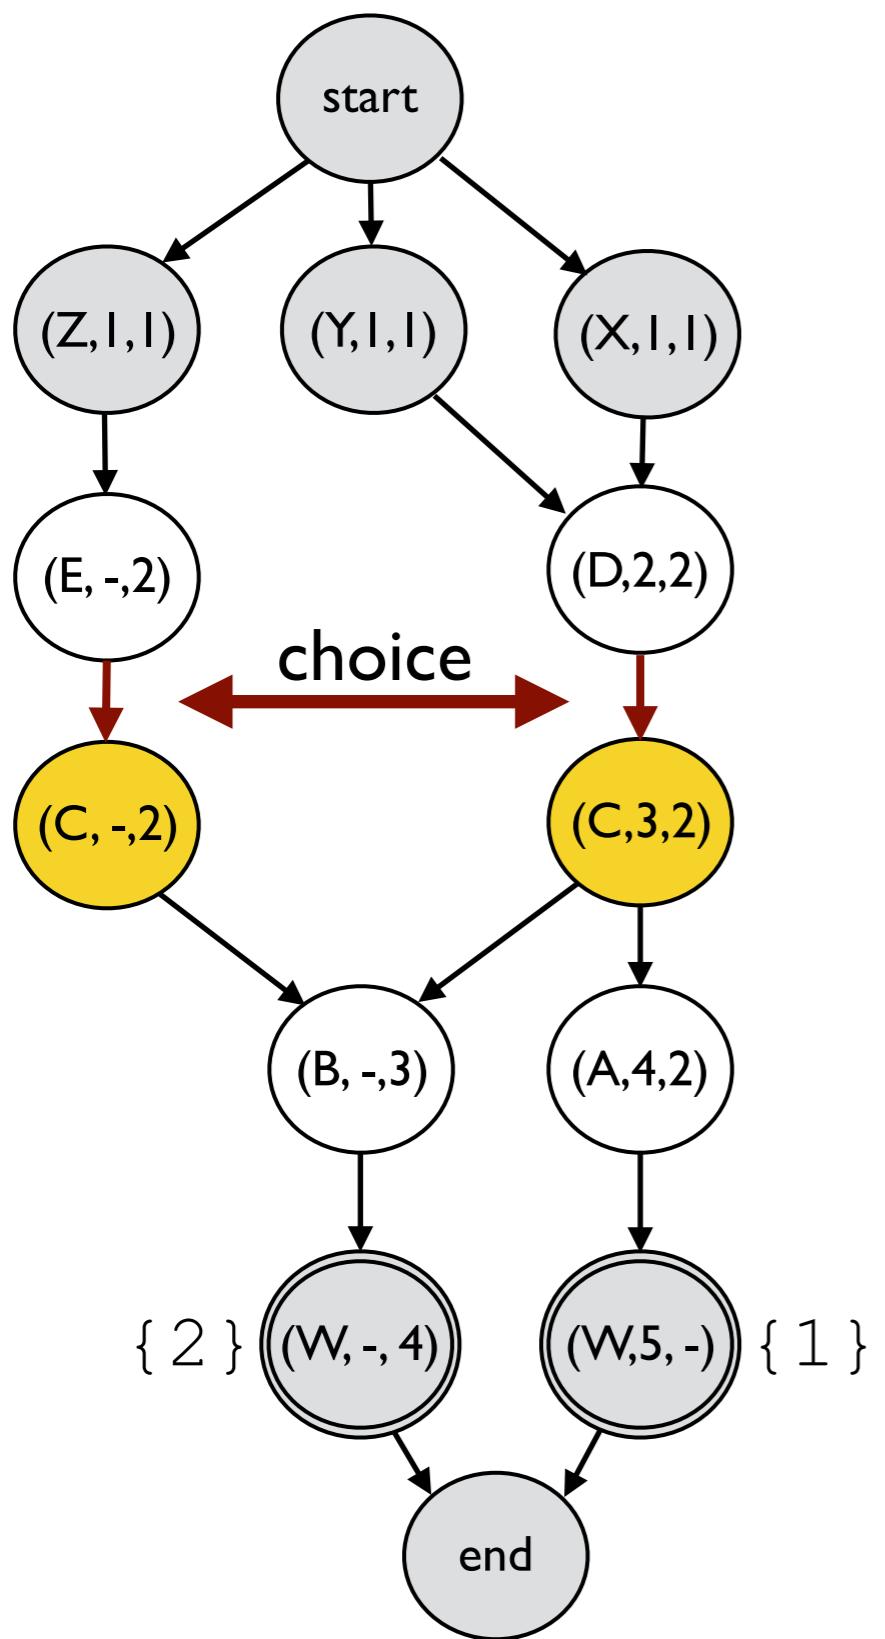
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```

match peer=C comm=(3,2)
export peer←W, comm←(4,2),
          comm←noexport, MED←80
    
```

Router B

```

match peer=C
export peer←W, comm←(-,3),
          comm←noexport, MED←81
    
```

Router C

```

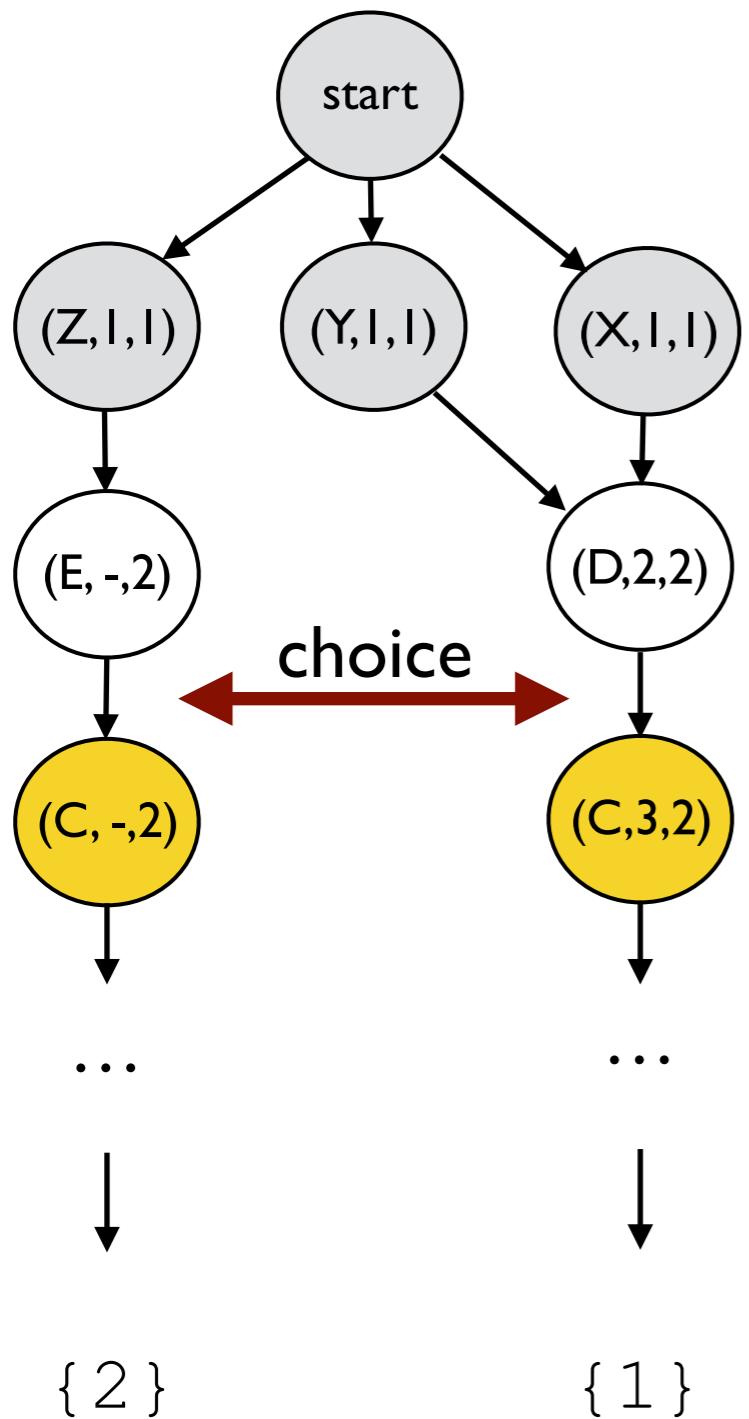
match[lp=99] peer=E, comm=(-,2)
export peer←B, comm←(-,2)
match[lp=100] peer=D, comm=(2,2)
export peer←A,B, comm←(3,2)
    
```

Router D

```

match regex=(X + Y)
export peer←C, comm←(2,2)
...
    
```

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
      comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
      comm←noexport, MED←81
```

Router C

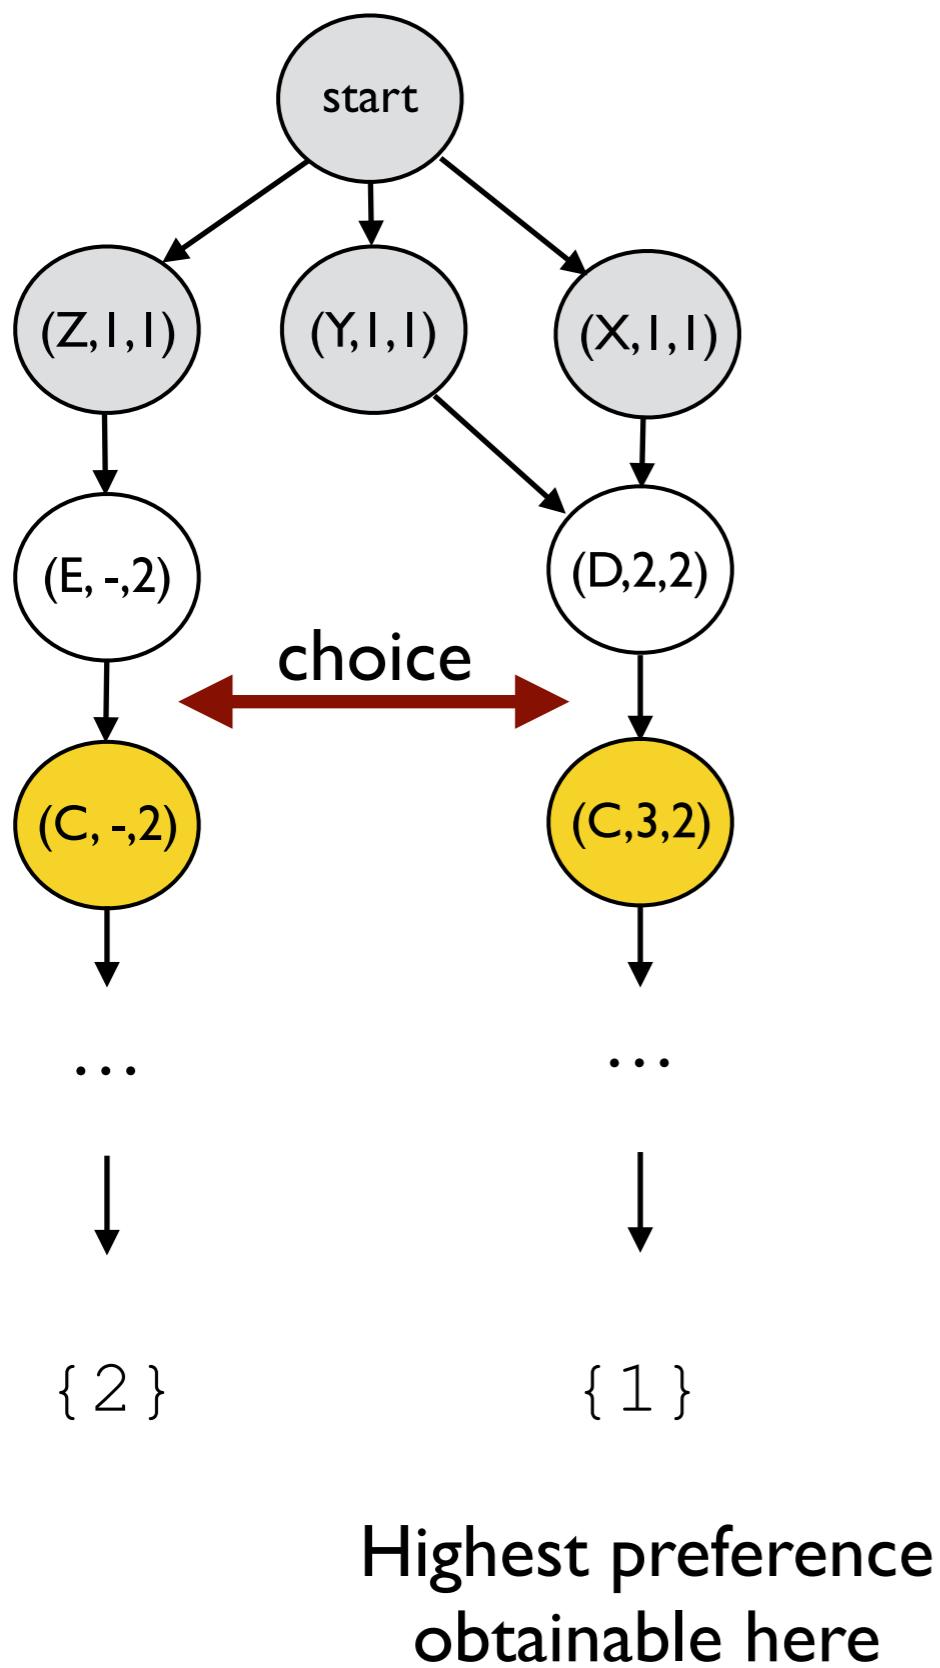
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

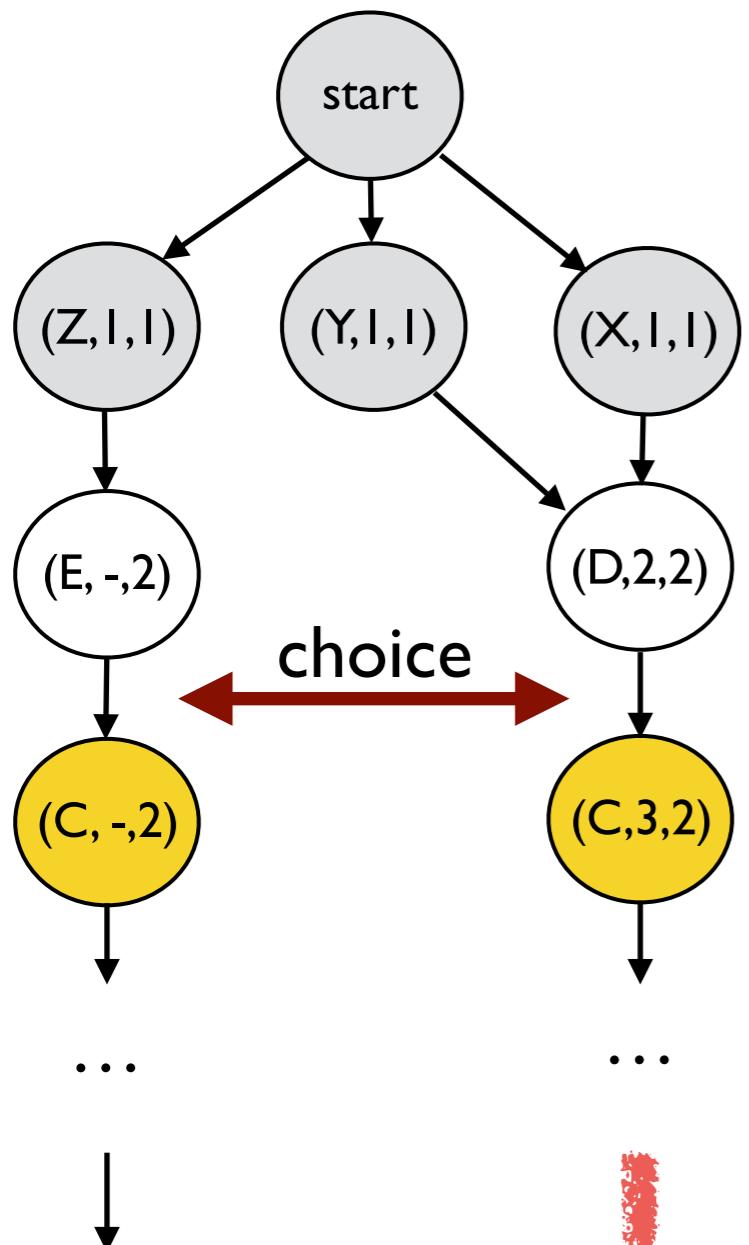
Router C

```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

Compilation to BGP:



{ 2 }

{ 1 }

But there
could be a
failure!

Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

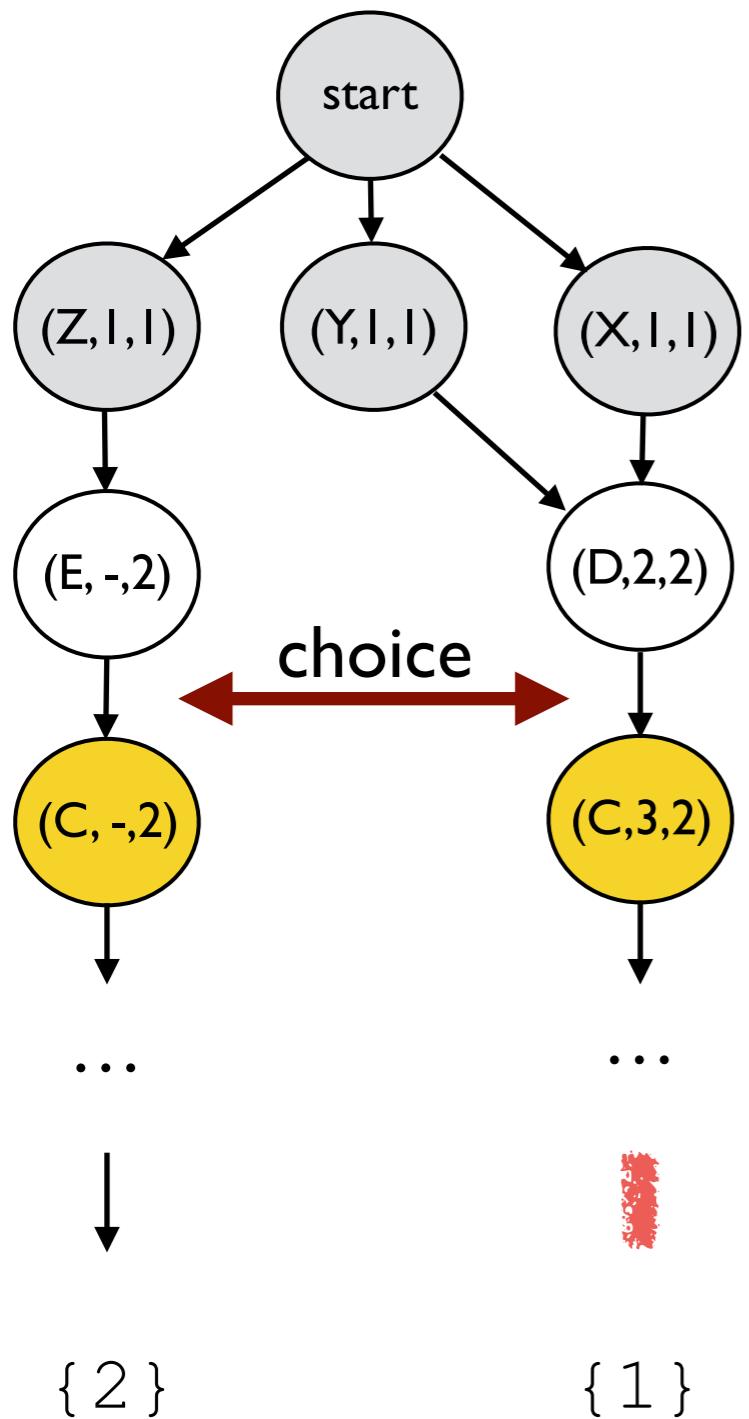
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



*Sometimes there is
no best choice*

Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm← noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

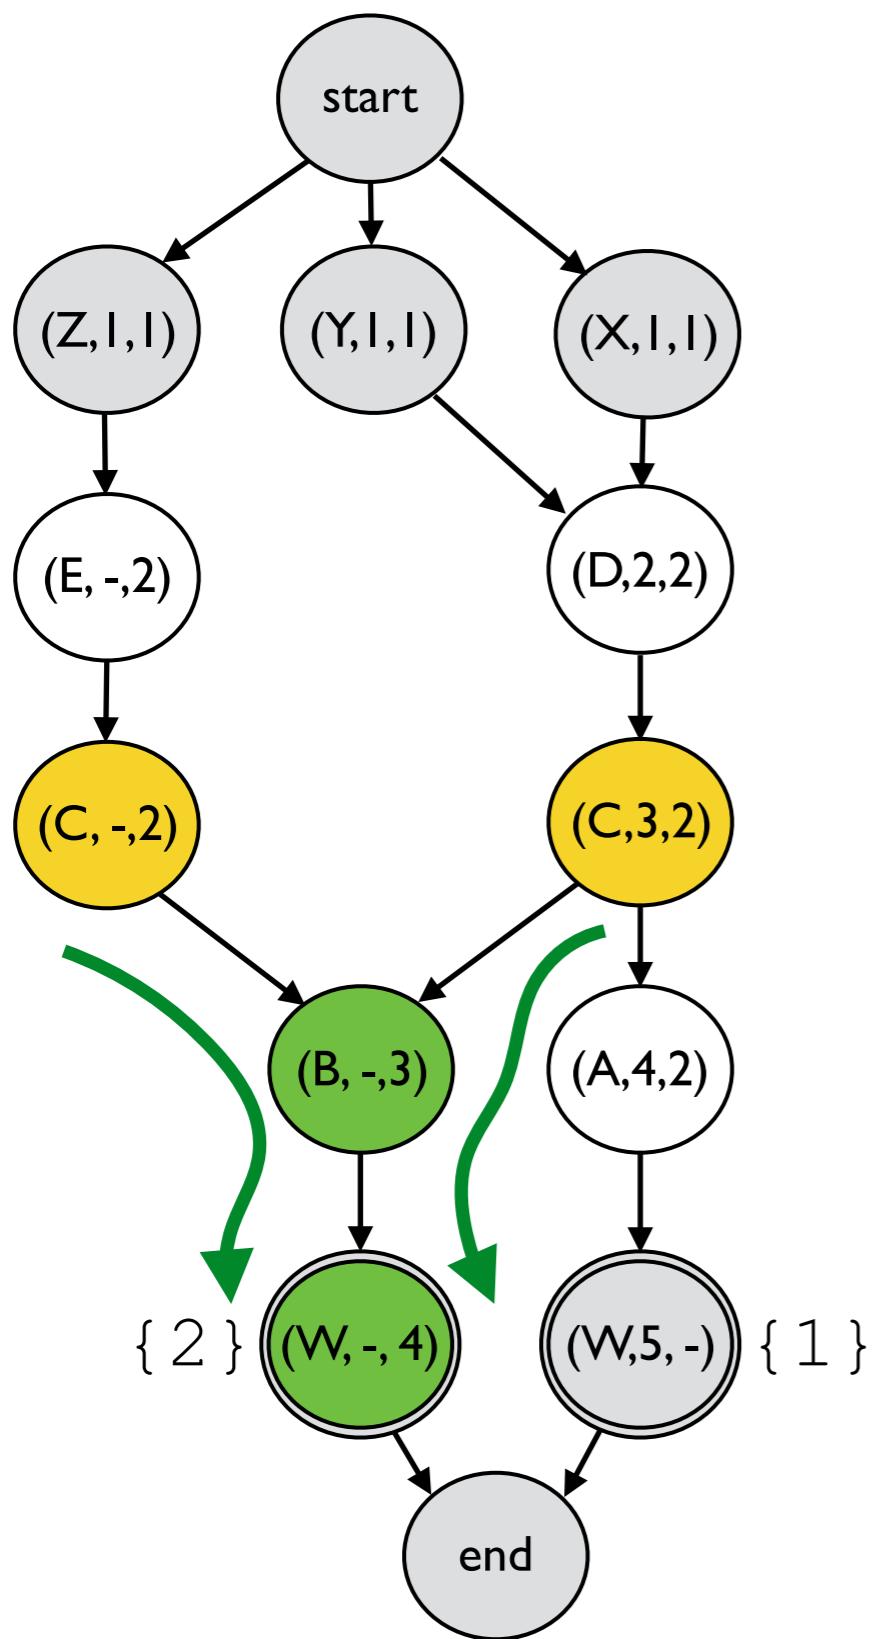
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```
match peer=C comm=(3,2)  
export peer←W, comm←(4,2),  
comm←noexport, MED←80
```

Router B

```
match peer=C  
export peer←W, comm←(-,3),  
comm←noexport, MED←81
```

Router C

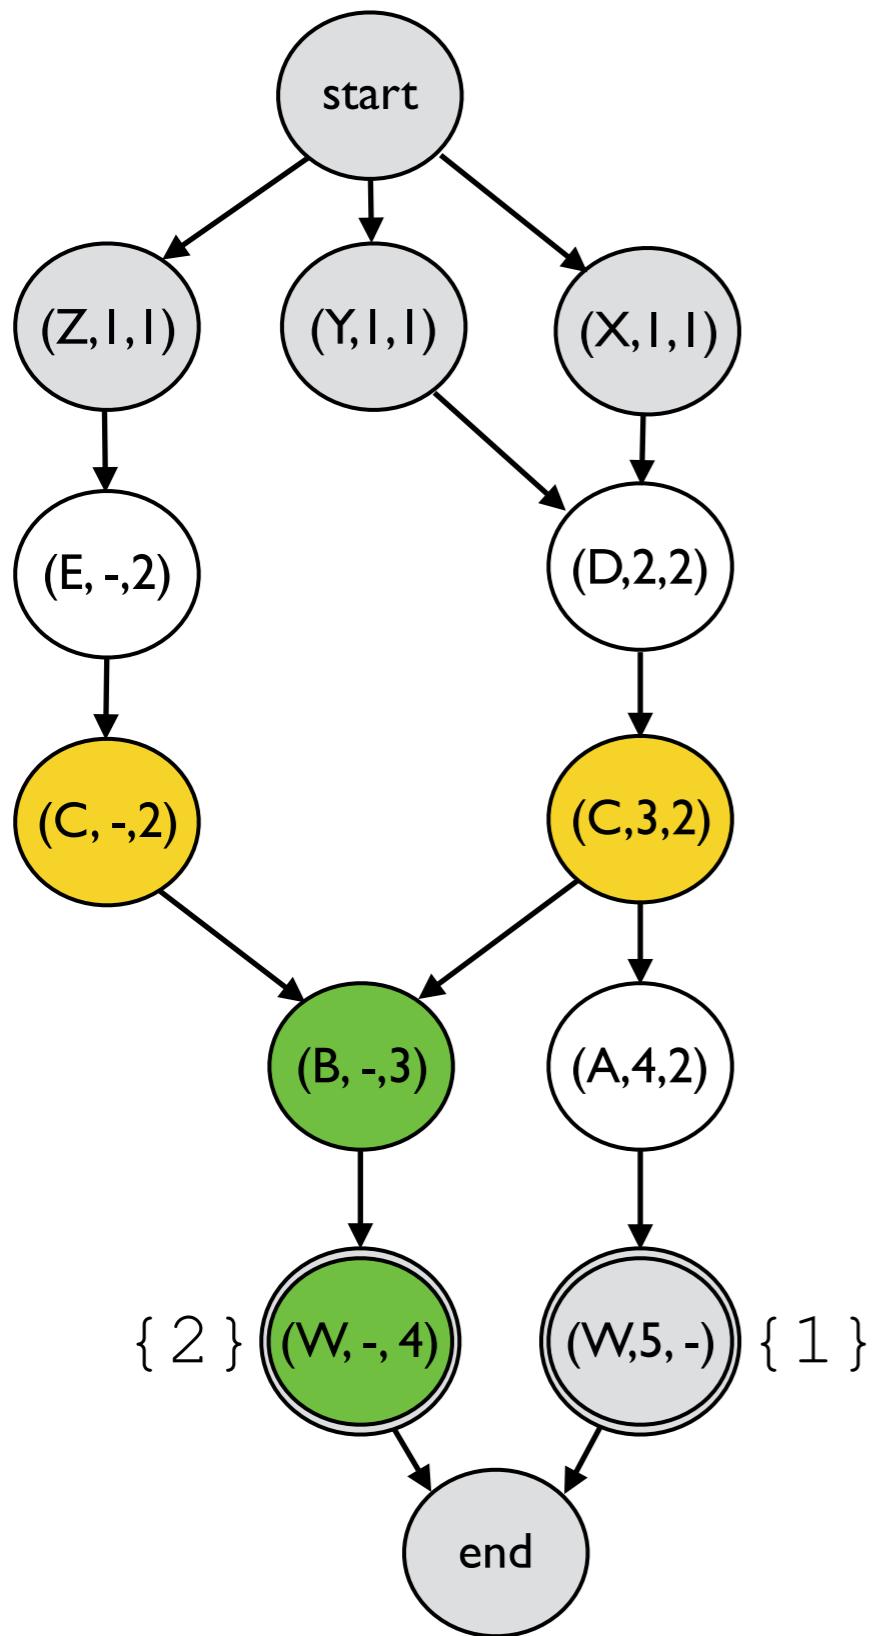
```
match[lp=99] peer=E, comm=(-,2)  
export peer←B, comm←(-,2)  
match[lp=100] peer=D, comm=(2,2)  
export peer←A,B, comm←(3,2)
```

Router D

```
match regex=(X + Y)  
export peer←C, comm←(2,2)
```

...

Compilation to BGP:



Router A

```

match peer=C comm=(3,2)
export peer←W, comm←(4,2),
          comm←noexport, MED←80
  
```

Router B

```

match peer=C
export peer←W, comm←(-,3),
          comm←noexport, MED←81
  
```

Router C

```

match[lp=99] peer=E, comm=(-,2)
export peer←B, comm←(-,2)
match[lp=100] peer=D, comm=(2,2)
export peer←A,B, comm←(3,2)
  
```

Prefer D's
announce

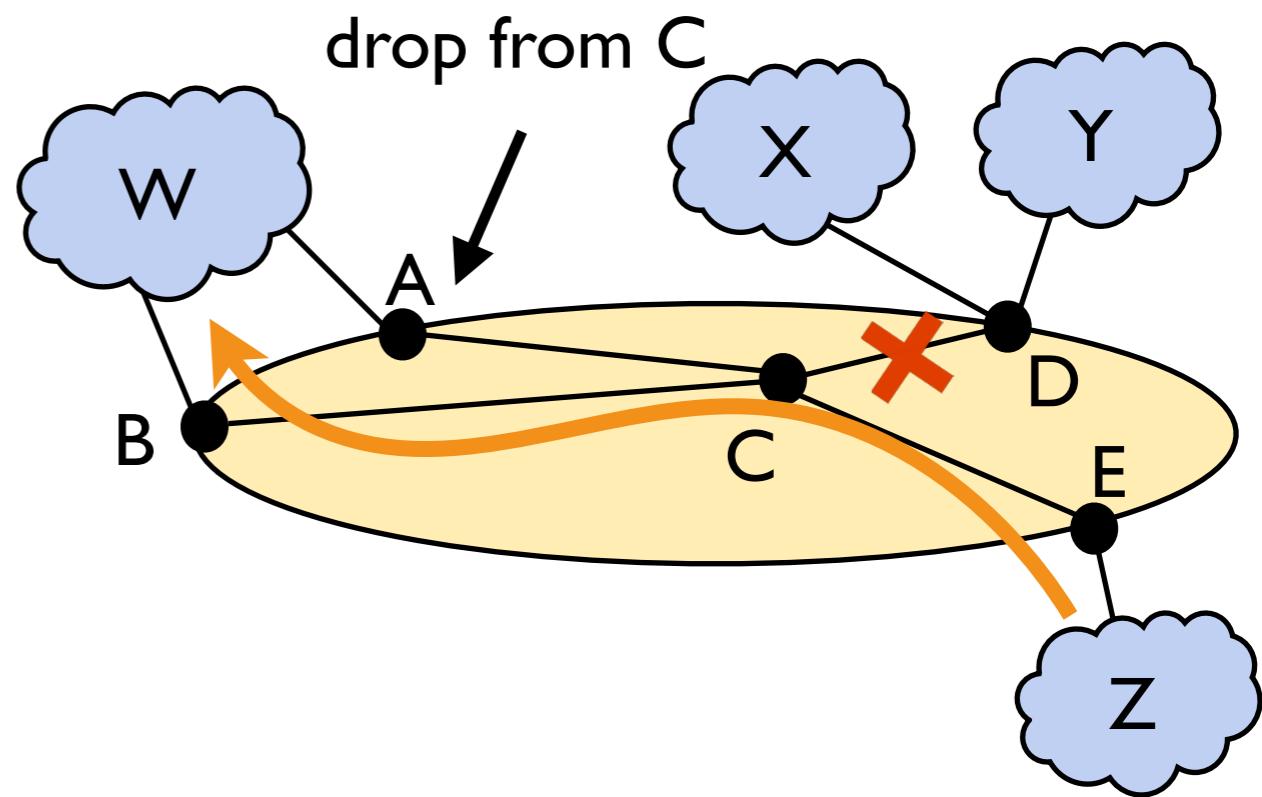
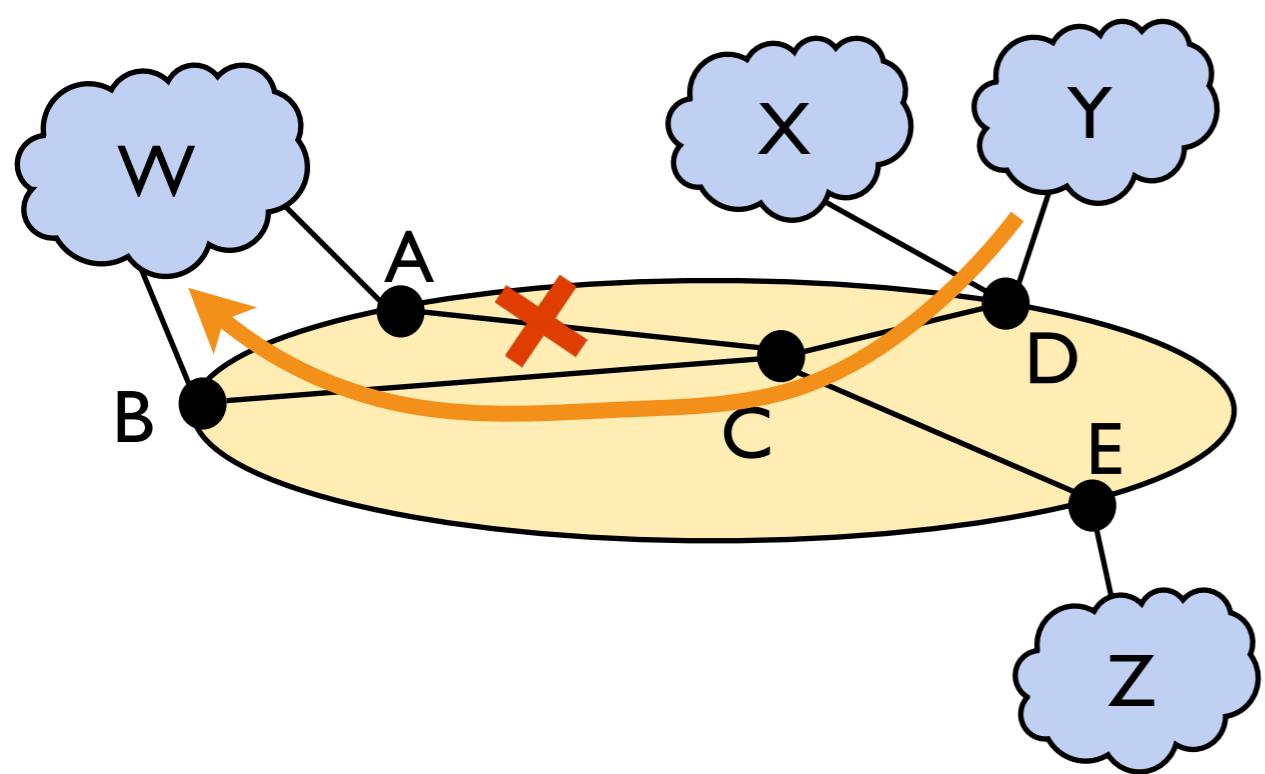
Router D

```

match regex=(X + Y)
export peer←C, comm←(2,2)
...
  
```

Compilation to BGP:

Policy: $(W.A.C.D.\text{.out}) \gg (W.B.\text{.in+}\text{.out})$



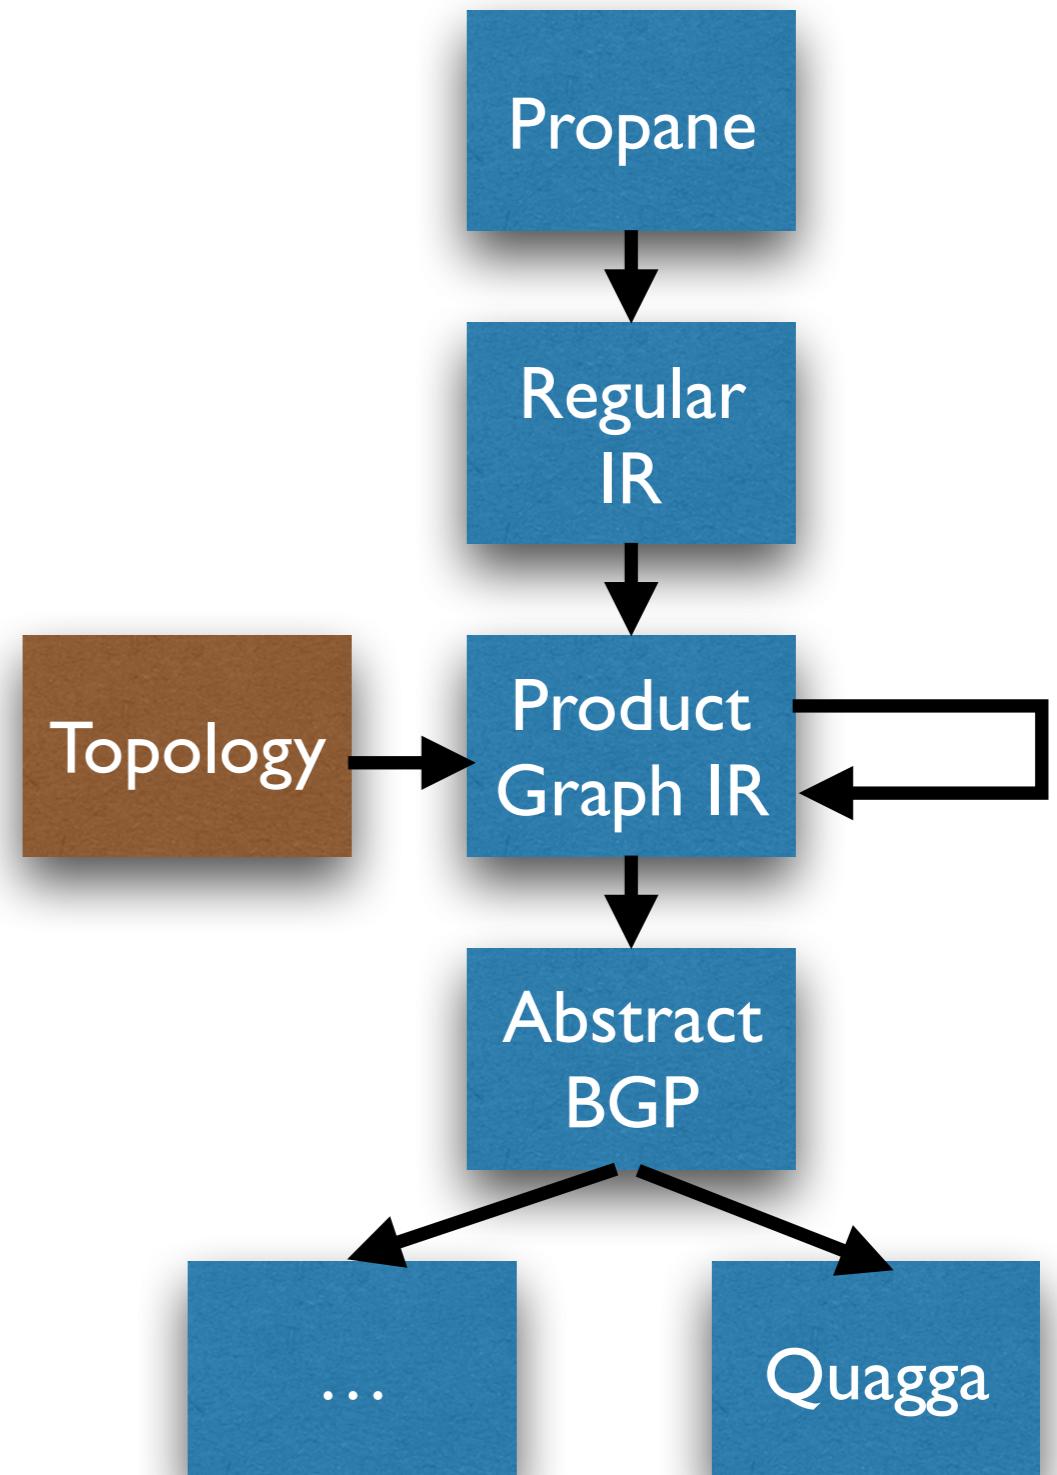
Implementation: C prefers D to E and tags accordingly

The implementation always uses the best available path

Implementation (5,500 lines of F#)

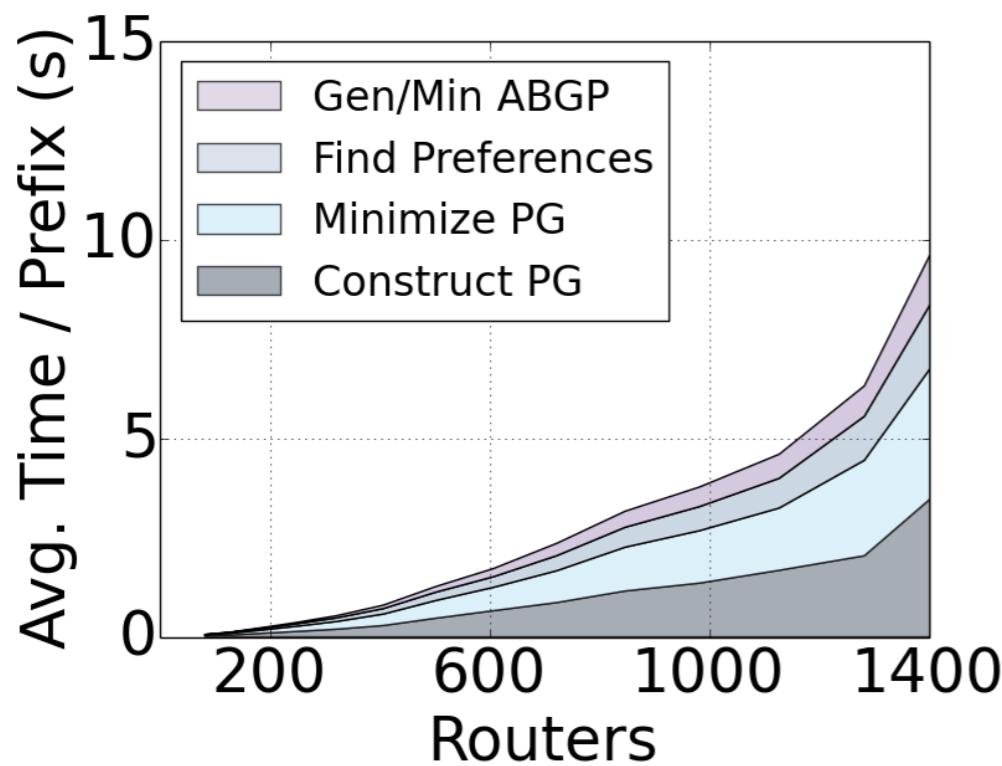
Benchmarks:

- Network configurations from Microsoft's networks (Policy from English docs)
- Data center policies (~1400 routers)
- Backbone policies (~200 routers, many peers/router)
- Ignoring prefix, customer group and ownership definitions:
 - 31 lines for data center
 - 43 lines for backbone

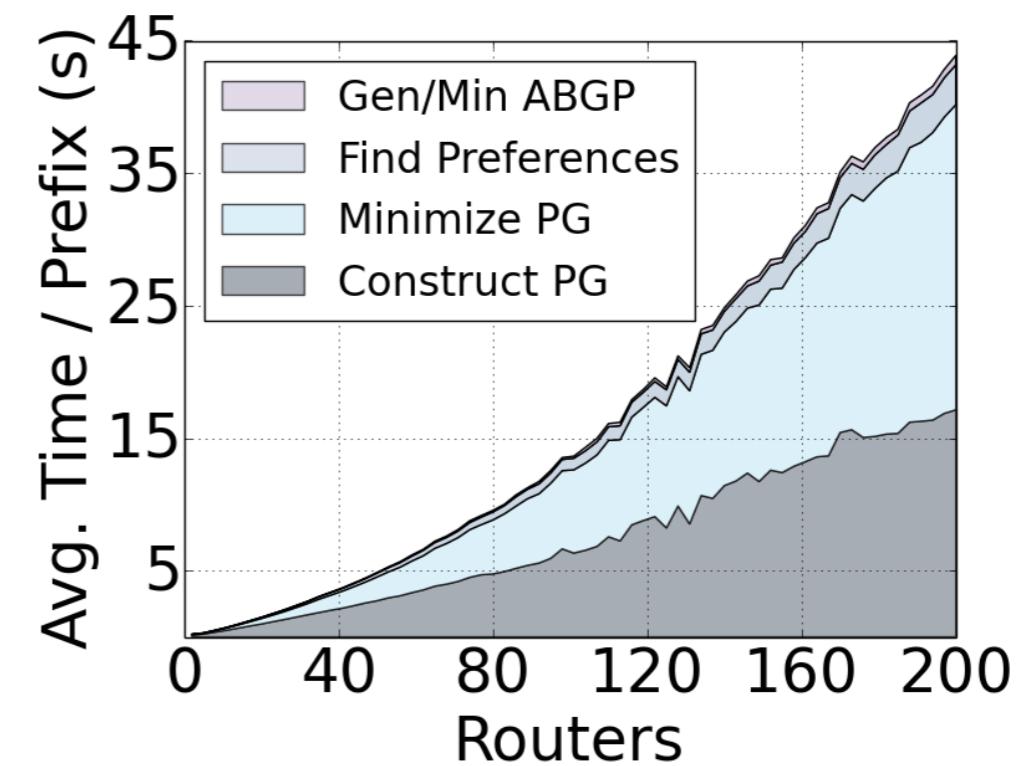


Compilation Time

- Parallelize compilation by prefix
- Use artificial topology to explore scalability



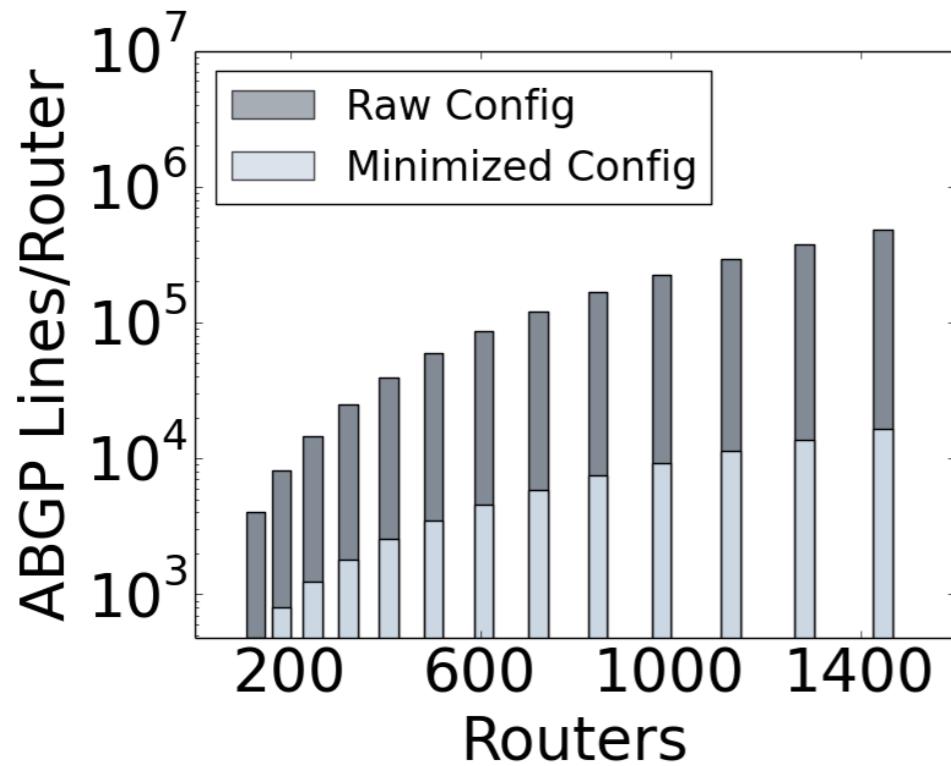
Data center (< 9 min)



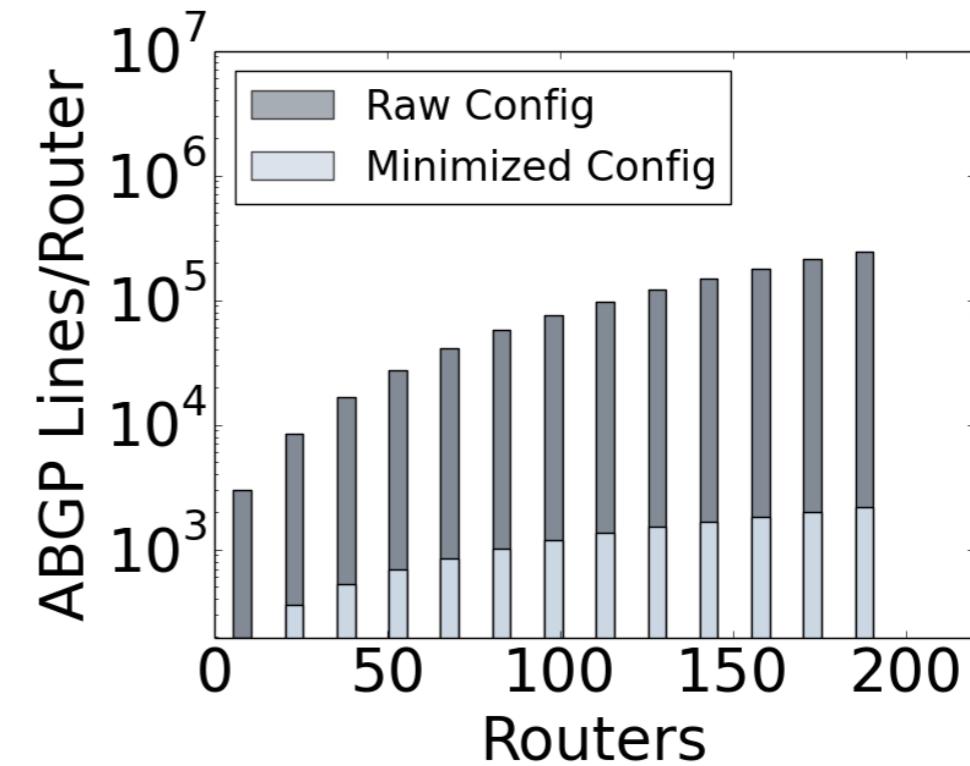
Backbone (< 3 min)

Configuration Size

- Perform configuration minimization during generation
- Avoid using community tags when choices are unambiguous
- Fall-through elimination of route maps



Data center



Backbone